



- Expert Verified, Online, **Free**.

Custom View Settings

Topic 1 - Question Set 1**Topic 1****Question #1****DRAG DROP -**

You have 100 chatbots that each has its own Language Understanding model.

Frequently, you must add the same phrases to each model.

You need to programmatically update the Language Understanding models to include the new phrases.

How should you complete the code? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all.

You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Values
AddPhraseListAsync
Phraselist
PhraselistCreateObject
Phrases
SavePhraselistAsync
UploadPhraseListAsync

Answer Area

```
var phraselistId = await client.Features.AddPhraseListAsync(appId, versionId, new PhraselistCreateObject
{
    EnabledForAllModels = false,
    IsExchangeable = true,
    Name = "PL1",
    Phrases = "item1,item2,item3,item4,item5"
});
```

Correct Answer:

Values
AddPhraseListAsync
Phraselist
PhraselistCreateObject
Phrases
SavePhraselistAsync
UploadPhraseListAsync

Answer Area

```
var phraselistId = await client.Features.AddPhraseListAsync(appId, versionId, new PhraselistCreateObject
{
    EnabledForAllModels = false,
    IsExchangeable = true,
    Name = "PL1",
    Phrases = "item1,item2,item3,item4,item5"
});
```

Box 1: AddPhraseListAsync -

Example: Add phraselist feature -

```
var phraselistId = await client.Features.AddPhraseListAsync(appId, versionId, new PhraselistCreateObject
{
    EnabledForAllModels = false,
    IsExchangeable = true,
    Name = "QuantityPhraselist",
    Phrases = "few,more,extra"
});
```

Box 2: PhraselistCreateObject -

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/client-libraries-rest-api>

 **sovofo**  3 months ago

The practice questions available for the AI-102 exam are 100% real and they allowed me to score 800/1000 on the test. Thank you
msmicrosoft.com
upvoted 34 times

 **cerexid**  1 month, 3 weeks ago

```
// Add phraselist feature
var phraselistId = await client.Features.AddPhraseListAsync(appId, versionId, new PhraselistCreateObject
{ dumps source: https://shorter.me/jqZ36
    EnabledForAllModels = false,
    IsExchangeable = true,
    Name = "QuantityPhraselist",
```

- Phrases = "few,more,extra"
});
upvoted 28 times
- ✉ **hidenori_music** Most Recent 1 day, 10 hours ago
I took this exam yesterday. The same case study questions were asked, although the ET has those questions scattered around instead of grouped together. Please note.
upvoted 1 times
- ✉ **nanaw770** 3 days, 18 hours ago
OMANKO!
upvoted 1 times
- ✉ **funny_penguin** 4 days ago
I've just passed! Around 60-70% of questions were from here. I had 6 case studies and NO simulation questions. Good luck everyone!
upvoted 3 times
- ✉ **michaelmorar** 1 week, 4 days ago
I've just passed! Most of the questions were in this exam topic. NOTE: There were NO simulation questions. Good luck all!
upvoted 6 times
- ✉ **funny_penguin** 6 days, 16 hours ago
thanks for letting us know and congrats.
upvoted 1 times
- ✉ **reiwanotora** 1 week, 1 day ago
I appreciate your reply, it is very helpful.
upvoted 2 times
- ✉ **Joel_** 2 weeks, 1 day ago
Really recommended to use this set of questions to study for the certification exam! Just passed the exam with 940/1000. I got some familiar questions on the exam :). There is no more "Simulation" questions on the exam. Good luck!
upvoted 2 times
- ✉ **TJ001** 2 weeks, 4 days ago
the given answers are correct
upvoted 1 times
- ✉ **joesatriani** 4 weeks, 1 day ago
Is this site dump ready for 29th-Apr English version of the exam?
upvoted 1 times
- ✉ **sivapolam90** 1 month, 2 weeks ago
Box 1: AddPhraseListAsync
Box 2: PhraselistCreateObject
upvoted 1 times
- ✉ **varinder82** 2 months ago
Final Answer:
Box 1: AddPhraseListAsync
Box 2: PhraselistCreateObject
upvoted 1 times
- ✉ **HemaKG** 2 months, 2 weeks ago
LUIS will be retired on October 1st 2025 and starting April 1st 2023 you will not be able to create new LUIS resources. We recommend migrating your LUIS applications to conversational language understanding to benefit from continued product support and multilingual capabilities.
Is the above change going to change this question for the exam point of view?
upvoted 3 times
- ✉ **Earth_Angel** 3 months, 3 weeks ago
What is the retake number for ai 102 in one year? maximum 3 or 5 exams?
upvoted 1 times
- ✉ **Florin83** 3 months, 4 weeks ago
// Add phraselist feature
var phraselistId = await client.Features.AddPhraseListAsync(appId, versionId, new PhraselistCreateObject
{
EnabledForAllModels = false,
IsExchangeable = true,
Name = "QuantityPhraselist",
Phrases = "few,more,extra"
});
upvoted 1 times
- ✉ **[Removed]** 5 months, 3 weeks ago
The practice questions available for the AI-102 exam are 100% real and they allowed me to score 900/1000 on the test. It is a user-friendly service
Thank you Dumps4azure.com

upvoted 15 times

✉️ **NullVoider_0** 4 weeks ago

They really are useful and I scored well, 922/1000 on the test with just only 50% of the questions from here.

upvoted 2 times

✉️ **QuangThang** 2 weeks, 6 days ago

Your exam was also having simulation or not? Thanks

upvoted 2 times

✉️ **dumpsfactory__com** 6 months, 1 week ago

This question was in the exam.

upvoted 5 times

✉️ **reachmymind** 2 years, 2 months ago

Box 1: AddPhraseListAsync

Box 2: PhraselistCreateObject

```
var phraselistId = await client.Features.AddPhraseListAsync(appId, versionId, new PhraselistCreateObject
{
    EnabledForAllModels = false,
    IsExchangeable = true,
    Name = "QuantityPhraselist",
    Phrases = "few,more,extra"
});
```

Mapping :

Model - Entity - Async

Feature - PhraseList - Async

upvoted 9 times

Question #2

DRAG DROP -

You plan to use a Language Understanding application named app1 that is deployed to a container.

App1 was developed by using a Language Understanding authoring resource named lu1.

App1 has the versions shown in the following table.

Version	Trained date	Published date
V1.2	<i>None</i>	<i>None</i>
V1.1	2020-10-01	<i>None</i>
V1.0	2020-09-01	2020-09-15

You need to create a container that uses the latest deployable version of app1.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions	Answer Area
Run a container that has version set as an environment variable.	
Export the model by using the Export as JSON option.	
Select v1.1 of app1.	
Run a container and mount the model file.	
Select v1.0 of app1.	
Export the model by using the Export for containers (GZIP) option.	
Select v1.2 of app1.	

Actions**Answer Area**

Run a container that has version set as an environment variable.

Export the model by using the Export for containers (GZIP) option.

Export the model by using the Export as JSON option.

Select v1.1 of app1.

Select v1.1 of app1.

Run a container and mount the model file.

Correct Answer: Run a container and mount the model file.

Select v1.0 of app1.

Export the model by using the Export for containers (GZIP) option.

Select v1.2 of app1.

Step 1: Export the model using the Export for containers (GZIP) option.

Export versioned app's package from LUIS portal

The versioned app's package is available from the Versions list page.

1. Sign on to the LUIS portal.
2. Select the app in the list.
3. Select Manage in the app's navigation bar.
4. Select Versions in the left navigation bar.
5. Select the checkbox to the left of the version name in the list.
6. Select the Export item from the contextual toolbar above the list.
7. Select Export for container (GZIP).
8. The package is downloaded from the browser.

Versions ?

		Rename	Clone	Export	All	Search for version(s)
		Version name		Created	Last modified	
<input checked="" type="checkbox"/>	0.1 (Active & Production)	Export as JSON		5/3/18	5/3/18	9/6/18
<input checked="" type="checkbox"/>		Export for container (GZIP)				

Step 2: Select v1.1 of app1.

A trained or published app packaged as a mounted input to the container with its associated App ID.

Step 3: Run a contain and mount the model file.

Run the container, with the required input mount and billing settings.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-container-howto>

✉ **kiassi1998** Highly Voted 2 years, 1 month ago

Select the version 1.1 before exporting the model
upvoted 25 times

✉ **Eltooth** Highly Voted 1 year, 10 months ago

Select v1.1
Export
Run and mount
upvoted 22 times

✉ **rdemontis** 6 months, 2 weeks ago

correct

<https://learn.microsoft.com/en-us/azure/ai-services/luis/luis-container-howto?tabs=v3>

upvoted 4 times

✉ **anto69** 4 months, 3 weeks ago

I think so too

upvoted 1 times

✉ **hidenori_music** Most Recent 1 day, 10 hours ago

I took this exam yesterday. The same case study questions were asked, although the ET has those questions scattered around instead of grouped together. Please note.

upvoted 1 times

✉ **demonite** 2 weeks, 2 days ago

Sign on to the LUIS portal.

Select the app in the list.

Select Manage in the app's navigation bar.

Select Versions in the left navigation bar.

Select the checkbox to the left of the version name in the list.

Select the Export item from the contextual toolbar above the list.

Select Export for container (GZIP).

The package is downloaded from the browser.

upvoted 1 times

✉ **CDL_Learner** 1 month, 2 weeks ago

Correct Sequence should be as mentioned below

Step-1 --> Select v1.1 of app1.(We need to use latest version which can be trained or published)

Step 2 --> Export the model by using the Export for container (GZIP) option

Step -3--> Run the Container and mount the model file.

Please refer - <https://learn.microsoft.com/en-us/azure/ai-services/luis/luis-container-howto?tabs=v3>

upvoted 3 times

✉ **sivapolam90** 1 month, 2 weeks ago

1. Export as GZIP

2. Select v1.1

3. Run and mount

upvoted 1 times

✉ **nokono8780** 2 months ago

I just passed the AI-102 exam at the beginning of 2024 (Mar/2024)!

<https://shorturl.at/LU158>

upvoted 1 times

✉ **varinder82** 2 months ago

Final Answer:

1. Select v1.1

2. Export as GZIP

3.. Run and mount

upvoted 2 times

✉ **varinder82** 2 months ago

Select v1.1

Export

Run and mount

upvoted 2 times

✉ **evangelist** 3 months, 2 weeks ago

you have to select first and then export and run:

Select v1.1 of app1: This step is necessary to ensure you are working with the correct version of the application that you intend to deploy.

Export the model by using the Export for containers (GZIP) option: This action is crucial because deploying an app to a container requires the model to be in a format that is optimized for container deployment, which is typically a GZIP file that includes all necessary components of the Language Understanding model.

Run a container and mount the model file: Once you have the model exported in the correct format, you run a container instance and mount the exported model file to it. This allows the container to use the Language Understanding model version 1.1 for processing natural language inputs.

upvoted 4 times

✉ **Nihilist11** 3 months, 3 weeks ago

It was in today's exam -

upvoted 2 times

✉ **costas99** 5 months ago

hi, hope it helps someone looking for the comments section. From the documentation, using a LUIS container it's possible to mount a gzip VERSIONED (just trained model) or PUBLISHED (trained and published) model. In the question it says "latest deployable version", so the app version to export as GZIP is v1.1. So just trained model is deployable, but also deployed model is deployable in a LUIS container

upvoted 3 times

✉ **Prodyna** 6 months, 2 weeks ago

appeared in november exam

upvoted 3 times

✉ **jameswatson211** 6 months, 4 weeks ago

I am feeling Extremely happy, today I just passed my AI-102 exam by scoring 917/1000. Found only 5 new questions. One of them was the question about case study. Thanks to ET and pass4surehub for helping me out. <https://shorturl.at/dhnyO>

upvoted 8 times

✉ **sepenth1** 9 months, 1 week ago

The Language Understanding (LUIS) container loads your trained or published

upvoted 1 times

✉ **sl_mslconsulting** 9 months, 3 weeks ago

I checked the link <https://learn.microsoft.com/en-us/azure/ai-services/luis/luis-container-howto?tabs=v3> and it says you can no longer create new LUIS resources. Wonder why it still appeared in the exam per acsoma's comment.

upvoted 1 times

✉ **rdemontis** 6 months, 2 weeks ago

thanks for the reference.

upvoted 1 times

✉ **Juvilnoz** 9 months, 3 weeks ago

<https://learn.microsoft.com/en-us/azure/ai-services/luis/luis-how-to-manage-versions#other-actions>

upvoted 3 times

✉ **rdemontis** 6 months, 4 weeks ago

thanks for the documentation

upvoted 1 times

Question #3

You need to build a chatbot that meets the following requirements:

- ⇒ Supports chit-chat, knowledge base, and multilingual models
- ⇒ Performs sentiment analysis on user messages
- ⇒ Selects the best language model automatically

What should you integrate into the chatbot?

- A. QnA Maker, Language Understanding, and Dispatch
- B. Translator, Speech, and Dispatch
- C. Language Understanding, Text Analytics, and QnA Maker
- D. Text Analytics, Translator, and Dispatch

Correct Answer: C

Language Understanding: An AI service that allows users to interact with your applications, bots, and IoT devices by using natural language.

QnA Maker is a cloud-based Natural Language Processing (NLP) service that allows you to create a natural conversational layer over your data.

It is used to find the most appropriate answer for any input from your custom knowledge base (KB) of information.

Text Analytics: Mine insights in unstructured text using natural language processing (NLP) no machine learning expertise required. Gain a deeper understanding of customer opinions with sentiment analysis. The Language Detection feature of the Azure Text Analytics REST API evaluates text input

Incorrect Answers:

A, B, D: Dispatch uses sample utterances for each of your bot's different tasks (LUIS, QnA Maker, or custom), and builds a model that can be used to properly route your user's request to the right task, even across multiple bots.

Reference:

<https://azure.microsoft.com/en-us/services/cognitive-services/text-analytics/> <https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/overview/overview>

Community vote distribution

C (83%)

A (17%)

✉ **hidenori_music** 1 day, 10 hours ago

I took this exam yesterday. The same case study questions were asked, although the ET has those questions scattered around instead of grouped together. Please note.

upvoted 1 times

✉ **reiwanotora** 1 week, 1 day ago

Selected Answer: C

C is right.

upvoted 1 times

✉ **JamesKJoker** 2 weeks ago

Selected Answer: C

As Luis is outdated (with its Dispatch option), I must be C

upvoted 1 times

✉ **fatso_567** 2 weeks, 6 days ago

The correct answer is A. QnA Maker, Language Understanding, Dispatch.

The reason is:

QnA Maker: Knowledge Base Support. Answers frequently asked questions and allows interaction.

Language understanding (LUIS): Understanding user intents and entities is essential for meaningful answers and message routing.

Dispatch: To automatically choose the best language model from user input. It manages numerous language models and selects one for a query.

Option B must be revised since translator and speech are unrelated to providing chit-chat, knowledge bases, multilingual models, sentiment analysis, and automatic language model selection.

The criteria already cover sentiment analysis; hence, Option C is inappropriate. Text analytics is mainly used to derive insights from text data. QnA Maker supports knowledge bases better than text analytics.

Option D must be corrected because text analytics is unrelated to the criteria, and the translator prioritises translation over language model selection.

upvoted 1 times

✉ **CDL_Learner** 1 month, 2 weeks ago

Selected Answer: C

Language Understanding: An AI service that allows users to interact with your applications, bots, and IoT devices by using natural language.

QnA Maker is a cloud-based Natural Language Processing (NLP) service that allows you to create a natural conversational layer over your data. It i

used to find the most appropriate answer for any input from your custom knowledge base (KB) of information.

Text Analytics: Mine insights in unstructured text using natural language processing (NLP)"no machine learning expertise required. Gain a deeper understanding of customer opinions with sentiment analysis. The Language Detection feature of the Azure Text Analytics REST API evaluates text input.

Dispatch uses sample utterances for each of your bota - so this is not required

upvoted 3 times

 **afriquiamarocissylesmoulineaux** 2 months, 1 week ago

You should integrate Language Understanding, Text Analytics, and QnA Maker into the chatbot.

upvoted 1 times

Question #4

Your company wants to reduce how long it takes for employees to log receipts in expense reports. All the receipts are in English. You need to extract top-level information from the receipts, such as the vendor and the transaction total. The solution must minimize development effort.

Which Azure service should you use?

- A. Custom Vision
- B. Personalizer
- C. Form Recognizer
- D. Computer Vision

Correct Answer: C

Azure Form Recognizer is a cognitive service that lets you build automated data processing software using machine learning technology.

Identify and extract text, key/value pairs, selection marks, tables, and structure from your documents—the service outputs structured data that includes the relationships in the original file, bounding boxes, confidence and more.

Form Recognizer is composed of custom document processing models, prebuilt models for invoices, receipts, IDs and business cards, and the layout model.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/form-recognizer>

Community vote distribution

C (100%)

✉️  **KingChuang** Highly Voted 1 year, 4 months ago

on my exam (2023-01-16)

My Answer:C

upvoted 9 times

✉️  **wiini** 1 year, 2 months ago

did you pass it? :D, how many questions did you get from here?

upvoted 2 times

✉️  **hidenori_music** Most Recent 1 day, 10 hours ago

I took this exam yesterday. The same case study questions were asked, although the ET has those questions scattered around instead of grouped together. Please note.

upvoted 1 times

✉️  **reiwanotora** 1 week, 1 day ago

Selected Answer: C

C is right.

upvoted 1 times

✉️  **p2006** 2 weeks, 6 days ago

<https://learn.microsoft.com/en-us/azure/ai-services/document-intelligence/overview?view=doc-intel-4.0.0#receipt>

> Form Recognizer is now Azure AI Document Intelligence!

upvoted 1 times

✉️  **CDL_Learner** 1 month, 2 weeks ago

Selected Answer: C

Form Recognizer is an AI-powered document extraction service that understands your forms, enabling you to extract text, tables, and key-value pairs from your documents, whether print or handwritten. It's designed to recognize and extract information from receipts, among other types of documents, which makes it the best choice for this scenario. It can easily extract top-level information such as the vendor and transaction total from receipts, thereby reducing the time employees spend on logging receipts in expense reports. This solution also minimizes development effort as it provides pre-built models for common extraction tasks.

upvoted 2 times

✉️  **CDL_Learner** 1 month, 2 weeks ago

the other options are not suitable:

- A. Custom Vision: This service is used to build custom image classification models. It could be used to recognize receipts, but it wouldn't extract the detailed information from them like Form Recognizer can.
- B. Personalizer: This is an AI service that delivers personalized user experiences. It's not designed for processing receipts or extracting information from documents.
- D. Computer Vision: This service can analyze visual features in images, but it's not specialized for extracting structured data from forms or receipts. It would require significant additional development to extract specific fields from a receipt compared to Form Recognizer.

upvoted 2 times

✉ **michaelmorar** 1 month, 2 weeks ago

Selected Answer: C

Form recogniser supports among other common document types, receipts and invoices.

upvoted 1 times

✉ **evangelist** 3 months, 2 weeks ago

Selected Answer: C

no doubt using Form recognizer to locate the content to be OCR

upvoted 2 times

✉ **ccampagna** 6 months, 2 weeks ago

The correct answer is C "Form Recognizer". This service had changed its name to Document Intelligence a few weeks ago

upvoted 4 times

✉ **jameswatson211** 6 months, 4 weeks ago

I am feeling Extremely happy, today I just passed my AI-102 exam by scoring 917/1000. Found only 5 new questions. One of them was the question about case study. Thanks to ET and pass4surehub for helping me out. <https://shorturl.at/dhnyO>

upvoted 4 times

✉ **rdemontis** 6 months, 4 weeks ago

Selected Answer: C

answer is correct

upvoted 1 times

✉ **dazdzadzadzaazd** 10 months, 1 week ago

Selected Answer: C

Form Recognizer

upvoted 2 times

✉ **james2033** 10 months, 2 weeks ago

Selected Answer: C

See example of Invoice with Azure Form Recognizer at <https://learn.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/overview?view=form-recog-3.0.0#invoice>

upvoted 2 times

✉ **zellick** 11 months ago

Selected Answer: C

C is the answer.

<https://learn.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/overview?view=form-recog-3.0.0>

Azure Form Recognizer is a cloud-based Azure Applied AI Service that enables you to build intelligent document processing solutions. Massive amounts of data, spanning a wide variety of data types, are stored in forms and documents. Form Recognizer enables you to effectively manage the velocity at which data is collected and processed and is key to improved operations, informed data-driven decisions, and enlightened innovation.

upvoted 3 times

✉ **zellick** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 3 times

✉ **ronan500** 10 months ago

Thanks how many questions did u get from ET ?

upvoted 1 times

✉ **kml2003** 1 year, 5 months ago

Selected Answer: C

Form recognizer

upvoted 2 times

✉ **NK0709** 1 year, 8 months ago

C. Form Recognizer

upvoted 1 times

✉ **Eltooth** 1 year, 10 months ago

Selected Answer: C

C is correct answer.

upvoted 2 times

✉ **boofin** 1 year, 11 months ago

Use Form recognizer

upvoted 1 times

Question #5

HOTSPOT -

You need to create a new resource that will be used to perform sentiment analysis and optical character recognition (OCR). The solution must meet the following requirements:

- Use a single key and endpoint to access multiple services.
- Consolidate billing for future services that you might use.
- Support the use of Computer Vision in the future.

How should you complete the HTTP request to create the new resource? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

https://management.azure.com/subscriptions/xxxxxxxxx-xxxx-
 PATCH
 POST
 PUT

xxxx-xxxx-
xxxxxxxxxx/resourceGroups/RG1/providers/Microsoft.CognitiveServices/
accounts/CS1?api-version=2017-04-18
{
 "location": "West US",
 "kind": " CognitiveServices
 ComputerVision
 TextAnalytics",
 "sku": {
 "name": "S0"
 },
 "properties": {},
 "identity": {
 "type": "SystemAssigned"
 }
}

Correct Answer:

Answer Area

```

https://management.azure.com/subscriptions/xxxxxxxxx-xxxx-
xxxx-xxxx-
xxxxxxxxxx/resourceGroups/RG1/providers/Microsoft.CognitiveServices/
accounts/CS1?api-version=2017-04-18
{
    "location": "West US",
    "kind": "CognitiveServices",
    "sku": {
        "name": "S0"
    },
    "properties": {},
    "identity": {
        "type": "SystemAssigned"
    }
}

```

Box 1: PUT -

Sample Request: PUT <https://management.azure.com/subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/test-rg/providers/>

[Microsoft.DeviceUpdate/accounts/contoso?api-version=2020-03-01-preview](#)

Incorrect Answers:

PATCH is for updates.

Box 2: CognitiveServices -

Microsoft Azure Cognitive Services provide us to use its pre-trained models for various Business Problems related to Machine Learning.

List of Different Services are:

- Decision
- Language (includes sentiment analysis)
- Speech
- Vision (includes OCR)
- Web Search

Reference:

<https://docs.microsoft.com/en-us/rest/api/deviceupdate/resourcemanager/accounts/create>

<https://www.analyticsvidhya.com/blog/2020/12/microsoft-azure-cognitive-services-api-for-ai-development/>

✉  **WillyMac**  2 years, 11 months ago

I think answer is correct.

PUT: puts a file or resource at a specific URI, and exactly at that URI.

If there's already a file or resource at that URI, PUT replaces that file or resource.

If there is no file or resource there, PUT creates one.

POST: POST sends data to a specific URI and expects the resource at that URI to handle the request.

upvoted 20 times

✉  **jeffangel28** 2 years, 11 months ago

It's seems correct, the link shows a similar example

<https://docs.microsoft.com/en-us/rest/api/deviceupdate/resourcemanager/accounts/create>

upvoted 1 times

✉  **YipingRuan** 2 years, 10 months ago

Yes, PUT

<https://management.azure.com/subscriptions/{subscriptionId}/resourceGroups/{resourceGroupName}/providers/Microsoft.CognitiveServices/>

/accounts/{accountName}?api-version=2021-04-30
upvoted 3 times

✉️ **AnonymousJhb** 8 months, 2 weeks ago

according to chatGPT, its not PUT.

In the context of HTTP requests and the Azure Resource Manager API, the choice between using a PUT request and a POST request depends on the specific operation you want to perform.

POST Request:

Typically used for resource creation.

When you send a POST request to the specified URL, it's typically used to create a new resource with a unique name.

If the resource specified in the URL already exists, a POST request might result in an error or create a new resource with a different name if Azure allows it.

PUT Request:

Typically used for resource updates or replacements.

When you send a PUT request to the specified URL, it's used to update the existing resource or replace it entirely.

In the context of Azure Resource Manager, if a resource with the specified name already exists at the specified location, a PUT request is used to update the configuration of that existing resource.

upvoted 1 times

✉️ **AnonymousJhb** 8 months, 2 weeks ago

You need to create a new resource that will be used to perform sentiment analysis and optical character recognition (OCR).

So...im presuming its POST...

upvoted 1 times

✉️ **AnonymousJhb** 8 months, 1 week ago

ignore this. there is no delete option!

upvoted 2 times

✉️ **Adedoyin_Simeon** Highly Voted 2 years, 11 months ago

Although in Web Programming, and API dev, PUT is an http(s) request method for an update operation. I can however create a resource when there is no resource to update. I don't know why precisely but the method used by Azure to actually make a REST API request to create a resource is actually "PUT". So, the answers are correct.

See Ref:

<https://docs.microsoft.com/en-us/rest/api/resources/resources/create-or-update>

upvoted 11 times

✉️ **hidenori_music** Most Recent 1 day, 10 hours ago

I took this exam yesterday. The same case study questions were asked, although the ET has those questions scattered around instead of grouped together. Please note.

upvoted 1 times

✉️ **rewanotora** 1 week, 1 day ago

Cognitive Services are "AI parts" that mimic human cognition (Cognitive) and are immediately available as WebAPI. Embedded applications can build cognitive solutions for vision, speech, language, decision making, and search without the need for technical knowledge in AI or data science

upvoted 1 times

✉️ **duyle2906** 1 month ago

When i differed with ChatGPT about PUT being correct, this is what I got

When creating a new resource in Azure, the PUT method is used to update or create the resource with the provided configuration. In the case of Azure Cognitive Services, you typically use the PUT method to provision a new instance of the service with the specified settings.

Therefore, the correct HTTP request to create the new resource should use the PUT method, not POST.

upvoted 1 times

✉️ **pwang009** 2 months, 1 week ago

Answer from ChatGPT, PUT and cognitiveservices

```
curl -X PUT -H "Authorization: Bearer {access_token}" -H "Content-Type: application/json" \
-d '{"kind":"TextAnalytics","sku":{"name":"F0"}}' \
"https://management.azure.com/subscriptions/{subscription_id}/resourceGroups/{resource_group_name}/providers/Microsoft.CognitiveServices/accounts/{api_name}?api-version=2021-04-30"
```

upvoted 1 times

✉️ **Murtuza** 2 months, 2 weeks ago

The given answers are CORRECT

upvoted 1 times

✉️ **Prodyna** 6 months, 2 weeks ago

appeared in november exam

upvoted 2 times

✉️ **sara_aras** 8 months ago

PUT is correct.

<https://learn.microsoft.com/en-us/rest/api/resources/resources/create-or-update>

upvoted 2 times

✉ **f2mrmlwr** 9 months, 1 week ago

Yes, PUT

<https://learn.microsoft.com/en-us/rest/api/cognitiveservices/accountmanagement/accounts/create?tabs=HTTP>

upvoted 1 times

✉ **RegTemp** 10 months, 3 weeks ago

1. PUT

2. CognitiveServices

upvoted 1 times

✉ **zellck** 11 months ago

1. PUT

2. CognitiveServices

<https://learn.microsoft.com/en-us/rest/api/resources/resources/create-or-update>

<https://learn.microsoft.com/en-us/azure/cognitive-services/cognitive-services-apis-create-account?tabs=multiservice%2Canomaly-detector%2Clanguage-service%2Ccomputer-vision%2Cwindows#types-of-cognitive-services-resources>

You can access Azure Cognitive Services through two different resources: A multi-service resource, or a single-service one.

- Multi-service resource:

Access multiple Azure Cognitive Services with a single key and endpoint.

Consolidates billing from the services you use.

upvoted 7 times

✉ **rdemontis** 6 months, 2 weeks ago

Thanks for the references

upvoted 1 times

✉ **Pixelmate** 11 months ago

This came in the exam 28/06/2023

upvoted 3 times

✉ **Eltooth** 1 year, 10 months ago

PUT

CognativeServices

upvoted 3 times

✉ **Contactfornitish** 2 years, 4 months ago

Was on exam 02/01/2022

upvoted 2 times

✉ **sumanshu** 2 years, 4 months ago

PATCH is eliminated (It is only used for update). I think we can use both POST and PUT (to create resources). But good to use PUT (just in case if A has been re-triggered, So it will not fail).

And 2nd answer is "Cognitive services" which provides a lot of models. (So we can use computer vision as well). if we select only computer vision, then we can't use Sentiment analysis and OCR (for which we are trying to create a resource).

upvoted 2 times

✉ **ashu789** 2 years, 5 months ago

was on exam 05/12/2021

upvoted 1 times

Question #6

You are developing a new sales system that will process the video and text from a public-facing website.

You plan to monitor the sales system to ensure that it provides equitable results regardless of the user's location or background.

Which two responsible AI principles provide guidance to meet the monitoring requirements? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. transparency
- B. fairness
- C. inclusiveness
- D. reliability and safety
- E. privacy and security

Correct Answer: BD

AI systems should treat all people fairly.

AI systems should perform reliably and safely.

Reference:

<https://docs.microsoft.com/en-us/azure/cloud-adoption-framework/strategy/responsible-ai>

Community vote distribution

BC (100%)

Eltooth Highly Voted 1 year, 10 months ago

Selected Answer: BC

B and C are correct answers.

upvoted 18 times

natgurulearning Highly Voted 5 months, 4 weeks ago

A and B 100%, key word is to Monitor

Transparency means you should be clear about how the system works and how it makes decisions, so you can monitor the results and ensure that they're not biased against any particular group of users. And fairness means that the system should be designed to treat all users equally, regardless of their location or background. By monitoring the system to ensure that it's transparent and fair, you can help to prevent any unintended biases or discrimination in the results. It's all about being ethical and responsible in the development and use of AI systems.

upvoted 6 times

hidenori_music Most Recent 1 day, 10 hours ago

I took this exam yesterday. The same case study questions were asked, although the ET has those questions scattered around instead of grouped together. Please note.

upvoted 1 times

nanaw770 3 days, 18 hours ago

Selected Answer: BC

It must be B and C.

upvoted 1 times

JamesKJoker 2 weeks ago

Selected Answer: BC

BC is obviously correct

upvoted 1 times

CDL_Learner 1 month, 2 weeks ago

Selected Answer: BC

B. Fairness: This principle is about ensuring that the AI system provides equitable results regardless of a user's location or background. It involves minimizing bias in the AI system's outcomes and providing equal opportunities for all users.

C. Inclusiveness: This principle is about designing AI systems that are accessible and usable by the widest possible range of people, regardless of their location or background. It involves considering the diverse characteristics of potential users during the design and deployment of the AI system.

upvoted 2 times

sivapolam90 1 month, 2 weeks ago

Selected Answer: BC

B and C

upvoted 1 times

✉️ **michaelmorar** 1 month, 2 weeks ago

Selected Answer: BC

The requirements to not mention reliability and safety, only that the users are treated equally regardless of their characteristics or background.
upvoted 1 times

✉️ **f2c587e** 2 months ago

Selected Answer: BC

En la pregunta claramente hablan sobre Equidad y la inclusión al abordar que debe ser independiente de la ubicación
upvoted 1 times

✉️ **evangelist** 3 months, 2 weeks ago

Selected Answer: BC

transparency (A) is important for making the operations of AI systems understandable to users, and reliability and safety (D), along with privacy and security (E), are crucial for ensuring the system's dependable operation, user safety, and data protection, fairness and inclusiveness are specifically targeted towards ensuring equitable results and addressing the monitoring requirements stated in your scenario.
upvoted 1 times

✉️ **rdemontis** 6 months, 4 weeks ago

Selected Answer: BC

Fairness (equity) and Inclusiveness (regardless locations and backgrounds)
upvoted 3 times

✉️ **zellck** 11 months ago

Selected Answer: BC

BC is the answer.

<https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/innovate/best-practices/trusted-ai#fairness>

Fairness is a core ethical principle that all humans aim to understand and apply. This principle is even more important when AI systems are being developed. Key checks and balances need to make sure that the system's decisions don't discriminate or run a gender, race, sexual orientation, or religion bias toward a group or individual.

<https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/innovate/best-practices/trusted-ai#inclusiveness>

Inclusiveness mandates that AI should consider all human races and experiences, and inclusive design practices can help developers to understand and address potential barriers that could unintentionally exclude people. Where possible, speech-to-text, text-to-speech, and visual recognition technology should be used to empower people with hearing, visual, and other impairments.

upvoted 5 times

✉️ **EliteAllen** 11 months, 3 weeks ago

Selected Answer: BC

B, C are correct

upvoted 3 times

✉️ **marti_tremblay000** 1 year, 2 months ago

Selected Answer: BC

ChatGPT answer :

The two responsible AI principles that provide guidance to meet the monitoring requirements are B. fairness and C. inclusiveness.

Fairness is the principle of ensuring that AI systems do not discriminate against any group or individual. In this scenario, monitoring the sales system to ensure equitable results regardless of the user's location or background helps to meet the fairness principle.

Inclusiveness is the principle of ensuring that AI systems are designed to be accessible to all users, regardless of their individual characteristics or abilities. Monitoring the sales system to ensure equitable results helps to meet the inclusiveness principle by ensuring that the system is not excluding or discriminating against any particular group or individual.

upvoted 4 times

✉️ **claps92** 1 year, 2 months ago

According to me, A and B

upvoted 2 times

✉️ **GigaCaster** 1 year, 7 months ago

The answer is B and D as the question stated that this is a public-facing website thus reliability and safety are a big must and the second part which asks for equitable results regardless of background and location gets covered by fairness. I believe inclusiveness isn't correct because it isn't the main focus of this website atm. For instance, if the question was worded a bit differently like the results should be equitable results regardless of background, location, and/or impairment/ disability then I would say yes inclusiveness is correct.

upvoted 4 times

✉️ **GigaCaster** 1 year, 7 months ago

After rereading the question I would like to correct myself in saying that D was incorrect and C was correct as the question doesn't ask for anything related to reliability or safety in the first place because you aren't monitoring the website but rather the sales system itself.

upvoted 2 times

✉️ **Radeke** 1 year, 8 months ago

I don't think C is correct in this context. Inclusiveness refers to features/solutions that consider people with disadvantages or disabilities of some sorts. E.g. adding speech -to-text features for the visually impaired. I think it's B and D.

<https://docs.microsoft.com/en-us/learn/modules/responsible-ai-principles/4-guiding-principles>

upvoted 2 times

 **dev2dev** 1 year, 4 months ago

A geo-location biased AI model will result in biased results. C is correct because without inclusiveness of variety of locations data in training, a trained model will can result in biased results.

upvoted 3 times

Question #7

DRAG DROP -

You plan to use containerized versions of the Anomaly Detector API on local devices for testing and in on-premises datacenters.

You need to ensure that the containerized deployments meet the following requirements:

- Prevent billing and API information from being stored in the command-line histories of the devices that run the container.
- Control access to the container images by using Azure role-based access control (Azure RBAC).

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

NOTE: More than one order of answer choices is correct. You will receive credit for any of the correct orders you select.

Select and Place:

Actions**Answer Area**

- | |
|--|
| Create a custom Dockerfile. |
| Pull the Anomaly Detector container image. |
| Distribute a docker run script. |
| Push the image to an Azure container registry. |
| Build the image. |
| Push the image to Docker Hub. |

Correct Answer:

Actions**Answer Area**

- | |
|--|
| Create a custom Dockerfile. |
| Pull the Anomaly Detector container image. |
| Distribute a docker run script. |
| Push the image to an Azure container registry. |
| Build the image. |
| Push the image to Docker Hub. |

- | |
|--|
| Pull the Anomaly Detector container image. |
| Create a custom Dockerfile. |
| Push the image to an Azure container registry. |
| Distribute a docker run script. |

Step 1: Pull the Anomaly Detector container image.

Step 2: Create a custom Dockerfile

Step 3: Push the image to an Azure container registry.

To push an image to an Azure Container registry, you must first have an image.

Step 4: Distribute the docker run script

Use the docker run command to run the containers.

Reference:

<https://docs.microsoft.com/en-us/azure/container-registry/container-registry-intro>

  **dinhhungitsoft**  2 years ago

I think the it should be:

1. Create a custom docker file
2. Pull the container image (in the dockerfile)
3. Build the image
- 4 Push to ACR

upvoted 47 times

  **sheldon73** 11 months, 3 weeks ago

This makes more sense, In the custom docker file you can define : FROM azure_image...

upvoted 3 times

  **mmaguero**  1 year ago

ChatGPT: To meet the requirements, perform the following actions in sequence:

1. Create a custom Dockerfile: This allows you to define the specific configuration and dependencies for your containerized deployment.
2. Build the image: Use the Dockerfile to build the container image with the necessary components and configurations.

3. Push the image to an Azure container registry: Store the container image in an Azure container registry to control access and ensure secure storage.

4. Distribute a docker run script: Provide a docker run script that can be used to run the containerized deployment on local devices or in on-premises datacenters. This script should include any necessary environment variables or configuration settings.

Note: The actions of pulling the Anomaly Detector container image and pushing the image to Docker Hub are not mentioned in the requirements. Therefore, they are not necessary for meeting the given requirements.

upvoted 18 times

✉️ **rdemontis** 6 months, 4 weeks ago

I agree with you except for the explanation in the endnote: for me your answer is correct, but the reason why the pull of the anomaly detector image is irrelevant is because you put it in the dockerfile. So it is not necessary to specify this instruction separately otherwise we would have to do the same for all commands in the dockerfile.

upvoted 5 times

✉️ **anto69** 4 months, 3 weeks ago

agree with u

upvoted 1 times

✉️ **hidenori_music** **Most Recent** 1 day, 10 hours ago

I took this exam yesterday. The same case study questions were asked, although the ET has those questions scattered around instead of grouped together. Please note.

upvoted 1 times

✉️ **nanaw770** 3 days, 18 hours ago

- 1.Create a custom Dockerfile
- 2.Build the image
- 3.Push the image to an Azure container registry
- 4.Distribute a docker run script

From Takedajuku perspective, if you study for 4 days and spend 2 days reviewing, you will have a better chance of passing the exam.

upvoted 1 times

✉️ **TJ001** 2 weeks, 4 days ago

it is important to note NOTE: More than one order of answer choices is correct.

upvoted 1 times

✉️ **TJ001** 2 weeks, 4 days ago

1. Pull the image
2. Create dockerfile
3. Build
4. Run
5. Push to ACR if required

upvoted 2 times

✉️ **tolliekk** 3 weeks, 5 days ago

1. Create.
2. Build.
3. Push
4. Run

Instruction example hello-world image

<https://learn.microsoft.com/en-us/azure/container-registry/container-registry-quickstart-task-cli>

upvoted 5 times

✉️ **varinder82** 2 months ago

Final Answer:

- 1.Create a custom Dockerfile
- 2.Build the image
- 3.Push the image to an Azure container registry
- 4.Distribute a docker run script:

upvoted 3 times

✉️ **Murtuza** 2 months, 2 weeks ago

I think the it should be:

1. Pull the container image (in the dockerfile)
2. Create a custom docker file
3. Build the image
- 4 Push to ACR

upvoted 2 times

✉️ **evangelist** 3 months, 2 weeks ago

check Azure official on youtube: <https://www.youtube.com/watch?v=XLQLNazid4I>

The steps should be

1. Pull the container image (in the dockerfile)
2. Create a custom docker file

3. Build the image
4 Push to ACR
upvoted 8 times

✉ **qs99** 4 months ago
Was on exam 22 Jan 2024
upvoted 2 times

✉ **Rsp_** 6 months ago
Its totally confusing, what should be exact answer
upvoted 3 times

✉ **acsoma** 10 months ago
appeared in august exam
upvoted 5 times

✉ **mgafar** 1 year, 1 month ago
1) Create a custom Docker file
2) Build the image
3) Push the image to an Azure container registry
4) Distribute a docker run script
upvoted 7 times

✉ **dev2dev** 1 year, 4 months ago
Pulling container has to be first step. Otherwise the given options doesn't have enough steps to complete the task.
upvoted 1 times

✉ **KingChuang** 1 year, 4 months ago
on my exam (2023-01-16 Passed)

My Answer:
1.Pull the container image
2.Create a custom dockerfile
3.Push the image to container registry
4.Distrbute a custom docker run script
upvoted 8 times

✉ **cbhjffgfrj** 4 months, 1 week ago
if we don't perform docker build command then image is built so how you can push it to the container registry??
upvoted 2 times

✉ **Canyu** 1 year, 6 months ago
It is more effective to refer to these two documents:
<https://learn.microsoft.com/en-us/azure/container-registry/container-registry-quickstart-task-cli>
<https://learn.microsoft.com/en-us/azure/cognitive-services/anomaly-detector/anomaly-detector-container-howto>
Getting the container image is essential.
upvoted 1 times

Question #8**HOTSPOT -**

You plan to deploy a containerized version of an Azure Cognitive Services service that will be used for text analysis.

You configure <https://contoso.cognitiveservices.azure.com> as the endpoint URI for the service, and you pull the latest version of the Text Analytics Sentiment Analysis container.

You need to run the container on an Azure virtual machine by using Docker.

How should you complete the command? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
docker run --rm -it -p 5000:5000 --memory 8g --cpus 1 \
```

http://contoso.blob.core.windows.net
https://contoso.cognitiveservices.azure.com
mcr.microsoft.com/azure-cognitive-services/textanalytics/keyphrase
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment

Eula=accept \

Billing=

http://contoso.blob.core.windows.net
https://contoso.cognitiveservices.azure.com
mcr.microsoft.com/azure-cognitive-services/textanalytics/keyphrase
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment

ApiKey=xxxxxxxxxxxxxxxxxxxxxx

Correct Answer:**Answer Area**

```
docker run --rm -it -p 5000:5000 --memory 8g --cpus 1 \
```

http://contoso.blob.core.windows.net
https://contoso.cognitiveservices.azure.com
mcr.microsoft.com/azure-cognitive-services/textanalytics/keyphrase
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment

Eula=accept \

Billing=

http://contoso.blob.core.windows.net
https://contoso.cognitiveservices.azure.com
mcr.microsoft.com/azure-cognitive-services/textanalytics/keyphrase
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment

ApiKey=xxxxxxxxxxxxxxxxxxxxxx

Box 1: <mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment>

To run the Sentiment Analysis v3 container, execute the following docker run command. docker run --rm -it -p 5000:5000 --memory 8g --cpus 1 \ <mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment> \

Eula=accept \

Billing={ENDPOINT_URI} \

ApiKey={API_KEY} is the endpoint for accessing the Text Analytics API. <https://<your-custom-subdomain>.cognitiveservices.azure.com>

Box 2: <https://contoso.cognitiveservices.azure.com>

{ENDPOINT_URI} is the endpoint for accessing the Text Analytics API: <https://<your-custom-subdomain>.cognitiveservices.azure.com>. The endpoint for accessing the Text

Analytics API. <zure.com> -

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/how-tos/text-analytics-how-to-install-containers?tabs=sentiment>

✉  **kolakone**  2 years, 11 months ago

Seems correct.
For first answer, you "pull the latest version of the Text Analytics Sentiment Analysis container", hence the "sentiment" in the endpoint.
For billing, other options are other services, hence the provided answer.
upvoted 31 times

✉  **zellck**  11 months ago

1. mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment
2. https://contoso.cognitiveservices.azure.com

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/sentiment-opinion-mining/how-to/use-containers#run-the-container-with-docker-run>

```
docker run --rm -it -p 5000:5000 --memory 8g --cpus 1 \
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment:{IMAGE_TAG} \
Eula=accept \
Billing={ENDPOINT_URI} \
ApiKey={API_KEY}
```

- {ENDPOINT_URI}

The endpoint for accessing the API. You can find it on your resource's Key and endpoint page, on the Azure portal.

<https://<your-custom-subdomain>.cognitiveservices.azure.com>

upvoted 8 times

✉  **rdemontis** 6 months, 2 weeks ago

agree with you, thanks for references

upvoted 1 times

✉  **hidenori_music**  1 day, 10 hours ago

I took this exam on May 26, 2024. The questions were case study questions and the ETs were scattered throughout the exam without being summarized. The same questions were asked. Carefully review the questions and answers for the actual exam.

upvoted 1 times

✉  **hidenori_music** 1 day, 10 hours ago

I took this exam yesterday. The same case study questions were asked, although the ET has those questions scattered around instead of grouped together. Please note.

upvoted 1 times

✉  **funny_penguin** 3 days, 23 hours ago

was on exam today 24/05/2024, zellck's response is correct and the one i selected as well

upvoted 1 times

✉  **audlindr** 2 months, 3 weeks ago

This was in 2-Mar-24 exam

upvoted 2 times

✉  **Nihilist11** 3 months, 3 weeks ago

It was in today's exam -

upvoted 1 times

✉  **james2033** 9 months, 2 weeks ago

```
docker run --rm -it -p 5000:5000 --memory 8g --cpus 1 \
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment \
Eula=accept \
Billing=http://contoso.cognitiveservices.azure.com \
ApiKey=xxxxxxxxxxxxxxxxxx
```

See <https://learn.microsoft.com/en-us/azure/ai-services/language-service/sentiment-opinion-mining/how-to/use-containers#run-the-container-with-docker-run>

upvoted 4 times

✉  **acsoma** 10 months ago

appeared in august exam

upvoted 3 times

✉  **Pixelmate** 11 months ago

This appeared in the exam 28/06/2023

upvoted 3 times

✉  **AusAv** 1 year, 10 months ago

Example here: <https://docs.microsoft.com/en-us/azure/cognitive-services/language-service/sentiment-opinion-mining/how-to/use-containers#run-the-container-with-docker-run>

upvoted 1 times

✉  **Eltooth** 1 year, 10 months ago

mcr/.../sentiment
contoso.cognitiveservices
upvoted 1 times

 **Contactfornitish** 2 years, 4 months ago

Was on exam 02/01/2022
upvoted 1 times

 **sumanshu** 2 years, 4 months ago

As per question, we have to do the sentiment analysis, So URL should be appended with /sentiment and in Billing we have to mention the provider URL , So given answer is correct
upvoted 2 times

 **Phong0411** 2 years, 5 months ago

Was on exam 30/11/2021
upvoted 1 times

 **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021
upvoted 1 times

 **Satvik1992** 2 years, 11 months ago

this question appeared in the exam
upvoted 2 times

Question #9

You have the following C# method for creating Azure Cognitive Services resources programmatically.

```
static void create_resource(CognitiveServicesManagementClient client, string
resource_name, string kind, string account_tier, string location)
{
    CognitiveServicesAccount parameters =
        new CognitiveServicesAccount(null, null, kind, location, resource_name,
new CognitiveServicesAccountProperties(), new Sku(account_tier));
    var result = client.Accounts.Create(resource_group_name, account_tier,
parameters);
}
```

You need to call the method to create a free Azure resource in the West US Azure region. The resource will be used to generate captions of images automatically.

Which code should you use?

- A. create_resource(client, "res1", "ComputerVision", "F0", "westus")
- B. create_resource(client, "res1", "CustomVision.Prediction", "F0", "westus")
- C. create_resource(client, "res1", "ComputerVision", "S0", "westus")
- D. create_resource(client, "res1", "CustomVision.Prediction", "S0", "westus")

Correct Answer: B

Many of the Cognitive Services have a free tier you can use to try the service. To use the free tier, use F0 as the SKU for your resource.

There are two tiers of keys for the Custom Vision service. You can sign up for a F0 (free) or S0 (standard) subscription through the Azure portal.

Incorrect Answers:

A: There is no free tier (F0) for ComputerVision.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/cognitive-services-apis-create-account-client-library?pivots=programming-language-csharp> <https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/limits-and-quotas>

Community vote distribution

A (83%)

B (18%)

✉ **dinesh_tng** Highly Voted 2 years, 9 months ago

Answer shall be A, as there is free tier available for Computer Vision service.

- Free - Web/Container
 - 20 per minute
 - 5,000 free transactions per month
- upvoted 54 times

✉ **sumanshu** 2 years, 4 months ago

But what Feature ? It's not mention in Pricing Tier. It could be normal Computer Vision i.e. Boundary detection etc
upvoted 1 times

✉ **durak** 2 years, 6 months ago

Why not c?
upvoted 1 times

✉ **abelarda** 1 week, 6 days ago

C uses the standard tier, which is not free.
upvoted 1 times

✉ **Messatsu** 2 years, 9 months ago

Agree. <https://azure.microsoft.com/en-us/pricing/details/cognitive-services/computer-vision/>
upvoted 1 times

✉ **Rob77** 1 year ago

"Caption" is in preview and seems to be S3 tier so not free - see this
<https://azure.microsoft.com/en-gb/pricing/details/cognitive-services/computer-vision/>
upvoted 1 times

✉ **KingChuang** Highly Voted 1 year, 4 months ago

on my exam (2023-01-16 Passed)

My Answer:A

upvoted 12 times

- ✉ **hidenorimusic** Most Recent 1 day, 10 hours ago

I took this exam on May 26, 2024. The questions were case study questions and the ETs were scattered throughout the exam without being summarized. The same questions were asked. Carefully review the questions and answers for the actual exam.

upvoted 1 times

- ✉ **reiwanotora** 1 week, 1 day ago

Selected Answer: A

A is right.

upvoted 1 times

- ✉ **harnoor24** 1 week, 5 days ago

Selected Answer: A

Computer Vision (Also Called AI Vision now) has free tier

upvoted 1 times

- ✉ **JamesKJoker** 1 week, 6 days ago

Selected Answer: A

Generating Caption is not Possible with Custom Vision

upvoted 1 times

- ✉ **demonite** 2 weeks, 2 days ago

Answer is A

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/concept-describe-images-40?tabs=image>

upvoted 1 times

- ✉ **AzureGC** 1 month ago

Correct Answer is B;

* While the discussion below is about the FREE v. for-fee tiers, the answer key phrase is "captions" for images; To generate captions, one must use the CustomVisionPredictionClient. ComputerVision generates descriptions, !captions! See the services options.

<https://learn.microsoft.com/en-us/azure/ai-services/multi-service-resource?tabs=windows&pivots=programming-language-csharp#choose-a-service-and-pricing-tier>

upvoted 1 times

- ✉ **varinder82** 2 months ago

Final Answer:

A

upvoted 1 times

- ✉ **audlindr** 3 months, 1 week ago

Selected Answer: A

There is Free tier available for Computer Vision

upvoted 1 times

- ✉ **evangelist** 3 months, 2 weeks ago

Selected Answer: A

Functionality Focus: The task of generating captions for images automatically aligns more closely with the features offered by the Computer Vision API, specifically designed to analyze images and provide information about them, including descriptive captions. This functionality is not a focus of the Custom Vision service, which is more about applying custom models to new images for classification or object detection purposes.

upvoted 2 times

- ✉ **suzanne_exam** 4 months ago

I don't know why the answer says there is no free tier for computer vision. if you go to the portal and create a new resource the F0 option is there

upvoted 3 times

- ✉ **rdemontis** 6 months, 4 weeks ago

Selected Answer: A

to me the correct answer is A. Free plan Computer Vision

upvoted 1 times

- ✉ **james2033** 9 months, 2 weeks ago

Selected Answer: A

See sample source code at here

<https://learn.microsoft.com/en-us/azure/ai-services/multi-service-resource?pivots=programming-language-csharp&tabs=windows#call-management-methods>

Choose where has F0, then choose where has Computer Vision, Computer Vision for inference caption of image.

upvoted 2 times

- ✉ **acsoma** 10 months ago

appeared in august exam

upvoted 3 times

 **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/cognitive-services-apis-create-account-client-library?pivots=programming-language-csharp#create-a-cognitive-services-resource-c>

To create and subscribe to a new Cognitive Services resource, use the Create method. This method adds a new billable resource to the resource group you pass in. When creating your new resource, you'll need to know the "kind" of service you want to use, along with its pricing tier (or SKU) and an Azure location. The following method takes all of these as arguments and creates a resource.

upvoted 5 times

 **Pixelmate** 11 months ago

This appeared in the exam 28/06/2023

upvoted 6 times

Question #10

You successfully run the following HTTP request.

```
POST https://management.azure.com/subscriptions/18c51a87-3a69-47a8-aedc-a54745f708a1/resourceGroups/RG1/providers/Microsoft.CognitiveServices/accounts/contoso1/regenerateKey?api-version=2017-04-18  
Body>{"keyName": "Key2"}
```

What is the result of the request?

- A. A key for Azure Cognitive Services was generated in Azure Key Vault.
- B. A new query key was generated.
- C. The primary subscription key and the secondary subscription key were rotated.
- D. The secondary subscription key was reset.

Correct Answer: B

Accounts - Regenerate Key regenerates the specified account key for the specified Cognitive Services account.

Syntax:

POST

```
https://management.azure.com/subscriptions/{subscriptionId}/resourceGroups/{resourceGroupName}/providers/Microsoft.CognitiveServices/accounts/{accountName}/regenerateKey?api-version=2017-04-18
```

Reference:

<https://docs.microsoft.com/en-us/rest/api/cognitiveservices/accountmanagement/accounts/regeneratekey>

Community vote distribution

D (91%)

9%

✉  **LKLK10**  2 years, 11 months ago

I'd've thought it's D?

upvoted 41 times

✉  **sumanshu** 2 years, 4 months ago

What is mean by "Reset" ? I think RESET generally means, not re-generating the new-keys, but resetting to previous version probably ? (just guessing by the term reset). I don;t think D is correct here.

upvoted 3 times

✉  **Pyguy** 1 year, 5 months ago

There is no option a key to resetting a previous version , like you say. Reset = Regenerate same in API keys..

upvoted 1 times

✉  **ninjia** 1 year, 9 months ago

It's D. In the exam on 8/24/2022. The wording is more precise:

The new secondary subscription key was created.

upvoted 15 times

✉  **ovokpus** 2 years, 6 months ago

Yes,D is the answer

upvoted 4 times

✉  **ziiuai** 2 years, 8 months ago

B is wrong. Query key is for search service. The Rest request of create query key is like POST

```
https://management.azure.com/subscriptions/{subscriptionId}/resourceGroups/{resourceGroupName}/providers/Microsoft.Search/searchServices/{searchServiceName}/createQueryKey/{name}?api-version=2021-04-01-preview
```

See in <https://docs.microsoft.com/en-us/rest/api/searchmanagement/2021-04-01-preview/query-keys/create>

upvoted 8 times

✉  **LPreethi**  2 years, 10 months ago

Answer is correct - Sample response will be,

```
{  
  "key1": "KEY1",  
  "key2": "KEY2"  
}
```

This shows, Key1 is already there and this JSON request will generate Key2.

upvoted 13 times

✉️ **ghoppa** 2 years, 2 months ago

Key2 has always been there. They are automatically created by default when you create a resource. You cannot delete them or add new ones, you can only regenerate/refresh them (poorly worded here as RESET). Still, correct answer is D in my opinion.

upvoted 4 times

✉️ **reiwanotora** **Most Recent** 1 week, 1 day ago

Selected Answer: D

D is right.

upvoted 1 times

✉️ **PeteColag** 1 week, 2 days ago

Selected Answer: D

D. The secondary subscription key was reset.

This response is indicated by the "keyName": "Key2" in the request body, specifying that the secondary subscription key (Key2) should be regenerated.

upvoted 1 times

✉️ **TJ001** 2 weeks, 4 days ago

indeed this API call will regenerate the secondary key for the Cognitive Service Account..key1 remains same. could have been clear if it is mentioned as 'account' instead of 'subscription'

upvoted 1 times

✉️ **kiina** 4 weeks ago

The HTTP POST request you mentioned is specifically for regenerating the key of an Azure Cognitive Services account, not the Azure subscription itself. Here's a breakdown to clarify:

Azure Cognitive Services Key: This is the key you're regenerating with the request. Cognitive Services keys are used to authenticate API requests for the various cognitive services provided by Azure, like computer vision, language understanding, and more.

Azure Subscription Key: This key is different; it's associated with your Azure subscription and is used to manage billing and service usage across all services within Azure.

So, in summary, the request regenerates the secondary key (Key2) for the specified Cognitive Services account (contoso1), allowing continued access to Cognitive Services APIs with the new key. It does not affect the Azure subscription key. This is important for security practices, ensuring that if a key is compromised, you can regenerate it without interrupting services or access.

upvoted 1 times

✉️ **AzureGC** 1 month ago

B; This is an "account key" for Cognitive Services REST API. The API call will regenerate key2.

The references to "subscription" are a distractor, based on the management.azure.com/subscription/..

upvoted 1 times

✉️ **V3rgil** 2 months, 2 weeks ago

Selected Answer: D

MS Learn video suggest its D

upvoted 1 times

✉️ **cp2323** 2 months, 3 weeks ago

I think answer is B.

on the answer D - I think because it says subscription, I dont think its resetting subscription key as this is more to do with the resource key

upvoted 1 times

✉️ **evangelist** 3 months, 2 weeks ago

Selected Answer: D

D. The secondary subscription key was reset.

The request is to the Azure Management API to regenerate a key for an Azure Cognitive Services account (contoso1). The body of the request specifies {"keyName": "Key2"}, which indicates that the operation is targeted at the secondary subscription key (commonly referred to as Key2 in Azure Cognitive Services management). The regenerateKey action causes the specified key to be reset, generating a new key value for it while invalidating the old one.

upvoted 2 times

✉️ **Ody** 2 months, 1 week ago

I think your explanation is the most accurate. With respect to Azure Management API the are found under Subscription.

This question would have been easier for me if the answer had referred to it as an "account" key.

I think the answer is D.

upvoted 1 times

✉️ **rdemontis** 6 months, 4 weeks ago

Selected Answer: B

I vote for B because to me this is a new access/account/query key generated (the key2). and subscription key is not pertinent here along with the reset operation. It's a new key creation

upvoted 1 times

✉ **rdemontis** 6 months, 2 weeks ago

Sorry, I need to revise my previous answer. The correct answer for me is D, because a query key cannot be regenerated and the command in question is regenerateKey, specifically it regenerates Key2. This type of command refers to access keys or admin keys which are precisely 2 and can be regenerated. Whereas queries can only be created or deleted and I can create as many as I want not just 2
upvoted 2 times

✉ **katrang** 7 months, 1 week ago

Selected Answer: D

<https://learn.microsoft.com/en-us/rest/api/cognitiveservices/accountmanagement/accounts/regenerate-key?tabs=HTTP>
upvoted 1 times

✉ **[Removed]** 9 months, 3 weeks ago

B: <https://learn.microsoft.com/en-us/rest/api/cognitiveservices/accountmanagement/accounts/regenerate-key?tabs=HTTP>

Regenerates the specified account key for the specified Cognitive Services account:

POST

<https://management.azure.com/subscriptions/{subscriptionId}/resourceGroups/{resourceGroupName}/providers/Microsoft.CognitiveServices/accounts/{accountName}/regenerateKey?api-version=2023-05-01>

This has nothing to do with a subscription key.

upvoted 1 times

✉ **rdemontis** 6 months, 4 weeks ago

Yes but it's not even a query key. It is an access key or API key. The answers are unfortunately misplaced

upvoted 1 times

✉ **acsoma** 10 months ago

appeared in august exam

upvoted 2 times

✉ **zellck** 11 months ago

Selected Answer: D

D is the answer.

<https://learn.microsoft.com/en-us/rest/api/cognitiveservices/accountmanagement/accounts/regenerate-key?tabs=HTTP>
Regenerates the specified account key for the specified Cognitive Services account.

upvoted 4 times

✉ **tranatrana** 9 months ago

Your document does not say anything about resetting..

upvoted 1 times

✉ **rveney** 11 months, 1 week ago

The correct answer is C.

The result of the request is that the primary subscription key and the secondary subscription key were rotated. This means that the existing keys were invalidated and new keys were generated, with the primary key being replaced by a new key and the secondary key being replaced by the previous primary key. The request specifies the keyName as "Key2", indicating that the new key will be associated with the name "Key2".

upvoted 2 times

✉ **Drummer** 11 months, 2 weeks ago

The correct answer cannot be D, as the request does not reset the secondary subscription key. It only regenerates the authentication key named "Key2" for the specified Cognitive Services account. Therefore, the correct answer should be B - a new query key was generated.

upvoted 2 times

Question #11

You build a custom Form Recognizer model.

You receive sample files to use for training the model as shown in the following table.

Name	Type	Size
File1	PDF	20 MB
File2	MP4	100 MB
File3	JPG	20 MB
File4	PDF	100 MB
File5	GIF	1 MB
File6	JPG	40 MB

Which three files can you use to train the model? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. File1
- B. File2
- C. File3
- D. File4
- E. File5
- F. File6

Correct Answer: ACF

Input requirements -

Form Recognizer works on input documents that meet these requirements:

Format must be JPG, PNG, PDF (text or scanned), or TIFF. Text-embedded PDFs are best because there's no possibility of error in character extraction and location.

File size must be less than 50 MB.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/form-recognizer/overview>

Community vote distribution

ACF (88%)

12%

 **Eltooth**  1 year, 10 months ago

Selected Answer: ACF

File 2 and 5 are excluded.

New service limits now goes up to 500MB so...

File 1, 3, and 6 are correct for "training the model", however if MSFT remove the word "training" from the question - be careful.

<https://docs.microsoft.com/en-gb/learn/modules/work-form-recognizer/3-get-started>

<https://docs.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/service-limits?tabs=v21>

upvoted 13 times

 **nanaw770**  3 days, 18 hours ago

Selected Answer: ACF

1 3 6 YES!

upvoted 1 times

 **RupRizal** 6 months, 1 week ago

For custom model training the total size is still 50MB. Answer is correct

<https://learn.microsoft.com/en-us/azure/ai-services/document-intelligence/concept-custom-classifier?view=doc-intel-4.0.0>

upvoted 4 times

 **rdemontis** 6 months, 2 weeks ago

Selected Answer: ACF

Correct but probably outdated answer. Because considering current limit of 500 MD we can consider correct the D answer too.

<https://docs.microsoft.com/en-gb/learn/modules/work-form-recognizer/3-get-started>

upvoted 3 times

✉  **Tazmania98** 7 months ago

the question precises custom model ---- size limit 50MB
see this link <https://learn.microsoft.com/en-us/azure/ai-services/document-intelligence/service-limits?view=doc-intel-3.1.0>

Correct answers : ACF

upvoted 1 times

✉  **kiro_kocha** 9 months ago

Selected Answer: ACF
I think A,C,F are correct
upvoted 1 times

✉  **james2033** 9 months, 2 weeks ago

Selected Answer: CDF
Correct answer: ACDF
<https://learn.microsoft.com/en-gb/training/modules/work-form-recognizer/3-get-started>

Understand Form Recognizer file input requirements

Form Recognizer works on input documents that meet these requirements:

- Format must be JPG, PNG, BMP, PDF (text or scanned), or TIFF.
- The file size must be less than 500 MB for paid (S0) tier and 4 MB for free (F0) tier.
- Image dimensions must be between 50 x 50 pixels and 10000 x 10000 pixels.
- The total size of the training data set must be 500 pages or less.

upvoted 3 times

✉  **acsoma** 10 months ago

appeared in august exam
upvoted 3 times

✉  **james2033** 10 months, 2 weeks ago

Selected Answer: ACF
Based on filetype: PDF, JPG; and less than 50 MB of file size.
upvoted 2 times

✉  **zellck** 11 months ago

Selected Answer: ACF
ACF is the answer.

<https://learn.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/concept-layout?view=form-recog-3.0.0#input-requirements>

Supported file formats:

- PDF
- Image: JPEG/JPG, PNG, BMP, and TIFF

upvoted 1 times

✉  **RAN_L** 1 year, 2 months ago

Selected Answer: ACF
The supported file types for training a Form Recognizer model are PDF, PNG, and JPEG. Therefore, you can use only the files that have one of these formats for training the model. Based on the information provided, the files that can be used for training the model are:

- A. File1 (PDF)
- C. File3 (JPG)
- F. File6 (JPEG)

Therefore, the correct answers are:

- A. File1
- C. File3
- F. File6

upvoted 1 times

✉  **AusAv** 1 year, 10 months ago

It should allow all except the GIF and MP4 on the paid tier, free tier supports file up to 4MB: <https://docs.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/concept-read#input-requirements>

upvoted 2 times

"For custom model training, the total size of training data is 50 MB for template model and 1G-MB for the neural model."

"The file size for analyzing documents must be less than 500 MB for paid (S0) tier and 4 MB for free (F0) tier." -> But question talks about training the model, so I think the 50MB applies, File 4 can't be used.

upvoted 1 times

✉  **DingDongSingSong** 2 years ago

Supported file formats: JPEG/JPG, PNG, BMP, TIFF, and PDF (text-embedded or scanned). So there is more than 3 correct responses. PDF and JPG files (2 each, therefore 4 possible responses even though question asks for 3)

upvoted 3 times

✉  **fux** 1 year, 10 months ago

The PDF must be capped at 50mb, so in reality, there are only 3 answers
upvoted 5 times

 **james2033** 9 months, 2 weeks ago

No, PDF size is limit 500 MB now, see <https://learn.microsoft.com/en-gb/training/modules/work-form-recognizer/3-get-started>.
upvoted 2 times

 **ovtchinnikov** 2 years ago

1 & 3 correct.
But GIFs are also allowed, if I remember right.
upvoted 1 times

Question #12

A customer uses Azure Cognitive Search.

The customer plans to enable a server-side encryption and use customer-managed keys (CMK) stored in Azure.

What are three implications of the planned change? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. The index size will increase.
- B. Query times will increase.
- C. A self-signed X.509 certificate is required.
- D. The index size will decrease.
- E. Query times will decrease.
- F. Azure Key Vault is required.

Correct Answer: ABE

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-security-overview>

Community vote distribution

ABF (94%)

6%

✉️  **methodidacte**  2 years, 9 months ago

Should be ABF (use AKV for customer managed key)
upvoted 48 times

✉️  **SuperPetey** 2 years, 9 months ago

correct : "Customer-managed keys require an additional billable service, Azure Key Vault, which can be in a different region, but under the same subscription, as Azure Cognitive Search. Enabling CMK encryption will increase index size and degrade query performance." same document also lists Azure Key Vault as a requirement:

<https://docs.microsoft.com/en-us/azure/search/search-security-overview#data-protection>
upvoted 13 times

✉️  **TanujitRoy** 2 years, 6 months ago

Bhai tame bahut katha kahucha....tame ki gangadhara
upvoted 2 times

✉️  **adarshsahoo** 2 years, 6 months ago

Na mu hauchi Lionel Toppo
upvoted 1 times

✉️  **Jyotiran** 10 months, 1 week ago

Chodi heigala
upvoted 1 times

✉️  **bhadrachagala** 2 months, 1 week ago

UHUN KI KATHA EIGUDIKA, LAJYA LAJYA
upvoted 1 times

✉️  **rdemontis** 6 months, 4 weeks ago

Thanks for the documentation attached
upvoted 1 times

✉️  **nanaw770**  3 days, 18 hours ago

Selected Answer: ABF
Coffee gift is ABF.
upvoted 1 times

✉️  **CDL_Learner** 1 month, 2 weeks ago

Based on the information from Microsoft's documentation¹², the implications of enabling server-side encryption and using customer-managed keys (CMK) stored in Azure for Azure Cognitive Search are:

- A. The index size will increase. Enabling CMK encryption increases the index size¹.
- B. Query times will increase. Enabling CMK encryption degrades query performance. You can expect to see an increase of 30-60 percent in query times¹.
- F. Azure Key Vault is required. CMK encryption depends on Azure Key Vault. You can create your own encryption keys and store them in a key vault

or you can use Azure Key Vault APIs to generate encryption keys¹.

So, the correct answers are A, B, and F. Please note that these changes might also lead to additional costs². It's always a good idea to review the Azure pricing details before making such changes.

upvoted 1 times

✉ **85e4c91** 3 months, 1 week ago

Based on the Azure documentation and best practices for Azure Cognitive Search with server-side encryption using customer-managed keys (CMKs), it appears that the correct answers to the question are B and F. Enabling CMKs is known to potentially increase query times due to the additional encryption/decryption steps required when data is requested (B). Azure Key Vault is required to store and manage these CMKs, making a necessary component of the configuration (F). There is no indication in the documentation that index size will be affected (A and D), nor that query times will decrease (E), and a self-signed X.509 certificate is not mentioned as a requirement (C).

upvoted 1 times

✉ **audlindr** 3 months, 1 week ago

Selected Answer: ABF

ABE can't be the right answer. B and E are conflicting

upvoted 1 times

✉ **suzanne_exam** 4 months ago

I think this answer highlights why you should come to the comments and do your own research. The answer here says the query time will both increase AND decrease, which doesn't make sense. The answer is ABF, not ABE

upvoted 2 times

✉ **AppSphere23** 5 months, 2 weeks ago

The query PERFORMANCE will decrease, a typo?

upvoted 1 times

✉ **sca88** 6 months ago

Selected Answer: ABF

The E is wrong! The query time will increase with CMK

upvoted 1 times

✉ **anewfriend** 6 months, 1 week ago

Selected Answer: ABF

Customer-managed keys require another billable service, Azure Key Vault, which can be in a different region, but under the same subscription, as Azure AI Search.

Enabling CMK encryption will increase index size and degrade query performance. Based on observations to date, you can expect to see an increase of 30-60 percent in query times, although actual performance will vary depending on the index definition and types of queries. Because of the negative performance impact, we recommend that you only enable this feature on indexes that really require it.

upvoted 1 times

✉ **rdemontis** 6 months, 4 weeks ago

Selected Answer: ABF

Correct answer to me are A,B,F. AKV needed to store CMK

upvoted 1 times

✉ **trashbox** 7 months ago

Selected Answer: ABF

A,B: "Enabling CMK encryption will increase index size and degrade query performance"

F: "Customer-managed keys require another billable service, Azure Key Vault, which can be in a different region, but under the same subscription, as Azure Cognitive Search."

<https://learn.microsoft.com/en-us/azure/search/search-security-overview#data-protection>

upvoted 1 times

✉ **[Removed]** 7 months ago

Why would the index size increase? You are not adding any data to it...

upvoted 1 times

✉ **JonHanes** 4 months, 2 weeks ago

Enforcing Customer-Managed Key (CMK) encryption can increase the index size due to the process of double encryption. Here's a brief explanation:

When CMK is enforced, Azure Search uses service-managed keys to encrypt the data first (this is the primary encryption). Then, it applies a second layer of encryption using the customer-managed keys (this is the secondary encryption). This double encryption process increases the size of the index because it adds additional metadata and requires more storage space to accommodate the encrypted data.

upvoted 2 times

✉ **katrang** 7 months, 1 week ago

Selected Answer: ABF

Enabling CMK encryption will increase index size and degrade query performance. And you need a key vault to store the keys

<https://learn.microsoft.com/en-us/azure/search/search-security-overview#data-protection>

upvoted 1 times

✉ **kiro_kocha** 9 months ago

Selected Answer: ABF

Index and query will increase because of the encryption.
Because customer wants to manage the keys he will need Azure key vault
upvoted 2 times

✉ **EliteAllen** 10 months, 1 week ago

Selected Answer: ABF

ABF seem correct
upvoted 2 times

✉ **zellck** 11 months ago

Selected Answer: ABF

ABF is the answer.

<https://learn.microsoft.com/en-us/azure/search/search-security-manage-encryption-keys>
Azure Cognitive Search automatically encrypts data at rest with service-managed keys. If more protection is needed, you can supplement default encryption with another encryption layer using keys that you create and manage in Azure Key Vault.

<https://learn.microsoft.com/en-us/azure/search/search-security-manage-encryption-keys?tabs=portal-pp%2Cmanaged-id-sys#work-with-encrypted-content>

With customer-managed key encryption, you'll notice latency for both indexing and queries due to the extra encrypt/decrypt work. Azure Cognitive Search doesn't log encryption activity, but you can monitor key access through key vault logging. We recommend that you enable logging as part of key vault configuration.

upvoted 2 times

✉ **zellck** 11 months ago

<https://learn.microsoft.com/en-us/azure/search/search-security-manage-encryption-keys?tabs=portal-pp%2Cmanaged-id-sys#full-double-encryption>

Enabling CMK encryption will increase index size and degrade query performance. Based on observations to date, you can expect to see an increase of 30-60 percent in query times, although actual performance will vary depending on the index definition and types of queries. Because of the negative performance impact, we recommend that you only enable this feature on indexes that really require it.

upvoted 1 times

✉ **RemcoGoy** 1 year ago

Selected Answer: ABF

ABF looks right, B and E are contradicting
upvoted 2 times

Question #13

You are developing a new sales system that will process the video and text from a public-facing website.

You plan to notify users that their data has been processed by the sales system.

Which responsible AI principle does this help meet?

- A. transparency
- B. fairness
- C. inclusiveness
- D. reliability and safety

Correct Answer: D

Reference:

<https://docs.microsoft.com/en-us/azure/cloud-adoption-framework/strategy/responsible-ai>

Community vote distribution



✉️ **SuperPeteY** Highly Voted 2 years, 9 months ago

The correct answer is A, transparency: "When an AI application relies on personal data, such as a facial recognition system that takes images of people to recognize them; you should make it clear to the user how their data is used and retained, and who has access to it." from: <https://docs.microsoft.com/en-us/learn/paths/prepare-for-ai-engineering/>

upvoted 71 times

✉️ **ayoitu** Highly Voted 2 years, 9 months ago

"Transparency: AI systems should be understandable."
"Reliability and safety: AI systems should perform reliably and safely."
so the answer is correct, reliability and safety

<https://docs.microsoft.com/en-us/azure/cloud-adoption-framework/strategy/responsible-ai>

upvoted 13 times

✉️ **TanujitRoy** 2 years, 6 months ago

Bhai mo ra net ta chalu ni asiki thk kari daba ki...sei azure ra gamucha tool ta nei asiba
upvoted 1 times

✉️ **PeteColag** Most Recent 3 days, 9 hours ago

Selected Answer: A

Transparency, because you are informing users that their data is being processed.
upvoted 1 times

✉️ **audlindr** 3 months, 1 week ago

Selected Answer: A

Transparency is the right answer.
Transparency: AI systems should be understandable.
upvoted 2 times

✉️ **evangelist** 3 months, 2 weeks ago

Selected Answer: A

Notifying users that their data has been processed by the sales system helps meet the A. transparency principle. Transparency involves informing users about how their data is used, processed, and the purpose behind it, thereby fostering trust and understanding between the technology providers and the users.

upvoted 3 times

✉️ **evangelist** 3 months, 2 weeks ago

Selected Answer: A

Notifying users that their data has been processed by the sales system helps meet the A. transparency principle. Transparency involves informing users about how their data is used, processed, and the purpose behind it, thereby fostering trust and understanding between the technology providers and the users.

upvoted 2 times

✉️ **orionduo** 5 months, 3 weeks ago

Selected Answer: A

Transparency.
AI systems should be understandable. Users should be made fully aware of the purpose of the system, how it works, and what limitations may be expected.
REF: <https://learn.microsoft.com/en-us/training/modules/prepare-to-develop-ai-solutions-azure/5-understand-considerations-for-responsible-ai>

upvoted 2 times

✉ **rdemontis** 6 months, 4 weeks ago

Selected Answer: A

Transparency --> need to notify the users

upvoted 2 times

✉ **NickYog** 8 months, 3 weeks ago

Selected Answer: B

System should treat everyone fairly

upvoted 1 times

✉ **Pinkshark** 8 months, 4 weeks ago

Selected Answer: A

correct for me A

upvoted 1 times

✉ **kiro_kocha** 9 months ago

Selected Answer: A

Transparency offcourse. Users should be notified what information will be processed by the ai solution.

upvoted 1 times

✉ **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 5 times

✉ **savetheplanet** 11 months, 1 week ago

A

it's clearly transparency, why this website selects the wrong answer?

upvoted 2 times

✉ **marti_tremblay000** 1 year, 2 months ago

Selected Answer: A

ChatGPT answer :

The responsible AI principle that notifying users that their data has been processed by the sales system helps meet is A. transparency.

Transparency is the principle of making AI systems understandable and providing clear explanations of their decisions and actions. By notifying users that their data has been processed by the sales system, you are being transparent about the fact that their data is being used and how it is being used. This helps users understand how their data is being used and can help build trust between users and the sales system.

upvoted 1 times

✉ **ap1234pa** 1 year, 4 months ago

Selected Answer: A

A is correct

upvoted 1 times

✉ **slcheng** 1 year, 5 months ago

Agreed with A

upvoted 1 times

✉ **HotDurian** 1 year, 5 months ago

Correct answer is Transparency.

upvoted 2 times

Question #14

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You create a web app named app1 that runs on an Azure virtual machine named vm1. Vm1 is on an Azure virtual network named vnet1.

You plan to create a new Azure Cognitive Search service named service1.

You need to ensure that app1 can connect directly to service1 without routing traffic over the public internet.

Solution: You deploy service1 and a public endpoint to a new virtual network, and you configure Azure Private Link.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

The Azure Private Link should use a private endpoint, not a public endpoint.

Private Link service can be accessed from approved private endpoints in any public region.

Reference:

<https://docs.microsoft.com/en-us/azure/private-link/private-link-overview>

Community vote distribution

B (88%)

13%

✉ **WynterTsai** Highly Voted 2 years, 9 months ago

Answer is no. you should create a private link with private endpoint
upvoted 31 times

✉ **MII1975** 2 years, 6 months ago

I agree
<https://docs.microsoft.com/en-us/azure/cognitive-services/cognitive-services-virtual-networks?tabs=portal#use-private-endpoints>
upvoted 1 times

✉ **nanaw770** Most Recent 3 days, 18 hours ago

Selected Answer: B

This is correct answer is no. All but M7, T20, and T25 graduates must give a silly answer of yes.
upvoted 1 times

✉ **demonite** 1 week ago

Answer is a big YES. Like other PaaS services in Azure, you can deploy your service with public access (i.e. in this case they just call it public endpoint) and configure your private access after with a private endpoint. Doing so, the private endpoint will block the public access.

Private endpoints for your search service enable you to:

Block all connections on the public endpoint for your search service.

<https://learn.microsoft.com/en-us/azure/search/service-create-private-endpoint#why-use-a-private-endpoint>

upvoted 1 times

✉ **natgurulearning** 5 months, 4 weeks ago

A is correct!

That solution would work perfectly. Here's why: By deploying the search service and the endpoint to the same virtual network, you create a private secure connection between the web app and the search service. And by using Azure Private Link, you're creating a private, dedicated connection between the two services, which means traffic doesn't have to go over the public internet. This not only increases security but also improves performance and reliability. So, deploying the service to a new virtual network and configuring Azure Private Link would definitely meet the goal of connecting app1 directly to service1 without using the public internet.

upvoted 3 times

✉ **Ody** 2 months, 1 week ago

I think you are correct. On the similar questions in the scenario the answer is definitely no.

The question just says you need to ensure you "can connect". It doesn't say anything at all about preventing connectivity from the Internet.

It does not appear that you can create the service without a public endpoint.

upvoted 1 times

✉ **Ody** 2 months, 1 week ago

I misread. I think B is correct. When deploying the search service it lets you pick public or private endpoint.

upvoted 1 times

 **rdemontis** 6 months, 4 weeks ago

Selected Answer: B

correct. If you need a private link you should create a private endpoint

upvoted 1 times

 **trashbox** 7 months ago

Selected Answer: B

Not a public endpoint but a private endpoint.

"You can use private endpoints for your Azure AI services resources to allow clients on a virtual network to securely access data over Azure Private Link."

<https://learn.microsoft.com/en-us/azure/ai-services/cognitive-services-virtual-networks?tabs=portal#use-private-endpoints>

upvoted 1 times

 **acsoma** 10 months ago

appeared in august exam

upvoted 3 times

 **zellck** 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/search/service-create-private-endpoint#why-use-a-private-endpoint-for-secure-access>

Private Endpoints for Azure Cognitive Search allow a client on a virtual network to securely access data in a search index over a Private Link. The private endpoint uses an IP address from the virtual network address space for your search service. Network traffic between the client and the search service traverses over the virtual network and a private link on the Microsoft backbone network, eliminating exposure from the public internet.

upvoted 3 times

 **Eltooth** 1 year, 10 months ago

Selected Answer: B

B is correct answer : No.

upvoted 2 times

 **mohamedba** 1 year, 11 months ago

Public endpoint means the traffic is going to Internet, so answer is a big 'NO'

upvoted 2 times

 **arpitexam** 2 years, 2 months ago

Selected Answer: B

private link with private endpoint required

upvoted 2 times

 **reachmymind** 2 years, 2 months ago

Selected Answer: A

Yes

Traffic between your virtual network and the service travels the Microsoft backbone network. Exposing your service to the public internet is no longer necessary. You can create your own private link service in your virtual network and deliver it to your customers. Setup and consumption using Azure Private Link is consistent across Azure PaaS, customer-owned, and shared partner services.

<https://docs.microsoft.com/en-us/azure/private-link/private-link-overview>

upvoted 1 times

 **Nexxiss** 2 years, 3 months ago

Selected Answer: B

I'd say no

upvoted 4 times

 **Deepusuraj** 2 years, 3 months ago

Selected Answer: A

A is correct answer

upvoted 1 times

 **sumanshu** 2 years, 4 months ago

Vote for 'B'

No Need for a Public Endpoint. Private Link between Service and Virtual network can handle it.

upvoted 1 times

 **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 1 times

 **mikegsm** 2 years, 6 months ago

SEEMS nO

upvoted 1 times

Question #15

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You create a web app named app1 that runs on an Azure virtual machine named vm1. Vm1 is on an Azure virtual network named vnet1.

You plan to create a new Azure Cognitive Search service named service1.

You need to ensure that app1 can connect directly to service1 without routing traffic over the public internet.

Solution: You deploy service1 and a public endpoint, and you configure an IP firewall rule.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

Instead deploy service1 and a private (not public) endpoint to a new virtual network, and you configure Azure Private Link.

Reference:

<https://docs.microsoft.com/en-us/azure/private-link/private-link-overview>

Community vote distribution

 B (100%)

✉  **acsoma**  10 months ago

appeared in august exam
upvoted 5 times

✉  **rdemontis**  6 months, 4 weeks ago

Selected Answer: B
correct answer
upvoted 1 times

✉  **zellick** 11 months ago

Selected Answer: B
B is the answer.

<https://learn.microsoft.com/en-us/azure/search/service-create-private-endpoint#why-use-a-private-endpoint-for-secure-access>

Private Endpoints for Azure Cognitive Search allow a client on a virtual network to securely access data in a search index over a Private Link. The private endpoint uses an IP address from the virtual network address space for your search service. Network traffic between the client and the search service traverses over the virtual network and a private link on the Microsoft backbone network, eliminating exposure from the public internet.

upvoted 4 times

✉  **Eltooth** 1 year, 10 months ago

Selected Answer: B
B is correct answer : No.
upvoted 2 times

✉  **mohamedba** 1 year, 11 months ago

Again NO
upvoted 1 times

✉  **timmayy54** 2 years, 4 months ago

Correct Answer is B. No. this scenario routes over public internet, to do this without touching public internet you would use a private endpoint on vnet then private link to access it.
upvoted 3 times

✉  **sumanshu** 2 years, 4 months ago

Vote for 'B'. Private Link is sufficient
upvoted 1 times

✉  **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021
upvoted 1 times

✉  **mikegsm** 2 years, 6 months ago

SEEMS NO

upvoted 2 times

 **dinesh_tng** 2 years, 8 months ago

Yes, is the answer.

upvoted 1 times

 **Ab_S** 2 years, 8 months ago

Answer is correct.

Instead created private end point and private link.

upvoted 4 times

Question #16

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You create a web app named app1 that runs on an Azure virtual machine named vm1. Vm1 is on an Azure virtual network named vnet1.

You plan to create a new Azure Cognitive Search service named service1.

You need to ensure that app1 can connect directly to service1 without routing traffic over the public internet.

Solution: You deploy service1 and a public endpoint, and you configure a network security group (NSG) for vnet1.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

Instead deploy service1 and a private (not public) endpoint to a new virtual network, and you configure Azure Private Link.

Reference:

<https://docs.microsoft.com/en-us/azure/private-link/private-link-overview>

Community vote distribution

B (100%)

✉  **acsoma**  10 months ago

appeared in august exam

upvoted 5 times

✉  **rdemontis**  6 months, 4 weeks ago

Selected Answer: B

correct answer

upvoted 1 times

✉  **trashbox** 7 months ago

Selected Answer: B

Use a private endpoint instead of a public endpoint to prevent from routing traffic over the public internet.

upvoted 3 times

✉  **zellck** 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/search/service-create-private-endpoint#why-use-a-private-endpoint-for-secure-access>

Private Endpoints for Azure Cognitive Search allow a client on a virtual network to securely access data in a search index over a Private Link. The private endpoint uses an IP address from the virtual network address space for your search service. Network traffic between the client and the search service traverses over the virtual network and a private link on the Microsoft backbone network, eliminating exposure from the public internet.

upvoted 4 times

✉  **mohamedba** 1 year, 11 months ago

Again NO

upvoted 2 times

✉  **DingDongSingSong** 2 years ago

The answer is B. Even if you configure vnet with a network security group, the app is still routing traffic to the public until it hits the vnet1 where then the network goes secure/private. That still does not address "no public routing of traffic" requirement

upvoted 2 times

✉  **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 1 times

✉  **mikegsm** 2 years, 6 months ago

SEEMS YES

upvoted 4 times

✉  **Derin_tade** 2 years, 8 months ago

I think this is the answer based on these two links.

<https://docs.microsoft.com/en-us/azure/virtual-network/network-security-groups-overview#network-security-groups>

<https://docs.microsoft.com/en-us/azure/virtual-network/virtual-networks-overview>

upvoted 1 times

 **Derin_tade** 2 years, 8 months ago

I mean I think this should be Yes instead of No. And the previous ones should be No.

upvoted 4 times

Question #17

You plan to perform predictive maintenance.

You collect IoT sensor data from 100 industrial machines for a year. Each machine has 50 different sensors that generate data at one-minute intervals. In total, you have 5,000 time series datasets.

You need to identify unusual values in each time series to help predict machinery failures.

Which Azure service should you use?

- A. Anomaly Detector
- B. Cognitive Search
- C. Form Recognizer
- D. Custom Vision

Correct Answer: A

Community vote distribution

A (100%)

✉  **zellck**  11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/anomaly-detector/overview>

Anomaly Detector is an AI service with a set of APIs, which enables you to monitor and detect anomalies in your time series data with little machine learning (ML) knowledge, either batch validation or real-time inference.

upvoted 5 times

✉  **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 2 times

✉  **reiwanotora**  1 week, 1 day ago

Selected Answer: A

A is right.

upvoted 1 times

✉  **CDL_Learner** 1 month, 2 weeks ago

Selected Answer: A

The best Azure service to use in this scenario is A. Anomaly Detector1.

upvoted 1 times

✉  **rdemontis** 6 months, 4 weeks ago

Selected Answer: A

correct Anomaly Detector

upvoted 1 times

✉  **josemiguelch** 1 year, 1 month ago

Azure Anomaly Detector is the recommended Azure service to use for this scenario.

upvoted 1 times

✉  **KingChuang** 1 year, 4 months ago

on my exam (2023-01-16 Passed)

My Answer:A

upvoted 4 times

✉  **kml2003** 1 year, 5 months ago

Selected Answer: A

Anomaly detection

upvoted 2 times

✉  **Eltooth** 1 year, 10 months ago

Selected Answer: A

A is correct answer : Anomaly Detector.

upvoted 1 times

✉️  **happychuks** 1 year, 10 months ago

Selected Answer: A

<https://azure.microsoft.com/en-us/services/cognitive-services/anomaly-detector/>
upvoted 2 times

✉️  **satishk4u** 2 years ago

Was on exam on 03-May-2022
upvoted 1 times

Question #18

HOTSPOT -

You are developing a streaming Speech to Text solution that will use the Speech SDK and MP3 encoding.

You need to develop a method to convert speech to text for streaming MP3 data.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
var audioFormat =   
  
  
  
  
var speechConfig = SpeechConfig.FromSubscription("18c51a87-3a69-47a8-aedc-a54745f708a1", "westus");  
  
var audioConfig = AudioConfig.FromStreamInput(pushStream, audioFormat);  
  
using (var recognizer = new   
  
  
  
(speechConfig, audioConfig))  
  
{  
  
    var result = await recognizer.RecognizeOnceAsync();  
  
    var text = result.Text;  
  
}
```

Correct Answer:

Answer Area

```
var audioFormat =   
  
  
  
  
var speechConfig = SpeechConfig.FromSubscription("18c51a87-3a69-47a8-aedc-a54745f708a1", "westus");  
  
var audioConfig = AudioConfig.FromStreamInput(pushStream, audioFormat);  
  
using (var recognizer = new   
  
  
  
(speechConfig, audioConfig))  
  
{  
  
    var result = await recognizer.RecognizeOnceAsync();  
  
    var text = result.Text;  
  
}
```

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-use-codec-compressed-audio-input-streams?tabs=debian&pivots=programming-language-csharp>

 **htolajide**  2 years, 8 months ago

The answer is correct

upvoted 23 times

 **MII1975** 2 years, 6 months ago

I agree

GetCompressedFormat

<https://docs.microsoft.com/en-us/dotnet/api/microsoft.cognitiveservices.speech.audio.audiostreamformat.getcompressedformat?view=azure-dotnet>

SpeechRecognizer

https://docs.microsoft.com/en-us/dotnet/api/microsoft.cognitiveservices.speech.speechrecognizer.-ctor?view=azure-dotnet#Microsoft_CognitiveServices_Speech_SpeechRecognizer_ctor_Microsoft_CognitiveServices_Speech_SpeechConfig_Microsoft_CognitiveServices_Speech_Audio_AudioConfig_

upvoted 7 times

✉  **rdemontis** 6 months, 4 weeks ago

correct answer. Thanks for the provided documentation

upvoted 2 times

✉  **sumanshu**  2 years, 4 months ago

We need to convert Streaming MP3 Data.

So, option 1 is eliminated, as it's not a Audio Streaming function. rest 3 contains the keyword Stream.

But, default audio stream format is WAV, and here we are passing MP3, So other then WAV, we need to pass compressed audio format, So correct answer is : `AudioStreamFormat.GetCompressedFormat`.

and we need to recognize the speech to convert into the text - so, Speech Recognizer.

upvoted 10 times

✉  **CDL_Learner**  1 month, 2 weeks ago

`AudioStreamFormat.GetCompressedFormat(AudioStreamContainerFormat.MP3)`: This option is selected because the audio data is in MP3 format, which is a compressed audio format. The `GetCompressedFormat` method is used to get the format of the compressed audio data.

`AudioConfig SetProperty`: This is used to set a property of the `AudioConfig` object, not to get the format of the audio data.

`AudioStreamFormat.GetWaveFormatPCM`: This is used to get the format of PCM audio data, not MP3.

`PullAudioInputStream`: This is used to create a pull audio input stream, not to get the format of the audio data.

upvoted 1 times

✉  **Prodyna** 6 months, 2 weeks ago

appeared in november exam

upvoted 4 times

✉  **acsoma** 10 months ago

appeared in august exam

upvoted 4 times

✉  **zellck** 11 months ago

1. `AudioStreamFormat.GetCompressedFormat`.

2. `SpeechRecognizer`

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-use-codec-compressed-audio-input-streams>

<https://learn.microsoft.com/en-us/dotnet/api/microsoft.cognitiveservices.speech.speechrecognizer?view=azure-dotnet>

Transcribes speech into text. Speech can arrive via microphone, audio file, or other audio input stream.

upvoted 5 times

✉  **DS_newb** 1 year, 1 month ago

It is no longer in exam

upvoted 1 times

✉  **DS_newb** 1 year, 1 month ago

wait the respond should in face related question, why appear in here?

upvoted 1 times

✉  **Eltooth** 1 year, 10 months ago

`AudioStream.GetCompressedFormat` (MP3 is compressed audio file)

`SpeechRecognizer`

upvoted 2 times

✉  **ashu789** 2 years, 5 months ago

was on exam 05/12/2021

upvoted 1 times

✉  **Phong0411** 2 years, 5 months ago

WAs on exam 30/11/2021

upvoted 1 times

✉  **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 1 times

Question #19

HOTSPOT -

You are developing an internet-based training solution for remote learners.

Your company identifies that during the training, some learners leave their desk for long periods or become distracted.

You need to use a video and audio feed from each learner's computer to detect whether the learner is present and paying attention. The solution must minimize development effort and identify each learner.

Which Azure Cognitive Services service should you use for each requirement? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

From a learner's video feed, verify whether the learner is present:

Face
Speech
Text Analytics

From a learner's facial expression in the video feed, verify whether the learner is paying attention:

Face
Speech
Text Analytics

From a learner's audio feed, detect whether the learner is talking:

Face
Speech
Text Analytics

Answer Area

From a learner's video feed, verify whether the learner is present:

Face
Speech
Text Analytics

Correct Answer:

From a learner's facial expression in the video feed, verify whether the learner is paying attention:

Face
Speech
Text Analytics

From a learner's audio feed, detect whether the learner is talking:

Face
Speech
Text Analytics

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/what-are-cognitive-services>

 **htolajide** Highly Voted 2 years, 8 months ago

The answer is correct

upvoted 19 times

 **zellck** Highly Voted 11 months ago

1. Face

2. Face

3. Speech

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/overview-identity#face-detection-and-analysis>
Face detection is required as a first step in all the other scenarios. The Detect API detects human faces in an image and returns the rectangle coordinates of their locations. It also returns a unique ID that represents the stored face data. This is used in later operations to identify or verify faces.

Optionally, face detection can extract a set of face-related attributes, such as head pose, age, emotion, facial hair, and glasses. These attributes are general predictions, not actual classifications. Some attributes are useful to ensure that your application is getting high-quality face data when users add themselves to a Face service. For example, your application could advise users to take off their sunglasses if they're wearing sunglasses.
upvoted 11 times

✉ **nanaw770** Most Recent 3 days, 18 hours ago

Face Face Speech
upvoted 1 times

✉ **rdemontis** 6 months, 4 weeks ago

correct answer
upvoted 3 times

✉ **trashbox** 7 months ago

The answer is correct. Face, Face, and Speech. What a terrible application :(
upvoted 2 times

✉ **Eltooth** 1 year, 10 months ago

Face
Face
Speech
upvoted 2 times

✉ **Coderhbti** 2 years, 3 months ago

From Video feed - Face
Facial Expression from - Face
Audio Feed is - Speech
upvoted 3 times

✉ **sumanshu** 2 years, 4 months ago

From Video feed - Face
Facial Expression from - Face
Audio Feed is - Speech
upvoted 3 times

✉ **mikegsm** 2 years, 6 months ago

Correct
upvoted 1 times

✉ **Adedoyin_Simeon** 2 years, 7 months ago

The answer is correct
upvoted 2 times

✉ **RomanXXX** 2 years, 8 months ago

Seems ok
upvoted 3 times

Question #20

You plan to provision a QnA Maker service in a new resource group named RG1.

In RG1, you create an App Service plan named AP1.

Which two Azure resources are automatically created in RG1 when you provision the QnA Maker service? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

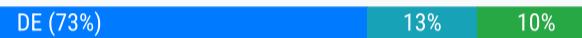
- A. Language Understanding
- B. Azure SQL Database
- C. Azure Storage
- D. Azure Cognitive Search
- E. Azure App Service

Correct Answer: DE

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/how-to/set-up-qnamaker-service-azure?tabs=v1#delete-azure-resources>

Community vote distribution



✉️ **methodidacte** Highly Voted 2 years, 8 months ago

Right answer : at the creation, we have to precise Azure Search an Azure Web App details.

"When you create a QnAMaker resource, you host the data in your own Azure subscription. Azure Search is used to index your data." & "When you create a QnAMaker resource, you host the runtime in your own Azure subscription. App Service is the compute engine that runs the QnA Maker queries for you."

upvoted 15 times

✉️ **ghoppa** Highly Voted 2 years, 2 months ago

Another one to memorise. A quick trick:

Q sounds like a C, for Cognitive

n

A for App Service

upvoted 12 times

✉️ **dev2dev** 1 year, 4 months ago

lol then there is a risk of selecting A and C as answers :)

upvoted 5 times

✉️ **nanaw770** Most Recent 3 days, 18 hours ago

Selected Answer: DE

DE is power.

upvoted 1 times

✉️ **rewanotora** 1 week, 1 day ago

Selected Answer: DE

D and E are right.

upvoted 1 times

✉️ **Waffel00** 4 weeks, 1 day ago

DE is correct because App not created before. Application PLAN was created

upvoted 1 times

✉️ **emiliocb4** 1 month ago

Selected Answer: CD

C and D, since E is already created as stated on the question.

upvoted 1 times

✉️ **CDL_Learner** 1 month, 2 weeks ago

Selected Answer: DE

Correct Answer Should be D & E

When we provision a QnA Maker service, it automatically creates several Azure resources¹. Specifically, it creates an Azure Cognitive Search resource and an Azure App Service

upvoted 1 times

✉️ **CDL_Learner** 1 month, 2 weeks ago

Language Understanding: This is not automatically created when provisioning a QnA Maker service

Azure SQL Database: While Azure SQL Database is a common Azure resource, it's not automatically created when you provision a QnA Maker service. The QnA Maker service doesn't require a SQL database to function

Azure Storage: Azure Storage is a service that provides scalable and secure storage for a variety of data objects, but it's not automatically created when you provision a QnA Maker service. The QnA Maker service doesn't require Azure Storage to function.

upvoted 1 times

✉️ **NullVoider_0** 1 month, 2 weeks ago

Selected Answer: CD

Final answer:

- C. Azure Storage
- D. Azure Cognitive Search

Explanation:

When you provision a QnA Maker service, the following two Azure resources are automatically created in the RG1 resource group:

C. Azure Storage: This storage account is used to store the knowledge base content and other metadata.

D. Azure Cognitive Search: This resource provides the search functionality for the QnA Maker service.

The other options are not automatically created:

A. Language Understanding (option A) is a separate service and is not part of the QnA Maker provisioning.

B. Azure SQL Database (option B) is not automatically created.

E. Azure App Service (option E) is not automatically created, as you have already created the App Service plan AP1 separately.

upvoted 1 times

✉️ **michaelmorar** 1 month, 2 weeks ago

I would hope this is not in the exam - I cannot even create a resource:

"QnA Maker service is being retired on 31st March, 2025. A newer version of this capability is now available as a part of Azure AI service for Language called question answering. To use this service, you need to provision a Language resource. For question answering capability within the Language service, see question answering and its pricing page. You can't create new QnA Maker resources anymore. For information on migrating your existing QnA Maker knowledge bases to question answering, consult the migration guide."

upvoted 1 times

✉️ **NullVoider_0** 1 month, 3 weeks ago

Selected Answer: CD

C. Azure Storage - Correct. When you create a QnA Maker service, an Azure Storage account is automatically created. This storage account is used for storing backups of the QnA Maker knowledge bases and other necessary data.

D. Azure Cognitive Search - Correct. Provisioning a QnA Maker service automatically creates an Azure Cognitive Search service. QnA Maker relies on Azure Cognitive Search for its knowledge base storage, indexing, and query capabilities.

upvoted 1 times

✉️ **rdemontis** 6 months, 4 weeks ago

Selected Answer: DE

answer is correct. When you deploy a new QnA Maker service an Azure Search is required to index your data and an App Service as compute engine that runs the QnA Maker queries for you

upvoted 1 times

✉️ **jolly90** 8 months, 1 week ago

Deprecated feature

upvoted 1 times

✉️ **ross_borissov** 9 months ago

This has to be outdated because QnA will be discontinued and can not be created at the moment. It becomes part of the Language resource that creates nothing additionally than a Search Service.

upvoted 2 times

✉️ **zellck** 11 months ago

Selected Answer: DE

DE is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/qnamaker/concepts/plan#azure-resources>

Each Azure resource created with QnA Maker has a specific purpose. Each resource has its own purpose, limits, and pricing tier. It's important to understand the function of these resources so that you can use that knowledge into your planning process.

- QnA Maker resource

Authoring and query prediction

- Cognitive Search resource

Data storage and search

- App Service resource and App Plan Service resource

Query prediction endpoint

- Application Insights resource

Query prediction telemetry

upvoted 5 times

✉️ **endeesa** 11 months, 2 weeks ago

Selected Answer: CD

Why is the answer not Azure storage and Cognitive search? I think the given answer is wrong
upvoted 1 times

✉️ **AnonymousJhb** 8 months, 1 week ago

common dude! make some effort and go deploy to verify.
upvoted 1 times

✉️ **EliteAllen** 1 year ago

Selected Answer: DE

When you provision the QnA Maker service, the following Azure resources are automatically created:

D. Azure Cognitive Search - Correct. Azure Cognitive Search is used to index the question and answer pairs and provide the underlying search capabilities for the QnA Maker service.

E. Azure App Service - Correct. QnA Maker creates an Azure App Service that hosts the QnA Maker runtime.
upvoted 3 times

✉️ **Pfffff** 1 year, 1 month ago

Selected Answer: BE

chatgpt says so
upvoted 1 times

Question #21

You are building a language model by using a Language Understanding (classic) service.

You create a new Language Understanding (classic) resource.

You need to add more contributors.

What should you use?

- A. a conditional access policy in Azure Active Directory (Azure AD)
- B. the Access control (IAM) page for the authoring resources in the Azure portal
- C. the Access control (IAM) page for the prediction resources in the Azure portal

Correct Answer: B

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-how-to-collaborate>

Community vote distribution

B (100%)

✉️  **zellck**  11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/luis/luis-how-to-collaborate#add-contributor-to-azure-authoring-resource>
In the Azure portal, find your Language Understanding (LUIS) authoring resource. It has the type LUIS.Authoring. In the resource's Access Control (IAM) page, add the role of contributor for the user that you want to contribute.

upvoted 6 times

✉️  **rdemontis** 6 months, 4 weeks ago

thanks for the provided documentation

upvoted 1 times

✉️  **nanaw770**  3 days, 18 hours ago

Selected Answer: B

I selected B, but it does not make sense. It is like there is nothing sadder than having a breakup with a lover.

upvoted 1 times

✉️  **reiwanotora** 1 week, 1 day ago

Selected Answer: B

B is right.

upvoted 1 times

✉️  **CDL_Learner** 1 month, 2 weeks ago

Selected Answer: B

The best answer is B. the Access control (IAM) page for the authoring resources in the Azure portal., to add more contributors to a Language Understanding (classic) resource, we should use the Access control (IAM) page for the authoring resources in the Azure portal. This allows us to manage access and permissions for the resource at a granular level, including adding new contributors.

upvoted 1 times

✉️  **CDL_Learner** 1 month, 2 weeks ago

Here's why the other options are not suitable:

- Option A: A conditional access policy in Azure Active Directory (Azure AD) is used to enforce certain conditions before granting access to resources. It doesn't directly allow us to add contributors to a specific resource.
- Option C: The Access control (IAM) page for the prediction resources in the Azure portal is used to manage access to prediction resources, not authoring resources. Since we're looking to add contributors to the Language Understanding resource we've created (which is an authoring resource), this option wouldn't be suitable.

upvoted 1 times

✉️  **evangelist** 3 months, 2 weeks ago

Selected Answer: B

In Azure's Language Understanding (LUIS) service, adding more contributors or managing permissions is usually done through the Access Control (IAM) page in the Azure portal. This is because, for the management and development of LUIS services, especially in the process of designing and building language models, the required permission management mainly focuses on authoring resources. Through the "Access Control (IAM)" page you can assign different roles to different users or user groups, such as "Contributor" or "Reader", to control their access and operation permissions to resources.

upvoted 2 times

✉️  **rdemontis** 6 months, 4 weeks ago

Selected Answer: B

correct

upvoted 4 times

  **ExamDev** 8 months, 2 weeks ago**Selected Answer: B**

The correct answer is "B"

- <https://learn.microsoft.com/en-us/azure/role-based-access-control/role-assignments-portal>

upvoted 2 times

  **RAN_L** 1 year, 2 months ago**Selected Answer: B**

B. the Access control (IAM) page for the authoring resources in the Azure portal is the correct option to add more contributors to a Language Understanding (classic) resource. IAM (Identity and Access Management) in Azure allows you to manage access to your Azure resources by assigning roles to users, groups, and applications. By using the Access control (IAM) page for the authoring resources in the Azure portal, you can add new users or groups to your Language Understanding (classic) resource and grant them the necessary permissions to edit and publish your language model.

upvoted 2 times

  **KingChuang** 1 year, 4 months ago

on my exam (2023-01-16 Passed)

My Answer:B

upvoted 4 times

  **ArchMelody** 1 year, 6 months ago**Selected Answer: B**

Correct Answer, indeed!

upvoted 3 times

Question #22

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure Cognitive Search service.

During the past 12 months, query volume steadily increased.

You discover that some search query requests to the Cognitive Search service are being throttled.

You need to reduce the likelihood that search query requests are throttled.

Solution: You migrate to a Cognitive Search service that uses a higher tier.

Does this meet the goal?

A. Yes

B. No

Correct Answer: A

A simple fix to most throttling issues is to throw more resources at the search service (typically replicas for query-based throttling, or partitions for indexing-based throttling). However, increasing replicas or partitions adds cost, which is why it is important to know the reason why throttling is occurring at all.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-performance-analysis>

Community vote distribution

A (81%)

B (19%)

✉️  **claps92** Highly Voted  1 year, 4 months ago

Using a higher tier does not imply scaling replicas or partitions.

The answer should be no

upvoted 7 times

✉️  **eham757** Highly Voted  2 years, 1 month ago

Selected Answer: A

answer is correct

upvoted 5 times

✉️  **takaimomoGcup** Most Recent  20 hours, 15 minutes ago

Selected Answer: B

Right answer is No.

upvoted 1 times

✉️  **nanaw770** 3 days, 18 hours ago

Selected Answer: B

This is correct to use a replica, so the answer is no. All but M7, T20, and T25 graduates must give a silly answer of yes.

upvoted 1 times

✉️  **MDWPartners** 2 weeks, 2 days ago

Selected Answer: B

I would say NO, the reason is that a higher tier does not always increase the resources, it depends from which tier you are increasing. However, this question later comes with the option of increase the replicas:

A simple fix to most throttling issues is to throw more resources at the search service (typically replicas for query-based throttling, or partitions for indexing-based throttling). However, increasing replicas or partitions adds cost, which is why it's important to know the reason why throttling is occurring at all. Investigating the conditions that cause throttling will be explained in the next several sections.

upvoted 1 times

✉️  **CDL_Learner** 1 month, 2 weeks ago

Selected Answer: A

The best answer is:

- A. Yes

Reason for choosing this option: Azure Cognitive Search service tiers determine the resources available for that service. Higher tiers provide more resources such as increased document counts, indexing speed, and query capacity. If the query volume has increased and requests are being throttled, it indicates that the current tier's resources are insufficient to handle the load. Migrating to a higher tier will provide more resources, reducing the likelihood of throttling.

upvoted 1 times

✉️  **CDL_Learner** 1 month, 2 weeks ago

Why the other option is not suitable: • B. No

This option is not suitable because it suggests that migrating to a higher tier wouldn't solve the problem. However, in Azure Cognitive Search, moving to a higher tier can indeed provide more resources and reduce the likelihood of request throttling. Therefore, this option is incorrect.

upvoted 1 times

 **evangelist** 3 months, 2 weeks ago

throttle is a simple performance tier issue, using higher tier will resolve the issue

upvoted 1 times

 **rdemontis** 6 months, 4 weeks ago

Selected Answer: A

to me the answer is yes because a higher tier increases the resources provided and then it could also increase the performance.

upvoted 1 times

 **james2033** 9 months, 2 weeks ago

Selected Answer: A

Choose higher tier, see <https://learn.microsoft.com/en-us/azure/search/search-sku-tier>

upvoted 4 times

 **RAN_L** 1 year, 2 months ago

Selected Answer: A

A. Yes.

Migrating to a higher tier in Azure Cognitive Search can provide more resources, such as increased storage, throughput, and replicas, which can help reduce the likelihood of search query requests being throttled.

upvoted 2 times

Question #23

DRAG DROP -

You need to develop an automated call handling system that can respond to callers in their own language. The system will support only French and English.

Which Azure Cognitive Services service should you use to meet each requirement? To answer, drag the appropriate services to the correct requirements. Each service may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Services**Answer Area**

Speaker Recognition

Speech to Text

Text Analytics

Text to Speech

Translator

Detect the incoming language:

Respond in the callers' own language:

Services**Answer Area**

Speaker Recognition

Speech to Text

Text Analytics

Text to Speech

Translator

Correct Answer: Detect the incoming language:

Text Analytics

Respond in the callers' own language:

Translator

Box 1: Text Analytics -

The Language Detection feature of the Azure Text Analytics REST API evaluates text input for each document and returns language identifiers with a score that indicates the strength of the analysis.

Incorrect Answers:

Speaker Recognition which accurately verifies and identifies speakers by their unique voice characteristics.

Box 2: Translator -

Translator is a cloud-based neural machine translation service that is part of the Azure Cognitive Services family of REST APIs. Translator can be used with any operating system and powers many Microsoft products and services used by thousands of businesses worldwide to perform language translation and other language-related operations.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/how-tos/text-analytics-how-to-language-detection>

<https://docs.microsoft.com/en-us/azure/cognitive-services/translator/translator-overview>

 **Khiem**  2 years ago

It should be:

- Speech to Text with AutoDetectSourceLanguageConfig. It can't be Text Analytics because the input is callers' voice.
- Text to Speech: the output is voice.

upvoted 63 times

 **PeteColag** 2 days, 22 hours ago

I agree with your answer. What is missing in this question is an explanation of how the response text (to feed the text to speech) is being generated. I am assuming this would be based on LUIS or CLU?

upvoted 1 times

✉️ **rdemontis** 6 months, 4 weeks ago

I agree with you
upvoted 1 times

✉️ **Eltooth** Highly Voted 1 year, 10 months ago

Speech-to-Text : <https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/speech-to-text>

Text-to-Speech : <https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/text-to-speech>

Both support common languages, including French.

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/language-support?tabs=speechtotext>

upvoted 13 times

✉️ **nanaw770** Most Recent 3 days, 18 hours ago

1. Speech to Text
2. Text to Speech

From Takedajuku perspective, if you study for 4 days and spend 2 days reviewing, you will have a better chance of passing the exam.

upvoted 1 times

✉️ **varinder82** 2 months ago

Final Answer:

1. Speech to Text
 2. Text to Speech
- upvoted 2 times

✉️ **dacchione** 2 months, 2 weeks ago

Language identification (LID) use cases include:

Speech to text recognition when you need to identify the language in an audio source and then transcribe it to text.

Speech translation when you need to identify the language in an audio source and then translate it to another language.

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/language-identification?tabs=once&pivots=programming-language-python>

upvoted 3 times

✉️ **Mehe323** 2 months, 4 weeks ago

People, read the descriptions, these are describing only parts of the process, not the whole process.

1) DETECT incoming language has got nothing to do with speech to text, but with text analytics. Yes, speech to text is a part of the whole process but this service TRANSCRIBES and does not do language detection.

2) Respond to caller's language should be Text to speech

upvoted 1 times

✉️ **Mehe323** 1 month, 4 weeks ago

IGNORE my comment! Speech to text also does language detection:

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/language-identification?tabs=once&pivots=programming-language-csharp>

upvoted 1 times

✉️ **audlindr** 3 months, 1 week ago

It can't be Text Analytics. It is an incoming call. It should be Speech to Text

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/language-identification?tabs=once&pivots=programming-language-csharp>

And then Text to Speech for the response

upvoted 1 times

✉️ **suzanne_exam** 4 months ago

Speech to text - as it's based on a voice call

Text to speech - it's not translator here, as the key thing is that the program needs to respond, not just translate

upvoted 1 times

✉️ **james2033** 9 months, 2 weeks ago

At this time, this question is out of date. See <https://learn.microsoft.com/en-us/azure/ai-services/language-service/language-detection/overview>.

In old time, Text Analytics is correct.

upvoted 3 times

✉️ **zellck** 11 months ago

1. Speech to Text
2. Text to Speech

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/language-identification?tabs=once&pivots=programming-language-csharp#speech-to-text>

You use Speech to text recognition when you need to identify the language in an audio source and then transcribe it to text.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/text-to-speech>

Text to speech enables your applications, tools, or devices to convert text into humanlike synthesized speech. The text to speech capability is also known as speech synthesis. Use humanlike prebuilt neural voices out of the box, or create a custom neural voice that's unique to your product or brand.

upvoted 6 times

✉ **endeesa** 11 months, 2 weeks ago

to detect the incoming language you need Text analytics, Speech to text does not have an option to detect language
upvoted 1 times

✉ **CauchyLee** 11 months, 1 week ago

POST /cognitiveservices/v1/speechtotext/recognition HTTP/1.1
Ocp-Apim-Subscription-Key: <subscription_key>
Content-Type: application/json

```
{  
  "config": {  
    "language": "auto",  
    "enableSeparation": true  
  },  
  "format": "audio-16khz-128kbitrate-mono-mp3",  
  "audio": <binary_audio_file>  
}
```

upvoted 1 times

✉ **ziggy1117** 11 months, 3 weeks ago

A. Speech to text: autodetectsourcelanguageconfig : <https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/language-identification?tabs=once&pivots=programming-language-csharp>

b. text to speech

upvoted 1 times

✉ **Mike19D** 1 year, 1 month ago

Speech to Text

Translator

upvoted 1 times

✉ **Anichebe** 1 year, 1 month ago

From ChatGPT:

To develop an automated call handling system that can respond to callers in their own language, you can use the following Azure Cognitive Services:

Speech-to-Text: This service can transcribe speech into text, which can then be analyzed to determine the language being spoken.

Text Translation: This service can translate text from one language to another.

Using these services together, you can create a system that can transcribe a caller's speech, detect their language, and then translate the transcript into the appropriate language for a response.

upvoted 2 times

✉ **AdarshKumarKhare** 1 year, 6 months ago

But caller is calling not writing text. So, audio should recognize the language not Text Analysis

upvoted 2 times

✉ **mathonno** 1 year, 7 months ago

<https://azure.microsoft.com/fr-fr/products/cognitive-services/translator/#overview>

translator and translator

upvoted 1 times

✉ **Ajose0** 1 year, 10 months ago

Translator is an AI service for real-time document and text translation

<https://azure.microsoft.com/en-us/services/cognitive-services/translator/#overview>

upvoted 1 times

Question #24

You have receipts that are accessible from a URL.

You need to extract data from the receipts by using Form Recognizer and the SDK. The solution must use a prebuilt model.

Which client and method should you use?

- A. the FormRecognizerClient client and the StartRecognizeContentFromUri method
- B. the FormTrainingClient client and the StartRecognizeContentFromUri method
- C. the FormRecognizerClient client and the StartRecognizeReceiptsFromUri method
- D. the FormTrainingClient client and the StartRecognizeReceiptsFromUri method

Correct Answer: D

To analyze receipts from a URL, use the StartRecognizeReceiptsFromUri method

Example code:

```
private static async Task AnalyzeReceipt(  
    FormRecognizerClient recognizerClient, string receiptUri)  
{  
    RecognizedFormCollection receipts = await recognizerClient.StartRecognizeReceiptsFromUri(new Uri(receiptUrl)).WaitForCompletionAsync();
```

Reference:

<https://docs.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/quickstarts/client-library>

Community vote distribution

C (100%)

✉ **ez4Zane** Highly Voted 2 years, 1 month ago

Should be C

```
private static async Task AnalyzeReceipt(  
    FormRecognizerClient recognizerClient, string receiptUri)  
{  
    RecognizedFormCollection receipts = await recognizerClient.StartRecognizeReceiptsFromUri(new Uri(receiptUrl)).WaitForCompletionAsync();
```

upvoted 28 times

✉ **Prodyna** Highly Voted 6 months, 2 weeks ago

was on november exam but it said "using Document Intelligence", the answer possibilities were the same

upvoted 10 times

✉ **PeteColag** Most Recent 2 days, 18 hours ago

Should be C. We are not doing training here, we are doing inference.

upvoted 1 times

✉ **nanaw770** 3 days, 18 hours ago

Selected Answer: C

C is right answer, but use word: Document Intelligence.

upvoted 1 times

✉ **AzureGC** 1 month ago

Selected Answer: C

C: Based on the comments in the example

upvoted 1 times

✉ **CDL_Learner** 1 month, 2 weeks ago

Selected Answer: C

The best answer is C. the FormRecognizerClient client and the StartRecognizeReceiptsFromUri method.

Reason for choosing option C: The FormRecognizerClient is the client used to interact with the service in the Azure.AI.FormRecognizer namespace. The StartRecognizeReceiptsFromUri method is specifically designed to recognize and extract data from receipts, which is exactly what the question asks for. This method uses a prebuilt model trained on receipts, making it the ideal choice for this scenario.

upvoted 2 times

✉ **CDL_Learner** 1 month, 2 weeks ago

Why other options are not suitable:

Option A: The StartRecognizeContentFromUri method is used to extract layout information such as tables, lines, words, and selection marks. It's not specifically designed for receipts.

Option B: The FormTrainingClient is used to train custom models, not to extract data from documents using prebuilt models. Also, the

StartRecognizeContentFromUri method, as mentioned above, is not specifically designed for receipts.

Option D: Similar to option B, the FormTrainingClient is not suitable for this scenario as it's used for training custom models. The StartRecognizeReceiptsFromUri method would be correct if used with FormRecognizerClient.

upvoted 2 times

✉ **audlindr** 3 months, 1 week ago

Selected Answer: C

Seems like a typo on the answer. the code in the explanation clearly shows that it is FormRecognizerclient

upvoted 2 times

✉ **evangelist** 3 months, 2 weeks ago

Selected Answer: C

C is Correct

upvoted 2 times

✉ **rdemontis** 6 months, 4 weeks ago

Selected Answer: C

Wrong answer but correct explanation. Anyway correct answer is obviously C.

upvoted 3 times

✉ **zellck** 11 months ago

Selected Answer: C

C is the answer.

<https://learn.microsoft.com/en-us/dotnet/api/azure.ai.formrecognizer.formrecognizerclient?view=azure-dotnet>

The client to use to connect to the Form Recognizer Azure Cognitive Service to recognize information from forms and images and extract it into structured data. It provides the ability to analyze receipts, business cards, and invoices, to recognize form content, and to extract fields from custom forms with models trained on custom form types.

upvoted 4 times

✉ **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 3 times

✉ **zellck** 11 months ago

<https://learn.microsoft.com/en-us/dotnet/api/azure.ai.formrecognizer.formrecognizerclient.startrecognizereceiptsfromuri?view=azure-dotnet>

Recognizes values from one or more receipts.

upvoted 2 times

✉ **rdemontis** 6 months, 4 weeks ago

thanks for the provided link

upvoted 1 times

✉ **nitz14** 11 months, 2 weeks ago

Selected Answer: C

To extract data from receipts using Form Recognizer and the SDK, while using a prebuilt model, you should use option C: the FormRecognizerClient client and the StartRecognizeReceiptsFromUri method.

Explanation:

Option A (the FormRecognizerClient client and the StartRecognizeContentFromUri method) is incorrect because the StartRecognizeContentFromUri method is used for general content recognition, not specifically for receipts.

Option B (the FormTrainingClient client and the StartRecognizeContentFromUri method) is incorrect because the FormTrainingClient client is used for training custom models, not for extracting data from prebuilt models.

Option D (the FormTrainingClient client and the StartRecognizeReceiptsFromUri method) is incorrect because the FormTrainingClient client is not used for extracting data; it is used for training custom models.

Therefore, the correct choice is option C: the FormRecognizerClient client and the StartRecognizeReceiptsFromUri method.

upvoted 4 times

✉ **Anichebe** 1 year, 1 month ago

The correct answer is option C

upvoted 1 times

✉ **RAN_L** 1 year, 2 months ago

Selected Answer: C

The StartRecognizeReceiptsFromUri method of the FormRecognizerClient client is used to extract data from receipts using a prebuilt model.

upvoted 2 times

✉ **ap1234pa** 1 year, 4 months ago

Selected Answer: C

C is correct

upvoted 1 times

✉️ **ap1234pa** 1 year, 4 months ago

Selected Answer: C

C is the answer

upvoted 1 times

✉️ **Pyguy** 1 year, 5 months ago

Selected Answer: C

prebuilt model = Nothing to do with "training" word

upvoted 3 times

✉️ **ArchMelody** 1 year, 6 months ago

Selected Answer: C

Obviously C, as you need to use a pre-built model, not train a new one...

upvoted 5 times

Question #25

You have a collection of 50,000 scanned documents that contain text.

You plan to make the text available through Azure Cognitive Search.

You need to configure an enrichment pipeline to perform optical character recognition (OCR) and text analytics. The solution must minimize costs.

What should you attach to the skillset?

- A. a new Computer Vision resource
- B. a free (Limited enrichments) Cognitive Services resource
- C. an Azure Machine Learning Designer pipeline
- D. a new Cognitive Services resource that uses the S0 pricing tier

Correct Answer: A

The Computer Vision API uses text recognition APIs to extract and recognize text information from images. Read uses the latest recognition models, and is optimized for large, text-heavy documents and noisy images.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/solution-ideas/articles/cognitive-search-with-skillsets>

Community vote distribution

D (91%) 6%

✉ Pyguy **Highly Voted** 1 year, 5 months ago

Selected Answer: D

Question says : ... "You need to configure an enrichment pipeline to perform optical character recognition (OCR) and "text analytics"... Just because of this second requirement Answer is a Multi Cognitive Service (D), alone vision service (=OCR wont make any text analytics, text extraction is not text analytics !).. Read here : <https://learn.microsoft.com/en-us/azure/search/cognitive-search-attach-cognitive-services?tabs=portal>

upvoted 23 times

✉ claps92 1 year, 4 months ago

i agree
upvoted 1 times

✉ rdemontis 6 months, 4 weeks ago

I totally agree with you.
upvoted 3 times

✉ evangelist **Highly Voted** 3 months, 2 weeks ago

Selected Answer: D

For processing 50,000 documents with OCR and text analytics in an enrichment pipeline, while still aiming to minimize costs, a more practical approach would be:

D. a new Cognitive Services resource that uses the S0 pricing tier

This option is more suitable for several reasons:

Capacity and Scalability: The S0 pricing tier offers a higher transaction limit, making it more suitable for processing a large volume of documents like 50,000. It can handle the scale of operations required for this task without hitting the cap as quickly as the free tier might.

Cost Efficiency at Scale: While the S0 tier incurs costs, it is designed to be cost-effective for medium to large-scale operations. By carefully managing the use of resources and optimizing the processing pipeline, it's possible to keep costs under control while meeting the project's requirements.

upvoted 5 times

✉ PeteColag **Most Recent** 1 day, 21 hours ago

Selected Answer: D

A Multiservice account is required since both OCR and text analytics are required.

Since a multiservice account does not support the free tier (F0), then B is not a valid option. So, the answer has to be D.

upvoted 1 times

✉ nanaw770 3 days, 18 hours ago

Selected Answer: D

Use S0 SKU
upvoted 1 times

✉ reiwanotora 1 week, 1 day ago

Selected Answer: D

D is right.

upvoted 1 times

✉ Alok1105 3 weeks, 2 days ago

Selected Answer: B

Given the goal of minimizing costs while performing OCR and text analytics on a collection of 50,000 scanned documents, and considering that the Limited enrichments tier offers basic enrichments at no additional cost, the most suitable choice is indeed:

B. a free (Limited enrichments) Cognitive Services resource.

upvoted 2 times

✉ brajesh8684 1 month ago

B. a free (Limited enrichments) Cognitive Services resource

Explanation:

The free (Limited enrichments) tier of Cognitive Services provides access to Optical Character Recognition (OCR) and basic text analytics capabilities.

Attaching this Cognitive Services resource to the skillset allows you to perform OCR and text analytics on the scanned documents.

Since the requirement is to minimize costs, choosing the free tier is the most cost-effective option.

This option aligns with the goal of making the text available through Azure Cognitive Search while keeping costs low.

upvoted 2 times

✉ CDL_Learner 1 month, 2 weeks ago

Selected Answer: D

Correct option is D - a new Cognitive Services resource that uses the S0 pricing tier

We can't do this with free tier, thus we can't select Option B

upvoted 2 times

✉ 3606c4a 2 months, 3 weeks ago

to minimize cost you would not use a S0 tier which is why the answer is A

upvoted 1 times

✉ Mehe323 2 months, 1 week ago

AI Vision OCR is only for very simple OCR tasks and not suitable for more complicated solutions. I don't have the exact source of below note I made about this, I think it was from Microsoft Learn:

If you want to extract simple words and text from a picture of a form or document, without contextual information, Azure AI Vision OCR is an appropriate service to consider. You might want to use this service if you already have your own analysis code, for example. However, Azure AI Document Intelligence includes a more sophisticated analysis of documents. For example, it can identify key/value pairs, tables, and context-specific fields. If you want to deploy a complete document analysis solution that enables users to both extract and understand text, consider Azure AI Document Intelligence.

upvoted 1 times

✉ sca88 6 months, 1 week ago

Selected Answer: D

OCR is a built-in skillset that required billable tier of Azure AI Search, so the D is correct

upvoted 1 times

✉ rdemontis 6 months, 4 weeks ago

Selected Answer: D

correct answer is D because you need to support a multiservice resource and it cannot be free.

upvoted 1 times

✉ zellck 11 months ago

Selected Answer: D

D is the answer.

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-attach-cognitive-services?tabs=portal>

When configuring an optional AI enrichment pipeline in Azure Cognitive Search, you can enrich a limited number of documents free of charge. For larger and more frequent workloads, you should attach a billable multi-service Cognitive Services resource.

A multi-service resource references "Cognitive Services" as the offering, rather than individual services, with access granted through a single API key. This key is specified in a skillset and allows Microsoft to charge you for using these APIs:

- Computer Vision for image analysis and optical character recognition (OCR)
- Language service for language detection, entity recognition, sentiment analysis, and key phrase extraction
- Translator for machine text translation

upvoted 3 times

✉ dd11121996 10 months, 1 week ago

what about option B ?

upvoted 1 times

✉ Mehe323 2 months, 1 week ago

Option B is not possible, the free tier is up to 10,000 documents.

upvoted 1 times

EliteAllen 11 months, 2 weeks ago

Selected Answer: D

Based on the information from the official Azure documentation, the most cost-effective way to configure an enrichment pipeline for OCR and text analytics in Azure Cognitive Search would be to use a billable Cognitive Services resource. This is because built-in skills, such as OCR and text analytics, are based on the Cognitive Services APIs and unless your content input is small, you would need to attach a billable Cognitive Services resource to run larger workloads.

Therefore, the correct answer to your question would be:

D. a new Cognitive Services resource that uses the S0 pricing tier

This option allows you to use the necessary built-in skills for OCR and text analytics and can handle larger workloads, which would be necessary for processing a large number of documents. The S0 pricing tier is a standard tier that provides a balance between cost and capabilities.

upvoted 4 times

Rob77 1 year ago

D - free tier is limited to 10,000 documents

<https://learn.microsoft.com/en-us/azure/search/search-limits-quotas-capacity#indexer-limits>

upvoted 2 times

dev2dev 1 year, 4 months ago

D is not possible because OCR requires s2 pricing tier and the options sys to use s0 so answer is A

upvoted 2 times

marti_tremblay000 1 year, 2 months ago

What about the text analytics part ? Computer Vision does not perform Text Analytics... You need a Language service to do so and therefore you must provision Cognitive Services as a whole to perform the 2 tasks.

upvoted 1 times

ap1234pa 1 year, 4 months ago

Selected Answer: D

D is valid

upvoted 2 times

ap1234pa 1 year, 4 months ago

Selected Answer: D

We need to consider both the services..OCR and Text Analytics. So answer is D

upvoted 3 times

Question #26

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure Cognitive Search service.

During the past 12 months, query volume steadily increased.

You discover that some search query requests to the Cognitive Search service are being throttled.

You need to reduce the likelihood that search query requests are throttled.

Solution: You add indexes.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

Instead, you could migrate to a Cognitive Search service that uses a higher tier.

Note: A simple fix to most throttling issues is to throw more resources at the search service (typically replicas for query-based throttling, or partitions for indexing-based throttling). However, increasing replicas or partitions adds cost, which is why it is important to know the reason why throttling is occurring at all.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-performance-analysis>

Community vote distribution

B (84%)

A (16%)

Eltooth Highly Voted 1 year, 10 months ago

Selected Answer: B

B is correct answer.

"How your search queries perform is directly connected to the size and complexity of your indexes. The smaller and more optimized your indexes, the fast Azure Cognitive Search can respond to queries.

If your index has been optimized but the performance still isn't where it needs to be, you can choose to scale up or scale out your search service.

If you've applied all of the above and still have individual queries that don't perform, you can scale out your index. Depending on the service tier you used to create your search solution, you can add up to 12 partitions. Partitions are the physical storage where your index resides. By default, a new search indexes are created with a single partition. If you add more partitions, the index is stored across them. For example, if your index is 200 GB and you've four partitions, each partition contains 50 GB of your index.

Adding extra partitions can help with performance as the search engine can run in parallel in each partition.

upvoted 9 times

Eltooth 1 year, 10 months ago

You can scale out with replicas if you need to scale because of an increase in load. You can also scale up your search service by using a higher tier.

<https://docs.microsoft.com/en-gb/learn/modules/maintain-azure-cognitive-search-solution/03-optimize-performance-of-azure-cognitive-search-solution>

upvoted 1 times

brajesh8684 Most Recent 1 month ago

Answer: B

The solution provided does not directly address the issue of search query requests being throttled. Adding indexes can improve search performance by optimizing the search process, but it may not directly address throttling issues.

upvoted 1 times

CDL_Learner 1 month, 2 weeks ago

Selected Answer: B

The best answer is B. No.

Reason for choosing this option: Adding indexes to an Azure Cognitive Search service does not directly reduce the likelihood of search query requests being throttled. Throttling occurs when the service is handling more requests than it is provisioned to handle. This is typically managed by scaling up the service to handle more requests, not by adding more indexes.

upvoted 1 times

✉️ **CDL_Learner** 1 month, 2 weeks ago

Why other option is not suitable: Option A is not suitable because adding indexes does not increase the capacity of the service to handle more requests. Indexes in Azure Cognitive Search are used to organize and retrieve data efficiently, but they do not increase the service's capacity to handle a larger volume of search query requests. Therefore, adding indexes would not reduce the likelihood of requests being throttled. Instead to handle increased query volume, we would need to scale up the service or optimize our queries for efficiency.

Effective solution is to Move toward Higher Tier

upvoted 1 times

✉️ **rdemontis** 6 months, 4 weeks ago

Selected Answer: B

"A simple fix to most throttling issues is to throw more resources at the search service (typically replicas for query-based throttling, or partitions for indexing-based throttling). However, increasing replicas or partitions adds cost, which is why it's important to know the reason why throttling is occurring at all!"

<https://learn.microsoft.com/en-us/azure/search/search-performance-analysis#throttling-behaviors>

upvoted 1 times

✉️ **sl_mslconsulting** 7 months ago

Selected Answer: A

Normally you won't scale out and scale up immediately if there are just some queries have performance issue. You will try to optimized your indexes or queries first. if you need to make a change to your index definition, such as editing a field or adding a new analyzer, you'll have to create a new search index because all search indexes are immutable. I would try to do this first and hence I pick A. <https://learn.microsoft.com/en-us/azure/search/search-how-to-alias?tabs=rest>

upvoted 1 times

✉️ **zellck** 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/search/search-performance-analysis#throttling-behaviors>

Throttling occurs when the search service is at capacity. Throttling can occur during queries or indexing. From the client side, an API call results in 503 HTTP response when it has been throttled. During indexing, there's also the possibility of receiving a 207 HTTP response, which indicates that one or more items failed to index. This error is an indicator that the search service is getting close to capacity.

A simple fix to most throttling issues is to throw more resources at the search service (typically replicas for query-based throttling, or partitions for indexing-based throttling). However, increasing replicas or partitions adds cost, which is why it's important to know the reason why throttling is occurring at all. Investigating the conditions that cause throttling will be explained in the next several sections.

upvoted 3 times

✉️ **rdemontis** 6 months, 4 weeks ago

thanks for explanation and provided link

upvoted 1 times

✉️ **RAN_L** 1 year, 2 months ago

Selected Answer: B

Adding indexes alone may not be sufficient to reduce the likelihood of search query requests being throttled. Increasing the number of replicas and scaling up to a higher tier can also help reduce throttling. Additionally, optimizing the search queries and reducing the number of requests can also help alleviate the issue. Therefore, without additional information, it is not possible to determine if adding indexes alone would meet the goal.

upvoted 2 times

✉️ **vmisirlis** 1 year, 3 months ago

The answer is B. Adding indexes will not reduce the likelihood that search query requests are being throttled.

Throttling occurs when the search service is at capacity and cannot handle the volume of incoming requests. Adding indexes will not necessarily increase the capacity of the search service.

upvoted 1 times

✉️ **STH** 1 year, 8 months ago

Selected Answer: A

"How your search queries perform is directly connected to the size and complexity of your indexes."

Question does not tell anything about the current state or indexes. There even may be none.

Without any context, yes adding indexes improves query performance

upvoted 2 times

✉️ **Rob77** 1 year ago

Unlikely, In Cognitive Search, you'll work with one index at a time, where all index-related operations target a single index. There is no concept of related indexes or the joining of independent indexes for either indexing or querying.

<https://learn.microsoft.com/en-us/azure/search/search-what-is-an-index#index-isolation>

upvoted 1 times

Question #27

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure Cognitive Search service.

During the past 12 months, query volume steadily increased.

You discover that some search query requests to the Cognitive Search service are being throttled.

You need to reduce the likelihood that search query requests are throttled.

Solution: You enable customer-managed key (CMK) encryption.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

Customer-managed key (CMK) encryption does not affect throttling.

Instead, you could migrate to a Cognitive Search service that uses a higher tier.

Note: A simple fix to most throttling issues is to throw more resources at the search service (typically replicas for query-based throttling, or partitions for indexing-based throttling). However, increasing replicas or partitions adds cost, which is why it is important to know the reason why throttling is occurring at all.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-performance-analysis>

Community vote distribution

 B (100%)

✉ CDL_Learner 1 month, 2 weeks ago

Selected Answer: B

The best answer is: • B. No

Reason for choosing this option: Enabling customer-managed key (CMK) encryption does not directly impact the throttling of search query requests in Azure Cognitive Search service. CMK encryption is used for data-at-rest encryption and does not affect the performance or capacity of the search service.

upvoted 2 times

✉ CDL_Learner 1 month, 2 weeks ago

Higher Tier or Optimizing the query is the way to go !!

upvoted 1 times

✉ rdemontis 6 months, 4 weeks ago

Selected Answer: B

CMK could probably reduce performance, not increase it for sure.

upvoted 1 times

✉ ManvalT 7 months, 3 weeks ago

Got this in Oct2023 exam. Ans here is correct

upvoted 3 times

✉ zellck 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/search/search-performance-analysis#throttling-behaviors>

Throttling occurs when the search service is at capacity. Throttling can occur during queries or indexing. From the client side, an API call results in a 503 HTTP response when it has been throttled. During indexing, there's also the possibility of receiving a 207 HTTP response, which indicates that one or more items failed to index. This error is an indicator that the search service is getting close to capacity.

A simple fix to most throttling issues is to throw more resources at the search service (typically replicas for query-based throttling, or partitions for indexing-based throttling). However, increasing replicas or partitions adds cost, which is why it's important to know the reason why throttling is occurring at all. Investigating the conditions that cause throttling will be explained in the next several sections.

upvoted 4 times

✉ RAN_L 1 year, 2 months ago

Selected Answer: B

No, enabling customer-managed key (CMK) encryption would not directly reduce the likelihood of search query requests being throttled. CMK encryption helps protect data at rest and in transit, but it does not directly impact the performance or throughput of the search service. To reduce the likelihood of search query requests being throttled, you might consider upgrading to a higher tier or making changes to the search indexing and query performance, such as optimizing search queries and filters, reducing query volume, or adding more indexes or replicas.

upvoted 1 times

Eltooth 1 year, 10 months ago

Selected Answer: B

B is correct answer : no.

Obvious.

upvoted 2 times

Question #28

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You create a web app named app1 that runs on an Azure virtual machine named vm1. Vm1 is on an Azure virtual network named vnet1.

You plan to create a new Azure Cognitive Search service named service1.

You need to ensure that app1 can connect directly to service1 without routing traffic over the public internet.

Solution: You deploy service1 and a private endpoint to vnet1.

Does this meet the goal?

A. Yes

B. No

Correct Answer: A

A private endpoint is a network interface that uses a private IP address from your virtual network. This network interface connects you privately and securely to a service powered by Azure Private Link. By enabling a private endpoint, you're bringing the service into your virtual network.

The service could be an Azure service such as:

- Azure Storage
- Azure Cosmos DB
- Azure SQL Database
- Your own service using a Private Link Service.

Reference:

<https://docs.microsoft.com/en-us/azure/private-link/private-endpoint-overview>

Community vote distribution

A (100%)

 CDL_Learner 1 month, 2 weeks ago

Selected Answer: A

The best answer is:

- A. Yes

Reason for choosing this option:

The solution meets the goal because deploying the Azure Cognitive Search service (service1) and a private endpoint to the Azure virtual network (vnet1) allows the web app (app1) on the virtual machine (vm1) to connect directly to service1. This setup ensures that the traffic between app1 and service1 does not route over the public internet, enhancing security and performance.

upvoted 1 times

 rdemontis 6 months, 4 weeks ago

Selected Answer: A

this is the right solution to the scenario

upvoted 2 times

 zellck 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/search/service-create-private-endpoint#why-use-a-private-endpoint-for-secure-access>

Private Endpoints for Azure Cognitive Search allow a client on a virtual network to securely access data in a search index over a Private Link. The private endpoint uses an IP address from the virtual network address space for your search service. Network traffic between the client and the search service traverses over the virtual network and a private link on the Microsoft backbone network, eliminating exposure from the public internet.

upvoted 4 times

 RAN_L 1 year, 2 months ago

Selected Answer: A

By deploying a private endpoint to vnet1, the traffic between app1 and service1 can be routed through the Azure backbone network instead of over the public internet, providing a more secure and reliable connection.

upvoted 2 times

 STH 1 year, 8 months ago

Selected Answer: A

Private Endpoint is the right solution to avoid public internet access

upvoted 3 times

 **not_a_robot** 1 year, 9 months ago

How would the solution work without Azure Private Link?

upvoted 2 times

 **Eltooth** 1 year, 10 months ago

Selected Answer: A

A is correct answer : yes.

upvoted 3 times

Question #29

You have a Language Understanding resource named lu1.
You build and deploy an Azure bot named bot1 that uses lu1.
You need to ensure that bot1 adheres to the Microsoft responsible AI principle of inclusiveness.
How should you extend bot1?

- A. Implement authentication for bot1.
- B. Enable active learning for lu1.
- C. Host lu1 in a container.
- D. Add Direct Line Speech to bot1.

Correct Answer: D

Inclusiveness: AI systems should empower everyone and engage people.

Direct Line Speech is a robust, end-to-end solution for creating a flexible, extensible voice assistant. It is powered by the Bot Framework and its Direct Line

Speech channel, that is optimized for voice-in, voice-out interaction with bots.

Incorrect:

Not B: The Active learning suggestions feature allows you to improve the quality of your knowledge base by suggesting alternative questions, based on user-submissions, to your question and answer pair. You review those suggestions, either adding them to existing questions or rejecting them.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/direct-line-speech>

Community vote distribution

D (94%) 6%

✉  **RamonKaus**  1 year, 10 months ago

Selected Answer: D

I disagree with JT Wang because Direct Line Speech lets the bot speak outloud. This supports Microsoft's goal of AI Inclusiveness because you can include people with disabilities.

upvoted 22 times

✉  **ninja** 1 year, 9 months ago

Agreed the answer should be D, which would people who don't text.

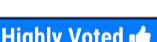
Active learning would improve the overall quality, but not necessarily make it more inclusive.

upvoted 2 times

✉  **rdemontis** 6 months, 4 weeks ago

Totally agree with you

upvoted 1 times

✉  **JTWang**  2 years ago

Answer is B: user can participate in the content of QnA knowledge base

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/how-to/use-active-learning>

The Active learning suggestions feature allows you to improve the quality of your knowledge base by suggesting alternative questions, based on user-submissions, to your question and answer pair. You review those suggestions, either adding them to existing questions or rejecting them.

upvoted 9 times

✉  **nanaw770**  3 days, 18 hours ago

Selected Answer: D

Add Direct Line Speech

upvoted 1 times

✉  **brajesh8684** 1 month ago

B. Enable active learning for lu1.

Explanation:

Active learning is a mechanism that allows the Language Understanding (LU) service to improve over time by actively seeking feedback from users. By enabling active learning for lu1, the bot can continuously learn from interactions with users, improving its understanding and inclusiveness over time.

While options like implementing authentication, hosting lu1 in a container, or adding Direct Line Speech may be relevant for other aspects of the bot's functionality, enabling active learning directly addresses the goal of improving inclusiveness through continuous learning and adaptation.

upvoted 1 times

✉ **trato** 1 month, 1 week ago

D. Add Direct Line Speech to bot1.

A, B, and C might be useful for other purposes, they do not directly contribute to the principle of inclusiveness. Authentication (A) is more about security, active learning (B) is about improving the bot's understanding over time, and hosting in a container. (C) is about deployment flexibility. None of these directly enhance the inclusiveness of the bot.

upvoted 1 times

✉ **evangelist** 3 months, 2 weeks ago

Selected Answer: D

Other options apparently have nothing to do with "inclusiveness"

upvoted 2 times

✉ **Nihilist11** 3 months, 3 weeks ago

It was in today's exam - Feb-04-2023

upvoted 1 times

✉ **rdemontis** 6 months, 4 weeks ago

Selected Answer: D

The correct answer is D because it can help include people who are unable to interact with limbs

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/direct-line-speech>

upvoted 1 times

✉ **zellick** 11 months ago

Selected Answer: D

D is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/direct-line-speech>

Direct Line Speech is a robust, end-to-end solution for creating a flexible, extensible voice assistant. It is powered by the Bot Framework and its Direct Line Speech channel, that is optimized for voice-in, voice-out interaction with bots.

Voice assistants listen to users and take an action in response, often speaking back. They use speech to text to transcribe the user's speech, then take action on the natural language understanding of the text. This action frequently includes spoken output from the assistant generated with text-to-speech.

upvoted 3 times

✉ **zellick** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 3 times

✉ **RAN_L** 1 year, 2 months ago

Selected Answer: D

Enabling Direct Line Speech for bot1 would improve accessibility by providing users with the ability to interact with the bot using natural language and speech. Therefore, D is a valid option to ensure that bot1 adheres to the Microsoft responsible AI principle of inclusiveness.

upvoted 2 times

✉ **Shaka711** 1 year, 6 months ago

Selected Answer: D

Direct line speech enables the bots to speak

upvoted 1 times

✉ **[Removed]** 1 year, 8 months ago

Answer is definitely D, other answers not make sense.

upvoted 1 times

✉ **Eltooth** 1 year, 10 months ago

Selected Answer: B

B is correct answer.

upvoted 1 times

so D is the right one

upvoted 2 times

Question is about how to include larger audience, like blind users or others, not how to improve bot quality

so D is the right one

upvoted 2 times

✉ **sdokmak** 1 year, 11 months ago

Selected Answer: B

from JT Wang's comment

upvoted 1 times

Question #30

HOTSPOT -

You are building an app that will process incoming email and direct messages to either French or English language support teams.

Which Azure Cognitive Services API should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area**https://**

api.cognitive.microsofttranslator.com
eastus.api.cognitive.microsoft.com
portal.azure.com

/text/analytics/v3.1/entities/recognition/general
/text/analytics/v3.1/languages
/translator/text/v3.0/translate?to=en
/translator/text/v3.0/translate?to=fr

Correct Answer:**Answer Area****https://**

api.cognitive.microsofttranslator.com
eastus.api.cognitive.microsoft.com
portal.azure.com

/text/analytics/v3.1/entities/recognition/general
/text/analytics/v3.1/languages
/translator/text/v3.0/translate?to=en
/translator/text/v3.0/translate?to=fr

Box 1: api/cognitive.microsofttranslator.com is used for translations.

Incorrect:

eastus.api.cognitive.microsoft.com is used for Face recognition.

Portal.azure.com is the URL of the Azure portal which is a web-based, unified console that provides an alternative to command-line tools. With the Azure portal, you can manage your Azure subscription using a graphical user interface. You can build, manage, and monitor everything from simple web apps to complex cloud deployments.

Box 2: /text/analytics/v3.1/entities/recognition/general

Named Entity Recognition -

The API returns a list of general named entities in a given document.

Request URL: [https://{endpoint}/text/analytics/v3.1/entities/recognition/general\[?model-version\]\[&showStats\]\[&loggingOptOut\]\[&stringIndexType\]](https://{endpoint}/text/analytics/v3.1/entities/recognition/general[?model-version][&showStats][&loggingOptOut][&stringIndexType])

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-translate>

<https://westcentralus.dev.cognitive.microsoft.com/docs/services/TextAnalytics-v3-1/operations/EntitiesRecognitionGeneral>

  **Adedoyin_Simeon**  1 year, 5 months ago

The answer is wrong.

Correct answer should be:

Box1:

<https://eastus.api.cognitive.microsoft.com>

Box2:

</text/analytics/v3.1/languages>

REF:

<https://learn.microsoft.com/en-us/rest/api/cognitiveservices-textanalytics/3.0/languages/languages?tabs=HTTP>.

NOTE:

Pay special attention to the Sample Request provided. Request to the API should be of the form:

POST {Endpoint}/text/analytics/v3.0/languages

Where the {Endpoint} as stated under the sub-heading "URI Parameters" was described as quoted here (see "Description" column of the table):

"Supported Cognitive Services endpoints (protocol and hostname, for example: <https://westus.api.cognitive.microsoft.com>)."

So the sample given shows the correct format of the *endpoint* as <https://{location}.api.cognitive.microsoft.com>

upvoted 35 times

  **p2006** 1 week, 2 days ago

<https://westus.dev.cognitive.microsoft.com/docs/services/TextAnalytics-v3-1/operations/Languages>

upvoted 1 times

- ✉ **ziggy1117** 11 months, 3 weeks ago
this is correct: <https://learn.microsoft.com/en-us/rest/api/cognitiveservices-textanalytics/3.1preview4/languages/languages?tabs=HTTP>
upvoted 1 times
- ✉ **rdemontis** 6 months, 4 weeks ago
correct! <https://eastus.dev.cognitive.microsoft.com/docs/services/TextAnalytics-v3-0/operations/Languages/console>
upvoted 3 times
- ✉ **mk1967** Highly Voted 1 year, 8 months ago
We only need to detect the language, so the Language Cognitive Service could be used, i.e.:
POST {Endpoint}/text/analytics/v3.0/languages
where Endpoint could be for example "westus"
<https://docs.microsoft.com/en-us/rest/api/cognitiveservices-textanalytics/3.0/languages/languages?tabs=HTTP>
upvoted 11 times
- ✉ **nanaw770** Most Recent 3 days, 18 hours ago
1. eastus.api.cognitive.microsoft.com
2. </text/analytics/v3.1/languages>
upvoted 1 times
- ✉ **nanaw770** 3 days, 18 hours ago
1. eastus.api.cognitive.microsoft.com
2. [text/analytics/v3.1/languages](https://eastus.dev.cognitive.microsoft.com/docs/services/TextAnalytics-v3-0/operations/Languages/console)
upvoted 1 times
- ✉ **funny_penguin** 3 days, 23 hours ago
on exam today 24/05/2024, Adedoyin_Simeon's response is the one I selected as well.
upvoted 1 times
- ✉ **reiwanotora** 1 week, 1 day ago
eastus.api.cognitive.microsoft.com and </text/analytics/v3.1/languages>
upvoted 1 times
- ✉ **varinder82** 2 months ago
Final Answer:
Box1:
<https://eastus.api.cognitive.microsoft.com>
Box2:
</text/analytics/v3.1/languages>
upvoted 1 times
- ✉ **audlindr** 2 months, 3 weeks ago
This was in 2-Mar-24 exam
upvoted 5 times
- ✉ **sca88** 6 months, 1 week ago
<https://eastus.dev.cognitive.microsoft.com/docs/services/TextAnalytics.V2.0/operations/56f30ceeeda5650db055a3c7>
upvoted 1 times
- ✉ **zellck** 11 months ago
1. eastus.api.cognitive.microsoft.com
2. </text/analytics/v3.1/languages>

<https://learn.microsoft.com/en-us/rest/api/cognitiveservices-textanalytics/3.0/languages/languages>
Detect Language
The API returns the detected language and a numeric score between 0 and 1. Scores close to 1 indicate 100% certainty that the identified language is true.
- POST {Endpoint}/text/analytics/v3.0/languages
- Endpoint
Supported Cognitive Services endpoints (protocol and hostname, for example: <https://westus.api.cognitive.microsoft.com>).
upvoted 8 times
- ✉ **Pixelmate** 11 months ago
This came in exam 28/06/2023
upvoted 3 times
- ✉ **zggy1117** 11 months, 3 weeks ago
The scenario needs a Language service. Specifically just a way to know if it's in English or French. No translation is needed. Hence:
Box1:
<https://eastus.api.cognitive.microsoft.com>
Box2:
</text/analytics/v3.1/languages>

Source: <https://learn.microsoft.com/en-us/rest/api/cognitiveservices-textanalytics/3.1preview4/languages/languages?tabs=HTTP>

upvoted 3 times

✉️ **RAN_L** 1 year, 2 months ago

api.cognitive.microsofttranslator.com
eastus.api.cognitive.microsoft.com
portal.azure.com

Out of these three options, only option 2 is a valid Azure Cognitive Services endpoint. The eastus.api.cognitive.microsoft.com endpoint is used for various Azure Cognitive Services, including the Text Analytics and Translator services.

Option 1, api.cognitive.microsofttranslator.com, is not a valid endpoint for Azure Cognitive Services. It appears to be a typo, as "cognitive" is misspelled.

Option 3, portal.azure.com, is the URL for the Azure portal, which is a web-based management interface for Azure services. It is not an endpoint for Azure Cognitive Services APIs.

-./text/analytics/v3.1/languages

One of the APIs that you can use to process incoming email and direct messages to either French or English language support teams is the Text Analytics API. The specific API endpoint you can use is /text/analytics/v3.1/languages. This API can detect the language of a given text and return the ISO 639-1 language code. You can then use this information to route the message to the appropriate language support team.

upvoted 3 times

✉️ **Ody** 2 months, 1 week ago

api.cognitive.microsofttranslator.com is a valid endpoint for text translation

<https://learn.microsoft.com/en-us/azure/ai-services/translator/quickstart-text-rest-api?tabs=csharp#prerequisites>

upvoted 1 times

✉️ **Nicoseal** 1 year, 2 months ago

mk1967 and Adedoyin_Simeon are correct

The correct answer should be:

Box1:

<https://eastus.api.cognitive.microsoft.com>

Box2:

/text/analytics/v3.1/languages

upvoted 4 times

✉️ **Adedoyin_Simeon** 1 year, 5 months ago

The answer is wrong. Correct answer should be:

<https://eastus.api.cognitive.microsoft.com>

&

/text/analytics/v3.1/languages

REF:

<https://learn.microsoft.com/en-us/rest/api/cognitiveservices-textanalytics/3.0/languages/languages?tabs=HTTP>.

NOTE:

Pay special attention to the Sample Request provided. Request to the API should be of the form:

POST {Endpoint}/text/analytics/v3.0/languages

Where the {Endpoint} as stated under the sub-heading "URI Parameters" was described as quoted here (see "Description" column of the table):

"Supported Cognitive Services endpoints (protocol and hostname, for example: <https://westus.api.cognitive.microsoft.com>)."

So the sample given shows the correct format of the *endpoint* as <https://{location}.api.cognitive.microsoft.com>

upvoted 2 times

✉️ **Adedoyin_Simeon** 1 year, 5 months ago

The answer is wrong. Correct answer should be:

<https://eastus.api.cognitive.microsoft.com>

&

/text/analytics/v3.1/languages

REF:

<https://learn.microsoft.com/en-us/rest/api/cognitiveservices-textanalytics/3.0/languages/languages?tabs=HTTP>.

NOTE:

Pay special attention to the Sample Request provided. Request to the API should be of the form:

POST {Endpoint}/text/analytics/v3.0/languages

Where the {Endpoint} as stated under the sub-heading "URI Parameters" was described as quoted here

"Supported Cognitive Services endpoints (protocol and hostname, for example: <https://westus.api.cognitive.microsoft.com>)."

So the sample given shows the correct format of the *endpoint* as <https://{location}.api.cognitive.microsoft.com>

upvoted 1 times

✉️ **HotDurian** 1 year, 5 months ago

Answer is incorrect. Language Cognitive Service is sufficient.

<https://learn.microsoft.com/en-us/rest/api/cognitiveservices-textanalytics/3.0/languages/languages?tabs=HTTP>

upvoted 1 times

Question #31

You have an Azure Cognitive Search instance that indexes purchase orders by using Form Recognizer.

You need to analyze the extracted information by using Microsoft Power BI. The solution must minimize development effort.

What should you add to the indexer?

- A. a projection group
- B. a table projection
- C. a file projection
- D. an object projection

Correct Answer: D

Projections are the physical tables, objects, and files in a knowledge store that accept content from a Cognitive Search AI enrichment pipeline.

If you're creating a knowledge store, defining and shaping projections is most of the work.

Objects is used when you need the full JSON representation of your data and enrichments in one JSON document. As with table projections, only valid JSON objects can be projected as objects, and shaping can help you do that.

Note: Form Recognizer analyzes your forms and documents, extracts text and data, maps field relationships as key-value pairs, and returns a structured JSON output. You quickly get accurate results that are tailored to your specific content without excessive manual intervention or extensive data science expertise.

Incorrect:

Not Tables: Tables is used for data that's best represented as rows and columns, or whenever you need granular representations of your data (for example, as data frames). Table projections allow you to define a schematized shape, using a Shaper skill or use inline shaping to specify columns and rows.

Not File: File is used when you need to save normalized, binary image files.

Reference:

<https://docs.microsoft.com/en-us/azure/search/knowledge-store-projection-overview>

Community vote distribution

B (90%)

10%

✉  **momentumhd**  1 year, 8 months ago

Selected Answer: B

Should be B . Its for Tables the Power BI

" Use Power BI for data exploration. This tool works best when the data is in Azure Table Storage. Within Power BI, you can manipulate data into new tables that are easier to query and analyze"

upvoted 19 times

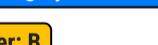
✉  **uira** 1 year, 1 month ago

You receive a JSON object, so ObjectProjection is the most appropriate way to explore:

Objects: "Used when you need the full JSON representation of your data and enrichments in one JSON document. As with table projections, only valid JSON objects can be projected as objects, and shaping can help you do that."

As per <https://learn.microsoft.com/en-us/azure/search/knowledge-store-projection-overview>

upvoted 4 times

✉  **Lion007**  7 months ago

Selected Answer: B

B is the correct answer. See below to understand the workflow:

Purchase Orders (POs) -> Form Recognizer -> OCR -> JSON (extracted info from POs) -> Shaper skill -> JSON -> Table Projection -> JSON -> Power BI

Ref: Define a table projection <https://learn.microsoft.com/en-us/azure/search/knowledge-store-projections-examples#define-a-table-projection>
upvoted 6 times

✉  **meluk**  2 days, 21 hours ago

The correct answer is B. a table projection.

Here's why:

A table projection in the indexer allows you to extract structured data from the indexed content.

By using a table projection, you can easily analyze the extracted information in Microsoft Power BI without extensive development effort.

It aligns with the requirement of minimizing development effort while enabling efficient data analysis.

Therefore, adding a table projection to the indexer is the most suitable choice for this scenario1.

upvoted 2 times

✉ **nanaw770** 3 days, 18 hours ago

Selected Answer: B

It must be B.

upvoted 1 times

✉ **reiwanotora** 1 week, 1 day ago

Selected Answer: B

B is right.

upvoted 1 times

✉ **anntv252** 3 weeks, 6 days ago

Its for Tables the Power BI

upvoted 1 times

✉ **Ronny05** 1 month, 2 weeks ago

Answer is Table Projection

Table projections are recommended for scenarios that call for data exploration, such as analysis with Power BI or workloads that consume data frames.

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-projections-examples#define-a-table-projection>

upvoted 1 times

✉ **evangelist** 3 months, 2 weeks ago

Selected Answer: B

B. a table projection

A table projection in the context of Azure Cognitive Search allows you to transform complex data structures into a flat tabular model. This is particularly useful when dealing with nested or complex data extracted from documents by Form Recognizer and indexed by Azure Cognitive Search. By projecting this data into a table format, you make it easier to import and analyze in Power BI, which excels at working with tabular data

upvoted 3 times

✉ **sca88** 6 months, 1 week ago

Selected Answer: B

I think is B. Look the exercise

<https://microsoftlearning.github.io/mslearn-knowledge-mining/Instructions/Labs/03-knowledge-store.html>.

It says:

"You may want to normalize index records into a relational schema of tables, for query analysis and reporting with tools such as Microsoft Power BI."

upvoted 3 times

✉ **kks0805** 6 months, 1 week ago

Selected Answer: D

Should be D, the given answer is correct.

upvoted 2 times

✉ **rdemontis** 6 months, 4 weeks ago

Selected Answer: B

To analyze the extracted information from the Azure Cognitive Search index with Microsoft Power BI, you should add a table projection to the indexing. This will allow you to present the data in a tabular format that can be easily imported and analyzed by Power BI with minimal development effort.

So, the correct answer is:

B. a table projection

(ChatGPT)

upvoted 1 times

✉ **rdemontis** 6 months, 2 weeks ago

Sorry but I have to revise my previous answer: Formrecognizer always returns results in json format

<https://learn.microsoft.com/en-us/azure/architecture/ai-ml/architecture/automate-document-processing-azure-form-recognizer>

so in that case it would be necessary to use Object Projection in Search.

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-projection-overview>

So Power BI can use the Json Connector to import the data in json format and parse it

<https://learn.microsoft.com/it-it/power-query/connectors/json>

upvoted 2 times

✉ **james2033** 9 months, 2 weeks ago

Selected Answer: B

Table projection: <https://learn.microsoft.com/en-us/azure/search/knowledge-store-projections-examples#define-a-table-projection>

File projection: for binary

Object projection: for tree structure

Table projection: for data records

upvoted 1 times

✉ **zellck** 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-connect-power-bi>

A knowledge store that's composed of tables in Azure Storage work best in Power BI. If the tables contain projections from the same skillset and projection group, you can easily "join" them to build table visualizations that include fields from related tables.

upvoted 4 times

✉ **ziggyl117** 11 months, 3 weeks ago

Selected Answer: B

Should be B. Power BI needs Table projections

upvoted 2 times

✉ **EliteAllen** 1 year ago

Selected Answer: B

B. a table projection - Correct. The table projection feature in Azure Cognitive Search allows you to flatten complex data structures into a format that can be easily indexed and queried. This is especially useful when you want to analyze the extracted information using Power BI, as Power BI works best with flattened data structures.

upvoted 2 times

✉ **Pfffff** 1 year, 1 month ago

Selected Answer: D

ChatGPT: Object projections are a better option than tables in this scenario because they minimize the amount of data that needs to be transferred from the search index to Power BI, reducing latency and improving performance. Additionally, object projections are simpler to set up and require less configuration than tables.

To add an object projection to your indexer, you can use the Azure Cognitive Search portal or the Azure Cognitive Search REST API. You will need to define a mapping that specifies which fields from the source document should be included in the object projection, and how they should be mapped to JSON properties.

upvoted 1 times

✉ **RAN_L** 1 year, 2 months ago

Selected Answer: B

To enable integration with Power BI, you should add a table projection to the indexer. This allows the data to be easily queried and visualized in Power BI.

A projection in Azure Cognitive Search is a way to specify the fields that should be included in the search index. A table projection specifies the columns of a table to be projected as fields in the search index.

By including a table projection in the indexer, you can extract the relevant data from the purchase orders and make it available for analysis in Power BI. This would minimize development effort, as Power BI has built-in integration with Azure Cognitive Search and can easily consume data from the search index.

Therefore, the correct answer is B. a table projection.

upvoted 3 times

Question #32

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure Cognitive Search service.

During the past 12 months, query volume steadily increased.

You discover that some search query requests to the Cognitive Search service are being throttled.

You need to reduce the likelihood that search query requests are throttled.

Solution: You add replicas.

Does this meet the goal?

A. Yes

B. No

Correct Answer: A

A simple fix to most throttling issues is to throw more resources at the search service (typically replicas for query-based throttling, or partitions for indexing-based throttling). However, increasing replicas or partitions adds cost, which is why it is important to know the reason why throttling is occurring at all.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-performance-analysis>

Community vote distribution

A (100%)

✉  **meluk** 2 days, 21 hours ago

The correct answer is B. No.

Here's why:

Adding replicas does not directly address the issue of query throttling in Azure Cognitive Search.

Query throttling occurs when the service limits the number of requests it can handle within a given time frame. It is typically related to the service's capacity, rate limits, and resource availability.

To reduce the likelihood of query throttling, consider other factors such as:

Service capacity: Ensure that your service tier (including replicas and partitions) can handle the query volume.

Optimize query design: Efficiently structure your queries to minimize resource usage.

Index composition: Smaller indexes tend to perform better, but also consider schema complexity and field attribution.

Therefore, adding replicas alone does not meet the goal of reducing query throttling. You need to address other aspects of your search service to achieve better performance and avoid throttling

upvoted 1 times

✉  **reiwanotora** 1 week, 1 day ago

Selected Answer: A

A is right.

upvoted 1 times

✉  **michaelmorar** 1 month, 2 weeks ago

Selected Answer: A

According to Microsoft's training material:

Use the best service tier for your search needs

You've seen that you can scale out service tiers by adding more partitions. You can scale out with replicas if you need to scale because of an increase in load. You can also scale up your search service by using a higher tier.

upvoted 1 times

✉  **evangelist** 3 months, 2 weeks ago

Selected Answer: A

add replicas can help read performance;

write performance needs adding partitions

upvoted 1 times

✉  **rdemontis** 6 months, 4 weeks ago

Selected Answer: A

correct

upvoted 1 times

✉️ **trashbox** 7 months ago

Selected Answer: A

The answer is correct.

upvoted 1 times

✉️ **james2033** 9 months, 2 weeks ago

Selected Answer: A

Quote "In Cognitive Search, replicas are copies of your index." at <https://learn.microsoft.com/en-us/azure/search/search-reliability>

upvoted 1 times

✉️ **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/search/search-performance-analysis#throttling-behaviors>

Throttling occurs when the search service is at capacity. Throttling can occur during queries or indexing. From the client side, an API call results in 503 HTTP response when it has been throttled. During indexing, there's also the possibility of receiving a 207 HTTP response, which indicates that one or more items failed to index. This error is an indicator that the search service is getting close to capacity.

A simple fix to most throttling issues is to throw more resources at the search service (typically replicas for query-based throttling, or partitions for indexing-based throttling). However, increasing replicas or partitions adds cost, which is why it's important to know the reason why throttling is occurring at all. Investigating the conditions that cause throttling will be explained in the next several sections.

upvoted 4 times

✉️ **eth1** 1 year, 3 months ago

Selected Answer: A

From : <https://learn.microsoft.com/en-us/azure/search/search-performance-tips#index-size-and-schema>

However, if the index is right-sized, the only other calibration you can make is to increase capacity: either by adding replicas or upgrading the service tier

So adding replicas and upgrading the sku is two valid answers for this question

upvoted 4 times

✉️ **HotDurian** 1 year, 5 months ago

Selected Answer: A

Answer is correct

upvoted 1 times

✉️ **practicewizards** 1 year, 8 months ago

Shouldn't it be Cognitive Search with a higher tier?

upvoted 4 times

✉️ **MDWPartners** 2 weeks, 2 days ago

a higher tier does not always increase the resources

upvoted 1 times

Question #33

SIMULATION -

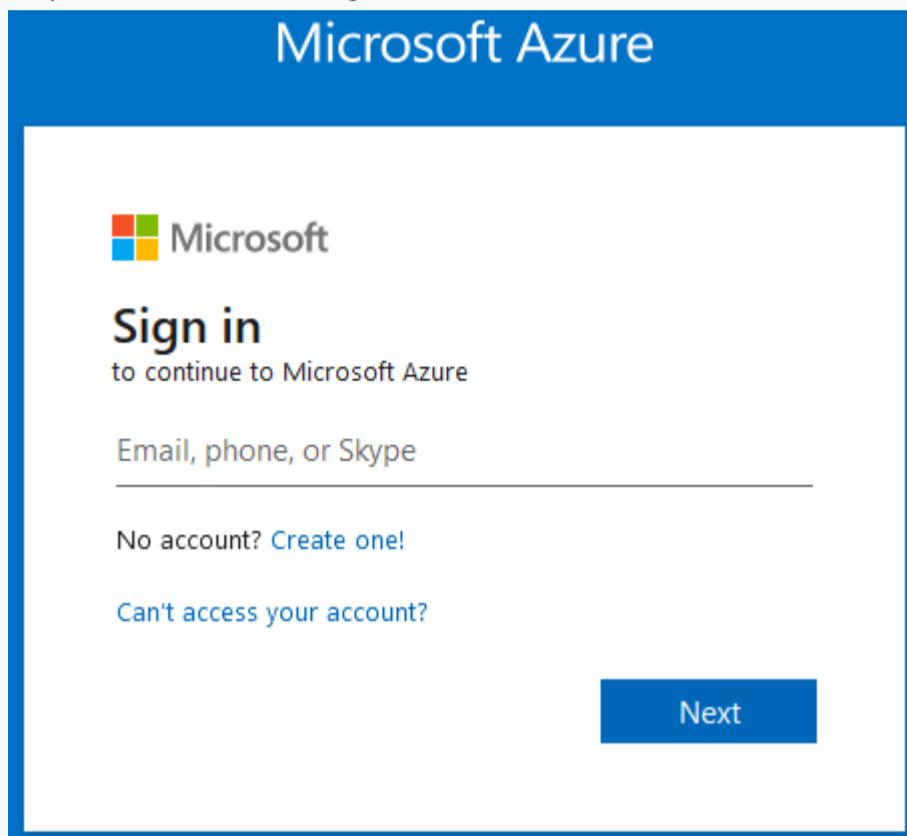
You need to create a Text Analytics service named Text12345678, and then enable logging for Text12345678. The solution must ensure that any changes to Text12345678 will be stored in a Log Analytics workspace.

To complete this task, sign in to the Azure portal.

Correct Answer: See explanation below.

Step 1: Sign in to the QnA portal.

Step 2: Create an Azure Cognitive multi-service resource:



Step 3: On the Create page, provide the following information.

Name: Text12345678 -

Create Cognitive Services

X

[Basics](#) [Tags](#) [Review + create](#)

Get access to Vision, Language, Search, and Speech Cognitive Services with a single API key. Quickly connect services together to achieve more insights into your content and easily integrate with other services like Azure Search. [Learn more](#)

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ

Resource group * ⓘ

[Create new](#)

Instance details

Region * ⓘ



Location specifies the region only for included regional services. This does not specify a region for included non-regional services. [Click here for more details.](#)

Name * ⓘ



Pricing tier * ⓘ

[View full pricing details](#)

By checking this box, I certify that use of this service is not by or for a police department in the United States.

I confirm I have read and understood the notice below.

[Review + create](#)

< Previous

Next : Tags >

Step 4: Configure additional settings for your resource as needed, read and accept the conditions (as applicable), and then select Review + create.

Step 5: Navigate to the Azure portal. Then locate and select The Text Analytics service resource Text12345678 (which you created in Step 4).

Step 6: Next, from the left-hand navigation menu, locate Monitoring and select Diagnostic settings. This screen contains all previously created diagnostic settings for this resource.

Step 7: Select + Add diagnostic setting.

Step 8: When prompted to configure, select the storage account and OMS workspace that you'd like to use to store your diagnostic logs. Note: If you don't have a storage account or OMS workspace, follow the prompts to create one.

Step 9: Select Audit, RequestResponse, and AllMetrics. Then set the retention period for your diagnostic log data. If a retention policy is set to zero, events for that log category are stored indefinitely.

Step 10: Click Save.

It can take up to two hours before logging data is available to query and analyze. So don't worry if you don't see anything right away.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/cognitive-services-apis-create-account> <https://docs.microsoft.com/en-us/azure/cognitive-services/diagnostic-logging>

✉ **puchuveryfatty** Highly Voted 4 months, 3 weeks ago

Will this "SIMULATION" type of questions come in exam? If yes, how to answer? It says to login to azure environment and then follow next steps.. How to perform those tasks in the exam?

upvoted 8 times

✉ **nanaw770** Most Recent 3 days, 18 hours ago

Simulation questions will not appear on the actual exam as of May 25, 2024; ET should remove this type of question. Ochinchin.

upvoted 1 times

✉ **chandiochan** 3 months, 2 weeks ago

Do we get simulation question in exams as of Feb 2024? Anyone please help

upvoted 2 times

✉ **ziggy1117** 11 months, 3 weeks ago

For any resource that needs this requirement:

1. First create a log analytics workspace.

2. Go to Diagnostic settings -> add diagnostic setting -> click audit, alllogs, allmetrics and choose the workspace created in 1 upvoted 4 times

✉️ **AnonymousJhb** 8 months, 1 week ago

#Step 1 - create your language service:
azure portal > search > azure ai services,
+ create language services,
continue to create your resource,
create your resource and accept the responsible AI notice,
Review & Create,
#Step 2 - add your LAW:
open your new language service,
search for diagnostic settings,
you may need to search across the top for your specific resource,
now + add your diagnostic setting,
name, desination log analytics workspace,
select your logs and metrics as required to minimize storage costs.
upvoted 2 times

✉️ **niru_1993** 4 months, 2 weeks ago

Is this type of answers will come in objective
upvoted 1 times

✉️ **practicewizards** 1 year, 7 months ago

This cannot be done anymore. Is not supported in azure since QnA is being retired on 31st March, 2025.
"QnA Maker service is being retired on 31st March, 2025. A newer version of this capability is now available as a part of Azure Cognitive Service for Language called question answering. To use this service, you need to provision a Language resource. For question answering capability within the Language service, see question answering and its pricing page. You can't create new QnA Maker resources anymore. For information on migrating your existing QnA Maker knowledge bases to question answering, consult the migration guide."

upvoted 3 times

Question #34

SIMULATION -

You need to create a search service named search12345678 that will index a sample Azure Cosmos DB database named hotels-sample. The solution must ensure that only English language fields are retrievable.

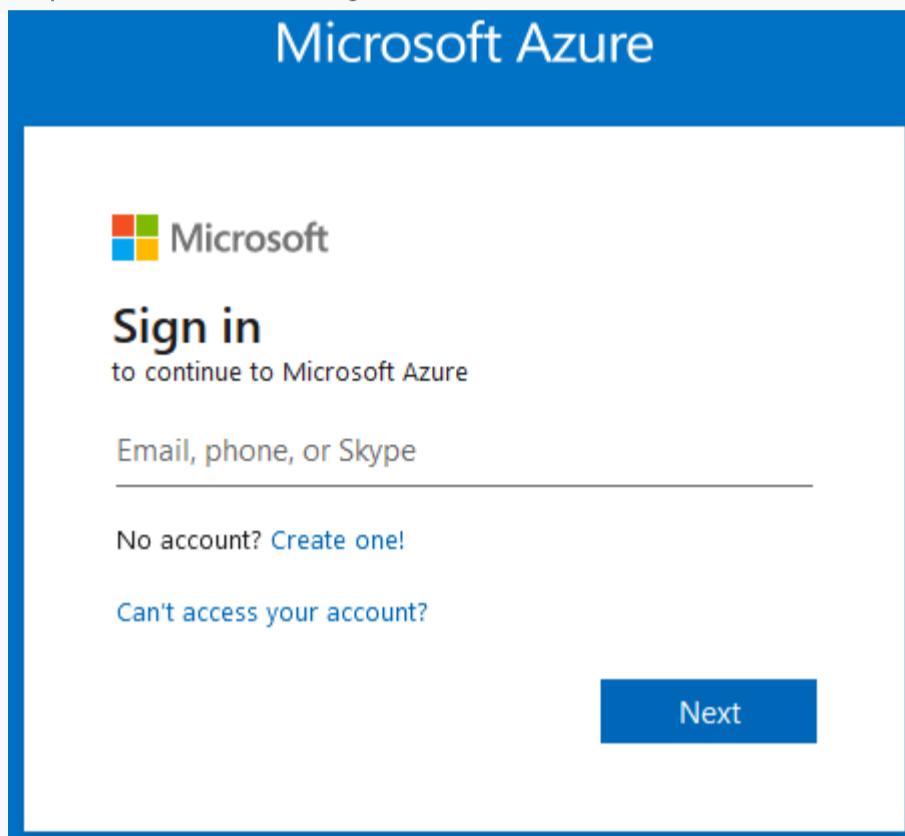
To complete this task, sign in to the Azure portal.

Correct Answer: See explanation below.

Part 1: Create a search service search12345678

Step 1: Sign in to the QnA portal.

Step 2: Create an Azure Cognitive multi-service resource:



Step 3: On the Create page, provide the following information.

Name: search12345678 -

Create Cognitive Services

X

[Basics](#) [Tags](#) [Review + create](#)

Get access to Vision, Language, Search, and Speech Cognitive Services with a single API key. Quickly connect services together to achieve more insights into your content and easily integrate with other services like Azure Search. [Learn more](#)

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ

Resource group * ⓘ

[Create new](#)

Instance details

Region * ⓘ

West US 2



Location specifies the region only for included regional services. This does not specify a region for included non-regional services. [Click here for more details.](#)

Name * ⓘ



Pricing tier * ⓘ

[View full pricing details](#)

By checking this box, I certify that use of this service is not by or for a police department in the United States.

I confirm I have read and understood the notice below.

[Review + create](#)

< Previous

Next : Tags >

Step 4: Click Review + create -

Part 2: Start the Import data wizard and create a data source

Step 5: Click Import data on the command bar to create and populate a search index.

+ Add index Import data Search explorer Refresh Delete Move

Step 6: In the wizard, click Connect to your data > Samples > hotels-sample. This data source is built-in. If you were creating your own data source, you would need to specify a name, type, and connection information. Once created, it becomes an "existing data source" that can be reused in other import operations.

Home > Microsoft.Search - Overview > my-new-search-service > Import data

Import data

Connect to your data Enrich content (Optional) Customize target index Create an indexer

Create and load a search index using data from an existing Azure data source in your current subscription. Azure Cognitive Search crawls the data structure you provide, extracts searchable content, optionally enriches it with cognitive skills, and loads it into an index. [Learn more](#)

Data Source	① <input type="button"/>	Samples
Type		Name
		realestate-us-sample
	② <input type="button"/>	hotels-sample

Step 7: Continue to the next page.

Step 8: Skip the "Enrich content" page

Step 9: Configure index.

Make sure English is selected for the fields.

Step 10: Continue and finish the wizard.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/cognitive-services-apis-create-account> <https://docs.microsoft.com/en-us/azure/search/search-get-started-portal>

✉ **nanaw770** 3 days, 18 hours ago

Simulation questions will not appear on the actual exam as of May 25, 2024; ET should remove this type of question.
upvoted 1 times

✉ **PeteColag** 1 day, 23 hours ago

Can you please provide an online reference from Microsoft which confirms that simulation exams will no longer appear in the exam. I have not been able to find any information on this.
upvoted 1 times

✉ **ArminZ11** 5 months ago

step 2 create an 'AI Search' resource (not 'Azure Cognitive multi-service resource')
step 3 shows a wrong screen
upvoted 3 times

✉ **PeteColag** 1 day, 23 hours ago

You are right. Many of the instructions provided in the examples are no longer up to date.
upvoted 1 times

✉ **GHill1982** 2 months, 2 weeks ago

Agree - you now create an AI Search since the services were renamed
upvoted 1 times

✉ **rdemontis** 6 months, 4 weeks ago

correct approach
upvoted 1 times

✉ **emilchrs** 6 months, 1 week ago

are these questions/simulations in the actual exam?
upvoted 4 times

✉ **audlindr** 2 months, 3 weeks ago

No simulation question as on 2-Mar-2024
upvoted 6 times

Question #35

SIMULATION -

You plan to create a solution to generate captions for images that will be read from Azure Blob Storage.

You need to create a service in Azure Cognitive Services for the solution. The service must be named captions12345678 and must use the Free pricing tier.

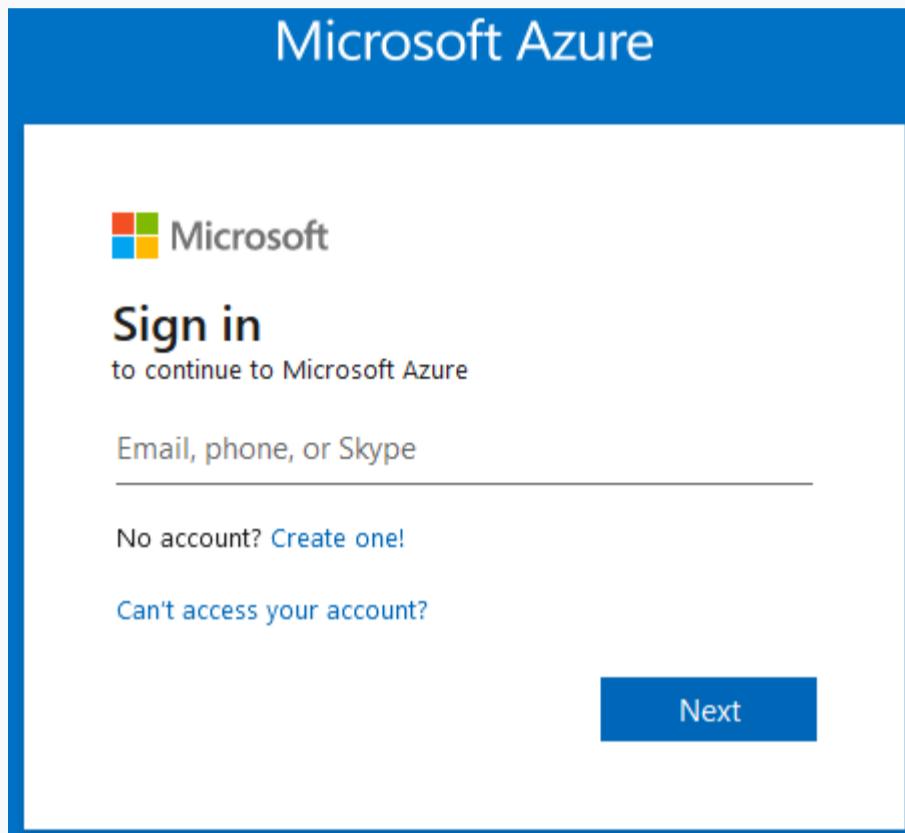
To complete this task, sign in to the Azure portal.

Correct Answer: See explanation below.

Part 1: Create a search service captions12345678

Step 1: Sign in to the QnA portal.

Step 2: Create an Azure Cognitive multi-service resource:



Step 3: On the Create page, provide the following information.

Name: captions12345678

Pricing tier: Free -

Create Cognitive Services

X

[Basics](#) [Tags](#) [Review + create](#)

Get access to Vision, Language, Search, and Speech Cognitive Services with a single API key. Quickly connect services together to achieve more insights into your content and easily integrate with other services like Azure Search. [Learn more](#)

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ

Resource group * ⓘ

[Create new](#)

Instance details

Region * ⓘ



Location specifies the region only for included regional services. This does not specify a region for included non-regional services. [Click here for more details.](#)

Name * ⓘ

 ✓

Pricing tier * ⓘ

[View full pricing details](#)

By checking this box, I certify that use of this service is not by or for a police department in the United States.

I confirm I have read and understood the notice below.

[Review + create](#)[< Previous](#)[Next : Tags >](#)

Step 4: Click Review + create -

(Step 5: Create a data source

In Connect to your data, choose Azure Blob Storage. Choose an existing connection to the storage account and container you created. Give the data source a name, and use default values for the rest.)

Microsoft Azure Search resources, services, and docs (G+) ...

Dashboard > Import data X

* **Connect to your data** Add cognitive skills (Optional) Customize target index Create an indexer

Create and load a search index using data from an external data source. Azure Cognitive Search crawls the data structure you provide, extracts searchable content, optionally enriches it with cognitive skills, and loads it into an index. [Learn more](#)

Data Source		Azure Blob Storage	✓
Data source name *		signs	✓
Data to extract ⓘ		Content and metadata	✓
Parsing mode		Default	✓
Connection string *		DefaultEndpointsProtocol=https;AccountName= ...	✓
		Choose an existing connection	
		<input type="checkbox"/> Authenticate using managed identity ⓘ	
Container name * ⓘ		signs	✓
Blob folder ⓘ		your/folder/here	
Description		(optional)	

Next: Add cognitive skills (Optional)

Reference:
<https://docs.microsoft.com/en-us/azure/search/search-create-service-portal> <https://docs.microsoft.com/en-us/azure/search/cognitive-search-quickstart-ocr>

⊕ **halfway** Highly Voted 1 year, 6 months ago

Create a 'Computer Vision' service and use it for image captioning:

upvoted 8 times

⊕ **rde montis** 6 months, 4 weeks ago

to me you are right. It is only requested to create a resource to generate image captions

upvoted 1 times

⊕ **nanaw770** Most Recent 3 days, 18 hours ago

Simulation questions will not appear on the actual exam as of May 25, 2024; ET should remove this type of question. Ochinchin.

upvoted 1 times

⊕ **Ghill1982** 2 months, 2 weeks ago

I think the steps are:

1. Create a Computer Vision resource in Azure AI services using the Free pricing tier.

2. Launch Vision Studio.

3. Select your Computer Vision resource.

4. Add new dataset.

5. Enter a dataset name and select Image classification as the model type.

6. Select an Azure blob storage container and Create dataset.

(I don't know in the exam simulation whether the Storage account with a Blob container already exists, if not this would need to be created first)

upvoted 4 times

⊕ **Ody** 2 months, 1 week ago

I tend to agree, but think the question is unclear.

upvoted 1 times

✉ **ziggy1117** 11 months, 4 weeks ago

I simulated this in my Azure Portal:

1. Create Cognitive Search Resource
2. Create a Blob Storage. Add a container and add the images with text.
3. In Cognitive Search, click Import Data then choose Data Source: Azure Blob Storage. Link the Blob Storage Connection String to the Container in #2.
4. In the Skillset Click Enable OCR then Click Image Cognitive Skills
5. Click Create till the very end and you are done

upvoted 2 times

✉ **propanther** 9 months, 3 weeks ago

Is simulation part of the exam?

upvoted 5 times

✉ **RupRizal** 6 months, 1 week ago

I want to know if Simulation is part of actual exam! Anyone?

upvoted 2 times

✉ **audlindr** 2 months, 3 weeks ago

It was not on 2-Mar-2024

upvoted 2 times

✉ **momentumhd** 1 year, 8 months ago

Should we use Cognitive Search for this?

upvoted 1 times

Question #36

SIMULATION -

You need to create a Form Recognizer resource named fr12345678.

Use the Form Recognizer sample labeling tool at <https://fott-2-1.azurewebsites.net/> to analyze the invoice located in the C:\Resources\Invoices folder.

Save the results as C:\Resources\Invoices\Results.json.

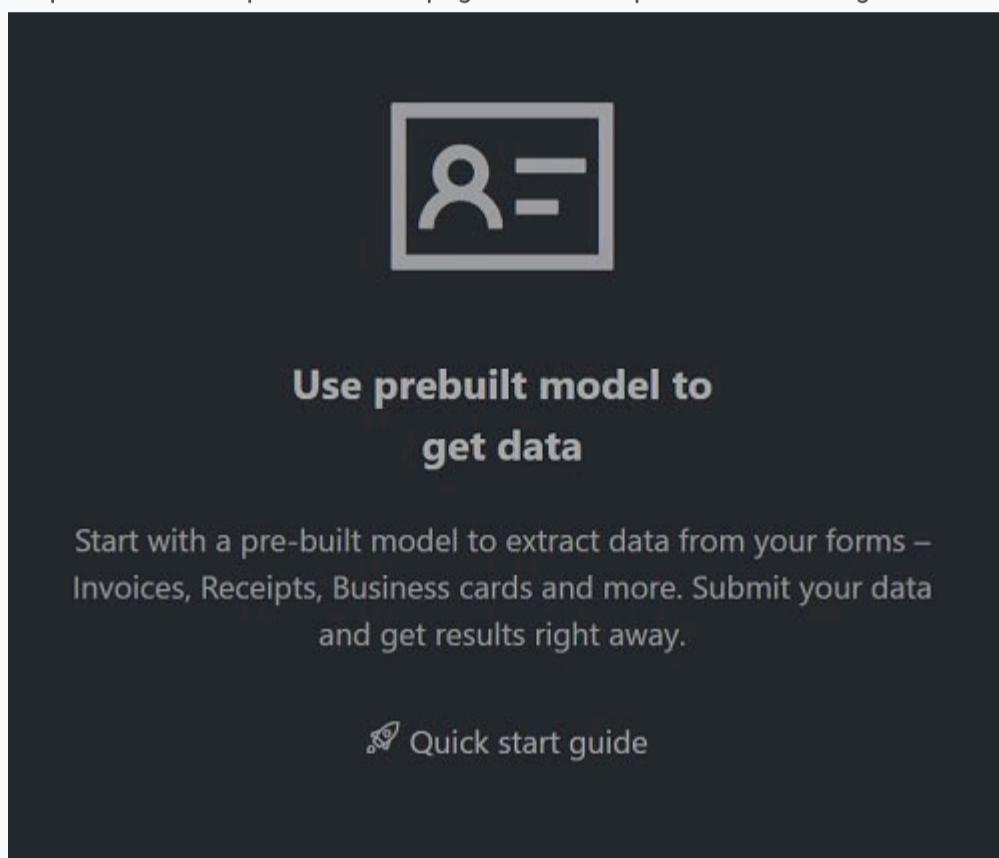
To complete this task, sign in to the Azure portal and open the Form Recognizer sample labeling tool.

Correct Answer: See explanation below.

Step 1: Sign in to the Azure Portal.

Step 2: Navigate to the Form Recognizer Sample Tool (at <https://fott-2-1.azurewebsites.net/>)

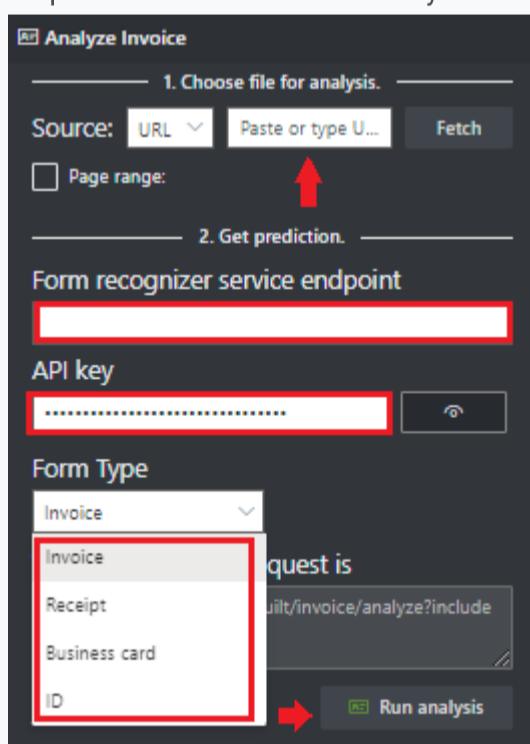
Step 3: On the sample tool home page select Use prebuilt model to get data.



Step 4: Select the Form Type you would like to analyze from the dropdown window.

Step 5: In the Source: URL field, paste the selected URL and select the Fetch button.

Step 6: In the Choose file for analysis use the file in the C:\Resources\Invoices folder and select the Fetch button.



Step 7: Select Run analysis. The Form Recognizer Sample Labeling tool will call the Analyze Prebuilt API and analyze the document.

Step 8: View the results - see the key-value pairs extracted, line items, highlighted text extracted and tables detected.

Step 9: Save the results as C:\Resources\Invoices\Results.json.

Reference:

<https://docs.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/quickstarts/try-sample-label-tool>

✉ **nanaw770** 3 days, 18 hours ago

Simulation questions will not appear on the actual exam as of May 25, 2024; ET should remove this type of question.

upvoted 1 times

✉ **michaelmorar** 1 month, 2 weeks ago

The issue I found when trying this is that you cannot choose the downloaded JSON file name. So one cannot ensure that the file is called Results.json - the studio names it for you.

upvoted 1 times

✉ **MonicaKarim** 3 weeks ago

don't choose download, instead press ctr+s then save as C:\Resources\Invoices\Results.json

upvoted 1 times

✉ **GHill1982** 2 months, 2 weeks ago

I think this would be the method as of March 2024:

1. Create a Document Intelligence resource in Azure AI services
2. Launch Document Intelligence Studio
3. Select Invoices from the Prebuilt models
4. Configure the service resource by selecting your Document Intelligence Resource
5. Drag & drop or browse for the invoice in C:\Resources\Invoices
6. Click Run analysis
7. Click on Result and the Download icon to save the JSON results file

upvoted 1 times

✉ **josebernabeo** 4 months, 1 week ago

Are these types of questions going to be on the exam?

upvoted 2 times

✉ **audlindr** 2 months, 3 weeks ago

It was not on 2-Mar-2024

upvoted 6 times

✉ **rdemontis** 6 months, 4 weeks ago

1. Create a form recognizer service as part of azure ai service
2. browse to <https://fott-2-1.azurewebsites.net/>
3. select prebuilt model for invoices
4. choose local file because the file is a local disk c: and insert the path
5. come back to the azure portal and copy endpoint and key from the relative page of the form recognizer service
6. come back to <https://fott-2-1.azurewebsites.net/prebuilt-analyze>
7. past endpoint and key
8. run analysis
9. download
10. choose json format and the destination indicated

upvoted 3 times

Question #37

You have a factory that produces food products.

You need to build a monitoring solution for staff compliance with personal protective equipment (PPE) requirements. The solution must meet the following requirements:

- * Identify staff who have removed masks or safety glasses.
- * Perform a compliance check every 15 minutes.
- * Minimize development effort.
- * Minimize costs.

Which service should you use?

- A. Face
- B. Computer Vision
- C. Azure Video Analyzer for Media (formerly Video Indexer)

Correct Answer: A

Face API is an AI service that analyzes faces in images.

Embed facial recognition into your apps for a seamless and highly secured user experience. No machine-learning expertise is required.

Features include face detection that perceives facial features and attributes such as a face mask, glasses, or face location in an image, and identification of a person by a match to your private repository or via photo ID.

Reference:

<https://azure.microsoft.com/en-us/services/cognitive-services/face/>

Community vote distribution

✉ **Davard** 1 year, 8 months ago

A. Face. The solution link explains:

Embed facial recognition into your apps for a seamless and highly secured user experience. No machine-learning expertise is required. Features include face detection that perceives facial features and attributes—such as a face mask, glasses, or face location—in an image, and identification of a person by a match to your private repository or via photo ID.

upvoted 8 times

✉ **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/overview-identity#face-detection-and-analysis>

Face detection is required as a first step in all the other scenarios. The Detect API detects human faces in an image and returns the rectangle coordinates of their locations. It also returns a unique ID that represents the stored face data. This is used in later operations to identify or verify faces.

Optionally, face detection can extract a set of face-related attributes, such as head pose, age, emotion, facial hair, and glasses. These attributes are general predictions, not actual classifications. Some attributes are useful to ensure that your application is getting high-quality face data when users add themselves to a Face service. For example, your application could advise users to take off their sunglasses if they're wearing sunglasses.

upvoted 5 times

✉ **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 3 times

✉ **PeteColag** 1 day, 21 hours ago

Selected Answer: C

According to both chatgpt (4o) and Gemini, the correct answer is C.

upvoted 1 times

✉ **juicingbliss** 1 week, 2 days ago

B is the answer.

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/intro-to-spatial-analysis-public-preview?tabs=sa#social-distancing-and-face-mask-detection>

upvoted 1 times

✉ **pramod80** 3 weeks, 5 days ago

Answer : A

upvoted 1 times

✉ **Jimmy1017** 1 month ago

answer B

upvoted 1 times

✉ **NullVoider_0** 1 month, 3 weeks ago

Selected Answer: C

C. Azure Video Analyzer for Media (formerly Video Indexer) - This service is designed to analyze video content and can be used to identify specific scenarios, such as whether individuals in a video are wearing masks or safety glasses. It supports a range of video analytics capabilities, including object detection, which can be leveraged to monitor PPE compliance. Using Azure Video Analyzer for Media can minimize development effort by providing built-in models for analyzing video content and can be set up to perform compliance checks at specified intervals, such as every 15 minutes.

upvoted 1 times

✉ **Mehe323** 2 months, 4 weeks ago

Selected Answer: B

I don't think Face alone can do a compliance check, Face does the recognition but the service needs to do something after that (determine it is compliant/not compliant, so you need to have the more comprehensive services of Computer Vision (AI vision) which includes Face.

upvoted 1 times

✉ **audlindr** 3 months ago

Selected Answer: B

Should be B Computer Vision. Spatial Analysis in Computer vision. Spatial Analysis can also be configured to detect if a person is wearing a protective face covering such as a mask.

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/intro-to-spatial-analysis-public-preview?tabs=sa>

upvoted 3 times

✉ **GHill1982** 1 month, 3 weeks ago

Thanks for the link. I agree Azure AI Vision Spatial Analysis seems correct.

upvoted 1 times

✉ **evangelist** 3 months, 2 weeks ago

Selected Answer: A

No doubt the correct answer is A: link is here

<https://learn.microsoft.com/en-us/legal/cognitive-services/face/characteristics-and-limitations>

upvoted 1 times

✉ **rdemontis** 6 months, 4 weeks ago

Selected Answer: A

I suppose the correct answer is A because every 15 minutes you can take a snapshot from the video surveillance system and process the image with the Face API. This is certainly less expensive and easier than using Video Indexer.

upvoted 2 times

✉ **sl_mslconsulting** 7 months, 1 week ago

Selected Answer: C

We are dealing with videos here and you need to check every 15 minutes. All these need to be done with minimum development and cost. Face won't work as it mainly deals with images and PPEs are not limited to just goggles. Check different video models available here:

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/video-indexer-overview>

upvoted 2 times

✉ **james2033** 9 months, 2 weeks ago

Selected Answer: A

Quote "For example, your application could advise users to take off their sunglasses if they're wearing sunglasses."

at <https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/overview-identity#face-detection-and-analysis>

upvoted 1 times

✉ **foexams** 11 months ago

Selected Answer: B

PPE does not include only face. Therefore, as chat GPT says, the correct answer is Computer Vision

upvoted 1 times

✉ **examworld** 1 year ago

azure Spatial Analysis is correct. This question is not up to date.

upvoted 2 times

✉ **KingChuang** 1 year, 4 months ago

on my exam (2023-01-16 Passed)

My Answer:A

upvoted 3 times

✉ **Davard** 1 year, 7 months ago

Selected Answer: A

Face includes face masks, glasses, or face location
upvoted 1 times

Question #38

You have an Azure Cognitive Search solution and a collection of blog posts that include a category field.

You need to index the posts. The solution must meet the following requirements:

- * Include the category field in the search results.
- * Ensure that users can search for words in the category field.
- * Ensure that users can perform drill down filtering based on category.

Which index attributes should you configure for the category field?

- A. searchable, sortable, and retrievable
- B. searchable, facetable, and retrievable
- C. retrievable, filterable, and sortable
- D. retrievable, facetable, and key

Correct Answer: C

Fields have data types and attributes. The check boxes across the top are index attributes controlling how the field is used.

- * Retrievable means that it shows up in search results list. You can mark individual fields as off limits for search results by clearing this checkbox, for example for fields used only in filter expressions.
- * Filterable, Sortable, and Facetable determine whether fields are used in a filter, sort, or faceted navigation structure.
- * Searchable means that a field is included in full text search. Strings are searchable. Numeric fields and Boolean fields are often marked as not searchable.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-get-started-portal>

Community vote distribution

 B (96%)

4%

✉  Internal_Koala  1 year, 8 months ago

Selected Answer: B

Vote for B

Retrievable: Include the category field in the search results.

Searchable: Ensure that users can search for words in the category field.

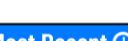
Facetable: Ensure that users can perform drill down filtering based on category.

upvoted 33 times

✉  qs99  4 months ago

Was on exam 22 Jan 2024

upvoted 6 times

✉  PeteColag  2 days, 8 hours ago

B. searchable, facetable, and retrievable

Here's why:

Searchable: This allows the category field to be included in full-text searches, meaning users can search for words within the category field.

Facetable: This allows the category field to be used in drill-down filtering, enabling users to filter search results based on the categories.

Retrievable: This ensures that the category field is included in the search results that are returned to the user.

These attributes together satisfy all the given requirements.

upvoted 1 times

✉  reiwanotora 1 week, 1 day ago

Selected Answer: B

B is right.

upvoted 1 times

✉  michaelmorar 1 month, 2 weeks ago

Selected Answer: B

Filterable - can be included in filter expressions

Facetable - can be used to determine values for facets (user interface elements used to filter results based on a list of known field values, including hit counts by category)

Retrievable - can be included in the search results - by default all fields are retrievable

upvoted 1 times

✉  f2c587e 2 months ago

Selected Answer: B

En la pregunta no es necesario que los resultados sean ordenables. Por eso la opcion es B
upvoted 2 times

✉ **audlindr** 3 months ago

Selected Answer: B

* Include the category field in the search results. - Retrievable: Fields returned in a query response.
* Ensure that users can search for words in the category field. - Searchable: Fields used in full text search. Strings are searchable. Numeric fields and Boolean fields are often marked as not searchable.
* Ensure that users can perform drill down filtering based on category. - Facetable: Fields used in a faceted navigation structure.
upvoted 2 times

✉ **evangelist** 3 months, 2 weeks ago

facetable and filterable are almost synonyms and they should not appear in the same options.
upvoted 3 times

✉ **evangelist** 3 months, 2 weeks ago

Selected Answer: B

For indexing blog posts in Azure Cognitive Search with the requirements specified, the correct attributes for the category field would be:

B. searchable, facetable, and retrievable

searchable: Allows users to search for words within the category field.

facetable: Enables drill down filtering based on category, which is essential for users to narrow down search results.

retrievable: Ensures that the category field is included in the search results, allowing users to see which category each blog post belongs to.

upvoted 2 times

✉ **rdemontis** 6 months, 4 weeks ago

Selected Answer: B

<https://learn.microsoft.com/en-us/azure/search/search-what-is-an-index>

upvoted 1 times

✉ **james2033** 9 months, 2 weeks ago

Selected Answer: B

What is facet query:

https://www.tutorialspoint.com/apache_solr/apache_solr_faceting.htm

<https://learn.microsoft.com/en-us/azure/search/search-faceted-navigation#faceted-navigation-in-a-search-page> See left side of image.

upvoted 1 times

✉ **msdfqwerfewf** 11 months ago

Selected Answer: C

!!! 'searchable' its based on full text search and not related to results. so it should be retrievable the field contents will show up in the results.
upvoted 2 times

✉ **zellick** 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/rest/api/searchservice/create-index#-field-definitions->

- retrievable

Indicates whether the field can be returned in a search result.

- searchable

Indicates whether the field is full-text searchable and can be referenced in search queries.

- facetable

Indicates whether to enable the field to be referenced in facet queries.

upvoted 4 times

✉ **zellick** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 4 times

✉ **zellick** 11 months ago

<https://learn.microsoft.com/en-us/azure/search/search-faceted-navigation>

Faceted navigation is used for self-directed drilldown filtering on query results in a search app, where your application offers form controls for scoping search to groups of documents (for example, categories or brands), and Azure Cognitive Search provides the data structures and filter to back the experience.

upvoted 1 times

✉ **EliteAllen** 11 months, 2 weeks ago

Selected Answer: B

"Searchable" attribute: This attribute allows users to search for words in the category field. It means that the field can be included in full-text searches.

"Facetable" attribute: This attribute allows users to perform drill-down filtering based on category. Faceting is used for self-directed drill-down

filtering on query results in a search app, where your application offers form controls for scoping search to groups of documents (for example, categories or brands), and Azure Cognitive Search provides the data structures and filters to back the experience.

"Retrievable" attribute: This attribute allows the category field to be included in the search results. It means that the field can be included in the search results returned by the search service.

upvoted 1 times

 **Sachz88** 1 year, 1 month ago

Chat GPT Response:

To meet the given requirements, you can configure the following index attributes for the "category" field in the Azure Cognitive Search solution:

"searchable": This attribute enables searching for words in the category field.

"filterable": This attribute enables filtering based on category.

"retrievable": This attribute ensures that the category field is included in the search results.

Therefore, the correct answer would be to configure the "searchable", "filterable", and "retrievable" attributes for the "category" field in the index schema of Azure Cognitive Search.

Hope it clarifies.

upvoted 1 times

 **NNU** 1 year, 3 months ago

The answer is correct C

<https://learn.microsoft.com/en-us/azure/search/search-get-started-dotnet>

upvoted 1 times

 **ap1234pa** 1 year, 4 months ago

Selected Answer: B

Drill down so facetable

upvoted 1 times

Question #39

SIMULATION -

Use the following login credentials as needed:

To enter your username, place your cursor in the Sign in box and click on the username below.

To enter your password, place your cursor in the Enter password box and click on the password below.

Azure Username: admin@abc.com -

Azure Password: XXXXXXXXXXXX -

The following information is for technical support purposes only:

Lab Instance: 12345678 -

Task -

You plan to build an API that will identify whether an image includes a Microsoft Surface Pro or Surface Studio.

You need to deploy a service in Azure Cognitive Services for the API. The service must be named AAA12345678 and must be in the East US Azure region. The solution must use the Free pricing tier.

To complete this task, sign in to the Azure portal.

Correct Answer: See explanation below.

Step 1: In the Azure dashboard, click Create a resource.

Step 2: In the search bar, type "Cognitive Services."

You'll get information about the cognitive services resource and a legal notice. Click Create.

Step 3: You'll need to specify the following details about the cognitive service (refer to the image below for a completed example of this page):

Subscription: choose your paid or trial subscription, depending on how you created your Azure account.

Resource group: click create new to create a new resource group or choose an existing one.

Region: choose the Azure region for your cognitive service. Choose: East US Azure region.

Name: choose a name for your cognitive service. Enter: AAA12345678

Pricing Tier: Select: Free pricing tier

Validation Passed

listed above; (b) authorize Microsoft to bill my current payment method for the fees associated with the offering(s), with the same billing frequency as my Azure subscription; and (c) agree that Microsoft may share my contact, usage and transactional information with the provider(s) of the offering(s) for support, billing and other transactional activities. Microsoft does not provide rights for third-party offerings. See the [Azure Marketplace Terms](#) for additional details.

Basics

Subscription	Visual Studio Enterprise Subscription
Resource group	ocr-rg
Region	West Europe
Name	ocr-cognitive-service
Pricing tier	Standard S0

Identity

Identity type	None
---------------	------

Create

< Previous

Next

[Download a template for automation](#)

Step 4: Review and create the resource, and wait for deployment to complete. Then go to the deployed resource.

Note: The Computer Vision Image Analysis service can extract a wide variety of visual features from your images. For example, it can determine whether an image contains adult content, find specific brands or objects, or find human faces.

Tag visual features -

Identify and tag visual features in an image, from a set of thousands of recognizable objects, living things, scenery, and actions. When the tags are ambiguous or not common knowledge, the API response provides hints to clarify the context of the tag. Tagging isn't limited to the main subject, such as a person in the foreground, but also includes the setting (indoor or outdoor), furniture, tools, plants, animals, accessories, gadgets, and so on.

Try out the image tagging features quickly and easily in your browser using Vision Studio.

Reference:

<https://docs.microsoft.com/en-us/learn/modules/analyze-images-computer-vision/3-analyze-images> <https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/overview-image-analysis>

✉️  **dev2dev**  1 year, 4 months ago

Computer Vision should be the resource we need to create.

upvoted 6 times

✉️  **rdemontis** 6 months, 4 weeks ago

agree with you

upvoted 1 times

✉️  **nanaw770**  3 days, 18 hours ago

Simulation questions will not appear on the actual exam as of May 25, 2024; ET should remove this type of question.

upvoted 2 times

✉️  **michaelmorar** 1 month ago

Both are Microsoft Products - so we don't need Brand recognition and it seems more like a Custom Vision solution.

HOWEVER, Computer Vision now seems to offer Product Recognition.

So, while I'm still leaning towards Custom Vision, it seems that Computer Vision might also be viable.

Question is probably out of date.

upvoted 1 times

✉️  **GHill1982** 2 months, 2 weeks ago

I think this is Custom Vision too. I don't believe Computer Vision is able to detect the difference between the two different devices which both have the Microsoft brand logo.

upvoted 3 times

✉️  **dimsok** 4 months, 2 weeks ago

I think it is a Custom Vision project because it is not about recognizing a brand, its about two laptops that need to be identified among tens of others....

upvoted 2 times

✉️  **ccampagna** 6 months, 1 week ago

This is a Computer Vision task, in computer vision you can extract the logo brands to identify if there is any. It also let you extract text to be analyzed if necessary. Custom Vision is useful to detect custom objects within an image or the image itself. More info here:

<https://stackoverflow.com/questions/52155632/difference-between-computer-vision-api-and-custom-vision-api>

upvoted 2 times

✉️  **ziggy1117** 11 months, 3 weeks ago

1. create cognitive service east-us, free
2. go to custom vision ai and choose #1
3. uploaded images of surface pro or studio
4. train the model
5. test
6. publish

upvoted 1 times

✉️  **AnonymousJhb** 5 months, 3 weeks ago

this is not CUSTOM vision because - Brand detection is a specialized mode of object detection that uses a database of thousands of global logos to identify commercial brands in images or video. You can use this feature, for example, to discover which brands are most popular on social media or most prevalent in media product placement.

The built-in logo database covers popular brands in consumer electronics, clothing, and more.

ONLY - If you find that the brand you're looking for is not detected by the Azure AI Vision service, you could also try creating and training your own logo detector using the Custom Vision service.

upvoted 2 times

Question #40

SIMULATION -

Use the following login credentials as needed:

To enter your username, place your cursor in the Sign in box and click on the username below.

To enter your password, place your cursor in the Enter password box and click on the password below.

Azure Username: admin@abc.com -

Azure Password: XXXXXXXXXXXX -

The following information is for technical support purposes only:

Lab Instance: 12345678 -

Task -

You need to build an API that uses the service in Azure Cognitive Services named AAA12345678 to identify whether an image includes a Microsoft Surface Pro or Surface Studio.

To achieve this goal, you must use the sample images in the C:\Resources\Images folder.

To complete this task, sign in to the Azure portal.

Correct Answer: See explanation below.

Step 1: In the Azure dashboard, click Create a resource.

Step 2: In the search bar, type "Cognitive Services."

You'll get information about the cognitive services resource and a legal notice. Click Create.

Step 3: You'll need to specify the following details about the cognitive service (refer to the image below for a completed example of this page):

Subscription: choose your paid or trial subscription, depending on how you created your Azure account.

Resource group: click create new to create a new resource group or choose an existing one.

Region: choose the Azure region for your cognitive service. Choose: East US Azure region.

Name: choose a name for your cognitive service. Enter: AAA12345678

Pricing Tier: Select: Free pricing tier

Step 4: Review and create the resource, and wait for deployment to complete. Then go to the deployed resource.

Note: The Computer Vision Image Analysis service can extract a wide variety of visual features from your images. For example, it can determine whether an image contains adult content, find specific brands or objects, or find human faces.

Tag visual features -

Identify and tag visual features in an image, from a set of thousands of recognizable objects, living things, scenery, and actions. When the tags are ambiguous or not common knowledge, the API response provides hints to clarify the context of the tag. Tagging isn't limited to the main subject, such as a person in the foreground, but also includes the setting (indoor or outdoor), furniture, tools, plants, animals, accessories, gadgets, and so on.

Try out the image tagging features quickly and easily in your browser using Vision Studio.

Reference:

<https://docs.microsoft.com/en-us/learn/modules/analyze-images-computer-vision/3-analyze-images> <https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/overview-image-analysis>

✉️  **t_isk**  1 year, 2 months ago

Is it possible to get questions like this on the exam?

upvoted 10 times

✉️  **ArminZ11** 4 months, 1 week ago

End of Dec 2023, no simulations in the exam

upvoted 13 times

✉️  **audlindr** 2 months, 3 weeks ago

Beginning of March 2024, no simulations in the exam

upvoted 4 times

✉️  **nanaw770**  3 days, 18 hours ago

Simulation questions will not appear on the actual exam as of May 25, 2024; ET should remove this type of question.

upvoted 2 times

✉ **chandiochan** 3 months, 2 weeks ago

Will these simulation questions appear in exams on 2024 as of Feb

upvoted 2 times

✉ **audlindr** 2 months, 3 weeks ago

No it was not on 2-Mar-2024

upvoted 5 times

✉ **ccampagna** 6 months, 1 week ago

This is a Computer Vision task, in computer vision you can extract the logo brands to identify if there is any. It also let you extract text to be analyzed if necessary. Custom Vision is useful to detect custom objects within an image or the image itself. More info here:

<https://stackoverflow.com/questions/52155632/difference-between-computer-vision-api-and-custom-vision-api>

upvoted 1 times

✉ **ziggy1117** 11 months, 4 weeks ago

steps are:

1. create cognitive resource
2. go to <https://www.customvision.ai/>.
3. upload images of each item. 5 minimum
4. train
5. verify via Quick test
6. publish

upvoted 2 times

✉ **marti_tremblay000** 1 year, 2 months ago

as per ChatGPT :

For identifying whether an image includes a Microsoft Surface Pro or Surface Studio, you can use the Azure Custom Vision service. Azure Custom Vision is a machine learning service that enables you to build, train, and deploy custom image classification models.

To use Azure Custom Vision for your API, you can follow these general steps:

Create an Azure Custom Vision project and upload a set of labeled images that include Microsoft Surface Pro or Surface Studio.

Train the model using the labeled images.

Publish the trained model as an API endpoint that can be used to classify new images.

Integrate the API endpoint into your own application to provide the image classification service.

upvoted 2 times

✉ **halfway** 1 year, 6 months ago

Image tagging won't work. It is either 'Brand Detection', if the sample images contain actual surface logos or 'Custom Vision' to train a model to detect surface pro/studios

upvoted 1 times

Question #41

SIMULATION -

Use the following login credentials as needed:

To enter your username, place your cursor in the Sign in box and click on the username below.

To enter your password, place your cursor in the Enter password box and click on the password below.

Azure Username: admin@abc.com -

Azure Password: XXXXXXXXXXXX -

The following information is for technical support purposes only:

Lab Instance: 12345678 -

Task -

You need to get insights from a video file located in the C:\Resources\Video\Media.mp4 folder.

Save the insights to the C:\Resources\Video\Insights.json folder.

To complete this task, sign in to the Azure Video Analyzer for Media at <https://www.videoindexer.ai/> by using admin@abc.com

Correct Answer: See explanation below.

Step 1: Login -

Browse to the Azure Video Indexer website and sign in.

URL: <https://www.videoindexer.ai/>

Login admin@abc.com -

Step 2: Create a project from your video

You can create a new project directly from a video in your account.

1. Go to the Library tab of the Azure Video Indexer website.

2. Open the video that you want to use to create your project. On the insights and timeline page, select the Video editor button.

Folder: C:\Resources\Video\Media.mp4

This takes you to the same page that you used to create a new project. Unlike the new project, you see the timestamped insights segments of the video, that you had started editing previously.

Step 3: Save the insights to the C:\Resources\Video\Insights.json folder.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-video-indexer/use-editor-create-project>

✉  **nanaw770** 3 days, 18 hours ago

Simulation questions will not appear on the actual exam as of May 25, 2024; ET should remove this type of question.

upvoted 2 times

✉  **joesatriani** 3 weeks, 6 days ago

Will these simulation questions appear in exams on 2024 as of May?

upvoted 3 times

✉  **rdemontis** 6 months, 2 weeks ago

the answer seems quite correct. After opened the video you can go on Insights and Timeline tab on the right and click Download --> Insights (Json)

upvoted 1 times

Question #42

SIMULATION -

Use the following login credentials as needed:

To enter your username, place your cursor in the Sign in box and click on the username below.

To enter your password, place your cursor in the Enter password box and click on the password below.

Azure Username: admin@abc.com -

Azure Password: XXXXXXXXXXXX -

The following information is for technical support purposes only:

Lab Instance: 12345678 -

Task -

You plan to analyze stock photography and automatically generate captions for the images.

You need to create a service in Azure to analyze the images. The service must be named caption12345678 and must be in the East US Azure region. The solution must use the Free pricing tier.

In the C:\Resources\Caption\Params.json folder, enter the value for Key 1 and the endpoint for the new service.

To complete this task, sign in to the Azure portal.

Correct Answer: / ¼ See explanation below.

Step 1: Provision a Cognitive Services resource

If you don't already have one in your subscription, you'll need to provision a Cognitive Services resource.

1. Open the Azure portal at <https://portal.azure.com>, and sign in using the Microsoft account associated with your Azure subscription.

2. Select the Create a resource button, search for cognitive services, and create a Cognitive Services resource with the following settings:

Subscription: Your Azure subscription

Resource group: Choose or create a resource group (if you are using a restricted subscription, you may not have permission to create a new resource group - use the one provided)

Region: East US Azure region -

Name: caption12345678 -

Pricing tier: Free F0 -

3. Select the required checkboxes and create the resource.

Wait for deployment to complete, and then view the deployment details.

4. When the resource has been deployed, go to it and view its Keys and Endpoint page. You will need the endpoint and one of the keys from this page in the next procedure.

Step 2: Save Key and Endpoint values in Params.json

Open the configuration file, C:\Resources\Caption\Params.json. and update the configuration values it contains to reflect the endpoint and an authentication key for your cognitive services resource. Save your changes.

Reference:

<https://microsoftlearning.github.io/AI-102-AIEngineer/Instructions/15-computer-vision.html>

✉ **nanaw770** 3 days, 18 hours ago

Simulation questions will not appear on the actual exam as of May 25, 2024; ET should remove this type of question.

upvoted 1 times

✉ **joesatriani** 3 weeks, 6 days ago

Will these simulation questions appear in exams on 2024 as of May?

upvoted 1 times

✉ **marti_tremblay000** 1 year, 2 months ago

create a Computer Vision service

upvoted 3 times

✉ **rde montis** 6 months, 4 weeks ago

exactly

upvoted 2 times

Question #43

SIMULATION -

Use the following login credentials as needed:

To enter your username, place your cursor in the Sign in box and click on the username below.

To enter your password, place your cursor in the Enter password box and click on the password below.

Azure Username: admin@abc.com -

Azure Password: XXXXXXXXXXXX -

The following information is for technical support purposes only:

Lab Instance: 12345678 -

Task -

You plan to build an application that will use caption12345678. The application will be deployed to a virtual network named VNet1.

You need to ensure that only virtual machines on VNet1 can access caption12345678.

To complete this task, sign in to the Azure portal.

Correct Answer: See explanation below.

Step 1: Create private endpoint for your web app

1. In the left-hand menu, select All Resources > caption12345678 - the name of your web app.
2. In the web app overview, select Settings > Networking.
3. In Networking, select Private endpoints.
4. Select + Add in the Private Endpoint connections page.
5. Enter or select the following information in the Add Private Endpoint page:

Name: Enter caption12345678.

Subscription Select your Azure subscription.

Virtual network Select VNet1.

Subnet: -

Integrate with private DNS zone: Select Yes.

6. Select OK.

Add Private Endpoint



Name *

mywebappendpoint

Subscription *

contoso subscription

Virtual network *

myVNet

Subnet *

mySubnet

If you have a network security group (NSG) enabled for the subnet above, it will be disabled for private endpoints on this subnet only. Other resources on the subnet will still have NSG enforcement.

Integrate with private DNS zone

No Yes

Your private endpoint will be integrated with the private DNS zone 'privatelink.azurewebsites.net' in the resource group of the selected subnet. If the private DNS zone does not exist, it will be created automatically. [Learn more](#)

Reference:

<https://docs.microsoft.com/en-us/azure/private-link/tutorial-private-endpoint-webapp-portal>

✉ **nanaw770** 3 days, 18 hours ago

Simulation questions will not appear on the actual exam as of May 25, 2024; ET should remove this type of question.

upvoted 2 times

✉ **michaelmorar** 1 month ago

Networking > Firewalls and virtual networks > Selected networks and private endpoints > Add existing virtual network

upvoted 1 times

✉ **AnonymousJhb** 4 months ago

This is a Custom vision resource > go to networking > you only need to configure the vnet1 only access.

- there is no mention to force traffic over the backbone only and deny internet bound traffic > which would require you then to go another step and enable private endpoints. (unless this question is incomplete, so practice both steps for fun)

upvoted 2 times

✉ **AnonymousJhb** 6 months, 2 weeks ago

portal > web app > networking > access restriction >

#1 Inbound Traffic > private endpoints > add > express > name + subs + vnet + subnet + private DNS > OK,

#2 Inbound Traffic > access restriction > uncheck Allow public access = uncheck > Save,

#3 Verify > Inbound Traffic > Access Restriction ON > Private Endpoint > ON

upvoted 1 times

✉ **AnonymousJhb** 4 months ago

ignore this. there is no delete option :(

upvoted 1 times

✉ **bull13** 10 months, 1 week ago

Select the resource -> Resource Management -> Networking -> Selected Networks and Private Endpoints

upvoted 1 times

✉ **dalones213** 1 year, 3 months ago

Create a virtual network by selecting Firewall and virtual networks

upvoted 1 times

✉ **halfway** 1 year, 5 months ago

Private endpoint is not the answer. Configure the service to allow access from selected network instead: <https://learn.microsoft.com/en-us/azure/cognitive-services/cognitive-services-virtual-networks?tabs=portal#managing-virtual-network-rules>

upvoted 3 times

✉ **rdemontis** 6 months, 4 weeks ago

agree with you

upvoted 1 times

Question #44

SIMULATION -

Use the following login credentials as needed:

To enter your username, place your cursor in the Sign in box and click on the username below.

To enter your password, place your cursor in the Enter password box and click on the password below.

Azure Username: admin@abc.com -

Azure Password: XXXXXXXXXXXX -

The following information is for technical support purposes only:

Lab Instance: 12345678 -

Task -

You need to ensure that a user named admin@abc.com can regenerate the subscription keys of AAA12345678. The solution must use the principle of least privilege.

To complete this task, sign in to the Azure portal.

Correct Answer: See explanation below.

Manually rotate subscription keys

1. (Update your application code to reference the secondary key for the Azure account and deploy.)
2. In the Azure portal, navigate to your Azure account.
3. Under Settings, select Authentication.
4. To regenerate the primary key for your Azure account, select the Regenerate button next to the primary key.
5. (Update your application code to reference the new primary key and deploy.)
6. Regenerate the secondary key in the same manner.

Reference:

<https://github.com/MicrosoftDocs/azure-docs/blob/main/articles/azure-maps/how-to-manage-authentication.md>

✉  **ziggy1117**  11 months, 3 weeks ago

Cognitive Services Contributor

Lets you create, read, update, delete and manage keys of Cognitive Services.

upvoted 10 times

✉  **rdmontis** 6 months, 2 weeks ago

agree with you

upvoted 1 times

✉  **nanaw770**  3 days, 18 hours ago

Simulation questions will not appear on the actual exam as of May 25, 2024; ET should remove this type of question.

upvoted 1 times

✉  **joesatriani** 3 weeks, 6 days ago

Will these simulation questions appear in exams on 2024 as of May?

upvoted 1 times

✉  **M25** 9 months ago

<https://learn.microsoft.com/en-us/azure/role-based-access-control/role-assignments-portal>

<https://learn.microsoft.com/en-us/azure/role-based-access-control/built-in-roles#cognitive-services-user>

Classic Storage Account Key Operator Service Role:

Classic Storage Account Key Operators are allowed to list and regenerate keys on Classic Storage Accounts

upvoted 2 times

✉  **mgafar** 1 year, 1 month ago

1. Sign in to the Azure portal (<https://portal.azure.com/>) using your account credentials.
2. In the left-hand navigation menu, click on "All services" and search for "Subscriptions." Click on the "Subscriptions" service to open the list of your Azure subscriptions.
3. Find the subscription with the ID "AAA12345678" and click on it to open the subscription details page.
4. In the left-hand navigation menu of the subscription details page, click on "Access control (IAM)."
5. Click on the "+ Add" button to add a new role assignment. This will open the "Add role assignment" pane.
6. In the "Role" dropdown menu, search for and select the "User Access Administrator" role. This role allows a user to manage access to Azure resources, including the ability to manage subscription keys, while adhering to the principle of least privilege.

7. In the "Select" field, type "admin@abc.com" and select the user from the list of suggestions.

8. Click on the "Save" button to complete the role assignment process.

upvoted 2 times

 **AnonymousJhb** 4 months ago

NEVER! this question asks for principle of least privilege.

you do not apply the RBAC IAM role at the subscription level.

you need to drill down past the resource group and directly on the resource, go to IAM > and add the user account as a Cognitive Services Contributor

upvoted 1 times

 **ziggy1117** 11 months, 3 weeks ago

user access administrator does not allow you to regenerate keys. but it allows you to add users

upvoted 1 times

 **Jo_Hannes** 10 months ago

Cognitive Services Contributor

Lets you create, read, update, delete and manage keys of Cognitive Services.

upvoted 1 times

Question #45

You have an Azure IoT hub that receives sensor data from machinery.

You need to build an app that will perform the following actions:

- Perform anomaly detection across multiple correlated sensors.
- Identify the root cause of process stops.
- Send incident alerts.

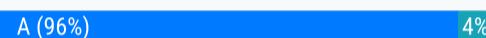
The solution must minimize development time.

Which Azure service should you use?

- A. Azure Metrics Advisor
- B. Form Recognizer
- C. Azure Machine Learning
- D. Anomaly Detector

Correct Answer: D

Community vote distribution



✉ **Pixelmate** Highly Voted 11 months ago

This came in the exam 28/06/2023

upvoted 10 times

✉ **zellck** Highly Voted 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/applied-ai-services/metrics-advisor/overview>

Metrics Advisor is a part of Azure Applied AI Services that uses AI to perform data monitoring and anomaly detection in time series data. The service automates the process of applying models to your data, and provides a set of APIs and a web-based workspace for data ingestion, anomaly detection, and diagnostics - without needing to know machine learning. Developers can build AIOps, predictive maintenance, and business monitor applications on top of the service.

upvoted 9 times

✉ **nanaw770** Most Recent 3 days, 18 hours ago

Selected Answer: A

A is right. From Takedajuku perspective, if you study for 4 days and spend 2 days reviewing, you will have a better chance of passing the exam.

upvoted 1 times

✉ **rewanotora** 1 week, 1 day ago

Why is A?

upvoted 2 times

✉ **Alok1105** 3 weeks, 2 days ago

Selected Answer: C

Given the requirements to perform anomaly detection across multiple correlated sensors, identify the root cause of process stops, and send incident alerts while minimizing development time, Azure Machine Learning seems to be the most appropriate choice.

So, I'm confident in option: C. Azure Machine Learning.

upvoted 1 times

✉ **AzureGC** 1 month ago

Selected Answer: A

Metrics Advisor is that includes "root cause",

upvoted 1 times

✉ **evangelist** 3 months, 2 weeks ago

Selected Answer: A

The answer must be A since metrics advisor is specialized in multivariate analysis which is required by the question

upvoted 2 times

✉ **evangelist** 3 months, 2 weeks ago

Selected Answer: A

note the type of detection should be multivariate

upvoted 1 times

✉ **josebernabeo** 4 months, 1 week ago

"Starting on the 20th of September, 2023 you won't be able to create new Metrics Advisor resources. The Metrics Advisor service is being retired on the 1st of October, 2026."

source: <https://learn.microsoft.com/en-us/azure/ai-services/metrics-advisor/overview>

upvoted 5 times

✉ **rdemontis** 6 months, 4 weeks ago

Selected Answer: A

correct answer is A <https://learn.microsoft.com/en-us/azure/ai-services/metrics-advisor/overview>

upvoted 1 times

✉ **james2033** 9 months, 2 weeks ago

Selected Answer: A

Quote

"Use Metrics Advisor to:

- Analyze multi-dimensional data from multiple data sources
- Identify and correlate anomalies
- Configure and fine-tune the anomaly detection model used on your data
- Diagnose anomalies and help with root cause analysis"

at <https://learn.microsoft.com/en-us/azure/ai-services/metrics-advisor/overview>

upvoted 4 times

✉ **M25** 8 months, 3 weeks ago

<https://azure.microsoft.com/en-us/products/ai-services/ai-metrics-advisor>

What's the difference between AI Metrics Advisor and AI Anomaly Detector?

AI Anomaly Detector is comprised of simple REST APIs with a code-first experience. It's the core engine of AI Metrics Advisor that detects anomalies in time-series data. It's best applied for ad-hoc data analysis, and it can be run in containers.

AI Metrics Advisor has additional time-series monitoring features, with pipeline APIs and a built-in user interface for managing the service. It's designed for live-streaming data and AI analytics, and it supports deploying in Azure.

upvoted 3 times

✉ **EliteAllen** 11 months, 2 weeks ago

Selected Answer: A

A. Azure Metrics Advisor

Azure Metrics Advisor is a service that provides an end-to-end anomaly detection platform, which includes data ingestion, anomaly detection, root cause analysis, and alerting. It is designed to monitor and detect anomalies in time-series data, diagnose incidents, and provide insights.

upvoted 1 times

✉ **ziggy1117** 11 months, 3 weeks ago

Selected Answer: A

Azure Metrics Advisor

- collect time series data
- detect anomalies
- send incident alerts
- analyze root cause

upvoted 1 times

✉ **RemcoGoy** 1 year ago

Selected Answer: A

Answer should indeed be A, I think.

<https://learn.microsoft.com/en-us/azure/applied-ai-services/metrics-advisor/overview> as mentioned in other comments has both 'Anomaly Detection' and 'Incident reporting'

upvoted 3 times

✉ **MaliSanFuu** 1 year, 1 month ago

Selected Answer: A

i agree with @Mike19D

minimal effort and root cause analysis:

<https://learn.microsoft.com/en-us/azure/applied-ai-services/metrics-advisor/overview>

upvoted 2 times

✉ **Mike19D** 1 year, 1 month ago

The Answer is A
upvoted 1 times

Question #46

You have an app that analyzes images by using the Computer Vision API.

You need to configure the app to provide an output for users who are vision impaired. The solution must provide the output in complete sentences.

Which API call should you perform?

- A. readInStreamAsync
- B. analyzeImagesByDomainInStreamAsync
- C. tagImageInStreamAsync
- D. describeImageInStreamAsync

Correct Answer: D

Community vote distribution

D (100%)

 **RAN_L**  1 year, 2 months ago

Selected Answer: D

The API call you should perform to provide an output in complete sentences for users who are vision impaired is describeImageInStreamAsync.

The describe feature of the Computer Vision API generates a human-readable sentence to describe the contents of an image. This is particularly useful for accessibility purposes, as it allows visually impaired users to understand what is in an image without needing to see it. The describe feature can also be customized to provide additional details or context, if desired.

Therefore, the correct answer is D. describeImageInStreamAsync.

upvoted 15 times

 **nanaw770**  3 days, 18 hours ago

Selected Answer: D

D. describeImageInStreamAsync is right answer.

upvoted 1 times

 **sivapolam90** 1 month, 2 weeks ago

Selected Answer: D

Answer is D. describeImageInStreamAsync.

upvoted 1 times

 **evangelist** 3 months, 2 weeks ago

Selected Answer: D

The correct answer is D. describeImageInStreamAsync.

To configure an app that provides output in complete sentences for users who are vision impaired, using the Computer Vision API, you should use the describeImageInStreamAsync call. This API call analyzes an image and provides a description in natural, human-readable language.

upvoted 1 times

 **rdemontis** 6 months, 4 weeks ago

Selected Answer: D

To provide output in complete sentences for users who are vision-impaired, you should perform the following API call:

D. describeImageInStreamAsync

The describeImageInStreamAsync API call is specifically designed to provide a description of the image content in a textual format that can be read aloud to users with visual impairments. This API call will generate a description of the image, making it accessible to those who rely on screen readers or other assistive technologies. (ChatGPT)

upvoted 2 times

 **trashbox** 7 months ago

Appeared on Oct/29/2023.

upvoted 3 times

 **james2033** 9 months, 2 weeks ago

Selected Answer: D

See sample source code

```
using (Stream imageStream = File.OpenRead(imagePath)) {
    ImageDescription descriptions = await computerVision.DescribeImageInStreamAsync(imageStream);
    Console.WriteLine(imagePath);
    DisplayDescriptions(descriptions);
}
```

at <https://github.com/Azure-Samples/cognitive-services-dotnet-sdk-samples/blob/master/samples/ComputerVision/DescribeImage/Program.cs#L75C20-L75C20>
upvoted 2 times

 **zellck** 11 months ago

Selected Answer: D

D is the answer.

<https://learn.microsoft.com/en-us/rest/api/computervision/3.1/describe-image-in-stream/describe-image-in-stream?tabs=HTTP>
This operation generates a description of an image in human readable language with complete sentences. The description is based on a collection of content tags, which are also returned by the operation. More than one description can be generated for each image. Descriptions are ordered by their confidence score. Descriptions may include results from celebrity and landmark domain models, if applicable. Two input methods are supported -- (1) Uploading an image or (2) specifying an image URL. A successful response will be returned in JSON. If the request failed, the response will contain an error code and a message to help understand what went wrong.

upvoted 3 times

Question #47

DRAG DROP

You have a Custom Vision service project that performs object detection. The project uses the General domain for classification and contains a trained model.

You need to export the model for use on a network that is disconnected from the internet.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

- Change the classification type.
- Export the model.
- Retrain the model.
- Change Domains to **General (compact)**.
- Create a new classification model.

Answer Area

Answer Area	
Change Domains to General (compact) .	
Retrain the model.	
Export the model.	

✉ **jimbojambo** 1 year, 2 months ago

The provided answer is correct. As reported here
<https://learn.microsoft.com/en-us/azure/cognitive-services/Custom-Vision-Service/export-your-model>
the model must be retrained after changing the domain to compact.

upvoted 16 times

✉ **rdemontis** 6 months, 4 weeks ago

correct answer and tested!
upvoted 2 times

✉ **zellck** 11 months ago

1. Change Domains to General (compact)
2. Retain model
3. Export model

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/export-your-model>

- In the Domains section, select one of the compact domains. Select Save Changes to save the changes.
- From the top of the page, select Train to retrain using the new domain.
- Go to the Performance tab and select Export.

upvoted 9 times

✉ **nanaw770** 3 days, 18 hours ago

1. Change domain to General(Compact) domain
2. Retrain model using new domain
3. Export model to desired export format

upvoted 1 times

✉ **funny_penguin** 3 days, 23 hours ago

on exam today 24/05/2024, zellck's response I selected as well. general compact, retrain, export
upvoted 1 times

✉ **aobakwemc** 2 months, 2 weeks ago

Appeared in exam 03/11/2024

upvoted 3 times

✉ **audlindr** 2 months, 3 weeks ago

This was in 2-Mar-24 exam

upvoted 2 times

✉ **Prodyna** 6 months, 2 weeks ago

was on november exam

upvoted 5 times

✉ **propanther** 9 months, 2 weeks ago

Correct sequence of answer is:

1. Change domain to General(Compact) domain
2. Retrain model using new domain
3. Export model to desired export format

Ref: <https://learn.microsoft.com/en-us/azure/ai-services/custom-vision-service/export-your-model>

upvoted 2 times

✉ **Pixelmate** 11 months ago

This appeared in the exam 28/06

upvoted 3 times

✉ **RAN_L** 1 year, 2 months ago

This sequence of actions is not correct. Changing the domain to General (compact) before retraining the model may result in reduced accuracy and performance, as the model would not be optimized for the specific domain it is intended to be used in.

Therefore, the correct sequence of actions should start with retraining the model to optimize it for the intended use case, followed by changing the domain to General (compact) to create a more compact version of the model, and then exporting it for use on a disconnected network.

So the correct sequence is:

Retrain the model

Change Domains to General (compact)

Export the model

upvoted 1 times

✉ **Mehe323** 2 months, 3 weeks ago

In Jimbo's link, the retraining is the last sub step in the 'Convert to a compact domain' step. The step before that is Changing domains. So the answer given by ExamTopics is correct.

upvoted 1 times

✉ **Rob77** 1 year ago

Unlikely.

1 - the model is already trained and 2 - after changing the domain you have to retrain the model see jimbo's link above...

upvoted 2 times

Question #48

You are building an AI solution that will use Sentiment Analysis results from surveys to calculate bonuses for customer service staff.

You need to ensure that the solution meets the Microsoft responsible AI principles.

What should you do?

- A. Add a human review and approval step before making decisions that affect the staff's financial situation.
- B. Include the Sentiment Analysis results when surveys return a low confidence score.
- C. Use all the surveys, including surveys by customers who requested that their account be deleted and their data be removed.
- D. Publish the raw survey data to a central location and provide the staff with access to the location.

Correct Answer: A

Community vote distribution

A (100%)

 **zellck** Highly Voted 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 7 times

 **nanaw770** Most Recent 3 days, 18 hours ago

Selected Answer: A

A is good answer. From Takedajuku perspective, if you study for 4 days and spend 2 days reviewing, you will have a better chance of passing the exam.

upvoted 1 times

 **reiwanotora** 1 week, 1 day ago

Selected Answer: A

A is right. This question is redundant.

upvoted 1 times

 **reiwanotora** 1 week, 1 day ago

Selected Answer: A

A is right.

upvoted 1 times

 **rdemontis** 6 months, 4 weeks ago

Selected Answer: A

To ensure that the AI solution meets the Microsoft responsible AI principles, you should:

- A. Add a human review and approval step before making decisions that affect the staff's financial situation.

This option aligns with the responsible AI principle of fairness and accountability. By adding a human review and approval step, you ensure that the decisions affecting staff bonuses are reviewed by humans who can consider factors beyond just the sentiment analysis results. It adds an element of transparency, accountability, and fairness to the process, reducing the risk of biased or unfair decisions. (ChatGPT)

upvoted 4 times

 **demonite** 1 week, 4 days ago

It adds an element of transparency, accountability - nonsense! Not unless the results are published for the staff to see. So D is the closest to an answer here.

fairness to the process, reducing the risk of biased or unfair decisions - nonsense again, a person, sitting there and reviewing answers could be even more biased than the AI.

<https://www.microsoft.com/en-in/ai/responsible-ai>

upvoted 1 times

 **kiro_kocha** 8 months, 4 weeks ago

Selected Answer: A

I also think is A

upvoted 2 times

 **973b658** 11 months, 3 weeks ago

Selected Answer: A

It is A.

upvoted 1 times

Question #49

You have an Azure subscription that contains a Language service resource named ta1 and a virtual network named vnet1.

You need to ensure that only resources in vnet1 can access ta1.

What should you configure?

- A. a network security group (NSG) for vnet1
- B. Azure Firewall for vnet1
- C. the virtual network settings for ta1
- D. a Language service container for ta1

Correct Answer: C

Community vote distribution

C (100%)

 **Rob77**  1 year ago

Correct

<https://learn.microsoft.com/en-us/azure/cognitive-services/cognitive-services-virtual-networks?tabs=portal>

upvoted 10 times

 **trashbox**  7 months ago

Appeared on Oct/29/2023.

upvoted 5 times

 **nanaw770**  3 days, 18 hours ago

Selected Answer: C

We can use virtual network.

upvoted 1 times

 **reowanotora** 1 week, 1 day ago

Selected Answer: C

C is right.

upvoted 1 times

 **sivapolam90** 1 month, 2 weeks ago

Selected Answer: C

C. the virtual network settings for ta1

upvoted 2 times

 **evangelist** 3 months, 2 weeks ago

Selected Answer: C

configure vent1 as "trustworthy network" for Language service Ta1, so no resource except from trustworthy network Vnet1 can access ta1.

upvoted 1 times

 **rdemontis** 6 months, 4 weeks ago

Selected Answer: C

correct answer

upvoted 1 times

 **zellck** 11 months ago

Selected Answer: C

C is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/cognitive-services-virtual-networks?tabs=portal#grant-access-from-a-virtual-network>
You can configure Cognitive Services resources to allow access only from specific subnets. The allowed subnets may belong to a VNet in the same subscription, or in a different subscription, including subscriptions belonging to a different Azure Active Directory tenant.

upvoted 4 times

Question #50

You are developing a monitoring system that will analyze engine sensor data, such as rotation speed, angle, temperature, and pressure. The system must generate an alert in response to atypical values.

What should you include in the solution?

- A. Application Insights in Azure Monitor
- B. metric alerts in Azure Monitor
- C. Multivariate Anomaly Detection
- D. Univariate Anomaly Detection

Correct Answer: C

Community vote distribution

C (88%) 12%

✉ **Pixelmate** Highly Voted 11 months ago

This appeared in exam 28/06
upvoted 9 times

✉ **nanaw770** Most Recent 3 days, 18 hours ago

Selected Answer: C
C is right answer.
upvoted 1 times

✉ **reiwanotora** 1 week, 1 day ago

Selected Answer: C
C is right.
upvoted 1 times

✉ **evangelist** 3 months, 1 week ago

Please do not choose A, this is a machine learning topic and has nothing to do with Azure Monitor
upvoted 2 times

✉ **evangelist** 3 months, 2 weeks ago

Selected Answer: C
if there is no metrics advisor, then choose Multivariate Anomaly detection as secondary option, Metrics Advisor is the best answer.
upvoted 4 times

✉ **evangelist** 3 months, 2 weeks ago

Selected Answer: C
C. Multivariate Anomaly Detection is the most comprehensive solution. It allows for the analysis of complex relationships between multiple metrics which is crucial for accurately identifying anomalies in a system as intricate as an engine. This approach can help detect situations where the anomaly is not in the individual metrics but in their unexpected patterns or combinations.
upvoted 2 times

✉ **rdemontis** 6 months, 4 weeks ago

Selected Answer: C
<https://learn.microsoft.com/en-us/azure/cognitive-services/anomaly-detector/overview#multivariate-anomaly-detection>
upvoted 2 times

✉ **zellck** 11 months ago

Selected Answer: C
C is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/anomaly-detector/overview#multivariate-anomaly-detection>
The Multivariate Anomaly Detection APIs further enable developers by easily integrating advanced AI for detecting anomalies from groups of metrics, without the need for machine learning knowledge or labeled data. Dependencies and inter-correlations between up to 300 different signals are now automatically counted as key factors. This new capability helps you to proactively protect your complex systems such as software applications, servers, factory machines, spacecraft, or even your business, from failures.
upvoted 4 times

✉ **Hisayuki** 11 months, 2 weeks ago

Selected Answer: C

Multivariate Anomaly Detection - If your goal is to detect system level anomalies from a group of time series data, use multivariate anomaly detection APIs.

upvoted 4 times

 **endeesa** 11 months, 2 weeks ago

Selected Answer: A
Anomaly detection does not support alerting though

upvoted 3 times

 **mVic** 1 year ago

Selected Answer: C
The Multivariate Anomaly Detection APIs further enable developers by easily integrating advanced AI for detecting anomalies from groups of metrics, without the need for machine learning knowledge or labeled data.

<https://learn.microsoft.com/en-us/azure/cognitive-services/anomaly-detector/overview#multivariate-anomaly-detection>

upvoted 4 times

Question #51

You have an app named App1 that uses an Azure Cognitive Services model to identify anomalies in a time series data stream.

You need to run App1 in a location that has limited connectivity. The solution must minimize costs.

What should you use to host the model?

- A. Azure Kubernetes Service (AKS)
- B. Azure Container Instances
- C. a Kubernetes cluster hosted in an Azure Stack Hub integrated system
- D. the Docker Engine

Correct Answer: C

Community vote distribution



✉️ **Rob77** 1 year ago

Docker - <https://learn.microsoft.com/en-us/azure/cognitive-services/cognitive-services-container-support>
upvoted 14 times

✉️ **rdemontis** 6 months, 3 weeks ago

thanks for the provided link
upvoted 3 times

✉️ **zellck** 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/cognitive-services-container-support>
upvoted 6 times

✉️ **Ody** 2 months, 1 week ago

I believe your link indicates that D is the answer:

You must satisfy the following prerequisites before using Azure AI containers:

Docker Engine: You must have Docker Engine installed locally. Docker provides packages that configure the Docker environment on macOS, Linux, and Windows. On Windows, Docker must be configured to support Linux containers. Docker containers can also be deployed directly to Azure Kubernetes Service or Azure Container Instances.

upvoted 2 times

✉️ **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 5 times

✉️ **nanaw770** 3 days, 18 hours ago

What is your selected answer? B or D? Docker is justice?

upvoted 1 times

✉️ **PeteColag** 9 hours, 43 minutes ago

Selected Answer: D

Azure Kubernetes Service (AKS) and Azure Container Instances both require a consistent and reliable internet connection to Azure, which make them not be suitable for a location with limited connectivity. A Kubernetes cluster hosted in an Azure Stack Hub integrated system could work, but it might be overkill for your needs and potentially more expensive.

Excluding these leaves only D as the correct solution.

upvoted 1 times

✉️ **nanaw770** 3 days, 18 hours ago

I would like to know the basis on which Azure Container Instances is the correct answer.

upvoted 1 times

✉️ **vovap0vovap** 6 days, 17 hours ago

I might be stupid, but "data stream" and "limited connectivity" rendering out to me any cloud base staff. Which lead to Docker Engine

upvoted 3 times

✉ emiliocb4 2 weeks, 6 days ago

Selected Answer: D

Docker is the correct response

upvoted 1 times

✉ AzureGC 1 month ago

Selected Answer: D

D for Docker; While you can run Azure HCI with containers, that does not minimize costs. running a container instance locally on docker minimize costs

upvoted 2 times

✉ NullVoider_0 1 month, 3 weeks ago

Selected Answer: C

C. a Kubernetes cluster hosted in an Azure Stack Hub integrated system is suitable for running applications like App1 in environments with limited connectivity. Azure Stack Hub allows you to run Azure services on-premises, providing a consistent hybrid cloud platform. By hosting the model on a Kubernetes cluster in Azure Stack Hub, you can process data locally, minimizing latency and operating costs associated with data transfer to and from the cloud. This setup supports scenarios where connectivity to Azure cloud services is intermittent or not available, ensuring that the application can still function effectively.

upvoted 3 times

✉ evangelist 3 months, 2 weeks ago

Selected Answer: B

has to be the Azure Container instances

upvoted 1 times

✉ Tactable 3 months, 3 weeks ago

Selected Answer: B

Azure Container Service (ACS) was indeed retired, but it's important to note that Azure Container Instances (ACI) and Azure Container Service (ACS) are two different services.

So in this context - hosting a model in a location with limited connectivity while minimizing costs - Azure Container Instances would still be a valid choice.

upvoted 2 times

✉ suzanne_exam 4 months ago

ACS has been deprecated. Docker engine is lower cost and will deal with the limited connectivity

upvoted 1 times

✉ HS999999 4 months, 3 weeks ago

The question is asking where to host the model, so the local container should host the model and the container will be run by the docker engine

upvoted 2 times

✉ ankitdhir 5 months ago

Considering the limited connectivity scenario, an option to host the model might involve:

C. a Kubernetes cluster hosted in an Azure Stack Hub integrated system.

Azure Stack Hub is designed for scenarios where cloud services need to be extended to on-premises or edge locations. It provides consistent Azure services and tooling, allowing you to run Azure services locally. Hosting a Kubernetes cluster within an Azure Stack Hub integrated system could enable you to manage and deploy containerized applications in an environment with limited connectivity while still leveraging Azure technology and services.

upvoted 2 times

✉ dimsok 4 months, 3 weeks ago

That's not cheap though, a simple docker is cheaper

upvoted 2 times

✉ rdemontis 6 months, 3 weeks ago

Selected Answer: D

agree with you. thanks for the provided link. Here we need not only to minimize the costs but minimize the usage of the network due to an app hosted in a location with limited connectivity. So it's preferable to host the model in the same location/network where the app is. And to do that the solution is to containerize the model, and host locally using a docker engine

upvoted 3 times

✉ rdemontis 6 months, 1 week ago

<https://learn.microsoft.com/en-us/azure/ai-services/cognitive-services-container-support#prerequisites>

upvoted 1 times

✉ sl_msiconsulting 7 months, 3 weeks ago

Selected Answer: D

ACI is still hosted on Azure and you need to have reasonable internet connectivity to make your solution work. Instead you should run the container on premise using Docker or other means and specify the API key and the endpoint while launching the container instance for billing purpose.

upvoted 1 times

✉️ **sl_msiconsulting** 7 months, 3 weeks ago

You also need to minimize your cost and hence one more reason not to use ACI.
upvoted 2 times

✉️ **nitz14** 11 months, 2 weeks ago

Selected Answer: B
To host the model for your App1 in a location with limited connectivity while minimizing costs, the recommended option would be:

B. Azure Container Instances

Azure Container Instances (ACI) is a serverless container runtime that allows you to quickly and easily run containers without managing the underlying infrastructure. It provides a simple and cost-effective way to host your application components.

ACI is suitable for scenarios with limited connectivity because it eliminates the need to set up and manage a complex infrastructure, such as a Kubernetes cluster or Azure Stack Hub integrated system. With ACI, you can deploy and run individual containers directly, without any additional overhead.

upvoted 4 times

✉️ **ziggy1117** 11 months, 3 weeks ago

Selected Answer: B
To host the Azure Cognitive Services model for running App1 in a location with limited connectivity while minimizing costs, the recommended option would be:

B. Azure Container Instances

Explanation:

Azure Container Instances (ACI) allows you to run containers without managing the underlying infrastructure. ACI is a suitable choice for hosting the model in a location with limited connectivity as it provides a serverless compute environment that is lightweight and can be deployed quickly. ACI is designed for running individual containers and is well-suited for scenarios with low traffic or intermittent usage.

upvoted 1 times

Question #52

HOTSPOT

You have an Azure Cognitive Search resource named Search1 that is used by multiple apps.

You need to secure Search1. The solution must meet the following requirements:

- Prevent access to Search1 from the internet.
- Limit the access of each app to specific queries.

What should you do? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

To prevent access from the internet:

- Configure an IP firewall.
- Create a private endpoint.
- Use Azure roles.

To limit access to queries:

- Create a private endpoint.
- Use Azure roles.
- Use key authentication.

Answer Area

To prevent access from the internet:

- Configure an IP firewall.
- Create a private endpoint.
- Use Azure roles.

Correct Answer:

To limit access to queries:

- Create a private endpoint.
- Use Azure roles.
- Use key authentication.

 Mike19D Highly Voted 1 year, 1 month ago

Create Private Endpoint
Use Azure Roles
upvoted 29 times

 zellck Highly Voted 11 months ago

1. Create a private endpoint
2. Use Azure roles

<https://learn.microsoft.com/en-us/azure/search/service-create-private-endpoint#why-use-a-private-endpoint-for-secure-access>
Private Endpoints for Azure Cognitive Search allow a client on a virtual network to securely access data in a search index over a Private Link. The private endpoint uses an IP address from the virtual network address space for your search service. Network traffic between the client and the search service traverses over the virtual network and a private link on the Microsoft backbone network, eliminating exposure from the public internet.

<https://learn.microsoft.com/en-us/azure/search/search-security-rbac?tabs=config-svc-portal%2Croles-portal%2Ctest-portal%2Ccustom-role-portal%2Cdisable-keys-portal#grant-access-to-a-single-index>
In some scenarios, you may want to limit application's access to a single resource, such as an index.

The portal doesn't currently support role assignments at this level of granularity, but it can be done with PowerShell or the Azure CLI.
upvoted 21 times

 **zellick** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 8 times

 **rdemontis** 6 months, 3 weeks ago

thanks for explanation

upvoted 2 times

 **demonite** Most Recent 1 week ago

you can configure IP firewall and only allow the apps in. Implementing private endpoint requires the apps to be on the same vnet or a s2s vpn which adds complexity.

<https://learn.microsoft.com/en-us/azure/search/service-configure-firewall>

<https://learn.microsoft.com/en-us/azure/search/search-security-api-keys?tabs=rest-use%2Cportal-find%2Cportal-query>

IP Firewall

Use Key authentication.

upvoted 1 times

 **reiwanotora** 1 week, 1 day ago

Create a private endpoint & Use key authentication.

upvoted 1 times

 **michaelmorar** 1 month, 2 weeks ago

- Private endpoint
- Azure Roles

There is only one Cognitive Search instance; keys will not control access at the correct level of granularity.

upvoted 1 times

 **NullVoider_0** 1 month, 3 weeks ago

Create Private Endpoint

Use Key authentication.

Azure roles, specifically Azure Role-Based Access Control (RBAC), are designed to manage who has access to Azure resources and what they can do with those resources. While RBAC is effective for controlling access at the resource level (e.g., who can manage the search service, indexes, or data sources), it does not provide the granularity needed to limit access to specific queries within Azure Cognitive Search.

upvoted 1 times

 **varinder82** 2 months ago

Final Answer:

1. Create a private endpoint
2. Use Azure roles

upvoted 1 times

 **Ody** 2 months, 1 week ago

In this article, learn how to secure an Azure AI Search service so that it can't be accessed over a public internet connection:

<https://learn.microsoft.com/en-us/azure/search/service-create-private-endpoint>

upvoted 1 times

 **evangelist** 3 months, 2 weeks ago

prevent Azure services to be accessed from internet? ==>create private endpoint of that service

limit the service to a specific queries?-->using Azure roles

upvoted 1 times

 **sca88** 6 months ago

Should be

1. Create Private Endpoint
2. Use Azure Roles

upvoted 1 times

 **shahnawazkhot** 8 months, 2 weeks ago

Folks, the correct answer is

- Create a Private Endpoint for Preventing access from Internet.
- Use Azure Roles to limit access.

upvoted 1 times

✉ **Hisayuki** 11 months, 1 week ago

Configuring IP Firewall is correct answer

<https://learn.microsoft.com/en-us/azure/search/service-configure-firewall>

upvoted 1 times

✉ **examworld** 1 year ago

1. Private Endpoint

<https://learn.microsoft.com/en-us/azure/search/service-create-private-endpoint>

2. Key authentication

<https://learn.microsoft.com/en-us/azure/search/search-security-api-keys?tabs=portal-use%2Cportal-find%2Cportal-query>

upvoted 5 times

✉ **sca88** 6 months ago

I totally disagree with 2.

Your link says: "Use role assignments to restrict access to API keys."

Moreover:

<https://learn.microsoft.com/en-us/azure/search/search-security-rbac?tabs=config-svc-portal%2Croles-portal%2Ctest-portal%2Ccustom-role-portal%2Cdisable-keys-portal#grant-access-to-a-single-index>

upvoted 1 times

✉ **ulloo** 1 year ago

As per ChatGPT:

To limit access to queries in an Azure Cognitive Search resource, you can use Azure Active Directory (Azure AD) authentication and authorization

upvoted 1 times

✉ **Rob77** 1 year ago

Pretty sure

1 should be Firewall

2 - not sure, can't find anything reasonable regarding the limitations - there is query key management in the keys though (?)

upvoted 2 times

✉ **Mehe323** 2 months, 3 weeks ago

A firewall would work for sure, but with these questions you also need to look at the most simple solution with the same results, and that is a private endpoint. The key is not the right answer, with every service you create there is a key to access it, it doesn't do any limitations with regard to limited access to queries. Once you have access, you can query as much as you like within the subscription limitation. So you need something more granular.

upvoted 1 times

Question #53

You are building a solution that will detect anomalies in sensor data from the previous 24 hours.

You need to ensure that the solution scans the entire dataset, at the same time, for anomalies.

Which type of detection should you use?

- A. batch
- B. streaming
- C. change points

Correct Answer: A

Community vote distribution

A (100%)

 **zellck**  11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/anomaly-detector/overview#univariate-anomaly-detection>

Batch detection

Use your time series to detect any anomalies that might exist throughout your data. This operation generates a model using your entire time series data, with each point analyzed with the same model.

upvoted 12 times

 **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 7 times

 **rdemontis** 6 months, 3 weeks ago

thanks for explanation

upvoted 1 times

 **nanaw770**  3 days, 18 hours ago

Selected Answer: A

A is right answer. From Takedajuku perspective, if you study for 4 days and spend 2 days reviewing, you will have a better chance of passing the exam.

upvoted 1 times

 **evangelist** 3 months, 1 week ago

Selected Answer: A

because it needs to process the whole 24 hours data, it has to be a batch

upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

to me A is the correct answer. <https://learn.microsoft.com/en-us/azure/cognitive-services/anomaly-detector/overview>

upvoted 1 times

 **Tin_Tin** 11 months, 2 weeks ago

Selected Answer: A

A is correct. see <https://learn.microsoft.com/en-us/azure/cognitive-services/anomaly-detector/overview>

upvoted 1 times

 **973b658** 11 months, 2 weeks ago

Selected Answer: A

A.

>solution scans the entire dataset

upvoted 1 times

Question #54

DRAG DROP

You are building an app that will scan confidential documents and use the Language service to analyze the contents.

You provision an Azure Cognitive Services resource.

You need to ensure that the app can make requests to the Language service endpoint. The solution must ensure that confidential documents remain on-premises.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Run the container and specify an App ID and Client Secret.	
Provision an on-premises Kubernetes cluster that is isolated from the internet.	
Pull an image from the Microsoft Container Registry (MCR).	(Up)
Run the container and specify an API key and the Endpoint URL of the Cognitive Services resource.	(Down)
Provision an on-premises Kubernetes cluster that has internet connectivity.	
Pull an image from Docker Hub.	
Provision an Azure Kubernetes Service (AKS) resource.	

Correct Answer:

Answer Area

- Pull an image from the Microsoft Container Registry (MCR).
- Provision an on-premises Kubernetes cluster that is isolated from the internet.
- Run the container and specify an API key and the Endpoint URL of the Cognitive Services resource.

 **zellck**  11 months ago

1. Provision on-premise k8 cluster that is isolated from Internet
2. Pull image from MCR
3. Run container and specify API key and endpoint URL of Cognitive Services Services

<https://learn.microsoft.com/en-us/azure/cognitive-services/containers/disconnected-containers>

Containers enable you to run Cognitive Services APIs in your own environment, and are great for your specific security and data governance requirements. Disconnected containers enable you to use several of these APIs disconnected from the internet.

<https://learn.microsoft.com/en-us/azure/cognitive-services/containers/disconnected-container-faq#how-do-i-download-the-disconnected-containers>

These containers are hosted on the Microsoft Container Registry and available for download on Microsoft Artifact Registry and Docker Hub. You won't be able to run the container if your Azure subscription has not been approved after completion of the request form.

upvoted 16 times

 **Mehe323** 2 months ago

You need permission from Microsoft to run the containers isolated from the internet and there are some very specific requirements (see your first link). Since this procedure is not included in the case study, I would go for the solution that uses internet connection for the billing. With this solution, the documents also remain on-premises.

upvoted 7 times

 **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 7 times

 **rdemontis** 6 months, 3 weeks ago

agree with you. ChatGPT confirm. Thanks for explanation

upvoted 1 times

 **sca88** 6 months ago

Your first link explain how to use Container without internet, but doesn't talk about AI Services.

So I think the correct link is this:

<https://learn.microsoft.com/en-us/training/modules/investigate-container-for-use-with-ai-services/3-use-ai-services-container>

"<https://learn.microsoft.com/en-us/training/modules/investigate-container-for-use-with-ai-services/3-use-ai-services-container>".

So the correct answer will be:

1. Provision on-premise k8 cluster that is connected to Internet
2. Pull image from MCR
3. Run container and specify API key and endpoint URL of Cognitive Services Services

upvoted 16 times

✉️ **dimsok** Highly Voted 4 months, 2 weeks ago

It needs to have internet connectivity to send billing information, so you cannot select an isolated K8s service.

1. K8s with internet
2. Image from MCR
3. AppID and Secret

upvoted 8 times

✉️ **nanaw770** Most Recent 3 days, 18 hours ago

1. Provision on-premise k8 cluster that is isolated from Internet
2. Pull image from MCR
3. Run container and specify API key and endpoint URL of Cognitive Services Services

upvoted 1 times

✉️ **JamesKJoker** 1 week, 4 days ago

1. Provision on-premise k8 cluster that is connected to Internet
2. Pull image from MCR
3. Run container and specify API key and endpoint URL of Cognitive Services Services

Internet Connectivity is required for billing

upvoted 1 times

✉️ **NullVoider_0** 1 month, 3 weeks ago

1. Provision an on-premises Kubernetes cluster that is isolated from the internet.
2. Pull an image from the Microsoft Container Registry (MCR).
3. Run the container and specify an API key and the Endpoint URL of the Cognitive Services resource.

upvoted 1 times

✉️ **varinder82** 2 months ago

Final Answer:

1. Provision on-premise k8 cluster that is isolated from Internet
2. Pull image from MCR
3. Run container and specify API key and endpoint URL of Cognitive Services Services

upvoted 1 times

✉️ **Ody** 2 months, 1 week ago

The Azure AI containers are required to submit metering information for billing purposes. Failure to allowlist various network channels that the Azure AI containers rely on will prevent the container from working.

<https://learn.microsoft.com/en-us/azure/ai-services/cognitive-services-container-support>

upvoted 3 times

✉️ **audlindr** 2 months, 3 weeks ago

This was in 2-Mar-24 exam

upvoted 2 times

✉️ **cloudguy1975** 4 months ago

The Image needs to be able to send billing info to the Azure AI Service. Also the question does not say anything about no internet traffic. It only states that the docs should remain on-premise.

Build a K8s cluster onprem with internet acecss.

Pull an image

specify the API key and URL

upvoted 3 times

✉️ **sl_mslconsulting** 7 months, 3 weeks ago

Request access to use containers in disconnected environments might not get approved what then? A on-premise k8 cluster with internet connectivity is the way to go for billing purpose. The documents still remain on premise.

upvoted 4 times

✉️ **sl_mslconsulting** 7 months, 3 weeks ago

Also it's meaningless to specify API key and endpoint URL of Cognitive Services Services while running the container if you don't have internet connectivity.

upvoted 2 times

✉️ **abcd9999** 9 months, 2 weeks ago

if your cluster is isolated from Internet, how can you pull image from MCR?

upvoted 2 times

✉️ **mhauber** 9 months, 2 weeks ago

via the microsoft network

upvoted 1 times

✉️ **msdfqwerfewf** 11 months ago

1. Provision an on-premises Kubernetes cluster that is isolated from the internet.
2. Pull an image from MCR (Microsoft Container Registry) or your private container registry.
3. Deploy the app container on the on-premises Kubernetes cluster, specifying the API Key and Endpoint URL of the Cognitive Services resource.

FROM CHATGPT

upvoted 5 times

 **Hisayuki** 11 months, 2 weeks ago

- Provision an on-prem kubernetes cluster that is isolated from Internet
- Pull an image from MCR
- Run the container and specify an API Key and Endpoint URL of the Cognitive Services resource

upvoted 4 times

 **973b658** 11 months, 2 weeks ago

It is true.

upvoted 2 times

Question #55

HOTSPOT

You have an Azure subscription that has the following configurations:

- Subscription ID: 8d3591aa-96b8-4737-ad09-00f9b1ed35ad
- Tenant ID: 3edfe572-cb54-3ced-ae12-c5c177f39a12

You plan to create a resource that will perform sentiment analysis and optical character recognition (OCR).

You need to use an HTTP request to create the resource in the subscription. The solution must use a single key and endpoint.

How should you complete the request? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

The screenshot shows the Azure Management Portal URL: <https://management.azure.com/>. Below it is a dropdown menu with the following options:

- /resourceGroups/OCRProject/providers/
- subscriptions/3edfe572-cb54-3ced-ae12-c5c177f39a12
- subscriptions/8d3591aa-96b8-4737-ad09-00f9b1ed35ad
- tenant/3edfe572-cb54-3ced-ae12-c5c177f39a12
- tenant/8d3591aa-96b8-4737-ad09-00f9b1ed35ad

Below this is another dropdown menu with the following options:

- /accounts/CS1?api-version=2021-10-01
- Microsoft.ApiManagement
- Microsoft.CognitiveServices
- Microsoft.ContainerService
- Microsoft.KeyVault

Correct Answer:

The screenshot shows the Azure Management Portal URL: <https://management.azure.com/>. Below it is a dropdown menu with the following options:

- /resourceGroups/OCRProject/providers/
- subscriptions/3edfe572-cb54-3ced-ae12-c5c177f39a12
- subscriptions/8d3591aa-96b8-4737-ad09-00f9b1ed35ad**
- tenant/3edfe572-cb54-3ced-ae12-c5c177f39a12
- tenant/8d3591aa-96b8-4737-ad09-00f9b1ed35ad

Below this is another dropdown menu with the following options:

- /accounts/CS1?api-version=2021-10-01
- Microsoft.CognitiveServices**
- Microsoft.ContainerService
- Microsoft.KeyVault

zellck Highly Voted 11 months ago

1. subscriptions/8d3591aa-96b8-4737-ad09-00f9b1ed35ad
2. Microsoft.CognitiveServices

<https://learn.microsoft.com/en-us/azure/cognitive-services/cognitive-services-apis-create-account?tabs=multiservice%2Canomaly-detector%2Clanguage-service%2Ccomputer-vision%2Cwindows#types-of-cognitive-services-resources>

You can access Azure Cognitive Services through two different resources: A multi-service resource, or a single-service one.

- Multi-service resource:

Access multiple Azure Cognitive Services with a single key and endpoint.

Consolidates billing from the services you use.

upvoted 16 times

zellck 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 7 times

rdemontis Most Recent 6 months, 3 weeks ago

correct answer

upvoted 3 times

kiro_kocha 8 months, 4 weeks ago

The solution must use a single key and endpoint. - Cognitive Services

You need to use an HTTP request to create the resource in the subscription. - obviously subscription

upvoted 1 times

Tin_Tin 11 months, 2 weeks ago

The answer is correct. <https://learn.microsoft.com/en-us/rest/api/resources/resources/create-or-update>

upvoted 3 times

 **973b658** 11 months, 2 weeks ago

It is true.

upvoted 2 times

Question #56

You have an Azure subscription that contains an Anomaly Detector resource.

You deploy a Docker host server named Server1 to the on-premises network.

You need to host an instance of the Anomaly Detector service on Server1.

Which parameter should you include in the docker run command?

- A. Fluentd
- B. Billing
- C. Http Proxy
- D. Mounts

Correct Answer: B

Community vote distribution

B (82%)

D (18%)

✉ **Tin_Tin** 11 months, 2 weeks ago

Selected Answer: B

The answer is correct. Important

The Eula, Billing, and ApiKey options must be specified to run the container; otherwise, the container won't start. For more information, see Billing. The ApiKey value is the Key from the Keys and Endpoints page in the LUIS portal and is also available on the Azure Cognitive Services resource key page.

<https://learn.microsoft.com/en-us/azure/cognitive-services/luis/luis-container-configuration#example-docker-run-commands>

upvoted 15 times

✉ **[Removed]** 11 months, 2 weeks ago

You are correct. Moreover, there is no "mounts" command there is a "mount" command.

upvoted 5 times

✉ **trato** 1 month, 1 week ago

Selected Answer: B

```
docker run --rm -it -p 5000:5000 --memory 4g --cpus 1 \
mcr.microsoft.com/azure-cognitive-services/decision/anomaly-detector:latest \
Eula=accept \
Billing={ENDPOINT_URI} \
ApiKey={API_KEY}
```

<https://learn.microsoft.com/en-us/azure/ai-services/anomaly-detector/anomaly-detector-container-howto#run-the-container-with-docker-run>

upvoted 1 times

✉ **michaelmorar** 1 month, 2 weeks ago

Selected Answer: B

Always about the billing with Microsoft! :)

upvoted 1 times

✉ **evangelist** 3 months, 2 weeks ago

no billing no use

upvoted 1 times

✉ **rdemontis** 6 months, 3 weeks ago

Selected Answer: D

In my opinion we are being asked here for a solution to host an instance of the Anomaly Detector on prem service. So although all 4 are configurable parameters (see <https://learn.microsoft.com/en-us/azure/ai-services/luis/luis-container-configuration>) only two of them are mandatory and only one, from my point of view, is relevant to the context of the question. And that is Mounts. So for me the correct answer is D.

upvoted 1 times

✉ **rdemontis** 6 months, 2 weeks ago

Sorry I have to revise my previous answer: Although both parameters Mounts and Billing are required in our case, actually in the context of the question Billing makes more sense because compared to what is required theoretically no volume needs to be mounted to the container.

Whereas on the other hand, Billing data is always required when hosting an Azure Cognitive Services solution in a container locally

upvoted 2 times

zellck 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/anomaly-detector/anomaly-detector-container-howto#run-the-container-with-docker>
run

The Eula, Billing, and ApiKey options must be specified to run the container; otherwise, the container won't start.

upvoted 3 times

zellck 11 months ago

<https://learn.microsoft.com/en-us/azure/cognitive-services/anomaly-detector/anomaly-detector-container-howto#billing-arguments>
- Billing

The endpoint of the Cognitive Services resource that's used to track billing information.

The value of this option must be set to the endpoint URI of a provisioned Azure resource.

upvoted 2 times

zellck 11 months ago

<https://learn.microsoft.com/en-us/azure/cognitive-services/anomaly-detector/anomaly-detector-container-howto#billing>
The Anomaly Detector containers send billing information to Azure, using an Anomaly Detector resource on your Azure account.

Queries to the container are billed at the pricing tier of the Azure resource that's used for the ApiKey parameter.

Azure Cognitive Services containers aren't licensed to run without being connected to the metering or billing endpoint. You must enable the containers to communicate billing information with the billing endpoint at all times. Cognitive Services containers don't send customer data, such as the image or text that's being analyzed, to Microsoft.

upvoted 3 times

Hisayuki 11 months, 1 week ago

Selected Answer: B

example)

```
$ docker run --rm -it -p 5000:5000 --memory 4g --cpus 2 --mount type=bind,src=c:\demo\container,target=/input --mount type=bind,src=C:\demo\container,target=/output  
mcr.microsoft.com/azure-cognitive-services/luis Eula=accept Billing=https://westus.api.cognitive.microsoft.com/luis/v2.0 ApiKey=__YOUR_API_KEY__
```

upvoted 3 times

user361836261 11 months, 2 weeks ago

Answer B is correct

<https://learn.microsoft.com/en-us/azure/cognitive-services/luis/luis-container-configuration#configuration-settings>

upvoted 2 times

endeesa 11 months, 2 weeks ago

Selected Answer: D

Billing is not a docker run parameter, answer is D

upvoted 2 times

nitz14 11 months, 2 weeks ago

Selected Answer: D

To host an instance of the Anomaly Detector service on Server1 using Docker, you should include the following parameter in the docker run command:

D. Mounts

Mounts allow you to attach a directory or a file from the Docker host server (in this case, Server1) to the container. By including the appropriate mount configuration, you can provide the necessary files or directories required to run the Anomaly Detector service.

upvoted 1 times

973b658 11 months, 2 weeks ago

Selected Answer: D

It is D.

upvoted 1 times

Question #57

You are building an app that will use the Speech service.

You need to ensure that the app can authenticate to the service by using a Microsoft Azure Active Directory (Azure AD), part of Microsoft Entra, token.

Which two actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Enable a virtual network service endpoint.
- B. Configure a custom subdomain.
- C. Request an X.509 certificate.
- D. Create a private endpoint.
- E. Create a Conditional Access policy.

Correct Answer: CE*Community vote distribution*

✉️ **radzio123** Highly Voted 4 months ago

Selected Answer: BD

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/role-based-access-control#authentication-with-keys-and-tokens>
upvoted 6 times

✉️ **nanaw770** Most Recent 3 days, 18 hours ago

Selected Answer: BD

Wrong again. It's frustrating.
upvoted 1 times

✉️ **reiwanotora** 1 week, 1 day ago

Selected Answer: BD

B and D.
upvoted 1 times

✉️ **GirishGururani** 1 month, 2 weeks ago

correct answer : B and D

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/role-based-access-control#authentication-with-keys-and-tokens>

To authenticate with a Microsoft Entra token, the Speech resource must have a custom subdomain and use a private endpoint. The Speech service uses custom subdomains with private endpoints only.
upvoted 2 times

✉️ **Murtuza** 1 month, 3 weeks ago

Selected Answer: BD

B and D is the correct answer
upvoted 2 times

✉️ **Murtuza** 2 months, 1 week ago

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/role-based-access-control#authentication-with-keys-and-tokens>
upvoted 3 times

✉️ **arcameon** 2 months, 3 weeks ago

Selected Answer: BD

The correct actions, according to Microsoft's documentation, are:

- B. Configure a custom subdomain.
Configuring a custom subdomain involves setting up a custom domain for your Speech service instance, which can help in integrating with Azure AD and managing authentication effectively.
- D. Create a private endpoint.
Creating a private endpoint allows your app to securely connect to the Speech service without exposing it to the public internet, maintaining network security and compliance.

upvoted 4 times

✉ **Murtuza** 3 months ago

Selected Answer: BD

B and D are the best choices its obvious

upvoted 4 times

✉ **Murtuza** 3 months ago

To authenticate with a Microsoft Entra token, the Speech resource must have a custom subdomain and use a private endpoint. The Speech service uses custom subdomains with private endpoints only.

upvoted 3 times

✉ **evangelist** 3 months, 1 week ago

Selected Answer: BE

The answer may not be so explicit and obvious, please memorize if you dont understand:

B. Configure a custom subdomain - This is essential for enabling Azure AD authentication for Cognitive Services, as it allows the service to be associated with a specific Azure AD identity.

E. Create a Conditional Access policy - This enhances security by defining conditions under which users or services can authenticate and access resources, aligning with the requirement to use Azure AD tokens for authentication.

These steps are critical for setting up Azure AD authentication, ensuring secure and compliant access to the Azure Speech Service within your application.

upvoted 2 times

✉ **Tactable** 3 months, 3 weeks ago

Selected Answer: DE

According to ChatGPT:

While configuring a custom subdomain might be part of setting up Azure Active Directory for authentication purposes, it's less directly related to securing service-to-service authentication, which is the requirement of the scenario.

Therefore, considering the emphasis on service-to-service authentication and the need to ensure secure access to the Speech service, options D and E (Create a private endpoint and Create a Conditional Access policy) are more directly aligned with the requirements of the scenario compared to B and E.

In summary, while B and E might be relevant for certain authentication scenarios, D and E are better suited for securing service-to-service authentication and access control in the context of the given scenario.

upvoted 2 times

✉ **dimsok** 4 months, 2 weeks ago

Selected Answer: BD

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/role-based-access-control#authentication-with-keys-and-token>

upvoted 4 times

✉ **HS999999** 4 months, 3 weeks ago

the link: <https://learn.microsoft.com/en-us/azure/ai-services/speech-service/role-based-access-control#authentication-with-keys-and-tokens> says The Speech service uses custom subdomains with private endpoints only. so the answer should not be BD?

upvoted 3 times

✉ **arcticliviv** 4 months, 3 weeks ago

Selected Answer: BE

To authenticate with a Microsoft Entra token, the Speech resource must have a custom subdomain and use a private endpoint. The Speech service uses custom subdomains with private endpoints only.

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/role-based-access-control#authentication-with-keys-and-tokens>

upvoted 3 times

✉ **vechiatto** 4 months, 3 weeks ago

Selected Answer: BD

To authenticate with Speech resource keys, all you need is the key and region. To authenticate with a Microsoft Entra token, the Speech resource must have a custom subdomain and use a private endpoint. The Speech service uses custom subdomains with private endpoints only.

upvoted 3 times

✉ **sca88** 6 months ago

Selected Answer: BE

To configure Microsoft AD Entra you need to provide a custom domain (B) and you need to assign Roles

upvoted 3 times

✉ **rdemontis** 6 months, 3 weeks ago

Selected Answer: BE

as for the microsoft documentation below it seems correct answers are B and E

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/how-to-configure-azure-ad-auth?tabs=portal&pivots=programming-language-csharp#:~:text=For%20Azure%20AD%20authentication%20with,the%20Azure%20portal%20or%20PowerShell>

upvoted 3 times

Question #58

HOTSPOT

You plan to deploy an Azure OpenAI resource by using an Azure Resource Manager (ARM) template.

You need to ensure that the resource can respond to 600 requests per minute.

How should you complete the template? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
{  
    "type": "Microsoft.CognitiveServices/accounts/deployments",  
    "apiVersion": "2023-05-01",  
    "name": "arm-aoai-sample-resource/arm-je-std-deployment",  
    "dependsOn": [  
        "[resourceId('Microsoft.CognitiveServices/accounts', 'arm-aoai-sample-resource')]"  
    ],  
    "sku": {  
        "name": "Standard",  
        :  
            "capacity"  
            "count"  
            "maxValue"  
            "size"  
        :  
            1  
            60  
            100  
            600  
    },  
    "properties": {  
        "model": {  
            "format": "OpenAI",  
            ...  
        }  
    }  
}
```

Answer Area

```
{
    "type": "Microsoft.CognitiveServices/accounts/deployments",
    "apiVersion": "2023-05-01",
    "name": "arm-aoai-sample-resource/arm-je-std-deployment",
    "dependsOn": [
        "[resourceId('Microsoft.CognitiveServices/accounts', 'arm-aoai-sample-resource')]"
    ],
    "sku": {
        "name": "Standard",
        "capacity": 100
    },
    "properties": {
        "model": {
            "format": "OpenAI",
            ...
        }
    }
}
```

Correct Answer:

✉ **chandiochan** Highly Voted 2 months, 2 weeks ago

Answer is correct:

Azure OpenAI allows you to manage how frequently your application can make inferencing requests. Your rate limits are based on Tokens-per-Minute (TPM). For example, if you have a capacity of 1, this equals 1,000 TPM, and the rate limit of requests you can make per minute (RPM) is calculated using a ratio. For every 1,000 TPM, you can make 6 RPM.

So, if you need to process 600 requests every minute, you'll require a TPM that supports that many RPM. Using the ratio, for 600 RPM, you need 100,000 TPM (because 600 divided by 6 equals 100, and 100 multiplied by 1,000 equals 100,000). In this scenario, you would set the capacity to 100, since each capacity unit equals 1,000 TPM.

upvoted 11 times

✉ **audlindr** Highly Voted 2 months, 4 weeks ago

Answer is correct

From here <https://learn.microsoft.com/en-us/azure/ai-services/openai/how-to/quota?tabs=rest>

When a deployment is created, the assigned TPM will directly map to the tokens-per-minute rate limit enforced on its inferencing requests. A Requests-Per-Minute (RPM) rate limit will also be enforced whose value is set proportionally to the TPM assignment using the following ratio:

6 RPM per 1000 TPM.

From here:

<https://learn.microsoft.com/en-us/azure/ai-services/openai/how-to/quota?tabs=rest>

capacity integer This represents the amount of quota you are assigning to this deployment. A value of 1 equals 1,000 Tokens per Minute (TPM). A value of 10 equals 10k Tokens per Minute (TPM).

So the math is 600 RPM needs 100000 TPM which translates to capacity 100

upvoted 5 times

✉ **audlindr** 2 months, 3 weeks ago

This was in 2-Mar-24 exam

upvoted 2 times

✉ **nanaw770** Most Recent 3 days, 18 hours ago

capacity and 100.

upvoted 1 times

✉ **nanaw770** 3 days, 18 hours ago

capacity and 100.

From Takedajuku perspective, if you study for 4 days and spend 2 days reviewing, you will have a better chance of passing the exam.

upvoted 1 times

✉ **funny_penguin** 3 days, 23 hours ago

on exam today 24/05/2024, I selected capacity and 600

upvoted 1 times

✉ **NullVoider_0** 1 month, 3 weeks ago

The accurate capacity value to meet the requirement of processing 600 requests per minute is indeed 100. I apologize for any confusion caused earlier

upvoted 3 times

 **NullVoider_0** 1 month, 3 weeks ago

In the provided options, the correct choice is:

"capacity" : 600

The capacity parameter in the SKU configuration of an Azure OpenAI resource specifies the maximum number of requests per minute (RPM) that the resource can handle. By setting "capacity": 600, you are configuring the resource to handle up to 600 requests per minute.

upvoted 1 times

Question #59

DRAG DROP

You have an app that manages feedback.

You need to ensure that the app can detect negative comments by using the Sentiment Analysis API in Azure AI Language. The solution must ensure that the managed feedback remains on your company's internal network.

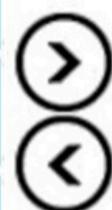
Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

NOTE: More than one order of answer choices is correct. You will receive credit for any of the correct orders you select.

Actions

- Identify the Language service endpoint URL and query the prediction endpoint.
- Provision the Language service resource in Azure.
- Run the container and query the prediction endpoint.
- Deploy a Docker container to an on-premises server.
- Deploy a Docker container to an Azure container instance.

Answer Area



Correct Answer:

Answer Area
Deploy a Docker container to an on-premises server.
Provision the Language service resource in Azure.
Run the container and query the prediction endpoint.

✉ **Jimmy1017** Highly Voted 1 month, 2 weeks ago

Provision the Language service resource in Azure.
Deploy a Docker container to an on-premises server.
Identify the Language service endpoint URL and query the prediction endpoint.

upvoted 7 times

✉ **Murtuza** Highly Voted 2 months ago

Remember that the order matters: provision the language service first, then Identify the Language Service Endpoint URL and Query the Prediction Endpoint
and finally deploy the container based on your chosen deployment target which is on-premises

upvoted 7 times

✉ **vovap0vovap** 12 hours, 39 minutes ago

Question stated that more than 1 correct order possible. You can provision resources first or deploy container first.
Now Identify the Language Service Endpoint URL and Query the Prediction Endpoint should not be correct as it indirectly assumed Azure Endpoint rather than local from Docker

upvoted 1 times

✉ **nanaw770** Most Recent 3 days, 18 hours ago

Deploy on-premises
Provision
Run

upvoted 1 times

✉ **Murtuza** 1 month, 3 weeks ago

Provision the Language service resource in Azure: This is the first step where you set up the Language service resource in Azure. This service will provide you with the Sentiment Analysis API.

Deploy a Docker container to an on-premises server: After provisioning the Language service, you should deploy a Docker container on an on-premises server. This container will host the Azure AI Language service and ensure that the managed feedback remains on your company's internal network.

Identify the Language service endpoint URL and query the prediction endpoint: Once the Docker container is running on your on-premises server, you can identify the Language service endpoint URL. You can then query the prediction endpoint to analyze the sentiment of the comments.

upvoted 5 times

 **NullVoider_0** 1 month, 3 weeks ago

1. Provision the Language service resource in Azure. This step involves creating the Azure Language service resource, which will provide you with the necessary credentials and endpoint URL to use the Sentiment Analysis API.
2. Deploy a Docker container to an on-premises server. By deploying a Docker container on-premises, you can run the Sentiment Analysis API locally, ensuring that the feedback data does not leave your internal network.
3. Run the container and query the prediction endpoint. Once the container is running on your on-premises server, you can start sending feedback data to the Sentiment Analysis API by querying the prediction endpoint provided by the Language service.

upvoted 4 times

 **Mehe323** 2 months, 3 weeks ago

That is correct, see prerequisites where you need to have 1) Docker installed and 2) provisioned a Language resource:
<https://learn.microsoft.com/en-us/azure/ai-services/language-service/sentiment-opinion-mining/how-to/use-containers>

upvoted 5 times

 **Ody** 2 months, 1 week ago

Right, but the question says Deploy a Docker Container. It is not saying setup a Docker host.

You can't run the container without having the API Key and Endpoint URI.

Provision the Language service (get the API Key and Endpoint URI)

Deploy a Docker container to an on-premise host

Run the container and query the prediction endpoint.

upvoted 6 times

 **Mehe323** 2 months ago

It says 'deploy a docker container to an on-premises server', you forgot the last part. A synonym for deploy is install.

upvoted 3 times

Question #60

HOTSPOT

You have an Azure OpenAI resource named AI1 that hosts three deployments of the GPT 3.5 model. Each deployment is optimized for a unique workload.

You plan to deploy three apps. Each app will access AI1 by using the REST API and will use the deployment that was optimized for the app's intended workload.

You need to provide each app with access to AI1 and the appropriate deployment. The solution must ensure that only the apps can access AI1.

What should you use to provide access to AI1, and what should each app use to connect to its appropriate deployment? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Provide access to AI1 by using:

- An API key
- A bearer token
- A shared access signature (SAS) token

Connect to the deployment by using:

- An API key
- A deployment endpoint
- A deployment name
- A deployment type

Answer Area

Provide access to AI1 by using:

- An API key
- A bearer token**
- A shared access signature (SAS) token

Correct Answer:

Connect to the deployment by using:

- An API key
- A deployment endpoint**
- A deployment name
- A deployment type

 **cacaflycloud** Highly Voted 2 months, 3 weeks ago

I think the answer should be
An API key
A deployment endpoint
upvoted 15 times

 **Harry300** Highly Voted 2 months, 3 weeks ago

API key
Deployment name (Different deployments can be configured on oai.azure.com)
upvoted 9 times

 **nanaw770** Most Recent 3 days, 18 hours ago

From Takedajuku perspective, if you study for 4 days and spend 2 days reviewing, you will have a better chance of passing the exam.
upvoted 1 times

 **funny_penguin** 3 days, 23 hours ago

on exam today, I selected Api key and deployment endpoint
upvoted 1 times

✉ **Jimmy1017** 1 month ago

An API key
A deployment endpoint
upvoted 3 times

✉ **wheebe** 1 month ago

chatgpt's answer was the bearer token and deployment name
upvoted 1 times

✉ **Jimmy1017** 1 month, 2 weeks ago

Provide access to All by using:

An API key
Connect to the deployment by using:

A deployment name
upvoted 1 times

✉ **NullVoider_0** 1 month, 3 weeks ago

To provide each app with access to the Azure OpenAI resource AI1 and ensure that only the apps can access it, you should use the following:

Provide access to AI1 by using:

An API key
Connect to the deployment by using:
A deployment endpoint

The API key is used for authentication, allowing the apps to access the Azure OpenAI resource. The deployment endpoint is a unique URL that each app will use to connect to its appropriate deployment, ensuring that the app uses the deployment optimized for its intended workload.

upvoted 3 times

✉ **Murtuza** 2 months ago

In summary:

Provide Access: Use an API key.

Connect to Deployment: Use the deployment endpoint along with the API key.

This approach ensures secure and controlled access to AI1 while allowing each app to connect to the appropriate deployment.

upvoted 2 times

✉ **Mehe323** 2 months ago

ChatGPT says a bearer token (hence why someone likely chose this as an answer) but it is incorrect. It should be an API key of course, very commonly used in Azure.

With regard to the second one, I think it should be endpoint. Deployment endpoint is usually referred to with an URL that you use to connect to it. Each deployment has its own endpoint URL, which contains the deployment name to make it unique. When you choose the name, the answer is only partially correct.

upvoted 1 times

✉ **varinder82** 2 months ago

Final Answer:
API Key
Deployment Name:
upvoted 5 times

✉ **chandiochan** 2 months, 2 weeks ago

API Key
Deployment Name:

```
curl $AZURE_OPENAI_ENDPOINT/openai/deployments/gpt-35-turbo-instruct/completions?api-version=2023-05-15 \
-H "Content-Type: application/json" \
-H "api-key: $AZURE_OPENAI_API_KEY" \
-d "{\"prompt\": \"Once upon a time\"}"
```

The format of your first line of the command with an example endpoint would appear as follows curl https://docs-test-001.openai.azure.com/openai/deployments/{YOUR-DEPLOYMENT_NAME_HERE}/completions?api-version=2023-05-15 \. If you encounter an error, double check to make sure that you don't have a doubling of the / at the separation between your endpoint and /openai/deployments.

upvoted 4 times

✉ **Mehe323** 2 months ago

But the deployment name is wrapped in an URL to make it unique, so endpoint would make more sense. You connect to an endpoint not to a name.

upvoted 2 times

✉ **vovap0vovap** 6 days, 15 hours ago

Yes. That sort of controversial. But if you're working with API - you will do one URL and 3 different deployment names from a code standpoint.
And that's likely what they mean.

upvoted 1 times

✉ **chandiochan** 2 months, 2 weeks ago

Provide access to AI1 by using:

An API key: API keys are used to authenticate applications to cloud services securely. They ensure that only authorized applications can access your AI1 resource.

Connect to the deployment by using:

A deployment endpoint: Each deployment will have a unique endpoint that the app can use to connect to it. This endpoint is specific to a deployment within the Azure OpenAI resource, which allows each app to use the deployment optimized for its intended workload.

upvoted 3 times

Question #61

Topic 1

You build a bot by using the Microsoft Bot Framework SDK.

You start the bot on a local computer.

You need to validate the functionality of the bot.

What should you do before you connect to the bot?

- A. Run the Bot Framework Emulator.
- B. Run the Bot Framework Composer.
- C. Register the bot with Azure Bot Service.
- D. Run Windows Terminal.

Correct Answer: A

Community vote distribution

A (100%)

✉ **Mehe323** Highly Voted 2 months, 3 weeks ago

The answer looks correct:

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-debug-emulator?view=azure-bot-service-4.0&tabs=csharp>

upvoted 7 times

✉ **nanaw770** Most Recent 3 days, 18 hours ago

Selected Answer: A

A is good answer.

upvoted 1 times

✉ **reowanotora** 1 week, 1 day ago

Selected Answer: A

A is right.

upvoted 1 times

✉ **NullVoider_0** 1 month, 3 weeks ago

Correct answer.

upvoted 1 times

✉ **Murtuza** 2 months ago

Selected Answer: A

Before connecting to the bot and validating its functionality, you should run the Bot Framework Emulator. The Bot Framework Emulator is a powerful tool that allows you to test and debug your bot locally

upvoted 1 times

✉ **Razvan_C** 2 months, 3 weeks ago

Answer seems to be correct

upvoted 2 times

Question #62

You have an Azure OpenAI model named AI1.

You are building a web app named App1 by using the Azure OpenAI SDK.

You need to configure App1 to connect to AI1.

What information must you provide?

- A. the endpoint, key, and model name
- B. the deployment name, key, and model name
- C. the deployment name, endpoint, and key
- D. the endpoint, key, and model type

Correct Answer: C

Community vote distribution

C (100%)

✉️  **anntv252** 3 weeks, 5 days ago

C. the deployment name, endpoint, and key .
Endpoint should include model name
upvoted 1 times

✉️  **michaelmorar** 1 month, 2 weeks ago

Selected Answer: C

Correct, remember that Model Name and Type aren't relevant, we only need the deployment details and key.
upvoted 1 times

✉️  **NullVoider_0** 1 month, 3 weeks ago

Selected Answer: C

Correct answer.
upvoted 2 times

✉️  **Mehe323** 2 months ago

Selected Answer: C

Correct.

When you access the model via the API you will need to refer to the deployment name rather than the underlying model name in API calls.
<https://learn.microsoft.com/en-us/azure/ai-services/openai/how-to/create-resource?pivots=web-portal>

upvoted 1 times

✉️  **Training** 1 month, 3 weeks ago

<https://learn.microsoft.com/en-us/azure/ai-services/openai/quickstart?tabs=command-line%2Cpython-new&pivots=rest-api>
upvoted 1 times

✉️  **Akber81** 2 months ago

Selected Answer: C

yes Azure OpenAI
has Endpoint (URL) , KEY , DEPLOYMENT NAME
upvoted 1 times

✉️  **GHill1982** 2 months, 3 weeks ago

Selected Answer: C

To connect to an Azure OpenAI model using the Azure OpenAI SDK, you need to provide:
The deployment name of the model that you want to use. This is the name that you assigned to the model when you deployed it.
The endpoint of your Azure OpenAI resource. This is the URL that you can find in the Overview section of your resource in the Azure portal or by using the Azure CLI.
The key of your Azure OpenAI resource. This is the API key that you can find in the Keys and Endpoint section of your resource in the Azure portal or by using the Azure CLI.
upvoted 4 times

✉️  **AlviraTony** 2 months ago

It is the name of the deployed model, that means it is "model name", not "deployment name"

upvoted 1 times

✉ **AlviraTony** 2 months ago

Ignore my comment

upvoted 1 times

✉ **Razvan_C** 2 months, 3 weeks ago

Selected Answer: C

Answer seems to be correct

upvoted 1 times

Question #63

Topic 1

You are building a solution in Azure that will use Azure Cognitive Service for Language to process sensitive customer data.

You need to ensure that only specific Azure processes can access the Language service. The solution must minimize administrative effort.

What should you include in the solution?

- A. IPsec rules
- B. Azure Application Gateway
- C. a virtual network gateway
- D. virtual network rules

Correct Answer: D

Community vote distribution

D (100%)

✉ **takaimomoGcup** 1 week ago

Selected Answer: D

I believe D is right answer.

upvoted 1 times

✉ **dragongoseki** 1 week, 1 day ago

Selected Answer: D

D is OK.

upvoted 1 times

✉ **michaelmorar** 1 month, 2 weeks ago

Selected Answer: D

This question is more relevant to AZ104 (Azure Administration) rather than AI, but anyway, I reckon the answer is D.

upvoted 1 times

✉ **Harry300** 2 months, 3 weeks ago

Selected Answer: D

Correct.

Source: <https://learn.microsoft.com/en-us/azure/ai-services/cognitive-services-virtual-networks?tabs=portal>

upvoted 4 times

Question #64

You plan to perform predictive maintenance.

You collect IoT sensor data from 100 industrial machines for a year. Each machine has 50 different sensors that generate data at one-minute intervals. In total, you have 5,000 time series datasets.

You need to identify unusual values in each time series to help predict machinery failures.

Which Azure service should you use?

- A. Azure AI Computer Vision
- B. Cognitive Search
- C. Azure AI Document Intelligence
- D. Azure AI Anomaly Detector

Correct Answer: D

Community vote distribution

D (100%)

✉  **nanaw770** 3 days, 18 hours ago

Selected Answer: D

D is OK.

upvoted 1 times

✉  **anntv252** 3 weeks, 5 days ago

Correct

D. predict machinery failures is using Anomaly detector in azure AI service

upvoted 2 times

✉  **sivapolam90** 1 month, 2 weeks ago

D. Azure AI Anomaly Detector

upvoted 1 times

✉  **GHill1982** 2 months, 3 weeks ago

Selected Answer: D

Azure AI Anomaly Detector is an AI service that enables you to monitor and detect anomalies in your time series data with little machine learning knowledge, either batch validation or real-time inference.

upvoted 2 times

✉  **Razvan_C** 2 months, 3 weeks ago

Selected Answer: D

Answer seems to be correct

upvoted 3 times

Question #65

HOTSPOT

You plan to deploy a containerized version of an Azure Cognitive Services service that will be used for sentiment analysis.

You configure https://contoso.cognitiveservices.azure.com as the endpoint URI for the service.

You need to run the container on an Azure virtual machine by using Docker.

How should you complete the command? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
docker run --rm -it -p 5000:5000 --memory 8g --cpus 1 \
```

```
http://contoso.blob.core.windows.net
https://contoso.cognitiveservices.azure.com
mcr.microsoft.com/azure-cognitive-services/textanalytics/keyphrase
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment
```

```
Eula=accept \
```

```
Billing=
```

```
http://contoso.blob.core.windows.net
https://contoso.cognitiveservices.azure.com
mcr.microsoft.com/azure-cognitive-services/textanalytics/keyphrase
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment
```

```
ApiKey=xxxxxxxxxxxxxxxxxxxxxx
```

Answer Area

```
docker run --rm -it -p 5000:5000 --memory 8g --cpus 1 \
```

```
http://contoso.blob.core.windows.net
https://contoso.cognitiveservices.azure.com
mcr.microsoft.com/azure-cognitive-services/textanalytics/keyphrase
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment
```

Correct Answer:

```
Eula=accept \
```

```
Billing=
```

```
http://contoso.blob.core.windows.net
https://contoso.cognitiveservices.azure.com
mcr.microsoft.com/azure-cognitive-services/textanalytics/keyphrase
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment
```

```
ApiKey=xxxxxxxxxxxxxxxxxxxxxx
```

 **carcasagon** Highly Voted 2 months, 3 weeks ago

Correct. <https://learn.microsoft.com/en-us/azure/ai-services/language-service/sentiment-opinion-mining/how-to/use-containers>
upvoted 5 times

 **dragongoseki** Most Recent 1 week, 1 day ago

I looked similar question before.
upvoted 1 times

 **michaelmorar** 1 month, 2 weeks ago

I agree with the answer. The first parameter indicates the service type, and the second is specifically named "billing" so should point to our service instance.
upvoted 1 times

Question #66

You are developing a system that will monitor temperature data from a data stream. The system must generate an alert in response to atypical values. The solution must minimize development effort.

What should you include in the solution?

- A. Multivariate Anomaly Detection
- B. Azure Stream Analytics
- C. metric alerts in Azure Monitor
- D. Univariate Anomaly Detection

Correct Answer: D

Community vote distribution



✉ **nanaw770** 3 days, 18 hours ago

Selected Answer: D

It must be D.

upvoted 1 times

✉ **AzureGC** 1 month ago

Selected Answer: D

D: Given this is an AI test, in general sway towards the AI service; Additionally, Stream Analytics gather the data, still need Azure Monitor to generate the Alert, which is two services; Univariate Anomaly Detection and emit the alert w/in the service, which meets the minimize effort

upvoted 1 times

✉ **Jimmy1017** 1 month ago

answer C

While options like Multivariate Anomaly Detection (option A) and Univariate Anomaly Detection (option D) offer more advanced anomaly detection techniques, implementing them would likely require more development effort, including building and training custom machine learning models.

upvoted 3 times

✉ **sivapolam90** 1 month, 2 weeks ago

Option B

upvoted 1 times

✉ **michaelmorar** 1 month, 2 weeks ago

Selected Answer: B

Only one data element needs monitoring, so univariate makes sense.

upvoted 1 times

✉ **TT924** 1 month, 2 weeks ago

Selected Answer: B

This is the similar example, I would vote for B.

Use case of Stream Analytics

Query: Alert to trigger a business workflow

Let's make our query more detailed. For every type of sensor, we want to monitor average temperature per 30-second window and display results only if the average temperature is above 100 degrees.

<https://learn.microsoft.com/en-us/azure/stream-analytics/stream-analytics-get-started-with-azure-stream-analytics-to-process-data-from-iot-devices>

upvoted 1 times

✉ **NullVoider_0** 1 month, 3 weeks ago

Selected Answer: D

D. Univariate Anomaly Detection

upvoted 1 times

✉ **kchinivar** 2 months ago

Option D

upvoted 1 times

✉ **Akber81** 2 months ago

according to GPT:

The most efficient solution to minimize development effort while monitoring temperature data from a data stream and generating alerts for atypical values would be:

C. Metric alerts in Azure Monitor.

Azure Monitor provides a straightforward way to set up metric alerts based on predefined or custom metrics. It allows you to monitor various metrics in real-time and trigger alerts based on specified conditions, such as threshold values. This solution requires minimal development effort compared to building custom anomaly detection models (options A and D) or setting up a streaming analytics solution (option B).

upvoted 3 times

ybermachine 2 months, 1 week ago

Selected Answer: D

according to gpt-4 the answer is D

upvoted 2 times

Murtuza 2 months, 1 week ago

Selected Answer: D

C is the best choice and explained in my previous responses

upvoted 2 times

Murtuza 2 months, 1 week ago

Metric Alerts in Azure Monitor: Azure Monitor provides metric-based alerts. You can define thresholds for temperature metrics (such as average temperature over a time window) and trigger alerts when values exceed or fall below those thresholds. This solution minimizes development effort and integrates seamlessly with Azure services.

Considering your requirement to minimize development effort, I recommend C. metric alerts in Azure Monitor. It provides a straightforward way to set up temperature-based alerts without complex modeling or extensive coding.

upvoted 1 times

pmllamas 2 months, 3 weeks ago

D is specifically build for that. USE IT

upvoted 2 times

GHill1982 2 months, 3 weeks ago

Selected Answer: B

Azure Stream Analytics is an easy-to-use, real-time analytics service that offers built-in machine learning based anomaly detection capabilities.

upvoted 1 times

SaSilar 2 months, 3 weeks ago

It's B. Azure Stream Analytics

upvoted 3 times

Mehe323 2 months, 3 weeks ago

Selected Answer: C

Considering the fact that the solution must minimize development effort, I would say the answer is C.

upvoted 3 times

Question #67

You have a Microsoft OneDrive folder that contains a 20-GB video file named File1.avi.

You need to index File1.avi by using the Azure Video Indexer website.

What should you do?

- A. Upload File1.avi to the www.youtube.com webpage, and then copy the URL of the video to the Azure AI Video Indexer website.
- B. Download File1.avi to a local computer, and then upload the file to the Azure AI Video Indexer website.
- C. From OneDrive, create a download link, and then copy the link to the Azure AI Video Indexer website.
- D. From OneDrive, create a sharing link for File1.avi, and then copy the link to the Azure AI Video Indexer website.

Correct Answer: C

Community vote distribution



✉️ **audlindr** Highly Voted 2 months, 3 weeks ago

Selected Answer: C

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/odrv-download>

Copy the embed code and extract only the URL part including the key. For example:

https://onedrive.live.com/embed?cid=5BC591B7C713B04F&resid=5DC518B6B713C40F%2110126&authkey=HnsodidN_50oA3ILfk

Replace embed with download. You will now have a url that looks like this:

https://onedrive.live.com/download?cid=5BC591B7C713B04F&resid=5DC518B6B713C40F%2110126&authkey=HnsodidN_50oA3ILfk

Now enter this URL in the Azure AI Video Indexer website in the URL field.

upvoted 8 times

✉️ **nanaw770** Most Recent 3 days, 18 hours ago

Selected Answer: C

C is right answer.

upvoted 1 times

✉️ **demonite** 1 week, 1 day ago

- A. could actually work, haven't tested
- B. Although plausible, 20GB is a lot to download and then upload
- C. Is the answer here, let Azure do the work
- D. Sharing it means someone or something else can view. Video indecer needs to download the video.

upvoted 1 times

✉️ **reiwanotora** 1 week, 1 day ago

Selected Answer: C

C is right.

upvoted 1 times

✉️ **anntv252** 3 weeks, 5 days ago

Selected Answer: C

Share ondrive link into Azure AI Video Indexer website

upvoted 1 times

✉️ **NullVoider_0** 1 month, 3 weeks ago

Selected Answer: D

D. From OneDrive, create a sharing link for File1.avi, and then copy the link to the Azure AI Video Indexer website.

This approach allows the Azure Video Indexer to access and index the video file directly from OneDrive without the need to download and re-upload the file, making the process more efficient and straightforward.

upvoted 1 times

✉️ **NullVoider_0** 1 month, 2 weeks ago

My previous response was wrong. I have tested and the given answer is correct. Apologies for the error from my end.

upvoted 4 times

 **SaSilar** 2 months, 3 weeks ago

D. From OneDrive, create a sharing link for File1.avi, and then copy the link to the Azure Video Indexer website.

Explanation:

Creating a sharing link for File1.avi in OneDrive allows you to generate a URL that can be accessed by other services or users.

Once you have the sharing link, you can copy it and paste it into the Azure Video Indexer website.

Azure Video Indexer will then use the provided link to access and process the video file for indexing.

upvoted 1 times

 **GHill1982** 2 months, 3 weeks ago

Azure AI Video Indexer requires a download link, not a sharing link, to access the file.

upvoted 2 times

 **Harry300** 2 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 3 times

Question #68

You have an Azure subscription that contains an Azure AI Service resource named CSAccount1 and a virtual network named VNet1. CSAccount1 is connected to VNet1.

You need to ensure that only specific resources can access CSAccount1. The solution must meet the following requirements:

- Prevent external access to CSAccount1.
- Minimize administrative effort.

Which two actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct answer is worth one point.

- A. In VNet1, enable a service endpoint for CSAccount1.
- B. In CSAccount1, configure the Access control (IAM) settings.
- C. In VNet1, modify the virtual network settings.
- D. In VNet1, create a virtual subnet.
- E. In CSAccount1, modify the virtual network settings.

Correct Answer: AD

Community vote distribution

AE (73%)

AB (27%)

 **chandiochan**  2 months, 2 weeks ago

Selected Answer: AE

A. In VNet1, enable a service endpoint for CSAccount1. This allows you to secure your Azure service resources to the virtual network.

E. In CSAccount1, modify the virtual network settings. This will allow you to configure CSAccount1 to accept connections only from the virtual network VNet1.

Enabling service endpoints and modifying the virtual network settings for the AI Service resource will limit access to the resources within VNet1, effectively fulfilling both requirements.

upvoted 6 times

 **nanaw770**  3 days, 18 hours ago

Selected Answer: AE

A and E.

upvoted 1 times

 **reiwanotora** 1 week, 1 day ago

Selected Answer: AE

I will also vote for AE.

upvoted 1 times

 **AzureGC** 1 month ago

Selected Answer: AE

AE ...

Do NOT think B is correct: The IAM controls do not necessarily help if the endpoint and key are compromised; Have to use VNET controls to gate the service endpoints;

upvoted 1 times

 **Jimmy1017** 1 month ago

- A. In VNet1, enable a service endpoint for CSAccount1.
- B. In CSAccount1, configure the Access control (IAM) settings.

Explanation:

A. Enabling a service endpoint for CSAccount1 in VNet1 allows traffic from the virtual network to reach CSAccount1 without traversing the public internet, thus preventing external access.

B. Configuring the Access control (IAM) settings in CSAccount1 allows you to specify which specific resources or identities have access to

CSAccount1. By configuring these settings, you can ensure that only specific resources can access CSAccount1, meeting the requirement to restrict access.

upvoted 1 times

✉️ **franceshuang** 1 month, 2 weeks ago

AE should be right

upvoted 1 times

✉️ **TT924** 1 month, 2 weeks ago

Selected Answer: AB

A. Enable a service endpoint for Azure AI services within the virtual network. The service endpoint routes traffic from the virtual network through an optimal path to the Azure AI service.

<https://learn.microsoft.com/en-us/azure/ai-services/openai/how-to/role-based-access-control>

B. Has default multiple Role-base access: Cognitive Services OpenAI User, Cognitive Services OpenAI Contributor, Cognitive Services Contributor, Cognitive Services Usages Reader

You can also set up Azure RBAC for whole resource groups, subscriptions, or management groups. Do this by selecting the desired scope level and then navigating to the desired item. For example, selecting Resource groups and then navigating to a specific resource group.

Select Access control (IAM) on the left navigation pane.

<https://learn.microsoft.com/en-us/azure/ai-services/openai/how-to/role-based-access-control>

upvoted 2 times

✉️ **TT924** 1 month, 2 weeks ago

A , should be <https://learn.microsoft.com/en-us/azure/search/service-create-private-endpoint>

upvoted 1 times

✉️ **NullVoider_0** 1 month, 3 weeks ago

Selected Answer: AE

A. In VNet1, enable a service endpoint for CSAccount1. Enabling a service endpoint for CSAccount1 in VNet1 will allow you to secure the Azure AI Service resource to a specific subset of networks. This means that only the applications requesting data over VNet1 will be able to access CSAccount1. It's a way to ensure that the resource is only accessible from within the virtual network.

E. In CSAccount1, modify the virtual network settings. By modifying the virtual network settings in CSAccount1, you can configure network rules that limit access to the resource. You would set the default network access rule to deny access to all networks, including the internet. Then, you can specify which virtual networks or subnets are allowed to access CSAccount1.

upvoted 3 times

✉️ **Training** 1 month, 3 weeks ago

Where does it mention that Azure AI services have support for Azure Vnet Service endpoints

upvoted 1 times

✉️ **Murtuza** 2 months, 1 week ago

Selected Answer: AB

the correct answers are A (enable a service endpoint for CSAccount1) and B (configure the Access control (IAM) settings). These actions provide a secure and efficient solution for restricting access to CSAccount1 while minimizing administrative overhead.

upvoted 2 times

✉️ **microsofter029** 2 months, 3 weeks ago

Selected Answer: AB

The other options are not suitable for achieving the desired security outcome in this scenario for the following reasons:

C. Modify the virtual network settings: This option alone does not directly restrict access to CSAccount1. While modifying certain settings like network security groups (NSGs) could potentially be used, it would require additional configuration and wouldn't address access control within the virtual network.

D. Create a virtual subnet: While creating a subnet within VNet1 could contribute to segmentation and potential security benefits, it's not sufficient on its own. You would still need to configure access control at the resource level (CSAccount1) and potentially implement additional security measures like NSGs within the subnet.

E. Modify the virtual network settings in CSAccount1: CSAccount1 is an Azure AI service resource, and modifying its virtual network settings wouldn't be applicable in this context. These resources are not directly managed like virtual networks and do not have settings for virtual network configuration.

upvoted 2 times

✉️ **Ghill1982** 2 months, 3 weeks ago

Selected Answer: AE

In VNet1, enable a service endpoint for CSAccount1. This will allow you to connect your virtual network to your Azure AI Service resource securely over the Azure backbone network.

In CSAccount1, modify the virtual network settings. This will allow you to configure virtual network rules that specify which subnets can access your Azure AI Service resource.

upvoted 4 times

Question #69

You are building an internet-based training solution. The solution requires that a user's camera and microphone remain enabled.

You need to monitor a video stream of the user and detect when the user asks an instructor a question. The solution must minimize development effort.

What should you include in the solution?

- A. speech-to-text in the Azure AI Speech service
- B. language detection in Azure AI Language Service
- C. the Face service in Azure AI Vision
- D. object detection in Azure AI Custom Vision

Correct Answer: A

Community vote distribution

A (100%)

 **reiwanotora** 1 week, 1 day ago

Selected Answer: A
user's camera and microphone remain enabled, so A is right.
upvoted 1 times

 **anntv252** 3 weeks, 5 days ago

Selected Answer: A
Because user's camera and microphone remain enabled. Azure AI Speech service is recommend for using
upvoted 1 times

 **Barry123456** 1 month, 1 week ago

It says video stream. It doesn't say the video stream has audio. I deal with video only streams all day. Don't assume.
upvoted 1 times

 **sivapolam90** 1 month, 2 weeks ago

Selected Answer: A
A. speech-to-text in the Azure AI Speech service
upvoted 1 times

 **Murtuza** 1 month, 2 weeks ago

Selected Answer: A
The best option for this scenario would be A. speech-to-text in the Azure AI Speech service.

This service can transcribe the user's spoken words into written text, which can then be analyzed to detect when a question is being asked. This would be more efficient and direct for detecting questions in a video stream, compared to the other options which focus on language detection, face recognition, and object detection. These other services might not be as effective for this specific use-case.

upvoted 2 times

 **NullVoider_0** 1 month, 3 weeks ago

Selected Answer: A
A. speech-to-text in the Azure AI Speech service

This service can transcribe the spoken words into text in real-time, which can then be analyzed to detect questions. It's an efficient way to monitor for specific verbal cues or keywords that indicate a question is being asked, without the need for extensive programming or manual review. This approach minimizes development effort while providing a robust solution for the requirement.

upvoted 1 times

 **Murtuza** 2 months, 1 week ago

The correct CHOICE is C. I made a silly typo but my explanations are right on point.
upvoted 2 times

 **Murtuza** 2 months, 1 week ago

The other options are not directly relevant to detecting user questions in a video stream:

Speech-to-text (Option A): Converts spoken language into text. While useful for transcribing audio, it doesn't directly address identifying user questions.

Language detection (Option B): Determines the language of text. It's not specifically designed for monitoring video streams or detecting question

Object detection (Option D): Identifies objects within images, but it's not suitable for detecting user interactions or questions.

Therefore, Option C (the Face service in Azure AI Vision) is the most appropriate choice for your scenario.

upvoted 2 times

Murtuza 2 months, 1 week ago

Selected Answer: A

Face Service (Azure AI Vision):

The Face service provides facial recognition capabilities, which can be used to identify when a user is facing the camera (e.g., looking at the instructor).

By analyzing facial features, expressions, and head movements, you can detect when a user is likely to be asking a question.

This approach minimizes development effort because it directly addresses the requirement of monitoring the video stream for user interactions.

upvoted 1 times

Mehe323 2 months ago

The user can talk, but it doesn't have to be a question. I think the focus should be on detecting whether something is a question or not and for that, you need speech to text first. Face doesn't make sense as identifying questions is not the purpose of that service: 'The Azure AI Face service provides AI algorithms that detect, recognize, and analyze human faces in images. Facial recognition software is important in many different scenarios, such as identification, touchless access control, and face blurring for privacy.'

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/overview-identity>

upvoted 3 times

chandiochan 2 months, 2 weeks ago

Selected Answer: A

speech-to-text in the Azure AI Speech service/

This service can transcribe spoken words into written text in real-time, allowing you to monitor the audio for specific triggers, like questions, which can then be further processed or flagged for response. This solution is efficient and requires minimal development effort for integrating audio streaming and speech recognition capabilities.

upvoted 2 times

AlviraTony 2 months ago

[ChatGPT]

A. Speech-to-text in the Azure AI Speech service.

Explanation:

Speech-to-text functionality can convert spoken words into text, allowing you to analyze the content of the speech.

By using speech-to-text, you can transcribe the user's spoken questions and then analyze the text to detect if a question is being asked to the instructor.

This option aligns with the requirement to monitor the user's speech in real-time without significant development effort.

upvoted 1 times

Topic 2 - Question Set 2**Question #1****Topic 2****HOTSPOT -**

You are developing an application that will use the Computer Vision client library. The application has the following code.

```
public async Task AnalyzeImage(ComputerVisionClient client, string localImage)
{
    List<VisualFeatureTypes> features = new List<VisualFeatureTypes>()
    {
        VisualFeatureTypes.Description,
        VisualFeatureTypes.Tags,
    };
    using (Stream imageStream = File.OpenRead(localImage))
    {
        try
        {
            ImageAnalysis results = await client.AnalyzeImageInStreamAsync(imageStream, features);

            foreach (var caption in results.Description.Captions)
            {
                Console.WriteLine($"{caption.Text} with confidence {caption.Confidence}");
            }

            foreach (var tag in results.Tags)
            {
                Console.WriteLine($"{tag.Name} {tag.Confidence}");
            }
        }
        catch (Exception ex)
        {
            Console.WriteLine(ex.Message);
        }
    }
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
-------------------	------------	-----------

The code will perform face recognition.

The code will list tags and their associated confidence.

The code will read a file from the local file system.

Answer Area

Statements	Yes	No
-------------------	------------	-----------

Correct Answer: The code will perform face recognition.

The code will list tags and their associated confidence.

The code will read a file from the local file system.

Box 1: No -

Box 2: Yes -

The ComputerVision.analyzeImageInStreamAsync operation extracts a rich set of visual features based on the image content.

Box 3: No -

Images will be read from a stream.

Reference:

<https://docs.microsoft.com/en-us/java/api/com.microsoft.azure.cognitiveservices.vision.computervision.computervision.analyzeimageinstreamasync>

✉️  **motu**  2 years, 11 months ago

Box 3 is Yes, the stream will be generated from a local image!

upvoted 75 times

✉️  **Adedoyin_Simeon**  2 years, 10 months ago

Box 3 should be Yes, a stream is only a pathway for data. and in this case the data actually comes from a local file. The correct answer would be No.

Yes, Yes.

upvoted 29 times

✉️  **anntv252**  3 weeks, 5 days ago

No Yes Yes

Image analysis

Return the category and confident score

Read local video streaming or streaming source

upvoted 2 times

✉️  **michaelmorar** 1 month, 2 weeks ago

The fact that the parameter is named 'local', might be misleading - it could be anything.

HOWEVER, I'm not a C# expert, but File.OpenRead tells us that the file is on a filesystem to which the device has access.

SO my vote is NYY

upvoted 1 times

✉️  **varinder82** 2 months ago

Final Answer:

N Y Y

upvoted 1 times

✉️  **rdemontis** 6 months, 3 weeks ago

To me correct answers are NYY. localImage parameter is a string, that's a file path

upvoted 3 times

✉️  **zellck** 11 months ago

NYY is the answer.

[https://learn.microsoft.com/en-us/java/api/com.microsoft.azure.cognitiveservices.vision.computervision.computervision?view=azure-java-legacy#com-microsoft-azure-cognitiveservices-vision-computervision-computervision-analyzeimageinstreamoptionalparameter\)](https://learn.microsoft.com/en-us/java/api/com.microsoft.azure.cognitiveservices.vision.computervision.computervision?view=azure-java-legacy#com-microsoft-azure-cognitiveservices-vision-computervision-computervision-analyzeimageinstreamasync(byte()-analyzeimageinstreamoptionalparameter))

This operation extracts a rich set of visual features based on the image content. Two input methods are supported (1) Uploading an image or (2) specifying an image URL.

upvoted 10 times

✉️  **Lion007** 7 months, 1 week ago

NYN

Box 3: No.... because the image (whether uploaded or a URL), the CODE will read it as a An image STEAM. Please continue reading the same URL you shared.

upvoted 1 times

✉️  **propanther** 7 months, 1 week ago

File.OpenRead is a System.IO method and reads file from the local storage path provided as a parameter hence last one is Y as it reads file from local storage

<https://learn.microsoft.com/en-us/dotnet/api/system.io.file.openread?view=net-7.0>

upvoted 4 times

✉️  **ninja** 1 year, 9 months ago

Box 1: No. The code generates description and tags. See line 3,4

Box 2: Yes. The code displays tag.Name and tag.Confidence

Box 3: Yes. File.OpenRead reads a local file. See <https://docs.microsoft.com/en-us/dotnet/api/system.io.file.openread?view=net-6.0>

upvoted 6 times

✉️  **Eltooth** 1 year, 10 months ago

No

Yes

Yes

upvoted 2 times

✉️  **mohamedba** 1 year, 11 months ago

Answer is NYY

upvoted 3 times

✉️  **gursimran_s** 2 years, 1 month ago

<https://docs.microsoft.com/en-us/java/api/com.microsoft.azure.cognitiveservices.vision.computervision.computervision.analyzeimageinstreamasync?view=azure-java-legacy>

upvoted 1 times

✉️  **sumanshu** 2 years, 4 months ago

Not A Face Recognition, But Analyze the Image (i.e. extract information from the images)

Console.write shows, it will return TAG and Confidence

File.Openread(locallImage) - Will read from local system

NO YES YES

upvoted 6 times

✉ **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 2 times

✉ **ovokpus** 2 years, 6 months ago

No, Yes, Yes.

upvoted 4 times

✉ **mikegsm** 2 years, 6 months ago

Seems to be: NO, YES, YES

upvoted 2 times

✉ **Duch003** 2 years, 11 months ago

I agree that C is technically correct, but they are still expecting us to answer No for that one, because, if we stick that much to the details, it will be read from stream, not as a picture directly. Tricky one, I do not like it as well.

upvoted 7 times

✉ **WillyMac** 2 years, 11 months ago

motu: I agree.

Box 3 should be Yes

upvoted 4 times

✉ **azurelearner666** 2 years, 11 months ago

Agree! File.OpenRead() can only accept a filepath so yes.

upvoted 4 times

You are developing a method that uses the Computer Vision client library. The method will perform optical character recognition (OCR) in images. The method has the following code.

```
public static async Task ReadFileUrl(ComputerVisionClient client, string urlFile)
{
    const int numberOfCharsInOperationId = 36;

    var txtHeaders = await client.ReadAsync(urlFile, language: "en");

    string opLocation = txtHeaders.OperationLocation;
    string operationId = opLocation.Substring(opLocation.Length -
numberOfCharsInOperationId);

    ReadOperationResult results;

    results = await client.GetReadResultAsync(Guid.Parse(operationId));

    var textUrlFileResults = results.AnalyzeResult.ReadResults;
    foreach (ReadResult page in textUrlFileResults)
    {
        foreach (Line line in page.Lines)
        {
            Console.WriteLine(line.Text);
        }
    }
}
```

During testing, you discover that the call to the GetReadResultAsync method occurs before the read operation is complete.

You need to prevent the GetReadResultAsync method from proceeding until the read operation is complete.

Which two actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Remove the Guid.Parse(operationId) parameter.
- B. Add code to verify the results.Status value.
- C. Add code to verify the status of the txtHeaders.Status value.
- D. Wrap the call to GetReadResultAsync within a loop that contains a delay.

Correct Answer: BD

Example code :

```
do
{
    results = await client.GetReadResultAsync(Guid.Parse(operationId));
}

while ((results.Status == OperationStatusCodes.Running ||
results.Status == OperationStatusCodes.NotStarted));
```

Reference:

<https://github.com/Azure-Samples/cognitive-services-quickstart-code/blob/master/dotnet/ComputerVision/ComputerVisionQuickstart.cs>

Community vote distribution

BD (100%)

✉  **sdokmak**  1 year, 11 months ago

Selected Answer: BD

as per link in solution

upvoted 13 times

✉  **sdokmak** 1 year, 11 months ago

and looking at what getReadAsync and getReadResultAsync methods return.

getReadResultAsync returns Observable<ReadOperationResult> object which contains as status() method.

upvoted 5 times

✉  **sdokmak** 1 year, 11 months ago

getReadAsync doesn't have status method. Answer is B and D

<https://docs.microsoft.com/en-us/dotnet/api/system.io.stream.readasync?view=net-6.0>

upvoted 4 times

✉  **Pixelmate**  11 months ago

this appeared in the exam 28/06
upvoted 6 times

✉  **nanaw770**  2 days, 20 hours ago

Selected Answer: BD

results.Status
GetReadResultAsync
upvoted 1 times

✉  **rdemontis** 6 months, 3 weeks ago

Selected Answer: BD

provided answer seems correct. The attached documentation demonstrate it
upvoted 3 times

✉  **PHD_CHENG** 1 year, 11 months ago

was on exam 7 Jun 2022
upvoted 3 times

✉  **SamedKia** 2 years, 1 month ago

C and D are the correct answers.
upvoted 3 times

✉  **ppo12** 1 year, 10 months ago

I don't think C is one of the answer based on <https://github.com/Azure-Samples/cognitive-services-quickstart-code/blob/master/dotnet/ComputerVision/ComputerVisionQuickstart.cs>.

It seems results.Status is part of the while condition, hence I agree with dokmak's B and D
upvoted 1 times

Question #3

HOTSPOT -

You have a Computer Vision resource named contoso1 that is hosted in the West US Azure region.

You need to use contoso1 to make a different size of a product photo by using the smart cropping feature.

How should you complete the API URL? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
curl -H "Ocp-Apim-Subscription-Key: xxx" /
-o "sample.png" -H "Content-Type: application/json" /
https://api.projectoxford.ai/vision/v3.1/
https://contoso1.cognitiveservices.azure.com/?width=100&height=100&smartCropping=true" /
https://westus.api.cognitive.microsoft.com/areaOfInterest
detect
generateThumbnail

-d "{\"url\":\"https://upload.litwareinc.org/litware/bicycle.jpg\"}"
```

Correct Answer:

Answer Area

```
curl -H "Ocp-Apim-Subscription-Key: xxx" /
-o "sample.png" -H "Content-Type: application/json" /
https://api.projectoxford.ai/vision/v3.1/
https://contoso1.cognitiveservices.azure.com/?width=100&height=100&smartCropping=true" /
https://westus.api.cognitive.microsoft.com/areaOfInterest
detect
generateThumbnail

-d "{\"url\":\"https://upload.litwareinc.org/litware/bicycle.jpg\"}"
```

Reference:

<https://westus.dev.cognitive.microsoft.com/docs/services/computer-vision-v3-2/operations/56f91f2e778daf14a499f21b>

<https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-generating-thumbnails#examples>

  **czmiel24**  2 years, 9 months ago

The second one should be generate Thumbnail imho.

upvoted 50 times

  **ziizai** 2 years, 8 months ago

yes, the question is exactly the sample here

<https://docs.microsoft.com/en-us/rest/api/computervision/3.1/generate-thumbnail/generate-thumbnail#examples>

upvoted 13 times

  **rdemontis** 6 months, 3 weeks ago

agree with you

upvoted 1 times

  **rdemontis** 6 months, 3 weeks ago

however, I don't understand how I do using the generic endpoint to meet the requirement, "You need to use contoso1 to make a different size of a product photo". I'm not so sure to use the generic endpoint.

upvoted 1 times

  **rdemontis** 6 months, 3 weeks ago

Clarified. If you go on the API documentation <https://westus.dev.cognitive.microsoft.com/docs/services/computer-vision-v3-2/operations/56f91f2e778daf14a499f20c/console> (here is the version 3.2 but it's the same) you can verify that both type of endpoint are supported. If you choose a custom endpoint named contoso1 you'll get the following url request:

<https://westus.dev.cognitive.microsoft.com/docs/services/computer-vision-v3-2/operations/56f91f2e778daf14a499f20c/console>.

So correct answers are

<https://contoso1.cognitiveservices.azure.com/>

generateThumbnail

upvoted 7 times

  **rdemontis** 6 months, 3 weeks ago

sorry i pasted twice the same url in the message before:
with the custom endpoint the url of the service is:
<https://contoso1.cognitiveservices.azure.com/vision/v3.2/generateThumbnail?width=300&height=200&smartCropping=true&model-version=latest>
upvoted 1 times

✉ **VulcanMXNY** Highly Voted 2 years, 7 months ago

Both answers are incorrect.

The correct answers are:

<https://contoso1.cognitiveservices.azure.com/>
AND
generateThumbnail

westus.dev.cognitive.microsoft.com wouldn't be a correct Computer Vision endpoint if the resource name is contoso1.

Also, per the documentation, areaOfInterest "returns a bounding box around the most important area of the image", it doesn't return a different size photo (<https://docs.microsoft.com/en-us/rest/api/computervision/3.1/get-area-of-interest>).

upvoted 42 times

✉ **MDawson** 1 year, 1 month ago

contoso1 is a Computer Vision resource, so you would not specify /vision in the URL. Therefore I think the correct answer must be westus.api.cognitive.microsoft.com

upvoted 2 times

✉ **AzureJobsTillRetire** 1 year, 3 months ago

I agree with both answers here. The example <https://westus.api.cognitive.microsoft.com> is just an example and it needs to be changed to use the source in real which is contoso1.

upvoted 1 times

✉ **dazdzadzadzaazd** 10 months, 1 week ago

Today (July 2023), both regional and resource endpoints are supported. So both are correct :

<https://contoso1.cognitiveservices.azure.com/>
AND
<https://westus.api.cognitive.microsoft.com>

Doc : <https://learn.microsoft.com/en-us/azure/ai-services/cognitive-services-custom-subdomains#what-if-an-sdk-asks-me-for-the-region-for-a-resource>

"Regional endpoints and custom subdomain names are both supported and can be used interchangeably."

I tested it by creating a custom vision resource and used it with both endpoints.

upvoted 6 times

✉ **Ody** 2 months, 1 week ago

Yes, I think the key thing is that the key is specified.

upvoted 1 times

✉ **ppo12** 1 year, 10 months ago

I agree with generateThumbnail, however first answer provided by ET should be correct <https://westus.api.cognitive.microsoft.com> as shown in <https://docs.microsoft.com/en-us/rest/api/computervision/3.1/generate-thumbnail/generate-thumbnail?tabs=HTTP#examples>

upvoted 6 times

✉ **varinder82** Most Recent 2 months ago

Final Answer:

1. <https://contoso1.cognitiveservices.azure.com/>
2. generateThumbnail

upvoted 1 times

✉ **Murtuza** 2 months, 1 week ago

API URL:

The base URL for the Analyze Image 4.0 API is typically:

<https://<region>.api.cognitive.microsoft.com/vision/v4.0/analyze>

Replace <region> with the appropriate Azure region (in this case, West US).

upvoted 1 times

✉ **ExamDev** 9 months ago

"You need to use contoso1 to make a different size of a product photo by using the smart cropping feature." -> You need to use contoso1 to make.... that is hosted in the west us....

upvoted 1 times

✉ **zellck** 11 months ago

1. <https://contoso1.cognitiveservices.azure.com/>
2. generateThumbnail

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-generating-thumbnails>

A thumbnail is a reduced-size representation of an image. Thumbnails are used to represent images and other data in a more economical, layout-friendly way. The Computer Vision API uses smart cropping to create intuitive image thumbnails that include the most important regions of an

image with priority given to any detected faces.

The Computer Vision thumbnail generation algorithm works as follows:

- Remove distracting elements from the image and identify the area of interest—the area of the image in which the main object(s) appears.
- Crop the image based on the identified area of interest.
- Change the aspect ratio to fit the target thumbnail dimensions.

upvoted 7 times

✉ **Pixelmate** 11 months ago

was on exam 28/06/2023

upvoted 4 times

✉ **ziggy1117** 11 months, 4 weeks ago

I simulated this in Azure Portal:

1. endpoint is <https://contoso1.cognitiveservices.azure.com/>
2. thumbnail

upvoted 2 times

✉ **Sachz88** 1 year, 1 month ago

<https://contoso1.cognitiveservices.azure.com/> is correct.

Context from ChatGPT:

westus.api.cognitive.microsoft.com is also a valid endpoint for the Cognitive Services APIs, including the Computer Vision API. However, it is important to note that this endpoint is deprecated and will be retired on October 31, 2024.

Therefore, it is recommended to use the newer endpoint format <https://<resource-name>.cognitiveservices.azure.com/> for any new development work. This endpoint format follows a more standard Azure resource URL pattern and is also more flexible in terms of geographic distribution and availability.

Hope it helps.

upvoted 4 times

✉ **NNU** 1 year, 3 months ago

The first is <https://contoso1.cognitiveservices.azure.com> the second is generateThumbnail

POST https://*.cognitiveservices.azure.com/vision/v3.2/generateThumbnail?width=100&height=100&smartCropping=true&model-version=latest
HTTP/1.1

Host: *.cognitiveservices.azure.com

Content-Type: application/json

{"url": "http://example.com/images/test.jpg"}

upvoted 1 times

✉ **KingChuang** 1 year, 4 months ago

on my exam (2023-01-16 Passed)

My Answer:

<https://westus.api.cognitive.microsoft.com>

But I think this is wrong. Because Question request use contoso1!

So correct answer is :

<https://contoso1.cognitiveservices.azure.com/>

upvoted 3 times

✉ **ap1234pa** 1 year, 4 months ago

Hello.. I have exam tomorrow. Can you suggest if ET questions were on exam?

upvoted 1 times

✉ **Eltooth** 1 year, 10 months ago

Westus

generateThumbnail

<https://docs.microsoft.com/en-gb/azure/cognitive-services/computer-vision/how-to/generate-thumbnail#call-the-generate-thumbnail-api>

```
curl -H "Ocp-Apim-Subscription-Key: <subscriptionKey>" -o <thumbnailFile> -H "Content-Type: application/json"
"https://westus.api.cognitive.microsoft.com/vision/v3.2/generateThumbnail?width=100&height=100&smartCropping=true" -d "
{"url": "https://upload.wikimedia.org/wikipedia/commons/thumb/5/56/Shorkie_Poo_Puppy.jpg/1280px-Shorkie_Poo_Puppy.jpg"}"
```

upvoted 3 times

✉ **RamonKaus** 1 year, 10 months ago

First one is contoso.cognitive services. Just checked my own script and cognitive services uses url name in the endpoint URI.

upvoted 1 times

✉ **RamonKaus** 1 year, 10 months ago

Second one is obv. generateThumbnail

upvoted 1 times

✉ **JDarshan** 1 year, 10 months ago

For first dropdown, 3rd option works with cognitive service key and computer vision key as well. whereas 2nd option works with only computer vision key. so answer 3rd works in both situation. therefore I'll go with <https://westus.api.cognitive.microsoft.com/vision/v3.1/generateThumbnail?width=500&height=500&smartCropping=True>

upvoted 1 times

✉ Eltooth 1 year, 10 months ago

Contoso1

GenerateThumbnail

upvoted 1 times

✉ Eltooth 1 year, 10 months ago

Correction: New resources created after July 1, 2019, will use custom subdomain names, therefore: westus and generateThumbnail are correct answers.

Exact copy here from MS docs:

```
curl -H "Ocp-Apim-Subscription-Key: <subscriptionKey>" -o <thumbnailFile> -H "Content-Type: application/json" "https://westus.api.cognitive.microsoft.com/vision/v3.2/generateThumbnail?width=100&height=100&smartCropping=true" -d "{\"url\":\"https://upload.wikimedia.org/wikipedia/commons/thumb/5/56/Shorkie_Poo_Puppy.jpg/1280px-Shorkie_Poo_Puppy.jpg\"}"
```

<https://docs.microsoft.com/en-gb/azure/cognitive-services/computer-vision/how-to/generate-thumbnail#call-the-generate-thumbnail-api>

upvoted 4 times

✉ satishk4u 2 years ago

Was on exam on 03-May-2022

upvoted 3 times

✉ 2ez4Zane 2 years, 1 month ago

the 2nd one should be generate thumbnail

Get Thumbnail

This operation generates a thumbnail image with the user-specified width and height. By default, the service analyzes the image, identifies the region of interest (ROI), and generates smart cropping coordinates based on the ROI

upvoted 1 times

Question #4

DRAG DROP -

You are developing a webpage that will use the Azure Video Analyzer for Media (previously Video Indexer) service to display videos of internal company meetings.

You embed the Player widget and the Cognitive Insights widget into the page.

You need to configure the widgets to meet the following requirements:

- Ensure that users can search for keywords.
- Display the names and faces of people in the video.
- Show captions in the video in English (United States).

How should you complete the URL for each widget? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Values	Answer Area
en-US	
false	
people,keywords	
people,search	
search	
true	

Cognitive Insights Widget

`https://www.videoindexer.ai/embed/insights/<accountId>/<videoId>/?widgets=` `controls=`

Player Widget

`https://www.videoindexer.ai/embed/player/<accountId>/<videoId>/? showcaptions=` `captions=`

Correct Answer:

Values	Answer Area
false	
people,search	

Cognitive Insights Widget

`https://www.videoindexer.ai/embed/insights/<accountId>/<videoId>/?widgets=` `controls=`

Player Widget

`https://www.videoindexer.ai/embed/player/<accountId>/<videoId>/? showcaptions=` `captions=`

Reference:

<https://docs.microsoft.com/en-us/azure/azure-video-analyzer/video-analyzer-for-media-docs/video-indexer-embed-widgets>

✉  **zellck**  11 months ago

1. people, keywords / search
2. true / en-US

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/video-indexer-embed-widgets#cognitive-insights-widget>

- widgets

Allows you to control the insights that you want to render.

- controls

Allows you to control the controls that you want to render.

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/video-indexer-embed-widgets#player-widget>

- showCaptions

Makes the player load with the captions already enabled.

- captions

Fetches the caption in the specified language during the widget loading to be available on the Captions menu

upvoted 14 times

✉  **rdemontis** 6 months, 3 weeks ago

thanks for explanation

upvoted 1 times

✉  **Eltooth**  1 year, 10 months ago

Answer is correct.

<https://docs.microsoft.com/en-us/azure/azure-video-indexer/video-indexer-embed-widgets>

upvoted 8 times

✉ **nanaw770** Most Recent 2 days, 20 hours ago

widgets=people,keywords
controls=search
showcaptions=true
captions=en-US
upvoted 1 times

✉ **takaimomoGcup** 1 week ago

widgets=people,keywords; controls=search; showcaptions=true; captions=en-US
upvoted 1 times

✉ **Daemon69** 1 year, 11 months ago

<https://docs.microsoft.com/en-us/azure/azure-video-indexer/video-indexer-embed-widgets>
upvoted 1 times

✉ **Daemon69** 1 year, 11 months ago

Cognitive Insights widget - answer is correct
Player widget - answer is correct
upvoted 2 times

Question #5

DRAG DROP -

You train a Custom Vision model to identify a company's products by using the Retail domain.

You plan to deploy the model as part of an app for Android phones.

You need to prepare the model for deployment.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions

Change the model domain.

Retrain the model.

Test the model.

Export the model.

Answer Area

Correct Answer:

Actions

Answer Area

Change the model domain.

Retrain the model.

Test the model.

Export the model.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/export-your-model>

czmiel24 Highly Voted 2 years, 9 months ago

Actually the model should be retrained prior to publishing:

"From the top of the page, select Train to retrain using the new domain."

So it should be:

1. Change the model domain

2. Retrain

3. Publish

upvoted 37 times

DingDongSingSong 2 years ago

Where is the Test step before publishing? After retraining you must test it before publishing it

upvoted 1 times

dinesh_tng 2 years, 8 months ago

Yep, Change the model to Retail (Compact). Exporting the Model is an optional step.

upvoted 4 times

ninja 1 year, 9 months ago

Agreed. see reference <https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/export-your-model>
upvoted 4 times

✉️ **dinesh_tng** 2 years, 8 months ago

Actually all four steps required in the sequence Change, Retrain, Test and Export. Export is also must as model has to be deployed on Android App. If I have to choose three options, I may drop "Test" as that is not mandatory to proceed, but good to have as part of process.
upvoted 15 times

✉️ **Mehe323** 2 months, 3 weeks ago

The question states: 'Which three actions should you perform in sequence?'

upvoted 1 times

✉️ **rdemontis** 6 months, 3 weeks ago

agree with you

upvoted 1 times

✉️ **reachmymind** Highly Voted 2 years, 2 months ago

Change the model domain {Retail(compact)}

Retrain the model

Export the model

<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/export-your-model>

upvoted 11 times

✉️ **nanaw770** Most Recent 2 days, 19 hours ago

1. Change the model domain

2. Retrain

3. Export

upvoted 1 times

✉️ **JamesKJoker** 1 week, 3 days ago

The Android App has Internet Connectivity, so there is no need for the Compact Version and also no need (and not even the option) for an Export Therefore.

1. Change the Model

2. Retrain the Model

3. Test the Model.

upvoted 2 times

✉️ **varinder82** 2 months ago

Final Answer:

Change the model domain {Retail(compact)}

Retrain the model

Export the model

upvoted 1 times

✉️ **zellck** 11 months ago

1. Change model domain

2. Retrain model

3. Test model

4. Export model

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/export-your-model>

upvoted 1 times

✉️ **dazdzadzadzaazd** 10 months, 1 week ago

Only three actions are requested. Test is optional imo

upvoted 5 times

✉️ **Pixelmate** 11 months ago

was on exam 28/06

upvoted 3 times

✉️ **RAN_L** 1 year, 2 months ago

Change the model domain: Since you trained the model using the Retail domain, you need to switch the domain to one that is optimized for mobile devices such as the General (compact) domain.

Retrain the model: After changing the domain, you need to retrain the model using the new domain settings.

Export the model: Once the model is retrained, you can export it in the format that is compatible with your Android app. The model can be exported as a TensorFlow or Core ML model for deployment on Android.

upvoted 4 times

✉️ **NNU** 1 year, 3 months ago

The model was trained, we must test it, change the model domain and export it.

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/test-your-model> (How-to guides) the first is test, second change domain in Export models finally export model

upvoted 1 times

✉️ **KingChuang** 1 year, 4 months ago
on my exam (2023-01-16 Passed)

My Answer:

1. Change the model domain
2. Retrain
3. Publish

upvoted 5 times

✉️ **AdarshKumarKhare** 1 year, 6 months ago
Export the Model must be excluded as it is asking for only three steps not four.
upvoted 1 times

✉️ **AdarshKumarKhare** 1 year, 6 months ago
But answer shows 4 steps
upvoted 2 times

✉️ **sumanshu** 2 years, 4 months ago
Change the Domain to compact (for mobiles, because of size)
Retrain
Export to Mobile
upvoted 3 times

✉️ **Phong0411** 2 years, 5 months ago
Was on the exam 30/11/2021
upvoted 1 times

✉️ **gs23mi** 2 years, 7 months ago
actions:
1 Change the model domain
2 Retrain the model
3 Export the model
Reference: <https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/export-your-model>
"Convert to a compact domain" for action #1 and #2
"Export your model" for action #3
upvoted 1 times

✉️ **SnowCheetah** 2 years, 8 months ago
Base on link provide answer
<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/export-your-model>

In user want to change to deploy offline model
1. Change model domain to compact model
2. Retrain compact model
3. Export model
upvoted 3 times

✉️ **SuperPetey** 2 years, 9 months ago
The stated question asks for three actions and the provided answer gives four, therefor I propose the following as the correct answer based off of the same documentation linked:

1. Change the model domain (to compact retail)
2. Export the model
3. Test the model

upvoted 2 times

Question #6

HOTSPOT -

You are developing an application to recognize employees' faces by using the Face Recognition API. Images of the faces will be accessible from a URI endpoint.

The application has the following code.

```
def add_face(subscription_key, person_group_id, person_id, image_uri):
    headers = {
        'Content-Type': 'application/json',
        'Ocp-Apim-Subscription-Key': subscription_key
    }
    body = {
        'url': image_uri
    }
    conn = httpplib.HTTPEConnection('westus.api.cognitive.microsoft.com')
    conn.request('POST',
    f'/face/v1.0/persongroups/{person_group_id}/persons/{person_id}/persistedFaces', f'{body}', headers)
    response = conn.getresponse()
    response_data = response.read()
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area**Statements**

Yes	No
<input type="radio"/>	<input type="radio"/>

The code will add a face image to a person object in a person group.

<input type="radio"/>	<input type="radio"/>
-----------------------	-----------------------

The code will work for up to 10,000 people.

<input type="radio"/>	<input type="radio"/>
-----------------------	-----------------------

add_face can be called multiple times to add multiple face images to a person object.

<input type="radio"/>	<input type="radio"/>
-----------------------	-----------------------

Correct Answer:**Answer Area****Statements**

Yes	No
<input checked="" type="radio"/>	<input type="radio"/>

The code will add a face image to a person object in a person group.

<input checked="" type="radio"/>	<input type="radio"/>
----------------------------------	-----------------------

The code will work for up to 10,000 people.

<input type="radio"/>	<input checked="" type="radio"/>
-----------------------	----------------------------------

add_face can be called multiple times to add multiple face images to a person object.

<input checked="" type="radio"/>	<input type="radio"/>
----------------------------------	-----------------------

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/face/face-api-how-to-topics/use-persondirectory>

✉  **Internal_Koala**  1 year, 8 months ago

Based on the subscription, I think, it could also be

Yes

Yes

Yes

"Free-tier subscription quota: 1,000 person groups. Each holds up to 1,000 persons.

S0-tier subscription quota: 1,000,000 person groups. Each holds up to 10,000 persons."

<https://docs.microsoft.com/en-us/rest/api/faceapi/person-group/create?tabs=HTTP>

upvoted 11 times

✉  **james2033** 9 months, 1 week ago

person groups, not persons. 2 choose No.

upvoted 2 times

✉  **AzureJobsTillRetire** 1 year, 3 months ago

The second box should be No. The given answers are correct. The second box states that the code will work for up to 10,000 people. While this is true for S0 tier, it is false for free-tier. Since the price tier is not given, we will have to say that it is not always true, and that means it is false

upvoted 9 times

✉  **surasahoo** 1 year, 3 months ago

Hi, have you passed the exam? Did you simulation questions?

upvoted 4 times

✉  **Adobe02** 1 year, 3 months ago

Following

upvoted 1 times

✉️ **Rob77** 1 year ago

2nd is "no". Nothing is stopping you from specifying another group using the code so even free tier is 1000x1000 = 1m people
upvoted 9 times

✉️ **rdemontis** 6 months, 3 weeks ago

agree with you
upvoted 1 times

✉️ **ziggy1117** Highly Voted 11 months, 3 weeks ago

Yes
No
Yes

"Free-tier subscription quota: 1,000 person groups. Each holds up to 1,000 persons. So in the code you can have 1000 person groups and 1000 persons each giving you 1,000,000 people

upvoted 9 times

✉️ **nanaw770** Most Recent 2 days, 20 hours ago

Yes
No
Yes
upvoted 1 times

✉️ **takaimomoGcup** 1 week ago

Yes No Yes
upvoted 1 times

✉️ **Murtuza** 1 month, 4 weeks ago

The code you've provided is intended to add a face image to a person object in a person group using the Azure Face API, so:

Yes, the code will add a face image to a person object in a person group, provided the code is corrected for syntax errors and proper API usage.
Yes, the code can work for up to 10,000 people, as long as the Azure Face API limits are adhered to and the appropriate subscription tier is used.
Yes, the add_face function can be called multiple times to add multiple face images to a person object, subject to the limits imposed by the Azure Face API.

upvoted 1 times

✉️ **Ody** 2 months, 1 week ago

Images sent to the service are not stored after analysis.

<https://learn.microsoft.com/en-us/legal/cognitive-services/face/data-privacy-security>

upvoted 1 times

✉️ **audlindr** 3 months ago

Is this a trick question?
From here: <https://westus.dev.cognitive.microsoft.com/docs/services/563879b61984550e40cbbe8d/operations/563879b61984550f3039523b>
No image will be stored. Only the extracted face feature(s) will be stored on server until PersonGroup PersonFace - Delete, PersonGroup Person - Delete or PersonGroup - Delete is called.
upvoted 1 times

✉️ **davidorti** 4 months, 2 weeks ago

Yes
No - As the code will keep working in other groups (for instance), and as AzureJobsTillRetire says, a statement that's not generally true is false
No - As the image is not really 'added', just their features

"No image will be stored. Only the extracted face feature will be stored on server until PersonGroup PersonFace - Delete, PersonGroup Person - Delete or PersonGroup - Delete is called." --see: <https://learn.microsoft.com/en-us/rest/api/faceapi/person-group/create?view=rest-faceapi-v1.0&tabs=HTTP>

The way it works, you have to update a face <https://learn.microsoft.com/en-us/rest/api/faceapi/person-group-person/update-face?view=rest-faceapi-v1.0&tabs=HTTP>. If you register a new pic for an existing user it will just create a new one and return a new persistedFacId.

upvoted 2 times

✉️ **Ody** 2 months, 1 week ago

If that's your reasoning, then you can't select Yes for the first one. They use the same verbiage in both.
upvoted 1 times

✉️ **sl_mslconsulting** 7 months, 1 week ago

You can vary the person group id and the person id while making the call, so even with the free tier is still way above the mention 10,000 people limitation. The only limitation that is defined for this API is this Response 403

Persisted face number reached limit, maximum is 248 per person.

Link here: <https://westus.dev.cognitive.microsoft.com/docs/services/563879b61984550e40cbbe8d/operations/563879b61984550f3039523b>

upvoted 1 times

✉️ **sl_mslconsulting** 7 months, 1 week ago

The 10,000 limitation would make more sense if they are asking about this API: POST {Endpoint}/face/v1.0/persongroups/{personGroupId}/persons
upvoted 1 times

 **momentumhd** 1 year, 8 months ago

Once you have the Person ID from the Create Person call, you can add up to 248 face images to a Person per recognition model.
They are all true, the limit is 75 milion persons per group

upvoted 4 times

Question #7

DRAG DROP -

You have a Custom Vision resource named acvdev in a development environment.

You have a Custom Vision resource named acvprod in a production environment.

In acvdev, you build an object detection model named obj1 in a project named proj1.

You need to move obj1 to acvprod.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions

Use the ExportProject endpoint on acvdev.

Use the GetProjects endpoint on acvdev.

Use the ImportProject endpoint on acvprod.

Use the ExportIteration endpoint on acvdev.

Use the GetIterations endpoint on acvdev.

Use the UpdateProject endpoint on acvprod.

Answer Area**Correct Answer:****Actions**

Use the ExportIteration endpoint on acvdev.

Use the GetIterations endpoint on acvdev.

Use the UpdateProject endpoint on acvprod.

Answer Area

Use the GetProjects endpoint on acvdev.

Use the ExportProject endpoint on acvdev.

**Reference:**<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/copy-move-projects> **snna4** Highly Voted 2 years, 5 months ago

1. GetProjects on acvDEV
 2. ExportProjects on acvDEV
 3. ImportProjects on acvPROD
- upvoted 30 times

 zellck Highly Voted 11 months ago

1. Use GetProjects endpoint on acvDEV
2. Use ExportProjects endpoint on acvDEV
3. Use ImportProjects endpoint on acvPROD

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/copy-move-projects#get-the-project-id>

First call GetProjects to see a list of your existing Custom Vision projects and their IDs. Use the training key and endpoint of your source account.

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/copy-move-projects#export-the-project>
Call ExportProject using the project ID and your source training key and endpoint.<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/copy-move-projects#import-the-project>
Call ImportProject using your target training key and endpoint, along with the reference token. You can also give your project a name in its new account.

upvoted 11 times

✉ **zellck** 11 months ago

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/copy-move-projects#process-overview>

The process for copying a project consists of the following steps:

- First, you get the ID of the project in your source account you want to copy.
- Then you call the ExportProject API using the project ID and the training key of your source account. You'll get a temporary token string.
- Then you call the ImportProject API using the token string and the training key of your target account. The project will then be listed under your target account.

upvoted 3 times

✉ **varinder82** Most Recent 2 months ago

Final Answer:

1. GetProjects on acvDEV
2. ExportProjects on acvDEV
3. ImportProjects on acvPROD

upvoted 2 times

✉ **rdemontis** 6 months, 3 weeks ago

correct answer as the documentation provided demonstrates

upvoted 1 times

✉ **Eltooth** 1 year, 10 months ago

Get on Dev

Export on Dev

Import on Prod

upvoted 1 times

✉ **Jzerpa_ccs** 1 year, 10 months ago

1. GetProjects on acvDEV
2. ExportProjects on acvDEV
3. ImportProjects on acvPROD

<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/copy-move-projects>

upvoted 3 times

✉ **Contactfornitish** 2 years, 4 months ago

Was on exam 02/01/2022

upvoted 2 times

✉ **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 3 times

✉ **Derin_tade** 2 years, 8 months ago

Given link proves this is correct.

upvoted 3 times

✉ **DS_sam2701** 2 years, 8 months ago

Here in this document it is clearly mentioned how can you move your resource from dev. to prod. : <https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-tutorial-pattern#what-did-this-tutorial-accomplish>

upvoted 2 times

✉ **fhqhfqh** 2 years, 7 months ago

Provided link is for LUIS. Incorrect Link.

upvoted 2 times

Question #8

DRAG DROP -

You are developing an application that will recognize faults in components produced on a factory production line. The components are specific to your business.

You need to use the Custom Vision API to help detect common faults.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

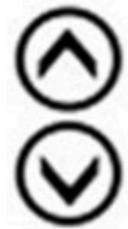
Select and Place:

Actions**Answer Area**

Train the classifier model.



Upload and tag images.



Initialize the training dataset.

Train the object detection model.

Create a project.

Correct Answer:

Actions**Answer Area**

Create a project.



Upload and tag images.



Initialize the training dataset.

Train the classifier model.

Step 1: Create a project -

Create a new project.

Step 2: Upload and tag the images

Choose training images. Then upload and tag the images.

Step 3: Train the classifier model.

Train the classifier -

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/getting-started-build-a-classifier>

arbest Highly Voted 1 year, 3 months ago

The answer is correct

Create a project

Upload and tag images

Train a classifier model

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/quickstarts/image-classification?tabs=visual-studio&pivots=programming-language-csharp>

For Type of model

<https://azure.microsoft.com/en-us/use-cases/defect-detection-with-image-analysis/>
upvoted 16 times

rdemontis Highly Voted 6 months, 3 weeks ago

The choice between Object Detection and Classification depends on the nature of your defect detection problem in the production line components. Here are some considerations:

Classification: You use classification when your goal is to determine whether a component is defective or non-defective. In this case, the Custom Vision model will be trained to classify the entire image as "good" or "defective." Classification is suitable if you want a binary answer.

Object Detection: You use object detection when you want to identify and locate specific defects or objects within an image. This is useful if you have multiple defect classes or if you want to identify the exact location of defects within a component.

So, if you only need to distinguish between good and defective components, classification may suffice. However, if you need to identify and locate specific defects within components, you should opt for object detection. The choice depends on the complexity of your use case and the level of detail you want to extract from the images. (Chat GPT)

Honestly I think Classifier is more appropriate in this case

upvoted 7 times

✉ **nanaw770** Most Recent 2 days, 20 hours ago

1. Create a project
2. Upload and tag images
3. Train a classifier model

upvoted 2 times

✉ **varinder82** 2 months ago

Final Answer:

- Create a project
Upload and tag images
Train a classifier model

upvoted 1 times

✉ **Ody** 2 months, 1 week ago

I think the key is recognizing it's a production line. That means the same component is coming down the line.

Then there are other lines producing other components.

Classification would be best suited for that scenario. If we were looking at an image that might have many different components in one image and want to find the location of different faults, then object would be more appropriate.

upvoted 1 times

✉ **suzanne_exam** 4 months ago

It's a classifier model because it's not detecting whether they objects are there or not, it's classifying them as faulty or not

upvoted 1 times

✉ **sl_msiconsulting** 7 months, 3 weeks ago

I would pick object detection. Custom Vision functionality can be divided into two features. Image classification applies one or more labels to an entire image. Object detection is similar, but it returns the coordinates in the image where the applied label(s) can be found. If I were the users, I would certainly want know where the faults are located to make sure and have a second look and it's quite useless by just telling me there are something wrong but can't tell you where they are as I need to do extra work to find them!

upvoted 1 times

✉ **NNU** 1 year, 3 months ago

Yes the anwser is correct

Create a project

Upload and tags images

Train the classifier model

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/quickstarts/image-classification?tabs=visual-studio&pivots=programming-language-csharp>

and for type of model

<https://azure.microsoft.com/en-us/use-cases/defect-detection-with-image-analysis/>

upvoted 3 times

✉ **Adedoyin_Simeon** 1 year, 5 months ago

Correct answer should be:

Create

Upload & Tag

Train the object detection model

The question was to help "detect" common faults. Detection means where the fault actually is in the image.

upvoted 2 times

✉ **cce1** 1 year, 4 months ago

Nope, answer should be

Create, Upload & Tag, and Train classifier (not a detection mode)

Bcz classifier has to classify whether the given component is faulty or not...

upvoted 5 times

✉ **Eltooth** 1 year, 10 months ago

Create

Upload

Train the classifier

upvoted 3 times

✉ **ppo12** 1 year, 10 months ago

Quite confusing on the questions, since Object Detection technically can be correct IMO
upvoted 2 times

✉ **momentumhd** 1 year, 8 months ago

You don't tag the detection images so by exclusion you could direct the answer to classification
upvoted 1 times

✉ **kiassi1998** 2 years, 1 month ago

Correct

upvoted 4 times

✉ **sdokmak** 1 year, 11 months ago

Agreed. Train the classifier, not object detection model because they make no mention of need to know the location of the detections, but they do mention detecting common faults. So can either classify as faulty, not faulty, or also classify different fault types.. not clear on that one but the answer is correct.

upvoted 4 times

Question #9

HOTSPOT -

You are building a model that will be used in an iOS app.

You have images of cats and dogs. Each image contains either a cat or a dog.

You need to use the Custom Vision service to detect whether the images is of a cat or a dog.

How should you configure the project in the Custom Vision portal? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area**Project Types:**

Classification
Object Detection

Classification Types:

Multiclass (Single tag per image)
Multilabel (Multiple tags per image)

Domains:

Audit
Food
General
General (compact)
Landmarks
Landmarks (compact)
Retail
Retail (compact)

Answer Area**Project Types:**

Classification
Object Detection

Classification Types:

Multiclass (Single tag per image)
Multilabel (Multiple tags per image)

Correct Answer:**Domains:**

Audit
Food
General
General (compact)
Landmarks
Landmarks (compact)
Retail
Retail (compact)

Box 1: Classification -

Incorrect Answers:

An object detection project is for detecting which objects, if any, from a set of candidates are present in an image.

Box 2: Multiclass -

A multiclass classification project is for classifying images into a set of tags, or target labels. An image can be assigned to one tag only.

Incorrect Answers:

A multilabel classification project is similar, but each image can have multiple tags assigned to it.

Box 3: General -

General: Optimized for a broad range of image classification tasks. If none of the other specific domains are appropriate, or if you're unsure of which domain to choose, select one of the General domains.

Reference:

<https://cran.r-project.org/web/packages/AzureVision/vignettes/customvision.html>

✉  **dinhhungitsoft**  2 years ago

The third choice should be General compact, in other that the model can be exported to be used in iOS device
upvoted 38 times

✉  **g2000** 2 years ago

it seems the general compact is for edge device not ios.
<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/select-domain#image-classification>
upvoted 3 times

✉  **sdokmak** 1 year, 11 months ago

So general compact is correct since ios device is an edge device.
upvoted 13 times

✉  **ExamDev** 8 months, 2 weeks ago

How about this article? <https://learn.microsoft.com/en-us/azure/ai-services/custom-vision-service/export-your-model>
upvoted 3 times

✉  **rdemontis** 6 months, 3 weeks ago

thanks for sharing the documentation.
upvoted 1 times

✉  **rdemontis** 6 months, 3 weeks ago

agree with you
upvoted 1 times

✉  **Eltooth**  1 year, 10 months ago

Classification
Multiclass
General (compact)
upvoted 13 times

✉  **nanaw770**  2 days, 20 hours ago

1. Classification
2. Multiclass
3. General (compact)
upvoted 2 times

✉  **[Removed]** 1 month, 2 weeks ago

1 - Classification
2 - Multiclass (single tag per image)
3 - General (compact)

Custom Vision Service only exports projects with compact domains.

The models generated by compact domains are optimized for the constraints of real-time classification on mobile devices.
Classifiers built with a compact domain may be slightly less accurate than a standard domain with the same amount of training data.
upvoted 1 times

✉  **varinder82** 2 months ago

Final Answer:
Classification
Multiclass
General (compact)
upvoted 1 times

✉  **zellck** 11 months ago

1. Classification
2. Multiclass
3. General (compact)

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/getting-started-build-a-classifier>

- Select Classification under Project Types. Then, under Classification Types, choose either Multilabel or Multiclass, depending on your use case. Multilabel classification applies any number of your tags to an image (zero or more), while multiclass classification sorts images into single categories (every image you submit will be sorted into the most likely tag). You'll be able to change the classification type later if you want to.
upvoted 4 times

✉️ **zellck** 11 months ago

- Next, select one of the available domains. Each domain optimizes the model for specific types of images, as described in the following table. You can change the domain later if you wish.

-- Generic

Optimized for a broad range of image classification tasks. If none of the other domains are appropriate, or you're unsure of which domain to choose, select the Generic domain.

-- Compact domains

Optimized for the constraints of real-time classification on mobile devices. The models generated by compact domains can be exported to run locally.

upvoted 1 times

✉️ **AiEngineerS** 1 year, 10 months ago

I also think that General(compact) <https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/export-your-model>

1. It can be running offline

2. Real time locally

upvoted 3 times

✉️ **SamedKia** 2 years, 1 month ago

Correct

upvoted 1 times

Question #10

You have an Azure Video Analyzer for Media (previously Video Indexer) service that is used to provide a search interface over company videos on your company's website.

You need to be able to search for videos based on who is present in the video.

What should you do?

- A. Create a person model and associate the model to the videos.
- B. Create person objects and provide face images for each object.
- C. Invite the entire staff of the company to Video Indexer.
- D. Edit the faces in the videos.
- E. Upload names to a language model.

Correct Answer: A

Video Indexer supports multiple Person models per account. Once a model is created, you can use it by providing the model ID of a specific Person model when uploading/indexing or reindexing a video. Training a new face for a video updates the specific custom model that the video was associated with.

Note: Video Indexer supports face detection and celebrity recognition for video content. The celebrity recognition feature covers about one million faces based on commonly requested data source such as IMDB, Wikipedia, and top LinkedIn influencers. Faces that aren't recognized by the celebrity recognition feature are detected but left unnamed. Once you label a face with a name, the face and name get added to your account's Person model. Video Indexer will then recognize this face in your future videos and past videos.

Reference:

<https://docs.microsoft.com/en-us/azure/media-services/video-indexer/customize-person-model-with-api>

Community vote distribution

A (94%) 6%

✉ **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

A seems the correct answer

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/customize-person-model-overview>

upvoted 5 times

✉ **nanaw770** 2 days, 20 hours ago

Selected Answer: A

A is right.

upvoted 2 times

✉ **propanther** 7 months, 1 week ago

Selected Answer: A

You can use a Person model to index your new video by assigning the Person model during the upload of the video.

You can use a Person model to index your new video by assigning the Person model during the upload of the video.

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/customize-person-model-with-website>

upvoted 3 times

✉ **sl_msiconsulting** 7 months, 3 weeks ago

Selected Answer: B

Customers can build custom Person models and enable Azure AI Video Indexer to recognize faces that aren't recognized by default. Customers can build these Person models by pairing a person's name with image files of the person's face. You can use the default model if you like but the point is to add persons and their associated image files to the model : <https://learn.microsoft.com/en-us/azure/azure-video-indexer/customize-person-model-with-website>.

upvoted 1 times

✉ **sl_msiconsulting** 7 months, 3 weeks ago

Azure AI Video Indexer can detect occurrences of this person in the future videos that you index and the current videos that you had already indexed, using the Person model to which you added this new face - you need to tell it what it should be comparing to when looking for a given face.

upvoted 1 times

✉ **sl_msiconsulting** 7 months, 3 weeks ago

Also If you don't need the multiple Person model support, don't assign a Person model ID to your video when uploading/indexing or reindexing. In this case, Azure AI Video Indexer will use the default Person model in your account. one more reason not to choose A.
upvoted 1 times

✉ **zellck** 11 months ago

Same as Question 1.

<https://www.examtopics.com/discussions/microsoft/view/55438-exam-ai-102-topic-2-question-1-discussion>

upvoted 1 times

✉ **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/customize-person-model-overview>

Azure Video Indexer supports celebrity recognition in your videos. The celebrity recognition feature covers approximately one million faces based on commonly requested data source such as IMDB, Wikipedia, and top LinkedIn influencers. Faces that are not recognized by Azure Video Indexer are still detected but are left unnamed. Customers can build custom Person models and enable Azure Video Indexer to recognize faces that are not recognized by default. Customers can build these Person models by pairing a person's name with image files of the person's face.

upvoted 2 times

✉ **Nebary** 1 year, 9 months ago

Selected Answer: A

Should be A

upvoted 4 times

✉ **g2000** 2 years ago

seems right

<https://docs.microsoft.com/en-us/azure/azure-video-indexer/customize-person-model-with-website>

upvoted 1 times

Question #11

You use the Custom Vision service to build a classifier.

After training is complete, you need to evaluate the classifier.

Which two metrics are available for review? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. recall
- B. F-score
- C. weighted accuracy
- D. precision
- E. area under the curve (AUC)

Correct Answer: AD

Custom Vision provides three metrics regarding the performance of your model: precision, recall, and AP.

Reference:

<https://www.tallan.com/blog/2020/05/19/azure-custom-vision/>

Community vote distribution

AD (100%)

✉️ **evangelist** 3 months, 1 week ago

Selected Answer: AD

Zellck provided an excellent answer with precise documentation and interpretation of precision and recall metrics. Precision refers to the positive predictive value - the proportion of true positive results among all positively identified outcomes. Meanwhile, recall represents sensitivity - the proportion of actual positive cases that are correctly identified as such.

Precision and recall form a fundamental pair of performance indicators that entail an inherent trade-off. As one metric is optimized, the other typically suffers as a consequence. Specifically, as the precision rate increases, the recall rate often correspondingly decreases. The optimal balance between precision and recall depends on the business context and specific needs of the use case. By clearly explaining the definitions and relationship between these two metrics, Zellck thoroughly addressed the concepts with clarity and accuracy.

upvoted 1 times

✉️ **rdemontis** 6 months, 3 weeks ago

Selected Answer: AD

Answer is correct

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/getting-started-build-a-classifier#evaluate-the-classifier>

upvoted 1 times

✉️ **trashbox** 7 months ago

Appeared on Oct/29/2023.

upvoted 3 times

✉️ **zellck** 11 months ago

Same as Question 2.

<https://www.examtopics.com/discussions/microsoft/view/55211-exam-ai-102-topic-2-question-2-discussion>

upvoted 1 times

✉️ **zellck** 11 months ago

Selected Answer: AD

AD is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/getting-started-build-a-classifier#evaluate-the-classifier>

After training has completed, the model's performance is estimated and displayed. The Custom Vision Service uses the images that you submitted for training to calculate precision and recall. Precision and recall are two different measurements of the effectiveness of a classifier:

- Precision indicates the fraction of identified classifications that were correct. For example, if the model identified 100 images as dogs, and 99 of them were actually of dogs, then the precision would be 99%.

- Recall indicates the fraction of actual classifications that were correctly identified. For example, if there were actually 100 images of apples, and the model identified 80 as apples, the recall would be 80%.

upvoted 3 times

✉️ **halfway** 1 year, 5 months ago

Selected Answer: AD

Precision and Recall: <https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/getting-started-build-a-classifier#evaluate-the-classifier>

upvoted 1 times

 **Eltooth** 1 year, 10 months ago

Selected Answer: AD

A and D are correct answers - as per PHD_CHENG

<https://docs.microsoft.com/en-us/learn/modules/cv-classify-bird-species/4-understand-results-test>

upvoted 2 times

 **PHD_CHENG** 1 year, 11 months ago

Was on exam 7 Jun 2022

upvoted 1 times

 **PHD_CHENG** 2 years ago

Selected Answer: AD

Answer is correct. You can find the metrics from Microsoft link <https://docs.microsoft.com/en-us/learn/modules/cv-classify-bird-species/4-understand-results-test>

upvoted 4 times

Question #12

DRAG DROP -

You are developing a call to the Face API. The call must find similar faces from an existing list named `employeefaces`. The `employeefaces` list contains 60,000 images.

How should you complete the body of the HTTP request? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Values

`"faceListId"`
`"LargeFaceListId"`
`"matchFace"`
`"matchPerson"`

Answer Area

```
{
  "faceId": "18c51a87-3a69-47a8-aedc-a54745f708a1",
  "LargeFaceListId": "employeefaces",
  "maxNumOfCandidatesReturned": 1,
  "mode": "matchFace"
}
```

Correct Answer:

Values

`"faceListId"`
`"LargeFaceListId"`
`"matchFace"`
`"matchPerson"`

Answer Area

```
{
  "faceId": "18c51a87-3a69-47a8-aedc-a54745f708a1",
  "LargeFaceListId": "employeefaces",
  "maxNumOfCandidatesReturned": 1,
  "mode": "matchFace"
}
```

Box 1: LargeFaceListID -

`LargeFaceList`: Add a face to a specified large face list, up to 1,000,000 faces.

Note: Given query face's `facelId`, to search the similar-looking faces from a `facelId` array, a face list or a large face list. A "`faceListId`" is created by `FaceList - Create` containing `persistedFacelIds` that will not expire. And a "`largeFaceListId`" is created by `LargeFaceList - Create` containing `persistedFacelIds` that will also not expire.

Incorrect Answers:

Not "`faceListId`": Add a face to a specified face list, up to 1,000 faces.

Box 2: `matchFace` -

`FindSimilar` has two working modes, "`matchPerson`" and "`matchFace`". "`matchPerson`" is the default mode that it tries to find faces of the same person as possible by using internal same-person thresholds. It is useful to find a known person's other photos. Note that an empty list will be returned if no faces pass the internal thresholds. "`matchFace`" mode ignores same-person thresholds and returns ranked similar faces anyway, even the similarity is low. It can be used in the cases like searching celebrity-looking faces.

Reference:

<https://docs.microsoft.com/en-us/rest/api/faceapi/face/findsimilar>

✉  **Jenny1**  2 years, 11 months ago

Correct.

upvoted 18 times

✉  **PHD_CHENG**  2 years ago

`Facelist ID` up to 1,000 faces; `LargeFaceListId` up to 1,000,000 faces

<https://docs.microsoft.com/en-us/rest/api/faceapi/large-face-list>

upvoted 9 times

✉  **nanaw770**  2 days, 19 hours ago

1. `LargeFaceListId`

2. `matchFace`

upvoted 1 times

✉  **varinder82** 2 months ago

Final Answer:

1. LargeFaceListId
2. matchFace (matchPerson could return an empty list, but matchFace will not.)

upvoted 3 times

✉ **Ody** 2 months, 1 week ago

The key is "must find". matchPerson could return an empty list, but matchFace will not.

upvoted 1 times

✉ **rdemontis** 6 months, 3 weeks ago

correct answer:

<https://learn.microsoft.com/en-us/rest/api/faceapi/face-list?view=rest-faceapi-v1.0>

upvoted 1 times

✉ **rdemontis** 6 months, 2 weeks ago

Sorry but I have to laugh off my previous answer:

"matchPerson" is the default mode that it tries to find faces of the same person as possible by using internal same-person thresholds. It is useful to find a known person's other photos. Note that an empty list will be returned if no faces pass the internal thresholds. "matchFace" mode ignores same-person thresholds and returns ranked similar faces anyway, even the similarity is low. It can be used in the cases like searching celebrity-looking faces"

<https://learn.microsoft.com/en-us/rest/api/faceapi/face/find-similar?view=rest-faceapi-v1.0&tabs=HTTP>

upvoted 1 times

✉ **rdemontis** 6 months, 2 weeks ago

in this case we are talking about finding the faces of people who are employees of an organization. We are not talking about famous people. So the recognition will have to be done by an internal set of images of the faces whose faces are to be matched in the corporate photo gallery. Also set maxNumOfCandidatesReturned = 1. Which means that it is also acceptable not to find matches. Which is not the case with matchFace because it also tolerates a low confidence coefficient. Considering what is reported in the article I think the second answer to the question is more correct matchPerson.

upvoted 1 times

✉ **rdemontis** 6 months, 2 weeks ago

Sorry, but after further investigation I am not sure that this is the case. Frankly, the Microsoft documentation in this case doesn't seem to be very satisfactory.

upvoted 1 times

✉ **Eltooth** 1 year, 10 months ago

I'm leaning towards: largeFaceListId and matchedPerson

<https://docs.microsoft.com/en-us/rest/api/faceapi/face/find-similar?tabs=HTTP#find-similar-results-example>

"matchPerson" is the default mode that it tries to find faces of the same person as possible by using internal same-person thresholds. It is useful to find a known person's other photos. Note that an empty list will be returned if no faces pass the internal thresholds.

"matchFace" mode ignores same-person thresholds and returns ranked similar faces anyway, even the similarity is low. It can be used in the cases like searching celebrity-looking faces.

upvoted 7 times

✉ **makciek** 9 months, 2 weeks ago

it says "The call must find similar faces from an existing list named employefaces", similar being the keyword here so matchFace i think is correct.

upvoted 1 times

✉ **ccie_pgh** 5 months ago

Agree since code also has "maxNumOfCandidatesReturned", which would be ignored by "matchFace"... both "matchFace" and "matchPerson" find similars but only "matchPerson" will use internal thresholds.

upvoted 2 times

✉ **PHD_CHENG** 1 year, 11 months ago

Was on exam 7 Jun 2022

upvoted 3 times

✉ **luishenriquesb** 1 year, 12 months ago

it's should be largeFaceListId not LargeFaceListId (Capitalized). It's wouldn't work in a http request...

upvoted 5 times

✉ **gursimran_s** 2 years, 1 month ago

The 1000 face parameter is for the facelids and not the faceListId. So, it could be faceListId as well.

upvoted 1 times

✉ **gursimran_s** 2 years, 1 month ago

Why not faceListId? Nothing specific mentioned on MS docs.

upvoted 1 times

✉ **klion** 2 years, 3 months ago

"matchPerson"

Find similar results example

Sample Request

HTTP

POST {Endpoint}/face/v1.0/findsimilar

Ocp-Apim-Subscription-Key: {API key}

Request Body

JSON

```
{  
    "facelId": "c5c24a82-6845-4031-9d5d-978df9175426",  
    "largeFaceListId": "sample_list",  
    "maxNumOfCandidatesReturned": 1,  
    "mode": "matchPerson"  
}
```

<https://docs.microsoft.com/en-us/rest/api/faceapi/face/find-similar>

upvoted 2 times

✉ **sumanshu** 2 years, 4 months ago

In Question it's given we have to find similar faces - So we have to use "matchFace" and because there are large list , so we have to use LargeFaceListID

upvoted 5 times

✉ **mikegsm** 2 years, 6 months ago

Seems Correct

upvoted 2 times

Question #13

DRAG DROP -

You are developing a photo application that will find photos of a person based on a sample image by using the Face API.

You need to create a POST request to find the photos.

How should you complete the request? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all.

You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Values	Answer Area
detect	POST {Endpoint}/face/v1.0/ <input type="text"/>
findsimilar	Request Body
group	
identify	
matchFace	
matchPerson	
verify	

```

POST {Endpoint}/face/v1.0/
{
  "faceId": "c5c24a82-6845-4031-9d5d-978df9175426",
  "largeFaceListId": "sample_list",
  "maxNumOfCandidatesReturned": 10,
  "mode": "  "
}

```

Correct Answer:

Values	Answer Area
detect	POST {Endpoint}/face/v1.0/ <input type="text"/> detect
findsimilar	Request Body
group	
identify	
matchFace	
matchPerson	
verify	

```

POST {Endpoint}/face/v1.0/
{
  "faceId": "c5c24a82-6845-4031-9d5d-978df9175426",
  "largeFaceListId": "sample_list",
  "maxNumOfCandidatesReturned": 10,
  "mode": "  matchPerson "
}

```

Box 1: detect -

Face - Detect With Url: Detect human faces in an image, return face rectangles, and optionally with faceIds, landmarks, and attributes.

POST {Endpoint}/face/v1.0/detect

Box 2: matchPerson -

Find similar has two working modes, "matchPerson" and "matchFace". "matchPerson" is the default mode that it tries to find faces of the same person as possible by using internal same-person thresholds. It is useful to find a known person's other photos. Note that an empty list will be returned if no faces pass the internal thresholds. "matchFace" mode ignores same-person thresholds and returns ranked similar faces anyway, even the similarity is low. It can be used in the cases like searching celebrity-looking faces.

Reference:

<https://docs.microsoft.com/en-us/rest/api/faceapi/face/detectwithurl> <https://docs.microsoft.com/en-us/rest/api/faceapi/face/findsimilar>

 **motu** Highly Voted  2 years, 11 months ago

Box 1 is "findsimilar", others do not match the given request body and make no sense anyway. <https://docs.microsoft.com/en-us/rest/api/faceapi/face/find-similar>

upvoted 69 times

 **leo822** 2 years, 11 months ago

cool. correct answer!

upvoted 2 times

 **idrisfl** 2 years, 11 months ago

definitely find-similar, as it is the only one whose body parameters correspond

upvoted 3 times

Eltooth **Highly Voted** 1 year, 10 months ago
findsimilar and matchPerson

<https://docs.microsoft.com/en-us/rest/api/faceapi/face/find-similar?tabs=HTTP#find-similar-results-example>

"matchPerson" is the default mode that it tries to find faces of the same person as possible by using internal same-person thresholds. It is useful to find a known person's other photos. Note that an empty list will be returned if no faces pass the internal thresholds.

"matchFace" mode ignores same-person thresholds and returns ranked similar faces anyway, even the similarity is low. It can be used in the cases like searching celebrity-looking faces.

upvoted 11 times

nanaw770 **Most Recent** 2 days, 19 hours ago
1. findsimilar
2. matchPerson
upvoted 1 times

varinder82 2 months ago

Final Answer:
findsimilar and matchPerson
upvoted 1 times

rdemontis 6 months, 3 weeks ago

Box 1 is findsimilar. Box 2 is matchPerson and the provided explanation is correct
upvoted 2 times

zellck 11 months ago

1. findsimilar
2. matchPerson

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/find-similar-faces?tabs=rest#find-and-print-matches>
The Find Similar operation does face matching between a target face and a set of candidate faces, finding a smaller set of faces that look similar to the target face. This is useful for doing a face search by image.

upvoted 10 times

reachmymind 2 years, 2 months ago

Box 1: findsimilar
Box 2: matchPerson

<https://dev.cognitive.azure.cn/docs/services/563879b61984550e40cbbe8d/operations/563879b61984550f30395237>
upvoted 5 times

bitcoin89 2 years, 3 months ago

FIRST BOX IDEBTIFY SECOND BOX NOTHING
POST {Endpoint}/face/v1.0/identify
Ocp-Apim-Subscription-Key: {API key}
{
"largePersonGroupId": "sample_group",
"facelids": [
"c5c24a82-6845-4031-9d5d-978df9175426",
"65d083d4-9447-47d1-af30-b626144bf0fb"
],
"maxNumOfCandidatesReturned": 1,
"confidenceThreshold": 0.5
}
upvoted 2 times

sumanshu 2 years, 4 months ago

Box 1 - FindSimilar
Box 2 - matchPerson (We have to find based on a sample photo)
upvoted 1 times

mikegsm 2 years, 6 months ago

Seems FIND SIMILAR AND MATCHPERSON
upvoted 1 times

DeBoer 2 years, 7 months ago

Looking at the ENTIRE document the answer has to be findsimilar: You cannot send the properties like faceListID and largeFaceListId to /detect
upvoted 2 times

nitkat 2 years, 8 months ago

The Answer is correct. The question asks to "find photos of a person based on a sample image". Key is "based on a sample image". Only detect does this : <https://docs.microsoft.com/en-us/rest/api/faceapi/face/detect-with-url>. Find Similar is used to search the similar-looking faces from a faceld array, a face list or a large face list

upvoted 4 times

Zoul 2 years, 7 months ago

detect does not take faceid. Cannot be detect !

upvoted 1 times

✉️ **Banye27** 2 years, 9 months ago

<https://docs.microsoft.com/en-us/rest/api/faceapi/face/find-similar>

upvoted 2 times

✉️ **azurelearner666** 2 years, 11 months ago

Correct!

upvoted 2 times

✉️ **Dalias** 2 years, 11 months ago

Motu is correct. verified the link too

<https://docs.microsoft.com/en-us/rest/api/faceapi/face/find-similar>

upvoted 4 times

✉️ **LKLK10** 2 years, 11 months ago

I think it's Verify. "Do two images of a face belong to the same person? This defines verification." <https://docs.microsoft.com/en-us/learn/modules/identify-faces-with-computer-vision/4-overview-of-face-recognition>

upvoted 4 times

Question #14

HOTSPOT -

You develop a test method to verify the results retrieved from a call to the Computer Vision API. The call is used to analyze the existence of company logos in images. The call returns a collection of brands named brands.

You have the following code segment.

```
for brand in image_analysis.brands:
    if brand_confidence >= 0.75:
        print(f"\nLogo of {brand_name} between {brand.rectangle_x}, {brand.rectangle.y} and
{brand.rectangle.w}, {brand.rectangle.h}")
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
The code will return the name of each detected brand with a confidence equal to or higher than 75 percent.	<input type="radio"/>	<input type="radio"/>
The code will return coordinates for the top-left corner of the rectangle that contains the brand logo of the displayed brands.	<input type="radio"/>	<input type="radio"/>
The code will return coordinates for the bottom-right corner of the rectangle that contains the brand logo of the displayed brands.	<input type="radio"/>	<input type="radio"/>

Answer Area

Statements	Yes	No
The code will return the name of each detected brand with a confidence equal to or higher than 75 percent.	<input checked="" type="radio"/>	<input type="radio"/>
The code will return coordinates for the top-left corner of the rectangle that contains the brand logo of the displayed brands.	<input checked="" type="radio"/>	<input type="radio"/>
The code will return coordinates for the bottom-right corner of the rectangle that contains the brand logo of the displayed brands.	<input type="radio"/>	<input checked="" type="radio"/>

Box 1: Yes -

Box 2: Yes -

Coordinates of a rectangle in the API refer to the top left corner.

Box 3: No -

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-brand-detection>

 **halfway**  1 year, 6 months ago

Maybe I take it too literally, but I think the third one is "NO": the response returns Width and Height, which can be used to calculate the coordinates of bottom right corner, but it does not include them directly.

upvoted 12 times

 **M25**  9 months ago

Y, Y, N

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/how-to/shelf-analyze#bounding-box-api-model>

x Left-coordinate of the top left point of the area, in pixels.

y Top-coordinate of the top left point of the area, in pixels.

w Width measured from the top-left point of the area, in pixels.

h Height measured from the top-left point of the area, in pixels.

upvoted 6 times

- ✉ **rdemontis** 6 months, 3 weeks ago
thanks for posting the documentation
upvoted 2 times
- ✉ **nanaw770** **Most Recent** 2 days, 20 hours ago
This is the same question as Topic2 #29.
upvoted 2 times
- ✉ **takaimomoGcup** 1 week ago
Yes Yes No
upvoted 2 times
- ✉ **takaimomoGcup** 1 week ago
Yes Yes No
upvoted 2 times
- ✉ **michaelmorar** 4 weeks ago
The brand_confidence variable is not declared in the snippet. Perhaps they meant brand.confidence? Same with rectangle_x . As it stands, the answers should be N,N,N.
upvoted 2 times
- ✉ **[Removed]** 1 month, 2 weeks ago
The Python syntax for working with Python dictionaries is wrong. Remember we are calling the API which returns a JSON object.

brand.rectangle_x - WRONG - it should be brand.rectangle.x
brand_confidence - WRONG - should be brand.confidence

So the answer should be N for all!
upvoted 1 times
- ✉ **varinder82** 2 months ago
Final Answer:
Yes
Yes
No
upvoted 2 times
- ✉ **RAN_L** 1 year, 2 months ago
The code will return the name of each detected brand with a confidence equal to or higher than 75 percent.
Yes, the code will return the name of each detected brand with a confidence equal to or higher than 75 percent.

2. The code will return coordinates for the top-left corner of the rectangle that contains the brand logo of the displayed brands.
Yes, the code will return the coordinates for the top-left corner of the rectangle that contains the brand logo of the displayed brands.

3. The code will return coordinates for the bottom-right corner of the rectangle that contains the brand logo of the displayed brands.
No, the code will not return coordinates for the bottom-right corner of the rectangle that contains the brand logo of the displayed brands. The code is printing the width and height of the rectangle instead.
upvoted 2 times
- ✉ **VinnieG** 1 year, 4 months ago
it could be a trap : so yes, no (it is rectangle.x and not _x) , no (should be x + w as the service returns the width and top left corner :

Console.WriteLine("Brands:");
foreach (var brand in results.Brands)
{
 Console.WriteLine(\$"Logo of {brand.Name} with confidence {brand.Confidence} at location {brand.Rectangle.X}, " +
 \$"{brand.Rectangle.X + brand.Rectangle.W}, {brand.Rectangle.Y}, {brand.Rectangle.Y + brand.Rectangle.H}");
}
Console.WriteLine();

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/call-analyze-image?tabs=csharp>
upvoted 2 times
- ✉ **TJ001** 1 week, 1 day ago
hope it is a typo. explanation is spot on
upvoted 1 times
- ✉ **rober13** 2 months, 4 weeks ago
it is true, it is a trick.
Yes, No, No
upvoted 1 times
- ✉ **rafael0** 1 year, 5 months ago
its' yes
yes
no
The coordinates are always regarding the top left point of the rectangle

upvoted 3 times

✉ **oliverio** 1 year, 7 months ago

Y

Y

Y

the code will return the coordinates for any position

upvoted 1 times

✉ **Anulf** 1 year, 7 months ago

yes

yes

yes

upvoted 1 times

✉ **Davard** 1 year, 7 months ago

What makes the third one "yes"?

upvoted 2 times

✉ **Anulf** 1 year, 7 months ago

According to the microsoft Document, I thought so. What makes you think it is No ?

upvoted 2 times

✉ **GigaCaster** 1 year, 7 months ago

if you look at the code it is _x and not .x

upvoted 2 times

✉ **AzureJobsTillRetire** 1 year, 3 months ago

I think that is a typo

upvoted 2 times

Question #15

HOTSPOT -

You develop an application that uses the Face API.

You need to add multiple images to a person group.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
Parallel.For(0, PersonCount, async i =>
{
    Guid personId = persons[i].PersonId;
    string personImageDir = $"{path}/path/to/person/{i}/images";
    foreach (string imagePath in Directory.GetFiles(personImageDir, "*.jpg"))
    {
        using (File t = File.OpenRead(imagePath))
        {
            await faceClient.PersonGroupPerson.
                (personGroupId, personId, t);
        }
    });
});
```

File
Stream
Uri
Url

AddFaceFromStreamAsync
AddFaceFromUrlAsync
CreateAsync
GetAsync

Correct Answer:

Answer Area

```
Parallel.For(0, PersonCount, async i =>
{
    Guid personId = persons[i].PersonId;
    string personImageDir = $"{path}/path/to/person/{i}/images";
    foreach (string imagePath in Directory.GetFiles(personImageDir, "*.jpg"))
    {
        using (Stream t = File.OpenRead(imagePath))
        {
            await faceClient.PersonGroupPerson.
                (personGroupId, personId, t);
        }
    });
});
```

File
Stream
Uri
Url

AddFaceFromStreamAsync
AddFaceFromUrlAsync
CreateAsync
GetAsync

Box 1: Stream -

The File.OpenRead(String) method opens an existing file for reading.

Example: Open the stream and read it back.

```
using (FileStream fs = File.OpenRead(path))
```

Box 2: CreateAsync -

Create the persons for the PersonGroup. Persons are created concurrently.

Example:

```
await faceClient.PersonGroupPerson.CreateAsync(personGroupId, personName);
```

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/face/face-api-how-to-topics/how-to-add-faces>

✉  **leo822**  2 years, 11 months ago

AddFaceFromStreamAsync. Step 5 on <https://docs.microsoft.com/en-us/azure/cognitive-services/face/face-api-how-to-topics/how-to-add-faces>
upvoted 63 times

✉  **rdemontis** 6 months, 3 weeks ago

thanks for the provided document. Clearly wrong the second answer. It should be AddFaceFromStreamAsync
upvoted 3 times

✉  **azurelearner666**  2 years, 11 months ago

Wrong!

A - Stream (this is correct)

B - AddFaceFromStreamAsync

(literally the same code from Step 5 at <https://docs.microsoft.com/en-us/azure/cognitive-services/face/face-api-how-to-topics/how-to-add-faces>)

upvoted 46 times

✉  **nanaw770**  2 days, 19 hours ago

1. Stream
 2. AddFaceFromStreamAsync
- upvoted 1 times

✉  **nanaw770** 3 days, 19 hours ago

Stream and CreateAsync

upvoted 1 times

✉  **varinder82** 2 months ago

Final Answer:

A - Stream (this is correct)

B - AddFaceFromStreamAsync

upvoted 2 times

✉  **chenglim** 7 months ago

1. Stream
 2. AddFaceFromStreamAsync
- <https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/how-to/add-faces#step-5-add-faces-to-the-persons>
upvoted 2 times

✉  **zellck** 11 months ago

1. Stream
2. AddFaceFromStreamAsync

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/add-faces#step-5-add-faces-to-the-persons>
upvoted 5 times

✉  **examworld** 12 months ago

Stream
AddFaceFromStreamAsync
upvoted 1 times

✉  **Mike19D** 1 year ago

Stream
AddFaceFromStreamAsync
upvoted 1 times

✉  **ninja** 1 year, 9 months ago

Box 1: Stream
Box 2: AddFaceFromStreamAsync

File.OpenRead() returns a Stream object.

```
using (Stream stream = File.OpenRead(imagePath))
{
    await faceClient.PersonGroupPerson.AddFaceFromStreamAsync(personGroupId, personId, stream);
}
```

ref: <https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/add-faces#step-5-add-faces-to-the-persons>

upvoted 3 times

Eltooth 1 year, 10 months ago

Stream and AddFaceFromStreamAsync are correct answers.

<https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/add-faces#step-5-add-faces-to-the-persons>

upvoted 3 times

Deepusuraj 2 years, 3 months ago

```
Parallel.For(0, PersonCount, async i =>
{
    Guid personId = persons[i].PersonId;
    string personImageDir = @"/path/to/person/{i}/images";

    foreach (string imagePath in Directory.GetFiles(personImageDir, "*.jpg"))
    {
        await WaitCallLimitPerSecondAsync();

        using (Stream stream = File.OpenRead(imagePath))
        {
            await faceClient.PersonGroupPerson.AddFaceFromStreamAsync(personGroupId, personId, stream);
        }
    }
});
```

upvoted 2 times

sumanshu 2 years, 4 months ago

Stream and AddFaceFromStreamAsync

upvoted 3 times

Happiness20 2 years, 8 months ago

```
Parallel.For(0, PersonCount, async i =>
{
    Guid personId = persons[i].PersonId;
    string personImageDir = @"/path/to/person/{i}/images";

    foreach (string imagePath in Directory.GetFiles(personImageDir, "*.jpg"))
    {
        await WaitCallLimitPerSecondAsync();

        using (Stream stream = File.OpenRead(imagePath))
        {
            await faceClient.PersonGroupPerson.AddFaceFromStreamAsync(personGroupId, personId, stream);
        }
    }
});
```

<https://docs.microsoft.com/en-us/azure/cognitive-services/face/face-api-how-to-topics/how-to-add-faces>

upvoted 3 times

Happiness20 2 years, 8 months ago

```
Parallel.For(0, PersonCount, async i =>
{
    Guid personId = persons[i].PersonId;
    string personImageDir = @"/path/to/person/{i}/images";

    foreach (string imagePath in Directory.GetFiles(personImageDir, "*.jpg"))
    {
        await WaitCallLimitPerSecondAsync();

        using (Stream stream = File.OpenRead(imagePath))
        {
            await faceClient.PersonGroupPerson.AddFaceFromStreamAsync(personGroupId, personId, stream);
        }
    }
});
```

<https://docs.microsoft.com/en-us/azure/cognitive-services/face/face-api-how-to-topics/how-to-add-faces>

upvoted 2 times

Happiness20 2 years, 8 months ago

```
Parallel.For(0, PersonCount, async i =>
{
    Guid personId = persons[i].PersonId;
    string personImageDir = @"/path/to/person/{i}/images";

    foreach (string imagePath in Directory.GetFiles(personImageDir, "*.jpg"))
    {
```

```
await WaitCallLimitPerSecondAsync();

using (Stream stream = File.OpenRead(imagePath))
{
    await faceClient.PersonGroupPerson.AddFaceFromStreamAsync(personGroupId, personId, stream);
}
};

upvoted 1 times
```

✉ **ramkinkarpandey** 2 years, 10 months ago

```
Parallel.For(0, PersonCount, async i =>
{
    Guid personId = persons[i].PersonId;
    string personImageDir = @"/path/to/person/{i}/images";

    foreach (string imagePath in Directory.GetFiles(personImageDir, "*.jpg"))
    {
        await WaitCallLimitPerSecondAsync();

        using (Stream stream = File.OpenRead(imagePath))
        {
            await faceClient.PersonGroupPerson.AddFaceFromStreamAsync(personGroupId, personId, stream);
        }
    };
}
```

upvoted 3 times

Question #16

Your company uses an Azure Cognitive Services solution to detect faces in uploaded images. The method to detect the faces uses the following code.

```
static async Task DetectFaces(string imagePath)
{
    HttpClient client = new HttpClient();
    DefaultRequestHeaders.Add("Ocp-Apim-Subscription-Key", subscriptionKey);
    string requestParameter = "detectionModel=detection_01&returnFaceId=true&returnFaceLandmarks=false";
    string uri = endpoint + "/face/v1.0/detect?" + requestParameters;
    HttpResponseMessage response;
    byte[] byteData = GetImagesAsByteArray(imagePath);
    using (ByteArrayContent content = new ByteArrayContent(byteData))
    {
        Headers.ContentType = new MediaTypeHeaderValue("application/octet-stream");
        response = await PostAsync(uri, content);
        string contentString = await Content.ReadAsStringAsync();
        ProcessDetection(contentString);
    }
}
```

You discover that the solution frequently fails to detect faces in blurred images and in images that contain sideways faces.

You need to increase the likelihood that the solution can detect faces in blurred images and images that contain sideways faces.

What should you do?

- A. Use a different version of the Face API.
- B. Use the Computer Vision service instead of the Face service.
- C. Use the Identify method instead of the Detect method.
- D. Change the detection model.

Correct Answer: D

Evaluate different models.

The best way to compare the performances of the detection models is to use them on a sample dataset. We recommend calling the Face-Detect API on a variety of images, especially images of many faces or of faces that are difficult to see, using each detection model. Pay attention to the number of faces that each model returns.

The different face detection models are optimized for different tasks. See the following table for an overview of the differences.

detection_01	detection_02	detection_03
Default choice for all face detection operations.	Released in May 2019 and available optionally in all face detection operations.	Released in February 2021 and available optionally in all face detection operations.
Not optimized for small, side-view, or blurry faces.	Improved accuracy on small, side-view, and blurry faces.	Further improved accuracy, including on smaller faces (64x64 pixels) and rotated face orientations.
Returns main face attributes (head pose, age, emotion, and so on) if they're specified in the detect call.	Does not return face attributes.	Returns mask and head pose attributes if they're specified in the detect call.
Returns face landmarks if they're specified in the detect call.	Does not return face landmarks.	Returns face landmarks if they're specified in the detect call.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/face/face-api-how-to-topics/specify-detection-model>

Community vote distribution

D (100%)

Eltooth [Highly Voted] 1 year, 10 months ago

Selected Answer: D

D is correct answer : change the detection model.

<https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/specify-detection-model#evaluate-different-models>
upvoted 12 times

takaimomoGcup [Most Recent] 1 week ago

Selected Answer: D

I use detection_02 or detection_03.
upvoted 1 times

evangelist 3 months, 1 week ago

Selected Answer: D

it contains sideway faces, so change face API detection model 03
upvoted 1 times

rdemontis 6 months, 3 weeks ago

Selected Answer: D

answer is correct
upvoted 2 times

zellck 11 months ago

Selected Answer: D

D is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/specify-detection-model#evaluate-different-models>
The different face detection models are optimized for different tasks.
- detection_02
Improved accuracy on small, side-view, and blurry faces.
upvoted 3 times

Question #17

You have the following Python function for creating Azure Cognitive Services resources programmatically.

```
def create_resource(resource_name, kind, account_tier, location):
    parameters = CognitiveServicesAccount(sku=Sku(name=account_tier), kind=kind, location=location, properties={})
    result = client.accounts.create(resource_group_name, resource_name, parameters)
```

You need to call the function to create a free Azure resource in the West US Azure region. The resource will be used to generate captions of images automatically.

Which code should you use?

- A. create_resource("res1", "ComputerVision", "F0", "westus")
- B. create_resource("res1", "CustomVision.Prediction", "F0", "westus")
- C. create_resource("res1", "ComputerVision", "S0", "westus")
- D. create_resource("res1", "CustomVision.Prediction", "S0", "westus")

Correct Answer: B

F0 is the free tier.

Custom Vision Service -

Upload images to train and customize a computer vision model for your specific use case. Once the model is trained, you can use the API to tag images using the model and evaluate the results to improve your classifier.

Incorrect:

Not C, not D: S0 is the standard tier, which isn't free.

Not A, not C: The Computer Vision service provides developers with access to advanced algorithms for processing images and returning information.

Computer Vision -

Returns information about visual content found in an image:

Use tagging, descriptions, and domain-specific models to identify content and label it with confidence.

Apply adult/racy settings to enable automated restriction of adult content.

Identify image types and color schemes in pictures.

Reference:

<https://docs.microsoft.com/en-us/python/api/overview/azure/cognitive-services?view=azure-python>

Community vote distribution

A (98%)

✉️  **ArchMelody**  1 year, 6 months ago

Selected Answer: A

Computer vision provide automatic vision solutions including captions. The key-phrase is "automatic". Therefore this answer should be obvious to everyone. I would expect more professionalism from people who request money for services like this one. Many questions here have incorrect and even contradictory answers... Shame!

upvoted 38 times

✉️  **ulloo** 1 year ago

I agree.

To me it looks like many answers are being deliberately set to incorrect answers. Not sure why, though.

upvoted 3 times

✉️  **nanaw770**  2 days, 20 hours ago

Selected Answer: A

ComputerVision and F0.

upvoted 2 times

✉️  **upliftinghut** 1 month ago

duplicate question

upvoted 1 times

✉️  **sca88** 6 months ago

Selected Answer: A

The answer is A

upvoted 1 times

✉ rdemontis 6 months, 3 weeks ago

Selected Answer: A

Seems to be duplicated question. Anyway to me the correct answer is A.

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/concept-describing-images>

upvoted 3 times

✉ trashbox 7 months ago

Appeared on Oct/29/2023.

upvoted 2 times

✉ bare 7 months, 4 weeks ago

Very confusion, Does Computer Vision provide captions in Free? I think Custom Vision Do:

Computer Vision Free tier:

Instance | Features | Price

Free - Web/Container 20 transactions per minute | - 5,000 transactions free per month

Custom Vision Free Tire:

Instance | Transactions Per Second (TPS) | Features | Price

Free | 2 TPS | Upload, training and prediction transactions, Up to 2 projects, Up to 1 hour training per month | 5,000 training images free per project, 10,000 predictions per month

upvoted 1 times

✉ zellck 11 months ago

Same as Topic 1 Question 9.

<https://www.examtopics.com/discussions/microsoft/view/57153-exam-ai-102-topic-1-question-9-discussion>

upvoted 3 times

✉ zellck 11 months ago

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/cognitive-services-apis-create-account-client-library?pivots=programming-language-python#create-a-cognitive-services-resource-python>

To create and subscribe to a new Cognitive Services resource, use the Create function. This function adds a new billable resource to the resource group you pass in. When you create your new resource, you'll need to know the "kind" of service you want to use, along with its pricing tier (or SKU) and an Azure location. The following function takes all of these arguments and creates a resource.

upvoted 2 times

✉ zellck 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 2 times

✉ Pixelmate 11 months ago

This was on exam 28/06

upvoted 2 times

✉ Pixelmate 11 months ago

Asked in 28/06/2023 exam

upvoted 2 times

✉ ziggy1117 11 months, 3 weeks ago

Selected Answer: A

computer vision can generate captions for images

upvoted 2 times

✉ ap1234pa 1 year, 4 months ago

Selected Answer: A

Computer Vision has generate captions feature

upvoted 1 times

✉ SSJA 1 year, 5 months ago

Selected Answer: B

This question was asked for free azure service. Do we have the generate caption feature supports this with free tier?

Reference - <https://azure.microsoft.com/en-us/pricing/details/cognitive-services/computer-vision/>

upvoted 1 times

✉ [Removed] 1 year, 8 months ago

Selected Answer: A

Should be A.

upvoted 4 times

✉ be_ml_team 1 year, 8 months ago

Selected Answer: A

A because computer vision provides tags

upvoted 2 times

✉ **goo1994** 1 year, 8 months ago

Answer Should be A.

upvoted 2 times

✉ **goo1994** 1 year, 8 months ago

one of the feature of Computer vision: The Image Analysis service extracts many visual features from images, such as objects, faces, adult content, and auto-generated text descriptions. Follow the Image Analysis quickstart to get started.

<https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/overview>

upvoted 3 times

Question #18

You are developing a method that uses the Computer Vision client library. The method will perform optical character recognition (OCR) in images.

The method has the following code.

```
def read_file_url(computervision_client, url_file):
    read_response = computervision_client.read(url_file, raw=True)
    read_operation_location = read_response.headers["Operation-Location"]
    operation_id = read_operation_location.split("/")[-1]
    read_result = computervision_client.get_read_result(operation_id)

    for page in read_result.analyze_result.read_results:
        for line in page.lines:
            print(line.text)
```

During testing, you discover that the call to the GetReadResultAsync method occurs before the read operation is complete.

You need to prevent the GetReadResultAsync method from proceeding until the read operation is complete.

Which two actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Remove the operation_id parameter.
- B. Add code to verify the read_results.status value.
- C. Add code to verify the status of the read_operation_location value.
- D. Wrap the call to get_read_result within a loop that contains a delay.

Correct Answer: BD

Community vote distribution

BD (100%)

 **RAN_L**  1 year, 2 months ago

Selected Answer: BD

- B. Add code to verify the read_results.status value.
- D. Wrap the call to get_read_result within a loop that contains a delay.

Explanation:

In order to prevent the GetReadResultAsync method from proceeding until the read operation is complete, we need to check the status of the read operation and wait until it's completed. To do this, we can add code to verify the status of the read_results.status value. If the status is not "succeeded", we can add a delay and then retry the operation until it's complete. This can be achieved by wrapping the call to get_read_result within a loop that contains a delay.

Removing the operation_id parameter or adding code to verify the status of the read_operation_location value will not solve the issue of waiting for the read operation to complete before proceeding with the GetReadResultAsync method.

upvoted 5 times

 **takaimomoGcup**  1 week ago

Selected Answer: BD

- memorize "read_results" words. So B and D.
- upvoted 1 times

 **takaimomoGcup** 1 week ago

Selected Answer: BD

- B and D.
- upvoted 1 times

 **[Removed]** 1 month, 2 weeks ago

Selected Answer: BD

The code snippet does not match the question. There is NO GetReadResultAsync call.

upvoted 1 times

 **anto69** 3 months, 4 weeks ago

where's GetReadResultAsync in the code? looks like C# method not python code

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: BD

duplicated question
upvoted 2 times

✉ **zellck** 11 months ago
Same as Question 2.
<https://www.examtopics.com/discussions/microsoft/view/74739-exam-ai-102-topic-2-question-2-discussion>
upvoted 2 times

✉ **Pixelmate** 11 months ago
This was on exam 28/06
upvoted 3 times

✉ **halfway** 1 year, 6 months ago

Selected Answer: BD

Duplicated with Topic 2, Question 2.
upvoted 1 times

✉ **michasacuer** 1 year, 8 months ago

Selected Answer: BD

Correct
upvoted 1 times

Question #19

HOTSPOT -

You are building an app that will enable users to upload images. The solution must meet the following requirements:

- * Automatically suggest alt text for the images.
- * Detect inappropriate images and block them.
- * Minimize development effort.

You need to recommend a computer vision endpoint for each requirement.

What should you recommend? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Generate alt text:

<https://westus.api.cognitive.microsoft.com/contentmoderator/moderate/v1.0/ProcessImage/Evaluate>
 <https://westus.api.cognitive.microsoft.com/customvision/v3.1/prediction/projectId/classify/iterations/publishedName/image>
 <https://westus.api.cognitive.microsoft.com/vision/v3.2/analyze/?visualFeatures=Adult,Description>

Detect inappropriate content:

<https://westus.api.cognitive.microsoft.com/contentmoderator/moderate/v1.0/ProcessImage/Evaluate>
 <https://westus.api.cognitive.microsoft.com/customvision/v3.1/prediction/projectId/classify/iterations/publishedName/image>
 <https://westus.api.cognitive.microsoft.com/vision/v3.2/analyze/?visualFeatures=Adult,Description>
 <https://westus.api.cognitive.microsoft.com/vision/v3.2/describe?maxCandidates=1>

Correct Answer:**Answer Area**

Generate alt text:

<https://westus.api.cognitive.microsoft.com/contentmoderator/moderate/v1.0/ProcessImage/Evaluate>
 <https://westus.api.cognitive.microsoft.com/customvision/v3.1/prediction/projectId/classify/iterations/publishedName/image>
 <https://westus.api.cognitive.microsoft.com/vision/v3.2/analyze/?visualFeatures=Adult,Description>

Detect inappropriate content:

<https://westus.api.cognitive.microsoft.com/contentmoderator/moderate/v1.0/ProcessImage/Evaluate>
 <https://westus.api.cognitive.microsoft.com/customvision/v3.1/prediction/projectId/classify/iterations/publishedName/image>
 <https://westus.api.cognitive.microsoft.com/vision/v3.2/analyze/?visualFeatures=Adult,Description>
 <https://westus.api.cognitive.microsoft.com/vision/v3.2/describe?maxCandidates=1>

Box 1: <https://westus.api.cognitive.microsoft.com/customvision/v3.1/prediction/projectid/classify/iterations/publishName/image>

Box 2: <https://westus.api.cognitive.microsoft.com/vision/v3.2/analyze/?visualFeatures=Adult,Description>

Computer Vision can detect adult material in images so that developers can restrict the display of these images in their software. Content flags are applied with a score between zero and one so developers can interpret the results according to their own preferences.

You can detect adult content with the Analyze Image API. When you add the value of Adult to the visualFeatures query parameter

Incorrect:

Use the Image Moderation API in Azure Content Moderator to scan image content. The moderation job scans your content for profanity, and compares it against custom and shared blocklists.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-detecting-adult-content> <https://docs.microsoft.com/en-us/azure/cognitive-services/content-moderator/try-image-api> <https://docs.microsoft.com/en-us/legal/cognitive-services/custom-vision/custom-vision-cvs-transparency-note>

  **Tanmay1178**  1 year, 6 months ago

I think it is vision/v3.2/analyze/?visualFeatures=Adult,Description for both
upvoted 27 times

  **zellck**  11 months ago

1. <https://westus.api.cognitive.microsoft.com/vision/v3.2/analyze/?visualFeatures=Adult,Description>
2. <https://westus.api.cognitive.microsoft.com/vision/v3.2/analyze/?visualFeatures=Adult,Description>

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-describing-images>

Computer Vision can analyze an image and generate a human-readable phrase that describes its contents. The algorithm returns several descriptions based on different visual features, and each description is given a confidence score. The final output is a list of descriptions ordered from highest to lowest confidence.

upvoted 9 times

✉  **josebernabeo** 4 months, 1 week ago

The first one says "descption", so I guess that's not the correct answer.

Does anybody know the correct one?

upvoted 3 times

✉  **zellick** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 4 times

✉  **TheCloudGuruu** 10 months, 3 weeks ago

zelleck, do you say this on every question?

upvoted 2 times

✉  **EliteAllen** 10 months, 1 week ago

He does not, he doesn't say that in the next question Topic 2 Q20

upvoted 4 times

✉  **zman_83** 9 months, 1 week ago

Thank you for confirming that it a active/valid question, and it might appear in the exam

upvoted 2 times

✉  **zellick** 11 months ago

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-detecting-adult-content>

Computer Vision can detect adult material in images so that developers can restrict the display of these images in their software. Content flags are applied with a score between zero and one so developers can interpret the results according to their own preferences.

You can detect adult content with the Analyze Image 3.2 API. When you add the value of Adult to the visualFeatures query parameter, the API returns three boolean properties—isAdultContent, isRacyContent, and isGoryContent—in its JSON response. The method also returns corresponding properties—adultScore, racyScore, and goreScore—which represent confidence scores between zero and one for each respective category.

upvoted 2 times

✉  **M25** 9 months ago

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/use-case-alt-text?source=recommendations>

Alt text, or alternative text, is an HTML attribute added to the tag that displays images on an application or web page.

Alt text enables website owners to describe an image in plain text. These image descriptions improve accessibility by enabling screen readers such as Microsoft Narrator, JAWS, and NVDA to accurately communicate image content to their visually impaired and blind users.

Alt text is also vital for image search engine optimization (SEO). It helps search engines understand the visual content in your images. The search engine is then better able to include and rank your website in search results when users search for the content in your website.

Auto-generate alt text with Image Analysis: Image Analysis offers image captioning models that generate one-sentence descriptions of image visual content.

Thank you @zellick!

upvoted 1 times

✉  **takaimomoGcup** Most Recent 1 week ago

Generate alt text should be "v3.2/analyze". Detect inappropriate content should be "v3.2/analyze".

upvoted 1 times

✉  **evangelist** 3 months, 1 week ago

description ==> Alt text

upvoted 1 times

✉  **propanther** 7 months, 1 week ago

Both are <https://westus.api.cognitive.microsoft.com/vision/v3.2/analyze/?visualFeatures=Adult,Description>

A string indicating what visual feature types to return. Multiple values should be comma-separated.

Valid visual feature types include:

Adult - detects if the image is pornographic in nature (depicts nudity or a sex act), or is gory (depicts extreme violence or blood). Sexually suggestive content (aka racy content) is also detected.

Description - describes the image content with a complete sentence in supported languages.

<https://westus.dev.cognitive.microsoft.com/docs/services/computer-vision-v3-2/operations/56f91f2e778daf14a499f21b>

upvoted 2 times

✉  **Tickxit** 1 year, 7 months ago

I think it is two times /analyze/?visualFeatures=Adult,Description

upvoted 4 times

✉️ **Anulf** 1 year, 7 months ago

Shouldn't it be option 2 in the first column ? "iteratio"

upvoted 1 times

✉️ **firewind** 1 year, 8 months ago

Generate alt text can use either analyze or describe. From the given option, I think it should be the analyze url too.

<https://westcentralus.dev.cognitive.microsoft.com/docs/services/computer-vision-v3-2/operations/56f91f2e778daf14a499f21f>

upvoted 2 times

✉️ **Rob77** 1 year ago

Agreed, analyze in both answers

upvoted 1 times

Question #20

You need to build a solution that will use optical character recognition (OCR) to scan sensitive documents by using the Computer Vision API. The solution must NOT be deployed to the public cloud. What should you do?

- A. Build an on-premises web app to query the Computer Vision endpoint.
- B. Host the Computer Vision endpoint in a container on an on-premises server.
- C. Host an exported Open Neural Network Exchange (ONNX) model on an on-premises server.
- D. Build an Azure web app to query the Computer Vision endpoint.

Correct Answer: B

One option to manage your Computer Vision containers on-premises is to use Kubernetes and Helm.

Three primary parameters for all Cognitive Services containers are required. The Microsoft Software License Terms must be present with a value of accept. An Endpoint URI and API key are also needed.

Incorrect:

Not D: This Computer Vision endpoint would be available for the public, unless it is secured.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/deploy-computer-vision-on-premises>

Community vote distribution

B (100%)

✉  takaimomoGcup 1 week ago

Selected Answer: B

B is right answer.

upvoted 1 times

✉  evangelist 3 months, 1 week ago

Selected Answer: B

model is hosted on-premise but the billing information has to be provided to the on-premise container to report the usage

upvoted 1 times

✉  rdemontis 6 months, 3 weeks ago

Selected Answer: B

correct answer

<https://learn.microsoft.com/en-us/azure/ai-services/cognitive-services-container-support>

upvoted 2 times

✉  zellck 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/deploy-computer-vision-on-premises>

One option to manage your Computer Vision containers on-premises is to use Kubernetes and Helm. Using Kubernetes and Helm to define a Computer Vision container image, we'll create a Kubernetes package. This package will be deployed to a Kubernetes cluster on-premises.

upvoted 4 times

✉  RAN_L 1 year, 2 months ago

Selected Answer: B

B. Host the Computer Vision endpoint in a container on an on-premises server.

Since the solution should not be deployed to the public cloud, option B is the correct answer. By hosting the Computer Vision endpoint in a container on an on-premises server, the solution can still leverage the capabilities of the Computer Vision API while keeping the processing and data within the on-premises environment. Option A and D both involve using a web app, which would likely require hosting in the public cloud. Option C involves hosting an exported ONNX model, which may not have the same capabilities as the Computer Vision API.

upvoted 2 times

✉  HotDurian 1 year, 5 months ago

Selected Answer: B

Answer is correct.

upvoted 2 times

Question #21

You have an Azure Cognitive Search solution and a collection of handwritten letters stored as JPEG files.

You plan to index the collection. The solution must ensure that queries can be performed on the contents of the letters.

You need to create an indexer that has a skillset.

Which skill should you include?

- A. image analysis
- B. optical character recognition (OCR)
- C. key phrase extraction
- D. document extraction

Correct Answer: B

Community vote distribution

B (100%)

 **RAN_L**  1 year, 2 months ago

Selected Answer: B

To ensure that queries can be performed on the contents of the letters, the skill that should be included in the indexer is optical character recognition (OCR).

Option B, optical character recognition (OCR), is a technology that can recognize text within an image and convert it into machine-readable text. This skill will enable the search engine to read the handwritten letters and convert them into searchable text that can be indexed by Azure Cognitive Search.

Option A, image analysis, is a useful skill for analyzing images to extract metadata, but it does not directly enable text recognition.

Option C, key phrase extraction, extracts important phrases and concepts from text, but it requires the text to be already recognized and extracted by OCR or other text extraction techniques.

Option D, document extraction, is a skill that extracts specific pieces of information from documents, but it does not address the challenge of recognizing and extracting text from handwritten letters.

upvoted 10 times

 **takaimomoGcup**  1 week ago

Why OCR?

upvoted 1 times

 **evangelist** 3 months, 1 week ago

Selected Answer: B

OCR to extract the text and then create an indexer on the text extracted

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

provided answer is correct. OCR to scan handwritten documents

upvoted 1 times

 **zellck** 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-skill-ocr>

The Optical character recognition (OCR) skill recognizes printed and handwritten text in image files.

upvoted 1 times

 **ExamPage** 11 months ago

Therefore, option B, optical character recognition (OCR), is the most suitable skill to include in the indexer for indexing the contents of handwritten letters and making them searchable.

upvoted 1 times

Question #22

HOTSPOT

You have a library that contains thousands of images.

You need to tag the images as photographs, drawings, or clipart.

Which service endpoint and response property should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Service endpoint:

- Computer Vision analyze images
- Computer Vision object detection
- Custom Vision image classification
- Custom Vision object detection

Property:

- categories
- description
- imageType
- metadata
- objects

Answer Area

Correct Answer:

Service endpoint:

- Computer Vision analyze images
- Computer Vision object detection
- Custom Vision image classification
- Custom Vision object detection

Property:

- categories
- description
- imageType
- metadata
- objects

 **jimbojambo** Highly Voted  1 year, 2 months ago

I think that the answers are wrong. They should be:

1 - Computer Vision analyze image

2 - imageType

According to <https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-detecting-image-types> Computer Vision can analyze the content type of images, indicating whether an image is clip art or a line drawing

upvoted 42 times

 **TJ001** 1 week, 1 day ago

out of box option this is the best bet if not custom vision and train the model for object detection - more work

upvoted 1 times

 **mmaguero** 12 months ago

Agree, see json example at: <https://westcentralus.dev.cognitive.microsoft.com/docs/services/computer-vision-v3-2/operations/56f91f2e778daf14a499f21b>

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

agree and thanks for posting the related documentation

upvoted 1 times

 **zellck** Highly Voted  11 months ago

1. Computer Vision analyze images

2. imageType

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-detecting-image-types>

With the Analyze Image API, Computer Vision can analyze the content type of images, indicating whether an image is clip art or a line drawing.
upvoted 11 times

 **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 6 times

 **nanaw770** Most Recent 2 days, 21 hours ago

1. Computer Vision analyze images

2. imageType

upvoted 1 times

 **takaimomoGcup** 1 week ago

Service endpoint should be "Computer Vision analyze images". Property should be "imageType".

upvoted 1 times

 **takaimomoGcup** 1 week ago

Service endpoint shoud be "Computer Vision analyze images". Property should be "imageType".

upvoted 1 times

 **trashbox** 7 months ago

Appeared on Oct/29/2023.

upvoted 5 times

 **sl_mslconsulting** 7 months, 3 weeks ago

I would say the answers are correct. Image type can only indicate whether an image is clip art or a line drawing. It can't tell you if it's a photograph or not - you can't just assume that if the image isn't a clip art or a line drawing will automatically be categorized as a photograph. It's a very slopp solution IMO. Besides you have thousands of images and it's a good reason to create your own model.

upvoted 2 times

 **HarshSharma786** 11 months, 2 weeks ago

To tag images as photographs, drawings, or clipart, you should use the following service endpoint and response property:

Service endpoint: Computer Vision image classification

Property: imageType

The Computer Vision image classification endpoint allows you to classify images into different categories, and the imageType property specifically provides information about the type of image, such as whether it is a photograph, drawing, or clipart.

upvoted 4 times

 **ulloo** 1 year ago

ChatGPT:

You can use the Microsoft Azure Computer Vision API to tag the images as photographs, drawings, or clipart.

You can call the "Describe Image" API endpoint and use the "imageType" property of the response to determine if the image is a photograph, a drawing, or clipart. The "imageType" property can have the following values:

"Clipart": Indicates that the image is a clipart.

"LineDrawing": Indicates that the image is a line drawing.

"Photograph": Indicates that the image is a photograph.

You can send an HTTP POST request to the API endpoint with the image file as the request body and specify the "imageType" in the "visualFeatures" parameter. The API will return a JSON response containing the "imageType" property along with other properties such as "tags", "description", and "categories".

upvoted 2 times

Question #23

You have an app that captures live video of exam candidates.

You need to use the Face service to validate that the subjects of the videos are real people.

What should you do?

- A. Call the face detection API and retrieve the face rectangle by using the FaceRectangle attribute.
- B. Call the face detection API repeatedly and check for changes to the FaceAttributes.HeadPose attribute.
- C. Call the face detection API and use the FaceLandmarks attribute to calculate the distance between pupils.
- D. Call the face detection API repeatedly and check for changes to the FaceAttributes.Accessories attribute.

Correct Answer: A

Community vote distribution



✉️ **zellck** Highly Voted 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/use-headpose#detect-head-gestures>
You can detect head gestures like nodding and head shaking by tracking HeadPose changes in real time. You can use this feature as a custom liveness detector.

Liveness detection is the task of determining that a subject is a real person and not an image or video representation. A head gesture detector could serve as one way to help verify liveness, especially as opposed to an image representation of a person.

upvoted 11 times

✉️ **rdemontis** 6 months, 2 weeks ago

thanks for explanation

upvoted 2 times

✉️ **takaimomoGcup** Most Recent 1 week ago

Selected Answer: B

FaceAttributes.HeadPose is used this solution.

upvoted 1 times

✉️ **1668f51** 1 month, 2 weeks ago

Selected Answer: A

"You need to use the Face service to validate that the subjects of the videos are real people." Never think too much into these questions. A can detect if it's real or not. Not asking anything else. It's A.

upvoted 1 times

✉️ **rdemontis** 6 months, 2 weeks ago

Selected Answer: B

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/use-headpose#detect-head-gestures>

upvoted 1 times

✉️ **sl_msiconsulting** 7 months, 2 weeks ago

Selected Answer: B

It can't be A. If you only try to detect the faces without tracking their position over time, the system can be easily fooled in this very specific scenario.

upvoted 1 times

✉️ **msdfqwerfewf** 11 months ago

Selected Answer: A

Option A is more appropriate for validating the presence of real people in the live video. By calling the face detection API and retrieving the face rectangle using the FaceRectangle attribute, you can detect and locate faces within the video frames. This helps in confirming the presence of actual human faces in the captured video.

Option B, on the other hand, suggests repeatedly calling the face detection API and checking for changes to the FaceAttributes.HeadPose attribute. While head pose analysis can provide information about the orientation of detected faces, it may not be the most reliable approach for validating the authenticity of the subjects as real people. Checking for changes in head pose alone may not be sufficient to differentiate between real people and other forms of visual representations.

upvoted 2 times

✉ **ziggy1117** 11 months, 3 weeks ago

Selected Answer: B

B: <https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/use-headpose#detect-head-gestures>

upvoted 2 times

✉ **ziggy1117** 11 months, 3 weeks ago

B: <https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/use-headpose#detect-head-gestures>

upvoted 1 times

✉ **Rob77** 1 year ago

B <https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/use-headpose#detect-head-gestures>

"Liveness detection is the task of determining that a subject is a real person and not an image or video representation"

upvoted 3 times

✉ **Mike19D** 1 year, 1 month ago

Selected Answer: B

The Answer is B. A could be a still picture

upvoted 3 times

✉ **marti_tremblay000** 1 year, 2 months ago

Selected Answer: B

The answer is B

Detect head gestures

You can detect head gestures like nodding and head shaking by tracking HeadPose changes in real time. You can use this feature as a custom liveness detector.

Reference <https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/use-headpose>

upvoted 2 times

Question #24

HOTSPOT

You make an API request and receive the results shown in the following exhibits.

HTTP request

```
POST https://facetesting.cognitiveservices.azure.com/face/v1.0/detect?returnFaceId=true&returnFaceLandmarks=false&returnFaceAttributes=qualityForRecognition&recognitionModel=recognition_04&returnRecognitionModel=false&detectionModel=detection_03&faceIdTimeToLive=86400 HTTP/1.1
Host: facetesting.cognitiveservices.azure.com
Content-Type: application/json
Ocp-Apim-Subscription-Key: *****

{
    "url": "https://news.microsoft.com/wp-content/uploads/prod/sites/68/2021/11/EDU19_HigherEdStudentsOnCampus_002-1536x1024.jpg"
}
```

Send

Response status

200 OK

Response content

```
x-envoy-upstream-service-time: 1292
apim-request-id: 8a3aa72f-5bad-45d0-b8a4-584312258f06
Strict-Transport-Security: max-age=31536000; includeSubDomains; preload
x-content-type-options: nosniff
CSP-Billing-Usage: CognitiveServices.Face.Transaction=1
Date: Sat, 04 Dec 2021 11:15:33 GMT
Content-Length: 655
Content-Type: application/json; charset=utf-8
```

```
[{
  "faceId": "d14d131c-76ba-43e9-9e3d-dcf6466e5022",
  "faceRectangle": {
    "top": 201,
    "left": 797,
    "width": 121,
    "height": 160
  },
  "faceAttributes": {
    "qualityForRecognition": "high"
  }
}, {
  "faceId": "a3a0f2ff-b015-464c-b87c-0dd09d0698da",
  "faceRectangle": {
    "top": 249,
    "left": 1167,
    "width": 103,
    "height": 159
  },
  "faceAttributes": {
    "qualityForRecognition": "medium"
  }
}, {
  "faceId": "45481ce8-dcc4-4564-a21c-3c15cdc9c4fa",
  "faceRectangle": {
    "top": 191,
    "left": 497,
    "width": 85,
    "height": 178
  },
  "faceAttributes": {
    "qualityForRecognition": "low"
  }
}, {
  "faceId": "eac17649-effd-42c9-9093-4dd60fd4cf07",
  "faceRectangle": {
    "top": 754,
    "left": 118,
    "width": 30,
    "height": 44
  },
  "faceAttributes": {
    "qualityForRecognition": "low"
  }
}]
```

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Answer Area

The API [answer choice] faces.

- detects
- finds similar
- recognizes
- verifies

A face that can be used in person enrollment is at position [answer choice] within the photo.

- 118, 754
- 497, 191
- 797, 201
- 1167, 249

Answer Area**Correct Answer:**

A face that can be used in person enrollment is at position [answer choice] within the photo.

The API [answer choice] faces.

detects
finds similar
recognizes
verifies

118, 754
497, 191
797, 201
1167, 249

✉ **zellck** Highly Voted 11 months ago

1. detects
2. 797, 201

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-face-detection#face-rectangle>

Each detected face corresponds to a faceRectangle field in the response. This is a set of pixel coordinates for the left, top, width, and height of the detected face. Using these coordinates, you can get the location and size of the face. In the API response, faces are listed in size order from largest to smallest.

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-face-detection#attributes>

- QualityForRecognition

The overall image quality regarding whether the image being used in the detection is of sufficient quality to attempt face recognition on. The value is an informal rating of low, medium, or high. Only "high" quality images are recommended for person enrollment, and quality at or above "medium" is recommended for identification scenarios.

upvoted 9 times

✉ **ExamDev** Highly Voted 8 months, 2 weeks ago

To answer the first question see the endpoint .../face/v1.0/detect?....

To answer the second question see the first object from the response "faceRectangle": {"TOP":201, "LEFT":797,...}

upvoted 5 times

✉ **rdemontis** 6 months, 3 weeks ago

exactly. Answer is correct

upvoted 1 times

✉ **takaimomoGcup** Most Recent 1 week ago

memorize. detects and 797, 201.

upvoted 1 times

✉ **varinder82** 2 months ago

Final Answer:

1. detects
2. 797, 201

upvoted 1 times

✉ **RAN_L** 1 year, 2 months ago

The API detects faces.

A face that can be used in person enrollment is at position 797, 201 within the photo.

This question provides information about an API request made to a face detection service. The request is sent to the endpoint "<https://facetesting.cognitiveservices.azure.com/face/v1.0/detect>" with the content of an image in the JSON format. The response from the API includes an array of detected faces, each with a unique facId, faceRectangle, and faceAttributes.

The first statement asks what the API does with faces. The correct answer is "detects" because the endpoint used in the request is "/detect," which implies that the API is used for face detection.

The second statement asks about the position of a face that can be used for person enrollment. The face's position is specified in the "faceRectangle" field of the JSON response. The correct answer is "118, 754" because that is the "left" and "top" position of the face rectangle for the fourth face in the response, which has a high enough quality for recognition to be used in person enrollment.

upvoted 4 times

✉ **uira** 1 year, 1 month ago

"118, 754" has low quality, isn't it?

upvoted 1 times

Question #25

You have an Azure subscription that contains an AI enrichment pipeline in Azure Cognitive Search and an Azure Storage account that has 10 GB of scanned documents and images.

You need to index the documents and images in the storage account. The solution must minimize how long it takes to build the index.

What should you do?

- A. From the Azure portal, configure parallel indexing.
- B. From the Azure portal, configure scheduled indexing.
- C. Configure field mappings by using the REST API.
- D. Create a text-based indexer by using the REST API.

Correct Answer: A

Community vote distribution

A (100%)

✉  **zellck**  11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/search/search-howto-large-index#run-indexers-in-parallel>

If you partition your data, you can create multiple indexer-data-source combinations that pull from each data source and write to the same search index. Because each indexer is distinct, you can run them at the same time, populating a search index more quickly than if you ran them sequentially.

upvoted 7 times

✉  **rdemontis** 6 months, 3 weeks ago

thanks for explanation

upvoted 1 times

✉  **takaimomoGcup**  1 week ago

Selected Answer: A

A is right answer.

upvoted 1 times

✉  **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

answer is correct

upvoted 1 times

✉  **nitz14** 11 months, 2 weeks ago

Selected Answer: A

To minimize the time it takes to build the index for the documents and images in the Azure Storage account, the best approach would be to use parallel indexing.

Therefore, the correct option is:

A. From the Azure portal, configure parallel indexing.

Configuring parallel indexing allows you to process multiple documents or images simultaneously

upvoted 2 times

✉  **ulloo** 1 year ago

<https://learn.microsoft.com/en-us/azure/search/search-howto-large-index>

upvoted 1 times

✉  **Mike19D** 1 year ago

seems logical

upvoted 1 times

Question #26

DRAG DROP

You need to analyze video content to identify any mentions of specific company names.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

- Add the specific company names to the exclude list.
- Add the specific company names to the include list.
- From Content model customization, select **Language**.
- Sign in to the Custom Vision website.
- Sign in to the Azure Video Analyzer for Media website.
- From Content model customization, select **Brands**.

Answer Area**Answer Area**

- Sign in to the Azure Video Analyzer for Media website.
- From Content model customization, select **Brands**.
- Add the specific company names to the include list.

Correct Answer:

✉️  **zellck**  11 months ago

1. Sign in to Azure Video Analyzer for Media website
2. From Content model customization, select Brands
3. Add specific company names to include list

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/customize-brands-model-with-website>

upvoted 20 times

✉️  **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 8 times

✉️  **takaimomoGcup**  1 week ago

Sign in website, From Brands, Add include list.

I memorized this line.

upvoted 1 times

✉️  **evangelist** 3 months, 1 week ago

login to azure video analyzer==>select brands tab to customize the analyzer==>add specific company names to the "include" list

upvoted 2 times

✉️  **josebernabeo** 4 months, 1 week ago

Now it's call "Azure AI Video Indexer website "

upvoted 1 times

✉️  **rdemontis** 6 months, 3 weeks ago

correct answer

upvoted 2 times

✉️  **mmaguero** 12 months ago

ChatGPT agree:

Answer Area:

Sign in to the Azure Video Analyzer for Media website.

From Content model customization, select Brands.

Add the specific company names to the include list.

upvoted 3 times

Question #27

You have a mobile app that manages printed forms.

You need the app to send images of the forms directly to Forms Recognizer to extract relevant information. For compliance reasons, the image files must not be stored in the cloud.

In which format should you send the images to the Form Recognizer API endpoint?

- A. raw image binary
- B. form URL encoded
- C. JSON

Correct Answer: A

Community vote distribution

A (100%)

✉  **rdemontis**  6 months, 3 weeks ago

Selected Answer: A

To send images to the Form Recognizer API endpoint without storing them in the cloud, you should send the images in the following format:

- A. raw image binary

Sending the images as raw image binary data allows you to transmit the image directly to the Form Recognizer API without the need to store them in the cloud or convert them into other formats. This format ensures compliance with your requirements. (ChatGPT)
upvoted 6 times

✉  **trashbox**  7 months ago

Appeared on Oct/29/2023

upvoted 5 times

✉  **nanaw770**  2 days, 21 hours ago

Selected Answer: A

A is right.

upvoted 1 times

✉  **takaimomoGcup** 1 week ago

Selected Answer: A

A is right answer.

upvoted 1 times

✉  **ankitdhir** 5 months ago

Selected Answer: A

its correct

upvoted 2 times

✉  **M25** 9 months ago

Selected Answer: A

A. raw image binary

<https://westus.dev.cognitive.microsoft.com/docs/services/form-recognizer-api-v2-1/operations/AnalyzeReceiptAsync>

Request body: Document containing the receipt image(s) to be analyzed. The POST body should be the raw image binary, or the image URL in JSON.

<https://ittichaicham.com/2020/03/call-azure-form-recognizer-api-on-sharepoint-document-image-url-in-power-automate/>

Power Automate (formerly Microsoft Flow) can call Azure Form Recognizer via the connector. Since Power Automate is a cloud solution, the natural choice is to use the image URL. This should work fine if the URL is accessible to the public or requires no authentication. Unfortunately, the company's SharePoint URL, most of the time, is not.

To solve this, we can add another flow step to move the SharePoint file to where it is accessible, or, better, instead of using file URL, we can pass binary content in the Form Recognizer API.

upvoted 1 times

✉  **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/faq?view=form-recog-3.0.0#does-form-recognizer-store-my-data->

For all features, Form Recognizer temporarily stores data and results in Azure storage in the same region as the request. Your data is then deleted within 24 hours from the time an analyze request was submitted.

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

thanks for the provided document

upvoted 2 times

 **ziggy1117** 11 months, 4 weeks ago

Should be A. When you send images to the endpoint, the images don't get stored anywhere.

upvoted 1 times

 **hens** 1 year ago

chat gpt "To send images directly to Forms Recognizer and extract relevant information without storing the image files in the cloud, you should use the raw image binary format."

upvoted 2 times

 **Rob77** 1 year ago

Looks like URL (not sure about the "encoded" part though!)

<https://learn.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/v3-migration-guide?view=form-recog-3.0.0#analyze-request-body>

upvoted 1 times

Question #28

You plan to build an app that will generate a list of tags for uploaded images. The app must meet the following requirements:

- Generate tags in a user's preferred language.
- Support English, French, and Spanish.
- Minimize development effort.

You need to build a function that will generate the tags for the app.

Which Azure service endpoint should you use?

- A. Content Moderator Image Moderation
- B. Custom Vision image classification
- C. Computer Vision Image Analysis
- D. Custom Translator

Correct Answer: B

Community vote distribution

C (100%)

✉️  **MaliSanFuu**  1 year, 1 month ago

Selected Answer: C

I think the answer should be C, because of the minimized development effort. Since the prebuilt model of C also fits the other two requirements, so there is no need to train a custom model.

source: <https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/how-to/call-analyze-image?tabs=rest>
upvoted 10 times

✉️  **rdemontis** 6 months, 3 weeks ago

agree with you
upvoted 1 times

✉️  **nanaw770**  2 days, 21 hours ago

Selected Answer: C

C is right.
upvoted 1 times

✉️  **takaimomoGcup** 1 week ago

Selected Answer: C

C. Computer Vision Image Analysis
upvoted 1 times

✉️  **evangelist** 3 months, 1 week ago

generate tags as per image is an Image Analysis operation rather than an image classification operation
upvoted 2 times

✉️  **evangelist** 4 months ago

Selected Answer: C

Here's why:

Multilingual Tag Generation: Azure's Computer Vision Image Analysis service can analyze images and provide a list of tags describing the content of the images. It also has the capability to return these tags in various languages, including English, French, and Spanish, which aligns with your requirement.

Minimizing Development Effort: This service offers a pre-built model, which means there is no need for you to collect data and train your own model. This significantly reduces the development effort and time. You simply need to call the API with your images, and it will return the tags.
upvoted 3 times

✉️  **ankitdhir** 5 months ago

Selected Answer: C

Verified with google bard
upvoted 2 times

- ✉️ **anto69** 3 months, 3 weeks ago
C, verified with Copilot too
upvoted 1 times
- ✉️ **TRUESON** 5 months ago
Selected Answer: B
French is not supported by computer vision
upvoted 2 times
- ✉️ **Mehe323** 2 months, 3 weeks ago
French is supported, I just tested it in Azure. See also this link (the table 'Analyze image':
<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/language-support#image-analysis>
upvoted 1 times
- ✉️ **josebernabeo** 4 months, 1 week ago
True.
- URL parameter Value Description
language en English
language es Spanish
language ja Japanese
language pt Portuguese
language zh Simplified Chinese
- Source: <https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/how-to/call-analyze-image?tabs=rest>
upvoted 1 times
- ✉️ **rdemontis** 6 months, 3 weeks ago
Selected Answer: C
C to minimize the effort:
<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/overview-image-analysis?tabs=4-0>
upvoted 1 times
- ✉️ **zellck** 11 months ago
Selected Answer: C
C is the answer.
- <https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-tagging-images>
Image Analysis can return content tags for thousands of recognizable objects, living beings, scenery, and actions that appear in images. Tags are not organized as a taxonomy and do not have inheritance hierarchies. A collection of content tags forms the foundation for an image description displayed as human readable language formatted in complete sentences. When tags are ambiguous or not common knowledge, the API response provides hints to clarify the meaning of the tag in context of a known setting.
upvoted 3 times
- ✉️ **zellck** 11 months ago
<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/language-support#image-analysis>
Some features of the Analyze - Image API can return results in other languages, specified with the language query parameter. Other actions return results in English regardless of what language is specified, and others throw an exception for unsupported languages. Actions are specified with the visualFeatures and details query parameters; see the Overview for a list of all the actions you can do with image analysis. Languages for tagging are only available in API version 3.2 or later.
upvoted 4 times
- ✉️ **rdemontis** 6 months, 2 weeks ago
thanks for explanation
upvoted 1 times
- ✉️ **Pixelmate** 11 months ago
Asked in 28/06/2023 exam
upvoted 4 times
- ✉️ **ziggy1117** 11 months, 3 weeks ago
C: <https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-tagging-images>
upvoted 1 times
- ✉️ **hens** 1 year ago
chat gpt "To generate a list of tags for uploaded images in multiple languages, you should use the Computer Vision Image Analysis service endpoint in Azure."
Computer Vision provides a pre-built model that can generate image tags based on a given image. It also supports multiple languages, including English, French, and Spanish, which meets the requirements of the app. Additionally, the service provides a REST API, which can be easily integrated into your app without requiring significant development effort."
upvoted 1 times

Question #29

HOTSPOT

You develop a test method to verify the results retrieved from a call to the Computer Vision API. The call is used to analyze the existence of company logos in images. The call returns a collection of brands named brands.

You have the following code segment.

```
foreach (var brand in brands)
{
    if (brand.Confidence >= .75)
        Console.WriteLine($"Logo of {brand.Name} between {brand.Rectangle.X}, {brand.Rectangle.Y} and {brand.Rectangle.W},
{brand.Rectangle.H}");
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
The code will display the name of each detected brand with a confidence equal to or higher than 75 percent.	<input type="radio"/>	<input type="radio"/>
The code will display coordinates for the top-left corner of the rectangle that contains the brand logo of the displayed brands.	<input type="radio"/>	<input type="radio"/>
The code will display coordinates for the bottom-right corner of the rectangle that contains the brand logo of the displayed brands.	<input type="radio"/>	<input type="radio"/>

Answer Area	
Correct Answer:	Statements
	The code will display the name of each detected brand with a confidence equal to or higher than 75 percent.
	<input checked="" type="radio"/> The code will display coordinates for the top-left corner of the rectangle that contains the brand logo of the displayed brands.
	<input type="radio"/> The code will display coordinates for the bottom-right corner of the rectangle that contains the brand logo of the displayed brands.

 zellck  11 months ago

YYN is the answer.

<https://learn.microsoft.com/en-us/rest/api/computervision/3.1/detect-objects/detect-objects?tabs=HTTP#boundingrect>
A bounding box for an area inside an image.

- X
X-coordinate of the top left point of the area, in pixels.
- y
Y-coordinate of the top left point of the area, in pixels.
- h
Height measured from the top-left point of the area, in pixels.
- w
Width measured from the top-left point of the area, in pixels.
upvoted 8 times

 takaimomoGcup  1 week ago

Yes Yes No
upvoted 1 times

 evangelist 4 months ago

answer is correct: Y Y N
response only display the top left corner and width and height from this origin
upvoted 2 times

- ✉ **rdemontis** 6 months, 3 weeks ago
YYN: correct answer and duplicated question
upvoted 3 times
- ✉ **jangotango** 8 months, 2 weeks ago
YYY - the last one is Y because all coordinates together give you top left and bottom right.
upvoted 1 times
- ✉ **AnonymousJhb** 8 months, 1 week ago
YYN. not the bottom right
upvoted 2 times
- ✉ **rdemontis** 6 months, 3 weeks ago
no, you can calculate the bottom right corner coordinates but they aren't displayed in the code provided
upvoted 1 times
- ✉ **ExamDev** 8 months, 2 weeks ago
Answer is correct
upvoted 1 times
- ✉ **zellck** 11 months ago
Same as Question 14.
<https://www.examtopics.com/discussions/microsoft/view/55050-exam-ai-102-topic-2-question-14-discussion>
upvoted 4 times
- ✉ **Rob77** 1 year ago
Correct YYN. The last one is width and height.
upvoted 1 times

Question #30

DRAG DROP

You have a factory that produces cardboard packaging for food products. The factory has intermittent internet connectivity.

The packages are required to include four samples of each product.

You need to build a Custom Vision model that will identify defects in packaging and provide the location of the defects to an operator. The model must ensure that each package contains the four products.

Which project type and domain should you use? To answer, drag the appropriate options to the correct targets. Each option may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Options	Answer Area
Food	Project type: <input type="text"/>
General	Domain: <input type="text"/>
General (compact)	
Image classification	
Logo	
Object detection	

Answer Area
Correct Answer: Project type: Object detection Domain: General (compact)

 **zellck**  11 months ago

1. Object detection
2. General (compact)

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/get-started-build-detector>
- Select Object Detection under Project Types.

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/select-domain#compact-domains>
The models generated by compact domains can be exported to run locally.
upvoted 13 times

 **dazdzadzadzaazd** 10 months, 1 week ago

1. I was about to say Classification, but they say "and provide the location of the defects" so it is definitively Object Detection.
 2. General (compact)
- upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

agree with you
upvoted 2 times

✉  **rveney**  11 months, 1 week ago

This was on my exam
upvoted 7 times

✉  **nanaw770**  2 days, 21 hours ago

1. Object detection
2. General(compact)
upvoted 1 times

✉  **takaimomoGcup** 1 week ago

Project type is Object detection. Domain is General(compact).
upvoted 1 times

✉  **evangelist** 3 months, 1 week ago

Icoally==>mean general(compact)
detect =>object detection
upvoted 2 times

✉  **evangelist** 4 months ago

The factory has intermittent internet connectivity:====> this means an edge deployment of the model without internet connectivity is needed
and then edge model using General(compact) domain suits the demands
upvoted 3 times

✉  **ExamDev** 8 months, 2 weeks ago

The answer is correct!
upvoted 1 times

✉  **Pixelmate** 11 months ago

Asked in 28/06/2023 exam
upvoted 5 times

✉  **Rob77** 1 year ago

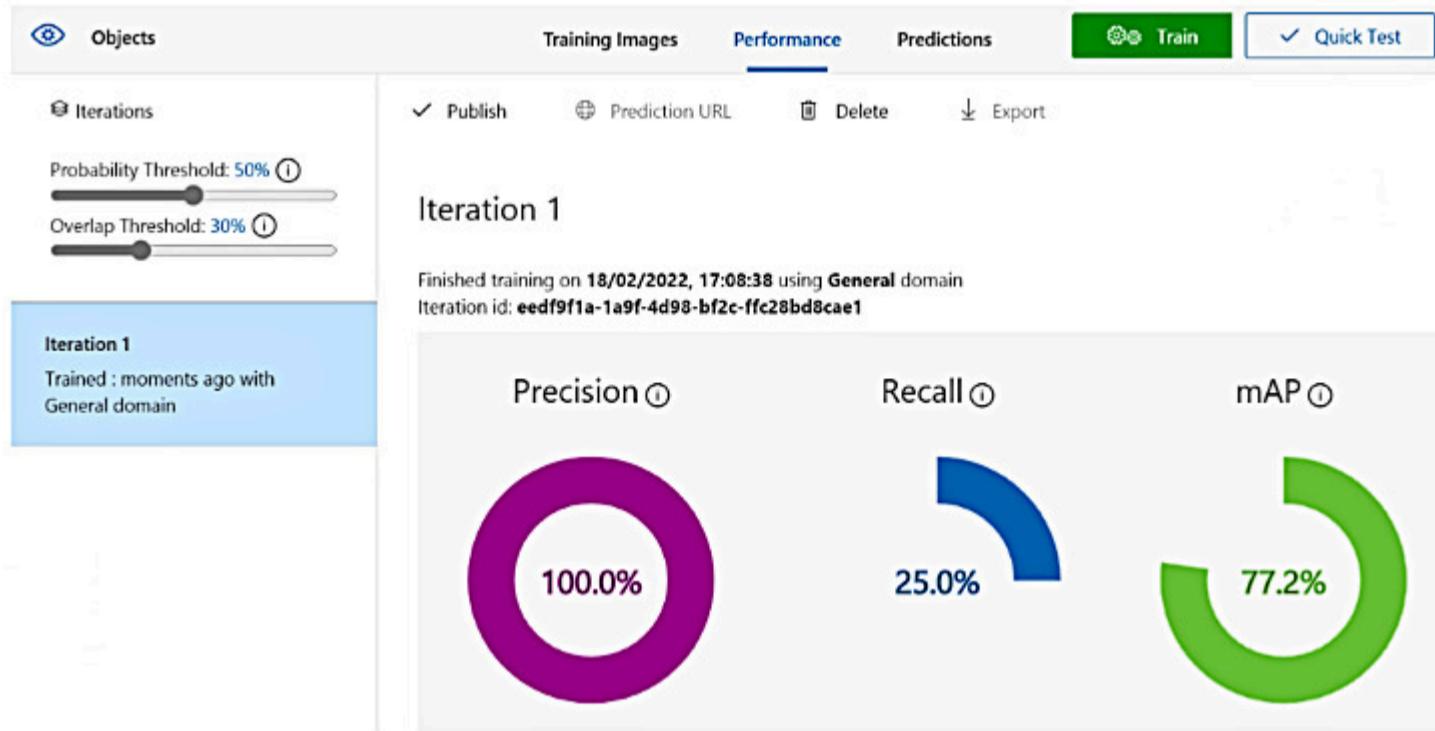
Correct - <https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/select-domain#compact-domains>
upvoted 3 times

Question #31

HOTSPOT

You are building a model to detect objects in images.

The performance of the model based on training data is shown in the following exhibit.



Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Answer Area

The percentage of false positives is [answer choice].

- 0
- 25
- 30
- 50
- 100

The value for the number of true positives divided by the total number of true positives and false negatives is [answer choice] %.

- 0
- 25
- 30
- 50
- 100

Answer Area

The percentage of false positives is [answer choice].

- 0
- 25
- 30
- 50
- 100

Correct Answer:

The value for the number of true positives divided by the total number of true positives and false negatives is [answer choice] %.

- 0
- 25
- 30
- 50
- 100

 zellck Highly Voted  11 months ago

1. 0
2. 25

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/custom-text-classification/concepts/evaluation-metrics>

- Precision: Measures how precise/accurate your model is. It's the ratio between the correctly identified positives (true positives) and all identified positives. The precision metric reveals how many of the predicted classes are correctly labeled.

$$\text{Precision} = \frac{\# \text{True Positive}}{\# \text{True Positive} + \# \text{False Positive}}$$

- Recall: Measures the model's ability to predict actual positive classes. It's the ratio between the predicted true positives and what was actually

tagged. The recall metric reveals how many of the predicted classes are correct.

Recall = #True_Positive / (#True_Positive + #False_Negatives)

upvoted 13 times

✉  **rdemontis** 6 months, 3 weeks ago

thanks for explanation

upvoted 2 times

✉  **Pixelmate**  11 months ago

Asked in 28/06/2023 exam

upvoted 5 times

✉  **takaimomoGcup**  2 days, 21 hours ago

1. 0

2. 25

upvoted 1 times

✉  **Tin_Tin** 11 months, 2 weeks ago

The answer is correct.

See <https://learn.microsoft.com/en-us/azure/cognitive-services/Custom-Vision-Service/get-started-build-detector>

upvoted 1 times

✉  **973b658** 11 months, 2 weeks ago

It is true.

#1:Precision = 100%

#2:recall = 25%

upvoted 1 times

Question #32

You are building an app that will include one million scanned magazine articles. Each article will be stored as an image file.

You need to configure the app to extract text from the images. The solution must minimize development effort.

What should you include in the solution?

- A. Computer Vision Image Analysis
- B. the Read API in Computer Vision
- C. Form Recognizer
- D. Azure Cognitive Service for Language

Correct Answer: A

Community vote distribution



✉️ **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

To me the correct answer is B. With the new Image Analysis API 4.0 (in preview) you could use OCR feature too, but as i said, it is in preview. And i don't think it is considered in the exam.

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/overview-image-analysis?tabs=4-0>

Instead the Read API is particularly adapted for text-heavy documents and it seems this is the case

<https://learn.microsoft.com/en-us/rest/api/computervision/read/read?view=rest-computervision-v3.1&tabs=HTTP>
upvoted 9 times

✉️ **jangotango** 7 months, 3 weeks ago

All answers should have a reference to prove the answer is true
upvoted 6 times

✉️ **takaimomoGcup** 2 days, 21 hours ago

Selected Answer: B

the Read API in Computer Vision is right.
upvoted 1 times

✉️ **Compras** 1 month, 1 week ago

Concordo com o @jangotango sobre as referências. Tem que evidenciar.
upvoted 1 times

✉️ **Murtuza** 1 month, 2 weeks ago

Selected Answer: B

Read API is particularly adapted for text-heavy documents and it seems this is the case
upvoted 1 times

✉️ **AlviraTony** 2 months ago

Selected Answer: C

Document Intelligence(Previously Form Recognizer) reads small to large volume of text from images and PDF documents. For example: receipts, articles, and invoices
upvoted 3 times

No I don't think so. Azure AI Document Intelligence is a more sophisticated solution. For example, it can identify key/value pairs, tables, and context-specific fields. If you want to deploy a complete document analysis solution for both extracting AND understanding text, Azure AI Document Intelligence is a good solution. In this case, Document Intelligence is a too advanced solution as the question doesn't provide any information about what to do with the extracted text, the focus here is on text extraction alone. So the answer is B.

upvoted 1 times

✉️ **devilsole** 7 months ago

i think it should be B
based on ChatGpt
using Azure Read API if:

Your primary focus is on text extraction from documents, forms, or images with printed text.

You have a batch processing requirement, and you need to process a large number of documents or images at once.

You need highly accurate text extraction with structured output.

upvoted 2 times

 **JDKJDKJDK** 7 months, 3 weeks ago

Selected Answer: B

i also think its B

Use this interface to get the result of a Read operation, employing the state-of-the-art Optical Character Recognition (OCR) algorithms optimized for text-heavy documents.

<https://learn.microsoft.com/en-us/rest/api/computervision/3.2preview2/read/read?tabs=HTTP>

upvoted 3 times

 **jangotango** 7 months, 3 weeks ago

Why not B?

upvoted 1 times

Question #33

You have a 20-GB video file named File1.avi that is stored on a local drive.

You need to index File1.avi by using the Azure Video Indexer website.

What should you do first?

- A. Upload File1.avi to an Azure Storage queue.
- B. Upload File1.avi to the Azure Video Indexer website.
- C. Upload File1.avi to Microsoft OneDrive.
- D. Upload File1.avi to the www.youtube.com webpage.

Correct Answer: B

Community vote distribution

C (78%)

B (22%)

✉️  **suryakalla**  7 months, 3 weeks ago

This question is part of free assessment given by Microsoft and the answer in that was C.
upvoted 13 times

✉️  **[Removed]** 1 month, 2 weeks ago

I saw it there too :)
upvoted 1 times

✉️  **Student2023** 7 months, 3 weeks ago

because that is the correct option
upvoted 1 times

✉️  **takaimomoGcup**  2 days, 21 hours ago

Selected Answer: C
C is right answer.
upvoted 1 times

✉️  **TJ001** 1 week, 1 day ago

Max file size for direct upload is 2 GB . 30 GB is through url.so answer C
upvoted 2 times

✉️  **Jimmy1017** 1 month ago

B. Upload File1.avi to the Azure Video Indexer website.

Explanation:

The Azure Video Indexer website is specifically designed to analyze and index video files, extracting insights such as keywords, faces, sentiments, and more.

Uploading File1.avi directly to the Azure Video Indexer website allows the platform to process the video file and generate the necessary metadata and insights.

Options A, C, and D are not relevant for indexing File1.avi using the Azure Video Indexer website. Uploading to Azure Storage queue (option A) is not appropriate for indexing videos. Microsoft OneDrive (option C) is a cloud storage service and doesn't provide video indexing capabilities like the Azure Video Indexer. Uploading to YouTube (option D) is also not relevant as the task is to index the video using the Azure Video Indexer website. Therefore, option B is the correct choice.

upvoted 1 times

✉️  **[Removed]** 1 month, 2 weeks ago

Selected Answer: C
Yep, uploading to OneDrive and creating a download link (if I recall) was the answer on the MS free assessment.
upvoted 1 times

✉️  **Murtuza** 1 month, 2 weeks ago

Selected Answer: C
C is correct
upvoted 2 times

✉️  **Mehe323** 2 months ago

Selected Answer: C

Upload file size and video duration

If uploading a file from your device, the file size limit is 2 GB.

If the video is uploaded from a URL, the file size limit is 30 GB. The URL must lead to an online media file with a media file extension (for example myvideo.MP4) and not a webpage such as <https://www.youtube.com>.

The file duration limit is 4 hours.

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/avi-support-matrix>

upvoted 3 times

 **varinder82** 2 months ago

Final Answer:

C

upvoted 1 times

 **schmoofed** 3 months, 1 week ago

Selected Answer: C

Looks like C is correct due to the large size of the file being 20GB in the question. See article here: <https://learn.microsoft.com/en-us/azure/azure-video-indexer/odrv-download>

upvoted 2 times

 **evangelist** 3 months, 1 week ago

Selected Answer: B

B is correct and B

upvoted 1 times

 **suzanne_exam** 4 months ago

C: There is a max byte array of 2gb when uploading direct to the indexer so it has to be via a url

upvoted 2 times

 **MelMac** 5 months, 2 weeks ago

Selected Answer: C

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/considerations-when-use-at-scale>

upvoted 1 times

 **ccampagna** 6 months ago

Selected Answer: C

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/odrv-download?tabs=With-classic-account>

As rdemontis explained , only 2GB files are accepted to upload directly to ACI. So the correct answer should be load the file in OneDrive and generate an URL

upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: C

<https://learn.microsoft.com/it-it/azure/azure-video-indexer/considerations-when-use-at-scale>

upvoted 3 times

 **rdemontis** 6 months, 3 weeks ago

When uploading videos consider using a URL over byte array.

Azure AI Video Indexer does give you the choice to upload videos from URL or directly by sending the file as a byte array, the latter comes with some constraints. For more information, see [uploading considerations and limitations](#))

First, it has file size limitations. The size of the byte array file is limited to 2 GB compared to the 30-GB upload size limitation while using URL.

<https://learn.microsoft.com/it-it/azure/azure-video-indexer/considerations-when-use-at-scale>

We can definitely say that the correct answer is C

upvoted 3 times

 **DEXTER1022** 7 months ago

B IS THE CORRECT ANSWER

upvoted 2 times

 **chenglim** 7 months ago

Selected Answer: C

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/upload-index-videos>

upvoted 2 times

Question #34

HOTSPOT

You are building an app that will share user images.

You need to configure the app to meet the following requirements:

- Uploaded images must be scanned and any text must be extracted from the images.
- Extracted text must be analyzed for the presence of profane language.
- The solution must minimize development effort.

What should you use for each requirement? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Text extraction:

Azure AI Language
Azure AI Computer Vision
Content Moderator
Azure AI Custom Vision
Azure AI Document Intelligence

Profane language detection:

Azure AI Language
Azure AI Computer Vision
Content Moderator
Azure AI Custom Vision
Azure AI Document Intelligence

Answer Area

Text extraction:

Azure AI Language
Azure AI Computer Vision
Content Moderator
Azure AI Custom Vision
Azure AI Document Intelligence

Correct Answer:

Profane language detection:

Azure AI Language
Azure AI Computer Vision
Content Moderator
Azure AI Custom Vision
Azure AI Document Intelligence

 takaimomoGcup 2 days, 21 hours ago

1. Azure AI Document Intelligence
2. Content Moderator

upvoted 1 times

✉️ **funny_penguin** 3 days, 23 hours ago

on exam but with AI Content Safety instead of content moderator. For the first, I chose Azure AI Document Intelligence
upvoted 1 times

✉️ **ShardulShende** 4 weeks ago

Why not Azure AI Document Intelligence for the first option?

upvoted 2 times

✉️ **ShardulShende** 4 weeks ago

Should be Azure AI Document Intelligence. Document Intelligence Read Optical Character Recognition (OCR) model runs at a higher resolution than Azure AI Vision Read and extracts print and handwritten text from PDF documents and scanned images.

<https://learn.microsoft.com/en-us/answers/questions/1512283/vision-studio-vs-document-intelligence-studio-ocr>

upvoted 1 times

✉️ **Murtuza** 1 month, 4 weeks ago

Text Extraction from Images: You can use an Optical Character Recognition (OCR) service to scan and extract text from the uploaded images. Computer Vision API are examples of services that provide OCR capabilities.

Profanity Check: Once the text is extracted, you can use a text analytics service to analyze the presence of profane language. Services Content Moderator API can help identify and filter out inappropriate content.

upvoted 3 times

✉️ **audlindr** 2 months, 4 weeks ago

I hope the Questions are revisited by Microsoft. Why would there be a question on deprecated services

<https://learn.microsoft.com/en-us/azure/ai-services/content-moderator/overview>

Azure Content Moderator is being deprecated in February 2024, and will be retired by February 2027. It is being replaced by Azure AI Content Safety, which offers advanced AI features and enhanced performance.

upvoted 4 times

✉️ **[Removed]** 2 months, 2 weeks ago

I agree, insert profanity here to show my frustration! AND YOU CAN USE AI Content Safety to detect it :D

upvoted 5 times

Question #35

You are building an app that will share user images.

You need to configure the app to perform the following actions when a user uploads an image:

- Categorize the image as either a photograph or a drawing.
- Generate a caption for the image.

The solution must minimize development effort.

Which two services should you include in the solution? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. object detection in Azure AI Computer Vision
- B. content tags in Azure AI Computer Vision
- C. image descriptions in Azure AI Computer Vision
- D. image type detection in Azure AI Computer Vision
- E. image classification in Azure AI Custom Vision

Correct Answer: CE

Community vote distribution



✉ **audlindr** 2 months, 4 weeks ago

Selected Answer: CD

I think Categorize the image as either a photograph or a drawing should be Image type detection
<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/concept-detecting-image-types>

Image Categorization doesn't identify image as photograph or drawing: See this: <https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/concept-categorizing-images>

Captions are generated using image descriptions in V3.2. However in V4.0 it is image captions
<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/concept-describing-images>
upvoted 8 times

✉ **takaimomoGcup** 2 days, 21 hours ago

Selected Answer: CD

image descriptions and image type detection.
upvoted 1 times

✉ **reiwanotora** 1 week, 1 day ago

Selected Answer: CD

It must be C and D.
upvoted 1 times

✉ **Jimmy1017** 1 month ago

A. Object detection in Azure AI Computer Vision
C. Image descriptions in Azure AI Computer Vision

Explanation:

Object Detection (option A): This service in Azure AI Computer Vision can be used to detect objects within an image. By analyzing the content of the image, it can identify whether the image contains elements typically found in photographs (e.g., people, landscapes) or drawings (e.g., sketches, illustrations). This information can help categorize the image accordingly.

Image Descriptions (option C): Azure AI Computer Vision can generate descriptions for images, providing textual summaries of the content. These descriptions can include details about the objects detected in the image, providing additional context that can aid in categorization and caption generation.

upvoted 1 times

✉ **Compras** 1 month, 1 week ago

C e E Sim! E para classificar como imagem e C para descrever a imagem.

upvoted 1 times

 **chandiochan** 1 month, 3 weeks ago

Selected Answer: CD

Must be C & D

upvoted 1 times

 **haverner** 2 months ago

Selected Answer: CE

CE, Categorize instantly becomes Classification

upvoted 2 times

 **Harry300** 2 months, 3 weeks ago

Selected Answer: CD

CD imageType for photograph/drawing

upvoted 1 times

 **Harry300** 2 months, 3 weeks ago

Selected Answer: CD

Should be CD

upvoted 1 times

Question #36

You are building an app that will use the Azure AI Video Indexer service.

You plan to train a language model to recognize industry-specific terms.

You need to upload a file that contains the industry-specific terms.

Which file format should you use?

- A. XML
- B. TXT
- C. XLS
- D. PDF

Correct Answer: B

Community vote distribution

 B (100%)

 **audlindr**  2 months, 4 weeks ago

Selected Answer: B

Answer is correct:

As per <https://learn.microsoft.com/en-us/azure/azure-video-indexer/customize-language-model-with-website>

This step creates the model and gives the option to upload text files to the model.

upvoted 6 times

 **takaimomoGcup**  2 days, 21 hours ago

Selected Answer: B

B is correct.

upvoted 1 times

 **reiwanotora** 1 week, 1 day ago

Selected Answer: B

B is right.

upvoted 1 times

 **pepe54362** 1 week, 4 days ago

Selected Answer: B

I agree, B is the option

upvoted 1 times

Question #37

DRAG DROP

You have an app that uses Azure AI and a custom trained classifier to identify products in images.

You need to add new products to the classifier. The solution must meet the following requirements:

- Minimize how long it takes to add the products.
- Minimize development effort.

Which five actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

- Label the sample images.
- From Vision Studio, open the project.
- Publish the model.
- From the Custom Vision portal, open the project.
- Retrain the model.
- Upload sample images of the new products.
- From the Azure Machine Learning studio, open the workspace.

Answer Area

Answer Area	
From Vision Studio, open the project.	
Upload sample images of the new products.	
Label the sample images.	
Retrain the model.	
Publish the model.	

Correct Answer:

Answer Area	
From Vision Studio, open the project.	
Upload sample images of the new products.	
Label the sample images.	
Retrain the model.	
Publish the model.	

Harry300 Highly Voted 2 months, 3 weeks ago

First step should be Custom Vision.
Then Upload sample images / Label / Retrain / Publish
Custom classifiers have to go through custom vision, obviously
upvoted 12 times

audlindr 2 months, 3 weeks ago

As per this <https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/concept-model-customization>

You can train a custom model using either the Custom Vision service or the Image Analysis 4.0 service with model customization.
upvoted 1 times

takaimomoGcup Most Recent 2 days, 21 hours ago

1. From the Custom Vision portal, open the project
2. Label
3. Upload
- 4 .Retrain
5. Publish

upvoted 1 times

reiwanotora 1 week, 1 day ago

Will this question be on the actual exam?
upvoted 1 times

aks_exam 2 weeks, 2 days ago

So I will choice,
1. From the Custom Vision portal, open the project
2. Upload sample images of the new products
3. Label the samples images
4. Retrain the model
5. Publish the model

labelling after uploading sample images.

upvoted 2 times

✉ **Murtuza** 1 month, 4 weeks ago

To add new products to the classifier while minimizing time and development effort, you should perform the following actions in sequence:

From the Custom Vision portal, open the project.

Upload sample images of the new products.

Label the sample images.

Retrain the model.

Publish the model.

This sequence ensures that the new product images are properly added, labeled, and incorporated into the existing model, and that the updated model is made available for use by your application.

upvoted 1 times

✉ **varinder82** 2 months ago

Final Answer:

1. From the Custom Vision portal, open the project

2. Label the samples images

3. Upload sample images of the new products

4 . Retrain the model

5. Publish the model

upvoted 1 times

✉ **Murtuza** 2 months, 1 week ago

4) Retrain the Model:

Trigger the retraining process for your custom classifier. This step involves:

Using the labeled samples to train the model.

Fine-tuning the existing classifier with the new data.

5) Publish the Model:

Once the retraining is complete and the model performs well on validation data, publish the updated classifier.

The published model will be ready for inference in your application, allowing it to identify the new products.

upvoted 2 times

✉ **Murtuza** 2 months, 1 week ago

Your proposed sequence of actions for adding new products to the classifier is a good start! Let's refine it a bit to ensure it aligns with best practices:

1) Open the Custom Vision Project:

Begin by accessing your project in the Custom Vision portal. This is where you'll manage your custom-trained classifier.

2) Label the Sample Images:

Next, label the sample images you've collected for the new products. Assign appropriate tags or classes to each image based on the product category.

Proper labeling is crucial for effective training.

3) Upload Sample Images:

Upload the labeled sample images to your project. These images will serve as the training data for your classifier.

Make sure you have a diverse set of samples to represent different variations of the new products.

upvoted 2 times

✉ **arcameon** 2 months, 3 weeks ago

According to chat GPT, the correct sequence is the following :

1. From the Custom Vision portal, open the project

2. Label the samples images

3. Upload sample images of the new products

4 . Retrain the model

5. Publish the model

Labeling the sample images before uploading them is crucial because it helps in structuring and organizing the dataset appropriately.

upvoted 3 times

✉ **varinder82** 2 months ago

2. Label the samples images

3. Upload sample images of the new products

Wrong, How you can label before upload so it should be Upload and then label

upvoted 2 times

Question #38

HOTSPOT

You are developing an application that will use the Azure AI Vision client library. The application has the following code.

```
def analyze_image(local_image):
    with open(local_image, "rb") as image_stream:
        image_analysis = client.analyze_image_in_stream(
            image=image_stream,
            visual_features=[
                VisualFeatureTypes.tags,
                VisualFeatureTypes.description
            ]
        )
        for caption in image_analysis.description.captions:
            print(f"\n{caption.text} with confidence {caption.confidence}")
        for tag in image_analysis.tags:
            print(f"\n{tag.name} with confidence {tag.confidence}")
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

	Statements	Yes	No
	The code will perform face recognition.	<input type="radio"/>	<input type="radio"/>
	The code will list tags and their associated confidence.	<input type="radio"/>	<input type="radio"/>
	The code will read an image file from the local file system.	<input type="radio"/>	<input type="radio"/>

Answer Area

	Statements	Yes	No
Correct Answer:	The code will perform face recognition.	<input type="radio"/>	<input checked="" type="checkbox"/>
	The code will list tags and their associated confidence.	<input checked="" type="checkbox"/>	<input type="radio"/>
	The code will read an image file from the local file system.	<input checked="" type="checkbox"/>	<input type="radio"/>

✉  **takaimomoGcup** 2 days, 21 hours ago

No

Yes

Yes

upvoted 1 times

✉  **reiwanotora** 1 week, 1 day ago

No Yes Yes

upvoted 1 times

✉  **Murtuza** 1 month, 4 weeks ago

Here are the answers to your questions:

The code will perform face recognition. No, the code does not include any features related to face recognition. It is only analyzing the image for tags and descriptions.

The code will list tags and their associated confidence. Yes, the code includes a loop that prints each tag and its associated confidence level.

The code will read an image file from the local file system. Yes, the code opens an image file from the local file system for analysis.

upvoted 2 times

✉  **Murtuza** 2 months, 1 week ago

The given answers by exam topics are CORRECT

upvoted 2 times

Topic 3 - Question Set 3**Question #1****Topic 3**

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution. After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen. You build a language model by using a Language Understanding service. The language model is used to search for information on a contact list by using an intent named FindContact.

A conversational expert provides you with the following list of phrases to use for training.

- Find contacts in London.
- Who do I know in Seattle?
- Search for contacts in Ukraine.

You need to implement the phrase list in Language Understanding.

Solution: You create a new pattern in the FindContact intent.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

Instead use a new intent for location.

Note: An intent represents a task or action the user wants to perform. It is a purpose or goal expressed in a user's utterance.

Define a set of intents that corresponds to actions users want to take in your application.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-intent>

Community vote distribution

A (57%) B (43%)

 **azurelearner666** Highly Voted  2 years, 11 months ago

Using a pattern could be a good solution IMHO...

- Find contacts in London.
- Who do I know in Seattle?
- Search for contacts in Ukraine.

Like

Where is {FormName}[]?

Who authored {FormName}[]?

{FormName} is published in French[]?

(taken from <https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-patterns>)

we could do:

- Find contacts in {CityOrCountry}.
- Who do I know in {CityOrCountry}[]?
- Search for contacts in {CityOrCountry}[].

So, to me a pattern is a Solution (A)

upvoted 26 times

 **YipingRuan** 2 years, 10 months ago

Agree, but Entity is also good <https://docs.microsoft.com/bs-cyrl-ba/azure/cognitive-services/luis/luis-concept-intent#intent-compared-to-entity>

upvoted 2 times

 **practicewizards** 1 year, 8 months ago

I agree. The intent here is "search for contact", for example. The location is an entity type "location", as you can see in the example What's the weather like in Seattle tomorrow? on the link given by @YipingRuan

upvoted 1 times

 **STH** Highly Voted  1 year, 8 months ago

Selected Answer: A

According to MS learn, answer should be yes (A)

<https://learn.microsoft.com/en-us/training/modules/create-language-understanding-app/5-use-patterns-to-differentiate-similar-utterances>

This is a FindContact intent with a location entity pattern
upvoted 9 times

✉ **nanaw770** Most Recent 3 days, 18 hours ago

Selected Answer: A

It MUST be A.
upvoted 1 times

✉ **nanaw770** 3 days, 18 hours ago

After all, would you choose Yes or No to be correct?
upvoted 1 times

✉ **Murtuza** 1 month, 2 weeks ago

Selected Answer: A

A is correct
upvoted 1 times

✉ **Murtuza** 1 month, 4 weeks ago

Yes, creating a new pattern in the FindContact intent with the provided phrases can help train the Language Understanding model to better recognize when the user is trying to find contacts in a specific location. This would meet the goal of implementing the phrase list in Language Understanding.

upvoted 1 times

✉ **f2c587e** 2 months ago

This same question (topic 3, question 18) seems to be indicated by the community that this question has two affirmative answers. I'm not sure if this is correct, normally in this type of question there is only one affirmative answer and the rest are negative. Is anyone clear on the real answer to this question?

upvoted 2 times

✉ **trysec** 4 months, 1 week ago

Selected Answer: B

B. No

The proposed solution of creating a new pattern in the FindContact intent does not fully meet the goal of implementing a phrase list in Language Understanding (often referred to as LUIS - Language Understanding Intelligent Service in Azure).

In Azure LUIS, patterns are used to identify specific sentence structures that indicate intents, and they are indeed a valuable part of intent recognition. However, patterns alone are not the same as a phrase list. A phrase list in LUIS is a feature that allows you to define a list of related words or phrases which can be used across various intents and utterances. It's more about providing synonyms or variations of words that help the model understand different ways a user might express the same concept.

upvoted 3 times

✉ **dimsok** 4 months, 2 weeks ago

Selected Answer: B

Its a NO, utterances are not similar
upvoted 1 times

✉ **sl_mslconsulting** 6 months, 2 weeks ago

Selected Answer: B

I am speaking for the newer CLU. The very purpose of using patterns is to differentiate similar utterances but with different intents.
upvoted 4 times

✉ **sl_mslconsulting** 6 months, 2 weeks ago

We only have one intent here. You need to add entity to provide context specific information location to your app to take proper action.
upvoted 1 times

✉ **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

IMHO we are in a case where the utterances are very similar. You are asking questions in a slightly different way to obtain the same result: a list of contact. So i think a pattern could be very useful because the phrase list here is also quite short. And this is a typical scenario where patterns fit the bill
<https://learn.microsoft.com/en-us/azure/ai-services/luis/concepts/patterns-features>
upvoted 1 times

✉ **sl_mslconsulting** 7 months, 2 weeks ago

Selected Answer: B

First off you can't even create a LUIS resource to experiment it. I doubt if this question will ever appear on the exam again. Also these phases are not similar at all so using pattern is out the question. You need entity to be extracted from the utterance to decide what to do next in your app. What do you need? The location so you can retrieve the contacts properly.
upvoted 1 times

✉ **james2033** 9 months, 1 week ago

Selected Answer: A

Intent <https://learn.microsoft.com/en-us/azure/ai-services/luis/concepts/intents>

upvoted 1 times

✉ **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/luis/concepts/patterns-features>

Patterns are designed to improve accuracy when multiple utterances are very similar. A pattern allows you to gain more accuracy for an intent without providing several more utterances.

upvoted 2 times

✉ **zellck** 11 months ago

<https://learn.microsoft.com/en-us/training/modules/build-language-understanding-model/4-use-patterns-differentiate-similar-utterances>

In some cases, a model might contain multiple intents for which utterances are likely to be similar. You can use the pattern of utterances to disambiguate the intents while minimizing the number of sample utterances.

upvoted 1 times

✉ **Pixelmate** 11 months ago

This usecase was asked in 28/06/2023 exam

upvoted 5 times

✉ **hens** 1 year ago

chat gpt "B. No.

Creating a new pattern in the FindContact intent is not the correct approach to implementing a phrase list in Language Understanding. Instead, you should use the Phrase List feature in Language Understanding. The Phrase List feature allows you to define a list of phrases that are relevant to a particular intent. These phrases can then be used to improve the accuracy of the language model for that intent."

upvoted 2 times

✉ **marti_tremblay000** 1 year, 2 months ago

Selected Answer: B

ChatGPT answer :

No, simply creating a new pattern in the FindContact intent is not enough to implement the phrase list in Language Understanding. You need to add the phrases to the intent as training examples so that the language model can learn to recognize them and correctly identify the FindContact intent when a user enters a similar phrase.

To implement the phrase list in Language Understanding, you should add the provided phrases as training examples to the FindContact intent. This will enable the language model to learn the different ways that users may express the intent to find contacts in different locations. You can do this by going to the Language Understanding service and adding the phrases to the FindContact intent's training data.

upvoted 1 times

Question #2

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You develop an application to identify species of flowers by training a Custom Vision model.

You receive images of new flower species.

You need to add the new images to the classifier.

Solution: You add the new images, and then use the Smart Labeler tool.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

The model needs to be extended and retrained.

Note: Smart Labeler generates suggested tags for images. This lets you label a large number of images more quickly when training a Custom Vision model.

Community vote distribution

B (100%)

✉ **TheB** Highly Voted 2 years, 10 months ago

The answer is correct.

upvoted 14 times

✉ **azurelearner666** Highly Voted 2 years, 11 months ago

correct! retraining is necessary!

upvoted 5 times

✉ **rdemontis** Most Recent 6 months, 3 weeks ago

Selected Answer: B

Answer and explanation are correct.

<https://learn.microsoft.com/en-us/azure/ai-services/custom-vision-service/suggested-tags>

upvoted 1 times

✉ **sl_mslconsulting** 7 months, 2 weeks ago

Selected Answer: B

The answer is B because the limitations of the smart labeler: You should only request suggested tags for images whose tags have already been trained on once. Don't get suggestions for a new tag that you're just beginning to train. You are given new images of species that have not been seen by the model how can you expect it to suggest what they are? Also you can train the model right in the smart labeler: check the workflow and the limitations in the doc. <https://learn.microsoft.com/en-us/azure/ai-services/custom-vision-service/suggested-tags>

upvoted 4 times

✉ **james2033** 9 months, 1 week ago

Selected Answer: B

Smart Labeler for suggestion <https://learn.microsoft.com/en-us/azure/ai-services/custom-vision-service/suggested-tags>

upvoted 1 times

✉ **zellck** 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/suggested-tags>

Smart Labeler will generate suggested tags for images. This lets you label a large number of images more quickly when you're training a Custom Vision model.

When you tag images for a Custom Vision model, the service uses the latest trained iteration of the model to predict the labels of new images. It shows these predictions as suggested tags, based on the selected confidence threshold and prediction uncertainty. You can then either confirm or change the suggestions, speeding up the process of manually tagging the images for training.

upvoted 1 times

✉ **hens** 1 year ago

chat gpt "

ChatGPT

A. Yes.

Using the Smart Labeler tool is a valid way to train a Custom Vision model with new images. It allows the user to label images more efficiently by using an active learning approach that selects images that will have the highest impact on the model's performance."

upvoted 2 times

 **ExamGuruBhai** 1 year, 9 months ago

retrain model so answer is B

upvoted 1 times

 **Eltooth** 1 year, 10 months ago

Selected Answer: B

B is correct answer : No.

Instead the model need to be extended and retrained (Udemy answer).

upvoted 4 times

 **hendriktytgatpwc** 2 years, 2 months ago

answer is correct

upvoted 1 times

 **sumanshu** 2 years, 4 months ago

Label + Retrain

upvoted 2 times

Question #3

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You develop an application to identify species of flowers by training a Custom Vision model.

You receive images of new flower species.

You need to add the new images to the classifier.

Solution: You add the new images and labels to the existing model. You retrain the model, and then publish the model.

Does this meet the goal?

A. Yes

B. No

Correct Answer: A

The model needs to be extended and retrained.

Community vote distribution

A (100%)

✉  **azurelearner666**  2 years, 11 months ago

Correct!

uploading, tagging, retraining and publishing the model

upvoted 18 times

✉  **PHD_CHENG**  1 year, 11 months ago

Was on exam 7 Jun 2022

upvoted 5 times

✉  **rdemontis**  6 months, 3 weeks ago

Selected Answer: A

correct answer

upvoted 1 times

✉  **james2033** 9 months, 1 week ago

Selected Answer: A

add the new images and labels to the existing model.

You retrain the model,

and then publish the model.

--> Perfect.

upvoted 1 times

✉  **Eltooth** 1 year, 10 months ago

Selected Answer: A

A is correct answer : Yes

Instead the model need to be extended and retrained (Udemy answer).

Note: Use Smart Labeler to generate suggested tags for images. This lets you label a large number of images more quickly when training a Custom Vision model.

upvoted 1 times

✉  **ppo12** 1 year, 10 months ago

Looks good to me!

upvoted 1 times

Question #4

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution. After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You develop an application to identify species of flowers by training a Custom Vision model.

You receive images of new flower species.

You need to add the new images to the classifier.

Solution: You create a new model, and then upload the new images and labels.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

The model needs to be extended and retrained.

Community vote distribution

B (100%)

✉  **nanaw770** 3 days, 19 hours ago

Selected Answer: B

The boy neither admits nor denies that he told a lie.

upvoted 1 times

✉  **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

Correct. Instead you need to add the new images and labels to the existing model. You retrain the model, and then publish the model

upvoted 2 times

✉  **sl_mslconsulting** 7 months, 2 weeks ago

Selected Answer: B

The answer is B is because the limitations of the smart labeler: You should only request suggested tags for images whose tags have already been trained on once. Don't get suggestions for a new tag that you're just beginning to train. You are given new images of species that have not been seen by the model how can you expect it to suggest what they are? Also you can train the model right in the smart labeler: check the workflow and the limitations in the doc. <https://learn.microsoft.com/en-us/azure/ai-services/custom-vision-service/suggested-tags>

upvoted 1 times

✉  **sl_mslconsulting** 7 months, 2 weeks ago

Oops I meant to answer the question 2 above this one.

upvoted 1 times

✉  **james2033** 9 months, 1 week ago

Selected Answer: B

Need training. Correct answer: No

upvoted 1 times

✉  **Eltooth** 1 year, 10 months ago

Selected Answer: B

B is correct answer : No.

The model needs to be extended and retrained. (Udemy answer)

Note: Use Smart Labeler to generate suggested tags for images. This lets you label a large number of images more quickly when training a Custom Vision model.

upvoted 1 times

✉  **htolajide** 2 years, 8 months ago

Answer is correct, no need to create a new model, the existing one should be extended and retrained

upvoted 4 times

✉  **Rdninja** 2 years, 10 months ago

You don't need to retrain because you created a brand new model

upvoted 1 times

✉  **Messatsu** 2 years, 10 months ago

No. If "You create a new model, and then upload the new images and labels." your model lacks previous images of other flowers. So the answer is correct.

upvoted 6 times

 **YipingRuan** 2 years, 10 months ago

If must, Create and upload the new model, not upload the image..

upvoted 1 times

 **azurelearner666** 2 years, 11 months ago

correct!

response lacks the model retraining...

upvoted 3 times

Question #5

HOTSPOT -

You are developing a service that records lectures given in English (United Kingdom).

You have a method named AppendToTranscriptFile that takes translated text and a language identifier.

You need to develop code that will provide transcripts of the lectures to attendees in their respective language. The supported languages are English, French,

Spanish, and German.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

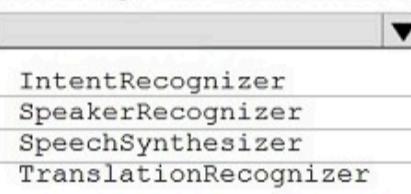
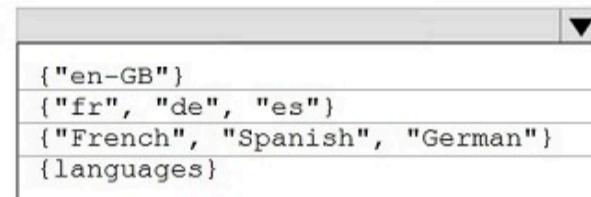
```
static async Task TranslateSpeechAsync()
{
    var config = SpeechTranslationConfig.FromSubscription("69cad5cc-0ab3-4704-bdff-afbf4aa07d85", "uksouth");

    var lang = new List<string>
    {
        {"en-GB"}
        {"fr", "de", "es"}
        {"French", "Spanish", "German"}
        {languages}
    }

    config.SpeechRecognitionLanguage = "en-GB";
    lang.ForEach(config.AddTargetLanguage);

    using var audioConfig = AudioConfig.FromDefaultMicrophoneInput();
    using var recognizer = new
    {
        IntentRecognizer
        SpeakerRecognizer
        SpeechSynthesizer
        TranslationRecognizer
    }

    var result = await recognizer.RecognizeOnceAsync();
    if (result.Reason == ResultReason.TranslatedSpeech)
```

**Correct Answer:****Answer Area**

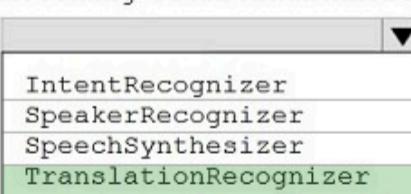
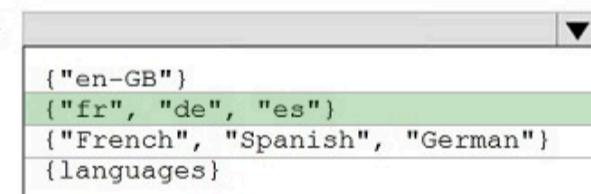
```
static async Task TranslateSpeechAsync()
{
    var config = SpeechTranslationConfig.FromSubscription("69cad5cc-0ab3-4704-bdff-afbf4aa07d85", "uksouth");

    var lang = new List<string>
    {
        {"en-GB"}
        {"fr", "de", "es"} // This row is highlighted in green
        {"French", "Spanish", "German"}
        {languages}
    }

    config.SpeechRecognitionLanguage = "en-GB";
    lang.ForEach(config.AddTargetLanguage);

    using var audioConfig = AudioConfig.FromDefaultMicrophoneInput();
    using var recognizer = new
    {
        IntentRecognizer
        SpeakerRecognizer
        SpeechSynthesizer
        TranslationRecognizer
    }

    var result = await recognizer.RecognizeOnceAsync();
    if (result.Reason == ResultReason.TranslatedSpeech)
```



Box 1: {"fr", "de", "es"}

A common task of speech translation is to specify target translation languages, at least one is required but multiples are supported. The following code snippet sets both French and German as translation language targets. static async Task TranslateSpeechAsync()

```
{
    var translationConfig =
        SpeechTranslationConfig.FromSubscription(SPEECH__SUBSCRIPTION__KEY, SPEECH__SERVICE__REGION);
    translationConfig.SpeechRecognitionLanguage = "it-IT";

    // Translate to languages. See, https://aka.ms/speech/sttt-languages
    translationConfig.AddTargetLanguage("fr");
    translationConfig.AddTargetLanguage("de");
}
```

Box 2: TranslationRecognizer -

After you've created a SpeechTranslationConfig, the next step is to initialize a TranslationRecognizer.

Example code:

```
static async Task TranslateSpeechAsync()
{
    var translationConfig =
        SpeechTranslationConfig.FromSubscription(SPEECH__SUBSCRIPTION__KEY, SPEECH__SERVICE__REGION);
    var fromLanguage = "en-US";
    var toLanguages = new List<string> { "it", "fr", "de" };
    translationConfig.SpeechRecognitionLanguage = fromLanguage;
    toLanguages.ForEach(translationConfig.AddTargetLanguage());
    using var recognizer = new TranslationRecognizer(translationConfig);
}
```

✉  **azurelearner666**  2 years, 11 months ago

Correct!

upvoted 12 times

✉  **idrisfl**  2 years, 10 months ago

Seems correct

<https://docs.microsoft.com/en-us/dotnet/api/microsoft.cognitiveservices.speech.translation.translationrecognizer?view=azure-dotnet>

upvoted 8 times

✉  **audlindr**  3 months ago

correct

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/how-to-translate-speech?tabs=terminal&pivots=programming-language-csharp>

upvoted 1 times

✉  **rdemontis** 6 months, 3 weeks ago

correct answer

<https://docs.microsoft.com/en-us/dotnet/api/microsoft.cognitiveservices.speech.translation.translationrecognizer?view=azure-dotnet>

upvoted 1 times

✉  **ManvalT** 7 months, 3 weeks ago

Correct,got this in Oct2023 exam

upvoted 5 times

✉  **Eltooth** 1 year, 10 months ago

Answer is correct.

("fr", "de", "es")
TranslationRecognizer

upvoted 6 times

Question #6

DRAG DROP -

You train a Custom Vision model used in a mobile app.

You receive 1,000 new images that do not have any associated data.

You need to use the images to retrain the model. The solution must minimize how long it takes to retrain the model.

Which three actions should you perform in the Custom Vision portal? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions	Answer Area
Upload the images by category.	
Get suggested tags.	
Upload all the images.	◀
Group the images locally into category folders.	▶
Review the suggestions and confirm the tags.	
Tag the images manually.	◀ ▶

Correct Answer:

Actions	Answer Area
	Group the images locally into category folders.
Get suggested tags.	Upload the images by category.
Upload all the images.	◀
	Tag the images manually.
Review the suggestions and confirm the tags.	▶

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/getting-started-build-a-classifier>

 SuperPeteY Highly Voted 2 years, 9 months ago

The given answer is incorrect. The question emphasizes two things - 1) the model has already been trained 2) the solution should be expedient. The given answer will be very slow to manually tag 1,000 images. instead:

- 1.) upload all the images
- 2.) Get suggested tags
- 3.) Review the suggestions and confirm the tags

reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/suggested-tags>
upvoted 113 times

 Derin_tade 2 years, 8 months ago

Thank you.
upvoted 3 times

 vominhtri854 2 years, 7 months ago

When you tag images for a Custom Vision model, the service uses the latest trained iteration of the model to predict the labels of untagged images
we need latest trained to predict the labels, but this isn NOT HAVE ANY ASSOCIATED DATA

upvoted 3 times

✉ **rdemontis** 6 months, 3 weeks ago

Exactly. Here we need to use Smart Labeler instead.

<https://learn.microsoft.com/en-us/azure/ai-services/custom-vision-service/suggested-tags>

upvoted 3 times

✉ **STH**  1 year, 8 months ago

Answer is correct.

When uploading all images from a same folder, you can tag all of them with the same value at the same time. Then you wont tag all 1000 images one by one, but only once by category (which is time saving as the question ask for).

Also, even if model is already trained, images are uploaded to workspace, and not to specific trained iteration. You then cannot get tag suggestion when importing an image. There is none, that feature simply does not exist.

Try by yourself :

<https://learn.microsoft.com/en-us/training/modules/classify-images/5-exercise-custom-vision>

upvoted 7 times

✉ **STH** 1 year, 8 months ago

my bad the feature is real :

<https://learn.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/suggested-tags>

so right answer is

- Upload all
- Get suggested tags
- Review and confirm tags

upvoted 9 times

✉ **nanaw770**  3 days, 19 hours ago

Group

Upload category

Tag

upvoted 1 times

✉ **9H3zmT6** 3 weeks, 6 days ago

This question was asked in the actual exam on April 30, 2024 (+9:00, Japan). I think SuperPetey's answer is CORRECT, because I passed the AI-102 exam with a score of 917/1000. Thank you very much.

upvoted 1 times

✉ **nanaw770** 3 days, 19 hours ago

So questions registered in 2021 will still be on the exam in April 2024? Japan is a scary country.

upvoted 1 times

✉ **varinder82** 2 months ago

Final Answer

- Upload all
- Get suggested tags
- Review and confirm tags

upvoted 1 times

✉ **evangelist** 3 months, 1 week ago

To minimize the time required for retraining the model, the correct three steps are:

Upload all images: First, you need to bulk upload the 1000 new images to the Custom Vision service. This is the foundational step for preparing the data.

Get suggested tags: Utilize Custom Vision's functionality to automatically suggest tags for the uploaded images. This can significantly reduce the workload of manual tagging.

Review and confirm suggested tags: Finally, manually review and confirm the tags suggested by the system to ensure their accuracy. Then, use these tagged images to retrain the model.

This process leverages the automation tools provided by Custom Vision to streamline and expedite the data preparation process, particularly effective when dealing with a large number of untagged images.

upvoted 4 times

✉ **tdctdc** 6 months ago

Well, it's a bit confusing. In both cases (ET answers and SuperPetey suggestion) - we will have to walk through the pictures manually if there is no info about them. If they are stored in class folders - the ET answer is less time consuming, if not - it's not possible to tell if separating them manually or manual check of suggested tags will take less time.

upvoted 1 times

✉ **sl_msiconsulting** 7 months, 2 weeks ago

The answer is correct - there is no magic here. You can't suggest any new tags based on the model you currently have. Read the limitations of the smart labeler carefully: When to use Smart Labeler
Keep the following limitations in mind:

You should only request suggested tags for images whose tags have already been trained on once. Don't get suggestions for a new tag that you're just beginning to train.

upvoted 1 times

 **josebernabeo** 4 months, 1 week ago

"When you tag images for a Custom Vision model, the service uses the latest trained iteration of the model to predict the labels of new images."

source: <https://learn.microsoft.com/en-us/azure/ai-services/custom-vision-service/suggested-tags>

upvoted 3 times

 **Eltooth** 1 year, 10 months ago

Answer given would be only option IF model had not already been trained with images, so...
I agree with SuperPetey et al...

Upload

Get suggested tags

Review and confirm tags

upvoted 3 times

 **Number00** 2 years ago

I agree with SuperPetey. The answer should be

- 1.) upload all the images
- 2.) Get suggested tags
- 3.) Review the suggestions and confirm the tags

Reason being that using the tools(suggested tags) would still applied to the new 1000 images item, even if those 1000 images doesn't associate with the original data pool. So, that means tagging even 1 less images using the suggested tags would still be faster than manually tagging them.
<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/suggested-tags>

upvoted 4 times

 **reachmymind** 2 years, 2 months ago

- 1.) Upload all the images
- 2.) Get suggested tags
- 3.) Review the suggestions and confirm the tags

If an image does not have any associated TAG, we can add a new one while reviewing

<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/getting-started-improving-your-classifier>

upvoted 1 times

 **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 2 times

 **EXCEL1177** 2 years, 8 months ago

@superpetey, kindly read through the article in the link you shared, I just did and confirmed from it that the provided answer by the platform is correct.

upvoted 3 times

 **GilEdwards** 2 years, 4 months ago

I disagree, the images are unlabeled, but there is nothing in the text of the question mentioning that there are new tags. I agree with SuperPetey.

upvoted 4 times

 **angie31** 2 years, 8 months ago

"You should only request suggested tags for images whose content has already been trained once. Don't get suggestions for a new tag that you're just beginning to train." And the question says RETRAINING of an existing model to which we are adding new images. So the response is actually wrong and @superpetey is correct

upvoted 3 times

 **angie31** 2 years, 8 months ago

AHHHH but the key word is 'DO NOT HAVE ANY ASSOCIATED DATA'. So the content of images is brand new!!! Therefore we cant use suggester and the response is correct!

upvoted 5 times

 **Mehe323** 2 months ago

The point of machine learning is that a model eventually LEARNS how to do things independently. Even though there is no associated data, there is previous learning done and existing labels can be used. I am not sure why we would need ML if we still have to do things manually all the time?

upvoted 2 times

 **ThomasKong** 2 years, 4 months ago

I support your highlighted point to the right point. So the given answer should be correct.

upvoted 1 times

Question #7

You are building a Conversational Language Understanding model for an e-commerce chatbot. Users can speak or type their billing address when prompted by the chatbot.

You need to construct an entity to capture billing addresses.

Which entity type should you use?

- A. machine learned
- B. Regex
- C. list
- D. Pattern.any

Correct Answer: B

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-entity-types>

Community vote distribution

A (100%)

✉ **Internal_Koala** Highly Voted 1 year, 8 months ago

Selected Answer: A

The link provided mentions addresses under 'ML Entities with Structure'. Will be hard to identify all possible international addresses with RegEx.
upvoted 16 times

✉ **practicewizards** Highly Voted 1 year, 8 months ago

Selected Answer: A

A - MAchine Learned

in documentation

<https://learn.microsoft.com/en-us/azure/cognitive-services/LUIS/concepts/entities#ml-entity-with-structure>

- ML Entity with Structure

An ML entity can be composed of smaller sub-entities, each of which can have its own properties. For example, an Address entity could have the following structure:

Address: 4567 Main Street, NY, 98052, USA

Building Number: 4567

Street Name: Main Street

State: NY

Zip Code: 98052

Country: USA

upvoted 8 times

✉ **Arun20212021** Most Recent 3 days, 16 hours ago

As per Gemini chatbot, the answer is B. Regex. The following is the explanation it gives:

Regex offers a powerful and efficient way to define patterns for capturing specific text formats, making it ideal for structured data like billing addresses. You can design a regex pattern to capture elements like street address, city, state, zip code, and potentially country code.

; and why the answer is not 'machine learned':

While machine learning can be powerful for complex entity recognition, it might be overkill for a well-defined structure like billing addresses. Training a machine learning model could be time-consuming and resource-intensive compared to a simpler solution.

upvoted 1 times

✉ **nanaw770** 3 days, 19 hours ago

Selected Answer: A

It must be A.

upvoted 1 times

✉ **funny_penguin** 3 days, 23 hours ago

Selected Answer: A

on exam, I selected machine learned

upvoted 1 times

✉ **Mehe323** 2 months ago

Selected Answer: A

Regex would be a good solution if the addresses were in a specific area (e.g, USA). But since that information is not given, you have to assume any type of address. So the correct answer in this case is A.

upvoted 3 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

<https://learn.microsoft.com/en-us/azure/ai-services/luis/concepts/entities#ml-entity-with-structure>

upvoted 2 times

 **[Removed]** 9 months, 2 weeks ago

all but literally all answers are deliberately false in here, thanks for having the discussions folks!

upvoted 4 times

 **Richi0907** 1 year, 8 months ago

Selected Answer: A

it should be A

upvoted 2 times

 **Sharks82** 1 year, 8 months ago

A. machine learned

upvoted 2 times

Question #8

You are building an Azure WebJob that will create knowledge bases from an array of URLs.

You instantiate a QnAMakerClient object that has the relevant API keys and assign the object to a variable named client.

You need to develop a method to create the knowledge bases.

Which two actions should you include in the method? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Create a list of FileDTO objects that represents data from the WebJob.
- B. Call the client.Knowledgebase.CreateAsync method.
- C. Create a list of QnADTO objects that represents data from the WebJob.
- D. Create a CreateKbDTO object.

Correct Answer: AC

Reference:

<https://docs.microsoft.com/en-us/rest/api/cognitiveservices-qnamaker/qnamaker4.0/knowledgebase/create>

Community vote distribution

BD (100%)

✉  **czmiel24**  2 years, 9 months ago

It should be BD.

upvoted 31 times

✉  **SuperPetey** 2 years, 9 months ago

Correct - see code example here: <https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/quickstarts/quickstart-sdk?tabs=v1%2Cversion-1&pivots=programming-language-csharp>

upvoted 6 times

✉  **rdemontis** 6 months, 3 weeks ago

agree with you. thank you for the provided documentation

upvoted 2 times

✉  **vominhtri854** 2 years, 7 months ago

A knowledge base stores question and answer pairs for the CreateKbDTO object from three sources:

- For editorial content, use the QnADTO object.

To use metadata and follow-up prompts, use the editorial context, because this data is added at the individual QnA pair level.

- For files, use the FileDTO object. The FileDTO includes the filename as well as the public URL to reach the file.

- For URLs, use a list of strings to represent publicly available URLs.

So I believe A and C correct

upvoted 3 times

✉  **reachmymind**  2 years, 2 months ago

Answer :: B & D

A. Create a list of FileDTO objects that represents data from the WebJob.

NO - as it is from URL - so optional

B. Call the client.Knowledgebase.CreateAsync method.

YES - Mandatory to Call the Method

C. Create a list of QnADTO objects that represents data from the WebJob.

NO - as it is from URL - so optional

D. Create a CreateKbDTO object.

YES - Mandatory to Create

Go through the lines starting line 92 at below URL:

https://github.com/Azure-Samples/cognitive-services-qnamaker-csharp/blob/master/documentation-samples/quickstarts/Knowledgebase_Quickstart/Program.cs

upvoted 5 times

✉  **ninjia** 1 year, 9 months ago

I agreed.

You are building an Azure WebJob that will create knowledge bases from an array of ##URLs##.

You could use FileDTO, QnADTO or urls to create the CreateKbDTO. Hence, FileDTO is not mandatory.

Code snippet from the link provided by reachmymind:

```
var createKbDto = new CreateKbDTO  
{
```

```
Name = "QnA Maker FAQ from c# quickstart",
QnaList = new List<QnADTO> { qna1 },
//Files = new List<FileDTO> { file1 },
Urls = urls
};
upvoted 3 times
```

✉ **AzureJobsTillRetire** 1 year, 2 months ago

DTO is data transfer object. You can use CreateKbDTO to create a knowledge base data transfer object that contains URLs and then call the client.Knowledgebase.CreateAsync method to create a knowledge base.

upvoted 1 times

✉ **nanaw770** **Most Recent** 3 days, 19 hours ago

Selected Answer: BD

client.Knowledgebase.CreateAsync and CreateKbDTO.

upvoted 1 times

✉ **evangelist** 3 months, 1 week ago

Selected Answer: BD

Create a CreateKbDTO object: First, you need to create a CreateKbDTO object, which contains the necessary information for creating a knowledge base, such as an array of URLs and the name of the knowledge base.

Call the client.Knowledgebase.CreateAsync method: Then, use the Knowledgebase.CreateAsync method of the QnAMakerClient object, passing in the previously created CreateKbDTO object, to asynchronously create the knowledge base.

upvoted 2 times

✉ **rdemontis** 6 months, 3 weeks ago

Selected Answer: BD

<https://learn.microsoft.com/en-us/azure/ai-services/qnamaker/quickstarts/quickstart-sdk?tabs=v1%2Cversion-1&pivots=programming-language-csharp>

upvoted 2 times

✉ **propanther** 7 months, 1 week ago

It should be BD

CreateKbDTO holds list of FileDto, QnADto. In given question, knowledge base accepts URLs which is one of the properties of FileDto. Hence, FileDtos will be created for each returned file URL from WebJob which can be used to create and hold within CreateKbDTO

[https://azuresdkdocs.blob.core.windows.net/\\$web/dotnet/Microsoft.Azure.CognitiveServices.Knowledge.QnAMaker/3.0.0-preview.1/api/Microsoft.Azure.CognitiveServices.Knowledge.QnAMaker.Models/Microsoft.Azure.CognitiveServices.Knowledge.QnAMaker.Models.CreateKbDTO.html](https://azuresdkdocs.blob.core.windows.net/$web/dotnet/Microsoft.Azure.CognitiveServices.Knowledge.QnAMaker/3.0.0-preview.1/api/Microsoft.Azure.CognitiveServices.Knowledge.QnAMaker.Models/Microsoft.Azure.CognitiveServices.Knowledge.QnAMaker.Models.CreateKbDTO.html)

upvoted 1 times

✉ **katrang** 7 months, 1 week ago

Selected Answer: BD

The client creates a knowledgebase from a CreateKbDTO object using the CreateAsync method. This object contains a list of FileDTOs and/or a list of QnADTOs. In this scenario we don't know which to use, but we definitely need the CreateAsync and the CreateKbDTO object.

upvoted 1 times

✉ **sl_mslconsulting** 7 months, 2 weeks ago

Selected Answer: BD

You can not even create QnA Maker resource any more. Just read though the code, memorize the answers, and then move on.

upvoted 2 times

✉ **nitz14** 11 months, 2 weeks ago

Selected Answer: BD

ANSWER : B & D

The options A and C are not necessary for creating knowledge bases using the QnAMakerClient object. They mention creating a list of FileDTO and QnADTO objects, which could potentially be used for populating the knowledge base with data, but they are not directly related to the process of creating the knowledge base itself.

upvoted 1 times

✉ **EliteAllen** 1 year ago

Selected Answer: BD

B. Call the client.Knowledgebase.CreateAsync method.

D. Create a CreateKbDTO object.

To create a knowledge base using the QnAMakerClient, you would need to create a CreateKbDTO object that contains the details of the knowledge base to be created. This object would include information such as the name of the knowledge base and the URLs of the documents to be included in the knowledge base.

After creating the CreateKbDTO object, you would then call the client.Knowledgebase.CreateAsync method, passing in the CreateKbDTO object as parameter. This method would create the knowledge base and return a response that includes the ID of the newly created knowledge base.

upvoted 1 times

 **SSJA** 1 year, 5 months ago

Selected Answer: BD

B & D is the correct answers.

upvoted 1 times

 **taer** 1 year, 8 months ago

Selected Answer: BD

It should be BD.

upvoted 1 times

 **RamonKaus** 1 year, 10 months ago

Selected Answer: BD

I agree with Eltooth & czmiel

code from MS:

```
var createOp = await client.Knowledgebase.CreateAsync(createKbDto);
```

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/quickstarts/quickstart-sdk?tabs=v1%2Cversion-1&pivots=programming-language-csharp#create-a-knowledge-base>

upvoted 2 times

 **Eltooth** 1 year, 10 months ago

Selected Answer: BD

B and D are correct answers.

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/quickstarts/quickstart-sdk?tabs=v1%2Cversion-1&pivots=programming-language-csharp#create-a-knowledge-base>

upvoted 1 times

 **torekx** 2 years, 7 months ago

should be bcd options based on

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/quickstarts/quickstart-sdk?tabs=v1%2Cversion-1&pivots=programming-language-csharp#create-a-knowledge-base>

upvoted 1 times

 **GMKanon** 2 years, 8 months ago

Should be BC rather.

upvoted 2 times

Question #9

HOTSPOT -

You are developing an application that includes language translation.

The application will translate text retrieved by using a function named `getTextToBeTranslated`. The text can be in one of many languages. The content of the text must remain within the Americas Azure geography.

You need to develop code to translate the text to a single language.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
    . . .
var endpoint =
    "https://api.cognitive.microsofttranslator.com/translate"
    "https://api.cognitive.microsofttranslator.com/transliterate"
    "https://api-apc.cognitive.microsofttranslator.com/detect"
    "https://api-nam.cognitive.microsofttranslator.com/detect"
    "https://api-nam.cognitive.microsofttranslator.com/translate"
;

var apiKey = "FF956C68B83B21B38691ABD200A4C606";
var text = getTextToBeTranslated();
var body = '[{"Text":"' + text + '"}]';
var client = new HttpClient();
client.DefaultRequestHeaders.Add("Ocp-Apim-Subscription-Key", apiKey);

var uri = endpoint + "?from=en";
var uri = endpoint + "?suggestedFrom=en";
var uri = endpoint + "?to=en";

HttpResponseMessage response;
var content = new StringContent(body, Encoding.UTF8, "application/json");
var response = await client.PutAsync(uri, content);
. . .
```

Correct Answer:

Answer Area

```

    . . .
var endpoint = "https://api.cognitive.microsofttranslator.com/translate";
var endpoint = "https://api.cognitive.microsofttranslator.com/transliterate";
var endpoint = "https://api-apc.cognitive.microsofttranslator.com/detect";
var endpoint = "https://api-nam.cognitive.microsofttranslator.com/detect";
var endpoint = "https://api-nam.cognitive.microsofttranslator.com/translate";

var apiKey = "FF956C68B83B21B38691ABD200A4C606";
var text = getTextToBeTranslated();
var body = '[{"Text":"' + text + '"}]';
var client = new HttpClient();
client.DefaultRequestHeaders.Add("Ocp-Apim-Subscription-Key", apiKey);

var uri = endpoint + "?from=en";
var uri = endpoint + "?suggestedFrom=en";
var uri = endpoint + "?to=en";

HttpResponseMessage response;
var content = new StringContent(body, Encoding.UTF8, "application/json");
var response = await client.PutAsync(uri, content);
. . .

```

 **Shourya007** Highly Voted  2 years, 4 months ago

Wrong , correct answer is both last option.

api-nam.cognitive.microsofttranslator.com

?to=en'

upvoted 22 times

 **ExamDev** Highly Voted  8 months, 2 weeks ago

Box 1: api-nam/translate : <https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-reference#base-urls>

Box 2: "?to=en";

upvoted 10 times

 **nanaw770** Most Recent  3 days, 19 hours ago

api-nam/translate

to=en

upvoted 1 times

 **funny_penguin** 3 days, 23 hours ago

on exam, /translate and to=en

upvoted 1 times

 **varinder82** 2 months ago

Final Answer

1. ./translate

2. to=en

upvoted 1 times

 **anto69** 3 months, 2 weeks ago

insane way to ask a question

upvoted 4 times

 **katrang** 7 months, 1 week ago

api-nam...../translate

?to=en

upvoted 2 times

 **zellck** 11 months ago

1. api-nam.cognitive.microsofttranslator.com

2. ?to=en'

<https://learn.microsoft.com/en-us/azure/cognitive-services/Translator/reference/v3-0-reference#base-urls>

Requests to Translator are, in most cases, handled by the datacenter that is closest to where the request originated. If there's a datacenter failure when using the global endpoint, the request may be routed outside of the geography.

To force the request to be handled within a specific geography, use the desired geographical endpoint. All requests are processed among the datacenters within the geography.

- United States

api-nam.cognitive.microsofttranslator.com

<https://learn.microsoft.com/en-us/azure/cognitive-services/translator/reference/rest-api-guide>

- translate

Translate specified source language text into the target language text.

upvoted 6 times

✉ **RAN_L** 1 year, 2 months ago

```
var endpoint = "https://api-nam.cognitive.microsofttranslator.com/translate";
var apiKey = "FF956C68883821838691A8D200A4C606";
var text = getTextToBeTranslated();
var body = "[{\\"Text\\":\"" + text + "\\"}]";
var client = new HttpClient();
client.DefaultRequestHeaders.Add("Ocp-Apim-Subscription-Key", apiKey);
var uri = endpoint + "?to=en";
var content = new StringContent(body, Encoding.UTF8, "application/json");
HttpResponseMessage response = await client.PutAsync(uri, content);
string translatedText = await response.Content.ReadAsStringAsync();
upvoted 1 times
```

✉ **Eltooth** 1 year, 10 months ago

Box 1: api-nam/translate : <https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-reference#base-urls>

Box 2: "?to=en";

upvoted 1 times

✉ **PHD_CHENG** 1 year, 11 months ago

Was on exam 7 Jun 2022

upvoted 2 times

✉ **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 4 times

✉ **Adedoyin_Simeon** 2 years, 7 months ago

SuperPetey is correct.

ref:

<https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-reference#base-urls>

upvoted 3 times

✉ **alan007** 2 years, 9 months ago

wrong for first choice, should be API.cognitive...../tanslate

upvoted 7 times

✉ **SuperPetey** 2 years, 9 months ago

incorrect - the question specifies it should be routed to the Americas region. The correct answer for box 1 is api-nam.cognitive.microsofttranslator.com/translate according to doc @

<https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-reference#base-urls>

I agree second drop-down is '?to=en'

upvoted 42 times

✉ **rdemontis** 6 months, 3 weeks ago

agree with you. thanks for the provided reference

upvoted 1 times

Question #10

You are building a conversational language understanding model.

You need to enable active learning.

What should you do?

- A. Add show-all-intents=true to the prediction endpoint query.
- B. Enable speech priming.
- C. Add log=true to the prediction endpoint query.
- D. Enable sentiment analysis.

Correct Answer: C

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-how-to-review-endpoint-utterances#log-user-queries-to-enable-active-learning>

Community vote distribution

C (100%)

✉ Eltooth Highly Voted 1 year, 10 months ago

Selected Answer: C

C is the correct answer.

"To enable active learning, you must log user queries. This is accomplished by calling the endpoint query with the log=true query string parameter value."

<https://docs.microsoft.com/en-us/azure/cognitive-services/LUIS/how-to/improve-application#log-user-queries-to-enable-active-learning>
upvoted 11 times

✉ nanaw770 Most Recent 3 days, 19 hours ago

Selected Answer: C

Add log=true
upvoted 1 times

✉ anto69 3 months ago

Selected Answer: C

C is correct
upvoted 1 times

✉ rdemontis 6 months, 3 weeks ago

Selected Answer: C

correct answer and reference
<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-how-to-review-endpoint-utterances#log-user-queries-to-enable-active-learning>
upvoted 1 times

✉ zellck 11 months ago

Selected Answer: C

C is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/LUIS/how-to/improve-application#log-user-queries-to-enable-active-learning>
To enable active learning, you must log user queries. This is accomplished by calling the endpoint query with the log=true query string parameter value.
upvoted 4 times

✉ zellck 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 4 times

✉ marti_tremblay000 1 year, 2 months ago

Log user queries to enable active learning
To enable active learning, you must log user queries. This is accomplished by calling the endpoint query with the log=true query string parameter value.
<https://learn.microsoft.com/en-us/azure/cognitive-services/LUIS/how-to/improve-application>
upvoted 1 times

✉ **Marilena96** 1 year, 3 months ago

To enable active learning in a conversational language understanding model, you should add show-all-intents=true to the prediction endpoint query. This will allow you to see all the intents that the model is predicting, including the None intent.[0] This information can be used to improve the model by adding more training data for the None intent or other intents that are not being predicted accurately.

upvoted 1 times

✉ **jekko** 2 years ago

Correct.

Reference: <https://docs.microsoft.com/en-us/azure/cognitive-services/LUIS/how-to/improve-application>

upvoted 1 times

Question #11

HOTSPOT -

You run the following command.

```
docker run --rm -it -p 5000:5000 --memory 10g --cpus 2 \
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment\
Eula=accept \
Billing={ENDPOINT_URI} \
ApiKey={API_KEY}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area**Statements****Yes****No**

Going to <http://localhost:5000/status> will query the Azure endpoint to verify whether the API key used to start the container is valid.

The container logging provider will write log data.

Going to <http://localhost:5000/swagger> will provide the details to access the documentation for the available endpoints.

Correct Answer:

Answer Area**Statements****Yes****No**

Going to <http://localhost:5000/status> will query the Azure endpoint to verify whether the API key used to start the container is valid.

The container logging provider will write log data.

Going to <http://localhost:5000/swagger> will provide the details to access the documentation for the available endpoints.

Box 1: Yes -

<http://localhost:5000/status> : Also requested with GET, this verifies if the api-key used to start the container is valid without causing an endpoint query.

Box 2: Yes -

The command saves container and LUIS logs to output mount at C:\output, located on container host

Box 3: Yes -

<http://localhost:5000/swagger> : The container provides a full set of documentation for the endpoints and a Try it out feature. With this feature, you can enter your settings into a web-based HTML form and make the query without having to write any code. After the query returns, an example CURL command is provided to demonstrate the HTTP headers and body format that's required.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-container-howto>

 Internal_Koala  1 year, 8 months ago

Yes

No

Yes

Log location is not mounted. The ET answer relates to an example provided on the given website which DOES mount a log location.

upvoted 24 times

✉ **zellck** Highly Voted 11 months ago

YNY is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/text-analytics-for-health/how-to/use-containers?tabs=language#validate-that-a-container-is-running>

- <http://localhost:5000/status>

Also requested with GET, this URL verifies if the api-key used to start the container is valid without causing an endpoint query. This request can be used for Kubernetes liveness and readiness probes.

- <http://localhost:5000/swagger>

The container provides a full set of documentation for the endpoints and a Try it out feature. With this feature, you can enter your settings into a web-based HTML form and make the query without having to write any code. After the query returns, an example CURL command is provided to demonstrate the HTTP headers and body format that's required.

upvoted 10 times

✉ **nanaw770** Most Recent 3 days, 19 hours ago

Yes No Yes

upvoted 1 times

✉ **funny_penguin** 3 days, 23 hours ago

on exam, YNY

upvoted 1 times

✉ **varinder82** 2 months ago

Fintal Answer:

Yes

No

Yes

upvoted 2 times

✉ **evangelist** 4 months ago

Going to <http://localhost:5000/status> will query the Azure endpoint to verify whether the API key used to start the container is valid.

Yes. Typically, Azure Cognitive Services containers provide a /status endpoint that can be used to check the status of the service, including the validity of the API key. Since the service is mapped to localhost:5000, accessing this URL should provide the status of the containerized service, including the API key's validity.

The container logging provider will write log data.

No (Assuming). This statement is somewhat ambiguous and depends on the configuration of the Docker container and the Azure Cognitive Services container.

Going to <http://localhost:5000/swagger> will provide the details to access the documentation for the available endpoints.

Yes. It is a common practice for web services and APIs, including those provided by Azure Cognitive Services, to offer a Swagger UI at a /swagger endpoint.

upvoted 4 times

✉ **Gvalli** 5 months, 1 week ago

A modified version of this was in the exam today.

upvoted 4 times

✉ **rdemontis** 6 months, 3 weeks ago

No --> Also requested with GET, this URL verifies if the api-key used to start the container is valid without causing an endpoint query.

<https://learn.microsoft.com/en-us/azure/ai-services/language-service/text-analytics-for-health/how-to/use-containers?tabs=language#validate-that-a-container-is-running>

No -- there isn't any Log location mounted (same link as above)

Yes --> correct, swagger show a full set of documentation for the endpoints

Yes.

upvoted 3 times

✉ **rdemontis** 6 months, 2 weeks ago

sorry the last yes is a typo

upvoted 2 times

✉ **M25** 8 months, 4 weeks ago

No ("without causing an endpoint query"), No, Yes

upvoted 2 times

✉ **ExamDev** 8 months, 2 weeks ago

Agree with you:

- <http://localhost:5000/status>

"Also requested with GET, this URL verifies if the api-key used to start the container is valid WITHOUT CAUSING AN ENDPOINT QUERY."
(According to documentation)

upvoted 1 times

✉ **mmaguero** 12 months ago

1 and 3, is true: <https://learn.microsoft.com/es-es/azure/cognitive-services/language-service/sentiment-opinion-mining/how-to/use-containers#validate-that-a-container-is-running>

2, i think, is yes... (by docker settings)

upvoted 4 times

✉ **odisor** 1 year, 1 month ago

I think the first one is YES

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/text-analytics-for-health/how-to/use-containers?tabs=language#validate-that-a-container-is-running>

upvoted 1 times

✉ **MDawson** 1 year, 1 month ago

Documentation says it will NOT cause an endpoint query, so I think the first one should be NO

upvoted 5 times

✉ **ExamDev** 8 months, 2 weeks ago

Correct!

upvoted 1 times

Question #12

You are building a Language Understanding model for an e-commerce platform.

You need to construct an entity to capture billing addresses.

Which entity type should you use for the billing address?

- A. machine learned
- B. Regex
- C. geographyV2
- D. Pattern.any
- E. list

Correct Answer: B

A regular expression entity extracts an entity based on a regular expression pattern you provide. It ignores case and ignores cultural variant.

Regular expression is best for structured text or a predefined sequence of alphanumeric values that are expected in a certain format. For example:

Entity	Regular expression	Example
Flight Number	flight [A-Z]{2} [0-9]{4}	flight AS 1234
Credit Card Number	[0-9]{16}	5478789865437632

Incorrect Answers:

C: The prebuilt geographyV2 entity detects places. Because this entity is already trained, you do not need to add example utterances containing GeographyV2 to the application intents. GeographyV2 entity is supported in English culture.

The geographical locations have subtypes:

Subtype	Purpose
poi	point of interest
city	name of city
countryRegion	name of country or region
continent	name of continent
state	name of state or province

D: Pattern.any is a variable-length placeholder used only in a pattern's template utterance to mark where the entity begins and ends.

E: A list entity represents a fixed, closed set of related words along with their synonyms. You can use list entities to recognize multiple synonyms or variations and extract a normalized output for them. Use the recommend option to see suggestions for new words based on the current list.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-entity-types>

Community vote distribution

A (79%)

B (15%)

6%

 **ExamPrep2021**  2 years, 11 months ago

My guess is A.

An ML entity can be composed of smaller sub-entities, each of which can have its own properties. For example, Address could have the following structure:

Address: 4567 Main Street, NY, 98052, USA
 Building Number: 4567
 Street Name: Main Street
 State: NY
 Zip Code: 98052
 Country: USA

upvoted 34 times

 **LKLK10** Highly Voted  2 years, 11 months ago

ML. Answer is A

upvoted 10 times

 **azurelearner666** 2 years, 11 months ago

Right! (the correct response is A, Machine Learned)

See

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-entity-types>

It is a Machine Learned Entity (check ML Entity with Structure in the link, as it is an Address example...)

upvoted 10 times

 **nanaw770** Most Recent  3 days, 18 hours ago

Selected Answer: A

A is right answer.

upvoted 1 times

 **evangelist** 4 months ago

Given these options, A. Machine Learned is the most appropriate choice for capturing billing addresses. Billing addresses are complex entities with a lot of variability in their format and structure. A machine-learned entity is capable of understanding and extracting such complex information from natural language inputs, which makes it suitable for this purpose. It can learn from examples and capture the billing address as an entity based on the context in which it appears, which is essential for handling the wide range of ways in which addresses can be presented.

upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

duplicated question

<https://learn.microsoft.com/en-us/azure/ai-services/LUIS/concepts/entities>

upvoted 1 times

 **jakespeed** 7 months, 3 weeks ago

Selected Answer: A

ML Entity with Structure

An ML entity can be composed of smaller sub-entities, each of which can have its own properties. For example, an Address entity could have the following structure:

Address: 4567 Main Street, NY, 98052, USA

Building Number: 4567

Street Name: Main Street

State: NY

Zip Code: 98052

Country: USA

upvoted 1 times

 **ExamDev** 8 months, 2 weeks ago

Selected Answer: A

Correct answer is A

upvoted 1 times

 **zellck** 11 months ago

Same as Question 7.

<https://www.examtopics.com/discussions/microsoft/view/60239-exam-ai-102-topic-3-question-7-discussion>

upvoted 1 times

 **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/LUIS/concepts/entities#machine-learned-ml-entity>

Machine learned entity uses context to extract entities based on labeled examples. It is the preferred entity for building LUIS applications. It relies on machine-learning algorithms and requires labeling to be tailored to your application successfully. Use an ML entity to identify data that isn't always well formatted but have the same meaning.

An ML entity can be composed of smaller sub-entities, each of which can have its own properties. For example, an Address entity could have the following structure:

Address: 4567 Main Street, NY, 98052, USA

Building Number: 4567

Street Name: Main Street

State: NY

Zip Code: 98052

Country: USA

upvoted 1 times

 **EliteAllen** 11 months, 2 weeks ago

Selected Answer: C

C. geographyV2

The geographyV2 prebuilt entity in Language Understanding (LUIS) is designed to recognize and label entities that are geographical locations, such as city, state, or country. This would be suitable for capturing billing addresses in an e-commerce platform.

upvoted 2 times

M25 8 months, 4 weeks ago<https://learn.microsoft.com/en-us/azure/ai-services/luis/luis-reference-prebuilt-geographyv2?tabs=V3>

The prebuilt geographyV2 entity detects places.

The geographical locations have subtypes:

poi point of interest

city name of city

countryRegion name of country or region

continent name of continent

state name of state or province

I guess you could charge a bill for the Statue of Liberty on Ellis Island as a (fixed) "poi", but a more generalized rule would rather look for an Address entity with sub-entities (variable) as an ML Entity with Structure type

upvoted 2 times

ap1234pa 1 year, 4 months ago**Selected Answer: A**

Wherever it is address it is ML

upvoted 3 times

David_ml 1 year, 7 months ago**Selected Answer: A**

A is correct

upvoted 3 times

Nebary 1 year, 9 months ago**Selected Answer: A**

It is 100% A

upvoted 3 times

ExamGuruBhai 1 year, 9 months ago**Selected Answer: A**<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/concepts/entities>

upvoted 2 times

RamonKaus 1 year, 10 months ago**Selected Answer: B**

Its regex. Udemy agrees.

upvoted 5 times

Eltooth 1 year, 10 months ago**Selected Answer: A**

A is correct answer.

upvoted 2 times

Saby2184 2 years, 1 month ago

ML Entity with Structure

An ML entity can be composed of smaller sub-entities, each of which can have its own properties. For example, an Address entity could have the following structure:

Address: 4567 Main Street, NY, 98052, USA

Building Number: 4567

Street Name: Main Street

State: NY

Zip Code: 98052

Country: USA

<https://docs.microsoft.com/en-us/azure/cognitive-services/LUIS/concepts/entities#ml-entity-with-structure>

upvoted 2 times

Question #13

You need to upload speech samples to a Speech Studio project for use in training.

How should you upload the samples?

- A. Combine the speech samples into a single audio file in the .wma format and upload the file.
- B. Upload a .zip file that contains a collection of audio files in the .wav format and a corresponding text transcript file.
- C. Upload individual audio files in the FLAC format and manually upload a corresponding transcript in Microsoft Word format.
- D. Upload individual audio files in the .wma format.

Correct Answer: B

To upload your data, navigate to the Speech Studio . From the portal, click Upload data to launch the wizard and create your first dataset. You'll be asked to select a speech data type for your dataset, before allowing you to upload your data.

The default audio streaming format is WAV

Use this table to ensure that your audio files are formatted correctly for use with Custom Speech:

Property	Value
File format	RIFF (WAV)
Sample rate	8,000 Hz or 16,000 Hz
Channels	1 (mono)
Maximum length per audio	2 hours
Sample format	PCM, 16-bit
Archive format	.zip
Maximum archive size	2 GB

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-custom-speech-test-and-train>

Community vote distribution

B (100%)

 evangelist Highly Voted 4 months ago

Selected Answer: B

, the best option is B. Upload a .zip file that contains a collection of audio files in the .wav format and a corresponding text transcript file. This method provides a balance of audio quality (with .wav files) and organization (having audio and transcripts together), which is essential for efficient and accurate training of speech recognition models.

upvoted 6 times

 Gvalli Highly Voted 5 months, 1 week ago

A modified version of this was in the exam today.

upvoted 5 times

 nanaw770 Most Recent 3 days, 19 hours ago

Selected Answer: B

wav zip

upvoted 1 times

 funny_penguin 3 days, 23 hours ago

Selected Answer: B

B is correct

upvoted 1 times

 orionduo 5 months, 2 weeks ago

Selected Answer: B

Answer is correct

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

Answer is correct:

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-custom-voice-training-data#types-of-training-data>

upvoted 2 times

 **ManvaiT** 7 months, 3 weeks ago

B is the Answer. Got this in Oct2023 exam

upvoted 3 times

 **JDKJDKJDK** 7 months, 3 weeks ago

just want to confirm, theres no labs included right?

upvoted 3 times

 **zellck** 11 months ago

Same as Question 2.

<https://www.examtopics.com/discussions/microsoft/view/55251-exam-ai-102-topic-3-question-2-discussion>

upvoted 1 times

 **zellck** 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-custom-voice-training-data#types-of-training-data>

A voice training dataset includes audio recordings, and a text file with the associated transcriptions. Each audio file should contain a single utterance (a single sentence or a single turn for a dialog system), and be less than 15 seconds long.

- Individual utterances + matching transcript

A collection (.zip) of audio files (.wav) as individual utterances. Each audio file should be 15 seconds or less in length, paired with a formatted transcript (.txt).

upvoted 3 times

 **rdemontis** 6 months, 3 weeks ago

thanks for explanation

upvoted 1 times

 **Eltooth** 1 year, 10 months ago

Selected Answer: B

B is correct answer.

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-custom-speech-test-and-train>

upvoted 1 times

Question #14

You are developing a method for an application that uses the Translator API.

The method will receive the content of a webpage, and then translate the content into Greek (el). The result will also contain a transliteration that uses the Roman alphabet.

You need to create the URI for the call to the Translator API.

You have the following URI.

<https://api.cognitive.microsofttranslator.com/translate?api-version=3.0>

Which three additional query parameters should you include in the URI? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

A. toScript=Cyril

B. from=el

C. textType=html

D. to=el

E. textType=plain

F. toScript=Latn

Correct Answer: CDF

C: textType is an optional parameter. It defines whether the text being translated is plain text or HTML text (used for web pages).

D: to is a required parameter. It specifies the language of the output text. The target language must be one of the supported languages included in the translation scope.

F: toScript is an optional parameter. It specifies the script of the translated text.

We use Latin (Roman alphabet) script.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-translate>

Community vote distribution

CDF (100%)

✉  **Pixelmate**  11 months ago

THIS WAS ON EXAM 28/06

upvoted 5 times

✉  **nanaw770**  3 days, 19 hours ago

Selected Answer: CDF

textType=html

to=el

toScript=Latn

upvoted 1 times

✉  **funny_penguin** 3 days, 23 hours ago

Selected Answer: CDF

correct answer

upvoted 1 times

✉  **evangelist** 4 months ago

Selected Answer: CDF

<https://api.cognitive.microsofttranslator.com/translate?api-version=3.0&textType=html&to=el&toScript=Latn>

upvoted 4 times

✉  **rdemontis** 6 months, 3 weeks ago

Selected Answer: CDF

correct answer and explanation

upvoted 3 times

✉  **zellck** 11 months ago

Same as Question 3.

<https://www.examtopics.com/discussions/microsoft/view/56391-exam-ai-102-topic-3-question-3-discussion>

upvoted 2 times

 **zellck** 11 months ago

Selected Answer: CDF

CDF is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-translate>

- to

Specifies the language of the output text. The target language must be one of the supported languages included in the translation scope. For example, use to=de to translate to German.

It's possible to translate to multiple languages simultaneously by repeating the parameter in the query string. For example, use to=de&to=it to translate to German and Italian.

- textType

Defines whether the text being translated is plain text or HTML text. Any HTML needs to be a well-formed, complete element. Possible values are: plain (default) or html.

upvoted 2 times

 **Aztek0403** 1 year, 6 months ago

Selected Answer: CDF

Correct!

upvoted 1 times

 **RamonKaus** 1 year, 10 months ago

Selected Answer: CDF

Agreed

upvoted 2 times

 **Eltooth** 1 year, 10 months ago

Selected Answer: CDF

C, D and F are correct answers.

textType=html

to=el

toScript=Latn

<https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-translate#translate-with-transliteration>

upvoted 4 times

Question #15

You have a chatbot that was built by using the Microsoft Bot Framework.

You need to debug the chatbot endpoint remotely.

Which two tools should you install on a local computer? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Fiddler
- B. Bot Framework Composer
- C. Bot Framework Emulator
- D. Bot Framework CLI
- E. ngrok
- F. nginx

Correct Answer: CE

Bot Framework Emulator is a desktop application that allows bot developers to test and debug bots, either locally or remotely. ngrok is a cross-platform application that "allows you to expose a web server running on your local machine to the internet." Essentially, what we'll be doing is using ngrok to forward messages from external channels on the web directly to our local machine to allow debugging, as opposed to the standard messaging endpoint configured in the Azure portal.

Reference:

<https://docs.microsoft.com/en-us/azure/bot-service/bot-service-debug-emulator>

Community vote distribution

CE (100%)

evangelist Highly Voted 4 months ago

Selected Answer: CE

C. Bot Framework Emulator: This is an essential tool for debugging Microsoft Bot Framework bots. It allows you to test and debug your bots on your local machine by emulating the Bot Framework's channels and activities. It can be very helpful in a local development environment but is less suited for remote debugging.

E. ngrok: ngrok is a tool that creates a secure tunnel to your localhost. This is very useful for remote debugging because it allows you to expose your local development server to the internet, which is necessary for testing and debugging interactions with services like the Microsoft Bot Framework.

upvoted 5 times

nanaw770 Most Recent 3 days, 19 hours ago

Selected Answer: CE

Bot Framework Emulator takedajuku ngrok are right answer.

upvoted 1 times

rdemontis 6 months, 3 weeks ago

Selected Answer: CE

correct answers

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-debug-channel-ngrok?view=azure-bot-service-4.0>

upvoted 2 times

zellck 11 months ago

Same as Question 4.

<https://www.examtopics.com/discussions/microsoft/view/56390-exam-ai-102-topic-3-question-4-discussion>

upvoted 3 times

zellck 11 months ago

Selected Answer: CE

CE is the answer.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-debug-emulator?view=azure-bot-service-4.0&tabs=csharp>

Bot Framework Emulator is a desktop application that allows bot developers to test and debug bots, either locally or remotely. Using the Emulator, you can chat with your bot and inspect the messages that your bot sends and receives. The Emulator displays messages as they would appear in a web chat UI and logs JSON requests and responses as you exchange messages with your bot. Before you deploy your bot to the cloud, run it locally and test it using the Emulator. You can test your bot using the Emulator even if you haven't yet created it with Azure Bot Service or configured it to run on any channels.

upvoted 1 times

 **zellck** 11 months ago

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-debug-channel-ngrok?view=azure-bot-service-4.0>

While your bot is in development, you can use an IDE and the Bot Framework Emulator to chat with your bot locally and inspect the messages your bot sends and receives. If your bot is in production, you can debug your bot from any channel using ngrok. The seamless connection of your bot to multiple channels is a key feature available in the Bot Framework.

upvoted 1 times

 **Eltooth** 1 year, 10 months ago

Selected Answer: CE

C and E are correct answers.

upvoted 2 times

 **PHD_CHENG** 2 years ago

Selected Answer: CE

Answer is correct.

<https://docs.microsoft.com/en-us/azure/bot-service/bot-service-debug-emulator?view=azure-bot-service-4.0&tabs=csharp>

upvoted 2 times

Question #16**DRAG DROP -**

You are building a retail chatbot that will use a QnA Maker service.

You upload an internal support document to train the model. The document contains the following question: "What is your warranty period?"

Users report that the chatbot returns the default QnA Maker answer when they ask the following question: "How long is the warranty coverage?"

The chatbot returns the correct answer when the users ask the following question: 'What is your warranty period?'

Both questions should return the same answer.

You need to increase the accuracy of the chatbot responses.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions**Answer Area**

Add a new question and answer (QnA) pair.

Retrain the model.

Add additional questions to the document.

Republish the model.

Add alternative phrasing to the question and answer (QnA) pair.

Correct Answer:**Actions****Answer Area**

Add a new question and answer (QnA) pair.

Retrain the model.

Add additional questions to the document.

Republish the model.

Add alternative phrasing to the question and answer (QnA) pair.

Step 1: Add alternative phrasing to the question and answer (QnA) pair.

Add alternate questions to an existing QnA pair to improve the likelihood of a match to a user query.

Step 2: Retrain the model.

Periodically select Save and train after making edits to avoid losing changes.

Step 3: Republish the model -

Note: A knowledge base consists of question and answer (QnA) pairs. Each pair has one answer and a pair contains all the information associated with that answer.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/how-to/edit-knowledge-base>

 **zellck**  11 months ago

1. Add alternative phrasing to the QnA pair.
2. Retrain model.
3. Republish model.

<https://learn.microsoft.com/en-us/azure/cognitive-services/qnamaker/how-to/edit-knowledge-base#question-and-answer-pairs>

The optional settings for a pair include:

- Alternate forms of the question

this helps QnA Maker return the correct answer for a wider variety of question phrasings

upvoted 17 times

✉ **nanaw770** Most Recent 3 days, 19 hours ago

Add alternative QnA

Retrain

Republish

upvoted 1 times

✉ **anto69** 3 months ago

correct answer

upvoted 1 times

✉ **evangelist** 4 months ago

1: Add alternate phrasings as follow-up prompts or as additional questions in the QnA pair:

2: Train the QnA Maker model:

3: Publish the updated QnA Maker model:

upvoted 4 times

✉ **Gvalli** 5 months, 1 week ago

A modified version of this was in the exam today.

upvoted 4 times

✉ **rdemontis** 6 months, 3 weeks ago

correct answer and explanation

upvoted 2 times

✉ **zellck** 11 months ago

Same as Question 5.

<https://www.examtopics.com/discussions/microsoft/view/55309-exam-ai-102-topic-3-question-5-discussion>

upvoted 1 times

✉ **RamonKaus** 1 year, 10 months ago

Correct

upvoted 4 times

✉ **Eltooth** 1 year, 10 months ago

Answer looks correct.

Add alternative phrasing

Retrain

Republish

upvoted 4 times

Question #17

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution. After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen. You build a language model by using a Language Understanding service. The language model is used to search for information on a contact list by using an intent named FindContact.

A conversational expert provides you with the following list of phrases to use for training.

- Find contacts in London.
- Who do I know in Seattle?
- Search for contacts in Ukraine.

You need to implement the phrase list in Language Understanding.

Solution: You create a new intent for location.

Does this meet the goal?

A. Yes

B. No

Correct Answer: A

An intent represents a task or action the user wants to perform. It is a purpose or goal expressed in a user's utterance.

Define a set of intents that corresponds to actions users want to take in your application.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-intent>

Community vote distribution

 B (100%)

□  **hiabcd**  2 years, 11 months ago

This should be "No"

upvoted 31 times

□  **iyoila_daniel** 2 years, 10 months ago

I think the answer should be 'No' too. The intent is for FindContact, not location really.

upvoted 6 times

□  **zubfaruq**  2 years, 8 months ago

The answer should be NO.

An utterance having wo intents? This is illogical.

The model should have an Entity "Location" that will help in finding the contacts.

Reference: <https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-intent>

upvoted 11 times

□  **evangelist**  4 months ago

The ANSWER is NO !

Instead of creating a new intent for location, you should:

Add these phrases to the "FindContact" intent: This helps the model understand that these are different ways a user might express the intent to find contacts.

Use Entities for Locations: In addition to training the intent with these phrases, you should use entities to capture the location information within the phrases. In LUIS, you can define an entity (like "Location") and then annotate locations in your training phrases. This way, the model learns to recognize and extract location details from user inputs.

By combining a well-trained "FindContact" intent with a location entity, your language model will be better equipped to understand and process user queries about finding contacts in specific locations.

upvoted 2 times

□  **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

IMHO is NO, you don't need to create an intent for each location. The intent is always the same: FindContact

upvoted 3 times

□  **Hisayuki** 11 months, 1 week ago

Selected Answer: B

Location should be for entity not for Intent

upvoted 1 times

 **zuggy1117** 11 months, 3 weeks ago

Should be NO. An intent represents a task or action the user wants to perform. It is a purpose or goal expressed in a user's utterance. In this case, the Intent is to find the Contact. And it is already created.

The utterances specified have the same Intent which is to find/search for contacts. Hence they should be added to the list and an Entity should be created to capture location.

upvoted 1 times

 **Marilena96** 1 year, 3 months ago

Think its Yes, creating a new intent for location would meet the goal of implementing the phrase list in Language Understanding. You can add the phrases to the new intent and label the location entity in each phrase. This will help the language model understand the user's intent to search for contacts in a specific location.

upvoted 3 times

 **ap1234pa** 1 year, 4 months ago

Selected Answer: B

This is "NO"

upvoted 1 times

 **Eltooth** 1 year, 10 months ago

Selected Answer: B

B is correct answer.

Needs Entity/type for location instead.

upvoted 2 times

 **zb99** 2 years ago

Selected Answer: B

Definitely no.

upvoted 1 times

 **[Removed]** 2 years, 4 months ago

This should be No.

upvoted 2 times

 **sumanshu** 2 years, 4 months ago

intent (means what task we need to perform) - So here task is same i.e. Find Contacts. So No need to add an new intent.

But need to add type/entity : location

upvoted 2 times

 **TanujitRoy** 2 years, 6 months ago

Selected Answer: B

This should be NO

upvoted 2 times

 **durak** 2 years, 6 months ago

Remember trick

E n (england) --> entity - noun

U s (USA) --> utterances - statement

I v (Ivory Coast) - Intent- Verb

So answer is B

upvoted 5 times

 **GMKanon** 2 years, 8 months ago

Location should be an entity. So "no" is correct

upvoted 2 times

 **YipingRuan** 2 years, 10 months ago

CheckWeather

"What's the weather like in Boston?"

"Show me the forecast for this weekend"

upvoted 1 times

 **vaskar** 2 years, 10 months ago

Answer should be no. Create one entity for location should be the correct approach.

upvoted 3 times

Question #18

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution. After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen. You build a language model by using a Language Understanding service. The language model is used to search for information on a contact list by using an intent named FindContact.

A conversational expert provides you with the following list of phrases to use for training.

⇒ Find contacts in London.

⇒ Who do I know in Seattle?

Search for contacts in Ukraine.

You need to implement the phrase list in Language Understanding.

Solution: You create a new entity for the domain.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

Instead use a new intent for location.

Note: An intent represents a task or action the user wants to perform. It is a purpose or goal expressed in a user's utterance.

Define a set of intents that corresponds to actions users want to take in your application.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-intent>

Community vote distribution

A (74%)

B (26%)

✉  **SamedKia**  2 years, 1 month ago

Wrong! We create a new location entity for domain to keep the location of FindContact intent.

upvoted 16 times

✉  **Nebary**  1 year, 9 months ago

Selected Answer: A

Agreed with SamedKia. Should be YES

The model should have an Entity "Location" that will help in finding the contacts. Reference: <https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-intent>

upvoted 5 times

✉  **f2c587e**  2 months ago

This same question (topic 3, question 1) seems to be indicated by the community that this question has two affirmative answers. I'm not sure if this is correct, normally in this type of question there is only one affirmative answer and the rest are negative. Is anyone clear on the real answer to this question?

upvoted 2 times

✉  **lesiris** 1 month ago

Indeed !!

upvoted 1 times

✉  **chandiochan** 2 months, 2 weeks ago

Selected Answer: B

must be B

upvoted 1 times

✉  **Mehe323** 2 months, 3 weeks ago

Selected Answer: B

The question says

"You need to implement the phrase list in Language Understanding.

Solution: You create a new entity for the domain.

Does this meet the goal?"

Creating an entity is not implementing a phrase list, so I vote NO.

upvoted 2 times

evangelist 4 months ago

The ANSWER is NO !

Instead of creating a new intent for location, you should:

Add these phrases to the "FindContact" intent: This helps the model understand that these are different ways a user might express the intent to find contacts.

Use Entities for Locations: In addition to training the intent with these phrases, you should use entities to capture the location information within the phrases. In LUIS, you can define an entity (like "Location") and then annotate locations in your training phrases. This way, the model learns to recognize and extract location details from user inputs.

By combining a well-trained "FindContact" intent with a location entity, your language model will be better equipped to understand and process user queries about finding contacts in specific locations.

upvoted 3 times

abbey127 5 months, 2 weeks ago

B. No

Creating a new entity for the domain does not directly address the goal of implementing the provided phrase list for the "FindContact" intent. Entities are typically used to extract specific pieces of information from user input, like names, locations, etc. In this case, the goal is to improve the understanding of the intent "FindContact" by providing training phrases related to searching for contacts in specific locations. Instead of creating a new entity, you should focus on training the language model within the existing "FindContact" intent and include the provided phrases to enhance its ability to recognize and understand user queries related to finding contacts in different locations.

upvoted 4 times

rdemontis 6 months, 3 weeks ago

Selected Answer: A

<https://learn.microsoft.com/en-us/azure/ai-services/luis/concepts/entities>

upvoted 2 times

zellck 11 months ago

Same as Question 8.

<https://www.examtopics.com/discussions/microsoft/view/60466-exam-ai-102-topic-3-question-8-discussion>

upvoted 1 times

zellck 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/luis/how-to/entities>

Create entities to extract key data from user utterances in Language Understanding (LUIS) apps. Extracted entity data is used by your client application to fulfill customer requests.

The entity represents a word or phrase inside the utterance that you want extracted. Entities describe information relevant to the intent, and sometimes they are essential for your app to perform its task.

upvoted 2 times

zippy1117 11 months, 3 weeks ago

Should be YES. An intent represents a task or action the user wants to perform. It is a purpose or goal expressed in a user's utterance. In this case, the Intent is to find the Contact. And it is already created.

The utterances specified have the same Intent which is to find/search for contacts. Hence they should be added to the list and an Entity should be created to capture location.

upvoted 2 times

mmaguero 12 months ago

NO, By defining the appropriate utterances within the FindContact intent, the language model will be trained to understand and recognize similar phrases, allowing it to effectively search for information on a contact list based on the given inputs.

upvoted 1 times

EliteAllen 1 year ago

Selected Answer: B

B. No

The solution does not meet the goal. In Language Understanding, entities are used to identify and categorize key data points in user utterances, such as names, places, or dates. While creating a new entity for the domain (e.g., city or country) could be part of the solution, it's not enough on its own.

The phrases provided are intended for training an intent, specifically the "FindContact" intent. Therefore, the solution should involve adding these phrases as example utterances under the "FindContact" intent in the Language Understanding service. This helps the model understand the kind of language users might use when they want to find a contact.

Additionally, an entity for locations (cities, countries) could be created and used in these utterances to help the model identify and extract the location information from user inputs. But merely creating an entity without associating it with the intent and utterances does not meet the goal.

upvoted 3 times

marti_tremblay000 1 year, 2 months ago

Selected Answer: A

ChatGPT answer :

To implement the phrase list in Language Understanding, you need to create a new intent called "FindContact" and add the following phrases as utterances:

Find contacts in London

Who do I know in Seattle?

Search for contacts in Ukraine

You also need to create a new entity for the domain, which could be something like "Location" or "City".

So, yes, creating a new entity for the domain and adding the given phrases to a new intent called "FindContact" would meet the goal of building a language model to search for information on a contact list.

upvoted 3 times

 ap1234pa 1 year, 4 months ago

Selected Answer: A

A is correct

upvoted 3 times

 SSJA 1 year, 5 months ago

Selected Answer: A

Correct answer is A

upvoted 2 times

Question #19

You are training a Language Understanding model for a user support system.

You create the first intent named GetContactDetails and add 200 examples.

You need to decrease the likelihood of a false positive.

What should you do?

- A. Enable active learning.
- B. Add a machine learned entity.
- C. Add additional examples to the GetContactDetails intent.
- D. Add examples to the None intent.

Correct Answer: A

Active learning is a technique of machine learning in which the machine learned model is used to identify informative new examples to label. In LUIS, active learning refers to adding utterances from the endpoint traffic whose current predictions are unclear to improve your model.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-glossary>

Community vote distribution

D (85%) A (15%)

✉  **Isidro**  2 years ago

I would say is D) as per the following: <https://docs.microsoft.com/en-us/azure/cognitive-services/language-service/conversational-language-understanding/concepts/none-intent#adding-examples-to-the-none-intent>

upvoted 26 times

✉  **mk1967** 1 year, 8 months ago

Agreed, as stated in the link:
"You should also consider adding false positive examples to the None intent."

upvoted 5 times

✉  **rdemontis** 6 months, 3 weeks ago

agree with you, thanks for the provided documentation
upvoted 2 times

✉  **evangelist**  4 months ago

False positive means =>
The model needs examples of what it should not classify as "GetContactDetails," which is the role of the "None" intent.
Therefore, the most effective approach is to add a diverse range of examples to the "None" intent, covering various phrases and queries that are outside the scope of "GetContactDetails." This helps create a clear boundary for the model, reducing the likelihood of it mistakenly classifying unrelated inputs as belonging to the "GetContactDetails" intent.

upvoted 6 times

✉  **nanaw770**  3 days, 19 hours ago

Selected Answer: D

It must be D.

upvoted 1 times

✉  **evangelist** 4 months ago

Selected Answer: D

The correct option to decrease the likelihood of a false positive in the Language Understanding model is to add additional None intent examples.

Option D is correct. By adding more varied examples that do not map to a valid intent to the None intent, the model can better learn when an utterance does not apply and avoid falsely matching invalid queries to a valid intent.

Options A, B, and C may improve the model in certain ways, but they do not directly address reducing false positives. Only adding additional out-of-scope examples to the None intent will help the model better distinguish when new utterances do not match any existing intent's patterns.

So out of the options, adding examples to the None intent is the way to decrease the likelihood of false positives.

upvoted 3 times

✉  **rdemontis** 6 months, 3 weeks ago

Selected Answer: D

to me the answer is D. Non intents have the purpose to reduce false positive too.

<https://docs.microsoft.com/en-us/azure/cognitive-services/language-service/conversational-language-understanding/concepts/none-intent#adding-examples-to-the-none-intent>

upvoted 2 times

✉ james2033 9 months, 1 week ago

Selected Answer: D

200 sample data. --> much false positive. --> Increase number of training data. --> Add example to the None intent, not active learning in this context.

upvoted 1 times

✉ zellck 11 months ago

Selected Answer: D

D is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/conversational-language-understanding/concepts/none-intent#adding-examples-to-the-none-intent>

The None intent is also treated like any other intent in your project. If there are utterances that you want predicted as None, consider adding similar examples to them in your training data. For example, if you would like to categorize utterances that are not important to your project as None, such as greetings, yes and no answers, responses to questions such as providing a number, then add those utterances to your intent.

You should also consider adding false positive examples to the None intent. For example, in a flight booking project it is likely that the utterance "want to buy a book" could be confused with a Book Flight intent. Adding "I want to buy a book" or "I love reading books" as None training utterances helps alter the predictions of those types of utterances towards the None intent instead of Book Flight.

upvoted 2 times

✉ Drummer 11 months, 2 weeks ago

A. Enable active learning.

By enabling active learning, the model can actively request feedback from users when it encounters uncertain or ambiguous queries. This feedback loop helps improve the model's understanding and reduces false positives by incorporating user input into its training process. Option A (Enable active learning) is the correct choice to decrease the likelihood of false positives.

upvoted 1 times

✉ ziggy1117 11 months, 3 weeks ago

Selected Answer: D

You should also consider adding false positive examples to the None intent. For example, in a flight booking project it is likely that the utterance "want to buy a book" could be confused with a Book Flight intent. Adding "I want to buy a book" or "I love reading books" as None training utterances helps alter the predictions of those types of utterances towards the None intent instead of Book Flight.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/conversational-language-understanding/concepts/none-intent#adding-examples-to-the-none-intent>

upvoted 3 times

✉ ziggy1117 11 months, 3 weeks ago

You should also consider adding false positive examples to the None intent. For example, in a flight booking project it is likely that the utterance "want to buy a book" could be confused with a Book Flight intent. Adding "I want to buy a book" or "I love reading books" as None training utterances helps alter the predictions of those types of utterances towards the None intent instead of Book Flight.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/conversational-language-understanding/concepts/none-intent#adding-examples-to-the-none-intent>

upvoted 1 times

✉ Marilena96 1 year, 3 months ago

To decrease the likelihood of a false positive, you can add additional examples to the GetContactDetails intent. This will help the model to better understand the intent and reduce the likelihood of false positive predictions.

upvoted 2 times

✉ Rob77 1 year ago

Nope, 20-30 examples per intent is recommended. See <https://learn.microsoft.com/en-us/azure/cognitive-services/LUIS/concepts/application-design#create-example-utterances-for-each-intent>

upvoted 1 times

✉ ap1234pa 1 year, 4 months ago

Selected Answer: D

As explained in MS Document
"false positive" = None intent

upvoted 4 times

✉ ap1234pa 1 year, 4 months ago

Selected Answer: D

Add examples to "None" intent

upvoted 2 times

✉ SSJA 1 year, 5 months ago

Selected Answer: D

Answer is D.

Reference - <https://learn.microsoft.com/en-us/azure/cognitive-services/LUIS/concepts/application-design>

upvoted 3 times

 **GigaCaster** 1 year, 7 months ago

Selected Answer: A
My reasoning for active learning is that it will be better in the long run whereas None intent is but a temporary fix to the current issue.
upvoted 3 times

 **Internal_Koala** 1 year, 8 months ago

Selected Answer: D
As described by Isidro and mk1967.
upvoted 2 times

 **nekkilodeon** 1 year, 9 months ago

Definitely D)
Adding examples to None reduces the chances of selecting the wrong intent.
upvoted 2 times

Question #20

DRAG DROP -

You are building a Language Understanding model for purchasing tickets.

You have the following utterance for an intent named PurchaseAndSendTickets.

Purchase [2 audit business] tickets to [Paris] [next Monday] and send tickets to [email@domain.com]

You need to select the entity types. The solution must use built-in entity types to minimize training data whenever possible.

Which entity type should you use for each label? To answer, drag the appropriate entity types to the correct labels. Each entity type may be used once, more than once, or not at all.

You may need to drag the split bar between panes or scroll to view content.

Select and Place:

Entity Types**Answer Area**

Paris:

email@domain.com:

2 audit business:

Entity Types**Answer Area**

Paris:

email@domain.com:

2 audit business:

Box 1: GeographyV2 -

The prebuilt geographyV2 entity detects places. Because this entity is already trained, you do not need to add example utterances containing GeographyV2 to the application intents.

Box 2: Email -

Email prebuilt entity for a LUIS app: Email extraction includes the entire email address from an utterance. Because this entity is already trained, you do not need to add example utterances containing email to the application intents.

Box 3: Machine learned -

The machine-learning entity is the preferred entity for building LUIS applications.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-reference-prebuilt-geographyv2> <https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-reference-prebuilt-email> <https://docs.microsoft.com/en-us/azure/cognitive-services/luis/reference-entity-machine-learned-entity>

 **Pixelmate**  11 months ago

THIS WAS ON EXAM 28/06

upvoted 7 times

✉ **nanaw770** Most Recent 3 days, 19 hours ago

Paris: GeographyV2
email@domain.com: Email
2 audit business: Machine learned
upvoted 1 times

✉ **evangelist** 4 months ago

The answer is correct
upvoted 3 times

✉ **rdemontis** 6 months, 3 weeks ago

correct answer
upvoted 4 times

✉ **zellck** 11 months ago

1. GeographyV2
2. Email
3. Machine learned

<https://learn.microsoft.com/en-us/azure/cognitive-services/luis/luis-reference-prebuilt-geographyv2?tabs=V3>

The prebuilt geographyV2 entity detects places. Because this entity is already trained, you do not need to add example utterances containing GeographyV2 to the application intents.

<https://learn.microsoft.com/en-us/azure/cognitive-services/luis/luis-reference-prebuilt-email?tabs=V3>

Email extraction includes the entire email address from an utterance. Because this entity is already trained, you do not need to add example utterances containing email to the application intents.

upvoted 4 times

✉ **Eltooth** 1 year, 10 months ago

Answer is correct :
Geography v2 :

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-reference-prebuilt-geographyv2?tabs=V3>

Email :

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-reference-prebuilt-email?tabs=V3-verbose>

Machine Learned :

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/reference-entity-machine-learned-entity?tabs=V3>

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/concepts/entities>

upvoted 4 times

✉ **sdokmak** 1 year, 11 months ago

Correct.

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-reference-prebuilt-email?tabs=V3>

upvoted 2 times

Question #21

You have the following C# method.

```
static void create_resource(string resource_name, string kind, string account_tier, string location)
{
    CognitiveServicesAccount parameters =
        new CognitiveServicesAccount(null, null, kind, location, resource_name, new CognitiveServicesAccountProperties(), new Sku(account_tier));
    var result = cog_svc_client.Accounts.Create(resource_group_name, account_tier, parameters);
}
```

You need to deploy an Azure resource to the East US Azure region. The resource will be used to perform sentiment analysis.

How should you call the method?

- A. create_resource("res1", "ContentModerator", "S0", "eastus")
- B. create_resource("res1", "TextAnalytics", "S0", "eastus")
- C. create_resource("res1", "ContentModerator", "Standard", "East US")
- D. create_resource("res1", "TextAnalytics", "Standard", "East US")

Correct Answer: B

To perform sentiment analysis, we specify TextAnalytics, not ContentModerator.

Possible SKU names include: 'F0'|'F1'|'S0'|'S1'|'S2'|'S3'|'S4'|'S5'|'S6'|'S7'|'S8'

Possible location names include: westus, eastus

Reference:

<https://docs.microsoft.com/en-us/powershell/module/az.cognitiveservices/new-azcognitiveservicesaccount>

Community vote distribution

B (100%)

nanaw770 3 days, 19 hours ago

Selected Answer: B

SKU is S0.

Region is eastus.

Sentiment analysis uses TextAnalytics

upvoted 1 times

evangelist 4 months ago

free tier==>S0, region eastus not East US, sentiment analysis=>TextAnalytics

upvoted 3 times

evangelist 4 months ago

Selected Answer: B

region has to be in format of: eastus

ContentModerator cannot do sentimental Analysis

upvoted 1 times

rdemontis 6 months, 3 weeks ago

Selected Answer: B

<https://learn.microsoft.com/en-us/azure/synapse-analytics/machine-learning/tutorial-text-analytics-use-mmlspark>

upvoted 1 times

SaviB 7 months, 2 weeks ago

Is there any way to get this entire Q & A for free? I'm unable to go past page #23, even though I've signed up for an account.

upvoted 2 times

Mobicid 5 months, 2 weeks ago

@SaviB, did you find a solution? Accessed more than page #23?

upvoted 1 times

zellck 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/sentiment-opinion-mining/overview>

Sentiment analysis and opinion mining are features offered by Azure Cognitive Service for Language, a collection of machine learning and AI algorithms in the cloud for developing intelligent applications that involve written language. These features help you find out what people think of your brand or topic by mining text for clues about positive or negative sentiment, and can associate them with specific aspects of the text.

upvoted 1 times

✉️ **ap1234pa** 1 year, 4 months ago

Selected Answer: B

B is correct

upvoted 1 times

✉️ **Eltooth** 1 year, 10 months ago

Selected Answer: B

Correct answer is B.

create_resource("res1", "TextAnalytics", "S0", "eastus")

Note TextAnalysis will be rebranded into Cognitive Services for Language Service

upvoted 3 times

✉️ **PHD_CHENG** 1 year, 11 months ago

Was on exam 7 Jun 2022

upvoted 2 times

✉️ **PHD_CHENG** 2 years ago

Selected Answer: B

Answer is correct

upvoted 3 times

Question #22

You build a Conversational Language Understanding model by using the Language Services portal.

You export the model as a JSON file as shown in the following sample.

```
{  
  "text": "average amount of rain by month at chicago last year",  
  "intent": "Weather.CheckWeatherValue",  
  "entities": [  
    {  
      "entity": "Weather.WeatherRange",  
      "startPos": 0,  
      "endPos": 6,  
      "children": []  
    },  
    {  
      "entity": "Weather.WeatherCondition",  
      "startPos": 18,  
      "endPos": 21,  
      "children": []  
    },  
    {  
      "entity": "Weather.Historic",  
      "startPos": 23,  
      "endPos": 30,  
      "children": []  
    }  
  ]  
}
```

To what does the Weather.Historic entity correspond in the utterance?

- A. by month
- B. chicago
- C. rain
- D. location

Correct Answer: A

Community vote distribution

A (100%)

 **reiwanotora** 3 days, 19 hours ago

Selected Answer: A

23 to 30 is by month.
upvoted 1 times

 **funny_penguin** 3 days, 23 hours ago

Selected Answer: A

on exam, by month
upvoted 1 times

 **anto69** 3 months ago

Selected Answer: A

"by month" seems correct
upvoted 1 times

 **evangelist** 4 months ago

Selected Answer: A

by month means to check history in the past by each month, no doubt answer is A
upvoted 2 times

 **evangelist** 4 months ago

Selected Answer: A

No doubt, A checked by different AI: chatgpt, claude 2 and Google bard
upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

correct
upvoted 1 times

 **chenglim** 7 months ago

Selected Answer: A

correct answer
upvoted 2 times

 **David_ml** 1 year, 7 months ago

Selected Answer: A

correct
upvoted 2 times

Question #23

You are examining the Text Analytics output of an application.

The text analyzed is: 'Our tour guide took us up the Space Needle during our trip to Seattle last week.'

The response contains the data shown in the following table.

Text	Category	ConfidenceScore
Tour guide	PersonType	0.45
Space Needle	Location	0.38
Trip	Event	0.78
Seattle	Location	0.78
Last week	DateTime	0.80

Which Text Analytics API is used to analyze the text?

- A. Entity Linking
- B. Named Entity Recognition
- C. Sentiment Analysis
- D. Key Phrase Extraction

Correct Answer: B

Named Entity Recognition (NER) is one of the features offered by Azure Cognitive Service for Language, a collection of machine learning and AI algorithms in the cloud for developing intelligent applications that involve written language. The NER feature can identify and categorize entities in unstructured text. For example: people, places, organizations, and quantities.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/language-service/named-entity-recognition/overview>

Community vote distribution

B (100%)

 **reiwanotora** 3 days, 19 hours ago

Selected Answer: B

NER is OK.

upvoted 1 times

 **evangelist** 4 months ago

Selected Answer: B

B. Named Entity Recognition

Named Entity Recognition (NER) is a process in natural language processing that identifies and classifies named entities mentioned in text into predefined categories such as the names of persons, organizations, locations, expressions of times, quantities, monetary values, percentages, etc.

Entity Linking typically involves linking entities to knowledge bases.

Sentiment Analysis is used to determine the sentiment expressed in the text.

Key Phrase Extraction identifies the main points or topics in a text but does not categorize them into entity types like NER does.

upvoted 3 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

<https://learn.microsoft.com/en-us/azure/ai-services/language-service/named-entity-recognition/overview>

upvoted 2 times

 **james2033** 9 months, 1 week ago

Selected Answer: B

Title "What is Named Entity Recognition (NER) in Azure AI Language?" at <https://learn.microsoft.com/en-us/azure/ai-services/language-service/named-entity-recognition/overview>. NER

upvoted 2 times

 **zellck** 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/named-entity-recognition/overview>

Named Entity Recognition (NER) is one of the features offered by Azure Cognitive Service for Language, a collection of machine learning and AI

algorithms in the cloud for developing intelligent applications that involve written language. The NER feature can identify and categorize entities in unstructured text. For example: people, places, organizations, and quantities.

upvoted 2 times

 ap1234pa 1 year, 4 months ago

Selected Answer: B

B is correct

upvoted 1 times

 halfway 1 year, 6 months ago

Selected Answer: B

From the link in the answer: "The NER feature can identify and categorize entities in unstructured text. For example: people, places, organizations, and quantities."

upvoted 2 times

Question #24

SIMULATION -

You need to configure and publish bot12345678 to support task management. The intent must be named TaskReminder. The LUDown for the intent is in the C:\Resources\LU folder.

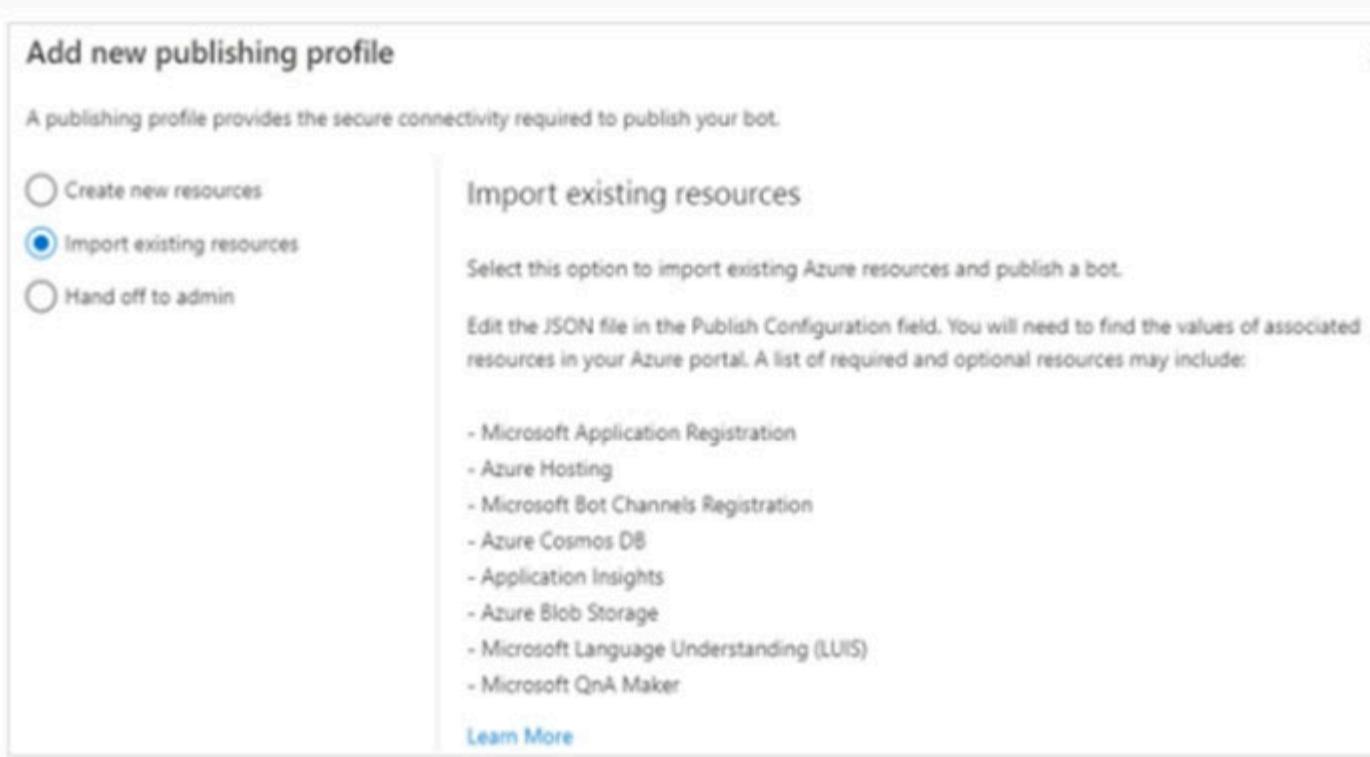
To complete this task, use the Microsoft Bot Framework Composer.

Correct Answer: See explanation below.

Step 1: Open Microsoft Bot Framework Composer

Step 2: Select the bot bot12345678

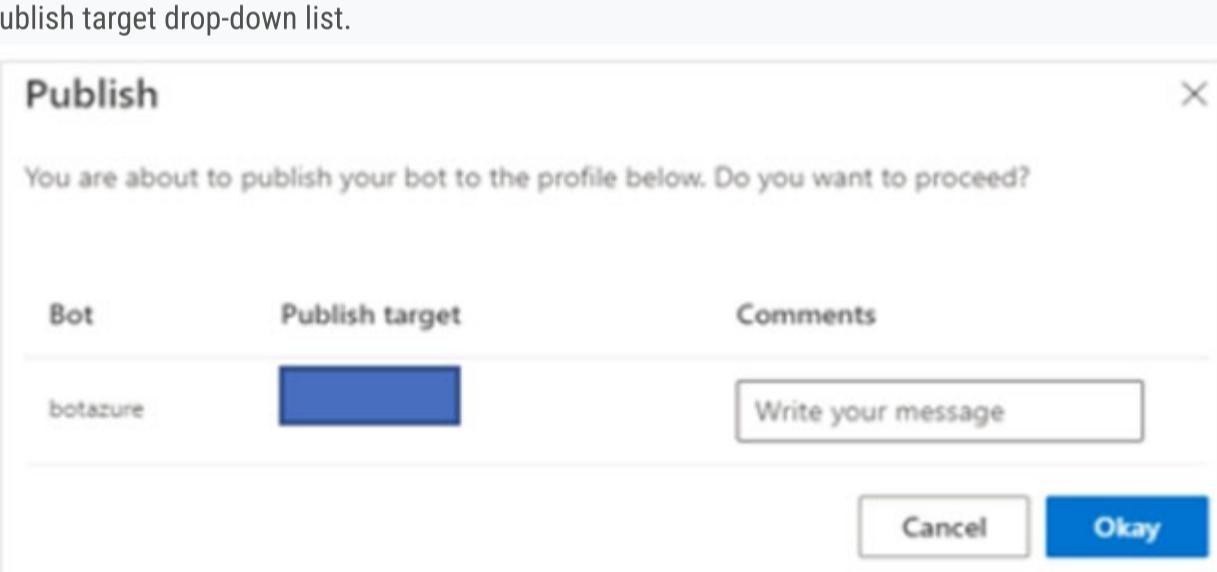
Step 3: Select Import existing resources. Read the instructions on the right side of the screen and select Next.



Step 4: Browse to the C:\Resources\LU folder and select the available .lu file

Step 5: In the pop-up window Importing existing resources, modify the JSON file content based on your resources information: Name the intent TaskReminder

Step 6: Select Publish from the Composer menu. In the Publish your bots pane, select the bot to publish (bot12345678), then select a publish profile from the Publish target drop-down list.



Reference:

<https://docs.microsoft.com/en-us/composer/how-to-publish-bot>

Question #25

SIMULATION -

You need to configure bot12345678 support the French (FR-FR) language.

Export the bot to C:\Resources\Bot\Bot1.zip.

To complete this task, use the Microsoft Bot Framework Composer.

Correct Answer: See explanation below.

Step 1: Open Microsoft Bot Framework Composer

Step 2: Select the bot bot12345678

Step 3: Select Configure.

Step 4: Select the Azure Language Understanding tab

Step 5: Select the Set up Language Understanding button. The Set up Language Understanding window will appear, shown below:

Set up Language Understanding X

To understand natural language input and direct the conversation flow, your bot needs a language understanding service. [Learn more](#)

- Use existing resources
- Create and configure new Azure resources
- Generate instructions for Azure administrator

Next

Cancel

Step 6: Select Use existing resources and then select Next at the bottom of the window.

Step 7: Now select the Azure directory, Azure subscription, and Language Understanding resource name (French).

Step 8: Select Next on the bottom. Your Key and Region will appear on the next on the next window, shown below:

Select Language Understanding resources



The following Language Understanding keys have been successfully added to your bot project:

Key

[REDACTED]

Region

[REDACTED]

Done

Step 9. Select Done -

Reference:

<https://docs.microsoft.com/en-us/composer/concept-language-understanding> <https://docs.microsoft.com/en-us/composer/how-to-add-luis>

ziggy1117 11 months, 3 weeks ago

Select the Configure page from the left and then select the Localization tab. Select Manage bot language to choose your bot's languages.
upvoted 4 times

halfway 1 year, 6 months ago

The answer is incorrect. The instructions in the following link should be used to add multi-language support to a bot:

<https://learn.microsoft.com/en-us/composer/how-to-use-multiple-language?tabs=v2x#update-language-settings>

upvoted 3 times

Question #26

SIMULATION -

You need to configure and publish bot12345678 to answer questions by using the frequently asked questions (FAQ) located at <https://docs.microsoft.com/en-us/azure/bot-service/bot-service-resources-bot-framework-faq>. The solution must use bot%@lab.LabInstance.Id-qna-qna%.

To complete this task, use the Microsoft Bot Framework Composer.

Correct Answer: See explanation below.

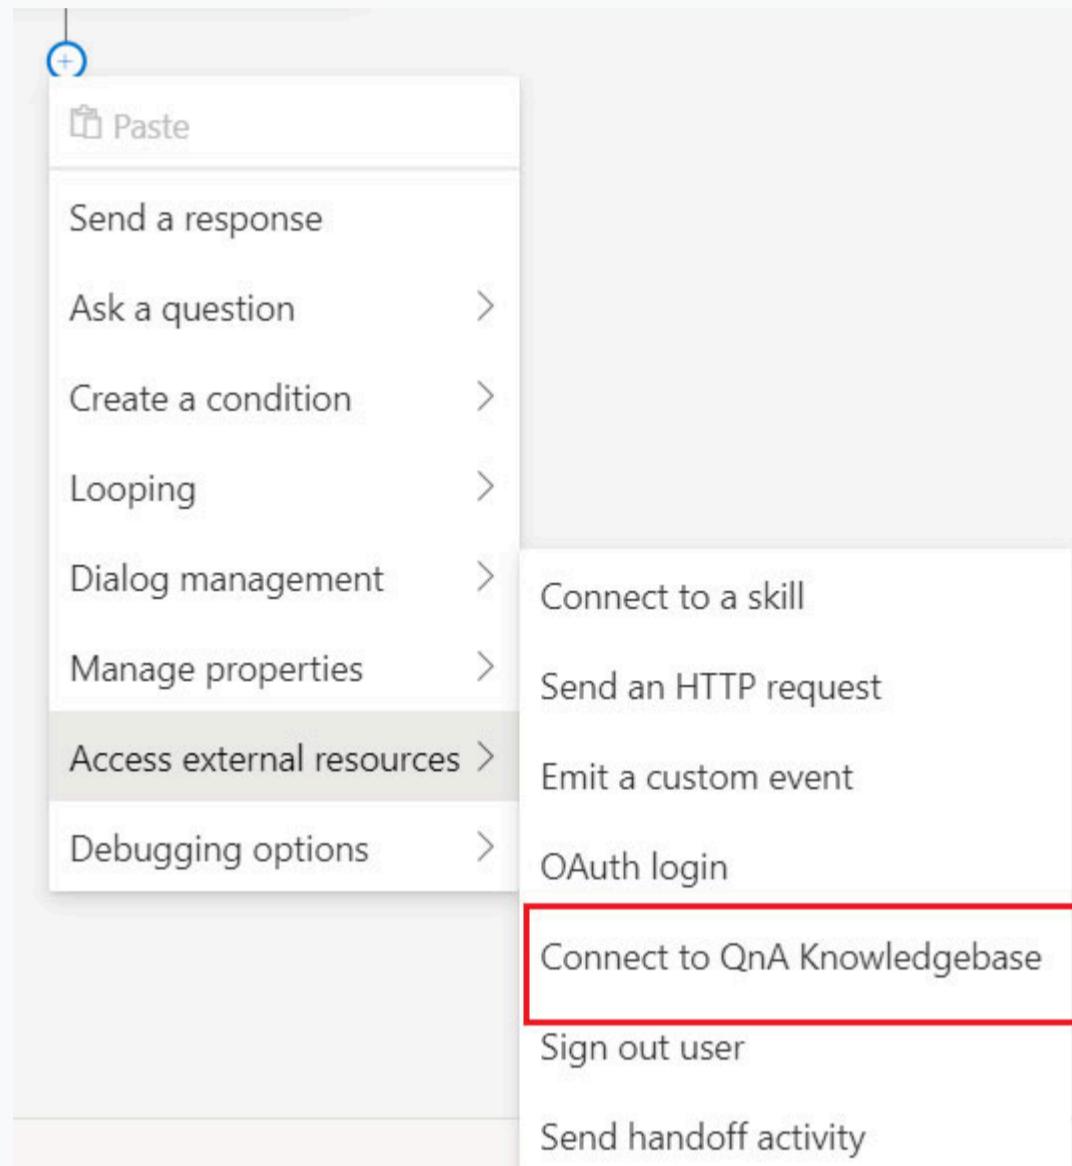
Step 1: Open Microsoft Bot Framework Composer

Step 2: Select the bot bot12345678

Step 3: Open the Configure page in Composer. Then select the Development resources, and scroll down to Azure QnA Maker.

Step 4: To access the Connect to QnA Knowledgebase action, you need to select + under the node you want to add the QnA knowledge base and then select

Connect to QnAKnowledgeBase from the Access external resources action menu.



Step 5: Review the QnA Maker settings panel after selecting the QnA Maker dialog.

Use:

Instance: bot%@lab.LabInstance.Id-qna-qna%

Reference:

<https://docs.microsoft.com/en-us/composer/how-to-create-qna-kb> <https://docs.microsoft.com/en-us/composer/how-to-add-qna-to-bot>

Question #27

You need to measure the public perception of your brand on social media by using natural language processing.

Which Azure service should you use?

- A. Language service
- B. Content Moderator
- C. Computer Vision
- D. Form Recognizer

Correct Answer: A

Azure Cognitive Service for Language is a cloud-based service that provides Natural Language Processing (NLP) features for understanding and analyzing text.

Use this service to help build intelligent applications using the web-based Language Studio, REST APIs, and client libraries.

Note: Natural language processing (NLP) has many uses: sentiment analysis, topic detection, language detection, key phrase extraction, and document categorization.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/language-service/overview>

Community vote distribution

A (100%)

✉  **reiwanotora** 3 days, 19 hours ago

Selected Answer: A

It must be A.

upvoted 1 times

✉  **anto69** 3 months ago

Selected Answer: A

A the only meaningful

upvoted 1 times

✉  **evangelist** 4 months ago

Selected Answer: A

no doubt A

upvoted 3 times

✉  **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

correct, sentiment analysis is the feature to use in this case

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/sentiment-opinion-mining/overview>

upvoted 1 times

✉  **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/sentiment-opinion-mining/overview>

Sentiment analysis and opinion mining are features offered by Azure Cognitive Service for Language, a collection of machine learning and AI algorithms in the cloud for developing intelligent applications that involve written language. These features help you find out what people think of your brand or topic by mining text for clues about positive or negative sentiment, and can associate them with specific aspects of the text.

upvoted 1 times

✉  **halfway** 1 year, 6 months ago

Selected Answer: A

Text Analytics, sentiment analysis: <https://azure.microsoft.com/en-us/products/cognitive-services/text-analytics/#features>

upvoted 2 times

Question #28

HOTSPOT -

You are developing an application that includes language translation.

The application will translate text retrieved by using a function named `get_text_to_be_translated`. The text can be in one of many languages. The content of the text must remain within the Americas Azure geography.

You need to develop code to translate the text to a single language.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
    . . .
    api_key = "FF956C68B83B21B38691ABD200A4C606"
    text = get_text_to_be_translated()
    headers = {
        'Content-Type': 'application/json',
        'Ocp-Apim-Subscription-Key': api_key
    }
    body = {
        'Text': text
    }
    conn = httpplib.HTTPSConnection
        ("api.cogninve.microsofttranslator.com")
        ("api-apc.cognitive.microsofttranslator.com")
        ("api-nam.cognitive.microsofttranslator.com")
    conn.request("POST",
        "/translate?fr=nn=en"
        "/translate?suggestedFrom=en"
        "/translate?to=en"
        "/detect?to=en"
        "/detect?from=en"
    str(body), headers)
    response = conn.getresponse()
    response_data = response.read()
    . . .
```

Correct Answer:**Answer Area**

```

    ...
    api_key = "FF956C68B83B21B38691ABD200A4C606"
    text = get_text_to_be_translated()
    headers = {
        'Content-Type': 'application/json',
        'Ocp-Apim-Subscription-Key': api_key
    }
    body = {
        'Text': text
    }
    conn = httpplib.HTTPSConnection
    conn.request("POST", "/translate?to=en", str(body), headers)
    response = conn.getresponse()
    response_data = response.read()
    ...

```

Box 1: ("api-nam.cognitive.microsofttranslator.com")
 Geography USA: api-nam.cognitive.microsofttranslator.com
 Datacenters: East US, South Central US, West Central US, and West US 2
 Box 2: "/translate?to=en"
 Must specify the language which it is being translated to. The 'to' parameter is required
 Reference:
<https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-reference> <https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-translate>

 **zellck**  11 months ago

1.api-nam.cognitive.microsofttranslator.com
 2. ./translate?to=en

<https://learn.microsoft.com/en-us/azure/cognitive-services/Translator/reference/v3-0-reference#base-urls>

Requests to Translator are, in most cases, handled by the datacenter that is closest to where the request originated. If there's a datacenter failure when using the global endpoint, the request may be routed outside of the geography.

To force the request to be handled within a specific geography, use the desired geographical endpoint. All requests are processed among the datacenters within the geography.

- United States
api-nam.cognitive.microsofttranslator.com

<https://learn.microsoft.com/en-us/azure/cognitive-services/translator/reference/rest-api-guide>

- translate
 Translate specified source language text into the target language text.
 upvoted 13 times

 **rdemontis** 6 months, 3 weeks ago

agree with you, answer is correct
 upvoted 3 times

 **reiwanotora**  3 days, 19 hours ago

api-nam.cognitive.microsofttranslator.com
 translate?to=en
 upvoted 1 times

 **evangelist** 4 months ago

"api-apc.cognitive.microsofttranslator.com" refers to the Asia Pacific endpoint. Using this endpoint will route requests to Azure data centers located in Asia Pacific for processing.

"api-nam.cognitive.microsofttranslator.com" refers to the North America endpoint. Using this endpoint will route requests to Azure data centers located in North America for processing.

So in summary:

"api-apc.cognitive.microsofttranslator.com" - Asia Pacific endpoint
"api-nam.cognitive.microsofttranslator.com" - North America endpoint
upvoted 3 times

 **Gvalli** 5 months, 1 week ago

A modified version of this was in the exam today.
upvoted 2 times

 **ddogg** 9 months, 1 week ago

I feel like the second one should be /detect?to=en since the input can be multiple unspecified languages.
upvoted 2 times

 **ziggy1117** 11 months, 3 weeks ago

answer is correct
upvoted 1 times

 **Davard** 1 year, 8 months ago

The answer seems correct to me. api.nam would keep it within the US.
upvoted 4 times

Question #29

You have the following data sources:

- Finance: On-premises Microsoft SQL Server database
- Sales: Azure Cosmos DB using the Core (SQL) API
- Logs: Azure Table storage
- HR: Azure SQL database

You need to ensure that you can search all the data by using the Azure Cognitive Search REST API.

What should you do?

- A. Migrate the data in HR to Azure Blob storage.
- B. Migrate the data in HR to the on-premises SQL server.
- C. Export the data in Finance to Azure Data Lake Storage.
- D. Ingest the data in Logs into Azure Sentinel.

Correct Answer: C

In Azure Cognitive Search, a data source is used with indexers, providing the connection information for ad hoc or scheduled data refresh of a target index, pulling data from supported Azure data sources.

Note: Supported data sources -

Indexers crawl data stores on Azure and outside of Azure.

Amazon Redshift (in preview)

Azure Blob Storage -

Azure Cosmos DB -

Azure Data Lake Storage Gen2 -

Azure MySQL (in preview)

Azure SQL Database -

Azure Table Storage -

Elasticsearch (in preview)

PostgreSQL (in preview)

Salesforce Objects (in preview)

Salesforce Reports (in preview)

Smartsheet (in preview)

Snowflake (in preview)

Azure SQL Managed Instance -

SQL Server on Azure Virtual Machines

Azure Files (in preview)

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-indexer-overview#supported-data-sources>

Community vote distribution

C (100%)

  **zellck**  11 months ago

Selected Answer: C

C is the answer.

<https://learn.microsoft.com/en-us/azure/search/search-data-sources-gallery>
upvoted 6 times

  **zellck** 11 months ago

<https://learn.microsoft.com/en-us/azure/search/search-indexer-overview#supported-data-sources>

Indexers crawl data stores on Azure and outside of Azure.

- Azure Blob Storage
- Azure Cosmos DB
- Azure Data Lake Storage Gen2
- Azure SQL Database
- Azure Table Storage
- Azure SQL Managed Instance
- SQL Server on Azure Virtual Machines
- Azure Files (in preview)
- Azure MySQL (in preview)
- SharePoint in Microsoft 365 (in preview)
- Azure Cosmos DB for MongoDB (in preview)
- Azure Cosmos DB for Apache Gremlin (in preview)

upvoted 4 times

✉️  **rdemontis** 6 months, 3 weeks ago

thanks for the reference

upvoted 1 times

✉️  **funny_penguin** Most Recent 3 days, 23 hours ago

Selected Answer: C

on exam, C is correct.

upvoted 1 times

✉️  **Mehe323** 2 months ago

Selected Answer: C

Trick question, make sure you read it well. All of the data but that from Finance is already in the cloud.

upvoted 2 times

✉️  **evangelist** 4 months ago

Finance data stays on-premises now, it has to be migrated to Azure Cloud first in order to be indexed and searchable by Cognitive Service

upvoted 2 times

✉️  **Gvalli** 5 months, 1 week ago

A modified version of this was in the exam today.

upvoted 1 times

✉️  **rdemontis** 6 months, 3 weeks ago

Selected Answer: C

<https://learn.microsoft.com/en-us/azure/search/search-indexer-overview#supported-data-sources>

upvoted 1 times

✉️  **halfway** 1 year, 6 months ago

Selected Answer: C

<https://learn.microsoft.com/en-us/azure/search/search-indexer-overview>

upvoted 3 times

Question #30

SIMULATION -

Use the following login credentials as needed:

To enter your username, place your cursor in the Sign in box and click on the username below.

To enter your password, place your cursor in the Enter password box and click on the password below.

Azure Username: admin@abc.com -

Azure Password: XXXXXXXXXXXX -

The following information is for technical support purposes only:

Lab Instance: 12345678 -

Task -

You need to create and publish a Language Understanding (classic) model named 1u12345678. The model will contain an intent of Travel that has an utterance of

Boat.

To complete this task, sign in to the Language Understanding portal at <http://www.luis-ai/>.

Correct Answer: See explanation below.

Create your LUIS model -

1. You should navigate to your LUIS.ai management portal and create a new application. In the portal create a model.

Model name: 1u12345678 -

2. Define one intent as "Travel" and add an example utterances of Boat.

The screenshot shows the LUIS.ai management portal interface. The top navigation bar includes 'Scheduling (V 0.1)~', 'DASHBOARD', 'BUILD' (which is selected), 'MANAGE', and tabs for 'Train', 'Test', and 'Publish'. On the left, a sidebar lists 'App Assets' (Intents, Entities), 'Improve app performance' (Review endpoint utterances, Phrase lists, Patterns), and 'PREVIEW Prebuilt Domains'. The main area is titled 'Schedule appointment' with a red border. It contains a text input field with placeholder 'Type about 5 examples of what a user might say and hit Enter' and a red box highlighting several example utterances: 'i want to schedule with my doctor', 'can you book an appointment next week for me ?', 'how do i make a new booking ?', 'i want to schedule an appointment', and 'how do i book an appointment ?'. To the right of these utterances are 'Labeled intent ?' buttons and dropdown menus all set to 'Schedule ap...'. Below this section is a table titled 'Entities used in this intent ?' with columns for 'Name' and 'Labeled utterances'. A note states 'There are no entities in use.' At the bottom left of the main area is a 'PREVIEW Prebuilt Domains' button.

3. Publish the model

In order to use your model, you have to publish it. This is as easy as hitting the Publish tab, selecting between the production or staging environments, and hitting

Publish. As you can see from this page, you can also choose to enable sentiment analysis, speech priming to improve speech recognition, or the spell checker.

For now, you can leave those unchecked.

Reference:

https://docs.microsoft.com/en-us/azure/health-bot/language_model_howto <https://www.codemag.com/article/1809021/Natural-Language-Understanding-with-LUIS>

✉️ **takaimomoGcup** 5 days, 19 hours ago

If the simulation question is not on the actual exam, delete this question.

upvoted 1 times

✉️ **Sheetalsns** 5 months, 2 weeks ago

will we get simulation questions in exam -ai-102

upvoted 2 times

✉️ **ArminZ11** 4 months, 1 week ago

End of Dec 2023, no simulations in the exam

upvoted 2 times

✉️ **AnonymousJhb** 6 months ago

will / is this querstion updated? Since LUIS is deprecating and already encouraging users to begin using CLU rather when accessing <https://www.luis.ai/>.

upvoted 2 times

✉️ **ziggy1117** 11 months, 3 weeks ago

1. go to <https://language.cognitive.azure.com/>
2. create an intent Travel
3. create sample utterances with boat
4. train, test, deploy

upvoted 3 times

✉️ **rdemontis** 6 months, 3 weeks ago

agree with you

upvoted 1 times

Question #31**SIMULATION -**

Use the following login credentials as needed:

To enter your username, place your cursor in the Sign in box and click on the username below.

To enter your password, place your cursor in the Enter password box and click on the password below.

Azure Username: admin@abc.com -

Azure Password: XXXXXXXXXXXX -

The following information is for technical support purposes only:

Lab Instance: 12345678 -

Task -

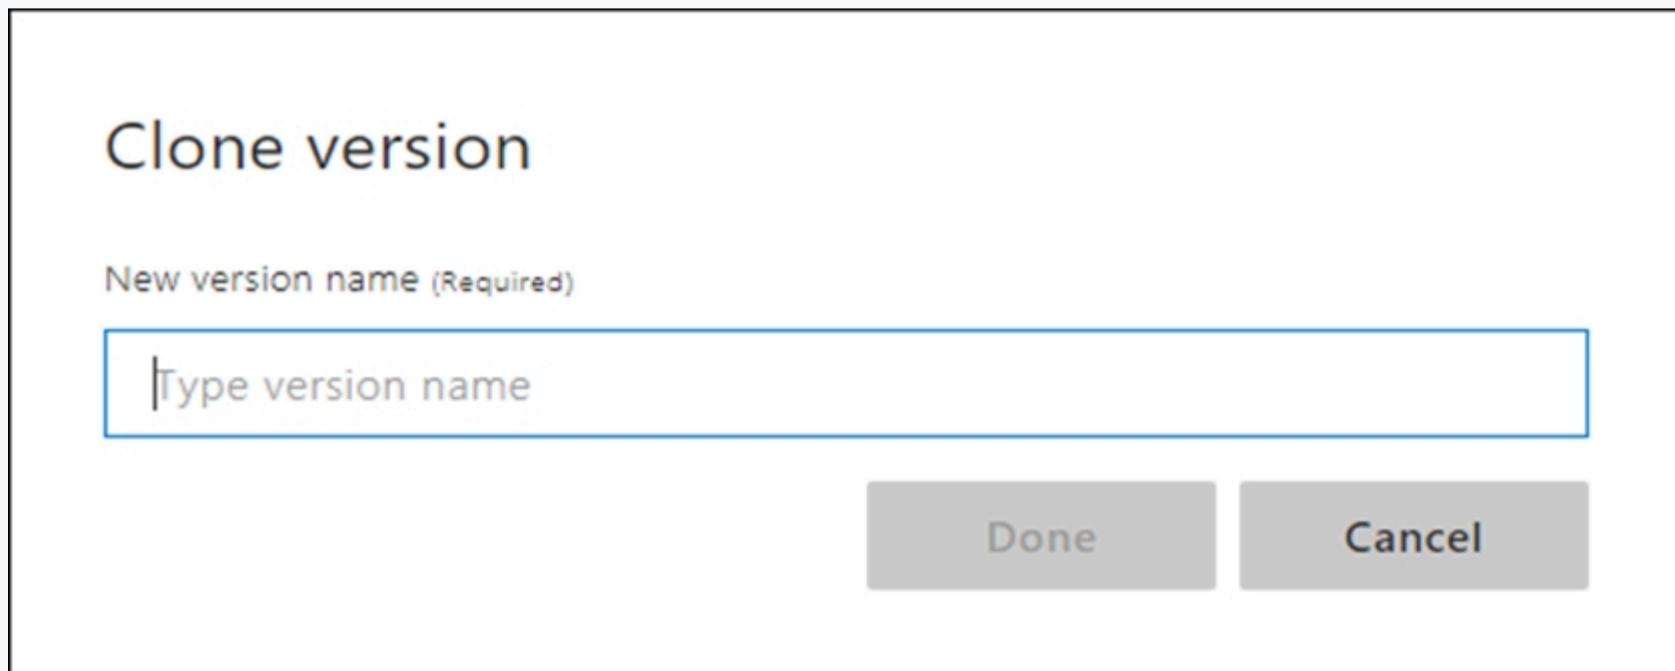
You need to create a version of the 1u12345678 Language Understanding (classic) model. The new version must have a version name of 1.0 and must be active.

To complete this task, sign in to the Language Understanding portal at <https://www.luis.ai/>.

Correct Answer: See explanation below.

Step 1: Clone a version -

1. Select the version you want to clone (1u12345678) then select Clone from the toolbar.
2. In the Clone version dialog box, type a name for the new version. Type 1.0



Step 2: Set active version -

Select a version from the list, then select Activate from the toolbar.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-how-to-manage-versions>

Question #32

You have a Language service resource that performs the following:

- Sentiment analysis
- Named Entity Recognition (NER)
- Personally Identifiable Information (PII) identification

You need to prevent the resource from persisting input data once the data is analyzed.

Which query parameter in the Language service API should you configure?

- A. model-version
- B. piiCategories
- C. showStats
- D. loggingOptOut

Correct Answer: D

Community vote distribution

D (100%)

✉️  **marti_tremblay000** Highly Voted 1 year, 2 months ago

Selected Answer: D

The LoggingOptOut parameter is true by default for the PII and health feature endpoints.

Reference <https://learn.microsoft.com/en-us/legal/cognitive-services/language-service/data-privacy>

upvoted 5 times

✉️  **marti_tremblay000** 1 year, 2 months ago

ChatGPT confirms :

To prevent the resource from persisting input data once the data is analyzed, you should configure the loggingOptOut query parameter in the Language service API. Setting the value of loggingOptOut to true will prevent the service from logging or storing the input data after analysis.

Therefore, the correct answer is D. loggingOptOut.

upvoted 3 times

✉️  **reiwanotora** Most Recent 3 days, 19 hours ago

Selected Answer: D

loggingOptOut Yes

upvoted 1 times

✉️  **evangelist** 3 months, 3 weeks ago

Selected Answer: D

no doubt answer is D

upvoted 2 times

✉️  **Gvalli** 5 months, 1 week ago

A modified version of this was in the exam today.

upvoted 1 times

✉️  **shuklabond007** 5 months, 1 week ago

Hi @Gavlli, May know how many simulation questions comes in AI 102 exam? also read through all the ET questions should be enough to clear the exam?

upvoted 2 times

✉️  **rdemontis** 6 months, 3 weeks ago

Selected Answer: D

Answer is correct

<https://learn.microsoft.com/en-us/legal/cognitive-services/language-service/data-privacy#how-is-data-retained-and-what-customer-controls-are-available>

upvoted 1 times

Question #33

You have an Azure Cognitive Services model named Model1 that identifies the intent of text input.

You develop an app in C# named App1.

You need to configure App1 to use Model1.

Which package should you add to App1?

- A. Universal.Microsoft.CognitiveServices.Speech
- B. SpeechServicesToolkit
- C. Azure.AI.Language.Conversations
- D. Xamarin.Cognitive.Speech

Correct Answer: A

Community vote distribution

C (100%)

 **evangelist**  3 months, 3 weeks ago

Selected Answer: C

Azure Conversational Language Understanding is part of Azure Cognitive Services and is designed to understand the intent of text input. The Azure.AI.Language.Conversations package provides C# developers with the client libraries needed to interact with these services, enabling you to send text to models and receive the results of intent and entity recognition.

Explanation of other options:

A. Universal.Microsoft.CognitiveServices.Speech is primarily used for speech services such as speech-to-text conversion, and not specifically for text intent recognition.

upvoted 5 times

 **reiwanotora**  3 days, 19 hours ago

Selected Answer: C

C is right answer.

Xamarin is done.

upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: C

<https://learn.microsoft.com/en-us/dotnet/api/overview/azure/ai.language.conversations-readme?view=azure-dotnet>

upvoted 2 times

 **zellck** 11 months ago

Selected Answer: C

C is the answer.

<https://learn.microsoft.com/en-us/dotnet/api/overview/azure/ai.language.conversations-readme?view=azure-dotnet>

Conversational Language Understanding - aka CLU for short - is a cloud-based conversational AI service which provides many language understanding capabilities like:

- Conversation App: It's used in extracting intents and entities in conversations

Start by importing the namespace for the ConversationAnalysisClient and related class:

- using Azure.AI.Language.Conversations;

upvoted 4 times

 **EliteAllen** 1 year ago

Selected Answer: C

C. Azure.AI.Language.Conversations

The Azure.AI.Language.Conversations package is part of the Azure SDK for .NET and is used for interacting with Azure Cognitive Services' Language Understanding (LUIS) models, which are used to identify the intent of text input. This makes it the correct choice for integrating Model1 into App1.

upvoted 2 times

 **Pfffff** 1 year, 1 month ago

Selected Answer: C

ChatGPT: The package you should add to App1 to use Model1 is C. Azure.AI.Language.Conversations.

Azure.AI.Language.Conversations is a package that provides the Language Understanding (LUIS) service, which can be used to identify the intent of text input. This package contains classes for authenticating with the LUIS service and sending text to the service to obtain intent and entity information.

upvoted 4 times

 **marti_tremblay000** 1 year, 2 months ago

ChatGPT says the correct answer isn't even listed :

To configure App1 to use Model1, you should add the Azure.AI.TextAnalytics package to the project. This package provides the necessary libraries and functionality to integrate with Azure Cognitive Services text analytics models such as Model1.

Therefore, the correct answer is not listed among the options. The correct answer is Azure.AI.TextAnalytics.

upvoted 2 times

 **marti_tremblay000** 1 year, 2 months ago

Selected Answer: C

The question is about the intent of text input. It has nothing to do with speech.

Therefore the Azure.AI.Language.Conversations is the answer :

Conversation Analysis is a cloud-based conversational AI service that applies custom machine-learning intelligence to a user's conversational, natural language text to predict overall meaning, and pull out relevant, detailed information.

Reference : <https://learn.microsoft.com/en-us/samples/azure/azure-sdk-for-net/azureailanguagconversations-samples/>

upvoted 2 times

Question #34

HOTSPOT

You are building content for a video training solution.

You need to create narration to accompany the video content. The solution must use Custom Neural Voice.

What should you use to create a custom neural voice, and which service should you use to generate the narration? To answer, select the appropriate options in the answer area.

NOTE: Each correct answer is worth one point.

Answer Area

Custom neural voice:

- Microsoft Bot Framework Composer
- The Azure portal
- The Language Understanding portal
- The Speech Studio portal

Narration:

- Language Understanding
- Speaker Recognition
- Speech-to-text
- Text-to-speech

Answer Area

Custom neural voice:

- Microsoft Bot Framework Composer
- The Azure portal
- The Language Understanding portal
- The Speech Studio portal

Correct Answer:

Narration:

- Language Understanding
- Speaker Recognition
- Speech-to-text
- Text-to-speech

  zellck  11 months ago

1. Speech Studio portal
2. Text-to-speech

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/custom-neural-voice#how-does-it-work>

To create a custom neural voice, use Speech Studio to upload the recorded audio and corresponding scripts, train the model, and deploy the voice to a custom endpoint.

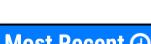
<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/text-to-speech>

Text to speech enables your applications, tools, or devices to convert text into humanlike synthesized speech. The text to speech capability is also known as speech synthesis. Use humanlike prebuilt neural voices out of the box, or create a custom neural voice that's unique to your product or brand.

upvoted 10 times

  rdemontis 6 months, 3 weeks ago

thanks for explanation
upvoted 1 times

  reiwanotora  3 days, 19 hours ago

The Speech Studio portal
Text-to-speech
upvoted 1 times

 **anto69** 2 months, 4 weeks ago

Absolutely correct
upvoted 2 times

 **evangelist** 3 months, 3 weeks ago

For Creating a Custom Neural Voice:

Use: The Speech Studio Portal

Explanation: The Speech Studio portal is part of Microsoft Azure's Cognitive Services and is specifically designed for speech-related applications, including the creation of Custom Neural Voices.

For Generating Narration:

Use: Text-to-Speech (TTS) Service

Explanation: Once you have created a Custom Neural Voice model in the Speech Studio portal, you can use Azure's Text-to-Speech service to generate narration. The Text-to-Speech service converts written text into spoken words in a natural-sounding voice, leveraging the custom voice model you've created.

upvoted 3 times

 **marti_tremblay000** 1 year, 2 months ago

The answer is correct. Speech Studio and text to speech.

upvoted 2 times

Question #35

HOTSPOT

You are building a call handling system that will receive calls from French-speaking and German-speaking callers. The system must perform the following tasks:

- Capture inbound voice messages as text.
- Replay messages in English on demand.

Which Azure Cognitive Services services should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

To capture messages:

- Speaker Recognition
- Speech-to-text
- Text-to-speech
- Translator

To replay messages:

- Speech-to-text only
- Speech-to-text and Language
- Speaker Recognition and Language
- Text-to-speech and Language
- Text-to-speech and Translator

Answer Area

To capture messages:

- Speaker Recognition
- Speech-to-text
- Text-to-speech
- Translator

Correct Answer:

To replay messages:

- Speech-to-text only
- Speech-to-text and Language
- Speaker Recognition and Language
- Text-to-speech and Language
- Text-to-speech and Translator

 **zellck** Highly Voted  11 months ago

1. Speech-to-text
2. Text-to-speech and Translator

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/speech-to-text>

With real-time speech to text, the audio is transcribed as speech is recognized from a microphone or file.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/text-to-speech>

Text to speech enables your applications, tools, or devices to convert text into humanlike synthesized speech. The text to speech capability is also known as speech synthesis. Use humanlike prebuilt neural voices out of the box, or create a custom neural voice that's unique to your product or brand.

upvoted 14 times

✉️ **rdemontis** 6 months, 3 weeks ago

agree. The answer is correct

upvoted 2 times

✉️ **takaimomoGcup** Most Recent ⓘ 5 days, 19 hours ago

Speech-to-text

Text-to-speech and Translator

upvoted 1 times

✉️ **Gvalli** 5 months, 1 week ago

A modified version of this was in the exam today.

upvoted 2 times

✉️ **ziggy1117** 11 months, 3 weeks ago

answer is correct

upvoted 3 times

Question #36

You are building a social media extension that will convert text to speech. The solution must meet the following requirements:

- Support messages of up to 400 characters.
- Provide users with multiple voice options.
- Minimize costs.

You create an Azure Cognitive Services resource.

Which Speech API endpoint provides users with the available voice options?

- A. <https://uksouth.api.cognitive.microsoft.com/speechtotext/v3.0/models/base>
- B. <https://uksouth.customvoice.api.speech.microsoft.com/api/texttospeech/v3.0/longaudiosynthesis/voices>
- C. <https://uksouth.tts.speech.microsoft.com/cognitiveservices/voices/list>
- D. <https://uksouth.voice.speech.microsoft.com/cognitiveservices/v1?deploymentId={deploymentId}>

Correct Answer: D

Community vote distribution

0 (96%) 4%

✉  **marti_tremblay000**  1 year, 2 months ago

Selected Answer: C

The correct answer is C

The question is about providing users with all the available voice options.

Get a list of voices

You can use the tts.speech.microsoft.com/cognitiveservices/voices/list endpoint to get a full list of voices for a specific region or endpoint.

Reference : <https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/rest-text-to-speech?tabs=streaming>

upvoted 11 times

✉  **MaliSanFuu** 1 year ago

Agreeing, as the only important question is: Which Speech API endpoint provides users with the available voice options?

Therefor answer C should be correct for this one

upvoted 2 times

✉  **zellck**  11 months ago

Selected Answer: C

C is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/rest-text-to-speech?tabs=streaming#get-a-list-of-voices>
ou can use the tts.speech.microsoft.com/cognitiveservices/voices/list endpoint to get a full list of voices for a specific region or endpoint. Prefix the voices list endpoint with a region to get a list of voices for that region. For example, to get a list of voices for the westus region, use the <https://westus.tts.speech.microsoft.com/cognitiveservices/voices/list> endpoint.

upvoted 7 times

✉  **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 4 times

✉  **reiwanotora**  3 days, 19 hours ago

Selected Answer: C

C is right answer.

upvoted 1 times

✉  **evangelist** 3 months, 3 weeks ago

Selected Answer: C

only C meets demand of minimum cost and support 400 characters with voice options

upvoted 1 times

✉  **Gvalli** 5 months, 1 week ago

A modified version of this was in the exam today.

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: C

C is the correct answer

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/rest-text-to-speech?tabs=streaming>

upvoted 1 times

 **james2033** 9 months, 2 weeks ago

Selected Answer: C

The answer is very clear at here <https://learn.microsoft.com/en-us/azure/ai-services/speech-service/rest-text-to-speech?tabs=streaming#list-voices>

You can use the tts.speech.microsoft.com/cognitiveservices/voices/list endpoint

For example, to get a list of voices for the westus region, use the <https://westus.tts.speech.microsoft.com/cognitiveservices/voices/list> endpoint.

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/regions>

Europe - UK South - uksouth

upvoted 1 times

 **Pfffff** 1 year, 1 month ago

Selected Answer: B

ChatGPT: The Speech API endpoint that provides users with the available voice options is B.

upvoted 1 times

Question #37

You develop a custom question answering project in Azure Cognitive Service for Language. The project will be used by a chatbot.

You need to configure the project to engage in multi-turn conversations.

What should you do?

- A. Add follow-up prompts.
- B. Enable active learning.
- C. Add alternate questions.
- D. Enable chit-chat.

Correct Answer: A

Community vote distribution

A (100%)

✉️  **zellck**  11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/question-answering/overview#multi-turn-conversations>
Question answering provides multi-turn prompts and active learning to help you improve your basic question and answer pairs.

Multi-turn prompts give you the opportunity to connect question and answer pairs. This connection allows the client application to provide a top answer and provides more questions to refine the search for a final answer.

upvoted 9 times

✉️  **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 4 times

✉️  **reiwanotora**  3 days, 19 hours ago

Selected Answer: A

A is right answer.

upvoted 1 times

✉️  **evangelist** 3 months, 3 weeks ago

Selected Answer: A

no doubt answer is A

upvoted 1 times

✉️  **evangelist** 4 months ago

Selected Answer: A

no doubt follow up prompt

upvoted 1 times

✉️  **evangelist** 4 months ago

Selected Answer: A

Multi-turn conversations in the context of chatbots or conversational AI refer to interactions where the dialogue between the user and the bot extends beyond a single question and response. In these conversations, the bot and the user exchange multiple messages, and the bot's responses are contextually dependent on the previous parts of the conversation.

upvoted 1 times

✉️  **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

To configure a project for multi-turn conversations in Azure Cognitive Services for Language, you should:

- A. Add follow-up prompts.

Adding follow-up prompts allows your question answering model to engage in multi-turn conversations by providing responses and asking for clarifications or additional information when a user's query is ambiguous or incomplete. This is a key feature for enabling conversational interactions with the chatbot. (ChatGPT)

upvoted 1 times

✉  **tzuyichao** 1 year, 1 month ago

Ref: <https://learn.microsoft.com/en-us/azure/cognitive-services/qnamaker/how-to/multi-turn>

upvoted 2 times

✉  **marti_tremblay000** 1 year, 2 months ago

Selected Answer: A

ChatGPT confirms :

To configure the project to engage in multi-turn conversations, you should add follow-up prompts. Follow-up prompts are a way to ask additional questions or provide more information to help the user clarify their intent. By adding follow-up prompts, the chatbot can engage in a back-and-forth conversation with the user to gather additional information and ultimately provide a better answer.

Therefore, the correct answer is A. Add follow-up prompts.

upvoted 1 times

Question #38

HOTSPOT

You are building a solution that students will use to find references for essays.

You use the following code to start building the solution.

```
using Azure;
using System;
using Azure.AI.TextAnalytics;

private static readonly AzureKeyCredential credentials = new AzureKeyCredential("<key>");
private static readonly Uri endpoint = new Uri("<endpoint>");

static void EntityLinker(TextAnalyticsClient client)
{
    var response = client.RecognizeLinkedEntities(
        "Our tour guide took us up the Space Needle during our trip to Seattle last week.");
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
The code will detect the language of documents.	<input type="radio"/>	<input type="radio"/>
The <code>url</code> attribute returned for each linked entity will be a Bing search link.	<input type="radio"/>	<input type="radio"/>
The <code>matches</code> attribute returned for each linked entity will provide the location in a document where the entity is referenced.	<input type="radio"/>	<input type="radio"/>

Answer Area		
Statements	Yes	No
The code will detect the language of documents.	<input type="radio"/>	<input checked="" type="radio"/>
Correct Answer: The <code>url</code> attribute returned for each linked entity will be a Bing search link.	<input checked="" type="radio"/>	<input type="radio"/>
The <code>matches</code> attribute returned for each linked entity will provide the location in a document where the entity is referenced.	<input checked="" type="radio"/>	<input type="radio"/>

 zellck Highly Voted 11 months ago

NNY is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/entity-linking/overview>

Entity linking is one of the features offered by Azure Cognitive Service for Language, a collection of machine learning and AI algorithms in the cloud for developing intelligent applications that involve written language. Entity linking identifies and disambiguates the identity of entities found in text. For example, in the sentence "We went to Seattle last week.", the word "Seattle" would be identified, with a link to more information on Wikipedia.

upvoted 15 times

 WhyWhyYellowYellow Highly Voted 1 year, 1 month ago

Y

N

Y

<https://learn.microsoft.com/en-us/rest/api/cognitiveservices-textanalytics/3.0/entities-linking/entities-linking?tabs=HTTP>
upvoted 11 times

✉ **ctrung** 11 months, 1 week ago

Agreed. The json file consists of language attributes indicating the detected language.
For the second statement, it's not always Bing search link.

upvoted 3 times

✉ **reiwanotora** **Most Recent** 3 days, 19 hours ago

No No Yes

upvoted 1 times

✉ **JamesKJoker** 1 week, 1 day ago

While the MS Learn specifically mention Wikipedia, the Class Doc for LinkedEnty refers to Bing. <https://learn.microsoft.com/en-us/java/api/com.azure.ai.textanalytics.models.linkedentity?view=azure-java-stable>

So Bing answers are correct.

=> NYY

upvoted 2 times

✉ **JamesKJoker** 1 week, 1 day ago

The code will detect the language of documents.

No. The provided code doesn't include any language detection functionality. Azure Text Analytics does have separate language detection capabilities, but they aren't used in this snippet.

The url attribute returned for each linked entity will be a Bing search link.

No. The RecognizeLinkedEntities function in Azure Text Analytics usually returns links to more authoritative sources like Wikipedia, not general Bing search results.

The matches attribute returned for each linked entity will provide the location in a document where the entity is referenced.

Yes. The matches attribute within the response from RecognizeLinkedEntities provides details about where each recognized entity occurs in the input text. This includes information like character offsets, allowing you to locate the entity within the original text.

upvoted 1 times

✉ **Murtuza** 1 month, 4 weeks ago

No, the code you provided does not detect the language of documents. It uses the RecognizeLinkedEntities method of the TextAnalyticsClient to identify and link entities in the text to more information on the web.

Yes, the url attribute returned for each linked entity will be a Bing search link. This link provides more information about the entity.

Yes, the matches attribute returned for each linked entity will provide the location in the document where the entity is referenced. It includes the text of the entity, its offset (the number of characters from the beginning of the document to the start of the entity), and its length (the number of characters of the entity).

upvoted 1 times

✉ **varinder82** 2 months ago

Final Answer:

N

N

Y

upvoted 1 times

✉ **audlindr** 2 months, 4 weeks ago

NNY

Tried out and the response does not detect language. it just returns the language code sent in the request and only english and Spanish is supported.

The version 3.1 returns Bing ID however all links returned are only wiki links

upvoted 3 times

✉ **evangelist** 3 months, 3 weeks ago

The code will detect the language of documents

No. The provided code uses the RecognizeLinkedEntities method, which recognizes and classifies named entities in a document and links them to more information on the web.

Yes. The RecognizeLinkedEntities method returns a list of recognized entities linked to more information on the web, typically in the form of a Bing search URL.

Yes. For each linked entity recognized by the RecognizeLinkedEntities method, the matches attribute contains information about occurrences of that entity within the input text.

upvoted 1 times

✉ **idcanymore** 3 months, 2 weeks ago

the question says "will", and not "typically". The question wants to know if it will ALWAYS be a bing link.

upvoted 1 times

✉ **Brandalf** 4 months ago

Incorrect image is displayed for me

upvoted 1 times

✉ **sismér** 6 months ago

N
N
Y

he language that the document is written in. If unspecified, this value will be set to the default language in DefaultLanguage in the request sent to the service. If set to an empty string, the service will apply a model where the language is explicitly set to "None".

upvoted 3 times

✉ **rdemontis** 6 months, 3 weeks ago

YNY

<https://learn.microsoft.com/en-us/dotnet/api/azure.ai.textanalytics.linkedentity?view=azure-dotnet>

upvoted 1 times

✉ **ExamDev** 8 months, 2 weeks ago

Properties of the Linked entity are: <https://learn.microsoft.com/en-us/dotnet/api/azure.ai.textanalytics.linkedentity?view=azure-dotnet>

upvoted 1 times

✉ **M25** 8 months, 4 weeks ago

Y, N, Y

<https://learn.microsoft.com/en-us/dotnet/api/azure.ai.textanalytics.textanalyticsclient.recognizelinkedentities?view=azure-dotnet>

Y – language String

The language that the document is written in. If unspecified, this value will be set to the default language in DefaultLanguage in the request sent to the service.

Default language value to use in all client calls. If no value is specified, "en" is set as default.

<https://learn.microsoft.com/en-us/rest/api/cognitiveservices-textanalytics/3.0/entities-linking/entities-linking?tabs=HTTP#linkedentity>

N – LinkedEntity dataSource string

Data source used to extract entity linking, such as Wiki/Bing etc.

<https://learn.microsoft.com/en-us/rest/api/cognitiveservices-textanalytics/3.0/entities-linking/entities-linking?tabs=HTTP#match>

Y – Match offset integer

Start position (in Unicode characters) for the entity match text.

upvoted 4 times

✉ **ziggy1117** 11 months, 3 weeks ago

Answer is correct.

N - code will not detect language

Y - Data source used to extract entity linking, such as Wiki/Bing etc. URL for the entity's page from the data source.

Y - matches: List of instances this entity appears in the text.

<https://learn.microsoft.com/en-us/rest/api/cognitiveservices-textanalytics/3.0/entities-linking/entities-linking?tabs=HTTP>

upvoted 1 times

✉ **MDawson** 1 year, 1 month ago

documentation says url CAN be bing, although the examples all show Wikipedia links. Confusing...

upvoted 2 times

Question #39

You train a Conversational Language Understanding model to understand the natural language input of users.

You need to evaluate the accuracy of the model before deploying it.

What are two methods you can use? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. From the language authoring REST endpoint, retrieve the model evaluation summary.
- B. From Language Studio, enable Active Learning, and then validate the utterances logged for review.
- C. From Language Studio, select Model performance.
- D. From the Azure portal, enable log collection in Log Analytics, and then analyze the logs.

Correct Answer: AC*Community vote distribution*

✉️ **ziggy1117** 11 months, 3 weeks ago

Selected Answer: AC

agree with SmallFire:

Active Learning cannot be initiated prior to the deployment of the model. The primary purpose of the 'Active Learning' feature is to leverage actual user interaction data to enhance the model's understanding capabilities. This is a continuous learning and optimization process that takes place after the model has been deployed and put into actual use.

so the answer is AC.

upvoted 8 times

✉️ **takaimomoGcup** 5 days, 19 hours ago

Selected Answer: AC

REST and model performance.

upvoted 1 times

✉️ **evangelist** 4 months ago

Selected Answer: AC

A. From the language authoring REST endpoint, retrieve the model evaluation summary.

This summary typically includes metrics like precision, recall, and accuracy, which are crucial for evaluating the effectiveness of a language understanding model.

C. From Language Studio, select Model performance.

In Language Studio, the Model performance section typically provides detailed analytics about the model's performance, including various metrics and possibly confusion matrices.

upvoted 3 times

✉️ **rdemontis** 6 months, 3 weeks ago

Selected Answer: AC

IMHO correct answers are:

A. --> <https://learn.microsoft.com/en-us/rest/api/language/conversational-analysis-authoring/get-model-evaluation-summary?view=rest-language-2023-04-01&tabs=HTTP>

C. --> <https://learn.microsoft.com/en-us/azure/ai-services/language-service/conversational-language-understanding/how-to/view-model-evaluation?tabs=Language-studio%2Cmodel-performance>

upvoted 2 times

✉️ **sl_mslconsulting** 7 months, 2 weeks ago

Selected Answer: AC

If you look closely, you can see that the Model performance feature in Language Studio is about evaluating the performance of the model using test data as opposed to active learning which is using the real data from users' interactions, which you couldn't do prior to the deployment. My only pet peeve is that rest endpoint never works for me - always giving me the 404 no matter what.

upvoted 1 times

✉️ **jangotango** 8 months, 2 weeks ago

AC are correct - <https://learn.microsoft.com/en-us/rest/api/language/2022-10-01-preview/text-analysis-authoring/get-model-evaluation-summary?tabs=HTTP>

upvoted 2 times

zellck 11 months ago

Selected Answer: AC

AC is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/conversational-language-understanding/how-to/view-model-evaluation?tabs=Language-studio%2Cmodel-performance>

upvoted 4 times

EliteAllen 11 months, 2 weeks ago

Selected Answer: BC

B. Active Learning in Language Studio is a feature that helps improve the performance of your model by suggesting utterances for you to review and label. This can help you evaluate the accuracy of your model by seeing how it performs on these suggested utterances.

C. The Model performance feature in Language Studio provides a detailed report on the performance of your model, including precision, recall, and F1 score. This can give you a good indication of the accuracy of your model.

Option A is incorrect because the language authoring REST endpoint does not provide a model evaluation summary.

upvoted 2 times

hawkzey 9 months, 1 week ago

B is not correct as you cannot do active learning on a model that is not yet deployed

upvoted 2 times

SmallFire 12 months ago

Active Learning cannot be initiated prior to the deployment of the model. The primary purpose of the 'Active Learning' feature is to leverage actual user interaction data to enhance the model's understanding capabilities. This is a continuous learning and optimization process that takes place after the model has been deployed and put into actual use.

so the answer is AC.

upvoted 3 times

sheldon73 1 year ago

Selected Answer: BC

Google Bard Answer : Sure, here are two methods you can use to evaluate the accuracy of a Conversational Language Understanding model before deploying it:

From Language Studio, select Model performance. This will show you a summary of the model's performance, including the F1 score, precision, and recall.

From Language Studio, enable Active Learning, and then validate the utterances logged for review. This will allow you to manually review utterances that the model has misclassified, and then retrain the model with the corrected data.

Here are the correct answers to your question:

C. From Language Studio, select Model performance.

B. From Language Studio, enable Active Learning, and then validate the utterances logged for review.

upvoted 1 times

Question #40

DRAG DROP

You develop an app in C# named App1 that performs speech-to-speech translation.

You need to configure App1 to translate English to German.

How should you complete the SpeechTranslationConfig object? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Values	Answer Area
addTargetLanguage	
speechSynthesisLanguage	
speechRecognitionLanguage	
voiceName	

```
var translationConfig = SpeechTranslationConfig.FromSubscription(SPEECH_SUBSCRIPTION_KEY, SPEECH_SERVICE_REGION);
    translationConfig.  = "en-US";
    translationConfig.  ("de");
```

Correct Answer:

```
var translationConfig = SpeechTranslationConfig.FromSubscription(SPEECH_SUBSCRIPTION_KEY, SPEECH_SERVICE_REGION);
    translationConfig.  = "en-US";
    translationConfig.  ("de");
```

✉ MaliSanFuu Highly Voted 1 year ago

Agreeing with @WinzigWeich

Answer should be:

- 1) SpeechRecognitionLanguage
 - 2) AddTargetLanguage
- upvoted 16 times

✉ ziggy1117 Highly Voted 11 months, 3 weeks ago

- 1) SpeechRecognitionLanguage
- 2) AddTargetLanguage

in the exercise of AI-102 online learning
upvoted 8 times

✉ takaimomoGcup Most Recent 5 days, 18 hours ago

SpeechRecognitionLanguage
AddTargetLanguage
upvoted 1 times

✉ sca88 6 months ago

- 1) SpeechRecognitionLanguage
- 2) AddTargetLanguage

<https://learn.microsoft.com/en-us/dotnet/api/microsoft.cognitiveservices.speech.speechtranslationconfig?view=azure-dotnet>
upvoted 1 times

✉ zellck 11 months ago

1. SpeechRecognitionLanguage
2. AddTargetLanguage

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-translate-speech?tabs=terminal&pivots=programming-language-csharp#change-the-source-language>
One common task of speech translation is specifying the input (or source) language. In your code, interact with the SpeechTranslationConfig instance by assigning it to the SpeechRecognitionLanguage property:

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-translate-speech?tabs=terminal&pivots=programming-language-csharp#add-a-translation-language>

Another common task of speech translation is to specify target translation languages. At least one is required, but multiples are supported. With every call to AddTargetLanguage, a new target translation language is specified. In other words, when speech is recognized from the source language, each target translation is available as part of the resulting translation operation.

upvoted 6 times

✉️ **rdemontis** 6 months, 3 weeks ago

thanks for explanation and the provided relevant documentation

upvoted 2 times

✉️ **WinzigWeich** 1 year ago

<https://microsoftlearning.github.io/AI-102-AIEngineer.de-de/Instructions/08-translate-speech.html>

C# Part

translationConfig.SpeechRecognitionLanguage

translationConfig.AddTargetLanguage

upvoted 5 times

Question #41

Topic 3

You have an Azure subscription that contains an Azure Cognitive Service for Language resource.

You need to identify the URL of the REST interface for the Language service.

Which blade should you use in the Azure portal?

- A. Identity
- B. Keys and Endpoint
- C. Networking
- D. Properties

Correct Answer: B

Community vote distribution

B (100%)

✉️ **973b658** Highly Voted 11 months, 3 weeks ago

Selected Answer: B

It is B.

upvoted 8 times

✉️ **takaimomoGcup** Most Recent 5 days, 19 hours ago

Selected Answer: B

B is right answer.

upvoted 1 times

✉️ **evangelist** 3 months, 3 weeks ago

Selected Answer: B

B. Keys and Endpoint blade in the Azure portal.

This blade provides the endpoint URL needed to access the Cognitive Services API, along with the keys required for authentication. The endpoint URL is essential for making API calls to the service, including those for Language features such as sentiment analysis, key phrase extraction, named entity recognition, and more.

upvoted 3 times

✉️ **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

correct

upvoted 1 times

✉️ **chenglim** 7 months ago

Selected Answer: B

ans correct

upvoted 1 times

Question #42

DRAG DROP

You are building a transcription service for technical podcasts.

Testing reveals that the service fails to transcribe technical terms accurately.

You need to improve the accuracy of the service.

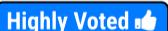
Which five actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

- Deploy the model.
- Create a Custom Speech project.
- Upload training datasets.
- Create a speech-to-text model.
- Create a Speaker Recognition model.
- Train the model.
- Create a Conversational Language Understanding model.

Answer Area

Answer Area	
Correct Answer:	<ul style="list-style-type: none">Create a Custom Speech project.Create a speech-to-text model.Upload training datasets.Train the model.Deploy the model.

✉  zellck  11 months ago

1. Create Custom Speech project
2. Create speech-to-text model
3. Upload training datasets
4. Train model
5. Deploy model

upvoted 17 times

✉  zellck 11 months ago

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/custom-speech-overview#how-does-it-work>
With Custom Speech, you can upload your own data, test and train a custom model, compare accuracy between models, and deploy a model to a custom endpoint.
- Create a project and choose a model. Use a Speech resource that you create in the Azure portal. If you will train a custom model with audio data, choose a Speech resource region with dedicated hardware for training audio data.
- Upload test data. Upload test data to evaluate the speech to text offering for your applications, tools, and products.
- Train a model. Provide written transcripts and related text, along with the corresponding audio data. Testing a model before and after training is optional but recommended.
- Deploy a model. Once you're satisfied with the test results, deploy the model to a custom endpoint. With the exception of batch transcription, you must deploy a custom endpoint to use a Custom Speech model.

upvoted 5 times

✉️ **rdemontis** 6 months, 3 weeks ago
thanks for the provided references
upvoted 1 times

✉️ **f2c587e** **Most Recent** 2 months ago
1. Create a Custom Voice Project
2. Create a speech-to-text model
3. Upload Training Datasets
4. Training Model
5. Implementation model
upvoted 2 times

✉️ **f2c587e** 2 months ago
According to the answer, then data should not be uploaded to train the model? Seriously? So how do you plan to train yourself if they're suppose to be technical words. I agree with zellck
upvoted 1 times

✉️ **rdemontis** 6 months, 3 weeks ago
correct answer
upvoted 1 times

✉️ **973b658** 11 months, 3 weeks ago
It is true.
upvoted 3 times

Question #43

You are building a retail kiosk system that will use a custom neural voice.

You acquire audio samples and consent from the voice talent.

You need to create a voice talent profile.

What should you upload to the profile?

- A. a .zip file that contains 10-second .wav files and the associated transcripts as .txt files
- B. a five-minute .flac audio file and the associated transcript as a .txt file
- C. a .wav or .mp3 file of the voice talent consenting to the creation of a synthetic version of their voice
- D. a five-minute .wav or .mp3 file of the voice talent describing the kiosk system

Correct Answer: C

Community vote distribution

C (75%)

A (25%)

 **Murtuza** 1 month, 2 weeks ago

Selected Answer: C

C is the correct answer
upvoted 1 times

 **Tactable** 3 months, 3 weeks ago

Selected Answer: C

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/how-to-custom-voice-talent>
upvoted 1 times

 **evangelist** 3 months, 3 weeks ago

Selected Answer: C

Based on the Azure AI documentation, the correct option for creating a voice talent profile for a custom neural voice is:

C. a .wav or .mp3 file of the voice talent consenting to the creation of a synthetic version of their voice.

This is because the documentation specifies the need for a recording of the voice talent's consent statement, acknowledging the use of their voice recordings by a specified company to create and use a synthetic version of their voice
upvoted 2 times

 **tdctdc** 5 months, 4 weeks ago

Selected Answer: C

The question is about the profile, not data.
upvoted 4 times

 **RupRizal** 5 months, 4 weeks ago

C is the correct answer as only wav & mp3 formats are allowed. Zip is not allowed.
upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: C

here we are requested to add a new voice talent profile. So a consent statement is needed.
<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/how-to-custom-voice-talent>
upvoted 1 times

 **katrang** 7 months, 1 week ago

Selected Answer: C

I initially thought A, but now think it is C
<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/how-to-custom-voice-talent>
upvoted 1 times

 **AnonymousJhb** 7 months, 3 weeks ago

Selected Answer: C

zip is not allowed. only wav & mp3.
And made synthetic. =C

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/how-to-custom-voice-talent>
follow the pictures very carefully.

upvoted 3 times

✉️ **ExamDev** 8 months, 2 weeks ago

Selected Answer: A

The correct answer is A! From the documentation: <https://learn.microsoft.com/en-us/azure/ai-services/speech-service/how-to-custom-voice-training-data>

Follow these guidelines when preparing audio.

Property Value

File format RIFF (.wav), grouped into a .zip file

File name File name characters supported by Windows OS, with .wav extension.

The characters \ / : * ? " < > | aren't allowed.

It can't start or end with a space, and can't start with a dot.

No duplicate file names are allowed.

Sampling rate When creating a custom neural voice, 24,000 Hz is required.

Sample format PCM, at least 16-bit

Audio length Shorter than 15 seconds

Archive format .zip

Maximum archive size 2048 MB

upvoted 1 times

✉️ **M25** 8 months, 4 weeks ago

Selected Answer: C

In Q#13 (How should you upload the samples?), the selected answer as option B (for speech samples to be used in training) corresponds to this Q#43's option A.

Q#43 (What should you upload to the profile?) though, refers specifically to the voice talent's profile, that should contain an audio sample of the given consent, and hereto "Only .wav and .mp3 files are accepted" (at the bottom of the screenshot provided in the link).

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/how-to-custom-voice-talent#add-voice-talent>

upvoted 1 times

✉️ **[Removed]** 9 months, 2 weeks ago

Selected Answer: C

I would have chosen A but it's actually C because of <https://learn.microsoft.com/en-us/azure/ai-services/speech-service/how-to-custom-voice-talent>

"A voice talent is an individual or target speaker whose voices are recorded and used to create neural voice models.

Before you can train a neural voice, you must submit a recording of the voice talent's consent statement. The voice talent statement is a recording of the voice talent reading a statement that they consent to the usage of their speech data to train a custom voice model. The consent statement also used to verify that the voice talent is the same person as the speaker in the training data."

upvoted 2 times

✉️ **zellck** 11 months ago

Selected Answer: C

C should be the correct answer instead.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-custom-voice-talent#add-voice-talent>

- On the Upload voice talent statement page, follow the instructions to upload the voice talent statement you've recorded beforehand. Make sure the verbal statement was recorded with the same settings, environment, and speaking style as your training data.

- Enter the voice talent name and company name. The voice talent name must be the name of the person who recorded the consent statement.

The company name must match the company name that was spoken in the recorded statement.

upvoted 4 times

✉️ **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-custom-voice-training-data#types-of-training-data>

A voice training dataset includes audio recordings, and a text file with the associated transcriptions. Each audio file should contain a single utterance (a single sentence or a single turn for a dialog system), and be less than 15 seconds long.

- Individual utterances + matching transcript

A collection (.zip) of audio files (.wav) as individual utterances. Each audio file should be 15 seconds or less in length, paired with a formatted transcript (.txt).

upvoted 2 times

✉️ **Pixelmate** 11 months ago

A is correct answer. Training data for custom-neural-voice format has to be in .wav and .zip files as per the below documentation.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-custom-voice-training-data>

upvoted 1 times

✉️ **EliteAllen** 11 months, 2 weeks ago

Selected Answer: A

A. a .zip file that contains 10-second .wav files and the associated transcripts as .txt files

When creating a voice talent profile for a custom neural voice in Azure Cognitive Services, you need to upload a .zip file that contains audio samples from the voice talent and the associated transcripts. The audio samples should be in .wav format and each sample should be approximately 10 seconds long. The transcripts should be in .txt format and should match the spoken content in the audio samples. This data is used to train the custom neural voice model to mimic the voice talent's unique speech patterns.

upvoted 4 times

 **sheldon73** 11 months, 3 weeks ago

Selected Answer: C

Source : <https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/custom-neural-voice>

upvoted 3 times

 **mVic** 1 year ago

Selected Answer: C

C is correct, as per below documentation.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-custom-voice-talent>

upvoted 4 times

Question #44

DRAG DROP

You have a Language Understanding solution that runs in a Docker container.

You download the Language Understanding container image from the Microsoft Container Registry (MCR).

You need to deploy the container image to a host computer.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

From the host computer, move the package file to the Docker input directory.

From the Language Understanding portal, export the solution as a package file.

From the host computer, build the container and specify the output directory.

From the host computer, run the container and specify the input directory.

From the Language Understanding portal, retrain the model.

Answer Area

Answer Area

From the Language Understanding portal, export the solution as a package file.

Correct Answer:

From the host computer, move the package file to the Docker input directory.

From the host computer, run the container and specify the input directory.

zellck Highly Voted 11 months ago

1. From portal, export solution as package file.
2. From host computer, move package file to Docker input directory.
3. From host computer, run container and specify input directory.

<https://learn.microsoft.com/en-us/azure/cognitive-services/luis/luis-container-howto?tabs=v3#how-to-use-the-container>

- Export package for container from LUIS portal or LUIS APIs.
- Move package file into the required input directory on the host computer. Do not rename, alter, overwrite, or decompress the LUIS package file.
- Run the container, with the required input mount and billing settings.

upvoted 30 times

M25 8 months, 4 weeks ago

@zellck: You're a role model for contributing to community discussions, love this style! Well-documented answers, plus the corresponding link to follow-up and form an opinion independently!

upvoted 15 times

rdemontis 6 months, 3 weeks ago

thanks for explanation

upvoted 2 times

takaimomoGcup Most Recent 5 days, 19 hours ago

export
move
run
upvoted 1 times

evangelist 3 months, 3 weeks ago

Export the LUIS application as a package file from the Azure portal: This involves downloading the LUIS model you've developed and want to run locally in a container.

Move the exported package file to the Docker input directory on the host computer: This step involves transferring the downloaded LUIS

application package to a specific directory that the Docker container will use as its input source.

Run the Docker container and specify the input directory: This involves using Docker commands to start the container with the necessary parameters, including the location of the LUIS application package file in the input directory.

Retraining the model is not mentioned as a step for deploying the container image to a host computer because the model should already be trained and exported from the LUIS portal before deployment.

upvoted 1 times

 **Tin_Tin** 11 months, 1 week ago

The answer seems correct.

<https://learn.microsoft.com/en-us/azure/cognitive-services/luis/luis-container-howto?tabs=v3>

upvoted 2 times

 **973b658** 11 months, 3 weeks ago

No.

Export,Docker,Output directory.

upvoted 1 times

 **ziggy1117** 11 months, 3 weeks ago

answer is correct

upvoted 3 times

Question #45

HOTSPOT

You are building a text-to-speech app that will use a custom neural voice.

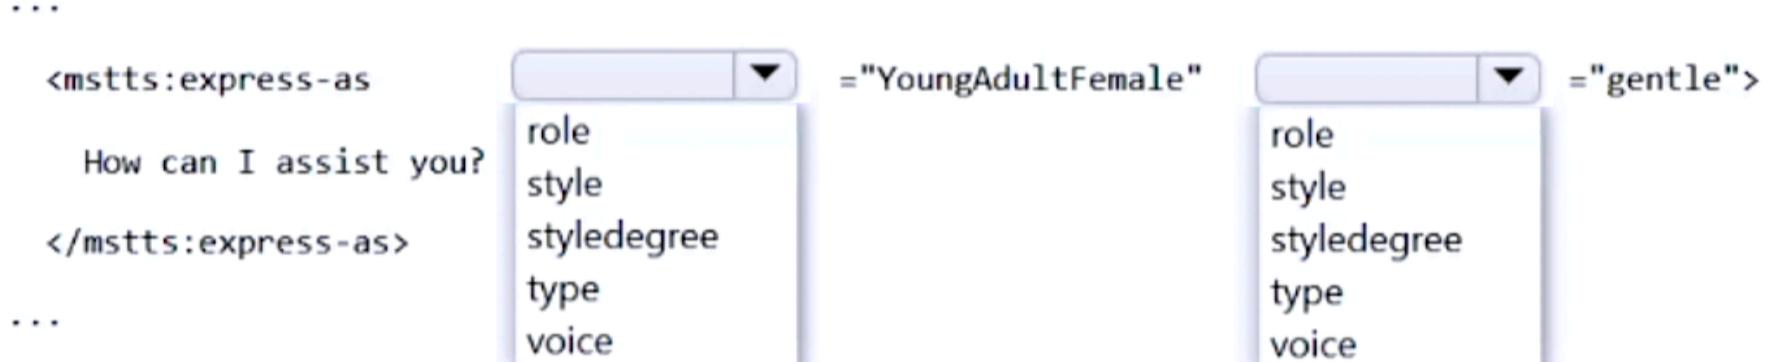
You need to create an SSML file for the app. The solution must ensure that the voice profile meets the following requirements:

- Expresses a calm tone
- Imitates the voice of a young adult female

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area



Correct Answer:

Answer Area

```
<mstts:express-as role="YoungAdultFemale" style="gentle">
    How can I assist you?
</mstts:express-as>
```

zellck 11 months ago

1. role
2. style

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/speech-synthesis-markup-voice#speaking-styles-and-roles>
By default, neural voices have a neutral speaking style. You can adjust the speaking style, style degree, and role at the sentence level.

The following table has descriptions of each supported style attribute.

- style="gentle"

Expresses a mild, polite, and pleasant tone, with lower pitch and vocal energy.

The following table has descriptions of each supported role attribute.

- role="YoungAdultFemale"

The voice imitates a young adult female.

upvoted 12 times

WinzigWeich 1 year ago

correct <https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/speech-synthesis-markup-voice>

upvoted 9 times

funny_penguin 3 days, 23 hours ago

on exam, role and style.

upvoted 1 times

✉  **takaimomoGcup** 5 days, 19 hours ago

role and style

upvoted 1 times

✉  **rdemontis** 6 months, 3 weeks ago

correct answer

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/speech-synthesis-markup-voice>

upvoted 2 times

✉  **james2033** 9 months, 1 week ago

role

style

see role at <https://learn.microsoft.com/en-us/azure/ai-services/speech-service/speech-synthesis-markup-voice#role-example> . May styles at <https://learn.microsoft.com/en-us/azure/ai-services/speech-service/speech-synthesis-markup-voice#speaking-styles-and-roles>

upvoted 2 times

Question #46

HOTSPOT

You have a collection of press releases stored as PDF files.

You need to extract text from the files and perform sentiment analysis.

Which service should you use for each task? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Extract text:

Azure Cognitive Search
Computer Vision
Form Recognizer

Perform sentiment analysis:

Azure Cognitive Search
Computer Vision
Form Recognizer
Language

Answer Area

Extract text:

Azure Cognitive Search
Computer Vision
Form Recognizer

Correct Answer:

Perform sentiment analysis:

Azure Cognitive Search
Computer Vision
Form Recognizer
Language

1. Computer Vision
2. Language

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/overview-ocr>

OCR or Optical Character Recognition is also referred to as text recognition or text extraction. Machine-learning based OCR techniques allow you to extract printed or handwritten text from images, such as posters, street signs and product labels, as well as from documents like articles, report forms, and invoices. The text is typically extracted as words, text lines, and paragraphs or text blocks, enabling access to digital version of the scanned text. This eliminates or significantly reduces the need for manual data entry.

upvoted 14 times

✉  **zellick** 11 months ago

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/sentiment-opinion-mining/overview>

Sentiment analysis and opinion mining are features offered by Azure Cognitive Service for Language, a collection of machine learning and AI algorithms in the cloud for developing intelligent applications that involve written language. These features help you find out what people think of your brand or topic by mining text for clues about positive or negative sentiment, and can associate them with specific aspects of the text.

upvoted 5 times

✉  **rdemontis** 6 months, 3 weeks ago

agree with you

upvoted 3 times

✉  **M25** 8 months, 4 weeks ago

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/overview-ocr#next-steps>

At the very bottom of the same link, under "Next steps":

- OCR for PDF, Office and HTML documents and document images: start with Document Intelligence Read.

upvoted 2 times

✉  **Pffffff**  1 year, 1 month ago

The service you should use to extract text from the PDF files is B. Computer Vision.

Computer Vision has the ability to extract text from images and PDF files, making it a suitable choice for this scenario. Once the text has been extracted, you can then use a text analytics service, such as the Azure Cognitive Services Text Analytics API, to perform sentiment analysis on the extracted text.

Azure Cognitive Search is a search-as-a-service solution that allows you to index and search structured and unstructured data. It can also extract text from PDF files, but it may not provide the level of accuracy required for sentiment analysis.

Form Recognizer is a service that is designed to extract structured data from forms, such as receipts, invoices, and business cards. It may not be the best choice for extracting text from press releases.

upvoted 6 times

✉  **973b658** 11 months, 3 weeks ago

But, Computer Vision < Form Recognizer in this case.

upvoted 1 times

✉  **AnonymousJhb** 7 months, 3 weeks ago

we are reading pdfs and not receipts so we use Computer Vision

upvoted 1 times

✉  **takaimomoGcup**  3 days, 19 hours ago

Computer Vision

Language

upvoted 1 times

✉  **Mehe323** 2 months, 2 weeks ago

Answers are correct. Form recognizer is the old name of Document Intelligence. About Document Intelligence on Microsoft Learn:

"Document Intelligence Read Optical Character Recognition (OCR) model runs at a higher resolution than Azure AI Vision Read and extracts print and handwritten text from PDF documents and scanned images."

<https://learn.microsoft.com/en-us/azure/ai-services/document-intelligence/concept-read?view=doc-intel-4.0.0>

upvoted 4 times

✉  **davidorti** 4 months, 1 week ago

Answer seems correct.

1. Form Recognizer (now Document Intelligence)
2. Language

In CV Read API docu for OCR it says clearly:

OCR for Images: "Optimized for general, *non-document images* with a performance-enhanced synchronous API"

Document Intelligence: "Optimized for text-heavy scanned and digital documents with an asynchronous API to help automate intelligent document processing at scale"

Here we're dealing with a collection of PDFs.

upvoted 1 times

✉  **davidorti** 4 months, 1 week ago

Link to the docu: <https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/how-to/call-read-api>
upvoted 1 times

✉️ **sismér** 6 months ago

The answer is correct:

For extracting text from PDF files, you can use Azure Cognitive Services specifically the Azure Form Recognizer service. Azure Form Recognizer is designed to extract key-value pairs, tables, and text from documents, including PDFs. It supports various document types, making it suitable for extracting text from press releases in PDF format.

For sentiment analysis, you can use the Azure Text Analytics service. Azure Text Analytics includes a sentiment analysis feature that can analyze the sentiment of text documents and provide a sentiment score. This service can help you determine whether the sentiment expressed in the press releases is positive, negative, or neutral.

upvoted 1 times

✉️ **shahnawazkhot** 8 months, 1 week ago

Answer is -

Azure Cognitive Search
Language

Yes, Azure Cognitive Search can be used to extract text from PDF files 1. The Azure Cognitive Search blob indexer can extract text from PDF and other document formats 2. However, extracting text from embedded images or tables is not yet integrated in Azure Search, but it is on the roadmap 3.

To extract text from PDF files using Azure Cognitive Search, you can use the Document Extraction cognitive skill 1. This skill extracts content from a file within the enrichment pipeline and can extract text and images with high accuracy 1. You can use this skill to extract text from PDF files and perform sentiment analysis on the extracted text using the Sentiment Analysis feature provided by Azure Cognitive Services.

upvoted 1 times

✉️ **AnonymousJhb** 7 months, 3 weeks ago

no its not. Azure Computer Vision API offers OCR scanning, which makes it much easier to extract text from PDF files.

Sentiment analysis is done by Language Services.

upvoted 2 times

✉️ **ziggy1117** 11 months, 3 weeks ago

answer is correct:

1. form recognizer -> READ.

Form Recognizer v3.0's Read Optical Character Recognition (OCR) model runs at a higher resolution than Computer Vision Read and extracts printed and handwritten text from PDF documents and scanned images. It also includes preview support for extracting text from Microsoft Word, Excel, PowerPoint, and HTML documents. It detects paragraphs, text lines, words, locations, and languages. The Read model is the underlying OCR engine for other Form Recognizer prebuilt models like Layout, General Document, Invoice, Receipt, Identity (ID) document, in addition to custom models.
<https://learn.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/overview?view=form-recog-3.0.0>
<https://learn.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/concept-read?view=form-recog-3.0.0>

2. Language obviously

upvoted 4 times

✉️ **kail85** 12 months ago

Azure Cognitive Search can be used to extract text from PDF files. It can ingest and index the content of various file formats, including PDFs, by using built-in document cracking capabilities or custom skills. The indexing process extracts text and metadata from the files, making the content searchable.

upvoted 1 times

✉️ **MaliSanFuu** 1 year ago

I think the answer is correct as the FormRecognizer supports the ability for document analysis. There you can easily use the read API to extract printed or handwritten text from images and documents.

upvoted 2 times

✉️ **MaliSanFuu** 1 year ago

easily use*

upvoted 1 times

Question #47

You have a text-based chatbot.

You need to enable content moderation by using the Text Moderation API of Content Moderator.

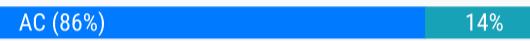
Which two service responses should you use? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. personal data
- B. the adult classification score
- C. text classification
- D. optical character recognition (OCR)
- E. the racy classification score

Correct Answer: AC

Community vote distribution



✉️ **zellick** 11 months ago

Selected Answer: AC

AC is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/content-moderator/text-moderation-api>

Use Content Moderator's text moderation models to analyze text content, such as chat rooms, discussion boards, chatbots, e-commerce catalogs, and documents.

The service response includes the following information:

- Profanity: term-based matching with built-in list of profane terms in various languages
 - Classification: machine-assisted classification into three categories
 - Personal data
 - Auto-corrected text
 - Original text
 - Language
- upvoted 9 times

✉️ **zellick** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 3 times

✉️ **M25** 8 months, 4 weeks ago

Correct! While A, C belong to Text Moderation

<https://learn.microsoft.com/en-us/azure/ai-services/content-moderator/text-moderation-api>,

B, E belong to Vision AI

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/concept-detecting-adult-content#content-flag-definitions>

upvoted 3 times

✉️ **DDD6** 1 year ago

Answer is correct.

The reference URL:

<https://learn.microsoft.com/en-us/azure/cognitive-services/content-moderator/api-reference>

upvoted 5 times

✉️ **Jimmy1017** 1 month ago

C. Text classification

A. Personal data

Here's why these two are important:

Text classification: This functionality allows the API to analyze the text for potentially harmful content like hate speech, bullying, threats, etc. This is crucial for ensuring a safe and positive environment in your chatbot.

Personal data: This helps identify and potentially mask sensitive information like names, addresses, phone numbers, etc., which users might accidentally or intentionally reveal during conversations. This protects user privacy.

Let's break down the other options:

B. The adult classification score: This functionality is not available in the Text Moderation API. It's likely part of the Content Moderator's Image Moderation API for identifying inappropriate visuals.

D. Optical character recognition (OCR): This is not relevant for text-based chatbots as OCR deals with converting images containing text into machine-readable format.

E. The racy classification score: Similar to adult classification score, this functionality is likely intended for image moderation and not directly applicable to text analysis.

upvoted 1 times

✉ **evangelist** 3 months, 3 weeks ago

Selected Answer: AC

The appropriate service responses for content moderation using the Text Moderation API of Content Moderator are:

A. Personal Data and C. Text Classification.

These features help identify sensitive information and categorize text content based on its potential appropriateness, including detecting profanity, personal data, and classifying text into categories related to potentially undesired content

PLEASE DO NOT SELECT A,C!!!

upvoted 1 times

✉ **rdemontis** 6 months, 3 weeks ago

Selected Answer: AC

Based on the official documentation correct answer seems to be AC

<https://learn.microsoft.com/en-us/azure/cognitive-services/content-moderator/text-moderation-api>

upvoted 1 times

✉ **ziggy1117** 11 months, 3 weeks ago

Selected Answer: AC

answer is correct:

<https://westus.dev.cognitive.microsoft.com/docs/services/57cf753a3f9b070c105bd2c1/operations/57cf753a3f9b070868a1f66f>

request parameters: PII and classify

upvoted 3 times

✉ **ziggy1117** 11 months, 3 weeks ago

answer is correct:

<https://westus.dev.cognitive.microsoft.com/docs/services/57cf753a3f9b070c105bd2c1/operations/57cf753a3f9b070868a1f66f>

request parameters: PII and classify

upvoted 2 times

✉ **EliteAllen** 12 months ago

Selected Answer: AC

Based on the information from the official Azure Cognitive Services documentation, it seems that the Text Moderation API indeed returns profanity terms and personal data (A), which can be used for content moderation. It also performs text classification (C), which can be used to categorize and filter content.

So, the correct answers according to the official Azure documentation are A. personal data and C. text classification.

upvoted 4 times

✉ **examworld** 12 months ago

Image Moderation API

Scan images and detect potential adult and racy content by using tags, confidence scores, and other extracted information.

Text Moderation API

Scan text content. Profanity terms and personal data are returned.

upvoted 1 times

✉ **EliteAllen** 1 year ago

Selected Answer: BE

B. the adult classification score

E. the racy classification score

To enable content moderation in a text-based chatbot using the Text Moderation API of Content Moderator, you should use the adult classification score (B) and the racy classification score (E). These scores will help you determine if the content is adult or racy in nature, enabling you to take appropriate action for moderation purposes.

upvoted 3 times

Question #48

HOTSPOT

You are developing a text processing solution.

You have the function shown below.

```
static void GetKeyWords(TextAnalyticsClient textAnalyticsClient, string text)
{
    var response = textAnalyticsClient.RecognizeEntities(text);
    Console.WriteLine("Key words:");

    foreach (CategorizedEntity entity in response.Value)
    {
        Console.WriteLine($"\\t{entity.Text}");
    }
}
```

For the second argument, you call the function and specify the following string.

Our tour of Paris included a visit to the Eiffel Tower

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

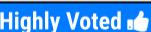
Answer Area

Statements	Yes	No
The output will include the following words: our and included.	<input type="radio"/>	<input type="radio"/>
The output will include the following words: Paris, Eiffel, and Tower.	<input type="radio"/>	<input type="radio"/>
The function will output all the key phrases from the input string to the console.	<input type="radio"/>	<input type="radio"/>

Answer Area

Correct Answer:

Statements	Yes	No
The output will include the following words: our and included.	<input type="radio"/>	<input checked="" type="radio"/>
The output will include the following words: Paris, Eiffel, and Tower.	<input checked="" type="radio"/>	<input type="radio"/>
The function will output all the key phrases from the input string to the console.	<input checked="" type="radio"/>	<input type="radio"/>

  StrateqEBS  10 months, 2 weeks ago

Should be NYN:

<https://learn.microsoft.com/en-us/dotnet/api/azure.ai.textanalytics.textanalyticsclient.recognizeentities?view=azure-dotnet>

Definition:

Runs a predictive model to identify a collection of named entities in the passed-in document, and categorize those entities into types such as person, location, or organization.

This method does not extract phrases.

upvoted 20 times

  rdemontis 6 months, 3 weeks ago

Agree with you. Particularly for the last point we are using the RecognizeEntities method that is used for NER purposes. And the we loop in to the list of entities.

https://github.com/Azure/azure-sdk-for-net/blob/main/sdk/textanalytics/Azure.AI.TextAnalytics/samples/Sample4_RecognizeEntities.md

For Key-Phrase extraction there is another method "ExtractKeyPhrases"

https://github.com/Azure/azure-sdk-for-net/blob/main/sdk/textanalytics/Azure.AI.TextAnalytics/samples/Sample3_ExtractKeyPhrases.md

for key-phrases

upvoted 6 times

✉ **ExamDev** 8 months, 2 weeks ago

The last one is clear "Will output all key phrases on the console" and we have on the example `Console.WriteLine($"\\t{entity.text}")` - With (\t) tabs

upvoted 1 times

✉ **M25** 8 months, 4 weeks ago

Correct! Examples:

Output NER: trip, Seattle, last week

<https://learn.microsoft.com/en-us/azure/ai-services/language-service/named-entity-recognition/quickstart?tabs=ga-api&pivots=programming-language-csharp#output>

Output Key phrase extraction: modern medical office, Dr. Smith, great staff

<https://learn.microsoft.com/en-us/azure/ai-services/language-service/key-phrase-extraction/quickstart?pivots=programming-language-csharp#output>

upvoted 3 times

✉ **zellck**  11 months ago

NYY is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/named-entity-recognition/overview>

Named Entity Recognition (NER) is one of the features offered by Azure Cognitive Service for Language, a collection of machine learning and AI algorithms in the cloud for developing intelligent applications that involve written language. The NER feature can identify and categorize entities in unstructured text. For example: people, places, organizations, and quantities.

upvoted 9 times

✉ **Mehe323** 2 months, 2 weeks ago

But entities are not the same as key phrases. A key phrase here could be 'Tour of Paris' which doesn't coincide with the entities. So I think the last one should be N.

upvoted 3 times

✉ **ProfessorZ** 1 month, 2 weeks ago

<https://learn.microsoft.com/en-us/azure/ai-services/language-service/named-entity-recognition/quickstart?tabs=ga-api&pivots=programming-language-csharp>

Example output shows "last week" as text output:

Named Entities:

Text: trip, Category: Event, Sub-Category:

Score: 0.74, Length: 4, Offset: 18

Text: Seattle, Category: Location, Sub-Category: GPE

Score: 1.00, Length: 7, Offset: 26

Text: **last week**, Category: DateTime, Sub-Category: DateRange

upvoted 1 times

✉ **funny_penguin**  3 days, 23 hours ago

on exam, NYN. For the last one I selected N because the method looks like it's extracting entities, not key phrases.

upvoted 1 times

✉ **takaimomoGcup** 5 days, 19 hours ago

NYN is right answer.

upvoted 1 times

✉ **Murtuza** 1 month, 4 weeks ago

The function will output all the key phrases from the input string to the console. No, the function will output the recognized entities, not all key phrases. Key phrases could include other important words or phrases in the text that are not necessarily entities. For key phrase extraction, a different method would be used.

upvoted 2 times

✉ **Murtuza** 2 months, 1 week ago

For the last choice that seems to be a topic of discussion

Console Output:

The code prints the header "Key words:" to the console.

It then iterates through the response.Value (presumably a collection of categorized entities)

upvoted 2 times

✉ **lastget** 4 months, 2 weeks ago

Agree NYN

upvoted 2 times

✉ **katrang** 7 months, 1 week ago

I think NYY. I could not see where all keywords would be output, but after checking the documentation they would also be identified as entities along with the Eiffel Tower (tour, paris and visit)

upvoted 1 times

 **973b658** 11 months, 2 weeks ago

It is true.

upvoted 1 times

Question #49

HOTSPOT

You are building an Azure web app named App1 that will translate text from English to Spanish.

You need to use the Text Translation REST API to perform the translation. The solution must ensure that you have data sovereignty in the United States.

How should you complete the URI? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area



Correct Answer:

Answer Area

https://api-nam.cognitive.microsofttranslator.com / ?api-version=3.0&to=es

api-nam.cognitive.microsofttranslator.com
api-nam.cognitiveservices.azure.com
eastus.api.cognitive.microsoft.com

detect
languages
text-to-speech
translate

zellck [Highly Voted] 11 months ago

1. api-nam.cognitive.microsofttranslator.com
2. translate

<https://learn.microsoft.com/en-us/azure/cognitive-services/Translator/reference/v3-0-reference#base-urls>

Requests to Translator are, in most cases, handled by the datacenter that is closest to where the request originated. If there's a datacenter failure when using the global endpoint, the request may be routed outside of the geography.

To force the request to be handled within a specific geography, use the desired geographical endpoint. All requests are processed among the datacenters within the geography.

- United States
- api-nam.cognitive.microsofttranslator.com

<https://learn.microsoft.com/en-us/azure/cognitive-services/translator/reference/rest-api-guide>

- translate

Translate specified source language text into the target language text.

upvoted 17 times

mpit 6 months, 3 weeks ago

Thank you for your valuable contribution.
upvoted 1 times

rdemontis 6 months, 3 weeks ago

thanks for your contribution
upvoted 1 times

takaimomoGcup [Most Recent] 5 days, 19 hours ago

api-nam.cognitive.microsofttranslator.com
translate
upvoted 2 times

abelarda 2 weeks ago

Why not api-nam.cognitiveservice.azure.com/translate?
upvoted 1 times

evangelist 3 months, 3 weeks ago

nam=North America
upvoted 3 times

 **sca88** 6 months ago

The response is correct

1. api-nam.cognitive.microsofttranslator.com

2. translate

upvoted 1 times

 **973b658** 11 months, 2 weeks ago

It is true.

upvoted 2 times

Question #50

DRAG DROP

You have a Docker host named Host1 that contains a container base image.

You have an Azure subscription that contains a custom speech-to-text model named model1.

You need to run model1 on Host1.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Retrain the model.	
Request approval to run the container.	
Export model1 to Host1.	
Run the container.	
Configure disk logging.	

Correct Answer:

Answer Area
Request approval to run the container.
Retrain the model.
Run the container.

 **zellck**  11 months ago

1. Request approval to run container
2. Export model1 to Host1
3. Run the container

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/speech-container-stt?tabs=container&pivots=programming-language-csharp>

upvoted 18 times

 **Mehe323** 2 months, 2 weeks ago

Okay, this link provides more information about the approval step, but approval is only necessary when you run the container in a disconnected environment. So if that is mentioned in a question, choosing this step will make more sense.

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/speech-container-overview#request-approval-to-run-the-container>

upvoted 2 times

 **973b658**  11 months, 2 weeks ago

- No.
1. Export model1 to Host1.
 2. Request approval to run the container.
 3. Run the container.

upvoted 7 times

 **takaimomoGcup**  5 days, 19 hours ago

Request
Export
Run

upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

Probably here we are requested to run the container in a disconnected environment so I think correct answer is :

Request approval to run container
Export model1 to Host1
Run the container

<https://learn.microsoft.com/en-us/azure/ai-services/containers/disconnected-containers>

upvoted 2 times

 **sl_msiconsulting** 7 months, 2 weeks ago

You only need to request approval if you plan to run the container in a completely disconnected environment. And you might not get approved at all as there are some requirements. The question does not indicate if we need to run the container in such an environment.

upvoted 3 times

✉️ **sl_msiconsulting** 7 months, 2 weeks ago

Based on what you can choose here I would say it's a disconnected environment. For a connected environment using docker run allows you to download the model at the same so there is no need to export the model manually and then copy it to the host. Zellck is right in the choices.

upvoted 2 times

✉️ **M25** 8 months, 4 weeks ago

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/speech-container-cstt?tabs=custom-model%2Ccontainer&pivots=programming-language-csharp>

In this article, you'll learn how to download [from Microsoft Container Registry (MCR)], install, and run a Custom speech to text container. Get the model ID (to use as the argument to the ModelId parameter of the docker run command): The custom model has to have been trained by using the Speech Studio.

N/a so far, excluding A (Retrain the model).

upvoted 1 times

✉️ **M25** 8 months, 4 weeks ago

Run the container with docker run:

B (Request approval precedes),

The docker run command will start the container when all three of the following options are provided with valid values: ApiKey, Billing, EULA
<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/speech-container-howto#billing-arguments>

D (Run container), C (Get model)

Here's an example docker run command with placeholder values. ... This command:

- Runs a custom speech to text container from the container image.
- Allocates 4 CPU cores and 8 GB of memory.
- Loads the custom speech to text model from the volume input mount, for example, C:\CustomSpeech.
- Exposes TCP port 5000 and allocates a pseudo-TTY for the container.
- Downloads the model given the ModelId (if not found on the volume mount).

upvoted 2 times

✉️ **Tin_Tin** 11 months, 1 week ago

Not sure.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/speech-container-overview#request-approval-to-run-the-container>

upvoted 1 times

Question #51

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You build a language model by using a Conversational Language Understanding. The language model is used to search for information on a contact list by using an intent named FindContact.

A conversational expert provides you with the following list of phrases to use for training.

- Find contacts in London.
- Who do I know in Seattle?
- Search for contacts in Ukraine.

You need to implement the phrase list in Conversational Language Understanding.

Solution: You create a new utterance for each phrase in the FindContact intent.

Does this meet the goal?

- A. Yes
- B. No

Correct Answer: B

Community vote distribution

A (56%) B (44%)

 evangelist Highly Voted 3 months, 3 weeks ago

Selected Answer: A

A. Yes

Creating a new utterance for each phrase in the FindContact intent is a correct approach to implement the phrase list in Conversational Language Understanding. This method trains the language model to recognize variations of how users might express the intent to find contacts in different locations, thereby improving the model's accuracy in identifying the FindContact intent.

upvoted 6 times

 evangelist 3 months, 3 weeks ago

Yes, creating a new utterance for each phrase in the FindContact intent aligns with the recommended practice for designing applications in Conversational Language Understanding, as detailed in the documentation. This approach helps in accurately capturing the intent by providing diverse examples of how users might express their request, thus enhancing the model's ability to understand and classify user queries correctly. For more detailed guidelines, refer to the section on creating example utterances for each intent in the documentation

upvoted 1 times

 nanaw770 Most Recent 3 days, 18 hours ago

Selected Answer: A

It MUST be A.

upvoted 1 times

 PCRamirez 3 months, 2 weeks ago

According to Windows Copilot: B. No

Creating a new utterance for each phrase in the FindContact intent is not the most efficient approach. Instead, you can use phrase lists in Conversational Language Understanding (LUIS) to group similar phrases together. By defining a phrase list, you can handle variations of the same intent more effectively. In this case, you can create a phrase list containing the cities (London, Seattle, Ukraine) and use it within the FindContact intent. This way, LUIS will recognize any variation of these cities as part of the same intent without creating individual utterances for each location.



upvoted 1 times

 Tactable 3 months, 3 weeks ago

Selected Answer: B

According to ChatGPT:

B. No

Explanation: While creating a new utterance for each phrase in the FindContact intent is a step in the right direction, it may not be sufficient to fully meet the goal. To effectively implement the phrase list in Conversational Language Understanding, it's essential to consider variations in how users might express the same intent. The provided phrases cover different scenarios (finding contacts in different locations), but there may be additional variations and nuances to consider. Therefore, merely creating a new utterance for each provided phrase might not capture all possible ways users could express the intent to find contacts. A more comprehensive approach to training the language model might involve incorporating synonyms, alternative phrasings, and potential variations that users might use when searching for contacts.

upvoted 1 times

 **dimsok** 4 months, 3 weeks ago

Selected Answer: A

It's "Yes" actually, not the more efficient way to do it, but it will work

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

B. No

Creating a new utterance for each phrase in the FindContact intent is not the most efficient approach for implementing the provided phrase list. Instead, you should use phrase list features or entities to capture variations of these phrases more effectively.

In Conversational Language Understanding, you can define a phrase list or entity that includes variations of location names like "London," "Seattle" and "Ukraine." By doing this, you allow the model to recognize these location names as entities, making your intent more flexible and capable of handling variations. This approach is much more scalable and less labor-intensive than creating individual utterances for each location.

The goal should be met by using phrase lists or entities effectively to capture variations in the input data and improve the model's performance. (ChatGPT)

upvoted 1 times

 **dimsok** 4 months, 3 weeks ago

It won't be efficient but it will do the job, I would vote for "Yes"

upvoted 2 times

 **sl_mslconsulting** 7 months, 2 weeks ago

Selected Answer: B

I picked B because you need the entity to retrieve the location from the utterances for your app to be able to know which contacts to retrieve from the store.

upvoted 1 times

 **sl_mslconsulting** 7 months, 2 weeks ago

The utterances provided are different enough. But if you're picky and insist that you should have at least 15 utterances that you would pick B anyway.

upvoted 1 times

 **zellck** 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/luis/concepts/application-design#create-example-utterances-for-each-intent>
To start, avoid creating too many utterances for each intent. Once you have determined the intents you need for your app, create 15 to 30 example utterances per intent. Each utterance should be different from the previously provided utterances. Include a variety of word counts, word choices, verb tenses, and punctuation.

upvoted 3 times

 **973b658** 11 months, 2 weeks ago

Selected Answer: B

B. Same question.

upvoted 1 times

Question #52

DRAG DROP

You have a question answering project in Azure Cognitive Service for Language.

You need to move the project to a Language service instance in a different Azure region.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
From the new Language service instance, train and publish the project.	
From the new Language service instance, import the project file.	
From the new Language service instance, enable custom text classification.	
From the original Language service instance, export the existing project.	
From the new Language service instance, regenerate the keys.	
From the original Language service instance, train and publish the model.	



Correct Answer:	Answer Area
	From the original Language service instance, export the existing project.
	From the new Language service instance, import the project file.
	From the new Language service instance, train and publish the project.

 **zellck**  11 months ago

1. From original instance, export existing project.
2. From new instance, import the project file.
3. From new instance, train and publish model.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/question-answering/how-to/migrate-knowledge-base>
upvoted 22 times

 **takaimomoGcup**  5 days, 19 hours ago

original, export existing project.
new, import the project file.
new, train and publish model.
upvoted 1 times

 **evangelist** 3 months, 3 weeks ago

Answer is correct
first from the source instance, export the source project;
at the new instance, import the exported source project;
to deploy, one has to train and publish the model from imported project, the training is needed before the model can work in new project from different Azure region
upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

the answer seems correct:

<https://learn.microsoft.com/en-us/azure/ai-services/language-service/question-answering/how-to/migrate-knowledge-base>
<https://learn.microsoft.com/en-us/azure/ai-services/qnamaker/quickstarts/create-publish-knowledge-base#publish-the-knowledge-base>
upvoted 3 times

 **973b658** 11 months, 2 weeks ago

It is true.
upvoted 2 times

Question #53

DRAG DROP

You are building a customer support chatbot.

You need to configure the bot to identify the following:

- Code names for internal product development
- Messages that include credit card numbers

The solution must minimize development effort.

Which Azure Cognitive Service for Language feature should you use for each requirement? To answer, drag the appropriate features to the correct requirements. Each feature may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Features	Answer Area
Custom named entity recognition (NER)	Identify code names for internal product development: <input type="text"/>
Key phrase extraction	Identify messages that include credit card numbers: <input type="text"/>
Language detection	
Named Entity Recognition (NER)	
Personally Identifiable Information (PII) detection	
Sentiment analysis	

Correct Answer:

Answer Area
Identify code names for internal product development: Custom named entity recognition (NER)
Identify messages that include credit card numbers: Personally Identifiable Information (PII) detection

zellck 11 months ago

1. Custom NER
2. PII detection

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/custom-named-entity-recognition/overview>
Custom NER enables users to build custom AI models to extract domain-specific entities from unstructured text, such as contracts or financial documents. By creating a Custom NER project, developers can iteratively label data, train, evaluate, and improve model performance before making it available for consumption. The quality of the labeled data greatly impacts model performance.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/personally-identifiable-information/overview>
PII detection is one of the features offered by Azure Cognitive Service for Language, a collection of machine learning and AI algorithms in the cloud for developing intelligent applications that involve written language. The PII detection feature can identify, categorize, and redact sensitive information in unstructured text. For example: phone numbers, email addresses, and forms of identification.

upvoted 14 times

rdemontis 6 months, 3 weeks ago
thank you for your great contribution
upvoted 3 times

takaimomoGcup 5 days, 19 hours ago
NER and PII.
upvoted 1 times

Tin_Tin 11 months, 2 weeks ago
seems correct.
<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/overview>
upvoted 2 times

973b658 11 months, 2 weeks ago
It is true.

upvoted 2 times

Question #54

Topic 3

HOTSPOT

You are building an app by using the Speech SDK. The app will translate speech from French to German by using natural language processing.

You need to define the source language and the output language.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
var speechTranslationConfig =  
    SpeechTranslationConfig.FromSubscription(speechKey, speechRegion);  
  
speechTranslationConfig.  
    AddTargetLanguage("fr")  
    .SpeechRecognitionLanguage  
    .SpeechSynthesisLanguage  
    .TargetLanguages  
    .VoiceName  
  
speech_translation_config.  
    AddTargetLanguage("de")  
    .SpeechRecognitionLanguage  
    .SpeechSynthesisLanguage  
    .TargetLanguages  
    .VoiceName
```

Answer Area

```
var speechTranslationConfig =  
    SpeechTranslationConfig.FromSubscription(speechKey, speechRegion);  
  
speechTranslationConfig.  
    AddTargetLanguage("fr")  
    .  
    .SpeechRecognitionLanguage  
    .SpeechSynthesisLanguage  
    .TargetLanguages  
    .VoiceName  
  
speech_translation_config.  
    AddTargetLanguage("de")  
    .  
    .SpeechRecognitionLanguage  
    .SpeechSynthesisLanguage  
    .TargetLanguages  
    .VoiceName
```

Correct Answer:

 **f2c587e** 1 month, 4 weeks ago

Las respuestas son B y C

upvoted 2 times

 **Florin83** 3 months, 3 weeks ago

correct

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

correct

upvoted 4 times

 **jangotango** 7 months, 2 weeks ago

Answer is correct

upvoted 4 times

Question #55

DRAG DROP

You have a collection of Microsoft Word documents and PowerPoint presentations in German.

You need to create a solution to translate the files to French. The solution must meet the following requirements:

- Preserve the original formatting of the files.
- Support the use of a custom glossary.

You create a blob container for German files and a blob container for French files. You upload the original files to the container for German files.

Which three actions should you perform in sequence to complete the solution? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Perform an asynchronous translation by using the list of files to be translated.	
Perform an asynchronous translation by using the document translation specification.	
Generate a list of files to be translated.	>
Upload a glossary file to the container for German files.	<
Upload a glossary file to the container for French files.	<
Define a document translation specification that has a French target.	<

Correct Answer:

Answer Area
Upload a glossary file to the container for German files.
Define a document translation specification that has a French target.
Perform an asynchronous translation by using the document translation specification.

 **rdemontis**  6 months, 3 weeks ago

IMHO the answer is:

1. Upload a Glossary file to the french files container
2. Define a document translation specification that has french target
3. Perform asynchronous translation by using the document specification

As you can see below the glossary is needed before the translation:

<https://learn.microsoft.com/en-us/azure/ai-services/translator/document-translation/how-to-guides/create-use-glossaries>

And DocumentTranslationInput is the class we can use for the translation:

<https://learn.microsoft.com/en-us/python/api/azure-ai-translation-document/azure.ai.translation.document.documenttranslationinput?view=azur>

upvoted 10 times

 **vovap0vovap**  5 days, 12 hours ago

Well, reality is that in MS example glossary in own container. But it is also in "targets" section of json file, so answer "French container" probably safer.

upvoted 1 times

 **takaimomoGcup** 5 days, 19 hours ago

1. Upload a Glossary file to the french files container
2. Define a document translation specification that has french target
3. Perform asynchronous translation by using the document specification

upvoted 1 times

 **evangelist** 3 months, 3 weeks ago

The given answer is correct : reason is below

Upload the terminology list: Begin by uploading a custom terminology list, essential for accurate translations. Ensure it's placed in a location

recognized by the translation service.

Define translation specifications: Specify source (German), target (French), and any special parameters in a translation specification.

Execute asynchronous translation: Initiate the translation process based on the specification, which will preserve original formatting and save translated files in the French document container.

These three steps form the core of the solution, although practical implementation may require additional configuration and permissions. Ensure the terminology list is uploaded to a location compatible with the translation service's requirements.

upvoted 3 times

✉ **dimsok** 4 months, 2 weeks ago

You need a French Glossary in the French container - how else can the translator know the French Jargon

upvoted 1 times

✉ **sl_msiconsulting** 7 months, 2 weeks ago

The glossary file can be in this own container and that's tricky as in theory you can put it anywhere. But since we need probably need have different permissions for the source and the target container putting it in the container where the source files would do. Also pay attention to what a glossary file its content looks like and how to specify the location in a request sent to the endpoint. <https://learn.microsoft.com/en-us/azure/ai-services/translator/document-translation/how-to-guides/create-use-glossaries>

upvoted 1 times

✉ **sl_msiconsulting** 7 months, 2 weeks ago

Although I can't find the term translation specification file in the doc, but you don't need to create a list of files that's for sure (pay attention to how the target and source are specified in the link I provided above).

upvoted 1 times

✉ **sl_msiconsulting** 7 months, 2 weeks ago

I changed my mind the glossary file should be in the target folder. Reason for that is because you can have multiple target languages each with its own glossary files. Remember the glossary provides the expected translation for each word in the file to help ensure accuracy. In other words you are providing the translation and should have done so for each target language.

upvoted 3 times

Question #56

You have the following C# function.

```
static void MyFunction(TextAnalyticsClient textAnalyticsClient, string text)
{
    var response = textAnalyticsClient.ExtractKeyPhrases(text);
    Console.WriteLine("Key phrases:");

    foreach (string keyphrase in response.Value)
    {
        Console.WriteLine($"{keyphrase}");
    }
}
```

You call the function by using the following code.

```
MyFunction(textAnalyticsClient, "the quick brown fox jumps over the lazy dog");
```

Which output will you receive?

- A. The quick -
The lazy
- B. the quick brown fox jumps over the lazy dog
- C. jumps over the
- D. quick brown fox
lazy dog

Correct Answer: D

Community vote distribution

D (100%)

✉  **takaimomoGcup** 5 days, 19 hours ago

Selected Answer: D

Key Phrases are "quick brown fox" and "lazy dog".

upvoted 2 times

✉  **evangelist** 3 months, 3 weeks ago

because the method is to extract "key Phrases" so article "the" will not be extracted.

upvoted 1 times

✉  **MelMac** 5 months ago

Selected Answer: D

I tried it out. D is correct. Key Phrases:

quick brown fox

lazy dog

upvoted 4 times

✉  **rdemontis** 6 months, 3 weeks ago

Selected Answer: D

the answer seems correct

upvoted 1 times

✉  **sl_msiconsulting** 7 months, 2 weeks ago

Selected Answer: D

Verified the answer using a simple console program

upvoted 3 times

✉ sl_mslconsulting 7 months, 2 weeks ago

Tried a couple examples and seems it will pick up the pattern: adjective + noun

upvoted 4 times

Question #57

Topic 3

You have the following Python method.

```
def create_resource(resource_name, kind, account_tier, location) :  
    parameters = CognitiveServicesAccount(sku=Sku(name=account_tier), kind=kind, location=location, properties={})  
    result = cogSvcClient.accounts.create(resource_group_name, resource_name, parameters)
```

You need to deploy an Azure resource to the East US Azure region. The resource will be used to perform sentiment analysis.

How should you call the method?

- A. create_resource("res1", "TextAnalytics", "Standard", "East US")
- B. create_resource("res1", "ContentModerator", "S0", "eastus")
- C. create_resource("res1", "ContentModerator", "Standard", "East US")
- D. create_resource("res1", "TextAnalytics", "S0", "eastus")

Correct Answer: D

Community vote distribution

D (100%)

✉ takaimomoGcup 5 days, 19 hours ago

Selected Answer: D

D is right answer.

upvoted 1 times

✉ evangelist 3 months, 3 weeks ago

In Azure services, "S0" represents a pricing tier or service level. For Azure Cognitive Services, including Text Analytics, the pricing tier determines the service's performance, request rate limits, features, and costs. Each pricing tier offers different service capabilities and quotas, and they come with varying price points.

upvoted 2 times

✉ rdemontis 6 months, 3 weeks ago

Selected Answer: D

correct

upvoted 2 times

✉ chenglim 7 months ago

Selected Answer: D

correct answer

upvoted 1 times

Question #58

DRAG DROP

You develop a Python app named App1 that performs speech-to-speech translation.

You need to configure App1 to translate English to German.

How should you complete the SpeechTranslationConfig object? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Values	Answer Area
add_target_language	
speech_synthesis_language	
speech_recognition_language	
voice_name	

```

def translate_speech_to_text():

    translation_config = speechsdk.translation.SpeechTranslationConfig(subscription=speech_key, region=service_region)

    translation_config.  Value = "en-US";
    translation_config.  Value ("de");

```

Correct Answer:

```

def translate_speech_to_text():

    translation_config = speechsdk.translation.SpeechTranslationConfig(subscription=speech_key, region=service_region)

    translation_config.  speech_recognition_language = "en-US";
    translation_config.  speech_synthesis_language ("de");

```

 **jangotango**  7 months, 3 weeks ago

Pretty sure the second one is add_target_language
upvoted 21 times

 **nastolgia** 5 months ago

you do speech-to-speech. so your output should be voice, but proposed answer is incomplete. You also need to specify the type of voice for the synthesizer
upvoted 3 times

 **upliftinghut** 4 weeks ago

Agree, the requirement should changed to translate from speech to text, look at the function name
upvoted 2 times

 **evangelist**  3 months, 3 weeks ago

The answer is WRONG!

Correct Answer is below:

```

def translate_speech_to_text():
    translation_config = speechsdk.translation.SpeechTranslationConfig(subscription=speech_key, region=service_region)
    translation_config.speech_recognition_language = "en-US"
    translation_config.add_target_language("de")

```

upvoted 12 times

 **takaimomoGcup**  5 days, 19 hours ago

It MUST be translation_config.speech_recognition_language = "en-US" and translation_config.add_target_language("de").
What exactly have you studied? Memorization is not enough.
upvoted 1 times

 **Florin83** 3 months, 4 weeks ago

Similar with Question #40 (language is C#, not Python)
I believe it's
1.translation_config.speech_recognition_value = "en-us"
2.translation_config.add_target_language("de")
the second one is a method, not a property, thus the "add"
upvoted 2 times

 **tdctdc** 5 months, 4 weeks ago

How is it even possible that ET provides us with SO MANY wrong answers? I agree with rdemontis and jangotango, the second one is add_target_language

upvoted 3 times

 **idcanymore** 3 months, 2 weeks ago

rdemontis should be getting paid

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

The second is wrong. it should be add_target_language

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/get-started-speech-translation?tabs=windows%2Cterminal&pivots=programming-language-python#translate-speech-from-a-microphone>

upvoted 3 times

 **jangotango** 7 months, 3 weeks ago

Proof - <https://learn.microsoft.com/en-us/azure/ai-services/speech-service/get-started-speech-translation?tabs=windows%2Cterminal&pivots=programming-language-python>

upvoted 3 times

Question #59

HOTSPOT

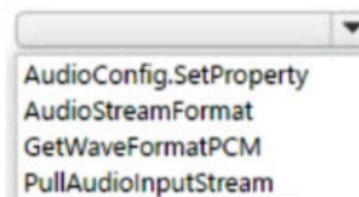
You are developing a streaming Speech to Text solution that will use the Speech SDK and MP3 encoding.

You need to develop a method to convert speech to text for streaming MP3 data.

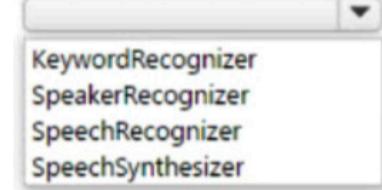
How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

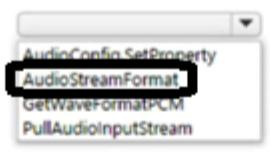
```
audio_format = speechsdk.audio.
(stream_compressed_stream_format=speechsdk.AudioStreamContainerFormat.MP3)

stream = speechsdk.audio.PullAudioInputStream(stream_format=audio_format, pull_stream_callback=callback)
speech_config = speechsdk.SpeechConfig("18c51a87-3a69-47a8-aedc-a54745f708a1", "westus")
audio_config = speechsdk.audio.AudioConfig(stream=stream)

recognizer = speechsdk.
(speech_config=speech_config, audio_config=audio_config)

result = recognizer.recognize_once()
text = result.text
```

Answer Area

```
audio_format = speechsdk.audio.
(stream_compressed_stream_format=speechsdk.AudioStreamContainerFormat.MP3)

stream = speechsdk.audio.PullAudioInputStream(stream_format=audio_format, pull_stream_callback=callback)
speech_config = speechsdk.SpeechConfig("18c51a87-3a69-47a8-aedc-a54745f708a1", "westus")
audio_config = speechsdk.audio.AudioConfig(stream=stream)

recognizer = speechsdk.
(speech_config=speech_config, audio_config=audio_config)

result = recognizer.recognize_once()
text = result.text
```

Correct Answer:

 evangelist  3 months, 3 weeks ago

first one has to configure the format as MP3, so "AudioStreamFormat" option is chosen. Second, since it is speech to text, the SpeechRecognition option is needed.

The answer is CORRECT

upvoted 7 times

 takaimomoGcup  5 days, 19 hours ago

```
audio_format = speechsdk.audio.AudioStreamFormat(compressed_stream_format=speechsdk.AudioStreamContainerFormat.MP3)
recognizer = speechsdk.SpeechRecognizer(speech_config, audio_config)
```

upvoted 1 times

 hkbnjos 4 months, 2 weeks ago

this question is same as Topic 1 Q18 but with different answer, which is correct?

upvoted 1 times

 Mehe323 2 months, 2 weeks ago

The answers are the same. Both chosen are AudioStreamFormat and SpeechRecognizer, but in question 18 there is an additional method for AudioStreamFormat (GetCompressedFormat).

upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

correct

upvoted 2 times

 **SCyrus** 7 months, 3 weeks ago

Answer correct

https://github.com/Azure-Samples/cognitive-services-speech-sdk/blob/master/samples/python/console/speech_sample.py

upvoted 4 times

Question #60

HOTSPOT

You are building a chatbot.

You need to use the Content Moderator API to identify aggressive and sexually explicit language.

Which three settings should you configure? To answer, select the appropriate settings in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Content Moderator - Moderate Text - Screen

The operation detects profanity in more than 100 languages and match against custom and shared blacklists.

Host

Name

Resource Name

Query parameters

autocorrect

[✖ Remove parameter](#)

PII

[✖ Remove parameter](#)

listId

[✖ Remove parameter](#)

classify

[✖ Remove parameter](#)

language

[✖ Remove parameter](#)[✚ Add parameter](#)

Headers

Content-Type

[✖ Remove header](#)

Ocp-Apim-Subscription-Key

Answer Area

Content Moderator - Moderate

Text - Screen

The operation detects profanity in more than 100 languages and match against custom and shared blacklists.

Host

Name

Resource Name

Query parameters

Correct Answer:

autocorrect

 Remove parameter

PII

 Remove parameter

listid

 Remove parameter

classify

 Remove parameter

language

 Remove parameter

 Add parameter

Headers

Content-Type

 Remove header

Ocp-Apim-Subscription-Key

 **audlindr**  2 months, 4 weeks ago

The answers should be

Resource Name

classify

ocp-Apim-Subscription-Key

autocorrect has nothing to do with the question

upvoted 11 times

 **audlindr** 2 months, 4 weeks ago

<https://westus.dev.cognitive.microsoft.com/docs/services/57cf753a3f9b070c105bd2c1/operations/57cf753a3f9b070868a1f66f>

upvoted 2 times

 **Training** 1 month, 3 weeks ago

<https://learn.microsoft.com/en-us/training/modules/classify-and-moderate-text-with-azure-content-moderator/4-exercise-use-the-api-console>

upvoted 2 times

 **takaimomoGcup**  5 days, 19 hours ago

Resource name

classify

Ocp-Api-Suscription-Key

upvoted 1 times

Question #61

Topic 3

You are developing an app that will use the Decision and Language APIs.

You need to provision resources for the app. The solution must ensure that each service is accessed by using a single endpoint and credential.

Which type of resource should you create?

- A. Language
- B. Speech
- C. Azure Cognitive Services
- D. Content Moderator

Correct Answer: C

Community vote distribution

C (100%)

 **Murtuza** 1 month, 3 weeks ago

Selected Answer: C

The correct answer is C. Azure Cognitive Services.

When you create an Azure Cognitive Services resource, you get access to a suite of services and APIs, including the Decision and Language APIs, under a single endpoint and credential. This simplifies the management of these services and enhances security by reducing the number of credentials you need to manage. Other options like Language, Speech, and Content Moderator are individual services within Azure Cognitive Services. They do not provide a single endpoint and credential for accessing multiple services. Therefore, they do not meet the requirement specified in the question

upvoted 1 times

 **Ghill1982** 2 months, 2 weeks ago

Selected Answer: C

A Cognitive Services multi-service resource allows you to access multiple Azure AI services with a single key and endpoint.

upvoted 3 times

Question #62

You are building a chatbot.

You need to ensure that the bot will recognize the names of your company's products and codenames. The solution must minimize development effort.

Which Azure Cognitive Service for Language service should you include in the solution?

- A. custom text classification
- B. entity linking
- C. custom Named Entity Recognition (NER)
- D. key phrase extraction

Correct Answer: C

Community vote distribution

C (100%)

 takaimomoGcup 5 days, 20 hours ago

Selected Answer: C

NER is right answer.

upvoted 1 times

 Murtuza 2 months, 1 week ago

Named Entity Recognition (NER) identifies and categorizes specific names or terms in text, such as names of people, organizations, places, and more. By customizing NER, you can tailor it to recognize your company-specific entities, making it an efficient solution for your chatbot's needs45

So, the correct answer is C. custom Named Entity Recognition (NER).

upvoted 1 times

 GHill1982 2 months, 2 weeks ago

Selected Answer: C

C. custom Named Entity Recognition (NER)

Custom NER allows you to train a model to recognize specific terms relevant to your business, such as product names and codenames.

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/concept-shelf-analysis>

upvoted 3 times

Question #63

Topic 3

You have an Azure subscription that contains an Azure App Service app named App1.

You provision a multi-service Azure Cognitive Services resource named CSAccount1.

You need to configure App1 to access CSAccount1. The solution must minimize administrative effort.

What should you use to configure App1?

- A. a system-assigned managed identity and an X.509 certificate
- B. the endpoint URI and an OAuth token
- C. the endpoint URI and a shared access signature (SAS) token
- D. the endpoint URI and subscription key

Correct Answer: D

✉  **michaelmorar** 1 month, 1 week ago

In general, you always need an endpoint and subscription key.

upvoted 1 times

✉  **Murtuza** 2 months, 1 week ago

By providing the endpoint URI and subscription key in your application, you can seamlessly connect to CSAccount1 without additional complexities or setup. This approach minimizes administrative overhead and ensures secure communication between your app and the cognitive services.

Therefore, the correct answer is D. the endpoint URI and subscription key.

upvoted 2 times

✉  **[Removed]** 2 months, 2 weeks ago

Answer correct, also what a beautiful night?!

upvoted 2 times

Question #64

Topic 3

You have an Azure subscription that contains a multi-service Azure Cognitive Services Translator resource named Translator1.

You are building an app that will translate text and documents by using Translator1.

You need to create the REST API request for the app.

Which headers should you include in the request?

- A. the access control request, the content type, and the content length
- B. the subscription key and the client trace ID
- C. the resource ID and the content language
- D. the subscription key, the subscription region, and the content type

Correct Answer: D

Community vote distribution

D (100%)

 **audlindr**  2 months, 4 weeks ago

Selected Answer: D

Answer is correct

<https://learn.microsoft.com/en-us/azure/ai-services/translator/reference/v3-0-translate>
content-type is required header

<https://learn.microsoft.com/en-us/azure/ai-services/translator/reference/v3-0-reference>

When you use a multi-service secret key, you must include two authentication headers with your request. There are two headers that you need to call the Translator.

Ocp-Apim-Subscription-Key The value is the Azure secret key for your multi-service resource.

Ocp-Apim-Subscription-Region The value is the region of the multi-service resource.

Region is required for the multi-service Text API subscription. The region you select is the only region that you can use for text translation when using the multi-service key. It must be the same region you selected when you signed up for your multi-service subscription through the Azure portal.

D is the only option with both the headers

upvoted 5 times

 **takaimomoGcup**  5 days, 19 hours ago

Selected Answer: D

the subscription key, the subscription region, and the content type

upvoted 1 times

 **kay1101** 2 weeks, 6 days ago

D.

<https://learn.microsoft.com/en-us/azure/ai-services/translator/quickstart-text-rest-api?tabs=csharp#headers>

upvoted 1 times

 **Murtuza** 2 months, 1 week ago

To create the REST API request for your app that will translate text and documents using the Azure Cognitive Services Translator resource (Translator1), you should include the following headers:

Subscription Key: This key serves as the authentication token to access the service.

Subscription Region: The region where your Translator resource is deployed.

Content Type: Specify the type of content you are sending (e.g., JSON, plain text).

Therefore, the correct answer is D. the subscription key, the subscription region, and the content type. When making the API request, ensure that you include these headers to connect your application to the Translator service.

upvoted 3 times

Question #65

You have a file share that contains 5,000 images of scanned invoices.

You need to analyze the images. The solution must extract the following data:

- Invoice items
- Sales amounts
- Customer details

What should you use?

- A. Custom Vision
- B. Azure AI Computer Vision
- C. Azure AI Immersive Reader
- D. Azure AI Document Intelligence

Correct Answer: D

Community vote distribution

D (100%)

✉  **takaimomoGcup** 5 days, 19 hours ago

Selected Answer: D

<https://learn.microsoft.com/en-us/azure/ai-services/?view=doc-intel-4.0.0>

To deal with this type of question, it is good to understand Microsoft's AI services in a nutshell. It is tough to do so by rote memorization.

upvoted 1 times

✉  **Murtuza** 2 months, 1 week ago

To efficiently extract data from scanned invoices, including invoice items, sales amounts, and customer details, you should use Azure AI Document Intelligence. Therefore, the correct answer is D.

upvoted 1 times

✉  **Ghill1982** 2 months, 2 weeks ago

Selected Answer: D

D. Azure AI Document Intelligence

This service is specifically designed for tasks like automated invoice processing and can extract key information from documents.

upvoted 3 times

✉  **Razvan_C** 2 months, 3 weeks ago

Selected Answer: D

Answer seems to be correct

upvoted 1 times

✉  **zoha_zohé** 2 months, 4 weeks ago

Answer is correct

upvoted 2 times

Question #66

HOTSPOT

You are developing a text processing solution.

You have the function shown below.

```
def get_key_words(textAnalyticsClient, text):

    response = textAnalyticsClient.recognize_entities(documents = [text])[0]
    print("Key Words:")
    for entity in response.entities:
        print("\t\t", entity.text)
```

For the second argument, you call the function and specify the following string.

Our tour of Paris included a visit to the Eiffel Tower

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Answer Area**Statements**

Yes	No
<input type="radio"/>	<input type="radio"/>

The output will include the following words: our and included.

<input type="radio"/>	<input type="radio"/>
-----------------------	-----------------------

The output will include the following words: Paris, Eiffel, and Tower.

<input type="radio"/>	<input type="radio"/>
-----------------------	-----------------------

The function will output all the key phrases from the input string to the console.

<input type="radio"/>	<input type="radio"/>
-----------------------	-----------------------

Answer Area**Statements**

Yes	No
<input type="radio"/>	<input checked="" type="radio"/>
<input checked="" type="radio"/>	<input type="radio"/>
<input checked="" type="radio"/>	<input type="radio"/>

Correct Answer: The output will include the following words: our and included.

The output will include the following words: Paris, Eiffel, and Tower.

The function will output all the key phrases from the input string to the console.

 **Harry300** Highly Voted  2 months, 3 weeks ago

Should be NYN.

Key phrases would be (tested with API): Eiffel Tower, tour, Paris, visit

Entities are: Eiffen Tower, Paris

upvoted 10 times

 **takaimomoGcup** Most Recent  5 days, 19 hours ago

NYN are right answer.

upvoted 1 times

 **Murtuza** 1 month, 2 weeks ago

There is a separate function to print key phrase to the console so the answer is NO. They are tricky you into thinking that key phrase is the same thing as key word which its not

upvoted 2 times

 **varinder82** 2 months ago

Final Answer:

N

Y

N

upvoted 2 times

✉  **Murtuza** 2 months, 1 week ago

The correct answers are NYY. The last one is yes because it will send info to the console as per the code snippet
upvoted 1 times

✉  **rober13** 2 months, 1 week ago

Same : Topic 3, Question 48. NYN.
upvoted 1 times

Question #67

HOTSPOT

You are developing a text processing solution.

You develop the following method.

```
def get_key_phrases(text_analytics_client, text):
    response = text_analytics_client.extract_key_phrases(text, language="en")
    print('Key phrases:')
    for keyphrase in response.key_phrases:
        print(f'\t{keyphrase}')
```

You call the method by using the following code.

```
get_key_phrases(text_analytics_client, "the cat sat on the mat")
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area**Statements**

	Yes	No
The call will output key phrases from the input string to the console.	<input type="radio"/>	<input type="radio"/>
The output will contain the following words: the, cat, sat, on, and mat.	<input type="radio"/>	<input type="radio"/>
The output will contain the confidence level for key phrases.	<input type="radio"/>	<input type="radio"/>

Answer Area**Statements**

Yes	No
<input checked="" type="checkbox"/>	<input type="radio"/>
<input type="radio"/>	<input checked="" type="checkbox"/>
<input checked="" type="checkbox"/>	<input type="radio"/>

Correct Answer: The call will output key phrases from the input string to the console.

The output will contain the following words: the, cat, sat, on, and mat.

The output will contain the confidence level for key phrases.

 **audlindr**  2 months, 3 weeks ago

Should be YNN

Refer: <https://learn.microsoft.com/en-us/azure/ai-services/language-service/key-phrase-extraction/quickstart?pivots=programming-language-csharp>

upvoted 8 times

 **takaimomoGcup**  5 days, 19 hours ago

YNN is right answer.

upvoted 1 times

 **varinder82** 2 months ago

Final Answer :

YNN

upvoted 2 times

 **f2c587e** 2 months ago

No se evidencia en el código que se muestre la confiabilidad de ninguna manera. YNN.

upvoted 1 times

 **anto69** 2 months, 1 week ago

Y-N-N for sure

upvoted 2 times

 **Harry300** 2 months, 3 weeks ago

Should be YNN. Seriously, where should the confidence level come from. keyphase is a string

upvoted 2 times

Question #68

HOTSPOT

You are developing a service that records lectures given in English (United Kingdom).

You have a method named append_to_transcript_file that takes translated text and a language identifier.

You need to develop code that will provide transcripts of the lectures to attendees in their respective language. The supported languages are English, French, Spanish, and German.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
speech_key = os.environ['SPEECH_SUBSCRIPTION_KEY']
service_region = os.environ['SPEECH_SERVICE_REGION']

def translate_speech():
    translation_config = speechsdk.translation.SpeechTranslationConfig(
        subscription=speech_key, region=service_region)
    translation_config.speech_recognition_language = "en-GB"
    languages = [
        ('en-GB')
        ('fr', 'de', 'es')
        ('French', 'Spanish', 'German')
        ('languages')]

    for language in languages: translation_config.add_target_language(language)
    audio_config = speechsdk.audio.AudioConfig(use_default_microphone=True)
    recognizer = speechsdk.translation.
        IntentRecognizer()
        SpeakerRecognizer()
        SpeechSynthesizer()
        TranslationRecognizer()

    translation_config=translation_config, audio_config=audio_config)

    result = recognizer.recognize_once()
    if result.reason == speechsdk.ResultReason.TranslatedSpeech:
        append_to_transcript_file(result.text, "en")
        for language in result.translations:
            append_to_transcript_file(result.translations[language], language)
```

Answer Area

```
speech_key = os.environ['SPEECH__SUBSCRIPTION__KEY']
service_region = os.environ['SPEECH__SERVICE__REGION']

def translate_speech():
    translation_config = speechsdk.translation.SpeechTranslationConfig(
        subscription=speech_key, region=service_region)
    translation_config.speech_recognition_language = "en-GB"
    languages = [
        ('en-GB')
        ([ 'fr', 'de', 'es' ])
        ([ 'French', 'Spanish', 'German' ])
        ([ 'languages' ])
```

Correct Answer:

```
for language in languages: translation_config.add_target_language(language)
audio_config = speechsdk.audio.AudioConfig(use_default_microphone=True)
recognizer = speechsdk.translation.  
IntentRecognizer()
SpeakerRecognizer()
SpeechSynthesizer()
TranslationRecognizer()  
translation_config=translation_config, audio_config=audio_config)
result = recognizer.recognize_once()
if result.reason == speechsdk.ResultReason.TranslatedSpeech:
    append_to_transcript_file(result.text, "en")
    for language in result.translations:
        append_to_transcript_file(result.translations[language], language)
```

✉  **takaimomoGcup** 6 days, 19 hours ago

fr,de,es and TranslationRecognizer

upvoted 1 times

✉  **Murtuza** 2 months, 1 week ago

The given answers are correct

upvoted 4 times

Question #69

You are developing an app that will use the text-to-speech capability of the Azure AI Speech service. The app will be used in motor vehicles.

You need to optimize the quality of the synthesized voice output.

Which Speech Synthesis Markup Language (SSML) attribute should you configure?

- A. the style attribute of the mstts:express-as element
- B. the effect attribute of the voice element
- C. the pitch attribute of the prosody element
- D. the level attribute of the emphasis element

Correct Answer: B

Community vote distribution

B (100%)

 takaimomoGcup 5 days, 20 hours ago

Selected Answer: B

voice element

upvoted 1 times

 chandiochan 2 months, 3 weeks ago

Selected Answer: B

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/speech-synthesis-markup-voice>

upvoted 2 times

 chandiochan 2 months, 3 weeks ago

Answer is correct:

The audio effect processor that's used to optimize the quality of the synthesized speech output for specific scenarios on devices.

For some scenarios in production environments, the auditory experience might be degraded due to the playback distortion on certain devices. For example, the synthesized speech from a car speaker might sound dull and muffled due to environmental factors such as speaker response, room reverberation, and background noise. The passenger might have to turn up the volume to hear more clearly. To avoid manual operations in such a scenario, the audio effect processor can make the sound clearer by compensating the distortion of playback.

The following values are supported:

eq_car – Optimize the auditory experience when providing high-fidelity speech in cars, buses, and other enclosed automobiles.

eq_telecomhp8k – Optimize the auditory experience for narrowband speech in telecom or telephone scenarios. You should use a sampling rate of kHz. If the sample rate isn't 8 kHz, the auditory quality of the output speech isn't optimized.

upvoted 4 times

 Harry300 2 months, 3 weeks ago

Selected Answer: B

correct

source: <https://learn.microsoft.com/en-us/azure/ai-services/speech-service/speech-synthesis-markup-voice>:

"Optimize the auditory experience when providing high-fidelity speech in cars, buses, and other enclosed automobiles."

upvoted 1 times

Question #70

You are designing a content management system.

You need to ensure that the reading experience is optimized for users who have reduced comprehension and learning differences, such as dyslexia. The solution must minimize development effort.

Which Azure service should you include in the solution?

- A. Azure AI Immersive Reader
- B. Azure AI Translator
- C. Azure AI Document Intelligence
- D. Azure AI Language

Correct Answer: A

Community vote distribution

A (100%)

-  **chandiochan** Highly Voted 2 months, 3 weeks ago
<https://learn.microsoft.com/en-us/azure/ai-services/immersive-reader/overview>
upvoted 5 times
-  **rober13** 2 months, 1 week ago
thanks for the reference!
upvoted 1 times
-  **michaelmorar** Most Recent 1 month, 1 week ago
Selected Answer: A
Immersive AI improves accessibility.
upvoted 1 times

Question #71

HOTSPOT

You are building an app that will answer customer calls about the status of an order. The app will query a database for the order details and provide the customers with a spoken response.

You need to identify which Azure AI service APIs to use. The solution must minimize development effort.

Which object should you use for each requirement? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Convert customer calls into text queries:

- SpeechRecognizer
- SpeechSynthesizer
- TranslationRecognizer
- VoiceProfileClient

Provide customers with the order details:

- SpeechRecognizer
- SpeechSynthesizer
- TranslationRecognizer
- VoiceProfileClient

Answer Area

Convert customer calls into text queries:

- SpeechRecognizer
- SpeechSynthesizer
- TranslationRecognizer
- VoiceProfileClient

Correct Answer:

Provide customers with the order details:

- SpeechRecognizer
- SpeechSynthesizer
- TranslationRecognizer
- VoiceProfileClient

 **Harry300** Highly Voted 2 months, 3 weeks ago

SpeechRecognizer / SpeechSynthesizer
doesn't say anything about translation recognizer, which needs more code
upvoted 9 times

 **audlindr** 2 months, 3 weeks ago

Agreed
<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/get-started-speech-to-text?tabs=windows%2Cterminal&pivots=programming-language-csharp>

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/how-to-speech-synthesis?tabs=browserjs%2Cterminal&pivots=programming-language-csharp>
upvoted 2 times

 **takaimomoGcup** Most Recent 5 days, 19 hours ago

It must be SpeechRecognizer and SpeechSynthesizer.
upvoted 1 times

✉  **anto69** 1 month, 3 weeks ago

SpeechRecognizer / SpeechSynthesizer
upvoted 1 times

✉  **Murtuza** 1 month, 3 weeks ago

To convert customer calls into text queries, you should use the SpeechRecognizer object. This is part of the Speech service in Azure Cognitive Services and it converts spoken language into written text.

To provide customers with the order details in a spoken response, you should use the SpeechSynthesizer object. This service converts text into lifelike speech, which can be used to deliver information to customers in a more interactive way.

So, the answer is:

Convert customer calls into text queries: SpeechRecognizer

Provide customers with the order details: SpeechSynthesizer

upvoted 2 times

✉  **varinder82** 2 months ago

Final Answer;

1. SpeechRecognizer
2. SpeechSynthesizer

upvoted 1 times

Question #72

You have an Azure AI service model named Model1 that identifies the intent of text input.

You develop a Python app named App1.

You need to configure App1 to use Model1.

Which package should you add to App1?

- A. azure-cognitiveservices-language-textanalytics
- B. azure-ai-language-conversations
- C. azure-mgmt-cognitiveservices
- D. azure-cognitiveservices-speech

Correct Answer: D

Community vote distribution

B (86%)

14%

✉ **warrior1234** Highly Voted 2 months, 3 weeks ago

The correct package for working with Azure AI service for text intent identification, like Model1, is:

- A. azure-cognitiveservices-language-textanalytics

Therefore, you should add the azure-cognitiveservices-language-textanalytics package to App1 for configuring it to use Model1. This package provides functionalities for working with the Language API in Azure Cognitive Services, which includes text analytics capabilities, such as identifying language and sentiment, and can be used for processing text input to determine intent.

upvoted 5 times

✉ **takaimomoGcup** Most Recent 5 days, 19 hours ago

Selected Answer: D

When identifying text input intent, it is usually aggregated to cognitiveservices. not just Microsoft, but mega cloud products and solutions have too many similar names and should be aggregated.

upvoted 1 times

✉ **QwertySunny** 1 month ago

Selected Answer: B

.....

upvoted 1 times

✉ **Murtuza** 1 month, 2 weeks ago

Selected Answer: B

B is correct

upvoted 1 times

✉ **chandiochan** 2 months ago

Selected Answer: B

Must be B

upvoted 1 times

✉ **9ae4fb5** 2 months, 1 week ago

answer is B.

"Conversation App: It's used in extracting intents and entities in conversations"

- <https://learn.microsoft.com/ko-kr/python/api/azure/ai-language-conversations-readme?view=azure-python>

"The (Text Analytics) API can be used to analyze unstructured text for tasks such as sentiment analysis, key phrase extraction and language detection."

<https://learn.microsoft.com/en-us/python/api/azure-cognitiveservices-language-textanalytics/azure.cognitiveservices.language.textanalytics.textanalyticsclient?view=azure-python-previous>

upvoted 2 times

✉ **anto69** 1 month, 3 weeks ago

Agree 100%

upvoted 1 times

✉️ **Ghill1982** 2 months, 2 weeks ago

Selected Answer: B

B. azure-ai-language-conversations

This package is part of the Azure AI SDK for Python and contains functionality for building applications that can understand user intent using language models.

upvoted 1 times

✉️ **Harry300** 2 months, 3 weeks ago

Selected Answer: B

Should be B.

Source: <https://pypi.org/project/azure-ai-language-conversations/>

"Conversation App: It's used in extracting intents and entities in conversations"

upvoted 2 times

Question #73

HOTSPOT

You are building an app that will automatically translate speech from English to French, German, and Spanish by using Azure AI service.

You need to define the output languages and configure the Azure AI Speech service.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
speech_key, service_region = os.environ['SPEECH__SERVICE__KEY'], os.environ['SPEECH__SERVICE__REGION']
languages = [
    ['en-GB'],
    {'en','fr','de','es'},
    ['fr','de','es'],
    {"French","Spanish","German"}
]

def translate_speech_to_text():
    translation_config = speechsdk.translation.SpeechTranslationConfig(subscription=speech_key, region=service_region)
    for lang in languages:
        translation_config.add_target_language(lang)
    for lang in languages:
        translation_config.add_target_language(lang)
    recognizer = speechsdk.translation. (translation_config=translation_config)
    ...
    IntentRecognizer
    SpeakerRecognizer
    SpeechSynthesizer
    TranslationRecognizer
```

Correct Answer:

```
Answer Area
speech_key, service_region = os.environ['SPEECH__SERVICE__KEY'], os.environ['SPEECH__SERVICE__REGION']
languages = [
    ['en-GB'],
    {'en','fr','de','es'},
    ['fr','de','es'],
    {"French", "Spanish", "German"}
]
def translate_speech_to_text():
    translation_config = speechsdk.translation.SpeechTranslationConfig(subscription=speech_key, region=service_region)
    for lang in languages:
        translation_config.add_target_language(lang)
    for lang in languages:
        translation_config.add_target_language(lang)
    recognizer = speechsdk.translation. (translation_config=translation_config)
    ...
    IntentRecognizer
    SpeakerRecognizer
    SpeechSynthesizer
    TranslationRecognizer
```

 anto69 1 month, 3 weeks ago

Duplicate

upvoted 1 times

 rober13 2 months, 1 week ago

Same Question 68

upvoted 2 times

 f2c587e 1 month, 4 weeks ago

The answer options for this question are slightly different from question 68 on topic 3

upvoted 1 times

 Razvan_C 2 months, 3 weeks ago

Answer seems to be correct.

upvoted 4 times

Question #74

DRAG DROP

You plan to implement an Azure AI Search resource that will use custom skill based on sentiment analysis.

You need to create a custom model and configure Azure AI Search use the model.

Which five actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

- Create an endpoint for the model.
- Rerun the indexer to enrich the index.
- Create an Azure Machine Learning workspace.
- Create and train the model in the Azure Machine Learning studio.
- Provision an Azure AI Services resource and obtain the endpoint.
- Connect the custom skill the endpoint.

Answer Area**Correct Answer:**

- Answer Area**
- Create an Azure Machine Learning workspace.
 - Provision an Azure AI Services resource and obtain the endpoint.
 - Connect the custom skill the endpoint.
 - Create and train the model in the Azure Machine Learning studio.
 - Rerun the indexer to enrich the index.

 **takaimomoGcup** 5 days, 19 hours ago

Create workspace
Provision
Connect
Create studio
Rerun
upvoted 1 times

 **3a0b61c** 2 months ago

<https://learn.microsoft.com/en-us/training/modules/build-azure-machine-learn-custom-skill-for-azure-cognitive-search/03-enrich-search-index-use-model>
1.Create an Azure Machine Learning workspace.
2.Create and train the model in the Azure Machine Learning studio.
3.Create an endpoint for the model.
4.Provision an Azure AI Services resource and obtain the endpoint.
5.Connect the custom skill to the endpoint.
upvoted 3 times

 **MonicaKarim** 3 weeks, 3 days ago

in the link you provided, last step is to update your existing cognitive search that means is to create the search resource first:
1- Provision an Azure AI Services resource and obtain the endpoint.
2- Create an Azure Machine Learning workspace.
3- Create and train the model in the Azure Machine Learning studio.
4- Create an endpoint for the model.
6- Connect the custom skill to the endpoint.
upvoted 1 times

 **Murtuza** 2 months, 1 week ago

Create an Azure Machine Learning workspace.
Create and train the model in the Azure Machine Learning studio.
Create an endpoint for the model.
Provision an Azure AI Services resource and obtain the endpoint.
Connect the custom skill to the endpoint.
upvoted 1 times

 **rober13** 2 months, 1 week ago

Create an endpoint for the model.
Create and train the model in the Azure Machine Learning studio.

Create an Azure Machine Learning workspace.
Connect the custom skill to the endpoint.

so, with the options:

Create an Azure Machine Learning
Create and train
Provision an Azure
Create an endpoint
Connect

upvoted 1 times

✉ **chandiochan** 2 months, 2 weeks ago

Create an Azure Machine Learning workspace.
Create and train the model in the Azure Machine Learning studio.
Create an endpoint for the model.
Provision an Azure AI Services resource and obtain the endpoint.
Connect the custom skill to the endpoint.

upvoted 2 times

✉ **Delta64** 2 months, 3 weeks ago

According to GPT4:

1. Create an Azure Machine Learning workspace.
2. Create and train the model in the Azure Machine Learning studio.
3. Provision an Azure AI Services resource and obtain the endpoint.
4. Create an endpoint for the model.

Connect the custom skill to the endpoint.

Rerun the indexer to enrich the index.

upvoted 1 times

✉ **Delta64** 2 months, 3 weeks ago

My bad, this is wrong information.

upvoted 1 times

Question #75

HOTSPOT

You have a collection of press releases stored as PDF files.

You need to extract text from the files and perform sentiment analysis.

Which service should you use for each task? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Extract text:

- Azure AI Search
- Azure AI Vision
- Azure AI Document Intelligence

Perform sentiment analysis:

- Azure Cognitive Search
- Azure AI Computer Vision
- Azure AI Document Intelligence
- Azure AI Language

Answer Area

Extract text:

- Azure AI Search
- Azure AI Vision
- Azure AI Document Intelligence

Correct Answer:

Perform sentiment analysis:

- Azure Cognitive Search
- Azure AI Computer Vision
- Azure AI Document Intelligence
- Azure AI Language

✉  **takaimomoGcup** 5 days, 19 hours ago

AI Document Intelligence and AI Language

upvoted 1 times

✉  **Ghill1982** 1 month ago

For extracting text from PDF files, Azure Document Intelligence would be the better choice. It's specifically optimized for text-heavy documents like PDFs and includes features such as higher-resolution scanning and paragraph detection. In addition, Azure Computer Vision API does not have direct PDF integration, and you would need to convert PDFs to images before text extraction, which adds an extra step to the process.

upvoted 2 times

✉  **Murtuza** 1 month, 2 weeks ago

Sentiment Analysis = AI Language

upvoted 1 times

✉  **Murtuza** 1 month, 3 weeks ago

the key word to answer this question is press release which implies computer vision

upvoted 1 times

✉  **Murtuza** 2 months, 1 week ago

Extract Text = AI Vision

upvoted 2 times

✉  **Mehe323** 2 months ago

Document Intelligence also does text extraction:

"Document Intelligence Read Optical Character Recognition (OCR) model runs at a higher resolution than Azure AI Vision Read and extracts print and handwritten text from PDF documents and scanned images."

<https://learn.microsoft.com/en-us/azure/ai-services/document-intelligence/concept-read?view=doc-intel-4.0.0>

upvoted 1 times

✉️ **fabrea** 2 months, 1 week ago

I disagree, the Computer Vision for text extraction is best suited for unstructured documents
IDP is the right choice

upvoted 1 times

✉️ **Razvan_C** 2 months, 2 weeks ago

Answer seems to be correct.

upvoted 1 times

Topic 4 - Question Set 4**Question #1****Topic 4****HOTSPOT -**

You are developing a text processing solution.

You develop the following method.

```
static void GetKeyPhrases(TextAnalyticsClient textAnalyticsClient, string text)
{
    var response = textAnalyticsClient.ExtractKeyPhrases(text);
    Console.WriteLine("Key phrases:");

    foreach (string keyphrase in response.Value)
    {
        Console.WriteLine($"\\t{keyphrase}");
    }
}
```

You call the method by using the following code.

```
GetKeyPhrases(textAnalyticsClient, "the cat sat on the mat");
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area**Statements****Yes****No**

The call will output key phrases from the input string to the console.

The output will contain the following words: the, cat, sat, on, and mat.

The output will contain the confidence level for key phrases.

Correct Answer:**Answer Area****Statements****Yes****No**

The call will output key phrases from the input string to the console.

The output will contain the following words: the, cat, sat, on, and mat.

The output will contain the confidence level for key phrases.

Box 1: Yes -

The Key Phrase Extraction API evaluates unstructured text, and for each JSON document, returns a list of key phrases.

Box 2: No -

'the' is not a key phrase.

This capability is useful if you need to quickly identify the main points in a collection of documents. For example, given input text "The food was delicious and there were wonderful staff", the service returns the main talking points: "food" and "wonderful staff".

Box 3: No -

Key phrase extraction does not have confidence levels.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/how-tos/text-analytics-how-to-keyword-extraction>

 **joecruz**  2 years, 11 months ago

I think the answer is correct

upvoted 30 times

 **zellick**  11 months ago

YNN is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/key-phrase-extraction/overview>

Key phrase extraction is one of the features offered by Azure Cognitive Service for Language, a collection of machine learning and AI algorithms in the cloud for developing intelligent applications that involve written language. Use key phrase extraction to quickly identify the main concepts in text. For example, in the text "The food was delicious and the staff were wonderful.", key phrase extraction will return the main topics: "food" and "wonderful staff".

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/key-phrase-extraction/quickstart?pivots=programming-language-csharp#output>

upvoted 5 times

✉  **zellick** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 4 times

✉  **reiwanotora** Most Recent 2 days, 5 hours ago

Yes

No

No

upvoted 1 times

✉  **rober13** 2 months, 1 week ago

Same Question 67, Topic 3

upvoted 1 times

✉  **rdemontis** 6 months, 3 weeks ago

correct answer

upvoted 2 times

✉  **Eltooth** 1 year, 10 months ago

Answer is correct.

Yes

No

No

upvoted 2 times

✉  **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 4 times

✉  **Adedoyin_Simeon** 2 years, 10 months ago

The answer is correct

upvoted 4 times

✉  **azurelearner666** 2 years, 11 months ago

The response is correct.

Check the sample code at https://github.com/Azure/azure-sdk-for-net/blob/master/sdk/textanalytics/Azure.AI.TextAnalytics/samples/Sample3_ExtractKeyPhrases.md

upvoted 4 times

✉  **WillyMac** 2 years, 11 months ago

The First option is also NO.

The code doesn't function, because "response" is of Type "KeyPhraseCollection", and this Class is a ReadOnlyCollection<string> and does not have a definition for: "Value"

upvoted 1 times

✉  **cavefish** 2 years, 11 months ago

response is of type Response<KeyPhraseCollection>, so value is the field with the result collection <https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/quickstarts/client-libraries-rest-api?tabs=version-3-1&pivots=programming-language-csharp#key-phrase-extraction>

upvoted 2 times

Question #2

You deploy a web app that is used as a management portal for indexing in Azure Cognitive Search. The app is configured to use the primary admin key.

During a security review, you discover unauthorized changes to the search index. You suspect that the primary access key is compromised.

You need to prevent unauthorized access to the index management endpoint. The solution must minimize downtime.

What should you do next?

- A. Regenerate the primary admin key, change the app to use the secondary admin key, and then regenerate the secondary admin key.
- B. Change the app to use a query key, and then regenerate the primary admin key and the secondary admin key.
- C. Regenerate the secondary admin key, change the app to use the secondary admin key, and then regenerate the primary key.
- D. Add a new query key, change the app to use the new query key, and then delete all the unused query keys.

Correct Answer: A

Regenerate admin keys.

Two admin keys are created for each service so that you can rotate a primary key, using the secondary key for business continuity.

1. In the Settings > Keys page, copy the secondary key.
2. For all applications, update the API key settings to use the secondary key.
3. Regenerate the primary key.
4. Update all applications to use the new primary key.

Note: Two admin api-keys, referred to as primary and secondary keys in the portal, are automatically generated when the service is created and can be individually regenerated on demand. Having two keys allows you to roll over one key while using the second key for continued access to the service.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-security-api-keys#regenerate-admin-keys>

Community vote distribution

C (86%)

10%

✉️  **Nouna**  2 years, 11 months ago

logicly it should be C. A can not be the answer as you regenerated the secndry key after your already add to your app. the right order to minize downtime is to regenerate the secondary, add to your app and then regenerate the primary
upvoted 54 times

✉️  **rde montis** 6 months, 3 weeks ago

totally agree with you
upvoted 1 times

✉️  **satishk4u** 2 years ago

wondering why should we regenerate the secondary key when primary key is compromised.
upvoted 2 times

✉️  **RamonKaus** 1 year, 10 months ago

In order to keep zero downtime on your application, you need to make sure it has A key so it can use. You must assume both keys are compromised and get new keys.
upvoted 4 times

✉️  **GustavoR10**  2 years, 9 months ago

If you change the application to use the secondary key and then you regenerate the key the application in not going to work. So C is the right one
upvoted 8 times

✉️  **anto69**  1 month, 3 weeks ago

Selected Answer: C
C for sure
upvoted 3 times

✉️  **Murtuza** 1 month, 3 weeks ago

Try to separate out the admin keys vs the query keys. The question is NOT ABOUT QUERY KEYS SO THAT rules out 2 answers
upvoted 2 times

✉️  **evangelist** 3 months, 3 weeks ago

for minimum to no downtime: The Answer is C,
Azure Cognitive Search provides two admin keys (primary and secondary) to facilitate key rotation without downtime. The existence of two keys is specifically designed to handle scenarios where one key may need to be changed (e.g., if it is compromised) without interrupting the service.

Immediate switch to the secondary admin key: By immediately switching the application to use the secondary admin key (which is assumed to be secure), you ensure that the application retains uninterrupted access to the management capabilities of Azure Cognitive Search.

Regenerate the primary admin key: Once the application is using the secondary admin key, regenerating the primary admin key does not cause downtime because the application is not currently using this key. The new primary key can be stored securely for future use or key rotation purposes.

upvoted 1 times

 **davidorti** 4 months, 1 week ago

Selected Answer: B

"The app is configured to use the *primary admin key*." ...

"You suspect that the primary access key is compromised. You need to prevent unauthorized access to the index management endpoint."

I think it's B: "Change the app to use a query key, and then regenerate the primary admin key and the secondary admin key."

So basically there are two admin keys. One of those is compromised. Both are reset for security. This way, the app is just using a query key with no enough privileges to make changes to the Cognitive Search service, so even if that key was leaked again unauthorized access to the management endpoint can be made. <https://learn.microsoft.com/en-us/azure/search/search-security-api-keys?tabs=portal-use%2Cportal-find%2Cportal-query>

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: C

To me the correct answer is C.

<https://learn.microsoft.com/en-us/azure/search/search-security-api-keys?tabs=portal-use%2Cportal-find%2Cportal-query>

upvoted 1 times

 **sl_mslconsulting** 7 months ago

Selected Answer: C

The moment you see "Regenerate the primary admin key" as the first action you already know it violates the minimum downtime requirement. Answer A makes it even worse - the web app will stop working all together in the end as the web app will be using an invalid secondary admin key. For answer C, regenerating the secondary admin key seems redundant as you suspect only the primary access key is compromised but it's safer and meet the minimum downtime requirement anyway. As for answer B, it won't provide the required permissions to manage the indexes hence the app won't be functioning and this violates the requirement. This question requires you to really think it through or you might be tricked easily

upvoted 1 times

 **katrang** 7 months, 1 week ago

Selected Answer: C

To avoid any downtime, regenerate the second key and use it (as a precaution, and only if it is not currently being used by anything else). Then regenerate the first key

upvoted 1 times

 **zellck** 11 months ago

Selected Answer: C

C is the answer.

<https://learn.microsoft.com/en-us/azure/search/search-security-api-keys?tabs=portal-use%2Cportal-find%2Cportal-query#regenerate-admin-key>
Two admin keys are created for each service so that you can rotate a primary key while using the secondary key for business continuity.

- Under Settings, select Keys, then copy the secondary key.
- For all applications, update the API key settings to use the secondary key.
- Regenerate the primary key.
- Update all applications to use the new primary key.

upvoted 2 times

 **Drummer** 11 months, 2 weeks ago

C. Regenerate the secondary admin key, change the app to use the secondary admin key, and then regenerate the primary key --- By regenerating the secondary admin key first, you can address the potential compromise of the primary key. Then, changing the app to use the regenerated secondary admin key ensures continued access to the index management endpoint. Finally, regenerating the primary key provides a new security key for future use

upvoted 1 times

 **kail85** 12 months ago

Selected Answer: C

C is correct

upvoted 1 times

 **EliteAllen** 12 months ago

Selected Answer: C

C is right.

upvoted 1 times

 **ap1234pa** 1 year, 4 months ago

Selected Answer: C

C is correct

upvoted 1 times

 **slcheng** 1 year, 5 months ago

Selected Answer: C

Vote for C

upvoted 1 times

 **marti_tremblay000** 1 year, 5 months ago

Selected Answer: C

C is the most logical. The order stated by the A answer doesn't make any sense.

upvoted 1 times

 **Josua2020** 1 year, 8 months ago

Selected Answer: C

Why not C?

upvoted 1 times

Question #3

You have an existing Azure Cognitive Search service.

You have an Azure Blob storage account that contains millions of scanned documents stored as images and PDFs.

You need to make the scanned documents available to search as quickly as possible.

What should you do?

- A. Split the data into multiple blob containers. Create a Cognitive Search service for each container. Within each indexer definition, schedule the same runtime execution pattern.
- B. Split the data into multiple blob containers. Create an indexer for each container. Increase the search units. Within each indexer definition, schedule a sequential execution pattern.
- C. Create a Cognitive Search service for each type of document.
- D. Split the data into multiple virtual folders. Create an indexer for each folder. Increase the search units. Within each indexer definition, schedule the same runtime execution pattern.

Correct Answer: D

Incorrect Answers:

A: Need more search units to process the data in parallel.

B: Run them in parallel, not sequentially.

C: Need a blob indexer.

Note: A blob indexer is used for ingesting content from Azure Blob storage into a Cognitive Search index.

Index large datasets -

Indexing blobs can be a time-consuming process. In cases where you have millions of blobs to index, you can speed up indexing by partitioning your data and using multiple indexers to process the data in parallel. Here's how you can set this up:

- ☞ Partition your data into multiple blob containers or virtual folders
- ☞ Set up several data sources, one per container or folder.
- ☞ Create a corresponding indexer for each data source. All of the indexers should point to the same target search index.
- ☞ One search unit in your service can run one indexer at any given time. Creating multiple indexers as described above is only useful if they actually run in parallel.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-howto-indexing-azure-blob-storage>

Community vote distribution

D (100%)

✉  **reiwanotora** 2 days, 5 hours ago

Selected Answer: D

FOCUS "virtual folders" word.

upvoted 1 times

✉  **Murtuza** 1 month, 3 weeks ago

Tricky question think of virtual folder AS blob containers and the answer will be obvious

upvoted 1 times

✉  **evangelist** 3 months, 3 weeks ago

Selected Answer: D

e, option D is the best choice because it leverages the scalability and parallel processing capabilities of Azure Cognitive Search to efficiently index large volume of documents. By organizing documents into virtual folders and creating an indexer for each folder, you can maximize the throughput of the indexing process. Increasing search units further supports this by allocating more resources to the task, thereby minimizing the time required to make the scanned documents searchable.

upvoted 1 times

✉  **rdemontis** 6 months, 3 weeks ago

Selected Answer: D

I think correct answer is D

<https://learn.microsoft.com/en-us/azure/search/search-howto-large-index#run-indexers-in-parallel>

upvoted 2 times

✉  **sl_mslconsulting** 7 months ago

Selected Answer: D

"One search unit in your service can run one indexer at any given time. Creating multiple indexers is only useful if they can run in parallel" so A and C are out. B is out as you are not running the indexers in parallel. Besides it's hard to image that with millions of scanned you don't have virtual folders in place to split the data already.

upvoted 1 times

✉ **zellck** 11 months ago

Selected Answer: D

D is the answer.

<https://learn.microsoft.com/en-us/azure/search/search-howto-large-index#run-indexers-in-parallel>

If you partition your data, you can create multiple indexer-data-source combinations that pull from each data source and write to the same search index. Because each indexer is distinct, you can run them at the same time, populating a search index more quickly than if you ran them sequentially.

Make sure you have sufficient capacity. One search unit in your service can run one indexer at any given time. Creating multiple indexers is only useful if they can run in parallel.

upvoted 2 times

✉ **zellck** 11 months ago

If your data source is an Azure Blob Storage container or Azure Data Lake Storage Gen 2, enumerating a large number of blobs can take a long time (even hours) until this operation is completed. This will cause that your indexer's documents succeeded count isn't increased during that time and it may seem it's not making any progress, when it is. If you would like document processing to go faster for a large number of blobs, consider partitioning your data into multiple containers and create parallel indexers pointing to a single index.

upvoted 3 times

✉ **Eltooth** 1 year, 10 months ago

Selected Answer: D

D is correct answer.

Also marked correct on Udemy course practice test.

upvoted 4 times

✉ **PHD_CHENG** 1 year, 11 months ago

Was on exam 7 Jun 2022

upvoted 2 times

✉ **prabhjot** 2 years, 3 months ago

correct ans

upvoted 1 times

✉ **azurelearner666** 2 years, 11 months ago

how to do this is defined here:

<https://docs.microsoft.com/en-us/azure/search/search-howto-indexing-azure-blob-storage#index-large-datasets>

The response is missing the data source creation for each virtual folder or blob container.

D is not correct, but the less wrong of a response...

So I give it a "pass", nowadays it is misleading and not fully correct...

upvoted 2 times

✉ **azurelearner666** 2 years, 11 months ago

seems to be correct

upvoted 2 times

Question #4

You need to implement a table projection to generate a physical expression of an Azure Cognitive Search index.

Which three properties should you specify in the skillset definition JSON configuration table node? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. tableName
- B. generatedKeyName
- C. dataSource
- D. dataSourceConnection
- E. source

Correct Answer: ABE

Defining a table projection.

Each table requires three properties:

- ☞ tableName: The name of the table in Azure Storage.
- ☞ generatedKeyName: The column name for the key that uniquely identifies this row.
- ☞ source: The node from the enrichment tree you are sourcing your enrichments from. This node is usually the output of a shaper, but could be the output of any of the skills.

Reference:

<https://docs.microsoft.com/en-us/azure/search/knowledge-store-projection-overview>

Community vote distribution

ABE (100%)

✉ PHD_CHENG Highly Voted 2 years ago

Selected Answer: ABE

Answer is correct. Below link shows the details

<https://docs.microsoft.com/en-us/azure/search/knowledge-store-projections-examples>

upvoted 9 times

✉ evangelist Most Recent 3 months, 3 weeks ago

Selected Answer: ABE

For implementing a table projection in a skillset definition for Azure Cognitive Search, the correct properties to specify in the JSON configuration are:

- A. tableName: Necessary for defining where the output data should be stored.
 - B. generatedKeyName: Important for scenarios requiring unique identification of rows, though its necessity can vary.
 - E. source: Essential for specifying the input data that will be processed and projected into the table.
- Therefore, the correct answers are A (tableName), B (generatedKeyName), and E (source), as they play direct roles in the configuration of a table projection within a skillset for Azure Cognitive Search.

upvoted 2 times

✉ rdemontis 6 months, 3 weeks ago

Selected Answer: ABE

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-projections-examples#define-a-table-projection>

upvoted 2 times

✉ Tin_Tin 11 months, 1 week ago

correct

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-projections-examples#define-a-table-projection>

upvoted 1 times

✉ Eltooth 1 year, 10 months ago

Selected Answer: ABE

A, B and E are correct answers.

To define a table projection, use the tables array in the projections property. A table projection has three required properties:

tableName: Determines the name of a new table created in Azure Table Storage.

generatedKeyName: Column name for the key that uniquely identifies each row. The value is system-generated. If you omit this property, a column will be created automatically that uses the table name and "key" as the naming convention.

source: A path to a node in an enrichment tree. The node should be a reference to a complex shape that determines which columns are created in the table.

<https://docs.microsoft.com/en-us/azure/search/knowledge-store-projections-examples#define-a-table-projection>

upvoted 1 times

Question #5

HOTSPOT -

You are creating an enrichment pipeline that will use Azure Cognitive Search. The knowledge store contains unstructured JSON data and scanned PDF documents that contain text.

Which projection type should you use for each data type? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

JSON data:

File projection
Object projection
Table projection

Scanned data:

File projection
Object projection
Table projection

Answer Area

JSON data:

File projection
Object projection
Table projection

Correct Answer:

Scanned data:

File projection
Object projection
Table projection

Box 1: Object projection -

Object projections are JSON representations of the enrichment tree that can be sourced from any node.

Box 2: File projection -

File projections are similar to object projections and only act on the normalized_images collection.

Reference:

<https://docs.microsoft.com/en-us/azure/search/knowledge-store-projection-overview>

 **azurelearner666**  2 years, 11 months ago

Correct!

upvoted 14 times

 **zellck**  11 months ago

1. Object
2. File

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-projections-examples#define-an-object-projection>

Object projections are JSON representations of the enrichment tree that can be sourced from any node. In comparison with table projections, object projections are simpler to define and are used when projecting whole documents. Object projections are limited to a single projection in a container and can't be sliced.

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-projections-examples#define-a-file-projection>

File projections are always binary, normalized images, where normalization refers to potential resizing and rotation for use in skillset execution. File projections, similar to object projections, are created as blobs in Azure Storage, and contain binary data (as opposed to JSON).

upvoted 5 times

✉ **rdemontis** Most Recent 6 months, 3 weeks ago

Correct answer.

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-projections-examples>

upvoted 3 times

✉ **sl_msconsulting** 7 months ago

we are dealing with unstructured JSON data here. Due to the potential irregularities in the data it will probably not be feasible using table projection. Using object projection is simpler and sometimes the only choice depending on the nature of the data.

upvoted 1 times

✉ **RamonKaus** 1 year, 10 months ago

java script OBJECT notation (JSON)

upvoted 2 times

✉ **Eltooth** 1 year, 10 months ago

Object
File

<https://docs.microsoft.com/en-us/azure/search/knowledge-store-projection-overview#types-of-projections-and-usage>

upvoted 1 times

✉ **reachmymind** 2 years, 2 months ago

Answer 1: Object Projection: Used when you need the full JSON representation of your data and enrichments in one JSON document. As with table projections, only valid JSON objects can be projected as objects, and shaping can help you do that.

Answer 2: File Projection: Used when you need to save normalized, binary image files, eg: Optical Character Recognition (OCR) extracts text from binary files

Not Answer: Table Projection: Used for data that's best represented as rows and columns, or whenever you need granular representations of your data (for example, as data frames). Table projections allow you to define a schematized shape, using a Shaper skill or use inline shaping to specify columns and rows.

upvoted 2 times

✉ **Deepusuraj** 2 years, 3 months ago

Types of projections and usage

A knowledge store is a logical construction that's physically expressed as a loose collection of tables, JSON objects, or binary image files in Azure Storage.

TYPES OF PROJECTIONS AND USAGE

Projection Storage Usage

Tables Azure Table Storage Used for data that's best represented as rows and columns, or whenever you need granular representations of your data (for example, as data frames). Table projections allow you to define a schematized shape, using a Shaper skill or use inline shaping to specify columns and rows. You can organize content into multiple tables based on familiar normalization principles. Tables that are in the same group are automatically related.

Objects Azure Blob Storage Used when you need the full JSON representation of your data and enrichments in one JSON document. As with table projections, only valid JSON objects can be projected as objects, and shaping can help you do that.

Files Azure Blob Storage Used when you need to save normalized, binary image files.

upvoted 1 times

Question #6

HOTSPOT -

You are building an Azure Cognitive Search custom skill.

You have the following custom skill schema definition.

```
{
  "@odata.type": "#Microsoft.Skills.Custom.WebApiSkill",
  "description": "My custom skill description",
  "uri": "https://contoso-webskill.azurewebsites.net/api/process",
  "context": "/document/organizations/*",
  "inputs": [
    {
      "name": "companyName",
      "source": "/document/organizations/*"
    }
  ],
  "outputs": [
    {
      "name": "companyDescription"
    }
  ]
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
------------	-----	----

CompanyDescription is available for indexing.

The definition calls a web API as part of the enrichment process.

The enrichment step is called only for the first organization under "/document/organizations".

Answer Area

Statements	Yes	No
------------	-----	----

Correct Answer: CompanyDescription is available for indexing.

The definition calls a web API as part of the enrichment process.

The enrichment step is called only for the first organization under "/document/organizations".

Box 1: Yes -

Once you have defined a skillset, you must map the output fields of any skill that directly contributes values to a given field in your search index.

Box 2: Yes -

The definition is a custom skill that calls a web API as part of the enrichment process.

Box 3: No -

For each organization identified by entity recognition, this skill calls a web API to find the description of that organization.

Reference:

<https://docs.microsoft.com/en-us/azure/search/cognitive-search-output-field-mapping>

  **zellick**  11 months ago

YYN is the answer.

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-custom-skill-web-api#skill-outputs>

There are no predefined outputs for this skill. Be sure to define an output field mapping in the indexer if the skill's output should be sent to a field in the search index.

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-defining-skillset#add-skills>

You can build complex skillsets with looping and branching using the Conditional skill to create the expressions. The syntax is based on the JSON Pointer path notation, with a few modifications to identify nodes in the enrichment tree. A "/" traverses a level lower in the tree and "*" acts as a for-each operator in the context.

upvoted 12 times

✉ **reiwanotora** Most Recent 2 days, 6 hours ago

No

Yes

No

upvoted 1 times

✉ **anto69** 1 month, 2 weeks ago

N, Y, N

upvoted 1 times

✉ **Murtuza** 1 month, 3 weeks ago

CompanyDescription is available for indexing: Yes. The companyDescription is defined in the outputs section of the custom skill, which means it will be produced by the skill and can be used in subsequent steps of the enrichment pipeline, including indexing.

The definition calls a web API as part of the enrichment process: Yes. The uri field in the custom skill definition is set to <https://contoso-webskill.azurewebsites.net/api/process>, which is a web API endpoint. This means that the custom skill will call this web API as part of the enrichment process.

The enrichment step is called only for the first organization under "/document/organizations": No. The context and source fields are set to /document/organizations/*, which means the custom skill will be applied to each organization under /document/organizations, not just the first one.

upvoted 2 times

✉ **varinder82** 2 months ago

Final Answer:

NYN

upvoted 1 times

✉ **Mehe323** 2 months, 2 weeks ago

The first answer should be no:

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-output-field-mapping?tabs=rest>

"Output field mappings are required if your indexer has an attached skillset that creates new information, such as text translation or key phrase extraction. During indexer execution, AI-generated information exists in memory only. To persist this information in a search index, you'll need to tell the indexer where to send the data."

According to my source, the code needs 'outputFieldMappings ', and that is not the case in this example.

upvoted 3 times

✉ **rdemontis** 6 months, 3 weeks ago

I think the answer is correct

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-custom-skill-web-api>

upvoted 3 times

✉ **Mehe323** 2 months, 2 weeks ago

The first answer should be no:

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-output-field-mapping?tabs=rest>

"Output field mappings are required if your indexer has an attached skillset that creates new information, such as text translation or key phrase extraction. During indexer execution, AI-generated information exists in memory only. To persist this information in a search index, you'll need to tell the indexer where to send the data."

According to my source, the code needs 'outputFieldMappings ', and here it is not the case.

upvoted 3 times

✉ **sl_msIconsulting** 7 months, 2 weeks ago

I think it's No for the first question. It's basically testing you if you know this: There are no predefined outputs for this skill. Be sure to define an output field mapping in the indexer if the skill's output should be sent to a field in the search index. So by default, it is not available for indexing. I can be available if you do whatever needs to be done as the above described.

upvoted 2 times

✉ **sl_msIconsulting** 7 months ago

Another way to look at this is that the output of a skill can be sent to a field in a search index, a knowledge store, or a downstream skill. You have choices not to send the output to a search index but use it for other purpose. It will be silly to throw it away then why bother to have this skill in the skillset to begin with.

upvoted 1 times

✉ **M25** 8 months, 4 weeks ago

N, Y, N

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-defining-skillset#send-output-to-a-destination>

Although skill output can be optionally cached for reuse purposes, it's usually temporary and exists only while skill execution is in progress.

- To send output to a field in a search index, create an output field mapping in an indexer.
- To send output to a knowledge store, create a projection.

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-output-field-mapping?tabs=rest>

In contrast with a fieldMappings definition that maps a path between two physical data structures, an outputFieldMappings definition maps in-memory data to fields in a search index.

upvoted 3 times

 **pid** 8 months ago

Agree with zellck's answer. First one is Y: You can define output field mapping after Output to index 'CompanyDescription'. Techncially it is available for indexing.

upvoted 2 times

 **halfway** 1 year, 6 months ago

The first answer is NO. This code snippet is a skill definition. The outputs need to be mapped before they are available for indexing:
<https://learn.microsoft.com/en-us/azure/search/cognitive-search-concept-annotations-syntax>

upvoted 2 times

 **AzureJobsTillRetire** 1 year, 3 months ago

I think the first answer is YES. The outputs are available for indexing once mapped. If we answer no, it means that the outputs are not available for indexing, that that would be incorrect.

upvoted 2 times

 **AzureJobsTillRetire** 1 year, 3 months ago

It would also be incorrect to say that CompayDescription can be indexed. The wording can be better, but in concept it is correct to say that the output column is available for future indexing.

upvoted 1 times

 **nekkilodeon** 1 year, 9 months ago

I'd say the box1 is NO. The property is available for mapping to an index property.

upvoted 3 times

Question #7

You have the following data sources:

- Finance: On-premises Microsoft SQL Server database
- Sales: Azure Cosmos DB using the Core (SQL) API
- Logs: Azure Table storage

HR: Azure SQL database -

You need to ensure that you can search all the data by using the Azure Cognitive Search REST API.

What should you do?

- A. Configure multiple read replicas for the data in Sales.
- B. Mirror Finance to an Azure SQL database.
- C. Ingest the data in Logs into Azure Data Explorer.
- D. Ingest the data in Logs into Azure Sentinel.

Correct Answer: B

On-premises Microsoft SQL Server database cannot be used as an index data source.

Note: Indexer in Azure Cognitive Search: : Automate aspects of an indexing operation by configuring a data source and an indexer that you can schedule or run on demand. This feature is supported for a limited number of data source types on Azure.

Indexers crawl data stores on Azure.

- Azure Blob Storage
- Azure Data Lake Storage Gen2 (in preview)
- Azure Table Storage
- Azure Cosmos DB
- Azure SQL Database
- SQL Managed Instance
- SQL Server on Azure Virtual Machines

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-indexer-overview#supported-data-sources>

Community vote distribution

 B (100%)

Eltooth **Highly Voted** 1 year, 10 months ago

Selected Answer: B

B is correct answer : Mirror Finance to an Azure SQL database.

<https://docs.microsoft.com/en-us/azure/search/search-indexer-overview#supported-data-sources>

upvoted 5 times

reiwanotora **Most Recent** 2 days, 6 hours ago

Selected Answer: B

B is right answer.

upvoted 1 times

Gvalli 5 months, 1 week ago

A modified version of this was in the exam today.

upvoted 2 times

rdemontis 6 months, 3 weeks ago

Selected Answer: B

correct, a similar question was already present before

upvoted 2 times

slcheng 1 year, 5 months ago

Selected Answer: B

Agreed is B

upvoted 2 times

Question #8

You are developing a solution to generate a word cloud based on the reviews of a company's products.

Which Text Analytics REST API endpoint should you use?

- A. keyPhrases
- B. sentiment
- C. languages
- D. entities/recognition/general

Correct Answer: A

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/overview>

Community vote distribution



✉️ **MII1975** Highly Voted 2 years, 6 months ago

I also thought it was B, but now I think given answer is correct:

Extracted from here

<https://docs.microsoft.com/en-us/azure/cognitive-services/language-service/key-phrase-extraction/tutorials/integrate-power-bi#create-the-word-cloud>

The key phrases provide us with the important words from our customer comments, not just the most common words. Also, word sizing in the resulting cloud isn't skewed by the frequent use of a word in a relatively small number of comments.

upvoted 21 times

✉️ **TamHas** Highly Voted 2 years, 3 months ago

Selected Answer: A

correct

upvoted 8 times

✉️ **reiwanotora** Most Recent 2 days, 5 hours ago

Selected Answer: A

It must be A.

upvoted 1 times

✉️ **PCRamirez** 3 months, 2 weeks ago

According to Windows Copilot:

The best choice for this scenario is **A. keyPhrases**. This endpoint extracts the main points or topics from a text document, which can be used to generate a word cloud that shows the most frequent or relevant terms in the reviews. The other endpoints are not suitable for this task because:

- B. sentiment analyzes the positive, negative, or neutral tone of a text document, which is not relevant for generating a word cloud.
- C. languages detects the language of a text document, which is not relevant for generating a word cloud.
- D. entities/recognition/general identifies and categorizes named entities in a text document, such as person names, locations, organizations, etc.

This might not capture the main points or topics of the reviews, and might include irrelevant or sensitive information in the word cloud.

upvoted 1 times

✉️ **sismier** 6 months, 1 week ago

Selected Answer: B

I think the correct answer is B:

Sentiment Analysis and opinion mining:

Sentiment analysis and opinion mining are preconfigured features that help you find out what people think of your brand or topic by mining text for clues about positive or negative sentiment, and can associate them with specific aspects of the text.

<https://learn.microsoft.com/en-us/azure/ai-services/language-service/overview#sentiment-analysis-and-opinion-mining>

upvoted 1 times

✉️ **tdctdc** 5 months, 4 weeks ago

Word Cloud has nothing to do with sentiment analysis.

upvoted 1 times

✉️ **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

correct answer. Here is an example:

<https://learn.microsoft.com/en-us/azure/ai-services/language-service/key-phrase-extraction/tutorials/integrate-power-bi#create-the-word-cloud>

upvoted 1 times

zellck 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/key-phrase-extraction/overview>

Key phrase extraction is one of the features offered by Azure Cognitive Service for Language, a collection of machine learning and AI algorithms in the cloud for developing intelligent applications that involve written language. Use key phrase extraction to quickly identify the main concepts in text. For example, in the text "The food was delicious and the staff were wonderful.", key phrase extraction will return the main topics: "food" and "wonderful staff".

upvoted 2 times

zellck 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 3 times

Pixelmate 11 months ago

This was on exam 28/06

upvoted 3 times

kail85 12 months ago

A. keyPhrases

The keyPhrases endpoint extracts key phrases from the given text, which can be useful for creating a word cloud. Key phrases represent the main points or topics discussed in the text, making them ideal for visualizing the most prominent themes in the reviews.

upvoted 1 times

slcheng 1 year, 5 months ago

Selected Answer: A

Agreed with A

upvoted 1 times

ExamGuruBhai 1 year, 9 months ago

Selected Answer: A

correct

upvoted 1 times

Eltooth 1 year, 10 months ago

Selected Answer: A

A is correct answer.

upvoted 1 times

ManAtWorkAtNight 2 years ago

Answer is A. keyPhrases

Word Cloud is a set of most frequently appeared words(actually an image)

<https://docs.microsoft.com/en-us/rest/api/cognitiveservices-textanalytics/3.0/key-phrases/key-phrases>

upvoted 3 times

PHD_CHENG 2 years ago

Selected Answer: A

I support "A" as answer as question is asking to develop a "Word cloud based" solution by reviewing of Company's product. (initially, I was thinking about "B" as an answer but I changed my mind.)

upvoted 2 times

prabhjot 2 years, 3 months ago

Sentiment is correct one

upvoted 1 times

SuperPetey 2 years, 9 months ago

Correct Answer should be "B":sentiment analysis because it is analyzing customer reviews. See this blog post from MSFT that uses sentiment analysis to create a word cloud application for the same purpose: <https://devblogs.microsoft.com/premier-developer/creating-word-cloud-for-sentiment-analysis-with-azure-cognitive-services-text-analytics/>

upvoted 3 times

astralyt 2 years, 9 months ago

It doesn't say analyzing it. Word clouds comprise of key phrases. So I think the answer is right.

upvoted 21 times

Question #9

DRAG DROP -

You have a web app that uses Azure Cognitive Search.

When reviewing billing for the app, you discover much higher than expected charges. You suspect that the query key is compromised.

You need to prevent unauthorized access to the search endpoint and ensure that users only have read only access to the documents collection.

The solution must minimize app downtime.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions

- Add a new query key.
- Regenerate the secondary admin key.
- Change the app to use the secondary admin key.
- Change the app to use the new key.
- Regenerate the primary admin key.
- Delete the compromised key.

Answer Area**Correct Answer:****Actions**

- Regenerate the secondary admin key.
- Change the app to use the secondary admin key.
- Regenerate the primary admin key.

Answer Area

- | |
|------------------------------------|
| Add a new query key. |
| Change the app to use the new key. |
| Delete the compromised key. |



Reference:

<https://docs.microsoft.com/en-us/azure/search/search-security-api-keys>

✉ **guruaws2021** 2 years, 9 months ago

The answer is wrong
regenerated secondary key
update the app to use the secondary key
regenerated the primary key
upvoted 33 times

✉ **ziggy1117** 11 months, 3 weeks ago

the app only uses a query key not secondary or primary keys. The answer is CORRECT
upvoted 14 times

✉ **satishk4u** 2 years ago

No, it is about query key not admin key.
upvoted 12 times

✉ **czmiel24** 2 years, 9 months ago

Nope, the answer is correct. It's all about query keys which can be up to 50.
upvoted 49 times

✉ **AnonymousJhb** 6 months, 3 weeks ago

The answer is correct. By default, only 1 query key is created, but as per the answer, a new query key (up to 50) can be generated.
More can be created on demand by a search service administrator.
Query keys are not admin keys! Dont get caught!
<https://learn.microsoft.com/en-us/azure/search/search-security-api-keys?tabs=portal-use%2Cportal-find%2Cportal-query>

upvoted 8 times

✉️ **rdemontis** 6 months, 3 weeks ago

exactly, we don't have to confuse Admin Keys with Query keys. They are different things and with a different management
upvoted 5 times

✉️ **Adedoyin_Simeon** Highly Voted 2 years, 7 months ago

The answer is correct. The question specifically mentioned a read-only access to the documents which is achieved only through query keys. And it is a query key that was compromised. Any of the admin keys (primary and secondary) allows read-write operations. Hence, solution is to create a new query key, switch to the new query key and delete the compromised query key.

upvoted 21 times

✉️ **sl_mslconsulting** Most Recent 7 months, 2 weeks ago

Suspecting unauthorized access doesn't give you a reason to interrupt legit users. Also by looking the read only access requirement to the document collections already tells you this is about query keys.

upvoted 1 times

✉️ **sl_mslconsulting** 7 months ago

Two key requirements are read-only access to the document collection and minimum app downtime. Answers are obvious.

upvoted 1 times

✉️ **zellck** 11 months ago

1. Add new query key
2. Change app to use new key
3. Delete compromised key

<https://learn.microsoft.com/en-us/azure/search/search-security-api-keys?tabs=portal-use%2Cportal-find%2Cportal-query#create-query-keys>

upvoted 7 times

✉️ **Pixelmate** 11 months ago

ChatGPT answer : The answer is correct but in a different order.

1. Revoke the compromised query key immediately to prevent further unauthorized access.
2. Create a new query key for the application to use for authorized access.
3. Update the search endpoint and application to use the new query key and ensure that it is properly secured.

upvoted 2 times

✉️ **M25** 8 months, 4 weeks ago

F (Delete compromised), A (Add new / "recreate"), D (Change / update app to use new)

<https://learn.microsoft.com/en-us/rest/api/searchmanagement/2021-04-01-preview/query-keys/Delete>

Deletes the specified query key. Unlike admin keys, query keys are not regenerated. The process for regenerating a query key is to delete and then recreate it.

upvoted 1 times

✉️ **aiml** 1 year ago

The answer is correct. As its about query key and not for admin key.

upvoted 3 times

✉️ **Eltooth** 1 year, 10 months ago

Answer is correct.

Add a new query key

Change the app to use the new key

Delete the compromised key

upvoted 4 times

✉️ **mohamedba** 1 year, 11 months ago

1. Regenerate 2nd key
2. Change the app to use the regenerated 2nd key
3. Regenerate the 1st key

upvoted 1 times

✉️ **Eltooth** 1 year, 10 months ago

Wrong. There is no mention of primary or secondary key in this question.

The provided answer is correct:

Add a new query key

Change the app to use the new key

Delete the compromised key

upvoted 1 times

✉️ **reachmymind** 2 years, 2 months ago

Given Answer is Correct:

Add a new query key

Change the app to use the new key

Delete the compromised key

The key here is to not confuse "query Keys" with "admin keys"

<https://docs.microsoft.com/en-us/azure/search/search-security-api-keys>

upvoted 5 times

 **angie31** 2 years, 8 months ago

The answer is correct. In order to have read only access when you call the search API you need to use query key.

upvoted 2 times

 **GMKanon** 2 years, 8 months ago

Answer is correct since the question did not specify which of the keys was compromised, then both primary and secondary keys must be regenerated and no option for that.

upvoted 3 times

 **SnowCheetah** 2 years, 8 months ago

The steps is for when ADMIN key is compromised. If it is only search API key compromised, the answer should be enough.

upvoted 2 times

Question #10

You are developing an application that will use Azure Cognitive Search for internal documents.

You need to implement document-level filtering for Azure Cognitive Search.

Which three actions should you include in the solution? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Send Azure AD access tokens with the search request.
- B. Retrieve all the groups.
- C. Retrieve the group memberships of the user.
- D. Add allowed groups to each index entry.
- E. Create one index per group.
- F. Supply the groups as a filter for the search requests.

Correct Answer: CDF

Your documents must include a field specifying which groups have access. This information becomes the filter criteria against which documents are selected or rejected from the result set returned to the issuer.

D: A query request targets the documents collection of a single index on a search service.

CF: In order to trim documents based on group_ids access, you should issue a search query with a group_ids/any(g:search.in(g, 'group_id1, group_id2,...')) filter, where 'group_id1, group_id2, ...' are the groups to which the search request issuer belongs.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-security-trimming-for-azure-search>

Community vote distribution

CDF (71%)

14%

14%

 **ninja**  1 year, 9 months ago

Selected Answer: CDF

D: Add allowed groups to each index entry.

Your documents must include a field specifying which groups have access.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-security-trimming-for-azure-search#create-security-field>

C: You need to get the membership of the user

F: Supply the groups as a filter for the search requests.

"In order to trim documents based on group_ids access, you should issue a search query with a group_ids/any(g:search.in(g, 'group_id1, group_id2,...')) filter, where 'group_id1, group_id2, ...' are the groups to which the search ##request issuer belongs##."

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-security-trimming-for-azure-search>

upvoted 10 times

 **anto69**  1 month, 2 weeks ago

The possible answers make no sense to me

upvoted 1 times

 **Murtuza** 2 months, 1 week ago

Selected Answer: CDF

C, D and F are the correct answers

upvoted 2 times

 **evangelist** 3 months, 3 weeks ago

The given answer CDF is CORRECT!

C. Retrieve the group memberships of the user. To understand which documents the user is allowed to access.

D. Add allowed groups to each index entry. To tag documents with the groups that can access them.

F. Supply the groups as a filter for the search requests. To ensure that search results are filtered according to the user's group memberships.

upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: CDF

The answer is correct as well as the explanation/reference provided

upvoted 1 times

 **sl_mslconsulting** 7 months, 2 weeks ago

A is wrong. Cognitive Search wouldn't know how to handle the token handed to it. All you app does is to come up with the right filter and send it to the Cognitive Search service via a request. You also need to index the documents with the group info so later the Cognitive Service can do a proper filtering on them.

upvoted 1 times

✉️ **sl_msiconsulting** 7 months, 2 weeks ago

C D F are my choices

upvoted 1 times

✉️ **zellck** 11 months ago

Selected Answer: CDF

CDF is the answer.

<https://learn.microsoft.com/en-us/azure/search/search-security-trimming-for-azure-search>

Cognitive Search doesn't provide document-level permissions and can't vary search results from within the same index by user permissions. As a workaround, you can create a filter that trims search results based on a string containing a group or user identity.

<https://learn.microsoft.com/en-us/azure/search/search-security-trimming-for-azure-search#apply-the-security-filter-in-the-query>

In order to trim documents based on group_ids access, you should issue a search query with a group_ids/any(g:search.in(g, 'group_id1, group_id2,...')) filter, where 'group_id1, group_id2,...' are the groups to which the search request issuer belongs.

upvoted 2 times

✉️ **M25** 8 months, 4 weeks ago

Right at the bottom of the page:

For an alternative pattern based on Azure Active Directory, or to revisit other security features, see the following links.

- Security filters for trimming results using Active Directory identities
- Security in Azure Cognitive Search

upvoted 1 times

✉️ **marti_tremblay000** 1 year, 2 months ago

Selected Answer: ACF

According to ChatGPT, the correct answers are ACF :

To implement document-level filtering for Azure Cognitive Search, the following three actions should be included in the solution:

A. Send Azure AD access tokens with the search request: This will allow Azure Cognitive Search to authenticate the user and retrieve their group memberships.

C. Retrieve the group memberships of the user: This information is necessary to determine which documents the user is allowed to access.

F. Supply the groups as a filter for the search requests: This will allow you to filter the search results based on the user's group memberships.

Therefore, the correct answers are A, C, and F.

upvoted 3 times

✉️ **M25** 8 months, 4 weeks ago

A, C, F

<https://learn.microsoft.com/en-us/azure/search/search-howto-aad?tabs=config-svc-portal%2Caad-dotnet>

A key benefit of using Azure AD is that your credentials and API keys no longer need to be stored in your code. Azure AD authenticates the security principal (a user, group, or service) running the application. If authentication succeeds, Azure AD returns the access token to the application, and the application can then use the access token to authorize requests to Azure Cognitive Search.

<https://learn.microsoft.com/en-us/azure/search/search-security-trimming-for-azure-search-with-aad>

This article demonstrates how to use Azure Active Directory (AD) security identities together with filters in Azure Cognitive Search to trim search results based on user group membership.

upvoted 1 times

✉️ **not_a_robot** 1 year, 9 months ago

Selected Answer: BDF

Shouldn't the first step be getting all the groups? As when you create a security field, you'll need the group ids. I'm not certain what the group membership of the user means.

upvoted 3 times

✉️ **mk1967** 1 year, 8 months ago

I guess this is the user using the application, therefore we must get groups to which he belongs.

upvoted 2 times

Question #11

You have an Azure Cognitive Search solution and an enrichment pipeline that performs Sentiment Analysis on social media posts.

You need to define a knowledge store that will include the social media posts and the Sentiment Analysis results.

Which two fields should you include in the definition? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. storageContainer
- B. storageConnectionString
- C. files
- D. tables
- E. objects

Correct Answer: DE

Knowledge store definition -

A knowledge store is defined inside a skillset definition and it has two components:

A connection string to Azure Storage

Projections that determine whether the knowledge store consists of tables, objects or files.

The projections element is an array. You can create multiple sets of table-object-file combinations within one knowledge store.

```
"knowledgeStore": {  
  "storageConnectionString": "<YOUR-AZURE-STORAGE-ACCOUNT-CONNECTION-STRING>",  
  "projections": [  
    {  
      "tables": [],  
      "objects": [],  
      "files": []  
    }  
  ]  
}
```

The type of projection you specify in this structure determines the type of storage used by knowledge store.

Objects - project JSON document into Blob storage. The physical representation of an object is a hierarchical JSON structure that represents an enriched document.

Tables - project enriched content into Table Storage. Define a table projection when you need tabular reporting structures for inputs to analytical tools or export as data frames to other data stores. You can specify multiple tables within the same projection group to get a subset or cross section of enriched documents. Within the same projection group, table relationships are preserved so that you can work with all of them.

Projected content is not aggregated or normalized. The following screenshot shows a table, sorted by key phrase, with the parent document indicated in the adjacent column. In contrast with data ingestion during indexing, there is no linguistic analysis or aggregation of content. Plural forms and differences in casing are considered unique instances.

Content.metadata_storage_name	Content.KeyPhrases
Cognitive Services and Content Intelligence.pptx	Computer Vision
10-K-FY16.html	computing device
10-K-FY16.html	computing devices
MSFT_FY17_10K.docx	computing devices
10-K-FY16.html	Computing segment
Cognitive Services and Bots (spanish).pdf	confianza

Incorrect:

Not C: files - project image files into Blob storage. A file is an image extracted from a document, transferred intact to Blob storage. Although it is named "files", it shows up in Blob Storage, not file storage.

Community vote distribution

BE (66%)

BD (17%)

Other

zellck 11 months ago

Selected Answer: BE

Typo, BE is the answer.

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-concept-intro?tabs=portal#knowledge-store-definition>

A knowledge store is defined inside a skillset definition and it has two components:

- A connection string to Azure Storage
- Projections that determine whether the knowledge store consists of tables, objects or files. The projections element is an array. You can create multiple sets of table-object-file combinations within one knowledge store.

upvoted 9 times

✉ **zelick** 11 months ago

The type of projection you specify in this structure determines the type of storage used by knowledge store, but not its structure. Fields in tables, objects, and files are determined by Shaper skill output if you're creating the knowledge store programmatically, or by the Import data wizard if you're using the portal.

- objects

project JSON document into Blob storage. The physical representation of an object is a hierarchical JSON structure that represents an enriched document.

upvoted 4 times

✉ **rdemontis** 6 months, 3 weeks ago

Agree with you, thanks for explanation

upvoted 1 times

✉ **reiwanotora** **Most Recent** 2 days, 6 hours ago

Selected Answer: BE

storageConnectionString and Object.

upvoted 1 times

✉ **evangelist** 3 months, 3 weeks ago

Selected Answer: DE

For the need to define a knowledge store that includes social media posts and the Sentiment Analysis results, the most relevant fields/types to include in your definition would likely be:

D. tables: For storing structured data such as the original posts along with their sentiment scores in a structured, queryable format.

E. objects: For preserving the full, enriched data set in its original format along with nested Sentiment Analysis results.

upvoted 3 times

✉ **dimsok** 4 months, 2 weeks ago

Selected Answer: BD

Enriched content is projected to tables:

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-concept-intro?tabs=portal#knowledge-store-definition>

upvoted 1 times

✉ **Gvalli** 5 months, 1 week ago

A modified version of this was in the exam today.

upvoted 2 times

✉ **rdemontis** 6 months, 3 weeks ago

Selected Answer: BE

To me correct answer is BE

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-concept-intro?tabs=portal#knowledge-store-definition>

upvoted 1 times

✉ **sl_mslconsulting** 7 months, 2 weeks ago

Selected Answer: BD

I would choose table projection over object as it will be easier to analyze the data later. A full JSON representation of your data and enrichments in one JSON document, in this case, is not doing you any good other than more work required for later analysis.

upvoted 2 times

✉ **sl_mslconsulting** 7 months, 2 weeks ago

Why would you put social media posts and the Sentiment Analysis results in one JSON document to make your later work more difficult?

upvoted 1 times

✉ **propanther** 7 months ago

Behind the scenes, table projection stored data in Azure Table Storage, where scalability and performance of the storage becomes vital to deal with when unstructured data like social media post is stored. By design, table storages should not be used to store objects which can be very big in size and can be extremely difficult to apply search pattern on. Check this out:

The scalability and performance targets listed here are high-end targets, but are achievable. In all cases, the request rate and bandwidth achieved by your storage account depends upon the size of objects stored, the access patterns utilized, and the type of workload your application performs.

ref: <https://learn.microsoft.com/en-us/azure/storage/tables/scalability-targets>

upvoted 1 times

✉️ **propanther** 7 months ago

<https://learn.microsoft.com/en-us/azure/storage/tables/scalability-targets>
upvoted 1 times

✉️ **zellck** 11 months ago

Selected Answer: BD

BD is the answer.

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-concept-intro?tabs=portal#knowledge-store-definition>
A knowledge store is defined inside a skillset definition and it has two components:

- A connection string to Azure Storage
- Projections that determine whether the knowledge store consists of tables, objects or files. The projections element is an array. You can create multiple sets of table-object-file combinations within one knowledge store.

upvoted 3 times

✉️ **zellck** 11 months ago

The type of projection you specify in this structure determines the type of storage used by knowledge store, but not its structure. Fields in tables, objects, and files are determined by Shaper skill output if you're creating the knowledge store programmatically, or by the Import data wizard if you're using the portal.

- objects
project JSON document into Blob storage. The physical representation of an object is a hierarchical JSON structure that represents an enriched document.

upvoted 2 times

✉️ **EliteAllen** 1 year ago

Selected Answer: AB

To define a knowledge store in Azure Cognitive Search that includes both the social media posts and the Sentiment Analysis results, you should include the following two fields in the definition:

A. storageContainer: This field specifies the Azure Blob Storage container where the enriched data will be stored. You need to define a storage container to hold the processed data from the enrichment pipeline.

B. storageConnectionString: This field provides the connection string to the Azure Storage account that will be used to store the enriched data. The connection string is necessary for Azure Cognitive Search to connect and store the data in the specified storage account.

Both the storageContainer and storageConnectionString fields are required to set up a knowledge store in Azure Cognitive Search that can store the social media posts and the Sentiment Analysis results.

upvoted 3 times

✉️ **MaliSanFuu** 1 year ago

I agree, as the task is to define a knowledge store and not the potential projections that define the format in which the enriched data is saved
upvoted 1 times

✉️ **claps92** 1 year, 3 months ago

why not "file"??
upvoted 1 times

✉️ **ap1234pa** 1 year, 4 months ago

Selected Answer: BE

BE is the answer
upvoted 4 times

✉️ **AzureJobsTillRetire** 1 year, 3 months ago

I agree. There seems to not have a need for tables.
upvoted 1 times

✉️ **SSJA** 1 year, 5 months ago

Selected Answer: BE

Correct Answer is B & E
upvoted 3 times

✉️ **marti_tremblay000** 1 year, 5 months ago

Selected Answer: BE

it should be BE. storageConnectionString and Objects must be included in the definition.
upvoted 2 times

✉️ **halfway** 1 year, 6 months ago

Selected Answer: BE

storageConnectionString and Tables. The projection types are mutually exclusive. There is no need to have 2 projections in this scenario.
upvoted 3 times

✉️ **AzureJobsTillRetire** 1 year, 3 months ago

This is incorrect.
Example is given as below
"knowledgeStore" : {

```
"storageConnectionString": "DefaultEndpointsProtocol=https;AccountName=<Acct Name>;AccountKey=<Acct Key>;",
"projections": [
{
"tables": [
{ "tableName": "ks-museums-main", "generatedKeyName": "ID", "source": "/document/tableprojection" },
{ "tableName": "ks-museumEntities", "generatedKeyName": "ID", "source": "/document/tableprojection/Entities/*" }
],
"objects": [
{ "storageContainer": "ks-museums", "generatedKeyName": "ID", "source": "/document/objectprojection" }
],
"files": [ ]
}
]
https://learn.microsoft.com/en-us/azure/search/knowledge-store-projection-overview
upvoted 1 times
```

✉️ **mk1967** 1 year, 8 months ago

Shouldn't it be BE? We need storageConnectionString (as seen in the original answer).

upvoted 2 times

✉️ **GigaCaster** 1 year, 7 months ago

No, the answer is correct as it wants the places it is stored and not how to get to the stored knowledge

upvoted 1 times

Question #12

SIMULATION -

Use the following login credentials as needed:

To enter your username, place your cursor in the Sign in box and click on the username below.

To enter your password, place your cursor in the Enter password box and click on the password below.

Azure Username: admin@abc.com -

Azure Password: XXXXXXXXXXXX -

The following information is for technical support purposes only:

Lab Instance: 12345678 -

Task -

You need to create an Azure resource named solution12345678 that will index a sample database named realestate-us-sample. The solution must ensure that users can search the index in English for people, organizations, and locations.

To complete this task, sign in to the Azure portal.

Correct Answer: See explanation below.

Step 1 - Start the Import data wizard and create a data source

1. Sign in to the Azure portal with your Azure account.

2. Find your search service and on the Overview page, click Import data on the command bar to create and populate a search index.

The screenshot shows the 'Import data' wizard in the Azure portal. At the top, there are tabs: 'Connect to your data' (which is underlined), 'Enrich content (Optional)', 'Customize target index', and 'Create an indexer'. Below these tabs, a note says: 'Create and load a search index using data from an existing Azure data source in your current subscription. Azure Cognitive Search crawls the data structure you provide, extracts searchable content, optionally enriches it with cognitive skills, and loads it into an index.' A 'Learn more' link is provided. The 'Data Source' section shows a dropdown menu with 'Samples' selected, circled with a red number 1. The 'Name' section shows two entries: 'realestate-us-sample' and 'hotels-sample', with 'realestate-us-sample' circled with a red number 2.

3. In the wizard, click Connect to your data, and select the sample database named realestate-us-sample

Step 2 - Skip the "Enrich content" page

The wizard supports the creation of an AI enrichment pipeline for incorporating the Cognitive Services AI algorithms into indexing.

We'll skip this step for now, and move directly on to Customize target index.

Step 3 - Configure index -

The solution must ensure that users can search the index in English for people, organizations, and locations.

Configure Searchable for the fields people, organizations, and locations.

Home > Microsoft.Search - Overview > my-new-search-service > Import data

Import data

Connect to your data Enrich content (Optional) **Customize target index *** Create an indexer

Info We provided a default index for you. You can delete the fields you don't need. Everything is editable, but once the index is built, deleting or changing existing fields will require re-indexing your documents.

Index name * hotels-sample-index

Key * HotelId

Suggerer name sg **Search mode**

Add field **Add subfield** **Delete**

Field name	Type	Retrievable	Filterable	Sortable	Facetable	Searchable	Analyzer	Suggerer
HotelId	Edm.String	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	English - Microsoft	...
HotelName	Edm.String	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	English - Microsoft	...
Description	Edm.String	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	English - Microsoft	...

Reference:
<https://docs.microsoft.com/en-us/azure/search/search-get-started-portal>

✉ **reiwanotora** 2 days, 6 hours ago

Simulation questions will not appear on the actual exam as of May 26, 2024.

upvoted 1 times

✉ **GHill1982** 2 months, 2 weeks ago

As of March 2024:

1. Create an AI Search resource under Azure AI services.
 2. On the Overview page select Import data.
 3. Select Samples from the Data Source drop-down box and the realestate-us-sample.
 4. On the Customize target index page ensure Searchable is checked and the Analyzer language is English for the relevant fields.
- upvoted 1 times

✉ **chandiochan** 3 months, 1 week ago

Do we get simulation question as of Feb 2024?

upvoted 1 times

✉ **rdemontis** 6 months, 3 weeks ago

1. Create the Cognitive Search Result
 2. Import data
 3. On Connect to your data choose Samples --> realestate_us_sample
 4. Add enrichments
 5. select Extract people, organization and location names
 6. Select English Microsoft for each one
 7. Create an indexer
- upvoted 4 times

Question #13

HOTSPOT

You create a knowledge store for Azure Cognitive Search by using the following JSON.

```
"knowledgeStore" : {
    "storageConnectionString": "DefaultEndpointsProtocol=https;AccountName=<Acct Name>;AccountKey=<Acct Key>;",
    "projections": [
        {
            "tables": [
                {
                    "tableName": "unrelatedDocument",
                    "generatedKeyName": "Documentid",
                    "source": "/document/pbiShape"
                },
                {
                    "tableName": "unrelatedKeyPhrases",
                    "generatedKeyName": "KeyPhraseid",
                    "source": "/document/pbiShape/keyPhrases"
                }
            ],
            "objects": [
                ],
                "files": []
            },
            {
                "tables": [],
                "objects": [
                    {
                        "storageContainer": "unrelatedocrtext",
                        "source": null,
                        "sourceContext": "/document/normalized_images/*/text",
                        "inputs": [
                            {
                                "name": "ocrText",
                                "source": "/document/normalized_images/*/text"
                            }
                        ]
                    },
                    {
                        "storageContainer": "unrelatedocrlayout",
                        "source": null,
                        "sourceContext": "/document/normalized_images/*/layoutText",
                        "inputs": [
                            {
                                "name": "ocrLayoutText",
                                "source": "/document/normalized_images/*/layoutText"
                            }
                        ]
                    }
                ],
                "files": []
            }
        ]
    }
}
```

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Answer Area

There will be [answer choice].

- no projection groups
- one projection group
- two projection groups
- four projection groups

Normalized images will [answer choice].

- not be projected
- be projected to Azure Blob storage
- be projected to Azure File storage
- be saved to an Azure Table storage

Answer Area

Correct Answer:

There will be [answer choice].

no projection groups
one projection group
two projection groups
four projection groups

Normalized images will [answer choice].

not be projected
be projected to Azure Blob storage
be projected to Azure File storage
be saved to an Azure Table storage

✉  **sl_msiconsulting**  7 months, 1 week ago

The answer for the second question is wrong. There will be no projection for normalized images. You have to remember object projection is for JSON document not binary raw images. Both projections don't have the files specified.

upvoted 16 times

✉  **dimsok** 4 months, 3 weeks ago

I agree with you, for images you should define a file projection but there is none (<https://learn.microsoft.com/en-us/azure/search/knowledge-store-projections-examples#define-a-table-projection>)

upvoted 1 times

✉  **Tin_Tin**  11 months, 1 week ago

The answer seems correct.

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-projections-examples#define-a-table-projection>

upvoted 10 times

✉  **M25** 8 months, 4 weeks ago

Correct!

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-projection-example-long#relationships-among-table-object-and-file-projections>

If you don't want the data related, define the projections in different projection groups. For example, the following snippet will result in the tables being related, but without relationships between the tables and the object (OCR text) projections.

<https://learn.microsoft.com/en-us/azure/search/knowledge-store-projections-examples#define-a-file-projection>

File projections are always binary, normalized images, where normalization refers to potential resizing and rotation for use in skillset execution.

File projections, similar to object projections, are created as blobs in Azure Storage, and contain binary data (as opposed to JSON).

upvoted 6 times

✉  **rdemontis** 6 months, 3 weeks ago

correct. thanks for your contribution

upvoted 2 times

✉  **Murtuza**  1 month, 3 weeks ago

The code you provided contains two projection groups. Here they are:

The first projection group includes the following tables:

unrelatedDocument

unrelatedKeyPhrases

The second projection group includes the following objects:

unrelatedocrtext

unrelatedocrlayout

upvoted 1 times

✉  **Murtuza** 1 month, 3 weeks ago

The correct answer is one projection group. In the provided JSON, there is one projections array, which means there is one projection group.

The correct answer is be projected to Azure Blob storage. The storageContainer fields "unrelatedocrtext" and "unrelatedocrlayout" indicate that the normalized images (both text and layout text) are being stored in Azure Blob storage containers. Azure Table storage is not used in this scenario.

upvoted 1 times

✉  **85e4c91** 3 months ago

The Answer is Correct, the normalized Images are defined in the object projection and object projections are stored in blob storage

upvoted 1 times

Question #14

You plan to create an index for an Azure Cognitive Search service by using the Azure portal. The Cognitive Search service will connect to an Azure SQL database.

The Azure SQL database contains a table named UserMessages. Each row in UserMessages has a field named MessageCopy that contains the text of social media messages sent by a user.

Users will perform full text searches against the MessageCopy field, and the values of the field will be shown to the users.

You need to configure the properties of the index for the MessageCopy field to support the solution.

Which attributes should you enable for the field?

- A. Sortable and Retrievable
- B. Filterable and Retrievable
- C. Searchable and Facetable
- D. Searchable and Retrievable

Correct Answer: D

Community vote distribution

D (100%)

zellick Highly Voted 11 months ago

Selected Answer: D

D is the answer.

<https://learn.microsoft.com/en-us/rest/api/searchservice/create-index#-field-definitions--retrievable>

Indicates whether the field can be returned in a search result.

- searchable

Indicates whether the field is full-text searchable and can be referenced in search queries.

upvoted 10 times

reiwanotora Most Recent 2 days, 6 hours ago

Selected Answer: D

D is the right answer.

upvoted 1 times

evangelist 3 months, 3 weeks ago

Selected Answer: D

only option D meets the demand

upvoted 1 times

rdemontis 6 months, 3 weeks ago

Selected Answer: D

correct

upvoted 2 times

MaliSanFuu 1 year ago

Selected Answer: D

Should be correct.

reference: <https://learn.microsoft.com/en-us/azure/search/search-what-is-an-index>

upvoted 2 times

Question #15

You have the following data sources:

- Finance: On-premises Microsoft SQL Server database
- Sales: Azure Cosmos DB using the Core (SQL) API
- Logs: Azure Table storage
- HR: Azure SQL database

You need to ensure that you can search all the data by using the Azure Cognitive Search REST API.

What should you do?

- A. Export the data in Finance to Azure Data Lake Storage.
- B. Configure multiple read replicas for the data in Sales.
- C. Ingest the data in Logs into Azure Data Explorer.
- D. Migrate the data in HR to Azure Blob storage.

Correct Answer: A

Community vote distribution

A (91%)

9%

✉ **reiwanotora** 2 days, 6 hours ago

Selected Answer: A

onpre -> cloud first.

upvoted 1 times

✉ **evangelist** 3 months, 3 weeks ago

Selected Answer: A

The same question has appeared multiple times in the exam, on-premise data has to move to the Cloud first before it can be made searchable on Cloud.

upvoted 2 times

✉ **Florin83** 3 months, 4 weeks ago

Selected Answer: A

All data sources except A are supported:

Supported data sources

Indexers crawl data stores on Azure and outside of Azure.

Azure Blob Storage

Azure Cosmos DB

Azure Data Lake Storage Gen2

Azure SQL Database

Azure Table Storage

Azure SQL Managed Instance

SQL Server on Azure Virtual Machines

Azure Files (in preview)

Azure MySQL (in preview)

SharePoint in Microsoft 365 (in preview)

Azure Cosmos DB for MongoDB (in preview)

Azure Cosmos DB for Apache Gremlin (in preview)

upvoted 2 times

✉ **Gvalli** 5 months, 1 week ago

A modified version of this was in the exam today.

upvoted 2 times

✉ **sismar** 6 months ago

Selected Answer: A

A is the correct

upvoted 2 times

✉ **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

the answer is correct. repeated question

<https://learn.microsoft.com/en-us/azure/search/search-indexer-overview#supported-data-sources>

upvoted 3 times

 **_LAW_** 7 months ago

Selected Answer: C

chat gpt

upvoted 1 times

Question #16

HOTSPOT

You plan to provision Azure Cognitive Services resources by using the following method.

You need to create a Standard tier resource that will convert scanned receipts into text.

```
static void provision_resource(CognitiveServicesManagementClient client, string name, string kind, string tier, string location)
{
    CognitiveServicesAccount parameters =
        new CognitiveServicesAccount(null, null, kind, location, name,
            new CognitiveServicesAccountProperties(), new Sku(tier));
    result = client.Accounts.Create(resource_group_name, tier, parameters);
}
```

How should you call the method? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

provision_resource("res1",

ComputerVision
CustomVision.Prediction
CustomVision.Training
FormRecognizer

"eastus", "S1")
"useast", "S1")
"S0", "eastus")
"S0", "useast")

Answer Area

Correct Answer:

provision_resource("res1",

ComputerVision
CustomVision.Prediction
CustomVision.Training
FormRecognizer

"eastus", "S1")
"useast", "S1")
"S0", "eastus")
"S0", "useast")

rewanotora 2 days, 6 hours ago

The current name for Form Recognizer (as of May 2024) is Document Intelligence.

1. Document Intelligence(Form Recognizer)
2. "S0", "eastus"

<https://azure.microsoft.com/en-us/pricing/details/ai-document-intelligence/>

upvoted 1 times

sergbs 1 month, 1 week ago

1. FormRecognizer -> Document Intelligences service.
2. S0, "eastus"

upvoted 2 times

Murtuza 1 month, 3 weeks ago

Forms recognizer comes in 2 tiers one is F0 and another is S0. Apologizes on listing out S1 Tiers as it doesn't exist with Document Intelligences service. Sorry for the confusion

upvoted 1 times

Murtuza 2 months ago

Tier: "S1"
Location: "eastus"
upvoted 1 times

✉  **Murtuza** 2 months ago

To create a Standard tier resource that converts scanned receipts into text using Form Recognizer, you should call the provision_resource method with the following parameters:

Resource Name: "res1"

Kind: FormRecognizer

upvoted 1 times

✉  **Murtuza** 2 months ago

Document Intelligence (formerly known as Form Recognizer) is an Azure AI service that specializes in extracting structured data from unstructured documents, including receipts.

upvoted 1 times

Question #17

HOTSPOT

You have an app named App1 that uses Azure AI Document Intelligence to analyze medical records and provide pharmaceutical dosage recommendations for patients.

You send a request to App1 and receive the following response.

```
{
  "status": "succeeded",
  "createdDateTime": "2023-09-14T21:01:02Z",
  "lastUpdatedDateTime": "2023-09-14T21:01:03Z",
  "analyzeResult": {
    "apiVersion": "2023-07-31",
    "modelId": "prebuilt-healthInsuranceCard.us",
    "stringIndexType": "utf16CodeUnit",
    "content": "Blood Pressure 118/72",
    "pages": [
      {
        ...
        "words": [
          {
            "content": "Blood",
            "polygon": [ ... ],
            "confidence": 0.766,
            "span": { ... }
          },
          {
            "content": "Pressure",
            "polygon": [ ... ],
            "confidence": 0.716,
            "span": { ... }
          },
          {
            "content": "118/72",
            "polygon": [ ... ],
            "confidence": 0.761,
            "span": { ... }
          }
        ],
        ...
      }
    ],
    "documents": [
      {
        "docType": "healthInsuranceCard.us",
        "boundingRegions": [ ... ]
      }
    ],
    "fields": {},
    "confidence": 1,
    "spans": [ ... ]
  }
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
The chosen model is suitable for the intended use case.	<input type="radio"/>	<input type="radio"/>
The text content was recognized with greater than 70 percent confidence.	<input type="radio"/>	<input type="radio"/>
The form elements were recognized with greater than 70 percent confidence.	<input type="radio"/>	<input type="radio"/>

Answer Area

Correct Answer:

Statements	Yes	No
The chosen model is suitable for the intended use case.	<input type="radio"/>	<input checked="" type="radio"/>
The text content was recognized with greater than 70 percent confidence.	<input checked="" type="radio"/>	<input type="radio"/>
The form elements were recognized with greater than 70 percent confidence.	<input checked="" type="radio"/>	<input type="radio"/>

 **reiwanotora** 2 days, 6 hours ago

No

Yes

Yes

upvoted 1 times

 **Murtuza** 1 month, 3 weeks ago

The form elements were recognized with greater than 70 percent confidence: No. The fields object in the response is empty, which suggests that no form elements were recognized. Therefore, we cannot say that they were recognized with greater than 70 percent confidence.

upvoted 3 times

 **Murtuza** 1 month, 3 weeks ago

The chosen model is suitable for the intended use case: No. The model used here is "prebuilt-health InsuranceCard.us", which is designed to extract information from health insurance cards in the US. However, the intended use case is to analyze medical records and provide pharmaceutical dosage recommendations for patients. A more suitable model would be one specifically trained for medical record analysis.

The text content was recognized with greater than 70 percent confidence: Yes. The confidence scores for the recognized words "Blood" (0.766), "Pressure" (0.716), and "118/72" (0.761) are all greater than 70 percent.

upvoted 2 times

 **[Removed]** 2 months, 2 weeks ago

Answer is correct, i drop my pants

upvoted 4 times

Question #18

HOTSPOT

You have an Azure subscription that contains an Azure AI Document Intelligence resource named DI1.

You build an app named App1 that analyzes PDF files for handwritten content by using DI1.

You need to ensure that App1 will recognize the handwritten content.

How should you complete the code? To answer, select the appropriate options in the answer area.

Answer Area

```
Uri fileUri = new Uri("<fileUri>");

AnalyzeDocumentOperation operation = await client.AnalyzeDocumentFromUriAsync(WaitUntil.Completed,
    AnalyzeResult result = operation.Value;
    foreach (DocumentStyle style in result.Styles)
    {
        bool isHandwritten = style.IsHandwritten.HasValue && style.IsHandwritten == true;
        if (isHandwritten && style.Confidence >  )
        {
            Console.WriteLine($"Handwritten content found:");
            foreach (DocumentSpan span in style.Spans)

```

"prebuilt-document"
 "prebuilt-contract"
 "prebuilt-read"

Answer Area

```
Uri fileUri = new Uri("<fileUri>");

AnalyzeDocumentOperation operation = await client.AnalyzeDocumentFromUriAsync(WaitUntil.Completed,
    AnalyzeResult result = operation.Value;
    foreach (DocumentStyle style in result.Styles)
    {
        bool isHandwritten = style.IsHandwritten.HasValue && style.IsHandwritten == true;
        if (isHandwritten && style.Confidence >  )
        {
            Console.WriteLine($"Handwritten content found:");
            foreach (DocumentSpan span in style.Spans)

```

"prebuilt-document"
 "prebuilt-contract"
 "prebuilt-read"

Correct Answer:

```
1
bool isHandwritten = style.IsHandwritten.HasValue && style.IsHandwritten == true;
if (isHandwritten && style.Confidence >  )
{
    Console.WriteLine($"Handwritten content found:");
    foreach (DocumentSpan span in style.Spans)
```

 Mehe323 Highly Voted  2 months ago

The first answer should be read:
Read OCR - Extract print and handwritten text including words, locations, and detected languages.

<https://learn.microsoft.com/en-us/azure/ai-services/document-intelligence/concept-model-overview?view=doc-intel-4.0.0>
upvoted 5 times

 reiwanotora Most Recent  2 days, 6 hours ago

1. "prebuilt-document"
2. 0.75
upvoted 1 times

 Murtuza 1 month, 2 weeks ago

The given answers by exam topics is CORRECT
upvoted 2 times

 AlviraTony 2 months ago

Prebuilt-read is correct which can classify the test extracted as handwritten or printed
upvoted 2 times

 Murtuza 2 months ago

It appears that the code snippet you've provided is in C#. Let's break it down

You're iterating through the DocumentStyle objects within the result.Styles

You're invoking an AnalyzeDocumentOperation from an AnalyzeDocumentFromUriAsync method call. The operation is awaited until completion

For each style, you're checking if it's handwritten (based on the IsHandwritten property) and if its confidence level is greater than 0.1
upvoted 1 times

✉️ **audlindr** 2 months, 4 weeks ago

Reference: <https://learn.microsoft.com/en-us/dotnet/api/overview/azure/ai.formrecognizer-readme?view=azure-dotnet#use-the-prebuilt-general-document-model>

upvoted 3 times

✉️ **Mehe323** 2 months ago

It is read:

Read OCR - Extract print and handwritten text including words, locations, and detected languages.

<https://learn.microsoft.com/en-us/azure/ai-services/document-intelligence/concept-read?view=doc-intel-4.0.0>

upvoted 2 times

✉️ **Harry300** 2 months, 3 weeks ago

why document and not the read model? the code appears in both examples on that website

upvoted 6 times

Question #19

Topic 4

You have an app named App1 that uses a custom Azure AI Document Intelligence model to recognize contract documents.

You need to ensure that the model supports an additional contract format. The solution must minimize development effort.

What should you do?

- A. Lower the confidence score threshold of App1.
- B. Create a new training set and add the additional contract format to the new training set. Create and train a new custom model.
- C. Add the additional contract format to the existing training set. Retrain the model.
- D. Lower the accuracy threshold of App1.

Correct Answer: C

Community vote distribution

C (100%)

✉️ **reiwanotora** 2 days, 6 hours ago

Selected Answer: C

It must be C.

upvoted 1 times

✉️ **Murtuza** 1 month, 2 weeks ago

Selected Answer: C

Given answer C is correct

upvoted 2 times

✉️ **Murtuza** 2 months ago

Add the additional contract format to the existing training set and retrain the model: This is a more efficient solution. By augmenting the existing training data with the new format, you can fine-tune the model without starting from scratch. Retraining the model with the updated dataset should help it adapt to the additional format

upvoted 1 times

✉️ **[Removed]** 2 months, 2 weeks ago

I C U P too.

upvoted 1 times

Question #20

HOTSPOT

You have an Azure subscription.

You need to deploy an Azure AI Document Intelligence resource.

How should you complete the Azure Resource Manager (ARM) template? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
"$schema": "https://schema.management.azure.com/schemas/2019-04-01/deploymentTemplate.json#",
"contentVersion": "1.0.0.0",
"parameters": {},
"variables": {},
"resources": [
{
  "type": "/accounts",
  "apiVersion": "2023-05-01",
  "name": "DocumentIntelligenceDemo",
  "location": "westeurope",
  "sku": {
    "name": "F0"
  },
  "kind": "CognitiveSearch",
  "properties": {}
}
]
```

Answer Area

```
"$schema": "https://schema.management.azure.com/schemas/2019-04-01/deploymentTemplate.json#",
"contentVersion": "1.0.0.0",
"parameters": {},
"variables": {},
"resources": [
{
  "type": "Microsoft.CognitiveServices/accounts",
  "name": "DocumentIntelligenceDemo",
  "location": "westeurope",
  "sku": {
    "name": "F0"
  },
  "kind": "FormRecognizer",
  "apiVersion": "2023-05-01"
}
]
```

Correct Answer:

"apiVersion": "2023-05-01",

"name": "DocumentIntelligenceDemo",

"location": "westeurope",

"sku": {

 "name": "F0"

},

"kind":

AiBuilder,

CognitiveSearch,

FormRecognizer,

OpenAI ,

}

]

}

✉  **pwang009** 2 months ago

Answer is correct. Validated.

1. Create a FormRecognizer resource
2. On the resource screen, click Export Template under automation
3. Check the template created

upvoted 4 times

✉  **Murtuza** 2 months ago

"kind": "FormRecognizer"

upvoted 1 times

✉  **Murtuza** 2 months ago

The template specifies parameters like cognitiveServiceName, location, and sku.

It creates a Cognitive Services account of type CognitiveServices with the specified properties.

upvoted 2 times

Question #21

You are building an app named App1 that will use Azure AI Document Intelligence to extract the following data from scanned documents:

- Shipping address
- Billing address
- Customer ID
- Amount due
- Due date
- Total tax
- Subtotal

You need to identify which model to use for App1. The solution must minimize development effort.

Which model should you use?

- A. custom extraction model
- B. contract
- C. invoice
- D. general document

Correct Answer: C

Community vote distribution

C (100%)

 **reiwanotora** 2 days, 6 hours ago

Selected Answer: C

C "invoice" is right answer.

upvoted 1 times

 **Murtuza** 2 months ago

To extract the specified data from scanned documents with minimal development effort, you should use the invoice model. The invoice model is specifically designed to handle structured information commonly found in invoices, including shipping and billing addresses, customer IDs, amounts due, due dates, total taxes, and subtotal. Therefore, the correct choice for App1 is C. invoice

upvoted 1 times

 **rober13** 2 months, 1 week ago

Selected Answer: C

Reference: <https://learn.microsoft.com/en-us/azure/ai-services/document-intelligence/concept-invoice?view=doc-intel-4.0.0#field-extraction>

upvoted 2 times

Topic 5 - Question Set 5

Question #1

Topic 5

You build a bot by using the Microsoft Bot Framework SDK and the Azure Bot Service.

You plan to deploy the bot to Azure.

You register the bot by using the Bot Channels Registration service.

Which two values are required to complete the deployment? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. botId
- B. tenantId
- C. appId
- D. objectId
- E. appSecret

Correct Answer: CE

Reference:

<https://github.com/MicrosoftDocs/bot-docs/blob/live/articles/bot-service-quickstart-registration.md>

Community vote distribution

CE (100%)

✉️  **zellck**  11 months ago

Selected Answer: CE

CE is the answer.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-quickstart-registration?view=azure-bot-service-4.0&tabs=multitenant>

The identity information you need to add depends on the bot's application type. Provide the following values in your configuration file.

- MicrosoftAppId
- The bot's app ID.
- MicrosoftAppPassword
- The bot's app password.

upvoted 5 times

✉️  **evangelist**  3 months, 3 weeks ago

Selected Answer: CE

The answer is correct

upvoted 1 times

✉️  **rdemontis** 6 months, 3 weeks ago

Selected Answer: CE

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-quickstart-registration?view=azure-bot-service-4.0&tabs=multitenant>

upvoted 1 times

✉️  **Eltooth** 1 year, 10 months ago

Selected Answer: CE

C and E are correct answers.

<https://docs.microsoft.com/en-gb/learn/modules/design-bot-conversation-flow/5-deploy-bot>

upvoted 4 times

✉️  **Eltooth** 1 year, 10 months ago

Create the Azure resources required to support your bot

Your will need to create an Azure application registration to give your bot an identity it can use to access resources, and a bot application service to host the bot.

Register an Azure app

You can create the application registration by using the `az ad app create` Azure command-line interface (CLI) command, specifying a display name and password for your app identity. This command registers the app and returns its registration information, including a unique application ID that you will need in the following step.

Create a bot application service

Your bot requires a Bot Channels Registration resource, along with associated application service and application service plan. To create these resources, you can use the Azure resource deployment templates provided with the Bot Framework SDK template you used to create your bot.

Just run the az deployment group create command, referencing the deployment template and specifying your bot application registration's ID (from the az ad app create command output) and the password you specified.

upvoted 1 times

✉ **reachmymind** 2 years, 2 months ago

appId
appSecret

```
az deployment sub create --template-file "<path-to-template-with-new-rg.json>" --location <region-location-name> --parameters  
appType="MultiTenant" appId="<app-id-from-previous-step>" appSecret="<password-from-previous-step>" botId="<id or bot-app-service-  
name>" botSku=F0 newAppServicePlanName="<new-service-plan-name>" newWebAppName="<bot-app-service-name>" groupName="<new-  
group-name>" groupLocation="<region-location-name>" newAppServicePlanLocation="<region-location-name>" --name "<bot-app-service-  
name>"
```

<https://docs.microsoft.com/en-us/azure/bot-service/bot-builder-deploy-az-cli?view=azure-bot-service-4.0&tabs=csharp%2Cmultitenant>

upvoted 4 times

✉ **JTWang** 2 years ago

The bot channels registration registers your web service as a bot with the Bot Framework, provided you have a Microsoft App Id and App password (client secret).

<https://github.com/OfficeDev/Microsoft-Teams-Samples/blob/main/samples/tab-adaptive-cards/nodejs/Wiki/azure-bot-channels-registration.md#create-the-bot-channels-registration>

upvoted 1 times

✉ **jayf** 2 years, 8 months ago

Answer should be correct

upvoted 3 times

✉ **MII1975** 2 years, 6 months ago

Agree

<https://github.com/MicrosoftDocs/bot-docs/blob/live/articles/bot-service-quickstart-registration.md#update-the-bot>

upvoted 2 times

Question #2

HOTSPOT -

You are building a chatbot by using the Microsoft Bot Framework Composer.

You have the dialog design shown in the following exhibit.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area**Statements**

Yes	No
<input type="radio"/>	<input type="radio"/>
<input type="radio"/>	<input type="radio"/>
<input type="radio"/>	<input type="radio"/>

`user.name` is an entity.

The dialog asks for a user name and a user age and assigns appropriate values to the `user.name` and `user.age` properties.

The chatbot attempts to take the first non-null entity value for `user.name` or `personName` and assigns the value to `user.name`.

Correct Answer:

Answer Area**Statements**

Yes	No
<input type="radio"/>	<input checked="" type="radio"/>
<input checked="" type="radio"/>	<input type="radio"/>
<input checked="" type="radio"/>	<input type="radio"/>

`user.name` is an entity.

The dialog asks for a user name and a user age and assigns appropriate values to the `user.name` and `user.age` properties.

The chatbot attempts to take the first non-null entity value for `user.name` or `personName` and assigns the value to `user.name`.

Box 1: No -

User.name is a property.

Box 2: Yes -

Box 3: Yes -

The coalesce() function evaluates a list of expressions and returns the first non-null (or non-empty for string) expression.

Reference:

<https://docs.microsoft.com/en-us/composer/concept-language-generation> <https://docs.microsoft.com/en-us/azure/data-explorer/kusto/query/coalesceffunction>

 **Duch003**  2 years, 11 months ago

I am not sure if the correct answer for third is yes - the fields are showing examples, not the actual values typed in the boxes.
upvoted 9 times

 **TuongNN** 2 years, 8 months ago

That's what I thinking about..
upvoted 5 times

 **rdemontis** 6 months, 3 weeks ago

I also had this doubt but I don't want to think that Microsoft wants to confuse people's ideas in this way, that would be shameful. At this point the test would not be a way to assess people's skills but would become mere garbage
upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

anyway an example is represented with "ex." at the beginning of the entered text. For the third field it does not, so I think the coalesce function is really applied
upvoted 1 times

 **TheB**  2 years, 10 months ago

The answer is correct.
upvoted 8 times

 **Murtuza**  2 months ago

The third answer in this question is YES
The chatbot indeed attempts to extract relevant information from the user's input. In this case, it looks for non-null values associated with the userName or personName entities and assigns those values to the user.name property. This process allows the chatbot to personalize its interactions and provide a more tailored experience for the user.
upvoted 2 times

 **zellck** 11 months ago

NYY is the answer.

<https://learn.microsoft.com/en-us/composer/how-to-ask-for-user-input?tabs=v2x#create-a-text-input-action>

<https://learn.microsoft.com/en-us/azure/bot-service/adaptive-expressions/adaptive-expressions-prebuilt-functions?view=azure-bot-service-4.0#coalesce>

Return the first non-null value from one or more parameters. Empty strings, empty arrays, and empty objects are not null.
upvoted 7 times

 **Pixelmate** 11 months ago

this was on exam 28/06
upvoted 6 times

 **2ez4Zane** 2 years, 1 month ago

The answer is correct. The 3rd one is yes
The coalesce() function evaluates a list of expressions and returns the first non-null (or non-empty for string) expression.
upvoted 3 times

 **Moody_L** 1 year, 12 months ago

Disagreed. The coalesce() function uses comma to separate out a list of expressions. However, in the question, the separator is a period.
``coalesce(tolong("not a number"), tolong("42"), 33) == 42``
Ref: <https://docs.microsoft.com/en-us/azure/data-explorer/kusto/query/coalesceffunction>
upvoted 1 times

 **ninja** 1 year, 9 months ago

I would agree with 2ez4Zane. Your link is a very good reference. I would consider the period in the question is a typo.
upvoted 4 times

 **RamonRW** 2 years, 4 months ago

Answers 1 and 2, I are correct in my eyes. But with the third answer, I am not sure. Especially it is referring to personName and I do not see that property anywhere. Usually that's a hint, if an answer should be wrong. What do you think or do I miss something here?
upvoted 1 times

✉️ **RamonRW** 2 years, 4 months ago

I missed something indeed. There is a coalesce function for username and personname. I am still not sure, if the correct answer should actually be no.

upvoted 1 times

✉️ **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 3 times

Question #3

Topic 5

You are building a multilingual chatbot.

You need to send a different answer for positive and negative messages.

Which two Language service APIs should you use? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Linked entities from a well-known knowledge base
- B. Sentiment Analysis
- C. Key Phrases
- D. Detect Language
- E. Named Entity Recognition

Correct Answer: BD

B: The Text Analytics API's Sentiment Analysis feature provides two ways for detecting positive and negative sentiment. If you send a Sentiment Analysis request, the API will return sentiment labels (such as "negative", "neutral" and "positive") and confidence scores at the sentence and document-level.

D: The Language Detection feature of the Azure Text Analytics REST API evaluates text input for each document and returns language identifiers with a score that indicates the strength of the analysis.

This capability is useful for content stores that collect arbitrary text, where language is unknown.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/how-tos/text-analytics-how-to-sentiment-analysis?tabs=version-3-1>
<https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/how-tos/text-analytics-how-to-language-detection>

Community vote distribution

BD (100%)

✉️ **rdemontis** 6 months, 3 weeks ago

Selected Answer: BD

correct answer

upvoted 2 times

✉️ **Lion007** 7 months ago

Selected Answer: BD

B. Sentiment Analysis <-- different answer for positive and negative messages

D. Detect Language <-- multilingual chatbot

upvoted 4 times

✉️ **am20** 1 year, 4 months ago

for sure B but not sure why D? although other options are not useful either!

upvoted 1 times

✉️ **AzureJobsTillRetire** 1 year, 3 months ago

You are building a "multilingual" chatbot, and that is why D (Detect Language) is correct

upvoted 6 times

✉️ **s1cheng** 1 year, 5 months ago

Selected Answer: BD

Agreed with B & D

upvoted 2 times

Question #4

DRAG DROP -

You plan to build a chatbot to support task tracking.

You create a Language Understanding service named lu1.

You need to build a Language Understanding model to integrate into the chatbot. The solution must minimize development time to build the model.

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions	Answer Area
---------	-------------

Train the application.

Publish the application.

Add a new application.

Add example utterances.

Add the prebuilt domain ToDo.

Actions**Answer Area**

Train the application.

Publish the application.

Add a new application.

Add example utterances.

Add the prebuilt domain ToDo.

Add a new application.

Add example utterances.

Train the application.

Publish the application.

Correct Answer:

Step 1: Add a new application -

Create a new app -

1. Sign in to the LUIS portal with the URL of <https://www.luis.ai>.

2. Select Create new app.

3. Etc.

Step 2: Add example utterances.

In order to classify an utterance, the intent needs examples of user utterances that should be classified with this intent.

Step 3: Train the application -

Step 4: Publish the application -

In order to receive a LUIS prediction in a chat bot or other client application, you need to publish the app to the prediction endpoint.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/tutorial-intents-only>

✉  **Nebary**  1 year, 9 months ago

As per Udemy it is:

- Add a new application.
- Add the prebuilt domain ToDo.
- Train the application.
- Publish the application

upvoted 17 times

✉  **rdemontis** 6 months, 3 weeks ago

to me you are right. The fastest way to build the model is to use a prebuilt domain. In this case we do not need to add example utterances, because they are already present in the prebuilt domain.

upvoted 3 times

Eltooth Highly Voted 1 year, 10 months ago

Given answer is correct.

(Taken from Udemy practice exam question)

The order should be :

1. Add a new app,
2. Add example utterances,
3. Train the app,
4. Publish the app.

upvoted 7 times

varinder82 Most Recent 2 months ago

Final Answer

1. Add a new application
2. Add the prebuilt domain ToDo
3. Publish the application

upvoted 1 times

evangelist 3 months, 3 weeks ago

Correction to answer:

- 1.add a new application
2. add the prebuilt domain ToDo
3. Train the application
4. Publish the application

Select a pre-built domain (e.g., "Todo"), which provides a predefined set of intents and entities designed specifically for scenarios like task tracking.

Train the application, even when using a pre-built domain. Training the application is necessary because it optimizes the model to better understand user queries specific to your data and use case.

Even if you choose a pre-built domain, training the application remains a necessary step. Pre-built domains provide a good starting point, including a set of related intents and entities, but through training, you can optimize the model's performance for your specific data and scenario.

upvoted 3 times

Moody_L 1 year, 12 months ago

The question requires minimal time to build a model. Should we use prebuilt model with a custom domain ("ToDo") here?

Ref: <https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-reference-prebuilt-domains#prebuilt-domains-per-language>

upvoted 7 times

sdokmak 1 year, 11 months ago

True, but then we only have either 3 steps (don't need to need train prebuilt domain models), or can end up with 5 steps (if we add example utterances and train model)... so I'm stumped. Maybe the question will be different in exam.

upvoted 3 times

sdokmak 1 year, 11 months ago

Ideally we'd do:

1. Add a new application
2. Add the prebuilt domain ToDo
3. Publish the application

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/howto-add-prebuilt-models>

upvoted 5 times

Pyguy 1 year, 3 months ago

It was in Real Exam , and yes actually were 3 options there (Question said : Chose 3 action, NOT 4 , I also made like yours. :

1. Add a new application
2. Add the prebuilt domain ToDo
3. Publish the application

upvoted 13 times

Question #5

You are building a bot on a local computer by using the Microsoft Bot Framework. The bot will use an existing Language Understanding model.

You need to translate the Language Understanding model locally by using the Bot Framework CLI.

What should you do first?

- A. From the Language Understanding portal, clone the model.
- B. Export the model as an .lu file.
- C. Create a new Speech service.
- D. Create a new Language Understanding service.

Correct Answer: B

You might want to manage the translation and localization for the language understanding content for your bot independently.

Translate command in the @microsoft/bf-lu library takes advantage of the Microsoft text translation API to automatically machine translate .lu files to one or more than 60+ languages supported by the Microsoft text translation cognitive service.

What is translated?

An .lu file and optionally translate

Comments in the lu file -

LU reference link texts -

List of .lu files under a specific path.

Reference:

<https://github.com/microsoft/botframework-cli/blob/main/packages/luis/docs/translate-command.md>

Community vote distribution

B (100%)

 **Heegbaryee** Highly Voted 2 years, 8 months ago

Correct

upvoted 6 times

 **evangelist** Most Recent 3 months, 3 weeks ago

Selected Answer: B

"B. Export the model as an .lu file.

When you need to translate a Language Understanding (LUIS) model using the Microsoft Bot Framework CLI locally, the first step is to export the existing LUIS model as a .lu file. The .lu file format (Language Understanding file format) allows you to represent intents, entities, and other model components in a text-based form, enabling you to process, modify, or translate them using the Bot Framework CLI. This is a crucial step for working with and integrating LUIS models in a local environment."

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

answer and explanation seem correct

upvoted 1 times

 **ExamDev** 8 months, 3 weeks ago

Selected Answer: B

The answer is correct, because of the word "Locally"... so export it

upvoted 1 times

 **Nebary** 1 year, 9 months ago

Selected Answer: B

Correct

upvoted 1 times

 **azurelearner666** 2 years, 11 months ago

correct

upvoted 3 times

Question #6

DRAG DROP -

You are using a Language Understanding service to handle natural language input from the users of a web-based customer agent.

The users report that the agent frequently responds with the following generic response: "Sorry, I don't understand that."

You need to improve the ability of the agent to respond to requests.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions**Answer Area**

Add prebuilt domain models as required.

Validate the utterances logged for review and modify the model.

Migrate authoring to an Azure resource authoring key.

Enable active learning.

Enable log collection by using Log Analytics.

Train and republish the Language Understanding model.

Correct Answer:

Actions**Answer Area**

Add prebuilt domain models as required.

Add prebuilt domain models as required.

Validate the utterances logged for review and modify the model.

Enable active learning.

Migrate authoring to an Azure resource authoring key.

Train and republish the Language Understanding model.

Enable active learning.

Enable log collection by using Log Analytics.

Train and republish the Language Understanding model.

Step 1: Add prebuilt domain models as required.

Prebuilt models provide domains, intents, utterances, and entities. You can start your app with a prebuilt model or add a relevant model to your app later.

Note: Language Understanding (LUIS) provides prebuilt domains, which are pre-trained models of intents and entities that work together for domains or common categories of client applications.

The prebuilt domains are trained and ready to add to your LUIS app. The intents and entities of a prebuilt domain are fully customizable once you've added them to your app.

Step 2: Enable active learning -

To enable active learning, you must log user queries. This is accomplished by calling the endpoint query with the log=true querystring parameter and value.

Step 3: Train and republish the Language Understanding model

The process of reviewing endpoint utterances for correct predictions is called Active learning. Active learning captures endpoint queries and selects user's endpoint utterances that it is unsure of. You review these utterances to select the intent and mark entities for these real-world utterances. Accept these changes into your example utterances then train and publish. LUIS then identifies utterances more accurately.

Incorrect Answers:

Enable log collection by using Log Analytics

Application authors can choose to enable logging on the utterances that are sent to a published application. This is not done through Log Analytics.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-how-to-review-endpoint-utterances#log-user-queries-to-enable-active-learning> <https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-prebuilt-model>

 **g2000** Highly Voted  2 years ago

1. Enable active learning
 2. Validate the utterances logged for review and modify the model
 3. Train and republish the language understanding model
- not sure how the pre built model can improve accuracy....

upvoted 47 times

 **Eltooth** 1 year, 10 months ago

Agreed

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

agree with you

upvoted 3 times

 **rdemontis** 6 months, 3 weeks ago

to improve the model it's better to review the utterances logged by user feedback and then change the model. In this way we could adapt i on a real utterances basis

upvoted 2 times

 **GigaCaster** 1 year, 6 months ago

But the question isn't asking for accuracy but rather making it able to give a response instead of the current "Sorry, I don't understand that." response the customers are receiving. Thus the answer wants a short-term quick fix so that the team can work on the accuracy later, leading to pre-built model being used.

upvoted 4 times

 **ninja** 1 year, 9 months ago

Agreed

upvoted 2 times

 **varinder82** Most Recent  2 months ago

Final Answer:

1. Enable active learning
 2. Validate the utterances logged for review and modify the model
 3. Train and republish the language understanding model
- upvoted 1 times

 **evangelist** 3 months, 3 weeks ago

given answer is NOT Correct!

why:

Enable Active Learning: LUIS's Active Learning feature automatically identifies vague or low-confidence user inputs that may require further training.

Review and Add Utterances Logged for Review: Utterances collected through Active Learning are marked for review. Regularly reviewing and categorizing these utterances into the relevant intents is a crucial step in enhancing model performance.

Train and Republish the Language Understanding Model: After adding new utterances and potentially adjusting intents and entities, you need to retrain the LUIS model to incorporate these updates.

upvoted 3 times

 **zellck** 11 months ago

Same as Question 4.

<https://www.examtopics.com/discussions/microsoft/view/57673-exam-ai-102-topic-5-question-4-discussion>

upvoted 1 times

 **Eltooth** 1 year, 10 months ago

1. Enable active learning
 2. Validate the utterances logged for review and modify the model
 3. Train and republish the language understanding model
- upvoted 4 times

Question #7

You build a conversational bot named bot1.

You need to configure the bot to use a QnA Maker application.

From the Azure Portal, where can you find the information required by bot1 to connect to the QnA Maker application?

- A. Access control (IAM)
- B. Properties
- C. Keys and Endpoint
- D. Identity

Correct Answer: C

Obtain values to connect your bot to the knowledge base

1. In the QnA Maker site, select your knowledge base.
2. With your knowledge base open, select the SETTINGS tab. Record the value shown for service name. This value is useful for finding your knowledge base of interest when using the QnA Maker portal interface. It's not used to connect your bot app to this knowledge base.
3. Scroll down to find Deployment details and record the following values from the Postman sample HTTP request:
4. POST /knowledgebases/<knowledge-base-id>/generateAnswer
5. Host: <your-host-url>
6. Authorization: EndpointKey <your-endpoint-key>

Reference:

<https://docs.microsoft.com/en-us/azure/bot-service/bot-builder-howto-qna>

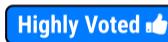
Community vote distribution

C (100%)

 **olowyinka**  2 years, 9 months ago

Correct 

upvoted 13 times

 **TheB**  2 years, 10 months ago

Given answer is correct!

upvoted 5 times

 **evangelist**  3 months, 3 weeks ago

Selected Answer: C

C. Keys and Endpoint

When configuring a conversational bot named bot1 to use the QnA Maker application, you need to obtain specific information from the QnA Maker service to establish the connection. This includes the service's endpoint address and the key used for authentication. You can find these details in the "Keys and Endpoint" section of the QnA Maker service in the Azure portal. Here, you will find the endpoint address and key for API calls, crucial for the bot to access and retrieve information from the QnA Maker service securely. These details enable your bot to interact safely with the QnA Maker service and retrieve answers based on your knowledge base content.

upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: C

given answer is correct

upvoted 2 times

 **Eltooth** 1 year, 10 months ago

Selected Answer: C

C is the correct answer : Keys and Endpoint.

upvoted 3 times

Question #8

HOTSPOT -

You are building a chatbot for a Microsoft Teams channel by using the Microsoft Bot Framework SDK. The chatbot will use the following code.

```
protected override async Task OnMembersAddedAsync(IList<ChannelAccount>
membersAdded, ITurnContext<IConversationUpdateActivity> turnContext,
CancellationToken cancellationToken)
{
    foreach (var member in membersAdded)
        if (member.Id != turnContext.Activity.Recipient.Id)
            await turnContext.SendActivityAsync($"Hi there - {member.Name}.
{WelcomeMessage}", cancellationToken: cancellationToken);
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
OnMembersAddedAsync will be triggered when a user joins the conversation.	<input type="radio"/>	<input type="radio"/>
When a new user joins the conversation, the existing users in the conversation will see the chatbot greeting.	<input type="radio"/>	<input type="radio"/>
OnMembersAddedAsync will be initialized when a user sends a message.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
OnMembersAddedAsync will be triggered when a user joins the conversation.	<input checked="" type="radio"/>	<input type="radio"/>
When a new user joins the conversation, the existing users in the conversation will see the chatbot greeting.	<input checked="" type="radio"/>	<input type="radio"/>
OnMembersAddedAsync will be initialized when a user sends a message.	<input type="radio"/>	<input checked="" type="radio"/>

Box 1: Yes -

The ActivityHandler.OnMembersAddedAsync method overrides this in a derived class to provide logic for when members other than the bot join the conversation, such as your bot's welcome logic.

Box 2: Yes -

membersAdded is a list of all the members added to the conversation, as described by the conversation update activity.

Box 3: No -

Reference:

<https://docs.microsoft.com/en-us/dotnet/api/microsoft.bot.builder.activityhandler.onmembersaddedasync?view=botbuilder-dotnet-stable>

 **guruaws2021**  2 years, 9 months ago

2nd one should be no, when a new member joins the chatbot will only greet the new member and not all existing members
upvoted 17 times

 **Viktor** 2 years, 8 months ago

I feel the answer is correct. The members are all in a "group". So when a new member is added, the old members will see the chatbot greet the new member ("Hello \${new.member}"), and not every other existing members.
upvoted 20 times

 **rdemontis** 6 months, 3 weeks ago

Agree with you.
"In the context of the OnMembersAddedAsync method, the membersAdded list contains information about the members who have just been added to a conversation. Here's how this iteration process works:

Method Parameters:

membersAdded: This is a list of ChannelAccount objects, each of which represents a member who has just been added to the conversation. ChannelAccount objects contain information about the member, such as their Id, Name, Role, and so on.
turnContext: The current context of the conversation and activity.
cancellationToken: A cancellation token, which can be used to cancel the request if necessary.

Iterating through Added Members: The code block starts with an iteration through the membersAdded list. This iteration is performed using a foreach loop. For each member in the list, the code within the loop will be executed." (ChatGPT)

upvoted 1 times

✉️ **rdemontis** 6 months, 3 weeks ago

if you look the code you can see that the method sends a Welcome message for each member added: "Hi there - {member.Name}". It would be very strange to give a welcome message to everyone in the chat, even member already joined, every time a new one is added. So I think ChatGPT is right and the second answer is wrong. It should be "NO".

Here another similar example:

<https://learn.microsoft.com/en-us/microsoftteams/platform/bots/how-to/conversations/subscribe-to-conversation-events?tabs=dotnet#members-added>

upvoted 2 times

✉️ **reachmymind** **Highly Voted** 2 years, 2 months ago

A1. Yes

<https://docs.microsoft.com/en-us/dotnet/api/microsoft.bot.builder.activityhandler.onmembersaddedasync?view=botbuilder-dotnet-stable>
ActivityHandler.OnMembersAddedAsync(IList<ChannelAccount>, ITurnContext<IConversationUpdateActivity>, CancellationToken)

Method invoked when members other than the bot join the conversation

A2. No

<https://docs.microsoft.com/en-us/dotnet/api/microsoft.bot.builder.turncontext.sendactivityasync?view=botbuilder-dotnet-stable>
Sends a message activity to the sender of the incoming activity in turncontext

A3: No

<https://docs.microsoft.com/en-us/dotnet/api/microsoft.bot.builder.activityhandler.onconversationupdateactivityasync?view=botbuilder-dotnet-stable>

ActivityHandler.OnConversationUpdateActivityAsync(ITurnContext<IConversationUpdateActivity>, CancellationToken)

Method invoked when a conversation update activity that indicates one or more users other than the bot are joining the conversation

upvoted 12 times

✉️ **ninjapunk** 1 year, 5 months ago

This is correct, as per the docs SendActivityAsync "sends an activity to the sender of the incoming activity", thus not a public message to be seen by everyone in the group.

upvoted 1 times

✉️ **Murtuza** **Most Recent** 1 month, 3 weeks ago

Yes, the OnMembersAddedAsync method is triggered when a new user joins the conversation.

Yes, when a new user joins the conversation, the existing users in the conversation will see the chatbot greeting. This is because the OnMembersAddedAsync method sends a welcome message to each new member.

No, the OnMembersAddedAsync method is not initialized when a user sends a message. It is only triggered when a new user joins the conversation. The method that handles incoming messages is typically OnMessageActivityAsync

upvoted 1 times

✉️ **idcanymore** 3 months, 2 weeks ago

Discord sends this exact message to people joining groups and everyone in the chat sees it. don't let chatgpt do your thinking for you.

upvoted 2 times

✉️ **PCRamirez** 3 months, 2 weeks ago

YNN: Asking Windows Copilot about second answer:

No, the existing users in the conversation will not see the chatbot greeting. The code snippet provided shows a method `OnMembersAddedAsync` that is triggered when new members are added to the conversation. Inside this method, there is a conditional statement `if (member.Id != turnContext.Activity.Recipient.Id)` which ensures that the welcome message is sent only to the new member who has joined, and not to the existing users or the bot itself. The welcome message is personalized with the new member's name and is sent directly to them.

upvoted 1 times

✉️ **evangelist** 3 months, 3 weeks ago

given answer is NOT Correct!

why:

Enable Active Learning: LUIS's Active Learning feature automatically identifies vague or low-confidence user inputs that may require further training.

Review and Add Utterances Logged for Review: Utterances collected through Active Learning are marked for review. Regularly reviewing and categorizing these utterances into the relevant intents is a crucial step in enhancing model performance.

Train and Republish the Language Understanding Model: After adding new utterances and potentially adjusting intents and entities, you need to retrain the LUIS model to incorporate these updates.

upvoted 1 times

✉️ **dimsok** 4 months, 3 weeks ago

YYN

- for the second one, everyone will see that the bot said hi to a new user , the message sent by the bot is not private, its posted in the channel

upvoted 3 times

✉️ **sl_mslconsulting** 7 months, 1 week ago

The answer to the second question is no. The membersAdded will have two guys : you and the bot. The OnMembersAddedAsync will only be triggered when non bot member joined the conversation which is you. The logic coded here will not send the welcome message to the bot which is also a existing user. You and the bot are both members. I think we are reading to much to this - it's not a chat room where a new guy joined everyone will be notified.

upvoted 3 times

✉ **sl_mslconsulting** 7 months, 1 week ago

You can actually write the code to verify this yourself and comment out the statement that check if a member is a bot to see what would happen.

upvoted 1 times

✉ **zellck** 11 months ago

YYN is the answer.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-activity-handler-concept?view=azure-bot-service-4.0&tabs=csharp#activity-handling>
- OnMembersAddedAsync

Non-bot members joined the conversation

upvoted 4 times

✉ **AzureJobsTillRetire** 1 year, 3 months ago

The given answers are correct.

On the second question:-

The sender is the Microsoft Teams channel.

"from" Identifies the sender of the message. example: from: ChannelAccount

The recipients are the members

According to the code,

If new member <> existing member, then send out welcome message

if (member.Id != turnContext.Activity.Recipient.Id)

upvoted 2 times

✉ **saadashraf** 1 year, 1 month ago

Thanks, spent quite a time on this question. Your answer is absolutely correct

upvoted 1 times

✉ **Adedoyin_Simeon** 2 years, 7 months ago

The answer is correct.

upvoted 3 times

✉ **TheB** 2 years, 10 months ago

The answer looks correct.

upvoted 7 times

Question #9

HOTSPOT -

You are reviewing the design of a chatbot. The chatbot includes a language generation file that contains the following fragment.

```
# Greet(user)
- ${Greeting()}, ${user.name}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
`\${user.name}` retrieves the user name by using a prompt.	<input type="radio"/>	<input type="radio"/>
Greet() is the name of the language generation template.	<input type="radio"/>	<input type="radio"/>
`\${Greeting()}` is a reference to a template in the language generation file.	<input type="radio"/>	<input type="radio"/>

Answer Area

Statements	Yes	No
`\${user.name}` retrieves the user name by using a prompt.	<input type="radio"/>	<input checked="" type="radio"/>
Greet() is the name of the language generation template.	<input type="radio"/>	<input checked="" type="radio"/>
`\${Greeting()}` is a reference to a template in the language generation file.	<input checked="" type="radio"/>	<input type="radio"/>

Box 1: No -

Example: Greet a user whose name is stored in `user.name`

```
- ${welcomeUser(user.name)}
```

Example: Greet a user whose name you don't know:

```
- ${welcomeUser()}
```

Box 2: No -

Greet(User) is a Send a response action.

Box 3: Yes -

Reference:

<https://docs.microsoft.com/en-us/composer/how-to-ask-for-user-input>

 **zellck** Highly Voted  11 months ago

NYY is the answer.

<https://learn.microsoft.com/en-us/azure/bot-service/file-format/bot-builder-lg-file-format?view=azure-bot-service-4.0#template-names>

<https://learn.microsoft.com/en-us/azure/bot-service/file-format/bot-builder-lg-file-format?view=azure-bot-service-4.0#references-to-templates>
Variation text can include references to another named template to aid with composition and resolution of sophisticated responses. References to other named templates are denoted using braces, such as \${<TemplateName>()}.

upvoted 14 times

 **rdemontis** 6 months, 3 weeks ago

To me you're right

<https://learn.microsoft.com/en-us/azure/bot-service/file-format/bot-builder-lg-file-format?view=azure-bot-service-4.0#templates>

upvoted 1 times

 **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 6 times

 **evangelist** Highly Voted  3 months, 3 weeks ago

Correct answer is : NYY

why:

`${user.name}` retrieves the user's name using a prompt.-->No.

In this context, `${user.name}` directly retrieves the user's name from the incoming user object. It doesn't obtain the user's name through a prompt but assumes that the user object already contains the name property.

Greet() is the name of the language generation template.-->Yes.

Greet(user) defines a language generation template named Greet. This template takes a user parameter, indicating that Greet is the template name used to generate personalized greeting messages.

`${Greeting()}` is a reference to a template in the language generation file.-->Yes.

`${Greeting()}` is a reference to another template in the language generation file. This syntax indicates that Greeting is a template that will be invoked to generate all or part of a greeting when messages are generated. This allows developers to reuse and combine templates in the language generation file to create more complex messages.

upvoted 6 times

✉️ **Lion007** Most Recent 7 months ago

`${user.name}` retrieves the "name" parameter from the "user" by using PARAMETER (NOT Prompt). Check out the Parametrization of Templates section <https://learn.microsoft.com/en-us/azure/bot-service/file-format/bot-builder-lg-file-format?view=azure-bot-service-4.0#parametrization-of-templates>

Greet() is indeed the NAME of the language generation template. Why not Greet(user) is the name, I hear you asking? Not including the parameter "user" in the question doesn't change the fact that what comes after the hashtag # is the NAME of the template, only without referring to its parameter. Again, check out Parametrization of Templates section <https://learn.microsoft.com/en-us/azure/bot-service/file-format/bot-builder-lg-file-format?view=azure-bot-service-4.0#parametrization-of-templates>

`${Greeting()}` is a reference to another template's name called Greeting() template.

upvoted 1 times

✉️ **Lion007** 7 months ago

N.Y.Y.

upvoted 1 times

✉️ **sl_msiconsulting** 7 months, 1 week ago

Just think of \$ as the eval in JavaScript. It evals an expression. A template is just like a function you can invoke. The answers are then quite evident

upvoted 1 times

✉️ **kail85** 12 months ago

No, Yes, Yes

upvoted 1 times

✉️ **NNU** 1 year, 3 months ago

the user.name retrieves from .log file, the answer is No

second is yes Greet() is the name of LG template

TemplateName

- You can say cheese and tomato \[toppings are optional]

third yes is a reference to a template

Variation text can include references to another named template to aid with composition and resolution of sophisticated responses. References to other named templates are denoted using braces, such as `${<TemplateName>()}``.

<https://learn.microsoft.com/en-us/azure/bot-service/file-format/bot-builder-lg-file-format?view=azure-bot-service-4.0>

upvoted 3 times

✉️ **Eltooth** 1 year, 10 months ago

No

No

Yes

upvoted 2 times

✉️ **kubikjakubik** 2 years, 11 months ago

First is true. You even have it in your reference.

<https://docs.microsoft.com/en-us/composer/how-to-ask-for-user-input>

"Hello \${user.name}, nice to talk to you!" ... so the part \${user.name} in fact does retrieve the name of the user.

upvoted 4 times

✉️ **AzureJobsTillRetire** 1 year, 3 months ago

One way of getting the user name is by using a prompt, but there are other ways as well

upvoted 1 times

✉️ **azurelearner666** 2 years, 11 months ago

it does retrieve the user name, but NOT USING A PROMPT. retrieves the already stored property.

But without a PROMPT.

so, the response is correct, it's a NO.

1. NO

2. NO

3. YES

upvoted 28 times

Question #10

HOTSPOT -

You are building a chatbot by using the Microsoft Bot Framework SDK.

You use an object named `UserProfile` to store user profile information and an object named `ConversationData` to store information related to a conversation.

You create the following state accessors to store both objects in state. `var userStateAccessors = _userState.CreateProperty<UserProfile>(nameof(UserProfile)); var conversationStateAccessors = _conversationState.CreateProperty<ConversationData>(nameof(ConversationData));`

The state storage mechanism is set to Memory Storage.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
------------	-----	----

The code will create and maintain the `UserProfile` object in the underlying storage layer.

The code will create and maintain the `ConversationData` object in the underlying storage layer.

The `UserProfile` and `ConversationData` objects will persist when the Bot Framework runtime terminates.

Answer Area

Statements	Yes	No
------------	-----	----

The code will create and maintain the `UserProfile` object in the underlying storage layer.

Correct Answer:

The code will create and maintain the `ConversationData` object in the underlying storage layer.

The `UserProfile` and `ConversationData` objects will persist when the Bot Framework runtime terminates.

Box 1: Yes -

You create property accessors using the `CreateProperty` method that provides a handle to the `BotState` object. Each state property accessor allows you to get or set the value of the associated state property.

Box 2: Yes -

Box 3: No -

Before you exit the turn handler, you use the state management objects' `SaveChangesAsync()` method to write all state changes back to storage.

Reference:

<https://docs.microsoft.com/en-us/azure/bot-service/bot-builder-howto-v4-state>

 **Voxo**  2 years, 6 months ago

Looks correct

upvoted 11 times

 **zellck**  11 months ago

YYN is the answer.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-state?view=azure-bot-service-4.0#state-property-accessors>
 State property accessors are used to actually read or write one of your state properties, and provide get, set, and delete methods for accessing your state properties from within a turn. To create an accessor, you must provide the property name, which usually takes place when you're initializing your bot. Then, you can use that accessor to get and manipulate that property of your bot's state.

The accessors allow the SDK to get state from the underlying storage, and update the bot's state cache for you. The state cache is a local cache maintained by your bot that stores the state object for you, allowing read and write operations without accessing the underlying storage. If it isn't already in the cache, calling the accessor's get method retrieves state and also places it in the cache. Once retrieved, the state property can be manipulated just like a local variable.

upvoted 10 times

 **zellick** 11 months ago

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-state?view=azure-bot-service-4.0#storage-layer>

The Bot Framework SDK includes some implementations for the storage layer:

- Memory storage implements in-memory storage for testing purposes. In-memory data storage is intended for local testing only as this storage is volatile and temporary. The data is cleared each time the bot is restarted.

upvoted 3 times

 **rdemontis** 6 months, 3 weeks ago

thanks for explanation

upvoted 1 times

 **varinder82** Most Recent 2 months ago

Final Answer

YYN

upvoted 2 times

 **Mehe323** 2 months ago

The first two are YES, the last is NO.

"Memory storage

The Bot Framework SDK allows you to store user inputs using in-memory storage. Since in-memory storage is cleared each time the bot is restarted, it's best suited for testing purposes and isn't intended for production use. Persistent storage types, such as database storage, are best for production bots."

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-howto-v4-storage?view=azure-bot-service-4.0&tabs=csharp>

upvoted 1 times

 **sl_msiconsulting** 7 months, 1 week ago

Ok seems everyone get this wrong. The answer are N, N, N. CreateProperty method provides a handle to the BotState object. BotState is part of the cache. Nothing will save to the storage layer unless you explicitly do so in the code. The storage layer is in memory in this case as a result nothing will be persisted once the bot is destroyed.

upvoted 1 times

 **dimsok** 4 months, 3 weeks ago

you are right, but i think since you dont the full code to check, save is implied

upvoted 1 times

 **kail85** 11 months, 4 weeks ago

Yes, No, No

upvoted 1 times

 **kail85** 11 months, 4 weeks ago

First one is Yes.

The code provided creates state accessors for both UserProfile and ConversationData objects. These state accessors are used to store and retrieve the respective objects in the underlying state storage mechanism, which in this case is Memory Storage. When you use the state accessors to set or get the UserProfile object, the Microsoft Bot Framework SDK will create and maintain the UserProfile object in the Memory Storage, allowing you to persist user profile information across multiple turns in the conversation.

upvoted 1 times

 **flutterb** 1 year, 6 months ago

Answer is :

No

No

No

If the user profile and conversation states are committed to memory, then they are in fact stored in cache memory and not on the underlying storage layer (disk). Hence, the user profile and conversation data will get deleted once the session runtime is over.

upvoted 4 times

 **ninjapunk** 1 year, 5 months ago

Wrong. The first two answers are true.

As per the documentation, the BotState Class "defines a state management object and automates the reading and writing of associated state properties to a storage layer."

CreateProperty is a method of the BotState Class, which "creates a named state property within the scope of a BotState and returns an accessor for the property."

Reference: <https://learn.microsoft.com/en-us/dotnet/api/microsoft.bot.builder.botstate?view=botbuilder-dotnet-stable>

upvoted 2 times

Question #11

HOTSPOT -

You are building a chatbot that will provide information to users as shown in the following exhibit.

Passengers

Sarah Hum
Jeremy Goldberg
Evan Litvak

2 Stops**Tue, May 30, 2017 10:25 PM**

San Francisco

Amsterdam



San Francisco

Amsterdam

SFO

AMS

SFO

AMS

Non-Stop**Fri, Jun 2, 2017 11:55 PM**

San Francisco

Amsterdam



San Francisco

Amsterdam

SFO

AMS

SFO

AMS

Total

\$4,032.54

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

The chatbot is showing [answer choice].

▼

an Adaptive Card
a Hero Card
a Thumbnail Card

The card includes [answer choice].

▼

an action set
an image
an image group
media

Answer Area

The chatbot is showing **[answer choice]**.

Correct Answer:

The card includes **[answer choice]**.

an Adaptive Card
a Hero Card
a Thumbnail Card
an action set
an image
an image group
media

Box 1: A Thumbnail card -

A Thumbnail card typically contains a single thumbnail image, some short text, and one or more buttons.

Incorrect Answers:

- an Adaptive card is highly customizable card that can contain any combination of text, speech, images, buttons, and input fields.
- a Hero card typically contains a single large image, one or more buttons, and a small amount of text.

Box 2: an image -

Reference:

<https://docs.microsoft.com/en-us/microsoftteams/platform/task-modules-and-cards/cards/cards-reference>

 **motu** Highly Voted 2 years, 11 months ago

ET answer is wrong. It's an adaptive card using an image.
upvoted 38 times

 **PaPaTee** 2 years, 10 months ago

Yes, the answer is Adaptive Card. The image is a combination of many texts, images, with date and currency fields.
upvoted 6 times

 **plalwa** 2 years, 5 months ago

its Thumbnail, Adaptive cards require buttons. Like in MS teams approve PR :)
upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

Agree.

zellck 4 months, 1 week ago
1. Adaptive Card
2. image

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-user-experience?view=azure-bot-service-4.0#cards>
upvoted 4 times

 **snilu** Highly Voted 2 years, 5 months ago

Answer is Adaptive Card
<https://docs.microsoft.com/en-us/composer/how-to-send-cards?tabs=v1x>
upvoted 29 times

 **Trumpenstein** 2 years, 4 months ago

Good Find - the image at the end is nearly identical with the screenshot!
upvoted 7 times

 **varinder82** Most Recent 2 months ago

Final Answer:
1- Adaptive Card
2- Image
upvoted 1 times

 **evangelist** 3 months, 3 weeks ago

The answer is adaptive card and it is the same on Microsoft Documentation:
<https://learn.microsoft.com/en-us/composer/how-to-send-cards?tabs=v1x#adaptive-card>
upvoted 2 times

 **sl_msiconsulting** 7 months, 1 week ago

Here is how you compose this card: <https://adaptivecards.io/samples/FlightItinerary.html>.

There are many text blocks and two images. Now you know what to choose.

upvoted 1 times

 **sl_msiconsulting** 7 months, 1 week ago

You really need to dig a bit deeper to find the right answer and get your hands dirty. It's very easy to show during your interviews that you are just guessing and don't really understand the topic.

upvoted 1 times

 **sl_msiconsulting** 7 months, 1 week ago

*Same image source but shows up twice <https://adaptivecards.io/content/airplane.png>

upvoted 1 times

 **james2033** 9 months, 1 week ago

1. is showing Adaptive card <https://learn.microsoft.com/en-us/composer/how-to-send-cards?tabs=v1x#adaptive-card>

2. image (just 1 image) of icon file.

upvoted 1 times

 **zellck** 11 months ago

1. Adaptive Card

2. image

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-user-experience?view=azure-bot-service-4.0#cards>

- AdaptiveCard

An open card exchange format rendered as a JSON object. Typically used for cross-channel deployment of cards. Cards adapt to the look and feel of each host channel.

upvoted 5 times

 **AzureJobsTillRetire** 1 year, 3 months ago

The given answers are correct.

First of all, everyone agrees that the answer to the second question is an image and not an action set, or an image group, or media. Hence the card includes an image and not many images. If there is only one image, the answer to the first question has to be a Thumbnail card, as an adaptive card is highly customizable and can contain any combination of text, speech, images, buttons, and input fields, which is not required.

upvoted 2 times

 **AzureJobsTillRetire** 1 year, 2 months ago

I was wrong. I think it is Adaptive Card. Although an adaptive card contains any combination of text, speech, images, buttons, and input fields, can also be rendered into a single PNG image.

<https://learn.microsoft.com/en-us/adaptive-cards/sdk/rendering-cards/net-image/getting-started>

upvoted 5 times

 **Eltooth** 1 year, 10 months ago

Answer is correct.

Adaptive Card

Image

<https://docs.microsoft.com/en-us/composer/how-to-send-cards?tabs=v1x>

upvoted 2 times

 **reachmymind** 2 years, 2 months ago

an Adaptive Card

an image

an Adaptive Card that is a JSON-serialized card object model rendered into a PNG image

<https://docs.microsoft.com/en-us/azure/bot-service/bot-service-design-user-experience?view=azure-bot-service-4.0>

<https://docs.microsoft.com/en-us/adaptive-cards/sdk/rendering-cards/net-image/render-a-card>

upvoted 2 times

 **RamonRW** 2 years, 4 months ago

The question is a bit confusing, as it is not entirely clear, if this is a sent screenshot. I think, if it is a screenshot Thumbnail and image makes absolutely sense, but if it is a "constructed" message, adaptive card would make sense. But due to the complexity of the adaptive card (images of the plane, information about passenger) several answers need to be right in the second element. Therefore, I believe the answer is correct. But the question is not really a good one.

upvoted 2 times

 **ParkXD** 1 year, 10 months ago

I agree with what you say.

And for the difference between cards: <https://docs.microsoft.com/en-us/microsoftteams/platform/task-modules-and-cards/cards/cards-reference>.

upvoted 1 times

 **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 2 times

 **SAANCODE** 2 years, 11 months ago

The answer is correct

upvoted 3 times

Question #12

HOTSPOT -

You are building a bot and that will use Language Understanding.

You have a LUDown file that contains the following content.

```
## Confirm
- confirm
- ok
- yes

## ExtractName
- call me steve !
- i am anna
- (i'm|i am) {@PersonName.Any}[]
- my name is {@PersonName.Any}[]

## Logout
- forget me
- log out

## SelectItem
- choose last
- choose the {@DirectionalReference=bottom left}
- choose {@DirectionalReference=top right}
- i like {@DirectionalReference=left} one

## SelectNone
- none

@ ml DirectionalReference
@ prebuilt personName
```

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

SelectItem is [answer choice].

a domain
an entity
an intent
an utterance

Choose {@DirectionalReference=top right} is [answer choice].

a domain
an entity
an intent
an utterance

Answer Area**SelectItem is [answer choice].**

a domain
an entity
an intent
an utterance

Correct Answer:**Choose {@DirectionalReference=top right} is [answer choice].**

a domain
an entity
an intent
an utterance

Reference:<https://github.com/solliancenet/tech-immersion-data-ai/blob/master/ai-exp1/README.md>**Zoul** Highly Voted 2 years, 7 months ago

Answer is correct. From the link provided:

SelectItem is the intent, and each item below it are example utterances that capture ways users can express this intent. Entities in .lu files are denoted using {<entityName>=<labeled value>} notation. Taking from our sample code once again, you can find the bottom left entity within the following utterance: choose the {DirectionalReference=bottom left}.

upvoted 15 times

zellck Highly Voted 11 months ago

1. intent
2. utterance

<https://learn.microsoft.com/en-us/azure/bot-service/file-format/bot-builder-lu-file-format?view=azure-bot-service-4.0#define-intents-using-sample-utterances>

Intents with their sample utterances are declared in the following way:

```
# <intent-name>
- <utterance1>
- <utterance2>
# <intent-name> describes a new intent definition section. Each line after the intent definition are example utterances that describe that intent using the - <utterance> format.
```

upvoted 9 times

rdemontis Most Recent 6 months, 3 weeks ago

correct answer and explanation. The provided reference demonstrates it

upvoted 2 times

ParkXD 1 year, 10 months ago

Answer is correct.

The LUIS model begins with categories of user intentions called intents. Each intent needs examples of user utterances (something a user should say). Each utterance can provide a variety of data that needs to be extracted with entities.

upvoted 1 times

Eltooth 1 year, 10 months ago

Intent

Utterance.

upvoted 2 times

sdokmak 1 year, 11 months ago

a trick stolen from previous questions
(U) (S) A, Utterance Statement
(E)(n)gland, Entity Noun
(I)(v)ory Coast, Intent Verb

upvoted 2 times

PHD_CHENG 2 years ago

Answer is correct

upvoted 1 times

Ravnit 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 1 times

Diem 2 years, 8 months ago

Based on the reference, it should be entity for the second question

upvoted 7 times

Question #13

HOTSPOT -

You are designing a conversation flow to be used in a chatbot.

You need to test the conversation flow by using the Microsoft Bot Framework Emulator.

How should you complete the .chat file? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
user=User1  
bot=watchbot  
user: I want a new watch.
```

```
bot: [ ] [Delay=3000]  
Attachment  
ConversationUpdate  
Typing
```

```
bot: I can help you with that! Let me see what I can find.
```

```
bot: Here's what I found.
```

```
bot:
```

```
[AttachmentLayout= [ ]  
adaptivecard  
carousel  
thumbnail
```

```
[Attachment=https://contoso.blob.core.windows.net/watch01.jpg]
```

```
[Attachment=https://contoso.blob.core.windows.net/watch02.jpg]
```

```
user: I like the first one.
```

```
bot: Sure, pulling up more information.
```

```
bot: [Attachment=cards\watchProfileCard.json
```

```
user: That's nice! Thank you.
```

```
bot: Sure, you are most welcome!
```

```
[ ]  
adaptivecard  
carousel  
list
```

Answer Area

```
user=User1  
bot=watchbot  
user: I want a new watch.
```

```
bot: [ ] [Delay=3000]  
Attachment  
ConversationUpdate  
Typing
```

```
bot: I can help you with that! Let me see what I can find.
```

```
bot: Here's what I found.
```

```
bot:
```

Correct Answer:

```
[AttachmentLayout=
```

```
[ ]  
adaptivecard  
carousel  
thumbnail
```

```
[Attachment=https://contoso.blob.core.windows.net/watch01.jpg]
```

```
[Attachment=https://contoso.blob.core.windows.net/watch02.jpg]
```

```
user: I like the first one.
```

```
bot: Sure, pulling up more information.
```

```
bot: [Attachment=cards\watchProfileCard.json
```

```
user: That's nice! Thank you.
```

```
bot: Sure, you are most welcome!
```

```
[ ]  
adaptivecard  
carousel  
list
```

Reference:

<https://docs.microsoft.com/en-us/azure/bot-service/bot-builder-howto-add-media-attachments?view=azure-bot-service-4.0&tabs=csharp>

  Diem  2 years, 8 months ago

The given answer is correct. Second box is Carousel which includes multiple attachments.

upvoted 10 times

 **zellck** Highly Voted 11 months ago

1. Typing
2. carousel
3. adaptivecard

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-user-experience?view=azure-bot-service-4.0#cards>
- CardCarousel

A horizontally scrollable collection of cards that allows your user to easily view a series of possible user choices.

- AdaptiveCard

An open card exchange format rendered as a JSON object. Typically used for cross-channel deployment of cards. Cards adapt to the look and feel of each host channel.

upvoted 5 times

 **rdemontis** Most Recent 6 months, 3 weeks ago

the answer seems correct

<https://docs.microsoft.com/en-us/azure/bot-service/bot-builder-howto-send-messages?view=azure-bot-service-4.0&tabs=csharp>

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-user-experience?view=azure-bot-service-4.0#cards>

upvoted 2 times

 **ziggy1117** 11 months, 3 weeks ago

answer is correct:

In the context of Azure Bot Service, a carousel refers to a visual component that displays a set of cards or items in a horizontally scrollable format. It is commonly used to present multiple options or pieces of information to the user in a visually appealing manner.

A carousel typically consists of multiple cards, where each card represents a specific item, option, or piece of content. Each card within the carousel can contain text, images, buttons, or other interactive elements. Users can scroll through the carousel horizontally to view and interact with different cards.

upvoted 1 times

 **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 2 times

 **GMKanon** 2 years, 8 months ago

Second box should be thumbnail

upvoted 2 times

 **SnowCheetah** 2 years, 8 months ago

I am not sure on the second answer is correct or not

1. Typing is correct (it's answer indicate on [delay]) <https://docs.microsoft.com/en-us/azure/bot-service/bot-builder-howto-send-messages?view=azure-bot-service-4.0&tabs=csharp>

2. Since attachment is attach with 2 images ==> it cannot be thumbnail, which can only a simple card as a proper response. however I am not sure adaptive card can be selected as well for this choice.

3. Adaptive card is correct (in context after user send selection, bot is sending detail of selected item) <https://docs.microsoft.com/en-us/azure/bot-service/bot-builder-howto-add-media-attachments?view=azure-bot-service-4.0&tabs=csharp>

upvoted 3 times

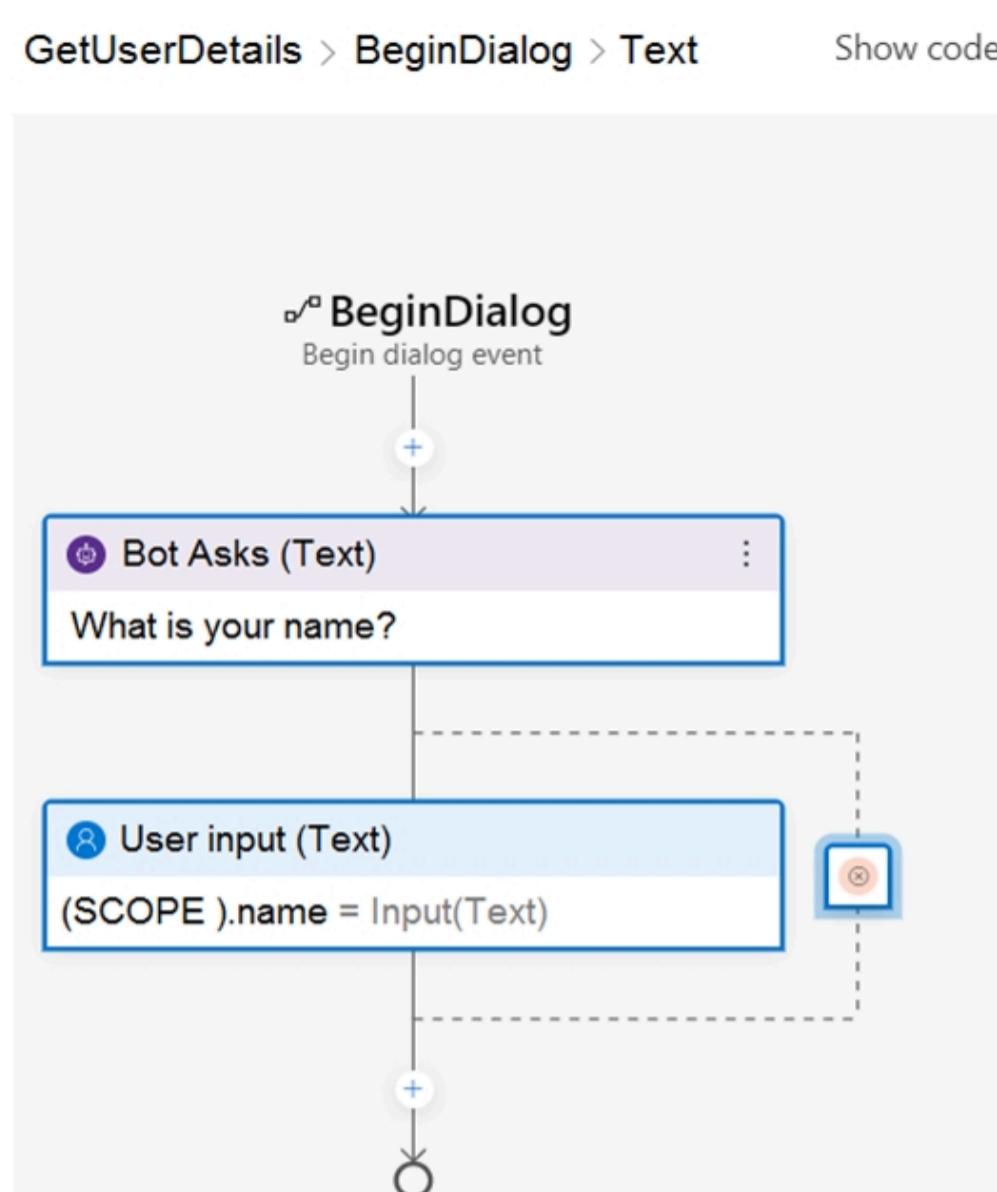
 **timmayy54** 2 years, 4 months ago

Fully Agree, for 2. having 2 attachments makes Carousel more right than Adaptive, plus Adaptive is last answer and double using the same one is quite rare.

upvoted 1 times

Question #14

You are building a chatbot by using the Microsoft Bot Framework Composer as shown in the exhibit. (Click the Exhibit tab.)

**Prompt for text**

Text Input

Collection information - Ask for a word or sentence.

[Learn more](#)**Bot Asks****User input****Other**

Property ⓘ

string

(SCOPE).name

Output format ⓘ

string

Value ⓘ

string

Expected responses (intent :

#TextInput_Response_FuvyF4)

The chatbot contains a dialog named GetUserDetails. GetUserDetails contains a TextInput control that prompts users for their name.

The user input will be stored in a property named name.

You need to ensure that you can dispose of the property when the last active dialog ends.

Which scope should you assign to name?

- A. dialog
- B. user
- C. turn
- D. conversation

Correct Answer: A

The dialog scope associates properties with the active dialog. Properties in the dialog scope are retained until the dialog ends.

Incorrect Answers:

A: The conversation scope associates properties with the current conversation. Properties in the conversation scope have a lifetime of the conversation itself.

These properties are in scope while the bot is processing an activity associated with the conversation (for example, multiple users together in a Microsoft Teams channel).

B: The user scope associates properties with the current user. Properties in the user scope do not expire. These properties are in scope while the bot is processing an activity associated with the user.

C: The turn scope associates properties with the current turn. Properties in the turn scope expire at the end of the turn.

Reference:

<https://docs.microsoft.com/en-us/composer/concept-memory?tabs=v2x>

Community vote distribution

A (100%)

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

<https://learn.microsoft.com/en-us/composer/concept-memory?tabs=v1x>

upvoted 2 times

 **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/composer/ref-memory-variables?tabs=v2x#the-scopes>

- Dialog

Properties associated with the active dialog. Properties in the dialog scope are kept until the dialog ends.

upvoted 4 times

 **Isidro** 2 years ago

Shouldn't it be USER? If Dialog is selected, once the last dialog is concluded, the information will be lost.

upvoted 1 times

 **abelarda** 1 week, 6 days ago

True, when it says "when the dialog ends", one could say that it means "after the dialog ends", or "until". Why are you so sure?

upvoted 1 times

 **Moody_L** 1 year, 12 months ago

Dialog scope are retained until the "last active dialog ends". That's what the question asked for.

Ref: Under "Composer v1.x" tab, <https://docs.microsoft.com/en-us/composer/concept-memory?tabs=v2x>

upvoted 4 times

 **ninja** 1 year, 9 months ago

Agreed. A. dialog

upvoted 2 times

Question #15

DRAG DROP -

You have a chatbot that uses a QnA Maker application.

You enable active learning for the knowledge base used by the QnA Maker application.

You need to integrate user input into the model.

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

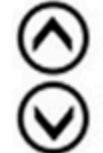
Actions**Answer Area**

Add a task to the Azure resource.

Approve and reject suggestions.

Publish the knowledge base.

Modify the automation task logic app to run an Azure Resource Manager template that creates the Azure Cognitive Services resource.



For the knowledge base, select Show active learning suggestions.

Save and train the knowledge base.

Select the properties of the Azure Cognitive Services resource.

Actions**Answer Area**

Add a task to the Azure resource.

For the knowledge base, select Show active learning suggestions.

Approve and reject suggestions.

Save and train the knowledge base.

Correct Answer:

Modify the automation task logic app to run an Azure Resource Manager template that creates the Azure Cognitive Services resource.



Publish the knowledge base.

Select the properties of the Azure Cognitive Services resource.

Step 1: For the knowledge base, select Show active learning suggestions.

In order to see the suggested questions, on the Edit knowledge base page, select View Options, then select Show active learning suggestions.

Step 2: Approve and reject suggestions.

Each QnA pair suggests the new question alternatives with a check mark, ✓, to accept the question or an x to reject the suggestions. Select the check mark to add the question.

λ"

add the question.

Step 3: Save and train the knowledge base.

Select Save and Train to save the changes to the knowledge base.

Step 4: Publish the knowledge base.

Select Publish to allow the changes to be available from the GenerateAnswer API.

When 5 or more similar queries are clustered, every 30 minutes, QnA Maker suggests the alternate questions for you to accept or reject.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/how-to/improve-knowledge-base>

✉️  **Eltooth** Highly Voted 1 year, 10 months ago

Answer provided is correct.

(Turn on active learning) - Done

1. Show active learning suggestions
2. Accept the question or reject the suggestions
3. Save and Train
4. Publish

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/how-to/improve-knowledge-base#view-suggested-questions>

upvoted 12 times

✉️  **ninja** 1 year, 9 months ago

Agreed.

upvoted 1 times

✉️  **zellck** Highly Voted 11 months ago

1. From knowledge base, select show active learning suggestions.
2. Approve and reject the suggestions.
3. Save and train knowledge base.
4. Publish knowledge base.

<https://learn.microsoft.com/en-us/azure/cognitive-services/qnamaker/how-to/improve-knowledge-base#view-suggested-questions>

- In order to see the suggested questions, on the Edit knowledge base page, select View Options, then select Show active learning suggestions. This option will be disabled if there are no suggestions present for any of the question and answer pairs.

- Each QnA pair suggests the new question alternatives with a check mark, ✓, to accept the question or an x to reject the suggestions. Select the check mark to add the question.

- Select Save and Train to save the changes to the knowledge base.

- Select Publish to allow the changes to be available from the GenerateAnswer API.

upvoted 11 times

✉️  **rdemontis** 6 months, 2 weeks ago

agree. thanks for explanation and references

upvoted 2 times

✉️  **evangelist** Most Recent 3 months, 3 weeks ago

1. Show active learning suggestions
2. Accept the question or reject the suggestions
3. Save and Train
4. Publish

upvoted 1 times

Question #16

You need to enable speech capabilities for a chatbot.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Enable WebSockets for the chatbot app.
- B. Create a Speech service.
- C. Register a Direct Line Speech channel.
- D. Register a Cortana channel.
- E. Enable CORS for the chatbot app.
- F. Create a Language Understanding service.

Correct Answer: ABC

You can use the Speech service to voice-enable a chat bot.

The Direct Line Speech channel uses the text-to-speech service, which has neural and standard voices.

You'll need to make a small configuration change so that your bot can communicate with the Direct Line Speech channel using web sockets.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/tutorial-voice-enable-your-bot-speech-sdk>

Community vote distribution

ABC (86%)

14%

✉ Eltooth **Highly Voted** 1 year, 10 months ago

Selected Answer: ABC

A, B and C are correct answers in order shown below.

B. Create a Speech service

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/tutorial-voice-enable-your-bot-speech-sdk#create-a-speech-service-resource>

A. Enable WebSockets for the chatbot app

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/tutorial-voice-enable-your-bot-speech-sdk#enable-web-sockets>

C. Register a Direct Line Speech channel

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/tutorial-voice-enable-your-bot-speech-sdk#register-the-direct-line-speech-channel>

upvoted 10 times

✉ evangelist **Most Recent** 3 months, 3 weeks ago

Selected Answer: ABC

Please do not choose BCE,

To enable speech capabilities for a chatbot, the correct actions include:

A. Enable WebSockets for the chatbot app: Necessary for the bot to communicate with the Direct Line Speech channel using web sockets.

B. Create a Speech service: Required to provide speech-to-text and text-to-speech capabilities for the chatbot.

C. Register a Direct Line Speech channel: Essential for connecting the chatbot with the Speech service to enable speech interactions.

For detailed steps, you can refer to the official Azure documentation: Tutorial: Voice-enable your bot.

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/tutorial-voice-enable-your-bot-speech-sdk#enable-web-sockets>

upvoted 1 times

✉ rdemontis 6 months, 3 weeks ago

Selected Answer: ABC

<https://learn.microsoft.com/en-us/azure/ai-services/speech-service/tutorial-voice-enable-your-bot-speech-sdk>

upvoted 1 times

✉ zellick 11 months ago

Selected Answer: ABC

ABC is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/tutorial-voice-enable-your-bot-speech-sdk#create-a-speech-service-resource>

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/tutorial-voice-enable-your-bot-speech-sdk#enable-web-sockets>

You need to make a small configuration change so that your bot can communicate with the Direct Line Speech channel by using web sockets.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/tutorial-voice-enable-your-bot-speech-sdk#register-the-direct-line-speech-channel>

This channel creates a connection between your bot and a client app compiled with the Speech SDK.

upvoted 4 times

 **Pfffff** 1 year, 1 month ago

Selected Answer: ABC

said ChatGPT

upvoted 3 times

 **marti_tremblay000** 1 year, 2 months ago

Selected Answer: BCE

According to ChatGPT, the correct answers are BCE :

The three actions that need to be performed to enable speech capabilities for an Azure chatbot are:

B. Create a Speech service: A Speech service is needed to process the audio input and output of the chatbot. Azure provides the Speech service that can be created from the Azure portal.

C. Register a Direct Line Speech channel: A Direct Line Speech channel needs to be registered in the Azure portal for the chatbot app. The Direct Line Speech channel enables the chatbot to receive audio input and provide audio output to the user.

E. Enable CORS for the chatbot app: Cross-Origin Resource Sharing (CORS) needs to be enabled for the chatbot app to allow the browser to access resources from a different domain. This is required when using the Direct Line Speech channel.

Therefore, the correct actions are B, C, and E.

upvoted 1 times

 **Pyguy** 1 year, 4 months ago

Selected Answer: BCE

instead of WebSockets you should enable CORS . Cross-Origin Resource Sharing , will allow your bot app to communicate with the Speech service. Nothing to do websockets here..

upvoted 2 times

 **AzureJobsTillRetire** 1 year, 3 months ago

Can you please share your reference?

Eltooth has provided link as below.

Enable web sockets

You need to make a small configuration change so that your bot can communicate with the Direct Line Speech channel by using web sockets.

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/tutorial-voice-enable-your-bot-speech-sdk#enable-web-sockets>

upvoted 1 times

Question #17

You use the Microsoft Bot Framework Composer to build a chatbot that enables users to purchase items.

You need to ensure that the users can cancel in-progress transactions. The solution must minimize development effort.

What should you add to the bot?

- A. a language generator
- B. a custom event
- C. a dialog trigger
- D. a conversation activity

Correct Answer: D

Handling interruptions is an important aspect of a robust bot. Users will not always follow your defined conversation flow, step by step. They may try to ask a question in the middle of the process, or simply want to cancel it instead of completing it.

Example:

If the user types "cancel", it calls CancelAllDialogsAsync on its inner dialog context, which clears its dialog stack and causes it to exit with a canceled status and no result value. To the MainDialog (shown later on), it will appear that the booking dialog ended and returned null, similar to when the user chooses not to confirm their booking.

```
private async Task<DialogTurnResult> InterruptAsync(DialogContext innerDc, CancellationToken cancellationToken)
```

```
{  
if (innerDc.Context.Activity.Type == ActivityTypes.Message)  
{  
var text = innerDc.Context.Activity.Text.ToLowerInvariant();  
switch (text)  
{  
case "cancel":  
case "quit":  
var cancelMessage = MessageFactory.Text(CancelMsgText, CancelMsgText, InputHints.IgnoringInput); await  
innerDc.Context.SendActivityAsync(cancelMessage, cancellationToken); return await innerDc.CancelAllDialogsAsync(cancellationToken);  
}  
}  
return null;  
}
```

Reference:

<https://docs.microsoft.com/en-us/azure/bot-service/bot-builder-howto-handle-user-interrupt>

Community vote distribution

zellck 11 months ago

Selected Answer: C

C is the answer.

<https://learn.microsoft.com/en-us/composer/concept-events-and-triggers?tabs=v2x>

In Bot Framework Composer, each dialog includes one or more event handlers called triggers. Each trigger contains one or more actions. Actions are the instructions that the bot will execute when the dialog receives any event that it has a trigger defined to handle. Once a given event is handled by a trigger, no further action is taken on that event. Some event handlers have a condition specified that must be met before it will handle the event and if that condition isn't met, the event is passed to the next event handler. If an event isn't handled in a child dialog, it gets passed up to its parent dialog to handle and this continues until it's either handled or reaches the bot's main dialog. If no event handler is found, it will be ignored and no action will be taken.

upvoted 14 times

Davard 1 year, 8 months ago

It seems like it should be C: a dialog trigger. Anyone else?

upvoted 13 times

Tickxit 1 year, 8 months ago

I agree, an interruption occurs when a trigger is fired. Context of adaptive dialogs.

upvoted 6 times

evangelist 3 months, 3 weeks ago

Selected Answer: C

The answer is C:

To ensure that users can cancel in-progress transactions with minimal development effort, you should add C. a dialog trigger to the bot. A dialog trigger allows the bot to respond to specific conditions or intents, such as a user's request to cancel a transaction, by invoking predefined dialogs that handle these scenarios efficiently.

upvoted 2 times

✉ **rdemontis** 6 months, 3 weeks ago

Selected Answer: C

To me the correct answer is C.

<https://learn.microsoft.com/en-us/composer/concept-events-and-triggers?tabs=v2x#dialog-events-triggers>

upvoted 4 times

✉ **propanther** 7 months ago

Selected Answer: C

You should use dialog triggers to:

Take actions immediately when the dialog starts, even before the recognizer is called.

Take actions when a "cancel" signal is detected.

Take actions on messages received or sent.

Evaluate the content of the incoming activity.

<https://learn.microsoft.com/en-us/composer/concept-events-and-triggers?tabs=v2x#dialog-events-triggers>

upvoted 3 times

✉ **sl_msiconsulting** 7 months, 1 week ago

Selected Answer: D

Tiggers are not a good fit for handling scenarios like these let along there are all 4 types of triggers for dialogs. Follow the link provided by the answer and read the code carefully and you will understand why it's D.

upvoted 2 times

✉ **sl_msiconsulting** 7 months, 1 week ago

*only 4 types... Dialog started (Begin dialog event)

Dialog cancelled (Cancel dialog event)

Error occurred (Error event)

Re-prompt for input (Reprompt dialog event). How are they appropriate in handling scenarios like this?

upvoted 2 times

✉ **EliteAllen** 11 months ago

Selected Answer: C

A dialog trigger in Bot Framework Composer allows you to define the conditions under which a specific dialog is started. For example, you can create a dialog trigger that listens for the user to say something like "cancel" or "stop", and then initiate a dialog that handles the cancellation of the in-progress transaction.

This approach doesn't require you to manually code the cancellation logic, and it leverages the built-in capabilities of the Bot Framework Composer, thus minimizing development effort.

upvoted 3 times

✉ **ArchMelody** 1 year, 6 months ago

Selected Answer: C

C seems to be the correct answer for me as well for the aforementioned reasons.

upvoted 5 times

✉ **DiegoGonL** 1 year, 7 months ago

Selected Answer: C

As Tickxit said: an interruption occurs when a trigger is fired in the context of adaptive dialogs.

upvoted 3 times

Question #18**SIMULATION -**

You need to create and publish a bot that will use Language Understanding and QnA Maker. The bot must be named bot12345678. You must publish the bot by using the User1-12345678@abc.com account.

NOTE: Complete this task first. It may take several minutes to complete the required deployment steps. While this is taking place, you can complete tasks 2-6 in this lab during the deployment.

To complete this task, use the Microsoft Bot Framework Composer.

Correct Answer: See explanation below.

Step 1: Sign in to the QnAMaker.ai portal with your Azure credentials. Use the User1-12345678@abc.com account

Step 2: Publish the knowledge base. In the QnA Maker portal, select Publish. Then to confirm, select Publish on the page.

The QnA Maker service is now successfully published. You can use the endpoint in your application or bot code.

Success! Your service has been deployed. What's next?

You can always find the deployment details in your service's settings.

Create Bot

[View all your bots on the Azure Portal.](#)

Use the below HTTP request to call your Knowledgebase. [Learn more.](#)

Postman Curl

```
POST /knowledgebases/ <knowledge-base-ID> /generateAnswer
Host: https://so-15indexes.azurewebsites.net/qnamaker
Authorization: EndpointKey <Authorization-key>
Content-Type: application/json
{"question":<Your question> "}
```

Need to fine-tune and refine? Go back and keep editing your service.

Edit Service

Step 3: In the QnA Maker portal, on the Publish page, select Create bot.

This button appears only after you've published the knowledge base.

After publishing the knowledge base, you can create a bot from the Publish page.

Success! Your service has been deployed. What's next?

You can always find the deployment details in your service's settings.

Create Bot

[View all your bots on the Azure Portal.](#)

Use the below HTTP request to call your Knowledgebase. [Learn more.](#)

Postman Curl

```
POST /knowledgebases/ <knowledge-base-ID> /generateAnswer
Host: https://so-15indexes.azurewebsites.net/qnamaker
Authorization: EndpointKey <Authorization-key>
Content-Type: application/json
{"question":<Your question> "}
```

Need to fine-tune and refine? Go back and keep editing your service.

Edit Service

Step 4: A new browser tab opens for the Azure portal, with the Azure Bot Service's creation page. Configure the Azure bot service.

Bot name: bot12345678 -

The bot will be created.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/quickstarts/create-publish-knowledge-base>

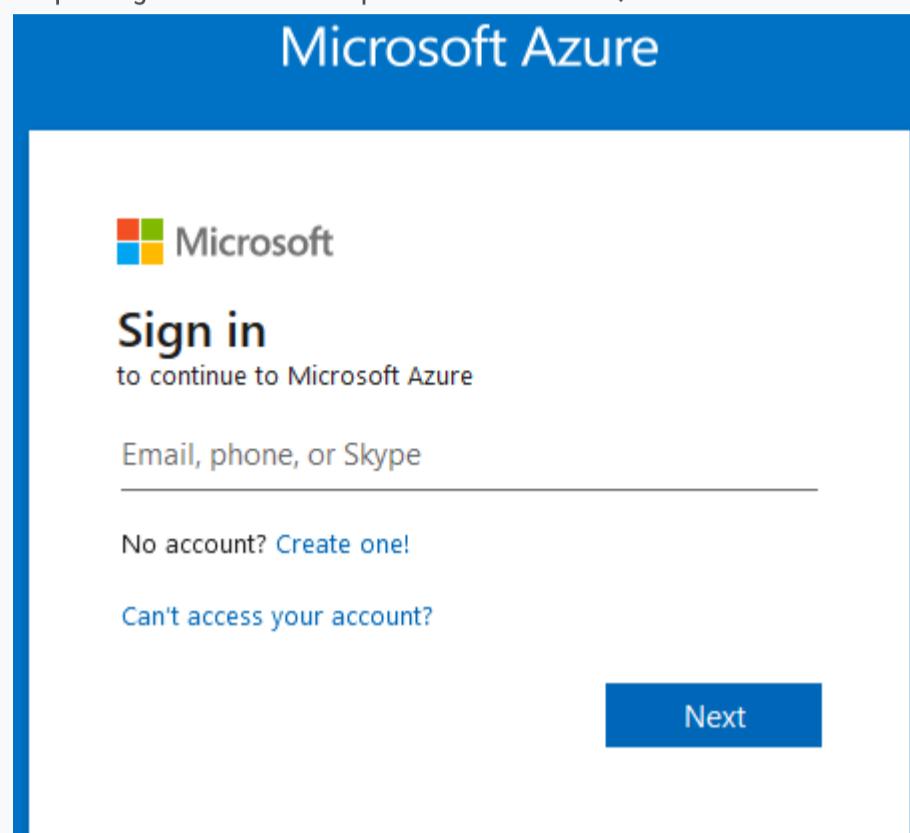
Question #19

SIMULATION -

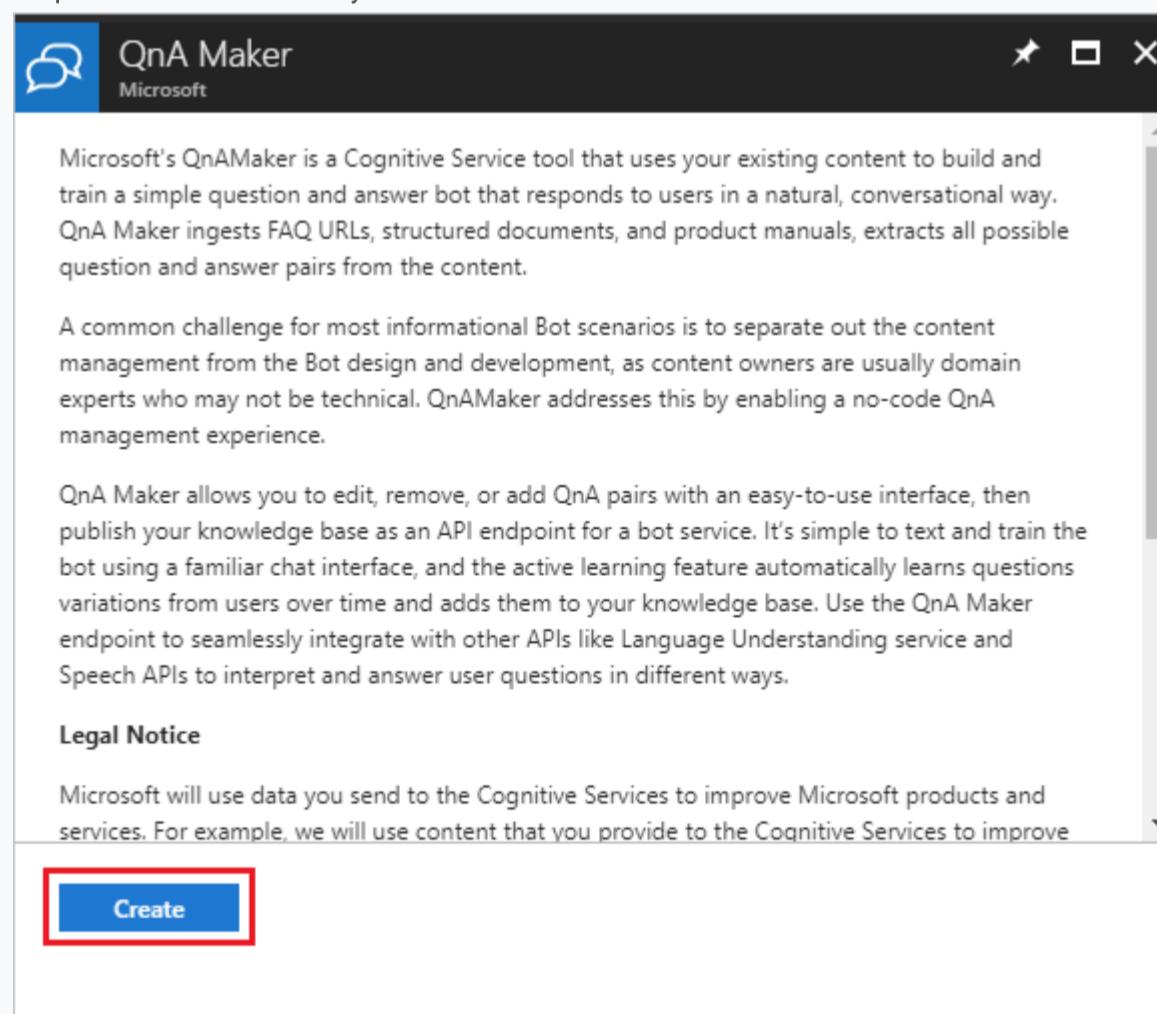
You need to create a QnA Maker service named QNA12345678 in the East US Azure region. QNA12345678 must contain a knowledge base that uses the questions and answers available at <https://support.microsoft.com/en-us/help/12435/windows-10-upgrade-faq>. To complete this task, sign in to the Azure portal and the QnA Maker portal.

Correct Answer: See explanation below.

Step 1: Sign in to the Azure portal create and a QnA Maker resource.



Step 2: Select Create after you read the terms and conditions:



Step 3: In QnA Maker, select the appropriate tiers and regions.

Name: QNA12345678 -

In the Name field, enter a unique name to identify this QnA Maker service. This name also identifies the QnA Maker endpoint that your knowledge bases will be associated with.

Resource Group Location: East US Azure

Create

QnA Maker

* Name
myqnamakerservice

* Subscription
team

* Pricing tier (View full pricing details)
F0 (3 managed documents per month, 3 tr...)

* Resource group
(New) myqnamakerservice

Create new

* Resource group location
(US) Central US

* Search pricing tier (View full pricing details)
B (15 Indexes)

* Search location
West US

* App name
myqnamakerservice

.azurewebsites.net

The App service plan currently defaults to standard(S1) tier. It can be modified by visiting the app service plan resource page once the resource has been created.

* Website location
West US

App insights
Enable **Disable**

Step 4: After all the fields are validated, select Create. The process can take a few minutes to complete.

After deployment is completed, you'll see the following resources created in your subscription:

myqnamakerservice																							
Overview	Subscription (change) My account	Subscription ID <subscription id>																					
Activity log	Tags (change) Click here to add tags	Deployments 2 Succeeded																					
Access control (IAM)																							
Tags																							
SETTINGS																							
Quickstart																							
Resource costs																							
Deployments																							
Policies																							
Properties																							
Locks																							
<input type="checkbox"/> Show hidden types																							
<table border="1"> <thead> <tr> <th>NAME</th> <th>TYPE</th> <th>LOCATION</th> </tr> </thead> <tbody> <tr> <td>Default1</td> <td>App Service plan</td> <td>East US</td> </tr> <tr> <td>Failure Anomalies - myqnamakerservice-ai</td> <td>microsoft.insights/alertrules</td> <td>East US</td> </tr> <tr> <td>myqnamakerservice</td> <td>Search service</td> <td>West US</td> </tr> <tr> <td>myqnamakerservice</td> <td>Cognitive Services</td> <td>West US</td> </tr> <tr> <td>myqnamakerservice</td> <td>App Service</td> <td>East US</td> </tr> <tr> <td>qnamakerpp-ai</td> <td>Application Insights</td> <td>East US</td> </tr> </tbody> </table>			NAME	TYPE	LOCATION	Default1	App Service plan	East US	Failure Anomalies - myqnamakerservice-ai	microsoft.insights/alertrules	East US	myqnamakerservice	Search service	West US	myqnamakerservice	Cognitive Services	West US	myqnamakerservice	App Service	East US	qnamakerpp-ai	Application Insights	East US
NAME	TYPE	LOCATION																					
Default1	App Service plan	East US																					
Failure Anomalies - myqnamakerservice-ai	microsoft.insights/alertrules	East US																					
myqnamakerservice	Search service	West US																					
myqnamakerservice	Cognitive Services	West US																					
myqnamakerservice	App Service	East US																					
qnamakerpp-ai	Application Insights	East US																					

Remember your Azure Active Directory ID, Subscription, QnA resource name you selected when you created the resource.

Step 5: When you are done creating the resource in the Azure portal, return to the QnA Maker portal, refresh the browser page.

Step 6: In the QnA Maker portal, select Create a knowledge base.

Step 7: Skip Step 1 as you already have your QnA Maker resource.

Step 8: In Step 2, select your Active directory, subscription, service (resource), and the language for all knowledge bases created in the service.

Azure QnA service: QNA12345678 -

STEP 2**Connect your QnA service to your KB.**

After you create an Azure QnA service, refresh this page and then select your Azure service using the options below

Refresh*** Microsoft Azure Directory ID**

Microsoft

*** Azure subscription name**

documentationteam

*** Azure QnA service**

qna-maker-10

*** Language**

English

Step 9: In Step 3, name your knowledge base

Step 10: In Step 4, configure the following setting:

+ Add URL: <https://support.microsoft.com/en-us/help/12435/windows-10-upgrade-faq>

Step 11: In Step 5, Select Create your KB.

The extraction process takes a few moments to read the document and identify questions and answers.

After QnA Maker successfully creates the knowledge base, the Knowledge base page opens.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/how-to/set-up-qnamaker-service-azure> <https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/quickstarts/create-publish-knowledge-base>**rewanotora** 1 day, 22 hours ago

Simulation question will not appear on the actual exam as of May 26, 2024.

upvoted 1 times

takaimomoGcup 5 days, 18 hours ago

If the simulation question is not on the real exam, delete this question and the related simulation question.

upvoted 1 times

Ghill1982 2 months, 2 weeks ago

The QnA Maker service is being retired and you can no longer create a new QnA Marker resource. Instead as of March 2024:

1. Create a Language service resource in Azure AI services with Custom question answering selected
2. Launch Language Studio, select Custom question answering
3. Create a new Project and add the provided URL as a data source

upvoted 1 times

rdemontis 6 months, 2 weeks ago

Deprecated service.

1. Create the Question and answering service on the azure portal with the information provided in the question (name and location)
2. Create a new project on the Language Studio portal
3. Choose in the project the Question and Answering resource created previously
4. Click on the project
5. + Add source
6. Choose URL
7. Enter the provided url in the question
8. Add

upvoted 4 times

Question #20

SIMULATION -

You need to add a question pair to the published knowledge base used by a QnA Maker service named QNA12345678. The question must be: 'What will be the next version of Windows?'
The answer must be: 'Windows 11'.
To complete this task, sign in to the QnA Maker portal.

Correct Answer: Answer: Windows 11

Step 1: Sign in to the QnA portal, then select the knowledge base to add the QnA pair to.

Step 2: On the EDIT page of the knowledge base, select Add QnA pair to add a new QnA pair.

Knowledge base

The screenshot shows the Microsoft QnA Maker knowledge base edit page. At the top, there is a search bar, a count of 81 QnA pairs, and two buttons: 'Add QnA pair' (which is highlighted with a red box) and 'View options'. Below the header, there is a navigation bar with page numbers 1 through 9 and 'Next >'. The main content area displays a table with two columns: 'Question' and 'Answer'. The first row shows a question 'How do I get the Surface Pro repaired?' with an answer 'Repairing the Surface Pro requires...'. Below the table, there are buttons for 'Add alternative phrasing' and 'Add follow-up prompt'.

Step 3: In the new QnA pair row, add the required question and answer fields. The other fields are optional. All fields can be changed at any time.

Question: What will be the next version of Windows?

Step 4: Select Save and train to see predictions including the new QnA pair.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/how-to/edit-knowledge-base>

✉ **reiwanotora** 1 day, 22 hours ago

Simulation question will not appear on the actual exam as of May 26 2024.

upvoted 1 times

✉ **AlviraTony** 2 months, 1 week ago

QnAMaker is no longer in the syllabus of Exam AI-102

upvoted 1 times

✉ **AnonymousJhb** 4 months ago

The given answer is deprecated.

This question follows on from pre-built question 5.19,

go back to your project QNA12345678 in language studio > edit knowledge base > + Add >

select the url as source > add your given q and a from the question > done > save >

edit knowledge base > click on test icon > select use deployed kb > paste your new question > What will be the next version of Windows? > scroll down to find a bold font Windows 11 answer

All done :)

upvoted 2 times

✉ **rdemontis** 6 months, 3 weeks ago

Now you need to use Language Studio <https://language.cognitive.azure.com>

upvoted 3 times

Question #21

SIMULATION -

Use the following login credentials as needed:

To enter your username, place your cursor in the Sign in box and click on the username below.

To enter your password, place your cursor in the Enter password box and click on the password below.

Azure Username: admin@abc.com -

Azure Password: XXXXXXXXXXXX -

The following information is for technical support purposes only:

Lab Instance: 12345678 -

Task -

You have a bot that was developed by using the Microsoft Bot Framework SDK. The bot is available at an endpoint of <https://bot.contoso.com/api/messages>.

You need to create an Azure Bot named bot12345678 that connects to the bot.

To complete this task, sign in to the Azure portal.

Correct Answer: See explanation below.

Create the resource -

Create the Azure Bot resource, which will allow you to register your bot with the Azure Bot Service.

1. Go to the Azure portal.
2. In the right pane, select Create a resource.
3. In the search box enter bot, then press Enter.
4. Select the Azure Bot card.



Azure Bot

Microsoft

Azure Service

Build enterprise-grade conversational
AI experiences with Bot Framework
Composer or SDK.

Create

5. Select Create.
6. Enter values in the required fields. Choose which type of app to create and whether to use existing or create new identity information.

Pricing

Select a pricing tier for your Azure Bot resource. You can change your selection later in the Azure portal's resource management. Learn more about available options, or request a pricing quote, by visiting the [Azure Bot Services pricing](#)

Pricing tier *

Standard
[Change plan](#)

Microsoft App ID

A Microsoft App ID is required to create an Azure Bot resource. If your bot app doesn't need to access resources outside of its home tenant and if your bot app will be hosted on an Azure resource that supports Managed Identities, then choose option User-Assigned Managed Identity so that Azure takes care of managing the App credentials for you. Otherwise, depending on whether your bot will be accessing resources only in its home tenant or not, choose either Single tenant or Multi tenant option respectively.

Type of App

User-Assigned Managed Identity

i Note: For User-Assigned Managed Identity and Single Tenant app, Azure Portal's "Open in Composer" link is not yet supported for bots with these app types. BotFramework SDK (C# or Javascript) version 4.15.0 or higher is needed for these app types.

A User-assigned managed identity can be automatically created below or you can manually create your own, then return to input your new App ID, tenant ID and MSI resource ID in the open fields.

[Manually create a User Managed Identity](#)

Creation type

Create new Microsoft App ID
 Use existing app registration

7. Select Review + create.

8. If the validation passes, select Create.

9. Once the deployment completes, select Go to resource. You should see the bot and related resources listed in the resource group you selected.

10. Enter the endpoint of the Bot Framework SDK: <https://bot.contoso.com/api/messages>

Reference:

<https://docs.microsoft.com/en-us/azure/bot-service/abs-quickstart?view=azure-bot-service-4.0&tabs=userassigned>

Question #22

You are designing a conversational interface for an app that will be used to make vacation requests. The interface must gather the following data:

- The start date of a vacation
- The end date of a vacation
- The amount of required paid time off

The solution must minimize dialog complexity.

Which type of dialog should you use?

- A. adaptive
- B. skill
- C. waterfall
- D. component

Correct Answer: D

Community vote distribution

C (95%) 5%

✉️  **evangelist**  3 months, 3 weeks ago

Selected Answer: C

The answer is C:

For designing a conversational interface that gathers data in a structured manner with minimal dialog complexity, you should use C. waterfall dialog. Waterfall dialogs are designed to guide the user through a series of steps or questions, one after another, which is ideal for collecting specific pieces of information sequentially, like vacation start and end dates, and the amount of required paid time off.

upvoted 5 times

✉️  **elizbeth2**  3 months, 1 week ago

Selected Answer: C

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-waterfall-dialogs?view=azure-bot-service-4.0#waterfall-dialogs>
upvoted 2 times

✉️  **trysec** 3 months, 4 weeks ago

Selected Answer: C

C. waterfall

A waterfall dialog is best suited for scenarios where a conversational flow requires a sequence of steps to collect information or guide the user through a set process. For gathering specific pieces of data like start date, end date, and the amount of required paid time off, a waterfall dialog allows you to prompt the user for each piece of information one after the other, which can help minimize the complexity of the dialog interaction:
upvoted 2 times

✉️  **dimsok** 4 months, 3 weeks ago

Selected Answer: C

Opinionated, but waterfall seems like a better UX than component

upvoted 1 times

✉️  **rdemontis** 6 months, 3 weeks ago

Selected Answer: C

I think correct answer is C.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-waterfall-dialogs?view=azure-bot-service-4.0#waterfall-dialogs>
upvoted 2 times

✉️  **sl_msiconsulting** 8 months, 2 weeks ago

Selected Answer: D

Should be D. You need to take a step further to reduce the complexity - <https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-compositcontrol?view=azure-bot-service-4.0&tabs=csharp>
upvoted 1 times

✉️  **ExamDev** 8 months, 3 weeks ago

Selected Answer: C

It's waterfall

upvoted 2 times

✉ M25 8 months, 3 weeks ago

Selected Answer: C

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-waterfall-dialogs?view=azure-bot-service-4.0>

The interface must gather ... date: Prompt dialogs (prompts) are dialogs designed to ask the user for specific types of information, such as a number, a date, or a name, and so on. Prompts are designed to work with waterfall dialogs in a component dialog.

A waterfall dialog (or waterfall) ... are designed to work within the context of a component dialog.

The component dialog provides a strategy for creating independent dialogs to handle specific scenarios, breaking a large dialog set into more manageable pieces. Each of these pieces has its own dialog set, and avoids any name collisions with the dialog set that contains it.

upvoted 3 times

✉ M25 8 months, 3 weeks ago

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-dialog?view=azure-bot-service-4.0#dialog-types>

adaptive dialog: A type of container dialog used by Composer to provide more natural conversational flows. Not intended to be used directly in an SDK-first bot.

skill dialog: Automates the management of one or more skill bots from a skill consumer. Composer directly supports skills as actions.

Important

Adaptive dialogs were first added in version 4.9 of the C# SDK. Adaptive dialogs support the Bot Framework Composer and are not intended to be used directly in an SDK-first bot.

Bot Framework Composer is a visual authoring tool for building Conversational AI applications.

upvoted 1 times

✉ tranatrana 9 months ago

What is it? C or D?

upvoted 1 times

✉ kail85 11 months, 4 weeks ago

C. waterfall

Waterfall dialogs are a good choice for this scenario because they are designed to guide users through a linear, step-by-step process. In this case, you can use a waterfall dialog to gather the vacation start date, end date, and required paid time off sequentially. This approach can help minimize dialog complexity while ensuring that all necessary data is collected in an orderly manner.

upvoted 1 times

✉ EliteAllen 12 months ago

Selected Answer: C

C. waterfall

A waterfall dialog is a type of dialog that guides the user through a series of steps or prompts in a specific order. This makes it a good choice for gathering a set of related data points, like the start date, end date, and amount of required paid time off for a vacation request. By using a waterfall dialog, you can ensure that all necessary information is collected in a structured and predictable manner, which can help minimize dialog complexity.

Option A, adaptive dialog, is a more flexible type of dialog that can handle more complex and dynamic conversation flows, but it might be overkill for this relatively straightforward data-gathering task. Option B, skill dialog, is used to manage the invocation of a bot skill, which is not relevant to this scenario. Option D, component dialog, is a reusable dialog that encapsulates its own state, but it doesn't inherently simplify the dialog structure.

upvoted 1 times

✉ odisor 1 year, 1 month ago

I think the answer could be C or D, but in order to minimize complexity, I think D is correct.

ChatGPT answer:

In Azure CLU, there are three main types of dialogs:

Root Dialog: The Root Dialog is the main entry point for the bot. It handles incoming messages and dispatches them to other dialogs as appropriate. The Root Dialog can be thought of as the top-level container for all other dialogs in the bot.

Component Dialog: A Component Dialog is a reusable dialog that can be embedded within other dialogs. It encapsulates a specific set of related functionality, such as a login dialog or a payment dialog. Component Dialogs can be used to create more modular and maintainable bot code.

Waterfall Dialog: A Waterfall Dialog is a type of dialog that guides the user through a series of steps or prompts to complete a task. Each step in the Waterfall Dialog can be implemented as a separate method, making it easy to create complex conversational flows. Waterfall Dialogs can be used for tasks like booking appointments, ordering food, or making reservations.

upvoted 3 times

✉ ptmk 1 year, 1 month ago

I think the answer is correct. The key requirement is "The solution must minimize dialog complexity".

My answer is C. Component

Waterfall dialog is used to manage linear and more complex conversation flows.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-dialog-manage-conversation-flow?view=azure-bot-service-4.0&tabs=csharp>

Component dialog is used to create independent dialogs to handle specific scenarios, breaking a large dialog set into more manageable pieces.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-compositcontrol?view=azure-bot-service-4.0&tabs=csharp>

upvoted 1 times

 **marti_tremblay000** 1 year, 2 months ago

Selected Answer: C

ChatGPT answer :

To minimize dialog complexity and gather the required data in a conversational interface for vacation requests, the best type of dialog to use would be a "waterfall" dialog.

A waterfall dialog is a simple, linear conversation flow where the bot prompts the user for information in a structured sequence. Each prompt is presented to the user one at a time, in a predefined order, until all the required information is collected. The waterfall dialog makes it easy to collect data in a conversational interface, while keeping the dialog simple and easy to follow.

In this case, the waterfall dialog can be designed to prompt the user for the start date, end date, and amount of required paid time off, in that order. The dialog can then confirm the details with the user and submit the request.

Skill and component dialogs are not relevant to this scenario, as they are used for integrating with external services and breaking down large bots into smaller, reusable parts, respectively.

upvoted 3 times

Question #23

DRAG DROP

You build a bot by using the Microsoft Bot Framework SDK.

You need to test the bot interactively on a local machine.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

NOTE: More than one order of answer choices is correct. You will receive credit for any of the correct orders you select.

Actions	Answer Area
Open the Bot Framework Composer.	1
Connect to the bot endpoint.	2
Register the bot with the Azure Bot Service.	3
Build and run the bot.	
Open the Bot Framework Emulator.	

Correct Answer:

Answer Area
1 Build and run the bot.
2 Open the Bot Framework Emulator.
3 Connect to the bot endpoint.

 **zellck** Highly Voted 11 months ago

1. Build and run the bot
2. Open Bot Framework Emulator
3. Connect to bot endpoint

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-debug-emulator?view=azure-bot-service-4.0&tabs=csharp#run-a-bot-locally>
Before connecting your bot to the Bot Framework Emulator, you need to run your bot locally.

upvoted 15 times

 **kail85** Highly Voted 11 months, 4 weeks ago

Answer is correct.

Build and run the bot
Open the Bot Framework Emulator
Connect to the bot endpoint
upvoted 5 times

 **evangelist** Most Recent 3 months, 3 weeks ago

Given answer is correct and below is why:

- 1: Build and run the bot: This step involves compiling the bot's code and starting it on your local machine to make it ready to accept connections.
- 2: Open the Bot Framework Emulator: The Emulator is a desktop application that allows developers to test and debug their bots on localhost or remotely through a tunnel.
- 3: Connect to the bot endpoint: Within the Bot Framework Emulator, you would connect to the bot's endpoint, usually something like <http://localhost:3978/api/messages>, to start interacting with your bot.

upvoted 4 times

 **rdemontis** 6 months, 3 weeks ago

I think the answer is correct. Because more than one order could be possible in this case I suppose another solution could be ordering the tasks inverting Build and Run with Open Bot Emulator:

1. Open Bot Framework Emulator
 2. Build and run the bot
 3. Connect to bot endpoint
- upvoted 2 times

Question #24

You create a bot by using the Microsoft Bot Framework SDK.

You need to configure the bot to respond to events by using custom text responses.

What should you use?

- A. a dialog
- B. an activity handler
- C. an adaptive card
- D. a skill

Correct Answer: B

Community vote distribution

B (100%)

 **marti_tremblay000**  1 year, 2 months ago

Selected Answer: B

B is the correct answer as explained by ChatGPT :

B. An activity handler is the correct choice for configuring the bot to respond to events by using custom text responses.

An activity handler is a class in the Bot Framework SDK that processes incoming activities (e.g., messages, events, etc.) from the user and generates outgoing activities (e.g., replies). By overriding the OnMessageActivityAsync method of the activity handler, you can provide custom logic for responding to user messages.

upvoted 5 times

 **evangelist**  3 months, 3 weeks ago

Selected Answer: B

To configure the bot to respond to events using custom text responses, you should use B. an activity handler. Activity handlers allow you to manage and respond to different types of activities and events in a bot's conversation, such as messages from users, with customized logic and responses.

upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

correct

upvoted 2 times

 **Tin_Tin** 11 months ago

Selected Answer: B

correct

<https://learn.microsoft.com/en-us/azure/bot-service/bot-activity-handler-concept?view=azure-bot-service-4.0&tabs=csharp>

upvoted 3 times

Question #25

HOTSPOT

You build a bot named app1 by using the Microsoft Bot Framework.

You prepare app1 for deployment.

You need to deploy app1 to Azure.

How should you complete the command? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

james2033 Highly Voted 9 months, 2 weeks ago

webapp

config-zip

See command

az webapp deployment source config-zip --resource-group "<resource-group-name>" --name "<name-of-app-service>" --src "<project-zip-path>"

at <https://learn.microsoft.com/en-us/azure/bot-service/provision-and-publish-a-bot?view=azure-bot-service-4.0&tabs=userassigned%2Csharp#publish-your-bot-to-azure>

upvoted 11 times

AnonymousJhb 7 months, 2 weeks ago

answer is correct, if you agree,simply click the like button.

upvoted 5 times

rdemontis 6 months, 3 weeks ago

thanks for the provided reference

upvoted 3 times

zellck Highly Voted 11 months ago

1. webapp
2. config-zip

<https://learn.microsoft.com/en-us/azure/bot-service/provision-and-publish-a-bot?view=azure-bot-service-4.0&tabs=userassigned%2Csharp#publish-your-bot-to-azure>

upvoted 6 times

evangelist Most Recent 3 months, 3 weeks ago

To deploy app1 to Azure using the Microsoft Bot Framework, the correct command components are az webapp deployment source config-zip --. This command specifies that you're deploying a web app to Azure, using a zip file as the source of the deployment. The config-zip argument is used to deploy the application from a zipped source file.

```
app1/
  ├── .env (or appsettings.json for configuration settings)
  ├── index.js (entry point of your bot)
  ├── package.json (lists package dependencies)
  ├── node_modules/ (dependencies, can be excluded if using Azure for installation)
  └── dialogs/ (or any other directories with additional bot logic)
```

upvoted 2 times

✉ **Tin_Tin** 11 months, 1 week ago

the answer seems correct

<https://learn.microsoft.com/en-us/azure/bot-service/provision-and-publish-a-bot?view=azure-bot-service-4.0&tabs=userassigned%2Cpython#publish-your-bot-to-azure>

upvoted 3 times

✉ **[Removed]** 11 months, 2 weeks ago

I did not find such command as "az bot deployment" there is only "az webapp deployment"

az webapp deployment source config-zip --resource-group "<resource-group-name>" --name "<name-of-app-service>" --src "<project-zip-path>"

az bot command is to manage a bot

<https://learn.microsoft.com/en-us/cli/azure/bot?view=azure-cli-latest>

upvoted 2 times

✉ **kail85** 12 months ago

first one should be bot

upvoted 3 times

✉ **odisor** 1 year, 1 month ago

the answers is correct, this is asking how to deploy the app for the bot , not the bot resource:

<https://learn.microsoft.com/en-us/azure/bot-service/provision-and-publish-a-bot?view=azure-bot-service-4.0&tabs=userassigned%2Cpython#publish-your-bot-to-azure>

to deploy a bot resource:

<https://learn.microsoft.com/en-us/azure/bot-service/provision-azure-bot?view=azure-bot-service-4.0&tabs=userassigned>

upvoted 3 times

✉ **MaliSanFuu** 1 year ago

i think its a bit tricky with the wording because the bot is named "app1", therefore i think bot should be correct

upvoted 3 times

✉ **Sachz88** 1 year, 1 month ago

The first box is incorrect. It should be "bot"

ChatGPT:

If you want to build a bot named app1 using the Microsoft Bot Framework, the deployment source you should use depends on your specific requirements and preferences.

If you want to deploy your bot on Microsoft Azure, you can use the "azbot" deployment source. This deployment source provides a pre-configure Azure Bot Service that allows you to easily deploy your bot to the cloud.

On the other hand, if you want to deploy your bot on a web server or hosting platform, you can use the "webapp" deployment source. This deployment source allows you to deploy your bot as a web application on a web server, using a programming language of your choice.

The question clearly states: "You need to deploy app1 to Azure."

Answer to the first box is "bot"

upvoted 2 times

Question #26

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have a chatbot that uses question answering in Azure Cognitive Service for Language.

Users report that the responses of the chatbot lack formality when answering spurious questions.

You need to ensure that the chatbot provides formal responses to spurious questions.

Solution: From Language Studio, you change the chitchat source to qna_chitchat_friendly.tsv, and then retrain and republish the model.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

Community vote distribution

B (83%)

A (17%)

 **9H3zmT6** 3 weeks, 3 days ago

Selected Answer: B

This file is used to provide responses suitable for a friendly or casual tone. You need to use qna_chitchat_professional.tsv.
upvoted 1 times

 **rober13** 2 months, 1 week ago

Selected Answer: A

Based on reference Professional = formal and Question 28.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/question-answering/how-to/chit-chat>
upvoted 2 times

 **evangelist** 3 months, 3 weeks ago

Selected Answer: B

The qna_chitchat_professional.tsv is specifically designed to provide responses with a formal tone compared to other chitchat sources. It contains curated set of question-and-answer pairs that are structured to reflect a professional and formal manner of communication. This makes it more suitable for environments or user interactions where formal responses are preferred, ensuring that the chatbot's replies to spurious or off-topic questions maintain the desired level of professionalism.

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

correct

upvoted 1 times

 **zellick** 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/question-answering/how-to/chit-chat>
upvoted 4 times

 **zellick** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 4 times

 **marti_tremblay000** 1 year, 2 months ago

Selected Answer: B

ChatGPT answer :

No, changing the chitchat source to "qna_chitchat_friendly.tsv" and retraining and republishing the model will not necessarily meet the goal of ensuring that the chatbot provides formal responses to spurious questions.

The "qna_chitchat_friendly.tsv" file is a source file for casual chitchat, which includes conversational responses for informal topics, such as hobbies and movies. This file is not designed to provide formal responses to spurious questions. Therefore, changing the source file to this file and retraining the model will not necessarily improve the formality of responses to spurious questions.

upvoted 2 times

Question #27

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have a chatbot that uses question answering in Azure Cognitive Service for Language.

Users report that the responses of the chatbot lack formality when answering spurious questions.

You need to ensure that the chatbot provides formal responses to spurious questions.

Solution: From Language Studio, you modify the question and answer pairs for the custom intents, and then retrain and republish the model.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

Community vote distribution



EliteAllen Highly Voted 11 months ago

Selected Answer: B

No, this solution does not meet the goal.

The formality of responses to spurious questions is not controlled by modifying the question and answer pairs for the custom intents in Language Studio. These pairs are used to train the model to understand and respond to specific intents, not to control the tone or formality of the response.

To ensure that the chatbot provides formal responses to spurious questions, you would need to adjust the chatbot's response templates or scripts, not the question and answer pairs for the custom intents. This might involve programming the chatbot to use more formal language in its responses, or to respond to unrecognized or spurious inputs with a standard, formal message.

upvoted 6 times

evangelist Most Recent 3 months, 3 weeks ago

Selected Answer: B

The qna_chitchat_professional.tsv is specifically designed to provide responses with a formal tone compared to other chitchat sources. It contains a curated set of question-and-answer pairs that are structured to reflect a professional and formal manner of communication. This makes it more suitable for environments or user interactions where formal responses are preferred, ensuring that the chatbot's replies to spurious or off-topic questions maintain the desired level of professionalism.

upvoted 1 times

rdemontis 6 months, 3 weeks ago

Selected Answer: B

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/question-answering/how-to/chit-chat>

upvoted 2 times

zellick 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/question-answering/how-to/chit-chat>

upvoted 3 times

zellick 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 4 times

Drummer 11 months, 1 week ago

No, this solution does not meet the goal. Modifying the question and answer pairs for custom intents in Language Studio and retraining and republishing the model would only affect the responses of the chatbot to specific intents. B NO

upvoted 3 times

✉ **ziggy1117** 11 months, 3 weeks ago

Selected Answer: B

Spurious questions are questions that are not relevant to the conversation or are not meant to be answered. In other words, these are chit-chat. So to make it professional, choose a professional chitchat

upvoted 4 times

✉ **kail85** 12 months ago

Selected Answer: B

you should use the Azure Cognitive Service for Language's Language Understanding (LUIS) feature to identify the intent of the user's input and filter out spurious questions or provide formal fallback responses. You can achieve this by training your LUIS model to recognize spurious questions as a separate intent and handle them accordingly in your chatbot's logic.

upvoted 2 times

✉ **marti_tremblay000** 1 year, 2 months ago

Selected Answer: A

ChatGPT answer :

Yes, modifying the question and answer pairs for the custom intents in Language Studio, and then retraining and republishing the model can meet the goal of ensuring that the chatbot provides formal responses to spurious questions.

The question and answer pairs in Language Studio are used to train the machine learning model for the chatbot. By modifying these pairs to provide more formal responses to spurious questions, and then retraining and republishing the model, the chatbot will have a better understanding of the user's input and will be able to provide more formal responses.

upvoted 2 times

Question #28

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have a chatbot that uses question answering in Azure Cognitive Service for Language.

Users report that the responses of the chatbot lack formality when answering spurious questions.

You need to ensure that the chatbot provides formal responses to spurious questions.

Solution: From Language Studio, you change the chitchat source to qna_chitchat_professional.tsv, and then retrain and republish the model.

Does this meet the goal?

A. Yes

B. No

Correct Answer: A

Community vote distribution

A (100%)

✉️ **evangelist** 3 months, 3 weeks ago

Selected Answer: A

The qna_chitchat_professional.tsv is specifically designed to provide responses with a formal tone compared to other chitchat sources. It contains a curated set of question-and-answer pairs that are structured to reflect a professional and formal manner of communication. This makes it more suitable for environments or user interactions where formal responses are preferred, ensuring that the chatbot's replies to spurious or off-topic questions maintain the desired level of professionalism.

upvoted 2 times

✉️ **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

Correct
<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/question-answering/how-to/chit-chat>
upvoted 2 times

✉️ **james2033** 9 months, 1 week ago

Selected Answer: A

Formal response, it is Personality = Professional, it is file qna_chitchat_professional.tsv . <https://learn.microsoft.com/en-us/azure/ai-services/language-service/question-answering/concepts/best-practices#choosing-a-personality>
upvoted 4 times

✉️ **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/question-answering/how-to/chit-chat>
upvoted 4 times

✉️ **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 4 times

✉️ **marti_tremblay000** 1 year, 2 months ago

Selected Answer: A

ChatGPT answer :

Yes, changing the chitchat source to "qna_chitchat_professional.tsv" and retraining and republishing the model can meet the goal of ensuring that the chatbot provides formal responses to spurious questions.

The "qna_chitchat_professional.tsv" file is a source file for professional chitchat, which includes conversational responses for formal topics, such as business and finance. This file is designed to provide more formal responses to chitchat questions, so changing the source file to this file and retraining the model can improve the formality of responses to spurious questions.

upvoted 2 times

Question #29

You create five bots by using Microsoft Bot Framework Composer.

You need to make a single bot available to users that combines the bots. The solution must support dynamic routing to the bots based on user input.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Create a composer extension.
- B. Change the Recognizer/Dispatch type.
- C. Create an Orchestrator model.
- D. Enable WebSockets.
- E. Create a custom recognizer JSON file.
- F. Install the Orchestrator package.

Correct Answer: BCF

Community vote distribution

BCF (100%)

✉️ **evangelist** 3 months, 3 weeks ago

Selected Answer: BCF

B. Change the Recognizer/Dispatch type: Adjusting the recognizer type to one that supports dispatching allows the bot to understand and route user input to the appropriate sub-bot.

=====

C. Create an Orchestrator model: The Orchestrator is an AI-based component that helps in making decisions on routing the conversation to the most appropriate bot based on the user's intent.

=====

F. Install the Orchestrator package: This is necessary to use the Orchestrator model within your bot, enabling it to dynamically route conversations to the correct bot based on the context and content of the user input.

upvoted 3 times

✉️ **rdemontis** 6 months, 3 weeks ago

Selected Answer: BCF

The answer is correct

upvoted 2 times

✉️ **zellck** 11 months ago

Selected Answer: BCF

BCF is the answer.

<https://learn.microsoft.com/en-us/composer/how-to-create-orchestrator-bot>

upvoted 3 times

✉️ **zellck** 11 months ago

<https://learn.microsoft.com/en-us/composer/how-to-create-orchestrator-bot#add-orchestrator-in-package-manager>
Orchestrator isn't a default feature in Composer. To use Orchestrator, first add the Orchestrator package to your bot. Composer will download the latest Orchestrator model and make it available to your bot.

On the Choose a recognizer type page, select Orchestrator (Intents only) and select Done to change the recognizer.

upvoted 5 times

✉️ **marti_tremblay000** 1 year, 2 months ago

Selected Answer: BCF

ChatGPT answers :

To make a single bot available to users that combines the bots and supports dynamic routing to the bots based on user input, the following three actions should be performed:

B. Change the Recognizer/Dispatch type: The Recognizer/Dispatch type should be changed to enable the bot to recognize user input and dispatch it to the appropriate bot.

C. Create an Orchestrator model: An Orchestrator model should be created to handle the routing of user input to the appropriate bot.

F. Install the Orchestrator package: The Orchestrator package should be installed to provide the bot with the necessary functionality to route user input to the appropriate bot.

Therefore, options B, C, and F are the correct answers.

upvoted 3 times

Question #30

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You are building a chatbot that will use question answering in Azure Cognitive Service for Language.

You have a PDF named Doc1.pdf that contains a product catalogue and a price list.

You upload Doc1.pdf and train the model.

During testing, users report that the chatbot responds correctly to the following question: What is the price of ?

The chatbot fails to respond to the following question: How much does cost?

You need to ensure that the chatbot responds correctly to both questions.

Solution: From Language Studio, you add alternative phrasing to the question and answer pair, and then retrain and republish the model.

Does this meet the goal?

A. Yes

B. No

Correct Answer: A

Community vote distribution

A (100%)

 **evangelist** 3 months, 3 weeks ago

Selected Answer: A

Yes

adding alternative phrasing to the question and answer pair in Azure Cognitive Service for Language, then retraining and republishing the model, meets the goal. This approach allows the model to understand and respond to variations in how questions are asked, improving its ability to correctly answer questions with similar meanings but different wordings.

upvoted 2 times

 **Gvalli** 5 months, 1 week ago

A modified version of this was in the exam today.

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

correct

<https://learn.microsoft.com/en-us/azure/ai-services/language-service/question-answering/concepts/best-practices#when-should-you-add-alternate-questions-to-question-and-answer-pairs>

upvoted 2 times

 **james2033** 9 months, 2 weeks ago

Selected Answer: A

"Question: What is the price of Microsoft Stock?"

"How much does a Microsoft share cost?"

at <https://learn.microsoft.com/en-us/azure/ai-services/language-service/question-answering/concepts/best-practices#when-should-you-add-alternate-questions-to-question-and-answer-pairs>

upvoted 3 times

 **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/question-answering/concepts/best-practices#when-should-you-add-alternate-questions-to-question-and-answer-pairs>

upvoted 4 times

 **973b658** 11 months, 3 weeks ago

Selected Answer: A

A from ChatGPT 3.5.

upvoted 4 times

Question #31

Topic 5

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You are building a chatbot that will use question answering in Azure Cognitive Service for Language.

You have a PDF named Doc1.pdf that contains a product catalogue and a price list.

You upload Doc1.pdf and train the model.

During testing, users report that the chatbot responds correctly to the following question: What is the price of ?

The chatbot fails to respond to the following question: How much does cost?

You need to ensure that the chatbot responds correctly to both questions.

Solution: From Language Studio, you enable chit-chat, and then retrain and republish the model.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

Community vote distribution

B (100%)

 **Murtuza** 1 month, 2 weeks ago

Selected Answer: B

clearly B, ChitChat is for pleasantries

upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

clearly B, ChitChat is for pleasantries

upvoted 3 times

 **james2033** 9 months, 1 week ago

Selected Answer: B

PDF file use for training, it is not rationale.

upvoted 1 times

Question #32

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You are building a chatbot that will use question answering in Azure Cognitive Service for Language.

You have a PDF named Doc1.pdf that contains a product catalogue and a price list.

You upload Doc1.pdf and train the model.

During testing, users report that the chatbot responds correctly to the following question: What is the price of ?

The chatbot fails to respond to the following question: How much does cost?

You need to ensure that the chatbot responds correctly to both questions.

Solution: From Language Studio, you create an entity for price, and then retrain and republish the model.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

Community vote distribution

B (100%)

 **Murtuza** 1 month, 2 weeks ago

Selected Answer: B

B is the answer.

upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

here the issue is not the price that could be an entity, but very different utterances that need to be considered

upvoted 4 times

 **zellck** 11 months ago

Selected Answer: B

B is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/question-answering/concepts/best-practices#when-should-you-add-alternate-questions-to-question-and-answer-pairs>

upvoted 3 times

 **kail85** 11 months, 4 weeks ago

B. No

Creating an entity for price and retraining the model in Language Studio is not the correct approach to solve the issue with Azure Cognitive Service for Language's question-answering capabilities.

Instead, you should use Language Studio to create and train a synonym for the term "price" or build a more comprehensive list of question variations that the chatbot should be able to handle. For example, you can include phrases like "How much does it cost?" or "What is the cost of?" to ensure the model can properly recognize and respond to different ways users might ask about the price. Retrain and republish the model after making these changes to improve the chatbot's ability to answer both questions correctly.

upvoted 4 times

Question #33

You have a Conversational Language Understanding model.

You export the model as a JSON file. The following is a sample of the file.

```
{  
  "text": "average amount of rain by month in Chicago last year",  
  "intent": "Weather.CheckWeatherValue",  
  "entities": [  
    {  
      "entity": "Weather.WeatherRange",  
      "startPos": 0,  
      "endPos": 6,  
      "children": []  
    },  
    {  
      "entity": "Weather.WeatherCondition",  
      "startPos": 18,  
      "endPos": 21,  
      "children": []  
    },  
    {  
      "entity": "Weather.Historic",  
      "startPos": 23,  
      "endPos": 30,  
      "children": []  
    }  
  ]  
}
```

What represents the Weather.Historic entity in the sample utterance?

- A. last year
- B. by month
- C. amount of
- D. average

Correct Answer: B

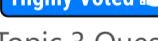
Community vote distribution

B (100%)

✉  **zellck**  10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

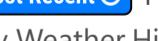
upvoted 8 times

✉  **zellck**  11 months ago

Same as Topic 3 Question 23.

<https://www.examtopics.com/discussions/microsoft/view/77619-exam-ai-102-topic-3-question-23-discussion>

upvoted 5 times

✉  **jv1**  1 month, 2 weeks ago

The entity Weather.Historic in the sample utterance is represented by the string that starts at character position 23 and ends at character position 30. In the given text "average amount of rain by month in Chicago last year", the substring "last year" corresponds to those positions. Therefore, the entity Weather.Historic represents "last year".

upvoted 1 times

✉  **rdemontis** 6 months, 2 weeks ago

Selected Answer: B

correct and repeated

upvoted 3 times

 **james2033** 9 months, 2 weeks ago

Selected Answer: B
"average amount of rain by month in Chicago last year" . The question has mistake, should start position - end position are 23 - 31.
upvoted 3 times

 **973b658** 11 months, 3 weeks ago

Selected Answer: B
B.
Same #22Topic 3
upvoted 3 times

Question #34

You are building a chatbot by using Microsoft Bot Framework Composer.

You need to configure the chatbot to present a list of available options. The solution must ensure that an image is provided for each option.

Which two features should you use? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. an entity
- B. an Azure function
- C. an utterance
- D. an adaptive card
- E. a dialog

Correct Answer: DE

Community vote distribution

DE (100%)

✉  **kail85**  11 months, 4 weeks ago

Selected Answer: DE

To configure the chatbot to present a list of available options with images, you should use the following features:

D. An adaptive card: Adaptive cards are a flexible, customizable way to present rich, interactive content within a chatbot conversation. You can use adaptive cards to display a list of available options, along with images and other interactive elements, such as buttons or inputs.

E. A dialog: Dialogs are used to manage and organize the conversation flow in your chatbot. You can create a dialog that is responsible for presenting the list of available options using an adaptive card. The dialog can be triggered when needed, and it can handle the user's selection of an option.

upvoted 7 times

✉  **evangelist**  3 months, 3 weeks ago

Selected Answer: DE

only card enables showing image in an option

upvoted 1 times

✉  **rdemontis** 6 months, 3 weeks ago

Selected Answer: DE

Correct answer

<https://learn.microsoft.com/en-us/composer/concept-dialog?tabs=v2x>

<https://learn.microsoft.com/en-us/composer/how-to-send-cards?tabs=v2x>

upvoted 2 times

✉  **james2033** 9 months, 2 weeks ago

Selected Answer: DE

Quote "Creating Adaptive Cards attachments, as seen above." at <https://learn.microsoft.com/en-us/composer/introduction?tabs=v2x#language-generation>. Dialog <https://learn.microsoft.com/en-us/composer/concept-dialog?tabs=v2x>

upvoted 2 times

✉  **zellick** 11 months ago

Selected Answer: DE

DE is the answer.

<https://learn.microsoft.com/en-us/composer/concept-dialog?tabs=v2x>

Modern conversational software has many different components. Bot Framework Composer integrates these pieces into dialogs, a single interface for constructing the building blocks of bot functionality.

Each dialog represents a portion of the bot's functionality and contains instructions for how the bot will react to the input. Dialogs can include custom business logic, calls to cloud APIs, training data for language processing systems, and importantly, the actual content used in conversation with the bot's end users. Simple bots will have just a few dialogs. Sophisticated bots might have dozens or hundreds of individual dialogs.

upvoted 4 times

✉  **zellick** 11 months ago

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-design-user-experience?view=azure-bot-service-4.0#cards>
Cards allow you to present your users with various visual, audio, and/or selectable messages and help to assist conversation flow. If a user needs to select from within a fixed set of items you can display a carousel of cards, each containing an image, a text description, and a single selection button. If a user has a set of choices for a single item, you can present a smaller single image and a collection of buttons with various options to choose between. Did they ask for more information on a subject? Cards can provide in-depth information using audio or video output, or a receipt that details their shopping experience. There's an incredibly wide range of uses for cards to help guide the conversation between your user and your bot. The type of card you use will be determined by the needs of your application.

upvoted 3 times

 **zellck** 11 months ago

<https://learn.microsoft.com/en-us/composer/how-to-send-cards?tabs=v2x>
Cards let you create bots that can communicate with users in various ways. You can think of a card as an object with a standard set of rich user controls that you can choose from to communicate with and gather input from users. There are times when you need messages that are plain text only, and there are times when you need richer message content, like images, animated GIFs, video clips, audio clips, and buttons.

upvoted 2 times

Question #35

You are building a chatbot.

You need to configure the bot to guide users through a product setup process.

Which type of dialog should you use?

- A. component
- B. action
- C. waterfall
- D. adaptive

Correct Answer: C

Community vote distribution

C (78%)

A (22%)

✉ **rdemontis** 6 months, 3 weeks ago

Selected Answer: C

To me the most suitable Dialog is Waterfall

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-dialog?view=azure-bot-service-4.0#dialog-types>
upvoted 2 times

✉ **sl_msiconsulting** 7 months, 1 week ago

Selected Answer: A

Setting up a product might not be a straight linear process in my view. You might need component dialog to coordinate among multiple waterfall dialogs which by design work within the context of a component dialog. <https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-waterfall-dialogs?view=azure-bot-service-4.0>

upvoted 2 times

✉ **james2033** 9 months, 2 weeks ago

Selected Answer: C

Quote "waterfall dialog: Defines a sequence of steps, allowing your bot to guide a user through a linear process. These are typically designed to work within the context of a component dialog." at <https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-dialog?view=azure-bot-service-4.0#dialog-types>

upvoted 1 times

✉ **abelarda** 1 week, 5 days ago

Shouldn't it be component to work in a "context of a component dialog"?

upvoted 1 times

✉ **zellck** 11 months ago

Selected Answer: C

C is the answer.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-dialog?view=azure-bot-service-4.0#dialog-types>
The dialogs library provides a few types of dialogs to make your bot's conversations easier to manage.

- waterfall dialog

Defines a sequence of steps, allowing your bot to guide a user through a linear process. These are typically designed to work within the context of a component dialog.

upvoted 3 times

✉ **973b658** 11 months, 3 weeks ago

Selected Answer: C

C

Like #22 Topic5

upvoted 1 times

Question #36

Topic 5

You have a chatbot that was built by using Microsoft Bot Framework and deployed to Azure.

You need to configure the bot to support voice interactions. The solution must support multiple client apps.

Which type of channel should you use?

- A. Cortana
- B. Microsoft Teams
- C. Direct Line Speech

Correct Answer: C

Community vote distribution

C (100%)

zellck Highly Voted 11 months ago

Selected Answer: C

C is the answer.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-channel-directline?view=azure-bot-service-4.0>

The Bot Framework offers multiple channels with the Direct Line branding. It's important that you select the version that best fits the conversation AI experience you're designing.

- Direct Line Speech. It provides text-to-speech and speech-to-text services within the channel. It allows a client to stream audio directly to the channel which will then be converted to text and sent to the bot. Direct Line Speech can also convert text messages from the bot into audio messages spoken by an AI-powered voice. Combined, this makes Direct Line Speech capable of having audio only conversations with clients.

upvoted 5 times

zellck 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 5 times

evangelist Most Recent 4 months ago

Selected Answer: C

Designed for Voice Interactions: ***Direct Line Speech*** is specifically designed to enable voice interactions for bots. It provides an integrated set of capabilities to convert speech to text (STT), perform natural language understanding, and then convert text back to speech (TTS). This is essential for creating a seamless voice interaction experience.

Cortana is tailored for interactions with Microsoft's virtual assistant and is limited to the Cortana ecosystem.

Microsoft Teams is focused on integrating bots within the Teams platform for collaboration and communication, but it's not specifically designed for voice-first experiences.

upvoted 3 times

rdemontis 6 months, 3 weeks ago

Selected Answer: C

Answer seems correct

<https://learn.microsoft.com/en-us/azure/bot-service/directline-speech-bot?view=azure-bot-service-4.0>

upvoted 2 times

sl_msiconsulting 7 months, 1 week ago

Selected Answer: C

Quoted from the link down below. Using DirectLine, not only can you define your own UI with your own features, but you can connect multiple applications to a registered Bot simultaneously. <https://blog.botframework.com/2017/06/26/speech-to-text/>

upvoted 2 times

james2033 9 months, 2 weeks ago

Selected Answer: C

<https://learn.microsoft.com/en-us/azure/bot-service/directline-speech-bot?view=azure-bot-service-4.0>

upvoted 3 times

Codia 1 year, 1 month ago

Direct Line Speech is a robust, end-to-end solution for creating a flexible, extensible voice assistant. It is powered by the Bot Framework and its Direct Line Speech channel, that is optimized for voice-in, voice-out interaction with bots.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/direct-line-speech>

upvoted 4 times

Question #37

You are building a bot by using Microsoft Bot Framework.

You need to configure the bot to respond to spoken requests. The solution must minimize development effort.

What should you do?

- A. Deploy the bot to Azure and register the bot with a Direct Line Speech channel.
- B. Integrate the bot with Cortana by using the Bot Framework SDK.
- C. Create an Azure function that will call the Speech service and connect the bot to the function.
- D. Deploy the bot to Azure and register the bot with a Microsoft Teams channel.

Correct Answer: B

Community vote distribution

A (83%) B (17%)

 **kail85**  12 months ago

- A. Deploy the bot to Azure and register the bot with a Direct Line Speech channel.

To enable the bot to respond to spoken requests with minimal development effort, you should deploy the bot to Azure and register the bot with a Direct Line Speech channel. Direct Line Speech provides an integrated speech and bot experience without requiring additional development. It combines both the speech-to-text and text-to-speech capabilities with the bot's logic to handle spoken requests and generate spoken responses.

upvoted 9 times

 **EliteAllen**  11 months ago

Selected Answer: A

- A. Deploy the bot to Azure and register the bot with a Direct Line Speech channel.

The Direct Line Speech channel in Azure Bot Service provides an integrated speech and bot experience. It combines the Bot Framework's Direct Line and Speech services into a single service that enables your bot to speak and listen to users. This is the simplest way to enable your bot to respond to spoken requests, as it doesn't require additional coding or integration with other services.

upvoted 6 times

 **evangelist**  3 months, 2 weeks ago

Selected Answer: A

checked with documentation and the answer is A

upvoted 1 times

 **evangelist** 4 months ago

Selected Answer: A

Direct Line Speech Channel: This channel is specifically designed to handle voice interactions. It provides integrated speech recognition and text-to-speech capabilities, which are essential for handling spoken requests. By using this channel, you can leverage Azure's built-in capabilities without the need for extensive custom development.

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

IMHO the correct answer is A.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-channel-directline?view=azure-bot-service-4.0>

upvoted 2 times

 **sl_mslconsulting** 7 months, 1 week ago

Selected Answer: B

Check the link here <https://blog.botframework.com/2017/06/26/speech-to-text/> and then decide the answer yourself. The direct line speech channel the development effort is not trivial.

upvoted 2 times

 **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-channel-directline?view=azure-bot-service-4.0>

The Bot Framework offers multiple channels with the Direct Line branding. It's important that you select the version that best fits the conversation AI experience you're designing.

- Direct Line Speech. It provides text-to-speech and speech-to-text services within the channel. It allows a client to stream audio directly to the channel which will then be converted to text and sent to the bot. Direct Line Speech can also convert text messages from the bot into audio messages spoken by an AI-powered voice. Combined, this makes Direct Line Speech capable of having audio only conversations with clients.
upvoted 5 times

 **ziggy1117** 11 months, 3 weeks ago

Selected Answer: B

B. Cortana.

Using Direct Line Speech needs more work because we still need to develop the speech recognition. however, cortana has some built in features already but is not going to be as flexible as Direct line speech

upvoted 2 times

 **973b658** 11 months, 3 weeks ago

Selected Answer: A

kail85 is correct.

upvoted 4 times

 **ziggy1117** 11 months, 3 weeks ago

B. Cortana.

Using Direct Line Speech needs more work because we still need to develop the speech recognition. however, cortana has some built in features already but is not going to be as flexible as Direct line speech

upvoted 2 times

 **Codia** 1 year, 1 month ago

minimize development effort -> Cortana (we can connect the Cortana channel and integrate a Bot Application to Cortana)

upvoted 2 times

Question #38

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have a chatbot that uses question answering in Azure Cognitive Service for Language.

Users report that the responses of the chatbot lack formality when answering spurious questions.

You need to ensure that the chatbot provides formal responses to spurious questions.

Solution: From Language Studio, you remove all the chit-chat question and answer pairs, and then retrain and republish the model.

Does this meet the goal?

- A. Yes
- B. No

Correct Answer: B

Community vote distribution

B (100%)

 **reiwanotora** 1 day, 22 hours ago

Selected Answer: B

No is right answer.
upvoted 1 times

 **jv1** 1 month, 2 weeks ago

The proposed solution to ensure that the chatbot provides formal responses to spurious questions is to remove all the chit-chat question and answer pairs, and then retrain and republish the model. Answer=B
upvoted 1 times

 **evangelist** 4 months ago

Selected Answer: B

"Professional personality" refers to a specific conversational style or "character" chosen when designing a chatbot. This personality is typically more formal and suitable for use in corporate and official environments. In contrast, there may be conversational styles that are friendlier, more casual, or humorous. By selecting different "personalities," you can adjust the way the chatbot engages with users to better fit its application scenarios and target audience. In Azure AI's chatbot service, various preset personalities are often provided, such as professional, friendly, humorous, etc. Developers can choose the most suitable one based on the chatbot's purpose and user expectations. This flexibility allows the chatbot to exhibit different interaction styles in various application scenarios, thereby enhancing the user experience.

upvoted 2 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: B

B. You need chitchat with professional personality
upvoted 2 times

 **Lion007** 7 months ago

Selected Answer: B

The given answer Correct.

Removing all chitchat QA pairs will not solve the problem of lack of formality.

To solve the issue, change the chitchat source to qna_chitchat_professional.tsv

Check out <https://learn.microsoft.com/en-us/azure/ai-services/language-service/question-answering/concepts/best-practices#choosing-a-personality>

upvoted 4 times

Question #39

HOTSPOT

You are building a chatbot.

You need to use the Content Moderator service to identify messages that contain sexually explicit language.

Which section in the response from the service will contain the category score, and which category will be assigned to the message? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Section:

Classification

pii

Terms

Category:

1

2

3

Answer Area

Correct Answer:

Section:

Classification

pii

Terms

Category:

1

2

3

 **zellck** Highly Voted  11 months ago

1. Classification
2. 1

<https://learn.microsoft.com/en-us/azure/cognitive-services/content-moderator/text-moderation-api#classification>

Content Moderator's machine-assisted text classification feature supports English only, and helps detect potentially undesired content. The flagged content may be assessed as inappropriate depending on context. It conveys the likelihood of each category. The feature uses a trained model to identify possible abusive, derogatory or discriminatory language. This includes slang, abbreviated words, offensive, and intentionally misspelled words.

Category1 refers to potential presence of language that may be considered sexually explicit or adult in certain situations.

upvoted 17 times

 **rdemontis** 6 months, 3 weeks ago

thanks for explanation
upvoted 2 times

 **tzuyichao** Highly Voted  1 year, 1 month ago

Correct

Ref: <https://learn.microsoft.com/en-us/azure/cognitive-services/content-moderator/text-moderation-api#classification>
upvoted 6 times

✉ **evangelist** Most Recent 3 months, 2 weeks ago

In the Content Moderator service response, the classification section can include different category scores that represent various types of content. Here's a summary of what each category generally signifies:

Category 1: Typically represents sexually explicit content. This includes images, text, or videos that contain nudity or sexual activities.

Category 2: Usually denotes sexually suggestive content. While not explicit, this content has sexual undertones or risqué material that might not be suitable for all audiences.

Category 3: Often related to offensive content. This can cover a broad range of materials, from language that might be considered profane or derogatory, to content that includes hate speech, violence, or other types of content that could be offensive or disturbing to some viewers.

upvoted 2 times

✉ **evangelist** 4 months ago

Category1 refers to potential presence of language that may be considered sexually explicit or adult in certain situations.

Category2 refers to potential presence of language that may be considered sexually suggestive or mature in certain situations.

Category3 refers to potential presence of language that may be considered offensive in certain situations.

upvoted 4 times

✉ **Gvalli** 5 months, 1 week ago

This was in the exam today.

upvoted 3 times

Question #40

You are building a chatbot for a travel agent. The bot will ask users for a destination and must repeat the question until a valid input is received, or the user closes the conversation.

Which type of dialog should you use?

- A. prompt
- B. input
- C. adaptive
- D. QnA Maker

Correct Answer: A

Community vote distribution

A (100%)

✉️  **zellick**  11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-dialog?view=azure-bot-service-4.0#dialog-types>

The dialogs library provides a few types of dialogs to make your bot's conversations easier to manage.

- prompt dialogs

Ask the user for input and return the result. A prompt will repeat until it gets valid input or it's canceled. They're designed to work with waterfall dialogs.

upvoted 6 times

✉️  **zellick** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 4 times

✉️  **PCRamirez**  3 months, 2 weeks ago

According to Windows Copilot:

I understand your reasoning, but option A is not the best choice for this scenario. A prompt dialog can validate the user's input, but it cannot handle complex conversational flows, such as interruptions, validations, and conditional logic. For example, if the user wants to change their destination, or ask for more information, or cancel the conversation, a prompt dialog cannot adapt to these situations. A prompt dialog can only ask the user for a single piece of information and return it to the parent dialog.

An adaptive dialog, on the other hand, can dynamically adjust to the user's input and context. It can handle multiple inputs, interruptions, validations, and conditional logic. It can also use language generation templates to create natural and varied responses. An adaptive dialog can provide a better user experience and a more robust chatbot for a travel agent.

I hope this clarifies why option C is the correct answer.

upvoted 2 times

✉️  **evangelist** 4 months ago

Selected Answer: A

In this scenario, where a chatbot needs to ask users for a destination and must repeat the question until a valid input is received, the most appropriate type of dialog to use is a "prompt" dialog. Here's why:

Prompt Dialog: The prompt dialog is specifically designed to obtain user input and can be configured to validate this input and reprompt the user if necessary. In the Microsoft Bot Framework, prompts are used to handle various types of user input such as text, numbers, dates, and confirmations. For your requirement, you can use a text prompt that keeps asking the user for a destination until a valid one is provided or the user exits the conversation.

upvoted 3 times

✉️  **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

correct

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-dialog?view=azure-bot-service-4.0#dialog-types>

upvoted 2 times

✉️  **Tin_Tin** 11 months, 1 week ago

answer seems correct.

<https://learn.microsoft.com/en-us/training/modules/create-bot-with-bot-framework-composer/4-dialogs>

upvoted 3 times

✉ 973b658 11 months, 3 weeks ago

Selected Answer: A

A. prompt
upvoted 3 times

Question #41

Topic 5

You are building a chatbot.

You need to configure the chatbot to query a knowledge base.

Which dialog class should you use?

- A. QnAMakerDialog
- B. AdaptiveDialog
- C. SkillDialog
- D. ComponentDialog

Correct Answer: A

Community vote distribution

A (100%)

✉ zellck Highly Voted 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-dialog?view=azure-bot-service-4.0#dialog-types>
The dialogs library provides a few types of dialogs to make your bot's conversations easier to manage.

- QnA Maker dialog

Automates access to a QnA Maker knowledge base. This dialog is designed to also work as an action within Composer.

upvoted 7 times

✉ zellck 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 5 times

✉ evangelist Most Recent 4 months ago

Selected Answer: A

Here's why:

QnAMakerDialog: This class is specifically designed for integration with QnA Maker, a cloud-based API service offered by Microsoft Azure that creates a question and answer layer over your data. QnAMakerDialog simplifies the process of querying a QnA Maker knowledge base and handling the responses within your chatbot. It manages the interaction pattern with the QnA Maker service, sending user queries and receiving answers that are then presented to the user.

upvoted 1 times

✉ rdemontis 6 months, 3 weeks ago

Selected Answer: A

Correct

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-dialog?view=azure-bot-service-4.0#dialog-types>
upvoted 1 times

✉ 973b658 11 months, 3 weeks ago

Selected Answer: A

A. QnAMakerDialog
upvoted 2 times

Question #42

HOTSPOT

You have a chatbot.

You need to ensure that the bot conversation resets if a user fails to respond for 10 minutes.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
...
if now_seconds != last_access and (now_seconds - last_access >= self.expire_after_seconds ):
    await turn_context.
        send_activity()
        send_trace_activity()
        update_activity()

    "Welcome back! Let's start over from the beginning."
)

await self.conversation_state.
    Delete_property_value()
    Save_changes()
    Set_property_value()
```

Answer Area

```
...
if now_seconds != last_access and (now_seconds - last_access >= self.expire_after_seconds ):
    await turn_context.
        send_activity()
        send_trace_activity()
        update_activity()

    "Welcome back! Let's start over from the beginning."
)

await self.conversation_state.
    Delete_property_value()
    Save_changes()
    Set_property_value()
```

Correct Answer:

  **zellck**  11 months ago

1. send_activity
2. clear_state

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-howto-expire-conversation?view=azure-bot-service-4.0&tabs=python#user-interaction-expiration>

Notify the user that the conversation is being restarted.

- await turn_context.send_activity()

Clear state.

- await self.conversation_state.clear_state(turn_context)

upvoted 13 times

  **rdemontis** 6 months, 3 weeks ago

thanks again for your contribution

upvoted 2 times

✉️  **evangelist** Most Recent 4 months ago

In summary, send_activities is the standard method for sending messages to users. on_send_activities is not typically used for sending messages directly. send_trace_activity is for debugging, and update_activity is for modifying existing messages.

upvoted 2 times

✉️  **Tin_Tin** 11 months, 1 week ago

The answer seems correct.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-howto-expire-conversation?view=azure-bot-service-4.0&tabs=python>

upvoted 2 times

✉️  **973b658** 11 months, 2 weeks ago

It is true.

upvoted 2 times

Question #43

You develop a Conversational Language Understanding model by using Language Studio.

During testing, users receive incorrect responses to requests that do NOT relate to the capabilities of the model.

You need to ensure that the model identifies spurious requests.

What should you do?

- A. Enable active learning.
- B. Add entities.
- C. Add examples to the None intent.
- D. Add examples to the custom intents.

Correct Answer: A

Community vote distribution

C (100%)

 **zellck**  11 months ago

Selected Answer: C

C is the answer.

<https://learn.microsoft.com/en-us/azure/cognitive-services/language-service/conversational-language-understanding/concepts/none-intent#adding-examples-to-the-none-intent>

The None intent is also treated like any other intent in your project. If there are utterances that you want predicted as None, consider adding similar examples to them in your training data. For example, if you would like to categorize utterances that are not important to your project as None, such as greetings, yes and no answers, responses to questions such as providing a number, then add those utterances to your intent.

You should also consider adding false positive examples to the None intent. For example, in a flight booking project it is likely that the utterance "want to buy a book" could be confused with a Book Flight intent. Adding "I want to buy a book" or "I love reading books" as None training utterances helps alter the predictions of those types of utterances towards the None intent instead of Book Flight.

upvoted 6 times

 **dd11121996** 10 months, 1 week ago

are you sure? active learning will not be the answer?

upvoted 1 times

 **evangelist**  4 months ago

Selected Answer: C

In a Conversational Language Understanding model, especially when using a tool like Language Studio, the "None" intent plays a crucial role in handling queries or requests that do not pertain to the defined capabilities of the model. Here's why adding examples to the None intent is the right approach:

None Intent: The None intent is designed to capture user inputs that fall outside of the other specified intents in your model. By training the None intent with examples of irrelevant or unrelated requests, you teach the model to recognize and appropriately handle spurious queries that do not align with its capabilities. This helps in reducing incorrect responses by ensuring that the model does not force these out-of-scope queries into one of the custom intents.

upvoted 2 times

 **Gvalli** 5 months, 3 weeks ago

Selected Answer: C

C - None intent.

Because of 'NOT' related utterances, and 'NOT' related utterances do not add value, hence 'Not active learning'.

Keyword - 'NOT' related.

During testing, users receive incorrect responses to requests that do NOT relate to the capabilities of the model.

upvoted 1 times

 **rdemontis** 6 months, 3 weeks ago

Selected Answer: C

None intent is where you place utterances not related to the model capabilities

upvoted 1 times

 **ExamDev** 8 months, 3 weeks ago

Selected Answer: C

I agree with zellck. C is the right answer.
upvoted 2 times

✉ **973b658** 11 months, 2 weeks ago

Selected Answer: C

C.
#22 topic3
upvoted 3 times

✉ **M25** 8 months, 3 weeks ago

Now #19 topic 3
upvoted 1 times

Question #44

You have a Speech resource and a bot that was built by using the Microsoft Bot Framework Composer.

You need to add support for speech-based channels to the bot.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Configure the language and voice settings for the Speech resource.
- B. Add the endpoint and key of the Speech resource to the bot.
- C. Add language understanding to dialogs.
- D. Add Orchestrator to the bot.
- E. Add Speech to the bot responses.
- F. Remove the setSpeak configuration.

Correct Answer: AEF

Community vote distribution



✉️ **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 6 times

✉️ **rdemontis** 6 months, 2 weeks ago

Selected Answer: ABE

To add support for speech-based channels to a bot built with the Microsoft Bot Framework Composer, you should perform the following actions:

Add the endpoint and key of the Speech resource to the bot (B): You need to configure the bot to connect to the Speech resource, which will enable the bot to convert text into speech and vice versa.

Add Speech to the bot responses (E): You should update your bot's responses to include speech output. This is important for speech-based channels, as the bot will need to generate spoken responses for users.

Configure the language and voice settings for the Speech resource (A): Configuring the language and voice settings in your Speech resource is important to ensure that the speech output is generated in the desired language and with the appropriate voice.

The other options (C, D, and F) are not directly related to adding support for speech-based channels and can be considered unrelated to this specific task. (ChatGPT)

upvoted 5 times

✉️ **JamesKJoker** 1 week, 1 day ago

A. Configure the language and voice settings for the Speech resource. You need to configure the Speech resource to match the language and voice settings of the bot.
B. Add the endpoint and key of the Speech resource to the bot. The bot needs to know where to connect to the Speech resource to process speech.
C. Add language understanding to dialogs. Language understanding allows the bot to understand the user's intent from spoken language. Here's why the other options are incorrect:
D. Add Orchestrator to the bot. Orchestrator is a component of Azure Logic Apps and is not directly related to speech-based channels in a bot.
E. Add Speech to the bot responses. While the bot will respond with speech, you don't need to explicitly add speech to the bot responses. It's handled by the Speech resource.
F. Remove the setSpeak configuration. The setSpeak configuration is likely related to speech functionality and should not be removed.

Gemini Pro 1.5

upvoted 1 times

✉️ **evangelist** 3 months, 2 weeks ago

A. Configure the language and voice settings for the Speech resource.

This step is crucial for ensuring that the Speech service can accurately recognize and synthesize speech in the desired language and voice.

B. Add the endpoint and key of the Speech resource to the bot.

To use the Azure Speech service with your bot, you need to authenticate your requests.

E. Add Speech to the bot responses.

To enhance the bot's interactions over speech-based channels, you should include speech-specific responses.

upvoted 2 times

✉️ **evangelist** 4 months ago

Selected Answer: ABE

A. Configure the language and voice settings for the Speech resource.

This step is important for ensuring that the Speech resource is set up to handle the specific languages and voice types that your bot will use.

B. Add the endpoint and key of the Speech resource to the bot.

This involves adding the endpoint URL and the key for your Speech resource to the bot's configuration.

E. Add Speech to the bot responses.

In the Bot Framework Composer, this typically involves specifying the speech output for each dialog or message.

upvoted 1 times

✉️ **sl_mslconsulting** 7 months, 1 week ago

Selected Answer: AEF

Check when you need to remove setSpeak <https://learn.microsoft.com/en-us/composer/concept-speech?tabs=v2x#add-speech-components-to-your-bot-responses>

upvoted 1 times

✉️ **sl_mslconsulting** 7 months, 1 week ago

Also learn how to connect a bot to a channel : <https://learn.microsoft.com/en-us/azure/bot-service/bot-service-channel-connect-directlinespeech?view=azure-bot-service-4.0>

upvoted 1 times

✉️ **ExamDev** 8 months, 3 weeks ago

Selected Answer: ABE

I also agree. Should be ABE.

upvoted 1 times

✉️ **Tin_Tin** 11 months ago

Selected Answer: ABE

should be ABE

upvoted 1 times

✉️ **EliteAllen** 11 months ago

Selected Answer: ABE

A. Configure the language and voice settings for the Speech resource.

This is necessary to ensure that the speech services can correctly interpret and generate speech in the desired language and voice.

B. Add the endpoint and key of the Speech resource to the bot.

This allows the bot to use the Speech resource for speech-to-text and text-to-speech capabilities.

E. Add Speech to the bot responses.

This is necessary to enable the bot to generate spoken responses to user input.

upvoted 3 times

✉️ **973b658** 11 months, 2 weeks ago

Selected Answer: ABE

ABE is.

upvoted 3 times

Question #45

DRAG DROP

You are building a bot.

You need to test the bot in the Bot Framework Emulator. The solution must ensure that you can debug the bot interactively.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

Run the bot app on a local host.
Use the input prompt object to send a trace activity.
Deploy the bot to Azure.
In the code for the bot, create a new trace activity.
In the code for the bot, send a trace activity.

Answer Area

In the code for the bot, create a new trace activity.
In the code for the bot, send a trace activity.
Run the bot app on a local host.

Drag the actions from the Actions list to the Answer Area. Use the arrows to rearrange the order of the actions in the Answer Area.

Correct Answer:

In the code for the bot, create a new trace activity.
In the code for the bot, send a trace activity.
Run the bot app on a local host.

zellck Highly Voted 11 months ago

1. In code, create new trace activity
2. In code, send a trace activity
3. Run bot app on local host

<https://learn.microsoft.com/en-us/azure/bot-service/using-trace-activities?view=azure-bot-service-4.0&tabs=csharp>

A trace activity is an activity that your bot can send to the Bot Framework Emulator. You can use trace activities to interactively debug a bot, as they allow you to view information about your bot while it runs locally.

Trace activities are sent only to the Emulator and not to any other client or channel. The Emulator displays them in the log but not the main chat panel.

upvoted 13 times

zellck 11 months ago

<https://learn.microsoft.com/en-us/azure/bot-service/using-trace-activities?view=azure-bot-service-4.0&tabs=csharp#to-use-a-trace-activity>

In order to see a trace activity in the Emulator, you need a scenario in which your bot will send a trace activity, such as throwing an exception and sending a trace activity from the adapter's on turn error handler.

To send a trace activity from your bot:

- Create a new activity.
- Create a new activity.

To view a trace activity in the Emulator:

- Run the bot locally on your machine.
- Test it using the Emulator.

upvoted 6 times

rdemontis 6 months, 3 weeks ago

thanks for your accurate explanation

upvoted 2 times

evangelist Most Recent 3 months, 2 weeks ago

The order should run bot app on local host first and then create trace activity and send trace activity

upvoted 2 times

evangelist 4 months ago

Connect with the Bot Framework Emulator:

Use the Bot Framework Emulator to connect to your bot running on localhost. The default URL is typically <http://localhost:3978/api/messages>.
Test Your Bot: Interact with your bot via the Emulator to see the responses and the trace activities.

upvoted 1 times

973b658 11 months, 2 weeks ago

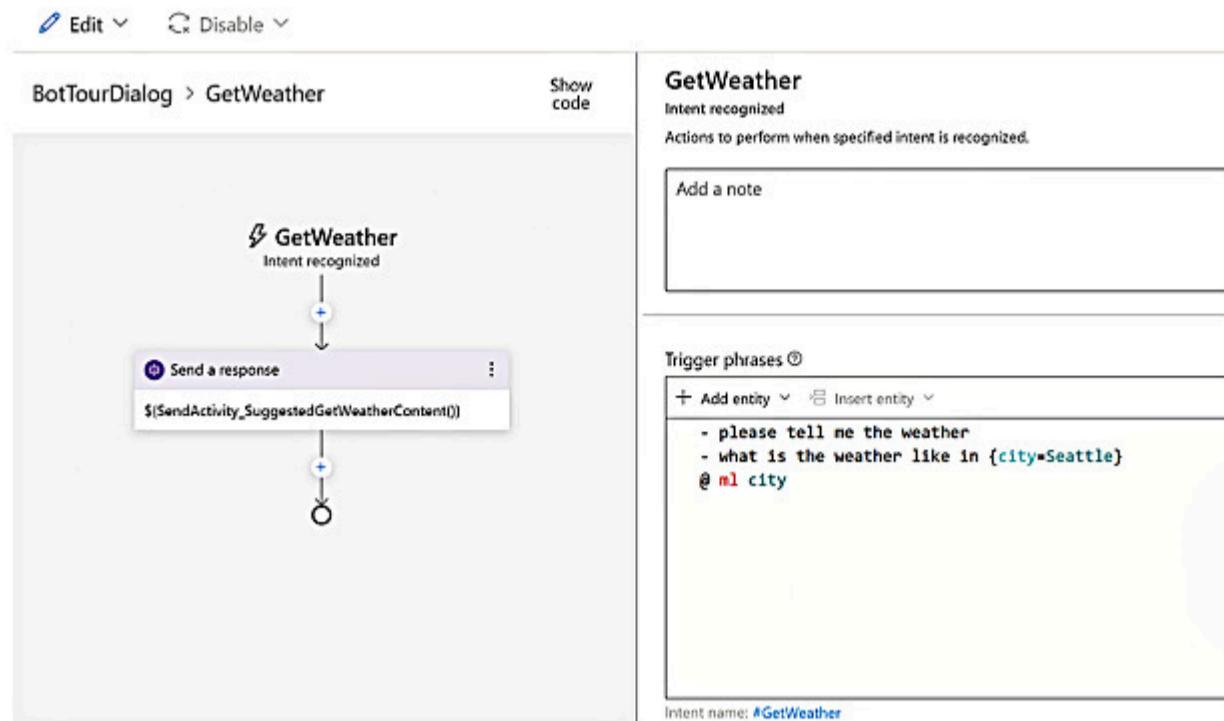
It is true.

upvoted 1 times

Question #46

HOTSPOT

You have a bot that was built by using the Microsoft Bot Framework composer as shown in the following exhibit.



Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Answer Area

If a user asks "what is the weather like in New York", the bot will [answer choice].

The GetWeather dialog uses a [answer choice] trigger.

[answer choice] dropdown options:

- change to a different dialog
- identify New York as a city entity
- identify New York as a state entity
- respond with the weather in Seattle

[answer choice] dropdown options:

- Custom events
- Dialog events
- Language Understanding Intent recognized
- QnA Intent recognized

Answer Area

If a user asks "what is the weather like in New York", the bot will [answer choice].

[answer choice] dropdown options:

- change to a different dialog
- identify New York as a city entity
- identify New York as a state entity
- respond with the weather in Seattle

[answer choice] dropdown options:

- Custom events
- Dialog events
- Language Understanding Intent recognized
- QnA Intent recognized

Correct Answer:

zellck 11 months ago

1. identify New York as city entity
2. Language Understanding intent recognised

<https://learn.microsoft.com/en-us/composer/concept-language-understanding?tabs=v2x#entities>

Entities are a collection of objects, each consisting of data extracted from an utterance such as places, times, and people. Entities and intents are both important data extracted from utterances. An utterance may include zero or more entities, while an utterance usually represents one intent. In Composer, all entities are defined and managed inline. Entities in the .lu file format are denoted using {<entityName>}=<labelled value> notation.

<https://learn.microsoft.com/en-us/composer/concept-events-and-triggers?tabs=v2x#intent-triggers>

upvoted 14 times

rdemontis 6 months, 3 weeks ago

agree with you. Thanks for the references
upvoted 3 times

✉️ **M25** 8 months, 3 weeks ago

(B) (3) Language Understanding Intent recognized

<https://learn.microsoft.com/en-us/azure/bot-service/adaptive-dialog/adaptive-dialog-prebuilt-triggers?view=azure-bot-service-4.0#recognize-event-triggers>

In this article

- Recognizer event triggers: Intent recognized | OnIntent | RecognizedIntent | Actions to perform when specified intent is recognized. [identical wording to the screenshot]
- Dialog event triggers
- Activity event triggers
- Message event triggers
- Custom event trigger

<https://learn.microsoft.com/en-us/composer/concept-events-and-triggers?tabs=v2x#intent-triggers>

There are four intent triggers in Composer:

- Intent recognized
- QnA Intent recognized
- Unknown intent
- Duplicated intents recognized

upvoted 3 times

✉️ **rdemontis** 6 months, 2 weeks ago

thanks for explanation and references

upvoted 2 times

✉️ **M25** 8 months, 3 weeks ago

(A) (2) identify New York as a city entity

<https://learn.microsoft.com/en-us/azure/bot-service/file-format/bot-builder-lu-file-format?view=azure-bot-service-4.0#machine-learned-entities>
The following example demonstrates a machine-learned entity named city (@ ml city) and a bookFlight intent with sample utterances with you entities labeled: [...].

When a user says something similar like "I need a flight booked from London to madrid", LUIS will detect the 'bookFlight` intent and extract both London and Madrid as city entities.

upvoted 3 times

✉️ **evangelist** Most Recent 3 months, 2 weeks ago

given answer is correct

upvoted 1 times

✉️ **evangelist** 4 months ago

Connect with the Bot Framework Emulator:

Use the Bot Framework Emulator to connect to your bot running on localhost. The default URL is typically <http://localhost:3978/api/messages>.

Test Your Bot: Interact with your bot via the Emulator to see the responses and the trace activities.

upvoted 1 times

✉️ **973b658** 11 months, 2 weeks ago

It is true.

upvoted 1 times

Question #47

You are building a flight booking bot by using the Microsoft Bot Framework SDK.

The bot will ask users for the departure date. The bot must repeat the question until a valid date is given, or the users cancel the transaction.

Which type of dialog should you use?

- A. prompt
- B. adaptive
- C. waterfall
- D. action

Correct Answer: A

Community vote distribution



✉️ **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-dialog?view=azure-bot-service-4.0#dialog-types>

The dialogs library provides a few types of dialogs to make your bot's conversations easier to manage.

- prompt dialogs

Ask the user for input and return the result. A prompt will repeat until it gets valid input or it's canceled. They're designed to work with waterfall dialogs.

upvoted 6 times

✉️ **evangelist** 3 months, 2 weeks ago

Selected Answer: A

The same question has appeared more than once, please choose "prompt"

upvoted 1 times

✉️ **PCRamirez** 3 months, 2 weeks ago

According to Windows Copilot:

To create a flight booking bot that repeatedly prompts users for the departure date until a valid date is provided or the transaction is canceled, you should use a waterfall dialog.

Here's why:

Waterfall Dialogs:

A waterfall dialog is a type of container dialog that guides the conversation through a series of steps or prompts.

It allows you to define a sequence of actions or questions that occur in a linear order.

You can easily manage the flow of conversation by chaining multiple steps together.

In your case, the bot can start by asking for the departure date, validate the input, and repeat the question if an invalid date is provided.

Once a valid date is obtained, the waterfall dialog can proceed to the next step or end the conversation.

upvoted 1 times

✉️ **PCRamirez** 3 months, 2 weeks ago

Other Options:

Prompt Dialogs: These are useful for asking a single question and collecting a specific type of input (e.g., asking for a date). However, they don't handle multi-turn conversations as effectively as waterfall dialogs.

Adaptive Dialogs: While adaptive dialogs provide more flexibility, they are better suited for complex scenarios involving dynamic branching and adaptive behavior. For a straightforward task like collecting a departure date, a waterfall dialog is simpler and more appropriate.

Action Dialogs: Action dialogs are not specifically designed for managing conversation flow. They are more focused on executing specific actions based on user input.

upvoted 1 times

✉️ **evangelist** 4 months ago

Selected Answer: A

for anything that needs repeated until a valid response is given, it leads to " Prompt"

upvoted 1 times

✉️ **rdemontis** 6 months, 3 weeks ago

Selected Answer: A

Correct but maybe repeated question?

upvoted 1 times

 **Tin_Tin** 11 months ago

Selected Answer: C

C. Waterfall

same question 22 of topic5

upvoted 1 times

 **973b658** 11 months, 2 weeks ago

Selected Answer: A

A.

Same Question.

upvoted 1 times

Question #48

HOTSPOT

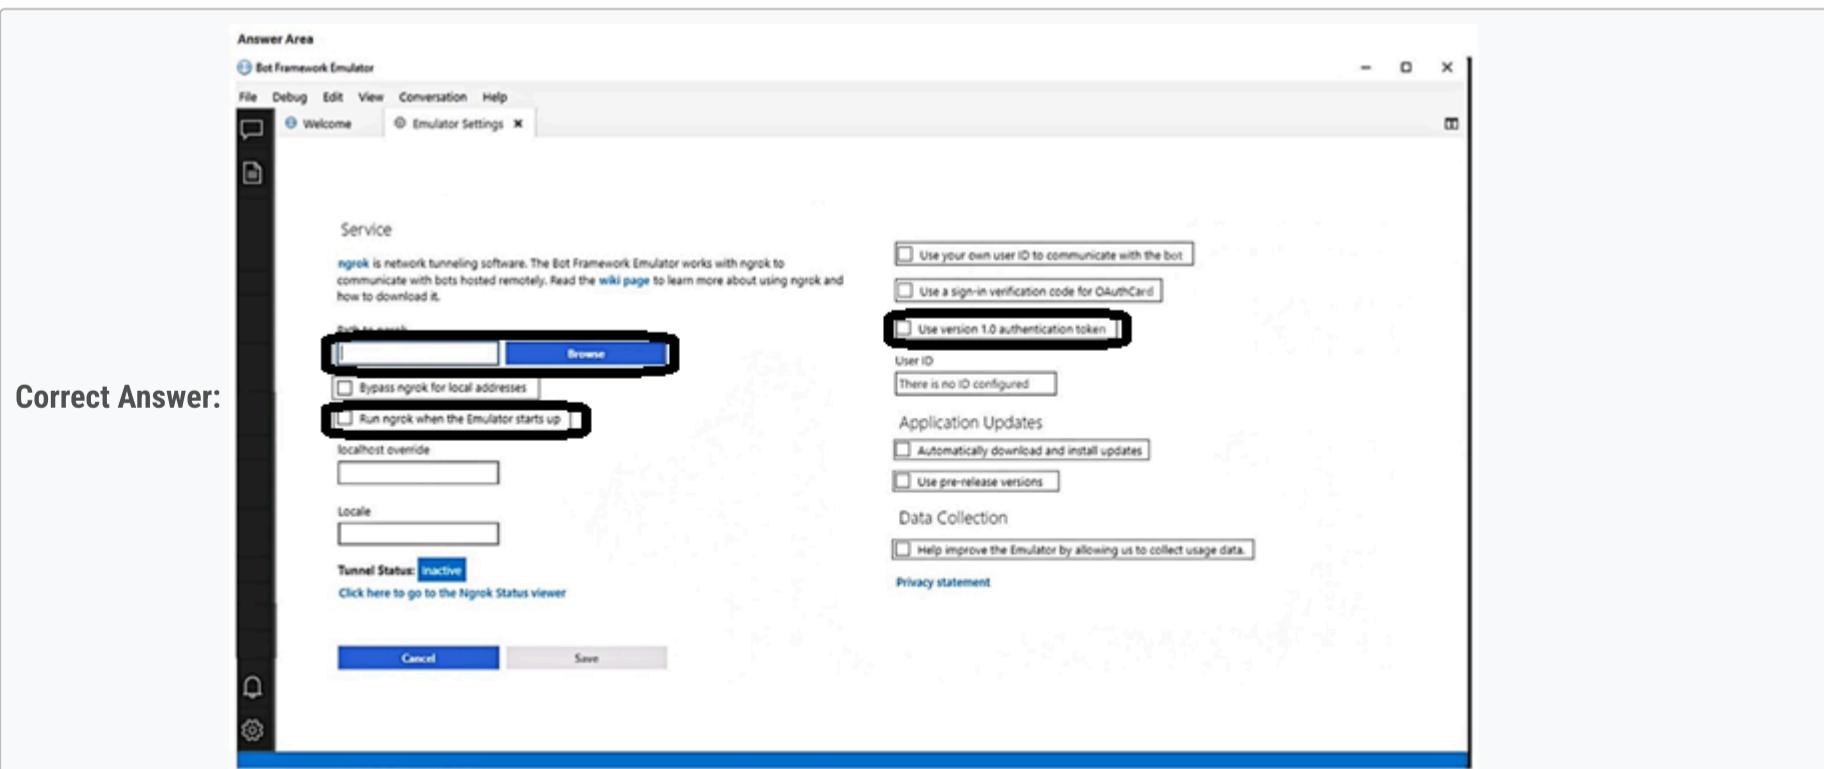
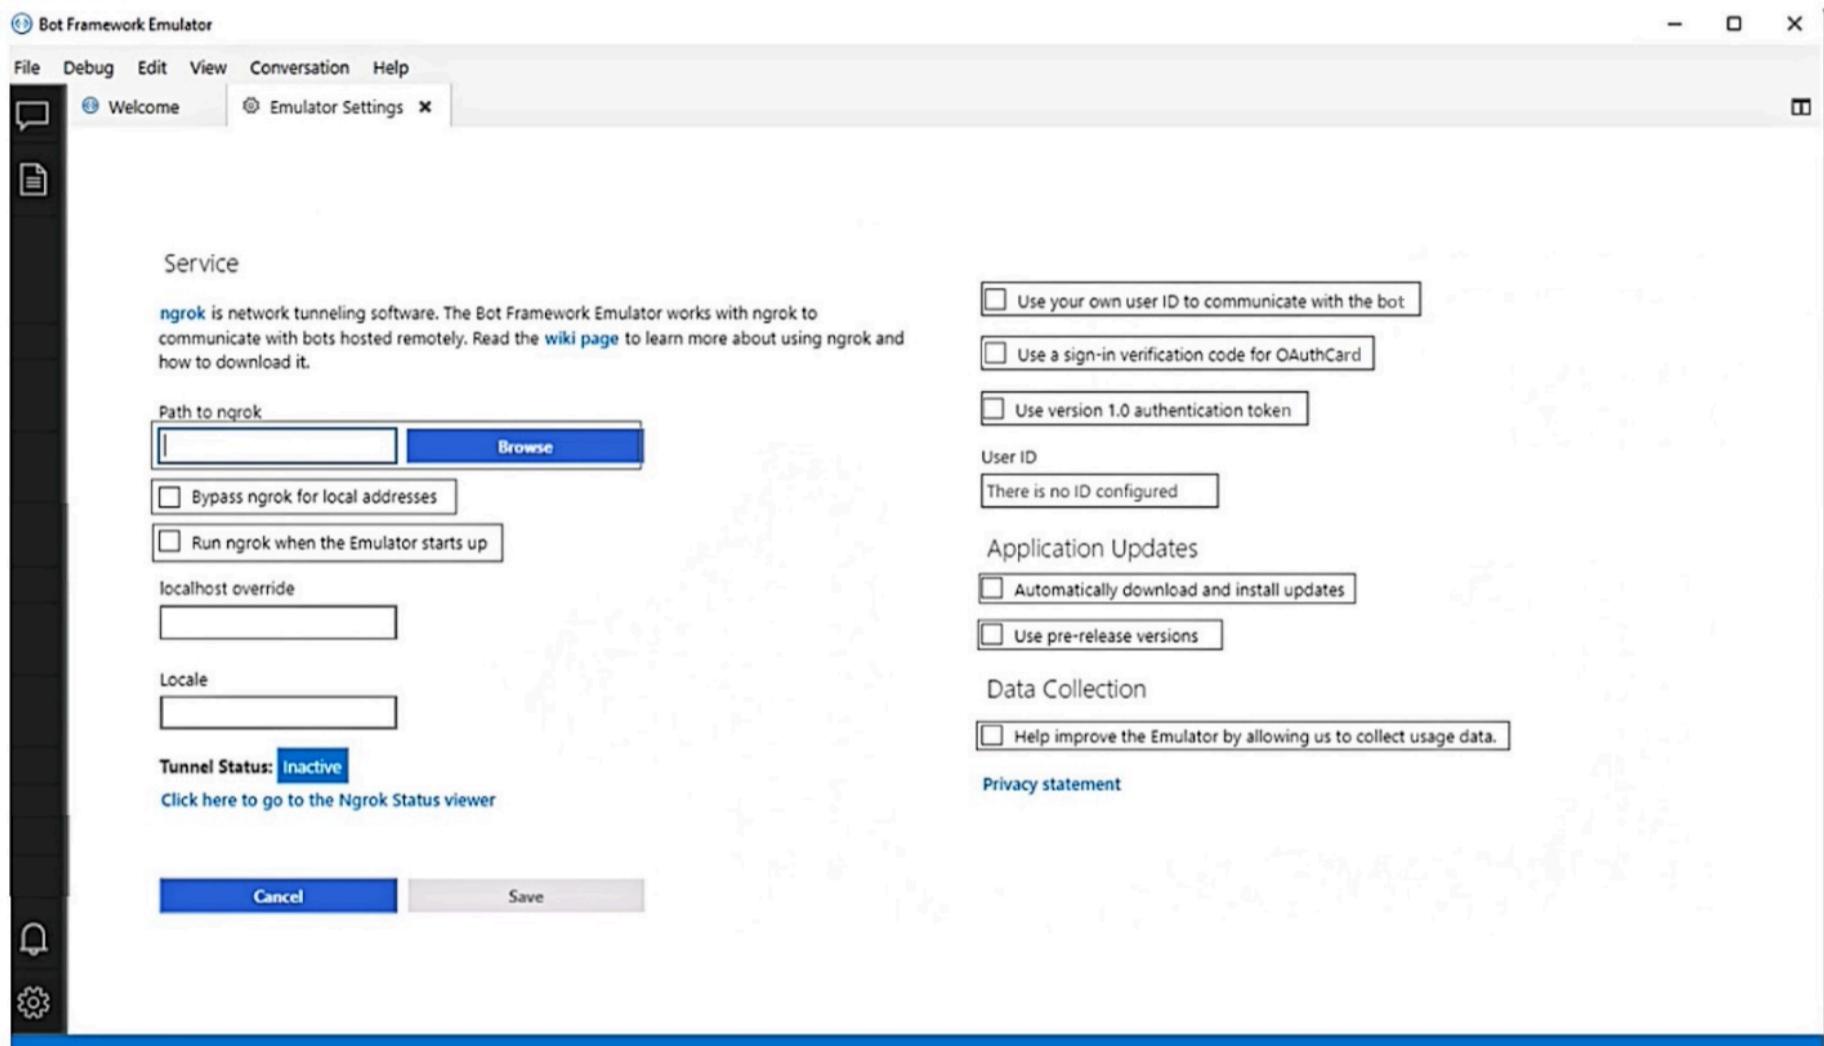
You have a chatbot.

You need to test the bot by using the Bot Framework Emulator. The solution must ensure that you are prompted for credentials when you sign in to the bot.

Which three settings should you configure? To answer, select the appropriate settings in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area



 **zellck**  11 months ago

1. Enter the local path to ngrok.
2. Enable Run ngrok when the Emulator starts up.
3. Enable Use version 1.0 authentication tokens.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-debug-emulator?view=azure-bot-service-4.0&tabs=csharp#using-authentication-tokens>

upvoted 12 times

✉️ **rdemontis** 6 months, 3 weeks ago

thanks for explanation

upvoted 4 times

✉️ **evangelist** **Most Recent** 3 months, 2 weeks ago

What is ngrok?

ngrok is a reverse proxy tool that creates a secure tunnel from a public endpoint to a locally running web service. It is widely used for testing local development versions of web applications, APIs, and bots, allowing them to be accessed over the internet. This is particularly useful for testing functionalities that require public accessibility, such as OAuth callbacks during authentication processes. By using ngrok, developers can simulate a production environment on their local machine, making it easier to test and debug features like sign-in flows that rely on external authentication services reaching the bot.

upvoted 2 times

✉️ **Tin_Tin** 11 months, 1 week ago

The answer seems correct.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-debug-emulator?view=azure-bot-service-4.0&tabs=csharp>

upvoted 1 times

✉️ **973b658** 11 months, 2 weeks ago

No.

1st is OK.

2nd is OK.

3rd is "Use a sing-in verification code for OAuthCard"

upvoted 1 times

✉️ **Tin_Tin** 11 months, 1 week ago

3rd should be enable "Use version 1.0 authentication tokens".

"Use a sing-in verification code for OAuthCard" is the setting for "Using a sign-in verification code"

upvoted 2 times

Question #49

HOTSPOT -

You are building a chatbot by using the Microsoft Bot Framework SDK.

You use an object named UserProfile to store user profile information and an object named ConversationData to store information related to a conversation.

You create the following state accessors to store both objects in state.

```
self.user_profile_accessor = self.user_state.create_property("UserProfile")
self.conversation_data_accessor = self.conversation_state.create_property("ConversationData")
```

The state storage mechanism is set to Memory Storage.

For each of the following statements, select Yes if the statement is true. Otherwise select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
The code will create and maintain the UserProfile object in the underlying storage layer.	<input type="radio"/>	<input type="radio"/>
The code will create and maintain the ConversationData object in the underlying storage layer.	<input type="radio"/>	<input type="radio"/>
The UserProfile and ConversationData objects will persist when the Bot Framework runtime terminates.	<input type="radio"/>	<input type="radio"/>

Answer Area		
Statements	Yes	No
The code will create and maintain the UserProfile object in the underlying storage layer.	<input checked="" type="checkbox"/>	<input type="radio"/>
The code will create and maintain the ConversationData object in the underlying storage layer.	<input checked="" type="checkbox"/>	<input type="radio"/>
The UserProfile and ConversationData objects will persist when the Bot Framework runtime terminates.	<input type="radio"/>	<input checked="" type="checkbox"/>

 **rdemontis** Highly Voted 6 months, 2 weeks ago

Answer is correct and probably this is a repeated question.

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-state?view=azure-bot-service-4.0>

<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-howto-v4-state?view=azure-bot-service-4.0&tabs=csharp>

upvoted 9 times

 **nanaw770** Most Recent 1 day, 19 hours ago

Yes

Yes

No

upvoted 1 times

 **evangelist** 3 months, 2 weeks ago

In summary, while the bot can indeed create and maintain UserProfile and ConversationData objects in memory, these objects will not persist beyond the lifetime of the bot's current runtime session when using Memory Storage. For persistent storage across bot restarts or runtime terminations, a different storage mechanism must be employed.

upvoted 1 times

Question #50

You build a bot.

You create an Azure Bot resource.

You need to deploy the bot to Azure.

What else should you create?

- A. only an app registration in Microsoft Azure Active Directory (Azure AD), part of Microsoft Entra, an Azure App Service instance, and an App Service plan
- B. only an app registration in Microsoft Azure Active Directory (Azure AD), part of Microsoft Entra, an Azure Kubernetes Service (AKS) instance, and a container image
- C. only an Azure App Service instance, and an App Service plan
- D. only an Azure Machine Learning workspace and an app registration in Microsoft Azure Active Directory (Azure AD), part of Microsoft Entra

Correct Answer: A

Community vote distribution

A (57%) C (43%)

 **anto69** 1 day, 21 hours ago

Selected Answer: A

No doubt is A
upvoted 1 times

 **JamesKJoker** 1 week, 1 day ago

Selected Answer: C

App registry is only required for OnPremise App
<https://learn.microsoft.com/en-us/azure/bot-service/bot-service-quickstart-registration?view=azure-bot-service-4.0&tabs=userassigned>
upvoted 1 times

 **michaelmorar** 3 weeks, 2 days ago

Selected Answer: C

Entra ID is surplus to requirement surely.
upvoted 1 times

 **Murtuza** 2 months ago

Selected Answer: A

Therefore, the correct answer is A. only an app registration in Microsoft Azure Active Directory (Azure AD), part of Microsoft Entra, an Azure App Service instance, and an App Service plan. These components work together to ensure your bot is up and running in Azure.
upvoted 1 times

 **Harry300** 2 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times

 **evangelist** 4 months ago

Selected Answer: A

To deploy a bot to Azure, at a minimum you need:

An app registration in Azure AD to represent the bot identity
An App Service instance to host the bot web service
An App Service plan which defines the pricing tier and scale for the App Service
upvoted 4 times

 **MelMac** 5 months ago

Selected Answer: A

<https://learn.microsoft.com/en-us/azure/bot-service/provision-and-publish-a-bot?view=azure-bot-service-4.0&tabs=userassigned%2Ccssharp>

Your bot will use these types of resources.

The Azure subscription that you'll use to provision, publish, and manage the bot
One or more resource groups
A user-assigned managed identity or an Microsoft Entra ID app registration
An App Service Plan resource
An App Service resource
An Azure Bot resource
upvoted 1 times

 **rdemontis** 6 months, 2 weeks ago

Selected Answer: C

I think it should be C the correct answer as you can choose USer-assigned Managed Identity to manage the identities of your bot. So in that case is not necessary to create an app reg

<https://learn.microsoft.com/en-us/azure/bot-service/provision-and-publish-a-bot?view=azure-bot-service-4.0&tabs=userassigned%2Csharp>
upvoted 2 times

 **Maccaoidh** 3 months, 3 weeks ago

I agree. However, you still have to have an identity of some sort. C does not include one, only A does.
upvoted 1 times

 **_LAW_** 6 months, 4 weeks ago

Selected Answer: C

AD it's not necessary, aks doesn't have sense here
upvoted 2 times

 **mh_63** 6 months, 4 weeks ago

I think the correct answer should be C. An Azure Service Plan and an Azure Service instance should are necessary to deploy the bot to Azure.
upvoted 2 times

Question #51

You are building a chatbot by using the Microsoft Bot Framework SDK. The bot will be used to accept food orders from customers and allow the customers to customize each food item.

You need to configure the bot to ask the user for additional input based on the type of item ordered. The solution must minimize development effort.

Which two types of dialogs should you use? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. adaptive
- B. action
- C. waterfall
- D. prompt
- E. input

Correct Answer: CD

Community vote distribution

CD (100%)

✉  **reiwanotora** 1 day, 22 hours ago

Selected Answer: CD

C and D are right answer.
upvoted 1 times

✉  **evangelist** 4 months ago

Selected Answer: CD

=====Waterfall Dialogs=====:
Guide users through a predefined, linear sequence of steps.
Each step prompts the user for specific input, and the bot moves to the next step based on the user's response. Useful for structured conversations where the bot needs to collect information in a specific order.
Ideal for gathering customization options for food items in a structured way.
=====Prompt Dialogs=====:
Simple dialogs that present a prompt to the user and wait for their response.
Can be used within waterfall dialogs or as standalone prompts.
Offer flexibility in how you ask questions and collect user input.

Can be incorporated within waterfall dialogs to ask for customization details for each item.
upvoted 3 times

✉  **rdemontis** 6 months, 2 weeks ago

Selected Answer: CD

correct
<https://learn.microsoft.com/en-us/azure/bot-service/bot-builder-concept-dialog?view=azure-bot-service-4.0#dialog-types>
upvoted 4 times

✉  **_LAW_** 6 months, 4 weeks ago

Selected Answer: CD

Correct
upvoted 4 times

Question #52

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You are building a chatbot that will use question answering in Azure Cognitive Service for Language.

You have a PDF named Doc1.pdf that contains a product catalogue and a price list.

You upload Doc1.pdf and train the model.

During testing, users report that the chatbot responds correctly to the following question: What is the price of ?

The chatbot fails to respond to the following question: How much does cost?

You need to ensure that the chatbot responds correctly to both questions.

Solution: From Language Studio, you create an entity for cost, and then retrain and republish the model.

Does this meet the goal?

A. Yes

B. No

Correct Answer: B

 **Murtuza** 2 months, 2 weeks ago

Solution: From Language Studio, you add alternative phrasing to the question and answer pair, and then retrain and republish the model.
Choice is B

upvoted 1 times

Topic 6 - Question Set 6**Topic 6****Question #1****DRAG DROP**

You have a monitoring solution that uses the Azure AI Anomaly Detector service.

You provision a server named Server1 that has intermittent internet access.

You need to deploy the Azure AI Anomaly Detector to Server1.

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

- Query the prediction endpoint on Server1.
- From Server1, run the docker push command.
- Install the Docker Engine on Server1.
- Query the prediction endpoint of the Azure AI Anomaly Detector in Azure.
- From Server1, run the docker run command.
- From Server1, run the docker pull command.

Answer Area

1	
2	
3	
4	

**Correct Answer:**

- | | |
|---|--|
| 1 | Install the Docker Engine on Server1. |
| 2 | From Server1, run the docker pull command. |
| 3 | From Server1, run the docker run command. |
| 4 | Query the prediction endpoint on Server1. |

3a0b61c 2 months ago

correct

<https://learn.microsoft.com/en-us/azure/ai-services/anomaly-detector/anomaly-detector-container-howto>

upvoted 3 times

AlviraTony 2 months, 1 week ago

Anomaly Detector is deprecated and not in syllabus for the exam

upvoted 1 times

Murtuza 2 months, 2 weeks ago

Given answers look correct to ME

upvoted 2 times

Topic 7 - Question Set 7

Question #1

Topic 7

You have an Azure subscription. The subscription contains an Azure OpenAI resource that hosts a GPT-4 model named Model1 and an app named App1. App1 uses Model1.

You need to ensure that App1 will NOT return answers that include hate speech.

What should you configure for Model1?

- A. the Frequency penalty parameter
- B. abuse monitoring
- C. a content filter
- D. the Temperature parameter

Correct Answer: C

Community vote distribution

C (100%)

✉️  **Murtuza** 1 month, 2 weeks ago

Selected Answer: C

C is correct

upvoted 2 times

✉️  **Murtuza** 2 months ago

To ensure that App1 does not return answers containing hate speech, you should configure Model1 with a content filter. The content filtering system in Azure OpenAI Service works alongside core models and aims to detect and prevent harmful content, including hate speech.

upvoted 1 times

✉️  **Murtuza** 2 months, 2 weeks ago

The content filtering models for the hate

upvoted 1 times

✉️  **Harry300** 2 months, 4 weeks ago

Selected Answer: C

Correct. <https://learn.microsoft.com/en-us/azure/ai-services/openai/concepts/content-filter?tabs=warning%2Cpython>

upvoted 4 times

Question #2

You have an Azure subscription. The subscription contains an Azure OpenAI resource that hosts a GPT-3.5 Turbo model named Model1.

You configure Model1 to use the following system message: "You are an AI assistant that helps people solve mathematical puzzles. Explain your answers as if the request is by a 4-year-old."

Which type of prompt engineering technique is this an example of?

- A. few-shot learning
- B. affordance
- C. chain of thought
- D. priming

Correct Answer: A

Community vote distribution

D (73%)

C (27%)

✉️ **River06** Highly Voted 2 months, 3 weeks ago

Selected Answer: D

Priming is utilised in this example because it involves setting up the context or role of the AI model explicitly in a system message. It instructs the model about its role ("You are an AI assistant that helps people solve mathematical puzzles.") and also provides directions about how it should respond ("Explain your answers as if the requestor is a 4-year-old."). This essentially 'primes' the model for the conversation, as it lets the model know the expected behavior and persona that it needs to take on throughout the dialogue.

upvoted 6 times

✉️ **9ae4fb5** 2 months, 2 weeks ago

ChatGPT gave the same answer.

upvoted 1 times

✉️ **bugimachi** Most Recent 4 weeks ago

Selected Answer: C

I'd go with "chain of thoughts", since it's about explaining the answers.

Priming is a different story: "This refers to including a few words or phrases at the end of the prompt to obtain a model response that follows the desired form. For example, using a cue such as "Here's a bulleted list of key points:\n- " can help make sure the output is formatted as a list of bullet points."

...that's not the case here!

<https://learn.microsoft.com/en-us/azure/ai-services/openai/concepts/advanced-prompt-engineering?pivots=programming-language-chat-completions#prime-the-output>

upvoted 2 times

✉️ **shorymor** 1 month, 1 week ago

Selected Answer: D

Priming is the answer

<https://learn.microsoft.com/en-us/azure/ai-services/openai/concepts/advanced-prompt-engineering?pivots=programming-language-chat-completions#system-message>

upvoted 1 times

✉️ **chandiochan** 2 months, 1 week ago

Selected Answer: D

Must be priming, here we setting the role for the AI model

upvoted 1 times

✉️ **warrior1234** 2 months, 3 weeks ago

D. Priming

Priming involves providing context or instructions to the model before it generates a response. In this case, the system message is priming the GPT-3.5 Turbo model by setting the expectation that it should provide explanations in a way that is understandable to a 4-year-old. This technique helps guide the model's behavior and output based on the given context or instruction.

upvoted 2 times

✉️ **Harry300** 2 months, 3 weeks ago

Selected Answer: C

chain of thought. Try it on OpenAI. It explains step by step for a formular like $7+5*3+8$

upvoted 1 times

Question #3

HOTSPOT

You build a chatbot by using Azure OpenAI Studio.

You need to ensure that the responses are more deterministic and less creative.

Which two parameters should you configure? To answer, select the appropriate parameters in the answer area.

NOTE: Each correct answer is worth one point.

Answer Area

Chat session

Clear chat Show raw JSON



Start chatting

Test your assistant by sending queries below. Then adjust your assistant setup to improve the assistant's responses.

Type user query here. (Shift + Enter for new line)

▶ ⏪

Configuration X

Parameters [Deployment](#)

Max response ⓘ	800
Temperature ⓘ	0.7
Top P ⓘ	0.9
Stop sequence ⓘ	Stop sequences
Frequency penalty ⓘ	0
Presence penalty ⓘ	0

[Learn more ↗](#)

Current token count ⓘ

Input tokens progress indicator
1/4000

Answer Area

Chat session

Clear chat View code Show raw JSON

Start chatting
Test your assistant by sending queries below. Then adjust your assistant setup to improve the assistant's responses.

Type user query here. (Shift + Enter for new line)

▶ ⏪

Configuration

Deployment Parameters

Max response 800

Temperature 0.7

Top P 0.9

Stop sequence Stop sequences

Frequency penalty 0

Presence penalty 0

Learn more

Current token count

Input tokens progress indicator 1/4000

Correct Answer:

✉️ **9H3zmT6** 3 weeks, 3 days ago

To make the responses more deterministic and less creative, it is recommended to lower the values of Temperature and TopP.

upvoted 1 times

✉️ **Murtuza** 2 months, 2 weeks ago

Given answers are correct

upvoted 3 times

Question #4

You are building a chatbot for a travel agent. The chatbot will use the Azure OpenAI GPT 3.5 model and will be used to make travel reservations.

You need to maximize the accuracy of the responses from the chatbot.

What should you do?

- A. Configure the model to include data from the travel agent's database.
- B. Set the Top P parameter for the model to 0.
- C. Set the Temperature parameter for the model to 0.
- D. Modify the system message used by the model to specify that the answers must be accurate.

Correct Answer: A

Community vote distribution



✉️ **GHill1982** 2 months, 2 weeks ago

Selected Answer: A

Could this not be A? Configuring the model to include data from the travel agent's database would allow the chatbot to access a source of relevant and specific information, which can significantly improve the accuracy of its responses. Fine-tuning the model with data specific to the travel agent's services and customer interactions can lead to more precise and contextually appropriate answers.

upvoted 6 times

✉️ **vovapOvovap** 3 days, 14 hours ago

It worse nothing, that even MS labs is used travel agency staff grounding information for a models.

upvoted 1 times

✉️ **chandiochan** 1 month ago

Selected Answer: D

Must be D

upvoted 1 times

✉️ **shorymor** 1 month, 1 week ago

Selected Answer: A

Answer is A.

Documentation is clear about providing with data to the model in order to avoid errors and therefore receive accurate answers.

<https://learn.microsoft.com/en-us/azure/ai-services/openai/concepts/advanced-prompt-engineering?pivot=programming-language-chat-completions#provide-grounding-context>

upvoted 3 times

✉️ **michaelmorar** 1 month, 1 week ago

Selected Answer: A

A system message is designed to give the model instructions, perspective from which to answer, and other helpful information.

It might include, Personality, Tone, Topics to exclude, and specifics such as formatting of answer.

It does NOTHING for accuracy - that is purely incumbent on the training data.

upvoted 1 times

✉️ **f2c587e** 1 month, 4 weeks ago

Selected Answer: U

Al especificarle al bot una fuente de datos que sirva como base de conocimiento sus respuestas seran mas precisas en cuanto a la base de conocimiento que se suministre.

upvoted 1 times

✉️ **Murtuza** 2 months ago

To maximize the accuracy of responses from the chatbot, you should modify the system message used by the model to explicitly specify that the answers must be accurate. By providing clear instructions to the model about the desired accuracy, you can guide it to generate more precise and reliable responses. This approach ensures that the chatbot focuses on accuracy while interacting with users. Therefore, the correct answer is D:
Modify the system message used by the model to specify that the answers must be accurate

upvoted 1 times

✉️ **michaelmorar** 1 month, 1 week ago

This is wrong.

upvoted 1 times

✉ **Mehe323** 2 months, 2 weeks ago

Selected Answer: D

I think it is D too.

<https://learn.microsoft.com/en-us/training/modules/apply-prompt-engineering-azure-openai/4-provide-context-to-improve-accuracy?pivot=python>

upvoted 3 times

✉ **Murtuza** 2 months, 2 weeks ago

While configuring the model with data from the travel agent's database (option A) could be beneficial,

upvoted 3 times

✉ **7afddb1** 2 months, 2 weeks ago

Selected Answer: D

D. Modify the system message used by the model to specify that the answers must be accurate.

This option involves priming the model with a specific instruction or context that emphasizes the importance of accuracy in the responses. By doing so, you set a clear expectation for the model to prioritize accuracy when generating responses. This method can be particularly effective when dealing with domain-specific queries where accurate information is crucial, such as travel dates, destinations, and reservation details.

upvoted 2 times

✉ **chandiochan** 2 months, 3 weeks ago

Must be D

upvoted 4 times

Question #5

You build a chatbot that uses the Azure OpenAI GPT 3.5 model.

You need to improve the quality of the responses from the chatbot. The solution must minimize development effort.

What are two ways to achieve the goal? Each correct answer presents a complete solution.

NOTE: Each correct answer is worth one point.

- A. Fine-tune the model.
- B. Provide grounding content.
- C. Add sample request/response pairs.
- D. Retrain the language model by using your own data.
- E. Train a custom large language model (LLM).

Correct Answer: AD

Community vote distribution

BC (100%)

 **chandiochan**  2 months, 3 weeks ago

Here are two ways to improve the quality of the responses from the chatbot with minimal development effort:

1. Provide grounding content: This involves feeding the chatbot with relevant domain-specific information and data. This can include documents, articles, FAQs, or any other content related to the chatbot's purpose. By providing this context, the chatbot can better understand the user's intent and respond in a more relevant and informative way.
2. Add sample request/response pairs: This technique involves providing the chatbot with a set of pre-defined questions and their corresponding answers. This helps the chatbot learn the conversation patterns and phrasing related to its specific domain. By analyzing these examples, the chatbot can improve its ability to generate natural and consistent responses to user queries.

Both options (A. Provide grounding content and C. Add sample request/response pairs) achieve the goal of improving response quality with minimal development effort, as they do not require extensive retraining or model building.

Therefore, the two correct answers are:

- B. Provide grounding content.
 - C. Add sample request/response pairs.
- upvoted 6 times

 **Murtuza**  2 months ago

Selected Answer: BC

Remember that while fine-tuning (A) and custom models (E) can yield high-quality results, they often require significant development effort and computational resources. In contrast, grounding content and sample pairs offer pragmatic improvements with minimal overhead

upvoted 4 times

 **chandiochan** 2 months, 1 week ago

Selected Answer: BC

Must be B & C

upvoted 3 times

 **Delta64** 2 months, 3 weeks ago

GPT 4 Answered:

- B. Provide grounding content.
 - C. Add sample request/response pairs.
- upvoted 3 times

Question #6**HOTSPOT**

You have an Azure subscription that contains an Azure OpenAI resource named AI1.

You build a chatbot that will use AI1 to provide generative answers to specific questions.

You need to ensure that the responses are more creative and less deterministic.

How should you complete the code? To answer, select the appropriate options in the answer area.

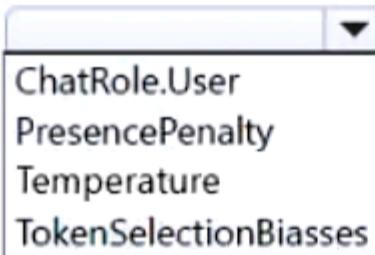
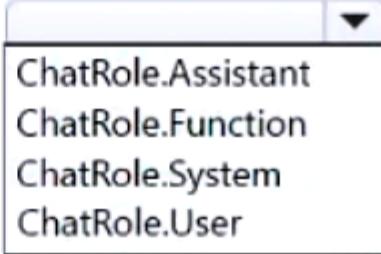
NOTE: Each correct selection is worth one point.

Answer Area

```
new ChatCompletionsOptions()
{
    Messages =
    {
        new ChatMessage(
            , @""),

        },
        = (float)1.0,
        ChatRole.User
        PresencePenalty
        Temperature
        TokenSelectionBiases
    }

    MaxTokens = 800,
});
```



Answer Area

Correct Answer:

```
new ChatCompletionsOptions()
{
    Messages =
    {
        new ChatMessage(
            },
            , @"""),
            ChatRole.Assistant
            ChatRole.Function
            ChatRole.System
            ChatRole.User
        = (float)1.0,
        ChatRole.User
        PresencePenalty
        Temperature
        TokenSelectionBiases
    MaxTokens = 800,
});
```

✉  **vovapOvovap** 2 days, 18 hours ago

Well, I think that should be ChatRole.System
ChatRole.User naturally make no sense - that user request
upvoted 1 times

✉  **Dhibi** 1 week, 5 days ago

Selecting ChatRole.User as one of the options would imply that the responses should be generated based on the user's input or perspective. However, in this scenario, the goal is to ensure that the responses from the chatbot are more creative and less deterministic. Including ChatRole.User might limit the creativity of the responses because the model would primarily consider the user's input rather than generating novel content.

By primarily focusing on ChatRole.Assistant, the responses will be predominantly generated from the perspective of the AI assistant, allowing for more creative and varied outputs. This approach ensures that the chatbot's responses are not overly influenced by the user's input, leading to more diverse and imaginative answers.

upvoted 1 times

✉  **shorymor** 1 month, 1 week ago

ChatRole.User (Assuming the input is an actual message from user)
Temperature

Documentation/Examples: <https://learn.microsoft.com/en-us/azure/ai-services/openai/how-to/reproducible-output?tabs=python>
upvoted 1 times

✉  **f2c587e** 1 month, 4 weeks ago

D.C. Con rol de user el chat es menos formal. con temperatura se ajusta la creatividad en las respuestas.
upvoted 1 times

✉  **chandiochan** 2 months ago

Yes, if the content you're inserting is the actual user message, then you should use ChatRole.User. This role signifies that the message is coming from the user, as opposed to the assistant, which would be generating the response.
upvoted 1 times

✉  **Ghill1982** 2 months, 2 weeks ago

The ChatRole that should be used is ChatRole.User. This role is assigned to the messages that come from the user, which the chatbot is responding to. The Temperature setting can be adjusted to increase creativity in the responses.
upvoted 4 times

Question #7

DRAG DROP

You have an Azure subscription that contains an Azure OpenAI resource named AI1.

You plan to build an app named App1 that will write press releases by using AI1.

You need to deploy an Azure OpenAI model for App1. The solution must minimize development effort.

Which three actions should you perform in sequence in Azure OpenAI Studio? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

- Create a deployment that uses the text-embedding-ada-002 model.
- Apply the Marketing Writing Assistant system message template.
- Apply the Default system message template.
- Create a deployment that uses the GPT-35 Turbo model.
- Deploy the solution to a new web app.

Answer Area

1	
2	
3	



Correct Answer:

Answer Area	
1	Create a deployment that uses the GPT-35 Turbo model.
2	Apply the Marketing Writing Assistant system message template.
3	Deploy the solution to a new web app.

✉ **Murtuza** 2 months, 2 weeks ago

In the Setup panel, under Use a system message template, select the Marketing Writing Assistant template and confirm that you want to update the system message.

<https://microsoftlearning.github.io/mslearn-openai/Instructions/Exercises/01-get-started-azure-openai.html>

upvoted 4 times

✉ **witkor** 2 months, 3 weeks ago

To deploy an Azure OpenAI model for the app "App1" with minimized development effort, follow these steps in Azure OpenAI Studio:

Create a deployment that uses the GPT-35 Turbo model.

Drag "Create a deployment that uses the GPT-35 Turbo model" to the answer area.

Apply the Default system message template.

Drag "Apply the Default system message template" to the answer area.

Deploy the solution to a new web app.

Drag "Deploy the solution to a new web app" to the answer area.

upvoted 3 times

✉ **Mehe323** 2 months, 2 weeks ago

Considering the fact that the question is about press releases, I would choose the marketing writing assistant template, like in the answer.

upvoted 7 times

✉ **TJ001** 4 days, 1 hour ago

<https://microsoftlearning.github.io/mslearn-openai/Instructions/Exercises/01-get-started-azure-openai.html>

upvoted 1 times

Question #8

HOTSPOT

You have an Azure subscription that contains an Azure OpenAI resource named AI1.

You build a chatbot that will use AI1 to provide generative answers to specific questions.

You need to ensure that the responses are more creative and less deterministic.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
response = openai.ChatCompletion.create(  
    engine="dgw-aoai-gpt35",  
    messages = [{"role": "assistant", "content":""}],  
    temperature=1,  
    Frequency_penalty  
    Presence_penalty  
    temperature  
    token_selection_biasses  
    max_tokens=800,  
    stop=None)
```

Answer Area

```
response = openai.ChatCompletion.create(  
    engine="dgw-aoai-gpt35",  
    messages = [{"role": "assistant", "content":""}],  
    temperature=1,  
    Frequency_penalty  
    Presence_penalty  
    temperature  
    token_selection_biasses  
    max_tokens=800,  
    stop=None)
```

Correct Answer:

 **chandiochan** Highly Voted 2 months ago

I think this must be user instead of system role. ChatRole.User identifies the text as coming from the user in the conversation. This is important because the AI model will use this information to understand the context of the prompt and tailor its response accordingly.

upvoted 5 times

✉ **chandiochan** Most Recent 1 month ago

In the role dropdown, you would select "user" because you are simulating a user's input to which the chatbot should respond. The API call will use this context to generate a response as if it were the assistant. You only need to specify the role for the input messages you're including in your API call, not the responses you're expecting to generate.

upvoted 2 times

✉ **TJ001** 4 days, 1 hour ago

it is the user message for which response is generated by this API function. so i will go with user as well.. the deterministic and creativity is controller by temperature value low or high

upvoted 1 times

✉ **TJ001** 4 days, 1 hour ago

the 'content' is blank here adds to the confusion :(.. so how can it be blank being a user message.. so could be system as well just to prime with the temperature settings. please suggest if anyone has clarity

upvoted 1 times

✉ **shorymor** 1 month, 1 week ago

Microsoft Copilot seems to be clear on this one. Based on that, answer is correct

It's response about this question (I did not ask about temperature because that one is obvious)

To ensure that your chatbot's responses are more creative and less deterministic, you should use the Chat.Role.Assistant. This role represents the AI language model (such as GPT-3.5 Turbo) and allows for imaginative and varied answers. By assigning messages to the assistant role, you encourage the model to generate creative and less predictable content.

upvoted 1 times

✉ **Murtuza** 1 month, 3 weeks ago

This one is SYSTEM here and not user. Looking at the code

Typically, a conversation is formatted with a system message first, followed by alternating user and assistant messages.

<https://platform.openai.com/docs/guides/text-generation/chat-completions-api>

upvoted 2 times

✉ **chandiochan** 2 months, 3 weeks ago

In this context, when configuring the chatbot's behavior, you would use the "system" role to provide initial settings or instructions that affect the overall behavior of the chatbot. This is not an actual message that's part of the conversation with the user but rather a directive to the AI on how to conduct itself in the conversation. For example, you might use this to set the chatbot's personality, instruct it to prioritize certain types of information, or follow specific conversational guidelines.

The "assistant" role, on the other hand, is used for messages that simulate responses from the AI as part of the conversation with the user. It represents the chatbot's side of the dialogue.

Since you want to ensure that the responses are more creative and less deterministic, and this is a setting affecting the AI's behavior, the "system" role is the correct choice. You might include instructions in the "system" message that tell the AI to be more creative or to use less strict adherence to certain conversational rules.

upvoted 2 times

✉ **Harry300** 2 months, 3 weeks ago

Should be system / temp

Without a system prompt, the responses are more creative

upvoted 2 times

Question #9

HOTSPOT

You have an Azure subscription that contains an Azure OpenAI resource.

You configure a model that has the following settings:

- Temperature: 1
- Top probabilities: 0.5
- Max response tokens: 100

You ask the model a question and receive the following response.

```
{  
  "choices": [  
    {  
      "finish_reason": "stop",  
      "index": 0,  
      "message": {  
        "content": "The founders of Microsoft are Bill Gates and Paul Allen. They co-founded the company in 1975.",  
        "role": "assistant"  
      }  
    }  
  ],  
  "created": 1679014554,  
  "id": "chatcmpl-6usfn2yyjkbmESe3G4jaQR6bDSc01",  
  "model": "gpt-3.5-turbo-0301",  
  "object": "chat.completion",  
  "usage": {  
    "completion_tokens": 86,  
    "prompt_tokens": 37,  
    "total_tokens": 123  
  }  
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area**Statements****Yes****No**

The subscription will be charged 86 tokens for the execution of the session.

The text completion was truncated because the Max response tokens value was exceeded.

The prompt_tokens value will be included in the calculation of the Max response tokens value.

Answer Area**Statements****Yes****No**

Correct Answer: The subscription will be charged 86 tokens for the execution of the session.

The text completion was truncated because the Max response tokens value was exceeded.

The prompt_tokens value will be included in the calculation of the Max response tokens value.

 michaelmorar 3 weeks, 2 days ago

N - the subscription is NOT charged for 86 tokens - the response does not contain 86 tokens. For reference, each token is roughly four characters for typical English text.

N - text completion is clearly under 86 and the sentence is not truncated. the finish_reason here is "stop". If the prompt had been cut off, the finish_reason would have been 'length'

N - max tokens is the maximum number to generate in the COMPLETION.

The token count of your prompt plus max_tokens can't exceed the model's context length. Most models have a context length of 2048 tokens (except for the newest models, which support 4096).

upvoted 1 times

✉ tk1828 1 month, 1 week ago

N/N/N

Subscriptions are charged for both the prompt and completion tokens.

Completion tokens is less than max response tokens.

It refers to max response tokens only, not max tokens.

<https://learn.microsoft.com/en-us/azure/ai-services/openai/how-to/manage-costs#understand-the-azure-openai-full-billing-model>

upvoted 4 times

✉ Murtuza 1 month, 3 weeks ago

The subscription will be charged 86 tokens for the execution of the session. Yes, that's correct. The completion_tokens value represents the number of tokens in the model's response, and this is what you're billed for.

The text completion was truncated because the Max response tokens value was exceeded. No, that's not correct. The response in this case wasn't truncated. The max_tokens parameter sets a limit on the length of the generated response. If the model's response had exceeded this limit, it would have been cut off, but in this case, the response is only 86 tokens long, which is less than the max_tokens value of 100.

The prompt_tokens value will be included in the calculation of the max_tokens value. Yes, that's correct. The max_tokens parameter includes both the prompt tokens and the completion tokens. So if your prompt is very long, it could limit the length of the model's response.

upvoted 1 times

✉ Murtuza 2 months ago

The session execution consumed 86 tokens is NO it should be total of 123 tokens which includes the prompt tokens

The text completion was truncated due to exceeding the Max response tokens value is Yes

The prompt_tokens value is included in the calculation of the Max response tokens value is YES

upvoted 2 times

✉ GHill1982 2 months, 2 weeks ago

I think it should be N/N/N.

upvoted 2 times

✉ GHill1982 1 month, 1 week ago

Changing my mind to Y/N/N

The subscription will be charged 86 tokens for the execution of the session.

Yes, the subscription will be charged for the completion_tokens used during the execution, which in this case is 86 tokens.

The text completion was truncated because the Max response tokens value was exceeded.

No, the text completion was not truncated due to exceeding the Max response tokens value. The finish_reason is listed as "stop," which indicates that the model stopped generating additional content because it reached a natural stopping point in the text, not because it hit the token limit.

The prompt_tokens value will be included in the calculation of the Max response tokens value.

No, the prompt_tokens value is not included in the calculation of the Max response tokens value. The Max response tokens setting only limits the length of the new content generated by the model in response to the prompt.

upvoted 1 times

✉ sergbs 1 month ago

You are wrong. First No. Azure OpenAI base series and Codex series models are charged per 1,000 tokens. <https://learn.microsoft.com/en-us/azure/ai-services/openai/how-to/manage-costs#understand-the-azure-openai-full-billing-model>

upvoted 2 times

Question #10

HOTSPOT

You have an Azure subscription that contains an Azure OpenAI resource named AI1.

You plan to develop a console app that will answer user questions.

You need to call AI1 and output the results to the console.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
openai.api_key = key
openai.api_base = endpoint
response =
    openai.ChatCompletion.create
    openai.Embedding.create
    openai.Image.create
    engine=deployment_name,
    prompt="What is Microsoft Azure?"
)
print
    (response.choices[0].text)
    (response.id)
    (response.text)
```

Answer Area

```
openai.api_key = key
openai.api_base = endpoint
response =
    openai.ChatCompletion.create
    openai.Embedding.create
    openai.Image.create
    engine=deployment_name,
    prompt="What is Microsoft Azure?"
)
print
    (response.choices[0].text)
    (response.id)
    (response.text)
```

Correct Answer:

carcasgon 1 month, 3 weeks ago

Correct. Under Completions > Example response
<https://learn.microsoft.com/en-us/azure/ai-services/openai/reference#completions>
upvoted 1 times

Murtuza 2 months ago

```
# Output the response to the console
print(response.choices[0].message.content)

response = client.chat.completions.create()
The given answers are CORRECT
upvoted 1 times
```

GHill1982 2 months, 2 weeks ago

Answer is correct.
upvoted 1 times

Question #11

Topic 7

HOTSPOT

You have an Azure subscription that contains an Azure OpenAI resource named AI1.

You plan to develop a console app that will answer user questions.

You need to call AI1 and output the results to the console.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
OpenAIclient client =  
    new OpenAIclient(new Uri(endpoint), new AzureKeyCredential(key));  
  
Response<Completions> response =  
  
    client. (deploymentName, "What is Microsoft Azure?");  
        GetCompletions  
        GetEmbeddings  
        GetImageGenerations  
  
Console.WriteLine  
    (response.Value.Choices[0].Text);  
    (response.Value.Id);  
    (response.Value.PromptFilterResults);
```

Answer Area

```
OpenAIclient client =  
    new OpenAIclient(new Uri(endpoint), new AzureKeyCredential(key));  
  
Response<Completions> response =  
    client. (deploymentName, "What is Microsoft Azure?");  
        GetCompletions  
        GetEmbeddings  
        GetImageGenerations  
  
Console.WriteLine  
    (response.Value.Choices[0].Text);  
    (response.value.Id);  
    (response.Value.PromptFilterResults);
```

Correct Answer:

 **Murtuza** 2 months ago

Agree answer is correct
upvoted 3 times

 **Ghill1982** 2 months, 2 weeks ago

Answer is correct.
upvoted 2 times

Question #12

HOTSPOT

You have an Azure subscription.

You need to create a new resource that will generate fictional stories in response to user prompts. The solution must ensure that the resource uses a customer-managed key to protect data.

How should you complete the script? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
az cognitiveservices account create -n myresource -g myResourceGroup --kind  AI  
           LanguageAuthoring  
           OpenAI  
  
          '--api-properties  
          '--assign-identity  
          '--encryption  
  
          "keySource": "Microsoft.KeyVault",  
          "keyVaultProperties": {  
              "keyName": "KeyName",  
              "keyVersion": "secretVersion",  
              "keyVaultUri": "https://issue23056kv.vault.azure.net/"  
          }  
      }'
```

Correct Answer:

```
Answer Area  
az cognitiveservices account create -n myresource -g myResourceGroup --kind  AI  
           LanguageAuthoring  
           OpenAI  
  
          '--api-properties  
          '--assign-identity  
           '--encryption  
  
          "keySource": "Microsoft.KeyVault",  
          "keyVaultProperties": {  
              "keyName": "KeyName",  
              "keyVersion": "secretVersion",  
              "keyVaultUri": "https://issue23056kv.vault.azure.net/"  
          }  
      }'
```

 **ProfessorZ** 1 month, 2 weeks ago

--encryption is correct:

<https://learn.microsoft.com/en-us/cli/azure/cognitiveservices/account?view=azure-cli-latest>
az cognitiveservices account create -n myresource -g myResourceGroup --assign-identity --kind TextAnalytics --sku S -l WestEurope --yes
--encryption '{
 "keySource": "Microsoft.KeyVault",
 "keyVaultProperties": {
 "keyName": "KeyName",
 "keyVersion": "secretVersion",
 "keyVaultUri": "https://issue23056kv.vault.azure.net/"
 }
}'
upvoted 1 times

 **Murtaza** 1 month, 3 weeks ago

az cognitiveservices account create \
--name myresource \
--resource-group myResourceGroup \
--kind OpenAI \
'

```
--sku S1 \
--location WestEurope \
upvoted 2 times
```

✉ **Ghill1982** 2 months, 2 weeks ago

Answer is correct.
upvoted 1 times

✉ **Murtuza** 2 months, 2 weeks ago

Azure OpenAI Service models can do everything from generating original stories to performing complex text an
upvoted 1 times

✉ **chandiochan** 2 months, 3 weeks ago

<https://learn.microsoft.com/en-us/azure/ai-services/openai/encrypt-data-at-rest>
<https://learn.microsoft.com/en-us/cli/azure/cognitiveservices/account?view=azure-cli-latest#az-cognitiveservices-account-create>
upvoted 2 times

Topic 8 - Testlet 1

Question #1

Topic 8

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

A company named Wide World Importers is developing an e-commerce platform.

You are working with a solutions architect to design and implement the features of the e-commerce platform. The platform will use microservices and a serverless environment built on Azure.

Wide World Importers has a customer base that includes English, Spanish, and Portuguese speakers.

Existing Environment -

Applications -

Wide World Importers has an App Service plan that contains the web apps shown in the following table.

Name	Description
Product Management	An app used by employees to create and manage products. The app and the expected inputs from the employees are in English.
Inventory Tracking	An app used by employees to manage inventory when dispatching orders, receiving refunds, and receiving consignments from suppliers.

Azure Resources -

You have the following resources:

An Azure Active Directory (Azure AD) tenant

- The tenant supports internal authentication.

- All employees belong to a group named AllUsers.

- Senior managers belong to a group named LeadershipTeam.

An Azure Functions resource

- A function app posts to Azure Event Grid when stock levels of a product change between OK, Low Stock, and Out of Stock. The function app uses the

Azure Cosmos DB change feed.

An Azure Cosmos DB account

- The account uses the Core (SQL) API.

- The account stores data for the Product Management app and the Inventory Tracking app.

An Azure Storage account

- The account contains blob containers for assets related to products.

- The assets include images, videos, and PDFs.

An Azure Cognitive Services resource named wwics

An Azure Video Analyzer for Media (previously Video Indexer) resource named wwivi

Requirements -

Business Goals -

Wide World Importers wants to leverage AI technologies to differentiate itself from its competitors.

Planned Changes -

Wide World Importers plans to start the following projects:

A product creation project: Help employees create accessible and multilingual product entries, while expediting product entry creation.

A smart e-commerce project: Implement an Azure Cognitive Search solution to display products for customers to browse.

A shopping on-the-go project: Build a chatbot that can be integrated into smart speakers to support customers.

Business Requirements -

Wide World Importers identifies the following business requirements for all the projects:

Provide a multilingual customer experience that supports English, Spanish, and Portuguese.

Whenever possible, scale based on transaction volumes to ensure consistent performance.

Minimize costs.

Governance and Security Requirements

Wide World Importers identifies the following governance and security requirements:

Data storage and processing must occur in datacenters located in the United States.

Azure Cognitive Services must be inaccessible directly from the internet.

Accessibility Requirements -

Wide World Importers identifies the following accessibility requirements:

All images must have relevant alt text.

All videos must have transcripts that are associated to the video and included in product descriptions.

Product descriptions, transcripts, and alt text must be available in English, Spanish, and Portuguese.

Product Creation Requirements -

Wide World Importers identifies the following requirements for improving the Product Management app:

Minimize how long it takes for employees to create products and add assets.

Remove the need for manual translations.

Smart E-Commerce Requirements -

Wide World Importers identifies the following requirements for the smart e-commerce project:

Ensure that the Cognitive Search solution meets a Service Level Agreement (SLA) of 99.9% availability for searches and index writes.

Provide users with the ability to search insight gained from the images, manuals, and videos associated with the products.

Support autocompletion and suggestion based on all product name variants.

Store all raw insight data that was generated, so the data can be processed later.

Update the stock level field in the product index immediately upon changes.

Update the product index hourly.

Shopping On-the-Go Requirements -

Wide World Importers identifies the following requirements for the shopping on-the-go chatbot:

Answer common questions.

Support interactions in English, Spanish, and Portuguese.

Replace an existing FAQ process so that all Q&A is managed from a central location.

Provide all employees with the ability to edit Q&As. Only senior managers must be able to publish updates.

Support purchases by providing information about relevant products to customers. Product displays must include images and warnings when stock levels are low or out of stock.

Product JSON Sample -

You have the following JSON sample for a product.

```
{
    "sku": "b1",
    "name": {
        "en": "Bicycle",
        "es": "Bicicleta",
        "pt": "Bicicleta"
    },
    "stocklevel": "Out of Stock",
    "description": {
        "en": "Bicycle",
        "es": "Bicicleta",
        "pt": "Bicicleta"
    },
    "image": {
        "uri": "https://upload.worldwideimporters.org/bicycle.jpg",
        "alttext": {
            "en": "Bicycle",
            "es": "Bicicleta",
            "pt": "Bicicleta"
        }
    },
    "createdUtc": "2020-02-14T06:08:39Z",
    "language": "en"
}
```

Question

DRAG DROP -

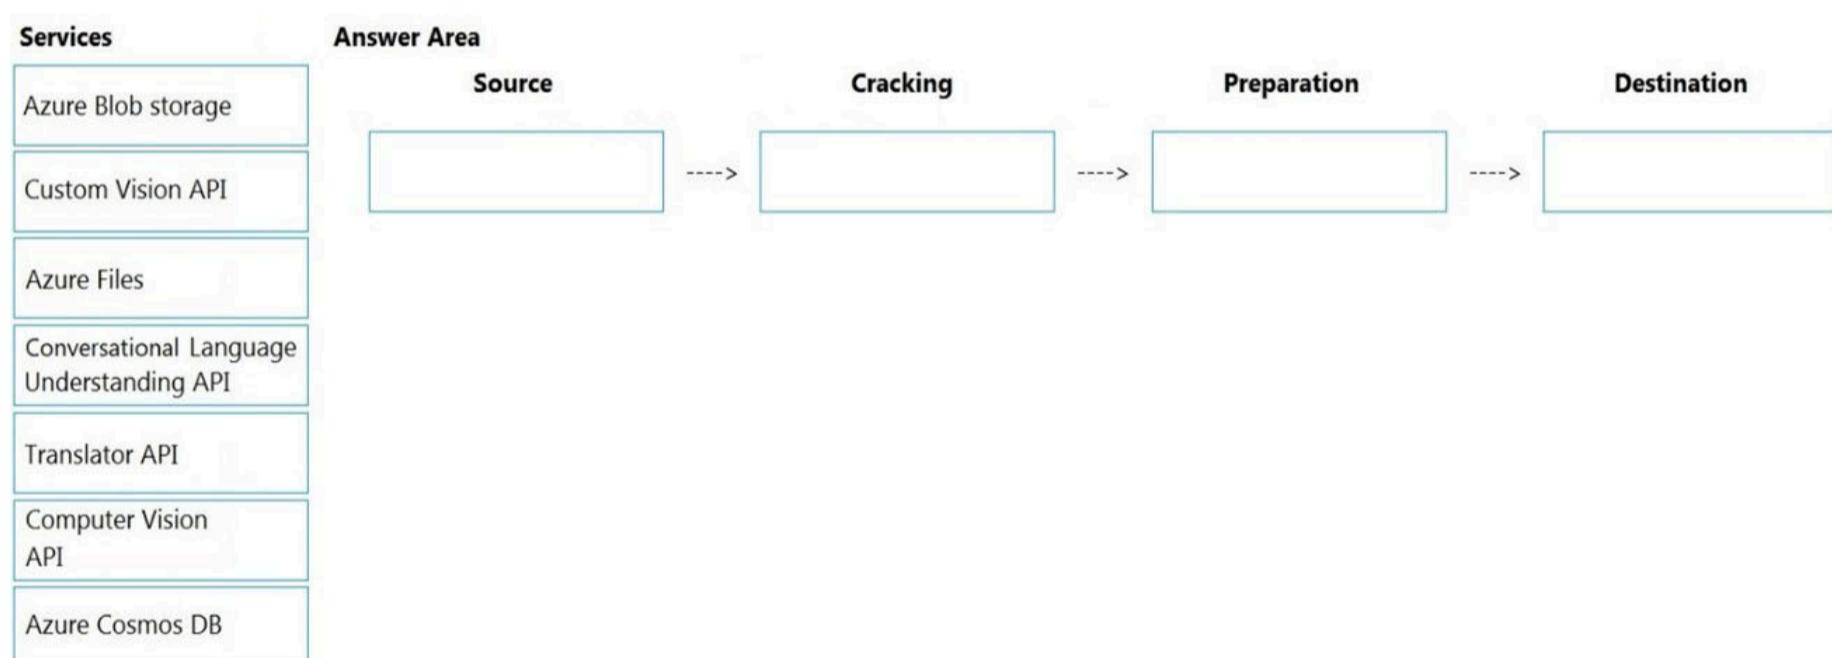
You are developing the smart e-commerce project.

You need to design the skillset to include the contents of PDFs in searches.

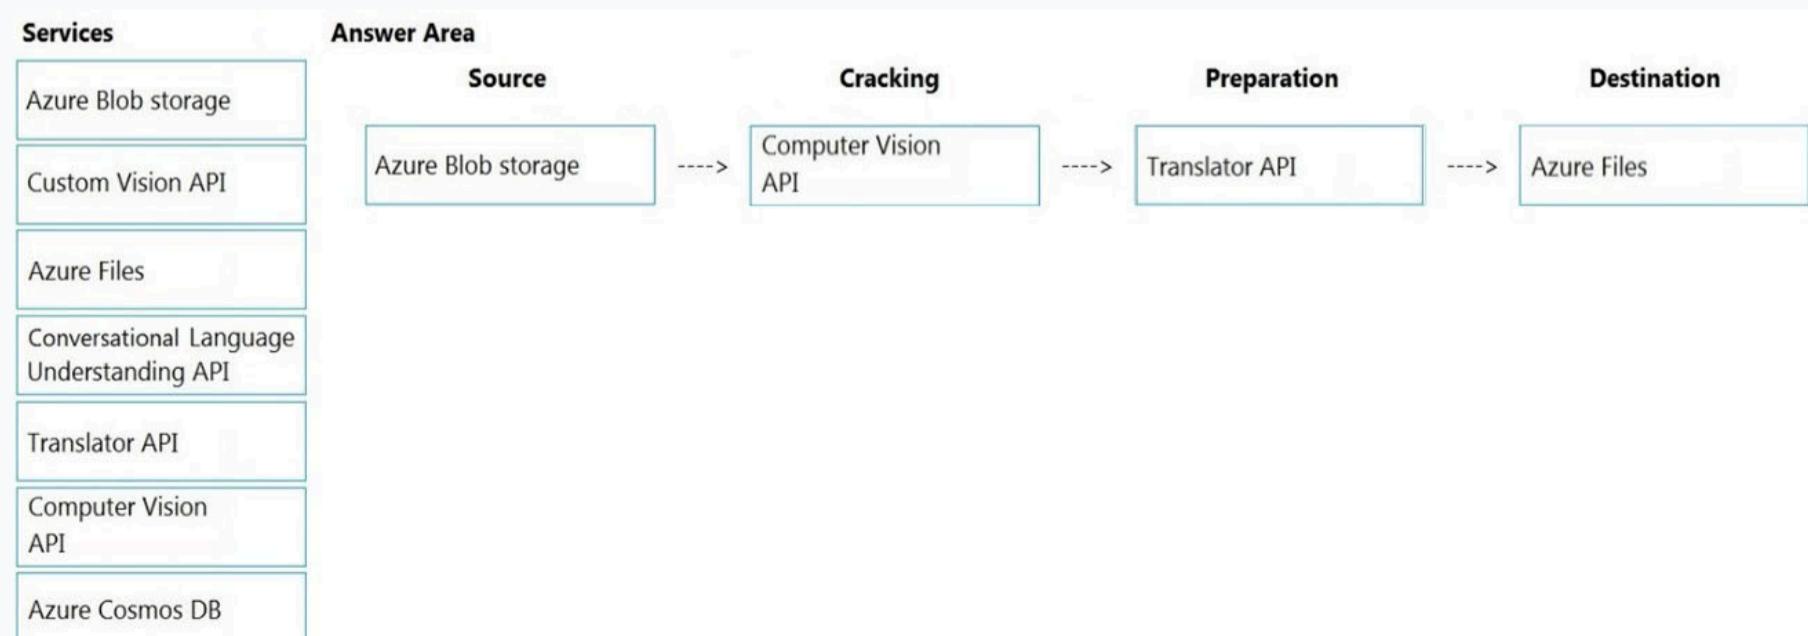
How should you complete the skillset design diagram? To answer, drag the appropriate services to the correct stages. Each service may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:



Correct Answer:



Box 1: Azure Blob storage -

At the start of the pipeline, you have unstructured text or non-text content (such as images, scanned documents, or JPEG files). Data must exist

in an Azure data storage service that can be accessed by an indexer.

Box 2: Computer Vision API -

Scenario: Provide users with the ability to search insight gained from the images, manuals, and videos associated with the products.

The Computer Vision Read API is Azure's latest OCR technology (learn what's new) that extracts printed text (in several languages), handwritten text (English only), digits, and currency symbols from images and multi-page PDF documents.

Box 3: Translator API -

Scenario: Product descriptions, transcripts, and alt text must be available in English, Spanish, and Portuguese.

Box 4: Azure Files -

Scenario: Store all raw insight data that was generated, so the data can be processed later.

Incorrect Answers:

The custom vision API from Microsoft Azure learns to recognize specific content in imagery and becomes smarter with training and time.

Reference:

<https://docs.microsoft.com/en-us/azure/search/cognitive-search-concept-intro> <https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/overview-ocr>

✉️  M25  8 months, 3 weeks ago

Source [Azure Blob Storage]

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-concept-intro>

Import is the first step. Here, the indexer connects to a data source and pulls content (documents) into the search service. Azure Blob Storage is the most common resource used in AI enrichment scenarios, but any supported data source can provide content.

<https://learn.microsoft.com/en-us/azure/search/search-indexer-overview#document-cracking>

- When the document is a file with embedded images, such as a PDF, the indexer extracts text, images, and metadata. Indexers can open files from Azure Blob Storage, Azure Data Lake Storage Gen2, and SharePoint.

upvoted 7 times

✉️  M25 8 months, 3 weeks ago

--> Cracking [Computer Vision API] --> Preparation [Translator API]

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-concept-intro>

Built-in skills are based on the Azure AI services APIs: Azure AI Computer Vision and Language Service.

- Image processing skills include Optical Character Recognition (OCR) and identification of visual features, such as facial detection, image interpretation, image recognition (famous people and landmarks), or attributes like image orientation. These skills create text representations of image content for full text search in Azure Cognitive Search.
- Machine translation is provided by the Text Translation skill, often paired with language detection for multi-language solutions.

upvoted 6 times

✉️  M25 8 months, 3 weeks ago

--> Destination [Azure Blob Storage]

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-concept-intro>

- Enriched content is generated during skillset execution, and is temporary unless you save it. You can enable an enrichment cache [Physically, the cache is stored in a blob container in your Azure Storage account, one per indexer.] to persist cracked documents and skill outputs for subsequent reuse during future skillset executions.

Exploration is the last step. Output is always a search index that you can query from a client app. Output can optionally be a knowledge store consisting of blobs and tables in Azure Storage that are accessed through data exploration tools or downstream processes. If you're creating a knowledge store, projections determine the data path for enriched content. The same enriched content can appear in both indexes and knowledge stores.

upvoted 8 times

✉️  rdemontis 6 months, 2 weeks ago

thanks for explanation

upvoted 3 times

✉️  Sharks82  1 year, 8 months ago

Given answer is correct

upvoted 4 times

Topic 9 - Testlet 2

Question #1

Topic 9

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

General Overview -

Contoso, Ltd. is an international accounting company that has offices in France, Portugal, and the United Kingdom.

Contoso has a professional services department that contains the roles shown in the following table.

Name	Position	Office
Accountant	Manager	United Kingdom, France, Portugal
Accountant	Consultant	United Kingdom, France, Portugal
Customer Service	Manager	United Kingdom
Customer Service	Agent	United Kingdom
Bookkeeper	Manager	United Kingdom, France, Portugal
Bookkeeper	Consultant	United Kingdom, France, Portugal

Existing environment -

Infrastructure -

Contoso has the following subscriptions:

Azure

Microsoft 365

Microsoft Dynamics 365

Azure Active (Azure AD) Directory

Contoso has Azure Active Directory groups for securing role-based access. The company uses the following group naming conventions:

[Country]-[Level]-[Role]

[Level]-[Role]

Intellectual Property -

Contoso has the intellectual property shown in the following table.

Content	Format	Language	Content store	Domain
Weekly webinars	Video	English	Azure Blob storage	Vid.contoso.com
Blogs	Text	English, French, Portuguese	WordPress	Pt-blog.contoso.com Blog.contoso.com Fr-blog.contoso.com
Wikis	Text	English, French, Portuguese	Azure Cosmos DB	Internal.contoso.com/wiki
Monthly conference recordings	Video	English	SharePoint Online	Contoso.sharepoint.com
Frequently asked questions (FAQs)	Text	English	SharePoint Online	Contoso.sharepoint.com

Text-based content is provided only in one language and is not translated.

Requirements -

Planned Projects -

Contoso plans to develop the following:

A document processing workflow to extract information automatically from PDFs and images of financial documents

A customer-support chatbot that will answer questions by using FAQs

A searchable knowledgebase of all the intellectual property

Technical Requirements -

Contoso identifies the following technical requirements:

All content must be approved before being published.

All planned projects must support English, French, and Portuguese.

All content must be secured by using role-based access control (RBAC).

RBAC role assignments must use the principle of least privilege.

RBAC roles must be assigned only to Azure Active Directory groups.

AI solution responses must have a confidence score that is equal to or greater than 70 percent.

When the response confidence score of an AI response is lower than 70 percent, the response must be improved by human input.

Chatbot Requirements -

Contoso identifies the following requirements for the chatbot:

Provide customers with answers to the FAQs.

Ensure that the customers can chat to a customer service agent.

Ensure that the members of a group named Management-Accountants can approve the FAQs.

Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.

Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.

Ensure that access to the customer service agents is managed by using Omnichannel for Customer Service.

When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

Document Processing Requirements

Contoso identifies the following requirements for document processing:

The document processing solution must be able to process standardized financial documents that have the following characteristics:

- Contain fewer than 20 pages.

- Be formatted as PDF or JPEG files.

- Have a distinct standard for each office.

The document processing solution must be able to extract tables and text from the financial documents.

The document processing solution must be able to extract information from receipt images.

Members of a group named Management-Bookkeeper must define how to extract tables from the financial documents.

Members of a group named Consultant-Bookkeeper must be able to process the financial documents.

Knowledgebase Requirements -

Contoso identifies the following requirements for the knowledgebase:

Supports searches for equivalent terms

Can transcribe jargon with high accuracy

▪ Can search content in different formats, including video

Provides relevant links to external resources for further research

Question

HOTSPOT -

You build a QnA Maker resource to meet the chatbot requirements.

Which RBAC role should you assign to each group? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Management-Accountants

<input type="checkbox"/>
Owner
Contributor
Cognitive Services User
Cognitive Services QnA Maker Read
Cognitive Services QnA Maker Editor

Consultant-Accountants

<input type="checkbox"/>
Owner
Contributor
Cognitive Services User
Cognitive Services QnA Maker Read
Cognitive Services QnA Maker Editor

Agent-CustomerServices

<input type="checkbox"/>
Owner
Contributor
Cognitive Services User
Cognitive Services QnA Maker Read
Cognitive Services QnA Maker Editor

Correct Answer:

Answer Area

Management-Accountants

Owner
Contributor
Cognitive Services User
Cognitive Services QnA Maker Read
Cognitive Services QnA Maker Editor

Consultant-Accountants

Owner
Contributor
Cognitive Services User
Cognitive Services QnA Maker Read
Cognitive Services QnA Maker Editor

Agent-CustomerServices

Owner
Contributor
Cognitive Services User
Cognitive Services QnA Maker Read
Cognitive Services QnA Maker Editor

Box 1: Cognitive Service User -

Ensure that the members of a group named Management-Accountants can approve the FAQs.

Approve=publish.

Cognitive Service User (read/write/publish): API permissions: All access to Cognitive Services resource except for ability to:

1. Add new members to roles.
2. Create new resources.

Box 2: Cognitive Services QnA Maker Editor

Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.

QnA Maker Editor: API permissions:

1. Create KB API
2. Update KB API
3. Replace KB API
4. Replace Alterations
5. "Train API" [in new service model v5]

Box 3: Cognitive Services QnA Maker Read

Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.

QnA Maker Read: API Permissions:

1. Download KB API
2. List KBs for user API
3. Get Knowledge base details
4. Download Alterations

Generate Answer -

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/concepts/role-based-access-control>

  **zellck**  10 months, 4 weeks ago

1. Cognitive Service User
2. QnA Maker Editor
3. QnA Maker Read

<https://learn.microsoft.com/en-us/azure/cognitive-services/qnamaker/concepts/role-based-access-control#access-is-provided-by-a-defined-role>

- Cognitive Service User (read/write/publish)
 - QnA Maker Editor (read/write)
 - QnA Maker Read (read)
- upvoted 13 times

✉️ **evangelist** Most Recent 3 months, 2 weeks ago

management account should use "Cognitive Service User" since only this role has permission to publish. Consultant should be able to edit and Agent should be only able to read

upvoted 2 times

✉️ **rdemontis** 6 months, 2 weeks ago

correct

upvoted 2 times

✉️ **ninja** 1 year, 9 months ago

The answer is correct, based on

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/concepts/role-based-access-control>

upvoted 4 times

Topic 10 - Testlet 3

Question #1

Topic 10

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

A company named Wide World Importers is developing an e-commerce platform.

You are working with a solutions architect to design and implement the features of the e-commerce platform. The platform will use microservices and a serverless environment built on Azure.

Wide World Importers has a customer base that includes English, Spanish, and Portuguese speakers.

Existing Environment -

Applications -

Wide World Importers has an App Service plan that contains the web apps shown in the following table.

Name	Description
Product Management	An app used by employees to create and manage products. The app and the expected inputs from the employees are in English.
Inventory Tracking	An app used by employees to manage inventory when dispatching orders, receiving refunds, and receiving consignments from suppliers.

Azure Resources -

You have the following resources:

An Azure Active Directory (Azure AD) tenant

- The tenant supports internal authentication.

- All employees belong to a group named AllUsers.

- Senior managers belong to a group named LeadershipTeam.

An Azure Functions resource

- A function app posts to Azure Event Grid when stock levels of a product change between OK, Low Stock, and Out of Stock. The function app uses the

Azure Cosmos DB change feed.

An Azure Cosmos DB account

- The account uses the Core (SQL) API.

- The account stores data for the Product Management app and the Inventory Tracking app.

An Azure Storage account

- The account contains blob containers for assets related to products.

- The assets include images, videos, and PDFs.

An Azure Cognitive Services resource named wwics

An Azure Video Analyzer for Media (previously Video Indexer) resource named wwivi

Requirements -

Business Goals -

Wide World Importers wants to leverage AI technologies to differentiate itself from its competitors.

Planned Changes -

Wide World Importers plans to start the following projects:

A product creation project: Help employees create accessible and multilingual product entries, while expediting product entry creation.

A smart e-commerce project: Implement an Azure Cognitive Search solution to display products for customers to browse.

A shopping on-the-go project: Build a chatbot that can be integrated into smart speakers to support customers.

Business Requirements -

Wide World Importers identifies the following business requirements for all the projects:

Provide a multilingual customer experience that supports English, Spanish, and Portuguese.

Whenever possible, scale based on transaction volumes to ensure consistent performance.

Minimize costs.

Governance and Security Requirements

Wide World Importers identifies the following governance and security requirements:

Data storage and processing must occur in datacenters located in the United States.

Azure Cognitive Services must be inaccessible directly from the internet.

Accessibility Requirements -

Wide World Importers identifies the following accessibility requirements:

All images must have relevant alt text.

All videos must have transcripts that are associated to the video and included in product descriptions.

Product descriptions, transcripts, and alt text must be available in English, Spanish, and Portuguese.

Product Creation Requirements -

Wide World Importers identifies the following requirements for improving the Product Management app:

Minimize how long it takes for employees to create products and add assets.

Remove the need for manual translations.

Smart E-Commerce Requirements -

Wide World Importers identifies the following requirements for the smart e-commerce project:

Ensure that the Cognitive Search solution meets a Service Level Agreement (SLA) of 99.9% availability for searches and index writes.

Provide users with the ability to search insight gained from the images, manuals, and videos associated with the products.

Support autocompletion and suggestion based on all product name variants.

Store all raw insight data that was generated, so the data can be processed later.

Update the stock level field in the product index immediately upon changes.

Update the product index hourly.

Shopping On-the-Go Requirements -

Wide World Importers identifies the following requirements for the shopping on-the-go chatbot:

Answer common questions.

Support interactions in English, Spanish, and Portuguese.

Replace an existing FAQ process so that all Q&A is managed from a central location.

Provide all employees with the ability to edit Q&As. Only senior managers must be able to publish updates.

Support purchases by providing information about relevant products to customers. Product displays must include images and warnings when stock levels are low or out of stock.

Product JSON Sample -

You have the following JSON sample for a product.

```
{  
    "sku": "b1",  
    "name": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "stocklevel": "Out of Stock",  
    "description": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "image":  
    {"uri": "https://upload.worldwideimporters.org/bicycle.jpg",  
        "alttext": {  
            "en": "Bicycle",  
            "es": "Bicicleta",  
            "pt": "Bicicleta"  
        }  
    },  
    "createdUtc": "2020-02-14T06:08:39Z",  
    "language": "en"  
}
```

Question

DRAG DROP -

You are planning the product creation project.

You need to recommend a process for analyzing videos.

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and

arrange them in the correct order.

Select and Place:

Actions

Index the video by using the Azure Video Analyzer for Media (previously Video Indexer) API.

Upload the video to blob storage.

Analyze the video by using the Computer Vision API.

Extract the transcript from Microsoft Stream.

Send the transcript to the Language Understanding API as an utterance.

Extract the transcript from the Azure Video Analyzer for Media (previously Video Indexer) API.

Translate the transcript by using the Translator API.

Upload the video to file storage.

Answer Area

Correct Answer:

Actions

Analyze the video by using the Computer Vision API.

Extract the transcript from Microsoft Stream.

Send the transcript to the Language Understanding API as an utterance.

Upload the video to file storage.

Answer Area

Upload the video to blob storage.

Index the video by using the Azure Video Analyzer for Media (previously Video Indexer) API.

Extract the transcript from the Azure Video Analyzer for Media (previously Video Indexer) API.

Translate the transcript by using the Translator API.

Scenario: All videos must have transcripts that are associated to the video and included in product descriptions.

Product descriptions, transcripts, and alt text must be available in English, Spanish, and Portuguese.

Step 1: Upload the video to blob storage

Given a video or audio file, the file is first dropped into a Blob Storage. T

Step 2: Index the video by using the Video Indexer API.

When a video is indexed, Video Indexer produces the JSON content that contains details of the specified video insights. The insights include: transcripts, OCRs, faces, topics, blocks, etc.

Step 3: Extract the transcript from the Video Indexer API.

Step 4: Translate the transcript by using the Translator API.

Reference:

<https://azure.microsoft.com/en-us/blog/get-video-insights-in-even-more-languages/> <https://docs.microsoft.com/en-us/azure/media-services/video-indexer/video-indexer-output-json-v2>

 **zellck** Highly Voted 10 months, 4 weeks ago

1. Upload video to blob storage.
2. Index video using Azure Video Indexer API.
3. Extract transcript from Azure Video Indexer API.
4. Translate transcript using Translator API.

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/transcription-translation-lid>
upvoted 17 times

 **ninja** Highly Voted 1 year, 9 months ago

1. Upload the video to blob storage. - choose blob over file
2. Index the video by using the Video Indexer API. - choose Video Indexer over Computer Vision API.
3. Extract the transcript from the Video Indexer API.
4. Translate the transcript by using the Translator API. - Support requirement "Remove the need for manual translations."

Reference:

<https://docs.microsoft.com/en-us/azure/azure-video-indexer/video-indexer-overview#what-can-i-do-with-azure-video-indexer>

upvoted 6 times

✉️ **am20** 1 year, 4 months ago

agree as the other options doesn't make sense. however videos are already in storage account

upvoted 2 times

✉️ **rdemontis** 6 months, 2 weeks ago

thanks for the reference

upvoted 1 times

✉️ **evangelist** **Most Recent** 4 months ago

(1) Upload the video to blob storage: Store the video in a reliable and scalable location like Azure Blob Storage. This enables access for subsequent analysis steps.

(2) Index the video by using the Azure Video Analyzer for Media (previously Video Indexer) API: Utilize Video Analyzer for Media to automatically extract key information like objects, actions, and emotions from the video. This provides valuable context for further analysis.

(3) Extract the transcript from the Azure Video Analyzer for Media (previously Video Indexer) API: Leverage the same Video Analyzer for Media API to generate a transcript of the spoken audio in the video. This enables textual analysis of the content.

(4) Translate the transcript by using the Translator API (optional): If your target audience uses languages other than the source language of the video, utilize the Translator API to translate the extracted transcript into the desired languages. This enables multilingual product descriptions or accessibility features.

upvoted 2 times

✉️ **tranatrana** 9 months ago

Scenario: All videos must have transcripts that are associated to the video and included in product descriptions.

Product descriptions, transcripts, and alt text must be available in English, Spanish, and Portuguese.

Otherwise, it does not make sense!

upvoted 2 times

✉️ **zellck** 10 months, 4 weeks ago

Same as Topic 7 Question 1.

<https://www.examtopics.com/discussions/microsoft/view/57566-exam-ai-102-topic-7-question-1-discussion>

upvoted 2 times

✉️ **ziggy1117** 11 months, 3 weeks ago

answer is correct

upvoted 3 times

Question #2

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

A company named Wide World Importers is developing an e-commerce platform.

You are working with a solutions architect to design and implement the features of the e-commerce platform. The platform will use microservices and a serverless environment built on Azure.

Wide World Importers has a customer base that includes English, Spanish, and Portuguese speakers.

Existing Environment -

Applications -

Wide World Importers has an App Service plan that contains the web apps shown in the following table.

Name	Description
Product Management	An app used by employees to create and manage products. The app and the expected inputs from the employees are in English.
Inventory Tracking	An app used by employees to manage inventory when dispatching orders, receiving refunds, and receiving consignments from suppliers.

Azure Resources -

You have the following resources:

An Azure Active Directory (Azure AD) tenant

- The tenant supports internal authentication.

- All employees belong to a group named AllUsers.

- Senior managers belong to a group named LeadershipTeam.

An Azure Functions resource

- A function app posts to Azure Event Grid when stock levels of a product change between OK, Low Stock, and Out of Stock. The function app uses the

Azure Cosmos DB change feed.

An Azure Cosmos DB account

- The account uses the Core (SQL) API.

- The account stores data for the Product Management app and the Inventory Tracking app.

An Azure Storage account

- The account contains blob containers for assets related to products.

- The assets include images, videos, and PDFs.

An Azure Cognitive Services resource named wwics

An Azure Video Analyzer for Media (previously Video Indexer) resource named wwivi

Requirements -

Business Goals -

Wide World Importers wants to leverage AI technologies to differentiate itself from its competitors.

Planned Changes -

Wide World Importers plans to start the following projects:

A product creation project: Help employees create accessible and multilingual product entries, while expediting product entry creation.

A smart e-commerce project: Implement an Azure Cognitive Search solution to display products for customers to browse.

A shopping on-the-go project: Build a chatbot that can be integrated into smart speakers to support customers.

Business Requirements -

Wide World Importers identifies the following business requirements for all the projects:

Provide a multilingual customer experience that supports English, Spanish, and Portuguese.

Whenever possible, scale based on transaction volumes to ensure consistent performance.

Minimize costs.

Governance and Security Requirements

Wide World Importers identifies the following governance and security requirements:

Data storage and processing must occur in datacenters located in the United States.

Azure Cognitive Services must be inaccessible directly from the internet.

Accessibility Requirements -

Wide World Importers identifies the following accessibility requirements:

All images must have relevant alt text.

All videos must have transcripts that are associated to the video and included in product descriptions.

Product descriptions, transcripts, and alt text must be available in English, Spanish, and Portuguese.

Product Creation Requirements -

Wide World Importers identifies the following requirements for improving the Product Management app:

Minimize how long it takes for employees to create products and add assets.

Remove the need for manual translations.

Smart E-Commerce Requirements -

Wide World Importers identifies the following requirements for the smart e-commerce project:

Ensure that the Cognitive Search solution meets a Service Level Agreement (SLA) of 99.9% availability for searches and index writes.

Provide users with the ability to search insight gained from the images, manuals, and videos associated with the products.

Support autocompletion and suggestion based on all product name variants.

Store all raw insight data that was generated, so the data can be processed later.

Update the stock level field in the product index immediately upon changes.

Update the product index hourly.

Shopping On-the-Go Requirements -

Wide World Importers identifies the following requirements for the shopping on-the-go chatbot:

Answer common questions.

Support interactions in English, Spanish, and Portuguese.

Replace an existing FAQ process so that all Q&A is managed from a central location.

Provide all employees with the ability to edit Q&As. Only senior managers must be able to publish updates.

Support purchases by providing information about relevant products to customers. Product displays must include images and warnings when stock levels are low or out of stock.

Product JSON Sample -

You have the following JSON sample for a product.

```
{  
    "sku": "bl",  
    "name": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "stocklevel": "Out of Stock",  
    "description": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "image":  
    {"uri": "https://upload.worldwideimporters.org/bicycle.jpg",  
        "alttext": {  
            "en": "Bicycle",  
            "es": "Bicicleta",  
            "pt": "Bicicleta"  
        }  
    },  
    "createdUtc": "2020-02-14T06:08:39Z",  
    "language": "en"  
}
```

Question

HOTSPOT -

You need to develop code to upload images for the product creation project. The solution must meet the accessibility requirements.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
public static async Task<string> SuggestAltText(ComputerVisionClient client,  
{  
    List<VisualFeatureTypes?> features = new List<VisualFeatureTypes?>()  
    {  
        VisualFeatureTypes.Description  
        VisualFeatureTypes.ImageType  
        VisualFeatureTypes.Objects  
        VisualFeatureTypes.Tags  
    };  
    ImageAnalysis results = await client.AnalyzeImageAsync(image, features);  
  
    var c = results.Brands.DetectedBrands[0]  
    var c = results.Description.Captions[0]  
    var c = results.Metadata[0]  
    var c = results.Objects[0]  
  
    if(c.Confidence>0.5) return(c.Text);  
}
```

Dictionary
stream
string

Answer Area

```

public static async Task<string> SuggestAltText(ComputerVisionClient client,
{
    List<VisualFeatureTypes?> features = new List<VisualFeatureTypes?>()
    {
        VisualFeatureTypes.Description
        VisualFeatureTypes.ImageType
        VisualFeatureTypes.Objects
        VisualFeatureTypes.Tags
    };
    ImageAnalysis results = await client.AnalyzeImageAsync(image, features);
}

var c = results.Brands.DetectedBrands[0]
var c = results.Description.Captions[0]
var c = results.Metadata[0]
var c = results.Objects[0]

if(c.Confidence>0.5) return(c.Text);
}

```

Dictionary
stream
string

image)

Correct Answer:**Reference:**

<https://github.com/Azure-Samples/cognitive-services-dotnet-sdk-samples/blob/master/documentation-samples/quickstarts/ComputerVision/Program.cs>

✉  **zellck**  10 months, 4 weeks ago

1. string
2. VisualFeatureTypes.Description
3. results.Description.Captions[0]

<https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-describing-images#use-the-api>

The image description feature is part of the Analyze Image API. You can call this API through a native SDK or through REST calls. Include Description in the visualFeatures query parameter. Then, when you get the full JSON response, parse the string for the contents of the "description" section.

upvoted 16 times

✉  **M25** 8 months, 3 weeks ago

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/use-case-alt-text#what-languages-are-image-captions-available-in>
Image captions are available in English, Chinese, Portuguese, Japanese, and Spanish in Image Analysis 3.2 API. In the Image Analysis 4.0 API (preview), image captions are only available in English.

upvoted 3 times

✉  **rdemontis** 6 months, 2 weeks ago

thanks for explanation

upvoted 1 times

✉  **evangelist** 3 months, 2 weeks ago

the correct choice between Stream and string depends on the source of the image:

If you have the image file and want to upload it directly from the client or server where the request is being made, you should use a Stream. This approach is direct and efficient for handling binary data without the need for encoding or hosting the image externally.

If the image is already hosted online and accessible via a URL, you use a string containing the URL to let the API fetch the image for analysis.
upvoted 5 times

✉  **evangelist**  3 months, 2 weeks ago

given answer is correct

upvoted 1 times

✉  **JTWang** 2 years ago

AnalyzeImageSample

<https://github.com/Azure-Samples/cognitive-services-dotnet-sdk-samples/blob/master/samples/ComputerVision/AnalyzeImage/Program.cs>
upvoted 2 times

✉  **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 4 times

✉  **Diem** 2 years, 8 months ago

Correct!

upvoted 4 times

Topic 11 - Testlet 4

Question #1

Topic 11

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

General Overview -

Contoso, Ltd. is an international accounting company that has offices in France, Portugal, and the United Kingdom.

Contoso has a professional services department that contains the roles shown in the following table.

Name	Position	Office
Accountant	Manager	United Kingdom, France, Portugal
Accountant	Consultant	United Kingdom, France, Portugal
Customer Service	Manager	United Kingdom
Customer Service	Agent	United Kingdom
Bookkeeper	Manager	United Kingdom, France, Portugal
Bookkeeper	Consultant	United Kingdom, France, Portugal

Existing environment -

Infrastructure -

Contoso has the following subscriptions:

Azure

Microsoft 365

Microsoft Dynamics 365

Azure Active (Azure AD) Directory

Contoso has Azure Active Directory groups for securing role-based access. The company uses the following group naming conventions:

[Country]-[Level]-[Role]

[Level]-[Role]

Intellectual Property -

Contoso has the intellectual property shown in the following table.

Content	Format	Language	Content store	Domain
Weekly webinars	Video	English	Azure Blob storage	Vid.contoso.com
Blogs	Text	English, French, Portuguese	WordPress	Pt-blog.contoso.com Blog.contoso.com Fr-blog.contoso.com
Wikis	Text	English, French, Portuguese	Azure Cosmos DB	Internal.contoso.com/wiki
Monthly conference recordings	Video	English	SharePoint Online	Contoso.sharepoint.com
Frequently asked questions (FAQs)	Text	English	SharePoint Online	Contoso.sharepoint.com

Text-based content is provided only in one language and is not translated.

Requirements -

Planned Projects -

Contoso plans to develop the following:

A document processing workflow to extract information automatically from PDFs and images of financial documents

A customer-support chatbot that will answer questions by using FAQs

A searchable knowledgebase of all the intellectual property

Technical Requirements -

Contoso identifies the following technical requirements:

All content must be approved before being published.

All planned projects must support English, French, and Portuguese.

All content must be secured by using role-based access control (RBAC).

RBAC role assignments must use the principle of least privilege.

RBAC roles must be assigned only to Azure Active Directory groups.

AI solution responses must have a confidence score that is equal to or greater than 70 percent.

When the response confidence score of an AI response is lower than 70 percent, the response must be improved by human input.

Chatbot Requirements -

Contoso identifies the following requirements for the chatbot:

Provide customers with answers to the FAQs.

Ensure that the customers can chat to a customer service agent.

Ensure that the members of a group named Management-Accountants can approve the FAQs.

Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.

Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.

Ensure that access to the customer service agents is managed by using Omnichannel for Customer Service.

When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

Document Processing Requirements

Contoso identifies the following requirements for document processing:

The document processing solution must be able to process standardized financial documents that have the following characteristics:

- Contain fewer than 20 pages.

- Be formatted as PDF or JPEG files.

- Have a distinct standard for each office.

The document processing solution must be able to extract tables and text from the financial documents.

The document processing solution must be able to extract information from receipt images.

Members of a group named Management-Bookkeeper must define how to extract tables from the financial documents.

Members of a group named Consultant-Bookkeeper must be able to process the financial documents.

Knowledgebase Requirements -

Contoso identifies the following requirements for the knowledgebase:

Supports searches for equivalent terms

Can transcribe jargon with high accuracy

Can search content in different formats, including video

Provides relevant links to external resources for further research

Question

DRAG DROP -

You are developing a solution for the Management-Bookkeepers group to meet the document processing requirements. The solution must contain the following components:

A Form Recognizer resource

An Azure web app that hosts the Form Recognizer sample labeling tool

The Management-Bookkeepers group needs to create a custom table extractor by using the sample labeling tool.

Which three actions should the Management-Bookkeepers group perform in sequence? To answer, move the appropriate cmdlets from the list of cmdlets to the answer area and arrange them in the correct order.

Select and Place:

Actions	Answer Area
Train a custom model	
Label the sample documents	>
Create a new project and load sample documents	<
Create a composite model	

Correct Answer:

Actions	Answer Area
	Create a new project and load sample documents
	>
	Label the sample documents
	<
Create a composite model	Train a custom model

Step 1: Create a new project and load sample documents

Create a new project. Projects store your configurations and settings.

Step 2: Label the sample documents

When you create or open a project, the main tag editor window opens.

Step 3: Train a custom model.

Finally, train a custom model.

Reference:

<https://docs.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/label-tool>

 **Eltooth**  1 year, 10 months ago

Answer is correct.

Create a project

Label the sample docs

Train the model

upvoted 13 times

- ✉️  **Eltooth** 1 year, 10 months ago
<https://docs.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/label-tool#create-a-new-project>
upvoted 3 times
- ✉️  **rdemontis** 6 months, 2 weeks ago
Agree with you. Thanks for the reference
upvoted 2 times
- ✉️  **takaimomoGcup** Most Recent 1 week ago
Is this question still available on May 21, 2024?
upvoted 1 times
- ✉️  **[Removed]** 3 months, 2 weeks ago
Exam support with great ease passing
Exam booking and clearance
All IT certifications possible
Payment after passing only
WhatsApp +1 9402685570
upvoted 1 times
- ✉️  **evangelist** 3 months, 2 weeks ago
The requireemnt has nothing to do with composite model given answer is correct
upvoted 2 times
- ✉️  **Gvalli** 5 months, 1 week ago
A modified version of this was in the exam today.
upvoted 2 times

Question #2

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

General Overview -

Contoso, Ltd. is an international accounting company that has offices in France, Portugal, and the United Kingdom.

Contoso has a professional services department that contains the roles shown in the following table.

Name	Position	Office
Accountant	Manager	United Kingdom, France, Portugal
Accountant	Consultant	United Kingdom, France, Portugal
Customer Service	Manager	United Kingdom
Customer Service	Agent	United Kingdom
Bookkeeper	Manager	United Kingdom, France, Portugal
Bookkeeper	Consultant	United Kingdom, France, Portugal

Existing environment -

Infrastructure -

Contoso has the following subscriptions:

Azure

Microsoft 365

Microsoft Dynamics 365

Azure Active (Azure AD) Directory

Contoso has Azure Active Directory groups for securing role-based access. The company uses the following group naming conventions:

[Country]-[Level]-[Role]

[Level]-[Role]

Intellectual Property -

Contoso has the intellectual property shown in the following table.

Content	Format	Language	Content store	Domain
Weekly webinars	Video	English	Azure Blob storage	Vid.contoso.com
Blogs	Text	English, French, Portuguese	WordPress	Pt-blog.contoso.com Blog.contoso.com Fr-blog.contoso.com
Wikis	Text	English, French, Portuguese	Azure Cosmos DB	Internal.contoso.com/wiki
Monthly conference recordings	Video	English	SharePoint Online	Contoso.sharepoint.com
Frequently asked questions (FAQs)	Text	English	SharePoint Online	Contoso.sharepoint.com

Text-based content is provided only in one language and is not translated.

Requirements -

Planned Projects -

Contoso plans to develop the following:

A document processing workflow to extract information automatically from PDFs and images of financial documents

A customer-support chatbot that will answer questions by using FAQs

A searchable knowledgebase of all the intellectual property

Technical Requirements -

Contoso identifies the following technical requirements:

All content must be approved before being published.

All planned projects must support English, French, and Portuguese.

All content must be secured by using role-based access control (RBAC).

RBAC role assignments must use the principle of least privilege.

RBAC roles must be assigned only to Azure Active Directory groups.

AI solution responses must have a confidence score that is equal to or greater than 70 percent.

When the response confidence score of an AI response is lower than 70 percent, the response must be improved by human input.

Chatbot Requirements -

Contoso identifies the following requirements for the chatbot:

Provide customers with answers to the FAQs.

Ensure that the customers can chat to a customer service agent.

Ensure that the members of a group named Management-Accountants can approve the FAQs.

Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.

Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.

Ensure that access to the customer service agents is managed by using Omnichannel for Customer Service.

When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

Document Processing Requirements

Contoso identifies the following requirements for document processing:

The document processing solution must be able to process standardized financial documents that have the following characteristics:

- Contain fewer than 20 pages.

- Be formatted as PDF or JPEG files.

- Have a distinct standard for each office.

The document processing solution must be able to extract tables and text from the financial documents.

The document processing solution must be able to extract information from receipt images.

Members of a group named Management-Bookkeeper must define how to extract tables from the financial documents.

Members of a group named Consultant-Bookkeeper must be able to process the financial documents.

Knowledgebase Requirements -

Contoso identifies the following requirements for the knowledgebase:

Supports searches for equivalent terms

Can transcribe jargon with high accuracy

▪

Can search content in different formats, including video

Provides relevant links to external resources for further research

Question

You are developing the knowledgebase.

You use Azure Video Analyzer for Media (previously Video indexer) to obtain transcripts of webinars.

You need to ensure that the solution meets the knowledgebase requirements.

What should you do?

- A. Create a custom language model
- B. Configure audio indexing for videos only
- C. Enable multi-language detection for videos
- D. Build a custom Person model for webinar presenters

Correct Answer: B

Can search content in different formats, including video

Audio and video insights (multi-channels). When indexing by one channel, partial result for those models will be available.

Keywords extraction: Extracts keywords from speech and visual text.

Named entities extraction: Extracts brands, locations, and people from speech and visual text via natural language processing (NLP).

Topic inference: Makes inference of main topics from transcripts. The 2nd-level IPTC taxonomy is included.

Artifacts: Extracts rich set of "next level of details" artifacts for each of the models.

Sentiment analysis: Identifies positive, negative, and neutral sentiments from speech and visual text.

Incorrect Answers:

C: Webinars Videos are in English.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-video-analyzer/video-analyzer-for-media-docs/video-indexer-overview>

Community vote distribution

A (100%)

✉  **james2033**  9 months, 1 week ago

Selected Answer: A

Keyword "jargon", so choose "A. Create a custom language model".

upvoted 12 times

✉  **zellick**  10 months, 4 weeks ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/customize-language-model-overview>

Azure Video Indexer supports automatic speech recognition through integration with the Microsoft Custom Speech Service. You can customize the Language model by uploading adaptation text, namely text from the domain whose vocabulary you'd like the engine to adapt to. Once you train your model, new words appearing in the adaptation text will be recognized, assuming default pronunciation, and the Language model will learn new probable sequences of words.

upvoted 10 times

✉  **zellick** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 5 times

✉  **evangelist** 3 months, 2 weeks ago

the correct choice between Stream and string depends on the source of the image:

If you have the image file and want to upload it directly from the client or server where the request is being made, you should use a Stream. This approach is direct and efficient for handling binary data without the need for encoding or hosting the image externally.

If the image is already hosted online and accessible via a URL, you use a string containing the URL to let the API fetch the image for analysis.
upvoted 1 times

 **takaimomoGcup** Most Recent 1 week ago

Is this question still available on May 21, 2024?

upvoted 1 times

 **NullVoider_0** 1 month, 3 weeks ago

Question is incomplete and needs more context to determine the answer.

upvoted 2 times

 **evangelist** 3 months, 2 weeks ago

Selected Answer: A

A. Create a custom language model: This option allows for the customization of the language model used for transcribing audio and video content. By creating a custom language model, you can train it to understand and transcribe specialized jargon and industry-specific terminology with high accuracy. This directly addresses the requirement to transcribe jargon with high accuracy and supports searches for equivalent terms by ensuring that the transcriptions are as accurate and relevant as possible.

upvoted 1 times

 **rdemontis** 6 months, 2 weeks ago

Selected Answer: A

A is the answer. We need to support jargon.

<https://learn.microsoft.com/en-us/azure/azure-video-indexer/customize-language-model-overview>

upvoted 3 times

 **mk1967** 1 year, 8 months ago

Selected Answer: A

The requirement "Can transcribe jargon with high accuracy" wouldn't be met with the answer B.

upvoted 3 times

 **Ajose0** 1 year, 9 months ago

Selected Answer: A

--- Can transcribe jargon with high accuracy

Video Indexer (VI), the AI service for Azure Media Services enables the customization of language models by allowing customers to upload examples of sentences or words belonging to the vocabulary of their specific use case. Since speech recognition can sometimes be tricky, VI enables you to train and adapt the models for your specific domain. Harnessing this capability allows organizations to improve the accuracy of the Video Indexer generated transcriptions in their accounts.

<https://azure.microsoft.com/en-us/blog/new-ways-to-train-custom-language-models-effortlessly/>

upvoted 3 times

 **Eltooth** 1 year, 10 months ago

Selected Answer: A

Leaning towards A on this.

<https://azure.microsoft.com/en-us/blog/new-ways-to-train-custom-language-models-effortlessly/>

<https://docs.microsoft.com/en-us/azure/azure-video-indexer/video-indexer-overview>

upvoted 1 times

 **AiEngineerS** 1 year, 10 months ago

I think this is B... because even A makes sense... but u have to extract this information from video somehow. So B

upvoted 1 times

 **ANIKI51419** 1 year, 11 months ago

should be A

upvoted 3 times

 **sdokmak** 1 year, 11 months ago

I think you're right

<https://azure.microsoft.com/en-us/blog/new-ways-to-train-custom-language-models-effortlessly/>

upvoted 2 times

Topic 12 - Testlet 5

Question #1

Topic 12

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

A company named Wide World Importers is developing an e-commerce platform.

You are working with a solutions architect to design and implement the features of the e-commerce platform. The platform will use microservices and a serverless environment built on Azure.

Wide World Importers has a customer base that includes English, Spanish, and Portuguese speakers.

Existing Environment -

Applications -

Wide World Importers has an App Service plan that contains the web apps shown in the following table.

Name	Description
Product Management	An app used by employees to create and manage products. The app and the expected inputs from the employees are in English.
Inventory Tracking	An app used by employees to manage inventory when dispatching orders, receiving refunds, and receiving consignments from suppliers.

Azure Resources -

You have the following resources:

An Azure Active Directory (Azure AD) tenant

- The tenant supports internal authentication.

- All employees belong to a group named AllUsers.

- Senior managers belong to a group named LeadershipTeam.

An Azure Functions resource

- A function app posts to Azure Event Grid when stock levels of a product change between OK, Low Stock, and Out of Stock. The function app uses the

Azure Cosmos DB change feed.

An Azure Cosmos DB account

- The account uses the Core (SQL) API.

- The account stores data for the Product Management app and the Inventory Tracking app.

An Azure Storage account

- The account contains blob containers for assets related to products.

- The assets include images, videos, and PDFs.

An Azure Cognitive Services resource named wwics

An Azure Video Analyzer for Media (previously Video Indexer) resource named wwivi

Requirements -

Business Goals -

Wide World Importers wants to leverage AI technologies to differentiate itself from its competitors.

Planned Changes -

Wide World Importers plans to start the following projects:

A product creation project: Help employees create accessible and multilingual product entries, while expediting product entry creation.

A smart e-commerce project: Implement an Azure Cognitive Search solution to display products for customers to browse.

A shopping on-the-go project: Build a chatbot that can be integrated into smart speakers to support customers.

Business Requirements -

Wide World Importers identifies the following business requirements for all the projects:

Provide a multilingual customer experience that supports English, Spanish, and Portuguese.

Whenever possible, scale based on transaction volumes to ensure consistent performance.

Minimize costs.

Governance and Security Requirements

Wide World Importers identifies the following governance and security requirements:

Data storage and processing must occur in datacenters located in the United States.

Azure Cognitive Services must be inaccessible directly from the internet.

Accessibility Requirements -

Wide World Importers identifies the following accessibility requirements:

All images must have relevant alt text.

All videos must have transcripts that are associated to the video and included in product descriptions.

Product descriptions, transcripts, and alt text must be available in English, Spanish, and Portuguese.

Product Creation Requirements -

Wide World Importers identifies the following requirements for improving the Product Management app:

Minimize how long it takes for employees to create products and add assets.

Remove the need for manual translations.

Smart E-Commerce Requirements -

Wide World Importers identifies the following requirements for the smart e-commerce project:

Ensure that the Cognitive Search solution meets a Service Level Agreement (SLA) of 99.9% availability for searches and index writes.

Provide users with the ability to search insight gained from the images, manuals, and videos associated with the products.

Support autocompletion and suggestion based on all product name variants.

Store all raw insight data that was generated, so the data can be processed later.

Update the stock level field in the product index immediately upon changes.

Update the product index hourly.

Shopping On-the-Go Requirements -

Wide World Importers identifies the following requirements for the shopping on-the-go chatbot:

Answer common questions.

Support interactions in English, Spanish, and Portuguese.

Replace an existing FAQ process so that all Q&A is managed from a central location.

Provide all employees with the ability to edit Q&As. Only senior managers must be able to publish updates.

Support purchases by providing information about relevant products to customers. Product displays must include images and warnings when stock levels are low or out of stock.

Product JSON Sample -

You have the following JSON sample for a product.

```
{
  "sku": "b1",
  "name": {
    "en": "Bicycle",
    "es": "Bicicleta",
    "pt": "Bicicleta"
  },
  "stocklevel": "Out of Stock",
  "description": {
    "en": "Bicycle",
    "es": "Bicicleta",
    "pt": "Bicicleta"
  },
  "image": {
    "uri": "https://upload.worldwideimporters.org/bicycle.jpg",
    "alttext": {
      "en": "Bicycle",
      "es": "Bicicleta",
      "pt": "Bicicleta"
    }
  },
  "createdUtc": "2020-02-14T06:08:39Z",
  "language": "en"
}
```

Question

HOTSPOT -

You are planning the product creation project.

You need to build the REST endpoint to create the multilingual product descriptions.

How should you complete the URI? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

api.cognitive.microsofttranslator.com
api-nam.cognitive.microsofttranslator.com
westus.tts.speech.microsoft.com
wwics.cognitiveservices.azure.com/translator

?api-version=3.0&to=es&to=pt
/detect
/languages
/text-to-speech
/translate

Correct Answer:**Answer Area**

api.cognitive.microsofttranslator.com
api-nam.cognitive.microsofttranslator.com
westus.tts.speech.microsoft.com
wwics.cognitiveservices.azure.com/translator

?api-version=3.0&to=es&to=pt
/detect
/languages
/text-to-speech
/translate

Box 1: api.cognitive.microsofttranslator.com

Translator 3.0: Translate. Send a POST request to:

<https://api.cognitive.microsofttranslator.com/translate?api-version=3.0>

Box 2: /translate -

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-translate>

 **SuperPeteY** Highly Voted  2 years, 9 months ago

first drop-down should instead be: api-nam.cognitive.microsofttranslator.com this is because the case study specifically states under Business Requirements "Data storage and processing must occur in datacenters located in the United States."

see reference documentation for base urls per geo region:

<https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-reference>
upvoted 55 times

✉️ **ninja** 1 year, 9 months ago

Agreed
upvoted 1 times

✉️ **rdemontis** 6 months, 2 weeks ago

agree with you
upvoted 2 times

✉️ **josebernabeo** 4 months ago

Where do you read that? This is the question I see:
"HOTSPOT -
You are planning the product creation project.
You need to build the REST endpoint to create the multilingual product descriptions.
How should you complete the URI? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.
Hot Area:"
upvoted 2 times

✉️ **zellck** Highly Voted 10 months, 4 weeks ago

1. api-nam.cognitive.microsofttranslator.com
2. /translate

<https://learn.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-reference#base-urls>

Requests to Translator are, in most cases, handled by the datacenter that is closest to where the request originated. If there's a datacenter failure when using the global endpoint, the request may be routed outside of the geography.

To force the request to be handled within a specific geography, use the desired geographical endpoint. All requests are processed among the datacenters within the geography.

- United States
api-nam.cognitive.microsofttranslator.com
upvoted 10 times

✉️ **rdemontis** 6 months, 2 weeks ago

agree with you
upvoted 1 times

✉️ **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.
upvoted 3 times

✉️ **takaimomoGcup** Most Recent 1 week ago

Is this question still available on May 21, 2024?
upvoted 1 times

✉️ **evangelist** 3 months, 2 weeks ago

1. api-nam.cognitive.microsofttranslator.com
2. /translate
upvoted 1 times

✉️ **Gvalli** 5 months, 1 week ago

A modified version of this was in the exam today.
upvoted 1 times

✉️ **chimex1** 11 months, 3 weeks ago

Does the exam for AI 102 includes labs or its multiple choice only.
upvoted 2 times

✉️ **Eltooth** 1 year, 10 months ago

api-nam
translate
upvoted 3 times

✉️ **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021
upvoted 3 times

✉️ **gs23mi** 2 years, 7 months ago

in the "existing environment" there is
An Azure Cognitive Services resource named wwics
so it should be "resource name".cognitiveservices.azure.com : wwics.cognitiveservices.azure.com/translator/
Ref. <https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-reference>

upvoted 4 times

 **vominhtri854** 2 years, 7 months ago

but this for Customers with a resource located in Switzerland North or Switzerland West

upvoted 1 times

Topic 13 - Testlet 6

Question #1

Topic 13

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

General Overview -

Contoso, Ltd. is an international accounting company that has offices in France, Portugal, and the United Kingdom.

Contoso has a professional services department that contains the roles shown in the following table.

Name	Position	Office
Accountant	Manager	United Kingdom, France, Portugal
Accountant	Consultant	United Kingdom, France, Portugal
Customer Service	Manager	United Kingdom
Customer Service	Agent	United Kingdom
Bookkeeper	Manager	United Kingdom, France, Portugal
Bookkeeper	Consultant	United Kingdom, France, Portugal

Existing environment -

Infrastructure -

Contoso has the following subscriptions:

Azure

Microsoft 365

Microsoft Dynamics 365

Azure Active (Azure AD) Directory

Contoso has Azure Active Directory groups for securing role-based access. The company uses the following group naming conventions:

[Country]-[Level]-[Role]

[Level]-[Role]

Intellectual Property -

Contoso has the intellectual property shown in the following table.

Content	Format	Language	Content store	Domain
Weekly webinars	Video	English	Azure Blob storage	Vid.contoso.com
Blogs	Text	English, French, Portuguese	WordPress	Pt-blog.contoso.com Blog.contoso.com Fr-blog.contoso.com
Wikis	Text	English, French, Portuguese	Azure Cosmos DB	Internal.contoso.com/wiki
Monthly conference recordings	Video	English	SharePoint Online	Contoso.sharepoint.com
Frequently asked questions (FAQs)	Text	English	SharePoint Online	Contoso.sharepoint.com

Text-based content is provided only in one language and is not translated.

Requirements -

Planned Projects -

Contoso plans to develop the following:

A document processing workflow to extract information automatically from PDFs and images of financial documents

A customer-support chatbot that will answer questions by using FAQs

A searchable knowledgebase of all the intellectual property

Technical Requirements -

Contoso identifies the following technical requirements:

All content must be approved before being published.

All planned projects must support English, French, and Portuguese.

All content must be secured by using role-based access control (RBAC).

RBAC role assignments must use the principle of least privilege.

RBAC roles must be assigned only to Azure Active Directory groups.

AI solution responses must have a confidence score that is equal to or greater than 70 percent.

When the response confidence score of an AI response is lower than 70 percent, the response must be improved by human input.

Chatbot Requirements -

Contoso identifies the following requirements for the chatbot:

Provide customers with answers to the FAQs.

Ensure that the customers can chat to a customer service agent.

Ensure that the members of a group named Management-Accountants can approve the FAQs.

Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.

Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.

Ensure that access to the customer service agents is managed by using Omnichannel for Customer Service.

When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

Document Processing Requirements

Contoso identifies the following requirements for document processing:

The document processing solution must be able to process standardized financial documents that have the following characteristics:

- Contain fewer than 20 pages.

- Be formatted as PDF or JPEG files.

- Have a distinct standard for each office.

The document processing solution must be able to extract tables and text from the financial documents.

The document processing solution must be able to extract information from receipt images.

Members of a group named Management-Bookkeeper must define how to extract tables from the financial documents.

Members of a group named Consultant-Bookkeeper must be able to process the financial documents.

Knowledgebase Requirements -

Contoso identifies the following requirements for the knowledgebase:

Supports searches for equivalent terms

Can transcribe jargon with high accuracy

•

Can search content in different formats, including video

Provides relevant links to external resources for further research

Question

You need to develop an extract solution for the receipt images. The solution must meet the document processing requirements and the technical requirements.

You upload the receipt images to the Form Recognizer API for analysis, and the API returns the following JSON.

```
"documentResults": [
    {
        "docType": "prebuilt:receipt",
        "pageRange": [
            1,
            1
        ],
        "fields": {
            "ReceiptType": {
                "type": "string",
                "valueString": "Itemized",
                "confidence": 0.672
            },
            "MerchantName": {
                "type": "string",
                "valueString": "Tailwind",
                "text": "Tailwind",
                "boundingBox": [],
                "page": 1,
                "confidence": 0.913,
                "elements": [
                    "#/readResults/0/lines/0/words/0"
                ]
            }
        },
        ...
    }
]
```

Which expression should you use to trigger a manual review of the extracted information by a member of the Consultant-Bookkeeper group?

- A. documentResults.docType == "prebuilt:receipt"
- B. documentResults.fields.*.confidence < 0.7
- C. documentResults.fields.ReceiptType.confidence > 0.7
- D. documentResults.fields.MerchantName.confidence < 0.7

Correct Answer: D

Need to specify the field name, and then use < 0.7 to handle trigger if confidence score is less than 70%.

Scenario:

- ☞ AI solution responses must have a confidence score that is equal to or greater than 70 percent.
- ☞ When the response confidence score of an AI response is lower than 70 percent the response must be improved by human input.

Reference:

<https://docs.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/api-v2-0/reference-sdk-api-v2-0>

Community vote distribution

B (100%)

✉️  **AusAv**  1 year, 10 months ago

Selected Answer: B
Answer is B as I just did the exam and got 100% for the section :)
upvoted 32 times

✉️  **zellck**  10 months, 3 weeks ago

Gotten this in Jul 2023 exam.
upvoted 6 times

✉️  **takaimomoGcup**  1 week ago

Selected Answer: B
Is this question still available on May 21, 2024?
upvoted 1 times

✉️  **evangelist** 3 months, 2 weeks ago

Selected Answer: B
The question:"AI solution responses must have a confidence score that is equal to or greater than 70 percent.
When the response confidence score of an AI response is lower than 70 percent, the response must be improved by human input."
so it has to be fields.*.confidence >0.7
upvoted 2 times

✉️  **rdemontis** 6 months, 2 weeks ago

Selected Answer: B
I think the correct answer is B because that expression evaluates whether the confidence score of any field in the extracted information is less than 70 percent. Option D focus on a specific field (MerchantName) but do not consider the overall confidence score across all fields
upvoted 2 times

✉️  **crunkiNhere** 1 year ago

Selected Answer: B
I can't imagine a world where the answer could be anything but B
upvoted 3 times

✉️  **Eltooth** 1 year, 10 months ago

Selected Answer: B
B is correct answer.
upvoted 2 times

✉️  **AiEngineerS** 1 year, 10 months ago

Selected Answer: B
I think is B.. is that?
upvoted 2 times

✉️  **ANIKI51419** 1 year, 11 months ago

should be B
upvoted 3 times

Topic 14 - Testlet 7

Question #1

Topic 14

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

A company named Wide World Importers is developing an e-commerce platform.

You are working with a solutions architect to design and implement the features of the e-commerce platform. The platform will use microservices and a serverless environment built on Azure.

Wide World Importers has a customer base that includes English, Spanish, and Portuguese speakers.

Existing Environment -

Applications -

Wide World Importers has an App Service plan that contains the web apps shown in the following table.

Name	Description
Product Management	An app used by employees to create and manage products. The app and the expected inputs from the employees are in English.
Inventory Tracking	An app used by employees to manage inventory when dispatching orders, receiving refunds, and receiving consignments from suppliers.

Azure Resources -

You have the following resources:

An Azure Active Directory (Azure AD) tenant

- The tenant supports internal authentication.

- All employees belong to a group named AllUsers.

- Senior managers belong to a group named LeadershipTeam.

An Azure Functions resource

- A function app posts to Azure Event Grid when stock levels of a product change between OK, Low Stock, and Out of Stock. The function app uses the

Azure Cosmos DB change feed.

An Azure Cosmos DB account

- The account uses the Core (SQL) API.

- The account stores data for the Product Management app and the Inventory Tracking app.

An Azure Storage account

- The account contains blob containers for assets related to products.

- The assets include images, videos, and PDFs.

An Azure Cognitive Services resource named wwics

An Azure Video Analyzer for Media (previously Video Indexer) resource named wwivi

Requirements -

Business Goals -

Wide World Importers wants to leverage AI technologies to differentiate itself from its competitors.

Planned Changes -

Wide World Importers plans to start the following projects:

A product creation project: Help employees create accessible and multilingual product entries, while expediting product entry creation.

A smart e-commerce project: Implement an Azure Cognitive Search solution to display products for customers to browse.

A shopping on-the-go project: Build a chatbot that can be integrated into smart speakers to support customers.

Business Requirements -

Wide World Importers identifies the following business requirements for all the projects:

Provide a multilingual customer experience that supports English, Spanish, and Portuguese.

Whenever possible, scale based on transaction volumes to ensure consistent performance.

Minimize costs.

Governance and Security Requirements

Wide World Importers identifies the following governance and security requirements:

Data storage and processing must occur in datacenters located in the United States.

Azure Cognitive Services must be inaccessible directly from the internet.

Accessibility Requirements -

Wide World Importers identifies the following accessibility requirements:

All images must have relevant alt text.

All videos must have transcripts that are associated to the video and included in product descriptions.

Product descriptions, transcripts, and alt text must be available in English, Spanish, and Portuguese.

Product Creation Requirements -

Wide World Importers identifies the following requirements for improving the Product Management app:

Minimize how long it takes for employees to create products and add assets.

Remove the need for manual translations.

Smart E-Commerce Requirements -

Wide World Importers identifies the following requirements for the smart e-commerce project:

Ensure that the Cognitive Search solution meets a Service Level Agreement (SLA) of 99.9% availability for searches and index writes.

Provide users with the ability to search insight gained from the images, manuals, and videos associated with the products.

Support autocompletion and suggestion based on all product name variants.

Store all raw insight data that was generated, so the data can be processed later.

Update the stock level field in the product index immediately upon changes.

Update the product index hourly.

Shopping On-the-Go Requirements -

Wide World Importers identifies the following requirements for the shopping on-the-go chatbot:

Answer common questions.

Support interactions in English, Spanish, and Portuguese.

Replace an existing FAQ process so that all Q&A is managed from a central location.

Provide all employees with the ability to edit Q&As. Only senior managers must be able to publish updates.

Support purchases by providing information about relevant products to customers. Product displays must include images and warnings when stock levels are low or out of stock.

Product JSON Sample -

You have the following JSON sample for a product.

```
{  
    "sku": "b1",  
    "name": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "stocklevel": "Out of Stock",  
    "description": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "image": {  
        "uri": "https://upload.worldwideimporters.org/bicycle.jpg",  
        "alttext": {  
            "en": "Bicycle",  
            "es": "Bicicleta",  
            "pt": "Bicicleta"  
        }  
    },  
    "createdUtc": "2020-02-14T06:08:39Z",  
    "language": "en"  
}
```

Question

You are developing the smart e-commerce project.

You need to implement autocompletion as part of the Cognitive Search solution.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Make API queries to the autocomplete endpoint and include suggesterName in the body.
- B. Add a suggester that has the three product name fields as source fields.
- C. Make API queries to the search endpoint and include the product name fields in the searchFields query parameter.
- D. Add a suggester for each of the three product name fields.
- E. Set the searchAnalyzer property for the three product name variants.
- F. Set the analyzer property for the three product name variants.

Correct Answer: ABF

Scenario: Support autocompletion and autosuggestion based on all product name variants.

A: Call a suggester-enabled query, in the form of a Suggestion request or Autocomplete request, using an API. API usage is illustrated in the following call to the

Autocomplete REST API.

POST /indexes/myxboxgames/docs/autocomplete?search&api-version=2020-06-30

```
{  
    "search": "minecraf",  
    "suggesterName": "sg"  
}
```

B: In Azure Cognitive Search, typeahead or "search-as-you-type" is enabled through a suggester. A suggester provides a list of fields that undergo additional tokenization, generating prefix sequences to support matches on partial terms. For example, a suggester that includes a City field with a value for "Seattle" will have prefix combinations of "sea", "seat", "seatt", and "seattl" to support typeahead.

F. Use the default standard Lucene analyzer ("analyzer": null) or a language analyzer (for example, "analyzer": "en.Microsoft") on the field.

Reference:

<https://docs.microsoft.com/en-us/azure/search/index-add-suggesters>

Community vote distribution

ABF (100%)

✉  **ziggy1117**  11 months, 3 weeks ago

Selected Answer: ABF

Let us eliminate the wrong answers:

- C. Make API queries to the search endpoint and include the product name fields in the searchFields query parameter. - as we need autocomplete endpoint, we rule this one out as this is for a search endpoint
- D. Add a suggester for each of the three product name fields. - we cannot have 3 suggesters.
- E. Set the searchAnalyzer property for the three product name variants. - searchAnalyzer helps in search not in autocomplete.

So A, B, F

upvoted 10 times

✉  **reachmymind**  2 years, 2 months ago

A , B , F

B. Add a suggester that has the three product name fields as source fields

F. Set the analyzer property for the three product name variants.

A. Make API queries to the autocomplete endpoint and include suggesterName in the body.

<https://docs.microsoft.com/en-us/azure/search/index-add-suggesters>

<https://docs.microsoft.com/en-us/azure/search/search-add-autocomplete-suggestions>

upvoted 8 times

✉  **takaimomoGcup**  1 week ago

Is this question still available on May 21, 2024?

upvoted 1 times

✉  **evangelist** 3 months, 2 weeks ago

let me explain why C D E are incorrect:

C. This action is more relevant to refining search results based on specific fields rather than implementing autocomplete. For autocomplete, the focus is on the autocomplete endpoint, not the search endpoint.

D. - Typically, you would add a single suggester that includes multiple fields as source fields rather than creating a separate suggester for each field. Therefore, this action is not as efficient or necessary if you can accomplish your goal with one suggester that encompasses all relevant fields.

E. Set the searchAnalyzer property for the three product name variants. - Setting the searchAnalyzer property is relevant for controlling how text is analyzed during search queries. it's not directly involved in setting up autocomplete.

upvoted 1 times

✉  **rdemontis** 6 months, 2 weeks ago

Selected Answer: ABF

provided answer and explanation are correct

upvoted 2 times

✉  **zellck** 10 months, 4 weeks ago

Selected Answer: ABF

ABF is the answer.

<https://learn.microsoft.com/en-us/azure/search/index-add-suggesters#use-a-suggester>

First, as with all queries, the operation is against the documents collection of an index and the query includes a "search" parameter, which in this case provides the partial query. Second, you must add "suggesterName" to the request. If a suggester isn't defined in the index, a call to autocomplete or suggestions will fail.

<https://learn.microsoft.com/en-us/azure/search/index-add-suggesters#choose-analyzers>

upvoted 3 times

✉  **KingChuang** 1 year, 4 months ago

on my exam. (2023-01-16 passed)

My Answer:ABF

upvoted 5 times

✉  **Eltooth** 1 year, 10 months ago

Selected Answer: ABF

A, B and F are correct answers.

upvoted 3 times

✉  **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021

upvoted 2 times

Topic 15 - Testlet 8

Question #1

Topic 15

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

General Overview -

Contoso, Ltd. is an international accounting company that has offices in France, Portugal, and the United Kingdom.

Contoso has a professional services department that contains the roles shown in the following table.

Name	Position	Office
Accountant	Manager	United Kingdom, France, Portugal
Accountant	Consultant	United Kingdom, France, Portugal
Customer Service	Manager	United Kingdom
Customer Service	Agent	United Kingdom
Bookkeeper	Manager	United Kingdom, France, Portugal
Bookkeeper	Consultant	United Kingdom, France, Portugal

Existing environment -

Infrastructure -

Contoso has the following subscriptions:

Azure

Microsoft 365

Microsoft Dynamics 365

Azure Active (Azure AD) Directory

Contoso has Azure Active Directory groups for securing role-based access. The company uses the following group naming conventions:

[Country]-[Level]-[Role]

[Level]-[Role]

Intellectual Property -

Contoso has the intellectual property shown in the following table.

Content	Format	Language	Content store	Domain
Weekly webinars	Video	English	Azure Blob storage	Vid.contoso.com
Blogs	Text	English, French, Portuguese	WordPress	Pt-blog.contoso.com Blog.contoso.com Fr-blog.contoso.com
Wikis	Text	English, French, Portuguese	Azure Cosmos DB	Internal.contoso.com/wiki
Monthly conference recordings	Video	English	SharePoint Online	Contoso.sharepoint.com
Frequently asked questions (FAQs)	Text	English	SharePoint Online	Contoso.sharepoint.com

Text-based content is provided only in one language and is not translated.

Requirements -

Planned Projects -

Contoso plans to develop the following:

A document processing workflow to extract information automatically from PDFs and images of financial documents

A customer-support chatbot that will answer questions by using FAQs

A searchable knowledgebase of all the intellectual property

Technical Requirements -

Contoso identifies the following technical requirements:

All content must be approved before being published.

All planned projects must support English, French, and Portuguese.

All content must be secured by using role-based access control (RBAC).

RBAC role assignments must use the principle of least privilege.

RBAC roles must be assigned only to Azure Active Directory groups.

AI solution responses must have a confidence score that is equal to or greater than 70 percent.

When the response confidence score of an AI response is lower than 70 percent, the response must be improved by human input.

Chatbot Requirements -

Contoso identifies the following requirements for the chatbot:

Provide customers with answers to the FAQs.

Ensure that the customers can chat to a customer service agent.

Ensure that the members of a group named Management-Accountants can approve the FAQs.

Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.

Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.

Ensure that access to the customer service agents is managed by using Omnichannel for Customer Service.

When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

Document Processing Requirements

Contoso identifies the following requirements for document processing:

The document processing solution must be able to process standardized financial documents that have the following characteristics:

- Contain fewer than 20 pages.

- Be formatted as PDF or JPEG files.

- Have a distinct standard for each office.

The document processing solution must be able to extract tables and text from the financial documents.

The document processing solution must be able to extract information from receipt images.

Members of a group named Management-Bookkeeper must define how to extract tables from the financial documents.

Members of a group named Consultant-Bookkeeper must be able to process the financial documents.

Knowledgebase Requirements -

Contoso identifies the following requirements for the knowledgebase:

Supports searches for equivalent terms

Can transcribe jargon with high accuracy

▪

Can search content in different formats, including video

Provides relevant links to external resources for further research

Question

You are developing the document processing workflow.

You need to identify which API endpoints to use to extract text from the financial documents. The solution must meet the document processing requirements.

Which two API endpoints should you identify? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. /vision/v3.1/read/analyzeResults
- B. /formrecognizer/v2.0/custom/models/{modelId}/analyze
- C. /formrecognizer/v2.0/prebuilt/receipt/analyze
- D. /vision/v3.1/describe
- E. /vision/v3.1/read/analyze

Correct Answer: CE

C: Analyze Receipt - Get Analyze Receipt Result.

Query the status and retrieve the result of an Analyze Receipt operation.

Request URL:

<https://{{endpoint}}/formrecognizer/v2.0-preview/prebuilt/receipt/analyzeResults/{{resultId}}>

E: POST {{Endpoint}}/vision/v3.1/read/analyze

Use this interface to get the result of a Read operation, employing the state-of-the-art Optical Character Recognition (OCR) algorithms optimized for text-heavy documents.

Scenario: Contoso plans to develop a document processing workflow to extract information automatically from PDFs and images of financial documents

☞ The document processing solution must be able to process standardized financial documents that have the following characteristics:

- Contain fewer than 20 pages.

- Be formatted as PDF or JPEG files.

- Have a distinct standard for each office.

☞ *The document processing solution must be able to extract tables and text from the financial documents.

The document processing solution must be able to extract information from receipt images.

▪

Reference:

<https://westus2.dev.cognitive.microsoft.com/docs/services/form-recognizer-api-v2-preview/operations/GetAnalyzeReceiptResult>

<https://docs.microsoft.com/en-us/rest/api/computervision/3.1/read/read>

Community vote distribution

BC (41%) BE (27%) B (22%) 11%

 **rupert1o1N**  1 year, 10 months ago

guys so what is the correct answer

upvoted 11 times

 **Moody_L**  1 year, 11 months ago

Selected Answer: B

Contoso have a distinct standard for each office. Is the customized form recognizer more appropriate?

upvoted 8 times

 **sdokmak** 1 year, 11 months ago

Agreed, also the receipt text extraction is separate to the financial documents, question is only about the financial documents.

"The document processing solution must be able to extract tables and text from the financial documents."

The document processing solution must be able to extract information from receipt images."

upvoted 2 times

 **takaimomoGcup** Most Recent 1 week ago

Is this question still available on May 21, 2024?

upvoted 1 times

 **takaimomoGcup** 1 week ago

Is this question still available on May 21, 2024?

upvoted 1 times

 **azure_bimonster** 1 month, 1 week ago

Selected Answer: BC

B and C are correct as they are part of Azure Form Recognizer.

upvoted 1 times

 **Murtuza** 1 month, 3 weeks ago

Please note that these endpoints are part of Azure's Form Recognizer service, which is designed for extracting text, key/value pairs, and tables from documents. It's a great fit for your document processing requirements. The other endpoints listed (Option A, D, and E) are part of the Computer Vision service and are more suited for different tasks such as analyzing results of Read operation, describing an image, and running Read operation respectively. They might not be the best fit for your specific document processing needs.

upvoted 1 times

 **Murtuza** 2 months ago

Selected Answer: BC

To meet Contoso's document processing requirements for standardized financial documents, consider the following two API endpoints:

/formrecognizer/v2.0/custom/models/{modelId}/analyze: This endpoint is ideal for extracting information from financial documents. It allows you to analyze custom models specifically designed for invoices and receipts. With this API, you can extract both tables and text from documents that adhere to distinct standards for each office1.

/formrecognizer/v2.0/prebuilt/receipt/analyze: Designed for receipt analysis, this endpoint simplifies the extraction of relevant data from receipt images. It's well-suited for handling financial paperwork in PDF or JPEG formats, especially when dealing with documents containing fewer than 20 pages1.

By utilizing these endpoints, you can efficiently process financial documents, extract essential information, and enhance your document processing workflow.

upvoted 2 times

 **evangelist** 3 months, 2 weeks ago

Selected Answer: BC

B for customization C for prebuilt receipt with minimum efforts

upvoted 2 times

 **PCRamirez** 3 months, 2 weeks ago

According to Windows Copilot:

B) /formrecognizer/v2.0/custom/models/{modelId}/analyze: This endpoint is part of the Form Recognizer service. It allows you to create custom models for extracting structured data from documents. By training a custom model with examples of financial documents, you can extract relevant information such as tables and text. Since the solution needs to process standardized financial documents with distinct standards for each office, using a custom model tailored to your specific requirements is a suitable choice.

C) /formrecognizer/v2.0/prebuilt/receipt/analyze: This endpoint is also part of the Form Recognizer service. It is specifically designed for extracting information from receipt images. Given that the solution must handle receipt data, this prebuilt receipt analysis endpoint is a good fit.

upvoted 1 times

 **lastget** 4 months, 2 weeks ago

Selected Answer: BC

Accroding to the problem, we need:

The document processing solution must be able to extract tables and text from the financial documents.

The document processing solution must be able to extract information from receipt images.

So the answer I will choose BC

upvoted 2 times

 **123aash** 4 months, 3 weeks ago

Selected Answer: BC

B - Members of a group named Management-Bookkeeper must define how to extract tables from the financial documents

C - Extract text for "receipt" images

upvoted 3 times

 **rdemontis** 6 months, 2 weeks ago

Selected Answer: BE

1. Here the requirement involves only Financial documents, receipts are not included.

2. For financial documents we need to extract tables and text.

3. Extract Tables --> Form recognizer
4. Extract Text --> Computer Vision

So according to me the correct answer is B to extract the data from the tables and E to extract from the texts

<https://westus.dev.cognitive.microsoft.com/docs/services/form-recognizer-api-v2/operations/GetCustomModel>

<https://learn.microsoft.com/en-us/rest/api/computervision/read/read?view=rest-computervision-v3.1&tabs=HTTP>
upvoted 1 times

✉ **sl_mslconsulting** 7 months ago

Selected Answer: BC

You need to analyze the receipt images and retrieve tables from the documents. E won't work. You can simply check what the outputs look like.

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/how-to/call-read-api>

upvoted 1 times

✉ **sl_mslconsulting** 7 months ago

Check the requirements carefully and verify what each solution can do.

upvoted 1 times

✉ **EliteAllen** 10 months, 1 week ago

Selected Answer: BE

Based on ChatGPT:

B. /formrecognizer/v2.0/custom/models/{modelId}/analyze

This endpoint is part of Azure's Form Recognizer service. It's particularly useful for extracting text and tables from financial documents because it can use custom models trained on your specific document types, enabling it to handle the distinct standards for each office as mentioned in the requirements.

E. /vision/v3.1/read/analyze

The Azure Computer Vision API's Read operation, which is executed by making a POST request to the /read/analyze endpoint, can analyze text in images, PDF documents, and TIFF files, recognizing both printed and handwritten text. It is designed to handle large documents and extracts the text and structure (such as tables) in the document.

upvoted 2 times

✉ **M25** 8 months, 3 weeks ago

Correct! It is specifically for fin. doc., not receipts (excl. C)

<https://westus.dev.cognitive.microsoft.com/docs/services/form-recognizer-api-v2/operations/GetCustomModel>

<https://thedavidmasters.com/2020/05/18/demo-azure-form-recognizers-without-labels/>

upvoted 1 times

✉ **rdemontis** 6 months, 2 weeks ago

exactly

upvoted 1 times

✉ **rdemontis** 6 months, 2 weeks ago

Agree with you

upvoted 1 times

✉ **ziggyp1117** 11 months, 3 weeks ago

Selected Answer: BC

B: <https://learn.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/concept-analyze-document-response?view=form-recog-3.0.0>

C: Form Recognizer Receipts

upvoted 4 times

✉ **ziggyp1117** 11 months, 3 weeks ago

B: <https://learn.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/concept-analyze-document-response?view=form-recog-3.0.0>

C: Form Recognizer Receipts

upvoted 2 times

✉ **Pyguy** 1 year, 4 months ago

Selected Answer: BE

B is sure and E is because :

The /vision/v3.1/read/analyze endpoint can be used to extract information from tables in images. The endpoint uses optical character recognition (OCR) technology to extract text from images. The extracted text includes information from tables, and the location of the text in the image can be used to identify where the text came from in the table.

upvoted 2 times

Question #2

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

General Overview -

Contoso, Ltd. is an international accounting company that has offices in France, Portugal, and the United Kingdom.

Contoso has a professional services department that contains the roles shown in the following table.

Name	Position	Office
Accountant	Manager	United Kingdom, France, Portugal
Accountant	Consultant	United Kingdom, France, Portugal
Customer Service	Manager	United Kingdom
Customer Service	Agent	United Kingdom
Bookkeeper	Manager	United Kingdom, France, Portugal
Bookkeeper	Consultant	United Kingdom, France, Portugal

Existing environment -

Infrastructure -

Contoso has the following subscriptions:

Azure

Microsoft 365

Microsoft Dynamics 365

Azure Active (Azure AD) Directory

Contoso has Azure Active Directory groups for securing role-based access. The company uses the following group naming conventions:

[Country]-[Level]-[Role]

[Level]-[Role]

Intellectual Property -

Contoso has the intellectual property shown in the following table.

Content	Format	Language	Content store	Domain
Weekly webinars	Video	English	Azure Blob storage	Vid.contoso.com
Blogs	Text	English, French, Portuguese	WordPress	Pt-blog.contoso.com Blog.contoso.com Fr-blog.contoso.com
Wikis	Text	English, French, Portuguese	Azure Cosmos DB	Internal.contoso.com/wiki
Monthly conference recordings	Video	English	SharePoint Online	Contoso.sharepoint.com
Frequently asked questions (FAQs)	Text	English	SharePoint Online	Contoso.sharepoint.com

Text-based content is provided only in one language and is not translated.

Requirements -

Planned Projects -

Contoso plans to develop the following:

A document processing workflow to extract information automatically from PDFs and images of financial documents

A customer-support chatbot that will answer questions by using FAQs

A searchable knowledgebase of all the intellectual property

Technical Requirements -

Contoso identifies the following technical requirements:

All content must be approved before being published.

All planned projects must support English, French, and Portuguese.

All content must be secured by using role-based access control (RBAC).

RBAC role assignments must use the principle of least privilege.

RBAC roles must be assigned only to Azure Active Directory groups.

AI solution responses must have a confidence score that is equal to or greater than 70 percent.

When the response confidence score of an AI response is lower than 70 percent, the response must be improved by human input.

Chatbot Requirements -

Contoso identifies the following requirements for the chatbot:

Provide customers with answers to the FAQs.

Ensure that the customers can chat to a customer service agent.

Ensure that the members of a group named Management-Accountants can approve the FAQs.

Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.

Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.

Ensure that access to the customer service agents is managed by using Omnichannel for Customer Service.

When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

Document Processing Requirements

Contoso identifies the following requirements for document processing:

The document processing solution must be able to process standardized financial documents that have the following characteristics:

- Contain fewer than 20 pages.

- Be formatted as PDF or JPEG files.

- Have a distinct standard for each office.

The document processing solution must be able to extract tables and text from the financial documents.

The document processing solution must be able to extract information from receipt images.

Members of a group named Management-Bookkeeper must define how to extract tables from the financial documents.

Members of a group named Consultant-Bookkeeper must be able to process the financial documents.

Knowledgebase Requirements -

Contoso identifies the following requirements for the knowledgebase:

Supports searches for equivalent terms

Can transcribe jargon with high accuracy

Can search content in different formats, including video

Provides relevant links to external resources for further research

Question

HOTSPOT -

You are developing the knowledgebase by using Azure Cognitive Search.

You need to build a skill that will be used by indexers.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

{

```
"@odata.type": "#Microsoft.Skills.Text.EntityRecognitionSkill",
{
    "categories": [],
    "categories": [ "Email", "Persons", "Organizations"],
    "categories": [ "Locations", "Persons", "Organizations"],

    "defaultLanguageCode": "en",
    "includeTypelessEntities": true,
    "minimumPrecision": 0.7,
    "inputs": [
        { "name": "text",
            "source": "/document/content"}
    ],
    "outputs": [
        { "name": "persons", "targetName": "people"},
        { "name": "locations", "targetName": "locations"},
        { "name": "organizations", "targetName": "organizations"},

        { "name": "entities"}
        { "name": "categories"}
        { "name": "namedEntities"}
    ]
}
```

Correct Answer:

Answer Area

{

```

"@odata.type": "#Microsoft.Skills.Text.EntityRecognitionSkill",
  "categories": [],
  "categories": [ "Email", "Persons", "Organizations"],
  "categories": [ "Locations", "Persons", "Organizations"],

  "defaultLanguageCode": "en",
  "includeTypelessEntities": true,
  "minimumPrecision": 0.7,
  "inputs": [
    { "name": "text",
      "source": "/document/content"}
  ],
  "outputs": [
    { "name": "persons", "targetName": "people"},
    { "name": "locations", "targetName": "locations"},
    { "name": "organizations", "targetName": "organizations"},

    { "name": "entities"}
    { "name": "categories"}
    { "name": "namedEntities"}
  ]
}

```

Box 1: "categories": ["Locations", "Persons", "Organizations"],

Locations, Persons, Organizations are in the outputs.

Scenario: Contoso plans to develop a searchable knowledgebase of all the intellectual property

Note: The categories parameter is an array of categories that should be extracted. Possible category types: "Person", "Location", "Organization", "Quantity",

"Datetime", "URL", "Email". If no category is provided, all types are returned.

Box 2: {"name": "entities"}

The include wikis, so should include entities in the outputs.

Note: entities is an array of complex types that contains rich information about the entities extracted from text, with the following fields name (the actual entity name. This represents a "normalized" form) wikipediaId wikipediaLanguage wikipediaUrl (a link to Wikipedia page for the entity) etc.

Reference:

<https://docs.microsoft.com/en-us/azure/search/cognitive-search-skill-entity-recognition>

✉  **Eltooth**  1 year, 10 months ago

Answer is correct however... the Entity Recognition skill is now discontinued replaced by Microsoft.Skills.Text.V3.EntityRecognitionSkill.

Also categories will only allow "Person" as valid category type - not "Persons".

This (old) version allows output to be either "entities" or "namedEntities"

New version only allows "namedEntities" from answer list.

<https://docs.microsoft.com/en-us/azure/search/cognitive-search-skill-entity-recognition>

upvoted 10 times

✉  **rober13** 2 months, 1 week ago

So, the second option should be nameEntities, isn't it ??

upvoted 1 times

✉  **rdemontis** 6 months, 2 weeks ago

thanks for your contribution

upvoted 2 times

✉  **az999999**  10 months, 1 week ago

I also got this on July 3 2023 exam, but it's changed to Skills.Text.V3.

upvoted 5 times

✉ **takaimomoGcup** Most Recent 1 week ago

Is this question still available on May 21, 2024?

upvoted 1 times

✉ **takaimomoGcup** 1 week ago

Is this question still available on May 21, 2024?

upvoted 1 times

✉ **Murtuza** 2 months ago

The document processing solution aims to extract information related to persons, locations, and organizations. Additionally, it should recognize named entities and assign them to appropriate categories.

upvoted 1 times

✉ **Gvalli** 5 months, 1 week ago

A modified version of this was in the exam today.

upvoted 3 times

✉ **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 5 times

✉ **RamonKaus** 1 year, 10 months ago

Agreed

upvoted 2 times

Question #3

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

General Overview -

Contoso, Ltd. is an international accounting company that has offices in France, Portugal, and the United Kingdom.

Contoso has a professional services department that contains the roles shown in the following table.

Name	Position	Office
Accountant	Manager	United Kingdom, France, Portugal
Accountant	Consultant	United Kingdom, France, Portugal
Customer Service	Manager	United Kingdom
Customer Service	Agent	United Kingdom
Bookkeeper	Manager	United Kingdom, France, Portugal
Bookkeeper	Consultant	United Kingdom, France, Portugal

Existing environment -

Infrastructure -

Contoso has the following subscriptions:

Azure

Microsoft 365

Microsoft Dynamics 365

Azure Active (Azure AD) Directory

Contoso has Azure Active Directory groups for securing role-based access. The company uses the following group naming conventions:

[Country]-[Level]-[Role]

[Level]-[Role]

Intellectual Property -

Contoso has the intellectual property shown in the following table.

Content	Format	Language	Content store	Domain
Weekly webinars	Video	English	Azure Blob storage	Vid.contoso.com
Blogs	Text	English, French, Portuguese	WordPress	Pt-blog.contoso.com Blog.contoso.com Fr-blog.contoso.com
Wikis	Text	English, French, Portuguese	Azure Cosmos DB	Internal.contoso.com/wiki
Monthly conference recordings	Video	English	SharePoint Online	Contoso.sharepoint.com
Frequently asked questions (FAQs)	Text	English	SharePoint Online	Contoso.sharepoint.com

Text-based content is provided only in one language and is not translated.

Requirements -

Planned Projects -

Contoso plans to develop the following:

A document processing workflow to extract information automatically from PDFs and images of financial documents

A customer-support chatbot that will answer questions by using FAQs

A searchable knowledgebase of all the intellectual property

Technical Requirements -

Contoso identifies the following technical requirements:

All content must be approved before being published.

All planned projects must support English, French, and Portuguese.

All content must be secured by using role-based access control (RBAC).

RBAC role assignments must use the principle of least privilege.

RBAC roles must be assigned only to Azure Active Directory groups.

AI solution responses must have a confidence score that is equal to or greater than 70 percent.

When the response confidence score of an AI response is lower than 70 percent, the response must be improved by human input.

Chatbot Requirements -

Contoso identifies the following requirements for the chatbot:

Provide customers with answers to the FAQs.

Ensure that the customers can chat to a customer service agent.

Ensure that the members of a group named Management-Accountants can approve the FAQs.

Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.

Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.

Ensure that access to the customer service agents is managed by using Omnichannel for Customer Service.

When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

Document Processing Requirements

Contoso identifies the following requirements for document processing:

The document processing solution must be able to process standardized financial documents that have the following characteristics:

- Contain fewer than 20 pages.

- Be formatted as PDF or JPEG files.

- Have a distinct standard for each office.

The document processing solution must be able to extract tables and text from the financial documents.

The document processing solution must be able to extract information from receipt images.

Members of a group named Management-Bookkeeper must define how to extract tables from the financial documents.

Members of a group named Consultant-Bookkeeper must be able to process the financial documents.

Knowledgebase Requirements -

Contoso identifies the following requirements for the knowledgebase:

Supports searches for equivalent terms

Can transcribe jargon with high accuracy

-

Can search content in different formats, including video

Provides relevant links to external resources for further research

Question

You are developing the knowledgebase by using Azure Cognitive Search.

You need to process wiki content to meet the technical requirements.

What should you include in the solution?

- A. an indexer for Azure Blob storage attached to a skillset that contains the language detection skill and the text translation skill
- B. an indexer for Azure Blob storage attached to a skillset that contains the language detection skill
- C. an indexer for Azure Cosmos DB attached to a skillset that contains the document extraction skill and the text translation skill
- D. an indexer for Azure Cosmos DB attached to a skillset that contains the language detection skill and the text translation skill

Correct Answer: C

The wiki contains text in English, French and Portuguese.

Scenario: All planned projects must support English, French, and Portuguese.

The Document Extraction skill extracts content from a file within the enrichment pipeline. This allows you to take advantage of the document extraction step that normally happens before the skillset execution with files that may be generated by other skills.

Note: The Translator Text API will be used to determine the from language. The Language detection skill is not required.

Incorrect Answers:

Not A, not B: The wiki is stored in Azure Cosmos DB.

Reference:

<https://docs.microsoft.com/en-us/azure/search/cognitive-search-skill-document-extraction> <https://docs.microsoft.com/en-us/azure/search/cognitive-search-skill-text-translation>

Community vote distribution

D (73%)

C (27%)

✉ g2000 Highly Voted 2 years ago

Selected Answer: D

Document extraction skill only works for files. It's less likely to work with CosmosDB.

upvoted 8 times

✉ sdukmak 1 year, 11 months ago

so A?..

upvoted 2 times

✉ sdukmak 1 year, 11 months ago

nvm i get it now

upvoted 1 times

✉ AzureJobsTillRetire 1 year, 2 months ago

This is not correct.

Azure Cosmos DB is a globally distributed multi-model database with support for multiple APIs.

The Get Document operation retrieves a document by its partition key and document key.

Method Request URI Description

GET <https://{{databaseaccount}}.documents.azure.com/dbs/{{db-id}}/colls/{{coll-id}}/docs/{{doc-id}}> Note that the {{databaseaccount}} is the name of the Azure Cosmos DB account created under your subscription. The {{db-id}} value is the user generated name/ID of the database, not the system generated ID (rid). The {{coll-id}} value is the name of the collection. The {{doc-id}} value is the ID of the document to be retrieved.

<https://learn.microsoft.com/en-us/rest/api/cosmos-db/get-a-document>

upvoted 2 times

✉ not_a_robot Highly Voted 1 year, 9 months ago

Selected Answer: C

The answer is correct.

It's to extract the content from the inputs. It supports a lot of formats of documents, including JSON.

<https://docs.microsoft.com/en-us/azure/search/cognitive-search-skill-document-extraction>

upvoted 6 times

✉ **takaimomoGcup** **Most Recent** 1 week ago

Selected Answer: D

Is this question still available on May 21, 2024?

upvoted 1 times

✉ **chandiochan** 1 month, 1 week ago

Selected Answer: D

must be D

upvoted 1 times

✉ **Murtuza** 2 months ago

Selected Answer: D

Therefore, A. an indexer for Cosmos DB attached to a skillset that contains the language detection skill and the text translation skill aligns with the technical requirements for content approval, multilingual support, and RBAC1

upvoted 2 times

✉ **chandiochan** 2 months, 1 week ago

Selected Answer: D

Must be D

upvoted 4 times

✉ **evangelist** 3 months, 1 week ago

Selected Answer: D

The correct answer is D. Use an indexer connected to Azure Cosmos DB that includes language detection and text translation skills. This solution meets the requirement for handling multiple languages and supports content transformation for English, French, and Portuguese by using the text translation skill, which aligns with the technical requirement for multilingual support. Additionally, using Azure Cosmos DB and associated skill set ensures content security and role-based access control (RBAC), meeting the requirements for content approval and security.

upvoted 6 times

✉ **PCRamirez** 3 months, 2 weeks ago

Windows Copilot:

B. An indexer for Azure Blob storage attached to a skillset that contains the language detection skill.

Here's the explanation:

Azure Blob Storage: You'll use Azure Blob storage to store your content, including wiki articles and other relevant documents.

Indexer: An indexer in Azure Cognitive Search automatically discovers new content from supported data sources (such as blob storage) and adds it to the search index. In your case, you'll create an indexer to process the wiki content.

Skillset: A skillset defines a set of cognitive skills (such as language detection, text translation, and more) that can be applied during indexing. Since your knowledgebase needs to support multiple languages (English, French, and Portuguese), you'll attach a skillset to the indexer.

...

upvoted 1 times

✉ **PCRamirez** 3 months, 2 weeks ago

Language Detection Skill: This skill detects the language of the content. It ensures that the appropriate language-specific processing (such as translation) can be applied.

By using the language detection skill, you'll be able to handle content in different languages and transcribe jargon accurately. Additionally, the Azure Cognitive Search service supports full-text search, semantic search, vector search, and hybrid search, which aligns with your requirement for searching content in various formats, including video.

Remember that the semantic search and vector search features are currently in public preview in Azure Cognitive Search. The solution will also provide relevant links to external resources for further research, enhancing the overall knowledgebase experience.   

For more details, you can refer to the official documentation.

upvoted 1 times

✉ **PCRamirez** 3 months, 2 weeks ago

- (1) Search and query an enterprise knowledge base by using Azure OpenAI or <https://learn.microsoft.com/en-us/azure/architecture/ai-ml/openai/idea/search-and-query-using-openai-service>.
- (2) Create a knowledge store with Azure AI Search - Training. <https://learn.microsoft.com/en-us/training/modules/create-knowledge-store-azure-cognitive-search/>.
- (3) Implement knowledge mining with Azure Cognitive Search. <https://learn.microsoft.com/en-us/training/parts/implement-knowledge-mining-azure-cognitive-search/>.
- (4) Knowledge store concepts - Azure AI Search | Microsoft Learn. <https://learn.microsoft.com/en-us/azure/search/knowledge-store-concept-intro>.

(5) Index file content and metadata by using Azure Cognitive Search. <https://learn.microsoft.com/en-us/azure/architecture/ai-ml/architecture/search-blob-metadata>.

upvoted 1 times

✉️ **lastget** 4 months, 2 weeks ago

Selected Answer: D
You only need to detect language and translate from wiki
upvoted 2 times

✉️ **dimsok** 4 months, 2 weeks ago

It is confusing because the table mentions wikis are in Text format, but on the other hand, wikis (and blog entries) usually contain HTML.
If they do contain HTML it should be C (extract text from html - don't translate tags)
If they don't it should be D

At the end I choose D because we shouldn't assume stuff, the question says "Text" not "Text/HTML"

upvoted 1 times

✉️ **rdemontis** 6 months, 2 weeks ago

Selected Answer: D
IMHO the correct answer should be D.
You don't need Document Extraction to extract text from files. You have already the text in wikis.

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-skill-document-extraction>

It's true that Text Translation could also detect the source language but reading the documentation seems it could be not so accurate.

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-skill-text-translation>

For sure you need two operations here: detect and translate.

upvoted 3 times

✉️ **sl_mslconsulting** 7 months ago

Selected Answer: D
Wikis are already in text format and available in three languages so what do you need the document extraction and translation normalize the text to a single language before indexing for? What you need is to detect the language of the text and translate it to English if you want to normalize the text to a single language before indexing for search. Only in this case it would make sense to use text translation skill.
upvoted 3 times

✉️ **sl_mslconsulting** 7 months ago

Correction: Wikis are already in text format and available in three languages so what do you need the document extraction and translation for?
What you need is to detect the language of the text and translate it to English if you want to normalize the text to a single language before indexing for search. Only in this case it would make sense to use text translation skill.

upvoted 1 times

✉️ **idcanymore** 1 month, 4 weeks ago

Where does it say wikis are available in three languages? All it says is 'All planned projects must support English, French, and Portuguese.'.
upvoted 1 times

✉️ **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 3 times

✉️ **ziggy1117** 11 months, 3 weeks ago

Selected Answer: C
Language Detection is not needed. The capability is especially useful when you need to provide the language of the text as input to other skills (for example, the Sentiment Analysis skill or Text Split skill). But it is not required for Text Translation
upvoted 3 times

✉️ **M25** 8 months, 3 weeks ago

Correct! Even if:

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-concept-intro>

A skillset that's assembled using built-in skills is well suited for the following application scenarios:

- Machine translation is provided by the Text Translation skill, often paired with language detection for multi-language solutions.

upvoted 1 times

✉️ **M25** 8 months, 3 weeks ago

Data Source vs. Destination (see also Topic 6 Q#1)

<https://learn.microsoft.com/en-us/azure/search/search-indexer-overview#document-cracking>

Depending on the data source, the indexer will try different operations to extract potentially indexable content:

- When the document is a record in Azure Cosmos DB, the indexer will extract non-binary content from fields and subfields from the Azure Cosmos DB document.

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-concept-intro>

Exploration is the last step. Output is always a search index that you can query from a client app. Output can optionally be a knowledge store consisting of blobs and tables in Azure Storage that are accessed through data exploration tools or downstream processes.

upvoted 1 times

✉️ **M25** 8 months, 3 weeks ago

Indexer

<https://learn.microsoft.com/en-us/azure/search/search-indexer-overview>

An indexer in Azure Cognitive Search is a crawler that extracts searchable content from cloud data sources and populates a search index using field-to-field mappings between source data and a search index. This approach is sometimes referred to as a 'pull model' because the search service pulls data in without you having to write any code that adds data to an index.

<https://learn.microsoft.com/en-us/azure/search/tutorial-multiple-data-sources#create-azure-cosmos-db-data-source-and-indexer>

After the data source is created, the program sets up an Azure Cosmos DB indexer named hotel-rooms-cosmos-indexer.

upvoted 1 times

✉️ **M25** 8 months, 3 weeks ago

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-skill-document-extraction>

The Document Extraction skill extracts content from a file within the enrichment pipeline. This allows you to take advantage of the document extraction step that normally happens before the skillset execution with files that may be generated by other skills.

This skill extracts text and images. Text extraction is free. Image extraction is metered by Azure Cognitive Search.

upvoted 1 times

✉️ **M25** 8 months, 3 weeks ago

But:

<https://learn.microsoft.com/en-us/azure/search/search-language-support>

A multilingual search application is one that provides a search experience in the user's own language. Language support is enabled through a language analyzer assigned to string field.

This article assumes you have translated strings in place. If that's not the case, you can attach Azure AI services to an enrichment pipeline, invoking text translation during data ingestion. Text translation takes a dependency on the indexer feature and Azure AI services, but all setup is done within Azure Cognitive Search.

upvoted 1 times

✉️ **ap1234pa** 1 year, 4 months ago

Selected Answer: C

C is correct

upvoted 1 times

✉️ **am20** 1 year, 4 months ago

Selected Answer: C

I'm assuming Cosmos DB will be used for output not as an input

upvoted 1 times

✉️ **am20** 1 year, 4 months ago

nevermind

upvoted 2 times

Question #4

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

General Overview -

Contoso, Ltd. is an international accounting company that has offices in France, Portugal, and the United Kingdom.

Contoso has a professional services department that contains the roles shown in the following table.

Name	Position	Office
Accountant	Manager	United Kingdom, France, Portugal
Accountant	Consultant	United Kingdom, France, Portugal
Customer Service	Manager	United Kingdom
Customer Service	Agent	United Kingdom
Bookkeeper	Manager	United Kingdom, France, Portugal
Bookkeeper	Consultant	United Kingdom, France, Portugal

Existing environment -

Infrastructure -

Contoso has the following subscriptions:

Azure

Microsoft 365

Microsoft Dynamics 365

Azure Active (Azure AD) Directory

Contoso has Azure Active Directory groups for securing role-based access. The company uses the following group naming conventions:

[Country]-[Level]-[Role]

[Level]-[Role]

Intellectual Property -

Contoso has the intellectual property shown in the following table.

Content	Format	Language	Content store	Domain
Weekly webinars	Video	English	Azure Blob storage	Vid.contoso.com
Blogs	Text	English, French, Portuguese	WordPress	Pt-blog.contoso.com Blog.contoso.com Fr-blog.contoso.com
Wikis	Text	English, French, Portuguese	Azure Cosmos DB	Internal.contoso.com/wiki
Monthly conference recordings	Video	English	SharePoint Online	Contoso.sharepoint.com
Frequently asked questions (FAQs)	Text	English	SharePoint Online	Contoso.sharepoint.com

Text-based content is provided only in one language and is not translated.

Requirements -

Planned Projects -

Contoso plans to develop the following:

A document processing workflow to extract information automatically from PDFs and images of financial documents

A customer-support chatbot that will answer questions by using FAQs

A searchable knowledgebase of all the intellectual property

Technical Requirements -

Contoso identifies the following technical requirements:

All content must be approved before being published.

All planned projects must support English, French, and Portuguese.

All content must be secured by using role-based access control (RBAC).

RBAC role assignments must use the principle of least privilege.

RBAC roles must be assigned only to Azure Active Directory groups.

AI solution responses must have a confidence score that is equal to or greater than 70 percent.

When the response confidence score of an AI response is lower than 70 percent, the response must be improved by human input.

Chatbot Requirements -

Contoso identifies the following requirements for the chatbot:

Provide customers with answers to the FAQs.

Ensure that the customers can chat to a customer service agent.

Ensure that the members of a group named Management-Accountants can approve the FAQs.

Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.

Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.

Ensure that access to the customer service agents is managed by using Omnichannel for Customer Service.

When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

Document Processing Requirements

Contoso identifies the following requirements for document processing:

The document processing solution must be able to process standardized financial documents that have the following characteristics:

- Contain fewer than 20 pages.

- Be formatted as PDF or JPEG files.

- Have a distinct standard for each office.

The document processing solution must be able to extract tables and text from the financial documents.

The document processing solution must be able to extract information from receipt images.

Members of a group named Management-Bookkeeper must define how to extract tables from the financial documents.

Members of a group named Consultant-Bookkeeper must be able to process the financial documents.

Knowledgebase Requirements -

Contoso identifies the following requirements for the knowledgebase:

Supports searches for equivalent terms

Can transcribe jargon with high accuracy

▪

Can search content in different formats, including video

Provides relevant links to external resources for further research

Question

You are developing the knowledgebase by using Azure Cognitive Search.

You need to meet the knowledgebase requirements for searching equivalent terms.

What should you include in the solution?

- A. synonym map
- B. a suggester
- C. a custom analyzer
- D. a built-in key phrase extraction skill

Correct Answer: A

Within a search service, synonym maps are a global resource that associate equivalent terms, expanding the scope of a query without the user having to actually provide the term. For example, assuming "dog", "canine", and "puppy" are mapped synonyms, a query on "canine" will match on a document containing "dog".

Create synonyms: A synonym map is an asset that can be created once and used by many indexes.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-synonyms>

Community vote distribution

A (100%)

 **takaimomoGcup** 1 week ago

Is this question still available on May 21, 2024?

upvoted 1 times

 **evangelist** 3 months, 2 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

 **rdemontis** 6 months, 2 weeks ago

Selected Answer: A

correct

<https://learn.microsoft.com/en-us/azure/search/search-synonyms>

upvoted 1 times

 **james2033** 9 months, 2 weeks ago

Selected Answer: A

Quote "In Azure Cognitive Search, a synonym map contains a list of rules for expanding or rewriting a search query to equivalent terms." at <https://learn.microsoft.com/en-us/rest/api/searchservice/create-synonym-map>.

upvoted 4 times

 **zellck** 11 months ago

Selected Answer: A

A is the answer.

<https://learn.microsoft.com/en-us/azure/search/search-synonyms>

Within a search service, synonym maps are a global resource that associate equivalent terms, expanding the scope of a query without the user having to actually provide the term. For example, assuming "dog", "canine", and "puppy" are mapped synonyms, a query on "canine" will match on a document containing "dog".

upvoted 2 times

 **zellck** 10 months, 3 weeks ago

Gotten this in Jul 2023 exam.

upvoted 3 times

 **Eltooth** 1 year, 10 months ago

Selected Answer: A

A is correct answer.

Synonym maps are a global resource that associate equivalent terms, expanding the scope of a query without the user having to actually provide the term.

You might create multiple synonym maps for different languages, such as English and French versions, or lexicons if your content includes technical or obscure terminology.

<https://docs.microsoft.com/en-us/azure/search/search-synonyms>

upvoted 2 times

Topic 16 - Testlet 9

Question #1

Topic 16

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

A company named Wide World Importers is developing an e-commerce platform.

You are working with a solutions architect to design and implement the features of the e-commerce platform. The platform will use microservices and a serverless environment built on Azure.

Wide World Importers has a customer base that includes English, Spanish, and Portuguese speakers.

Existing Environment -

Applications -

Wide World Importers has an App Service plan that contains the web apps shown in the following table.

Name	Description
Product Management	An app used by employees to create and manage products. The app and the expected inputs from the employees are in English.
Inventory Tracking	An app used by employees to manage inventory when dispatching orders, receiving refunds, and receiving consignments from suppliers.

Azure Resources -

You have the following resources:

An Azure Active Directory (Azure AD) tenant

- The tenant supports internal authentication.

- All employees belong to a group named AllUsers.

- Senior managers belong to a group named LeadershipTeam.

An Azure Functions resource

- A function app posts to Azure Event Grid when stock levels of a product change between OK, Low Stock, and Out of Stock. The function app uses the -

Azure Cosmos DB change feed.

An Azure Cosmos DB account

- The account uses the Core (SQL) API.

- The account stores data for the Product Management app and the Inventory Tracking app.

An Azure Storage account

- The account contains blob containers for assets related to products.

- The assets include images, videos, and PDFs.

An Azure Cognitive Services resource named wwics

An Azure Video Analyzer for Media (previously Video Indexer) resource named wwivi

Requirements -

Business Goals -

Wide World Importers wants to leverage AI technologies to differentiate itself from its competitors.

Planned Changes -

Wide World Importers plans to start the following projects:

A product creation project: Help employees create accessible and multilingual product entries, while expediting product entry creation.

A smart e-commerce project: Implement an Azure Cognitive Search solution to display products for customers to browse.

A shopping on-the-go project: Build a chatbot that can be integrated into smart speakers to support customers.

Business Requirements -

Wide World Importers identifies the following business requirements for all the projects:

Provide a multilingual customer experience that supports English, Spanish, and Portuguese.

Whenever possible, scale based on transaction volumes to ensure consistent performance.

Minimize costs.

Governance and Security Requirements

Wide World Importers identifies the following governance and security requirements:

Data storage and processing must occur in datacenters located in the United States.

Azure Cognitive Services must be inaccessible directly from the internet.

Accessibility Requirements -

Wide World Importers identifies the following accessibility requirements:

All images must have relevant alt text.

All videos must have transcripts that are associated to the video and included in product descriptions.

Product descriptions, transcripts, and alt text must be available in English, Spanish, and Portuguese.

Product Creation Requirements -

Wide World Importers identifies the following requirements for improving the Product Management app:

Minimize how long it takes for employees to create products and add assets.

Remove the need for manual translations.

Smart E-Commerce Requirements -

Wide World Importers identifies the following requirements for the smart e-commerce project:

Ensure that the Cognitive Search solution meets a Service Level Agreement (SLA) of 99.9% availability for searches and index writes.

Provide users with the ability to search insight gained from the images, manuals, and videos associated with the products.

Support autocompletion and suggestion based on all product name variants.

Store all raw insight data that was generated, so the data can be processed later.

Update the stock level field in the product index immediately upon changes.

Update the product index hourly.

Shopping On-the-Go Requirements -

Wide World Importers identifies the following requirements for the shopping on-the-go chatbot:

Answer common questions.

Support interactions in English, Spanish, and Portuguese.

Replace an existing FAQ process so that all Q&A is managed from a central location.

Provide all employees with the ability to edit Q&As. Only senior managers must be able to publish updates.

Support purchases by providing information about relevant products to customers. Product displays must include images and warnings when stock levels are low or out of stock.

Product JSON Sample -

You have the following JSON sample for a product.

```
{  
    "sku": "b1",  
    "name": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "stocklevel": "Out of Stock",  
    "description": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "image":  
    {"uri": "https://upload.worldwideimporters.org/bicycle.jpg",  
        "alttext": {  
            "en": "Bicycle",  
            "es": "Bicicleta",  
            "pt": "Bicicleta"  
        }  
    },  
    "createdUtc": "2020-02-14T06:08:39Z",  
    "language": "en"  
}
```

Question

HOTSPOT -

You are developing the shopping on-the-go project.

You need to build the Adaptive Card for the chatbot.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

```
{  
  "$schema": "http://adaptivecards.io.schemas/adaptive-card.json",  
  "type": "AdaptiveCard",  
  "version": "1.3",  
  "body": [  
    {  
      "type": "TextBlock",  
      "size": "Medium",  
      "weight": "Bolder",  
      "text": "${  
        if(language == 'en', 'en', name)  
        name  
        name.en  
        name[language]  
      }"  
    },  
    {  
      "type": "TextBlock",  
      "$when": "${stockLevel != 'OK' }"  
      "$when": "${stockLevel == 'OK' }"  
      "$when": "${stockLevel.OK}"  
      "text": "${stockLevel},  
      "color": "Attention"  
    },  
    {  
      "type": "Image",  
      "url": "${image.uri}",  
      "size": "Medium",  
      "altText": "${  
        image.altText.en  
        image.altText.language  
        image.altText.[\"language\"]  
        image.altText.[language]  
      }"  
    }  
  ]  
}
```

```

{
  "$schema": "http://adaptivecards.io.schemas/adaptive-card.json",
  "type": "AdaptiveCard",
  "version": "1.3",
  "body": [
    {
      "type": "TextBlock",
      "size": "Medium",
      "weight": "Bolder",
      "text": "${"
    }
  ],
  {
    "type": "TextBlock",
    "$when": "${stockLevel != 'OK'}"
    "$when": "${stockLevel == 'OK'}"
    "$when": "${stockLevel.OK}"
    "text": "${stockLevel}",
    "color": "Attention"
  },
  {
    "type": "Image",
    "url": "${image.uri}",
    "size": "Medium",
    "altText": "${"
  }
}
]
}

```

Box 1: name [language]

if(language == 'en', 'en', name)
name
name.en
name[language]

Box 2: "\$when:\${stockLevel != 'OK'}"

"\$when": "\${stockLevel != 'OK'}"
"\$when": "\${stockLevel == 'OK'}"
"\$when": "\${stockLevel.OK}"

Box 3: image.altText[language]

image.altText.en
image.altText.language
image.altText.[language]
image.altText.[language]

Correct Answer:

Box 1: name [language]

Chatbot must support interactions in English, Spanish, and Portuguese.

Box 2: "\$when:\${stockLevel != 'OK'}"

Product displays must include images and warnings when stock levels are low or out of stock.

Box 3: image.altText[language]

✉ **KingChuang** Highly Voted 1 year, 4 months ago

on my exam. (2023-01-16 passed)

My Answer:

1. name[language]
2. != OK
3. image.altText.[language]

upvoted 19 times

✉ **takaimomoGcup** Most Recent 1 week ago

Is this question still available on May 21, 2024?

upvoted 1 times

✉ **rdemontis** 6 months, 2 weeks ago

answer seems to be correct

upvoted 3 times

✉ **rdemontis** 6 months, 2 weeks ago

although there is probably a typo on the last answer image.altText.[language]. Probably the correct one should be image.altText[language]

upvoted 3 times

✉ **AzureJobsTillRetire** 1 year, 2 months ago

The correct answers are those that are given in answers to question 3.

The correct answer pairs for box1 and box3 could be-

- 1) name.en and image.altText.en
- 2) name["language"] and image.altText["language"]

Since name["language"] does not exist as an option in box1, we will have to choose the first pair as the correct answer.

https://www.w3schools.com/js/js_json_syntax.asp

upvoted 1 times

AzureJobsTillRetire 1 year, 2 months ago

Looking at the given JSON sample for a product again, I actually cannot find either the data for name["language"]/name.language or image.altText["language"]/image.altText.language.

upvoted 1 times

AzureJobsTillRetire 1 year, 2 months ago

Since name["language"]/name.language and image.altText["language"]/image.altText.language do not exist, we will have to choose the first pa as the answers

upvoted 1 times

halfway 1 year, 5 months ago

My last comment was wrong. There is no correct answer for the last question. The correct answers should be:

1. name[language] 2. "\$when": "\${stockLevel} != 'OK'" 3. image.altText[language]

upvoted 2 times

halfway 1 year, 5 months ago

1. name[language] 2. "\$when": "\${stockLevel} != 'OK'" 3. image.altText.language

upvoted 1 times

AdarshKumarKhare 1 year, 5 months ago

Repeat Question. Answers are different. I don't know which one is correct

upvoted 3 times

Question #2

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

A company named Wide World Importers is developing an e-commerce platform.

You are working with a solutions architect to design and implement the features of the e-commerce platform. The platform will use microservices and a serverless environment built on Azure.

Wide World Importers has a customer base that includes English, Spanish, and Portuguese speakers.

Existing Environment -

Applications -

Wide World Importers has an App Service plan that contains the web apps shown in the following table.

Name	Description
Product Management	An app used by employees to create and manage products. The app and the expected inputs from the employees are in English.
Inventory Tracking	An app used by employees to manage inventory when dispatching orders, receiving refunds, and receiving consignments from suppliers.

Azure Resources -

You have the following resources:

An Azure Active Directory (Azure AD) tenant

- The tenant supports internal authentication.

- All employees belong to a group named AllUsers.

- Senior managers belong to a group named LeadershipTeam.

An Azure Functions resource

- A function app posts to Azure Event Grid when stock levels of a product change between OK, Low Stock, and Out of Stock. The function app uses the -

Azure Cosmos DB change feed.

An Azure Cosmos DB account

- The account uses the Core (SQL) API.

- The account stores data for the Product Management app and the Inventory Tracking app.

An Azure Storage account

- The account contains blob containers for assets related to products.

- The assets include images, videos, and PDFs.

An Azure Cognitive Services resource named wwics

An Azure Video Analyzer for Media (previously Video Indexer) resource named wwivi

Requirements -

Business Goals -

Wide World Importers wants to leverage AI technologies to differentiate itself from its competitors.

Planned Changes -

Wide World Importers plans to start the following projects:

A product creation project: Help employees create accessible and multilingual product entries, while expediting product entry creation.

A smart e-commerce project: Implement an Azure Cognitive Search solution to display products for customers to browse.

A shopping on-the-go project: Build a chatbot that can be integrated into smart speakers to support customers.

Business Requirements -

Wide World Importers identifies the following business requirements for all the projects:

Provide a multilingual customer experience that supports English, Spanish, and Portuguese.

Whenever possible, scale based on transaction volumes to ensure consistent performance.

Minimize costs.

Governance and Security Requirements

Wide World Importers identifies the following governance and security requirements:

Data storage and processing must occur in datacenters located in the United States.

Azure Cognitive Services must be inaccessible directly from the internet.

Accessibility Requirements -

Wide World Importers identifies the following accessibility requirements:

All images must have relevant alt text.

All videos must have transcripts that are associated to the video and included in product descriptions.

Product descriptions, transcripts, and alt text must be available in English, Spanish, and Portuguese.

Product Creation Requirements -

Wide World Importers identifies the following requirements for improving the Product Management app:

Minimize how long it takes for employees to create products and add assets.

Remove the need for manual translations.

Smart E-Commerce Requirements -

Wide World Importers identifies the following requirements for the smart e-commerce project:

Ensure that the Cognitive Search solution meets a Service Level Agreement (SLA) of 99.9% availability for searches and index writes.

Provide users with the ability to search insight gained from the images, manuals, and videos associated with the products.

Support autocompletion and suggestion based on all product name variants.

Store all raw insight data that was generated, so the data can be processed later.

Update the stock level field in the product index immediately upon changes.

Update the product index hourly.

Shopping On-the-Go Requirements -

Wide World Importers identifies the following requirements for the shopping on-the-go chatbot:

Answer common questions.

Support interactions in English, Spanish, and Portuguese.

Replace an existing FAQ process so that all Q&A is managed from a central location.

Provide all employees with the ability to edit Q&As. Only senior managers must be able to publish updates.

Support purchases by providing information about relevant products to customers. Product displays must include images and warnings when stock levels are low or out of stock.

Product JSON Sample -

You have the following JSON sample for a product.

```
{  
    "sku": "b1",  
    "name": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "stocklevel": "Out of Stock",  
    "description": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "image":  
    {"uri": "https://upload.worldwideimporters.org/bicycle.jpg",  
        "alttext": {  
            "en": "Bicycle",  
            "es": "Bicicleta",  
            "pt": "Bicicleta"  
        }  
    },  
    "createdUtc": "2020-02-14T06:08:39Z",  
    "language": "en"  
}
```

Question

HOTSPOT -

You are developing the shopping on-the-go project.

You are configuring access to the QnA Maker (classic) resources.

Which role should you assign to AllUsers and LeadershipTeam? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

AllUsers:

Cognitive Service User
Contributor
Owner
QnA Maker Editor
QnA Maker Read

LeadershipTeam:

Cognitive Service User
Contributor
Owner
QnA Maker Editor
QnA Maker Read

Answer Area

AllUsers:

Cognitive Service User
Contributor
Owner
QnA Maker Editor
QnA Maker Read

Correct Answer:

LeadershipTeam:

Cognitive Service User
Contributor
Owner
QnA Maker Editor
QnA Maker Read

Box 1: QnA Maker Editor -

Scenario: Provide all employees with the ability to edit Q&As.

The QnA Maker Editor (read/write) has the following permissions:

- Create KB API
- Update KB API
- Replace KB API
- Replace Alterations
- "Train API" [in new service model v5]

Box 2: Contributor -

Scenario: Only senior managers must be able to publish updates.

Contributor permission: All except ability to add new members to roles

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/reference-role-based-access-control>✉  **g2000**  2 years ago

How about Cognitive Service User? the link explicitly says publish is given to this role.

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/concepts/role-based-access-control#access-is-provided-by-a-defined-role>
upvoted 17 times✉  **AusAv** 1 year, 10 months ago

I agree it is: Cognitive Service Userm Contributor can also create new resources

upvoted 1 times

✉  **firewind** 1 year, 8 months ago

Agreed, from a least privilege principle perspective, it should be Cognitive Service User.

upvoted 1 times

✉  **SuperPetey**  2 years, 9 months ago

The Owner role would give the leadership team the ability to publish changes, but also the additional capability to give others access to Qna maker, which they do not require. security best practices (in case of a compromised account etc) dictate the principles of least privilege; leadership team should only be given access to the capabilities they require to do their job and nothing more.

Since the problem only mentions publishing, the leadership team requires the role contributor. In other RBAC exam questions for other certs MSF remembers to preface such a question with "bearing in mind the principle of least privilege..." but they neglected that hint here.

upvoted 15 times

✉  **ziggy1117** 11 months, 3 weeks ago

Should be Cognitive Service User -> All access to Cognitive Services resource except for ability to:

1. Add new members to roles.
2. Create new resources.

Contributor role provides Leadership team access to create new resources. And they shouldnt be allowed to do that given that they arent technical people

upvoted 7 times

✉  **rdemontis** 6 months, 2 weeks ago

agree with you
upvoted 3 times

✉️ **takaimomoGcup** Most Recent 1 week ago

Is this question still available on May 21, 2024?
upvoted 1 times

✉️ **evangelist** 3 months, 2 weeks ago

The leadership should have permission to publish!!! here is why
AllUsers: QnA Maker Editor
Leadership team: Cognitive Service User

Key Differences in Privilege Scope

Specificity: The QnA Maker Editor role is specific to the QnA Maker service with a focus on content editing, while the Cognitive Service User role is more general and can encompass a broader range of tasks across Azure Cognitive Services.

Publishing the Knowledge Base: The ability to publish the knowledge base to make it live is a critical task. The QnA Maker Editor might have restricted access in this regard, requiring administrative approval, whereas a Cognitive Service User with the right permissions could directly publish changes.

upvoted 2 times

✉️ **james2033** 9 months, 2 weeks ago

1. Cognitive Services QnA Maker Editor
2. Cognitive Services User

See <https://learn.microsoft.com/en-us/azure/ai-services/qnamaker/concepts/role-based-access-control#access-is-provided-by-a-defined-role>
upvoted 5 times

✉️ **zellck** 11 months ago

1. QnA Maker Editor
2. Cognitive Service User

<https://learn.microsoft.com/en-us/azure/cognitive-services/qnamaker/concepts/role-based-access-control#access-is-provided-by-a-defined-role>
- QnA Maker Editor (read/write)
- Cognitive Service User (read/write/publish)
upvoted 12 times

✉️ **ziggy1117** 11 months, 3 weeks ago

AllUsers: QnA Maker Editor
(read/write)
Leadership Team: Cognitive Service User
(read/write/publish)

Leadership team does not need to add resources
upvoted 2 times

✉️ **KingChuang** 1 year, 4 months ago

on my exam. (2023-01-16 passed)
My Answer:
1. QnA Maker Editor
2. Cognitive Service User
upvoted 7 times

✉️ **ninja** 1 year, 9 months ago

To address the requirement: Provide all employees with the ability to edit Q&As. Only senior managers must be able to publish updates.

AllUsers: QnA Maker Editor
(read/write)
Leadership Team: Cognitive Service User
(read/write/publish)

Cognitive Service User has less permissions than Contributor.
upvoted 7 times

✉️ **ninja** 1 year, 9 months ago

Reference: <https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/concepts/role-based-access-control#access-is-provided-by-a-defined-role>
upvoted 1 times

✉️ **Eltooth** 1 year, 10 months ago

AllUsers: QnA Maker Editor
LeadershipTeam: Contributor
upvoted 1 times

✉️ **Ravnit** 2 years, 6 months ago

Was on exam 27/11/2021
upvoted 3 times

✉️ **AnonymousJhb** 7 months ago

it would be in the exam? its listed as question under ai-102. what a pointless comment to make
upvoted 1 times

✉️ **josebernabeo** 4 months ago

Pointless is your comment.
upvoted 2 times

✉️ **rikku33** 2 years, 7 months ago

from case study: Provide all employees with the ability to edit Q&As. Only senior managers must be able to publish updates, so the given answer correct
upvoted 2 times

✉️ **aakash_0086** 2 years, 10 months ago

Shouldn't Leadership team have owner rights which has right to publish the updates ?
upvoted 3 times

✉️ **aakash_0086** 2 years, 10 months ago

Shouldn't Leadership eam have owner rights which has right to publish the updates ?
upvoted 2 times

Question #3

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

A company named Wide World Importers is developing an e-commerce platform.

You are working with a solutions architect to design and implement the features of the e-commerce platform. The platform will use microservices and a serverless environment built on Azure.

Wide World Importers has a customer base that includes English, Spanish, and Portuguese speakers.

Existing Environment -

Applications -

Wide World Importers has an App Service plan that contains the web apps shown in the following table.

Name	Description
Product Management	An app used by employees to create and manage products. The app and the expected inputs from the employees are in English.
Inventory Tracking	An app used by employees to manage inventory when dispatching orders, receiving refunds, and receiving consignments from suppliers.

Azure Resources -

You have the following resources:

An Azure Active Directory (Azure AD) tenant

- The tenant supports internal authentication.

- All employees belong to a group named AllUsers.

- Senior managers belong to a group named LeadershipTeam.

An Azure Functions resource

- A function app posts to Azure Event Grid when stock levels of a product change between OK, Low Stock, and Out of Stock. The function app uses the -

Azure Cosmos DB change feed.

An Azure Cosmos DB account

- The account uses the Core (SQL) API.

- The account stores data for the Product Management app and the Inventory Tracking app.

An Azure Storage account

- The account contains blob containers for assets related to products.

- The assets include images, videos, and PDFs.

An Azure Cognitive Services resource named wwics

An Azure Video Analyzer for Media (previously Video Indexer) resource named wwivi

Requirements -

Business Goals -

Wide World Importers wants to leverage AI technologies to differentiate itself from its competitors.

Planned Changes -

Wide World Importers plans to start the following projects:

A product creation project: Help employees create accessible and multilingual product entries, while expediting product entry creation.

A smart e-commerce project: Implement an Azure Cognitive Search solution to display products for customers to browse.

A shopping on-the-go project: Build a chatbot that can be integrated into smart speakers to support customers.

Business Requirements -

Wide World Importers identifies the following business requirements for all the projects:

Provide a multilingual customer experience that supports English, Spanish, and Portuguese.

Whenever possible, scale based on transaction volumes to ensure consistent performance.

Minimize costs.

Governance and Security Requirements

Wide World Importers identifies the following governance and security requirements:

Data storage and processing must occur in datacenters located in the United States.

Azure Cognitive Services must be inaccessible directly from the internet.

Accessibility Requirements -

Wide World Importers identifies the following accessibility requirements:

All images must have relevant alt text.

All videos must have transcripts that are associated to the video and included in product descriptions.

Product descriptions, transcripts, and alt text must be available in English, Spanish, and Portuguese.

Product Creation Requirements -

Wide World Importers identifies the following requirements for improving the Product Management app:

Minimize how long it takes for employees to create products and add assets.

Remove the need for manual translations.

Smart E-Commerce Requirements -

Wide World Importers identifies the following requirements for the smart e-commerce project:

Ensure that the Cognitive Search solution meets a Service Level Agreement (SLA) of 99.9% availability for searches and index writes.

Provide users with the ability to search insight gained from the images, manuals, and videos associated with the products.

Support autocompletion and suggestion based on all product name variants.

Store all raw insight data that was generated, so the data can be processed later.

Update the stock level field in the product index immediately upon changes.

Update the product index hourly.

Shopping On-the-Go Requirements -

Wide World Importers identifies the following requirements for the shopping on-the-go chatbot:

Answer common questions.

Support interactions in English, Spanish, and Portuguese.

Replace an existing FAQ process so that all Q&A is managed from a central location.

Provide all employees with the ability to edit Q&As. Only senior managers must be able to publish updates.

Support purchases by providing information about relevant products to customers. Product displays must include images and warnings when stock levels are low or out of stock.

Product JSON Sample -

You have the following JSON sample for a product.

```
{  
    "sku": "b1",  
    "name": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "stocklevel": "Out of Stock",  
    "description": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "image":  
    {"uri": "https://upload.worldwideimporters.org/bicycle.jpg",  
        "alttext": {  
            "en": "Bicycle",  
            "es": "Bicicleta",  
            "pt": "Bicicleta"  
        }  
    },  
    "createdUtc": "2020-02-14T06:08:39Z",  
    "language": "en"  
}
```

Question

HOTSPOT -

You are developing the shopping on-the-go project.

You need to build the Adaptive Card for the chatbot.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
version": "1.3",
"body": [
  {

    "type": "TextBlock",
    "size": "Medium",
    "weight": "Bolder",
    "text": "${if(language == 'en', 'en', name)}  
name  
name.en  
name[language]"

  },
  {
    "type": "TextBlock",
    "$when": "${stockLevel != 'OK'}"  
"$when": "${stockLevel == 'OK'}"  
"$when": "${stockLevel.OK}"

    color : Attention
  },
  {
    "type": "Image",
    "url": "${image.uri}",
    "size": "Medium",
    "altText": "${image.altText.en}  
image.altText.language  
image.altText["language"]  
image.altText[language]"
  }
]
```

Answer Area

```

    "version": "1.3",
    "body": [
        {

            "type": "TextBlock",
            "size": "Medium",
            "weight": "Bolder",
            "text": "${

                if(language == 'en', 'en', name)
                name
                name.en
                name[language]
            }"
        },
        {
            "type": "TextBlock",
            "color": Attention
        }
    ]
}

```

Correct Answer:

```

    "when": "${stockLevel != 'OK'}"
    "when": "${stockLevel == 'OK'}"
    "when": "${stockLevel.OK}"

    color : Attention

}

]

}

```

Box 1: name.en -

Box 2: "\$when": "\${stockLevel != 'OK'}"

Product displays must include images and warnings when stock levels are low or out of stock.

Box 3:image.altText.en

halfway Highly Voted 1 year, 5 months ago

1. name[language] 2. "\$when": "\${stockLevel} != 'OK'" 3. image.altText[language]
upvoted 27 times

james2033 Highly Voted 9 months, 2 weeks ago

File JSON

```
{
    "sku": "b1",
    "name": {
        "en": "Bicycle",
        "es": "Bicicleta",
        "pt": "Bicicleta"
    },
    "stocklevel": "Out of Stock",
    "description": {
        "en": "Bicycle",
        "es": "Bicicleta",
        "pt": "Bicicleta"
    },
    "image": {
        "uri": "https://upload.wordwideimporters.org/bicycle.jpg",
        "alttext": {
            "en": "Bicycle",
            "es": "Bicicleta",
            "pt": "Bicicleta"
        }
    }
},
```

```
"createdUtc": "2020-0214T06:08:39Z",
"language": "en"
}
```

1. name.en
2. "\$when": "\${stockLevel != 'OK'}"
3. image.altText.en

upvoted 5 times

✉️ **upliftinghut** 3 weeks, 4 days ago

This is correct because the structure of original requirement Json is reflected and replaced by \$<parameter name>

upvoted 1 times

✉️ **takaimomoGcup** Most Recent 1 week ago

Is this question still available on May 21, 2024?

upvoted 1 times

✉️ **sl_mslconsulting** 8 months, 1 week ago

You don't want code it in a way that the code just works with en - it's a bad coding practice.

upvoted 3 times

✉️ **AzureJobsTillRetire** 1 year, 2 months ago

The given answers are correct.

The correct answer pairs for box1 and box3 could be-
name.en and image.altText.en
name["language"] and image.altText["language"]

since name["language"] does not exist as an option in box3, we will have to choose the first pair as the correct answer.

https://www.w3schools.com/js/js_json_syntax.asp

upvoted 3 times

✉️ **Canyu** 1 year, 6 months ago

This is a repeat question, but this time the answer does not give an explanation or reference document. Is this answer correct?

upvoted 2 times

✉️ **ParkXD** 1 year, 10 months ago

it seems a repeat question, but with a slight different option. The earlier one with altText.[language], but this altText[language] without . which one is right?

upvoted 4 times

✉️ **sdokmak** 1 year, 11 months ago

If there was an option for name.[language] instead of name[language] then I'd use that for more than one language option, otherwise we can't. For name[language], language would have to be an integer but it's a string, wouldn't work.

But.. for name.[language], that works.

<https://stackoverflow.com/questions/882727/is-there-a-way-to-use-variable-keys-in-a-javascript-object-literal>

upvoted 2 times

✉️ **RamonKaus** 1 year, 10 months ago

This is a repeat of an earlier question but they changed the answer? Weird.

upvoted 7 times

Topic 17 - Testlet 10

Question #1

Topic 17

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

General Overview -

Contoso, Ltd. is an international accounting company that has offices in France, Portugal, and the United Kingdom.

Contoso has a professional services department that contains the roles shown in the following table.

Name	Position	Office
Accountant	Manager	United Kingdom, France, Portugal
Accountant	Consultant	United Kingdom, France, Portugal
Customer Service	Manager	United Kingdom
Customer Service	Agent	United Kingdom
Bookkeeper	Manager	United Kingdom, France, Portugal
Bookkeeper	Consultant	United Kingdom, France, Portugal

Existing environment -

Infrastructure -

Contoso has the following subscriptions:

Azure

Microsoft 365

Microsoft Dynamics 365

Azure Active (Azure AD) Directory

Contoso has Azure Active Directory groups for securing role-based access. The company uses the following group naming conventions:

[Country]-[Level]-[Role]

[Level]-[Role]

Intellectual Property -

Contoso has the intellectual property shown in the following table.

Content	Format	Language	Content store	Domain
Weekly webinars	Video	English	Azure Blob storage	Vid.contoso.com
Blogs	Text	English, French, Portuguese	WordPress	Pt-blog.contoso.com Blog.contoso.com Fr-blog.contoso.com
Wikis	Text	English, French, Portuguese	Azure Cosmos DB	Internal.contoso.com/wiki
Monthly conference recordings	Video	English	SharePoint Online	Contoso.sharepoint.com
Frequently asked questions (FAQs)	Text	English	SharePoint Online	Contoso.sharepoint.com

Text-based content is provided only in one language and is not translated.

Requirements -

Planned Projects -

Contoso plans to develop the following:

A document processing workflow to extract information automatically from PDFs and images of financial documents

A customer-support chatbot that will answer questions by using FAQs

A searchable knowledgebase of all the intellectual property

Technical Requirements -

Contoso identifies the following technical requirements:

All content must be approved before being published.

All planned projects must support English, French, and Portuguese.

All content must be secured by using role-based access control (RBAC).

RBAC role assignments must use the principle of least privilege.

RBAC roles must be assigned only to Azure Active Directory groups.

AI solution responses must have a confidence score that is equal to or greater than 70 percent.

When the response confidence score of an AI response is lower than 70 percent, the response must be improved by human input.

Chatbot Requirements -

Contoso identifies the following requirements for the chatbot:

Provide customers with answers to the FAQs.

Ensure that the customers can chat to a customer service agent.

Ensure that the members of a group named Management-Accountants can approve the FAQs.

Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.

Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.

Ensure that access to the customer service agents is managed by using Omnichannel for Customer Service.

When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

Document Processing Requirements

Contoso identifies the following requirements for document processing:

The document processing solution must be able to process standardized financial documents that have the following characteristics:

- Contain fewer than 20 pages.

- Be formatted as PDF or JPEG files.

- Have a distinct standard for each office.

The document processing solution must be able to extract tables and text from the financial documents.

The document processing solution must be able to extract information from receipt images.

Members of a group named Management-Bookkeeper must define how to extract tables from the financial documents.

Members of a group named Consultant-Bookkeeper must be able to process the financial documents.

Knowledgebase Requirements -

Contoso identifies the following requirements for the knowledgebase:

Supports searches for equivalent terms

Can transcribe jargon with high accuracy

-

Can search content in different formats, including video

Provides relevant links to external resources for further research

Question

You are developing the chatbot.

You create the following components:

- A QnA Maker resource
- A chatbot by using the Azure Bot Framework SDK

You need to integrate the components to meet the chatbot requirements.

Which property should you use?

- A. QnAMakerOptions.StrictFilters
- B. QnADialogResponseOptions.CardNoMatchText
- C. QnAMakerOptions.RankerType
- D. QnAMakerOptions.ScoreThreshold

Correct Answer: C

Scenario: When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

When no good match is found by the ranker, the confidence score of 0.0 or "None" is returned and the default response is "No good match found in the KB". You can override this default response in the bot or application code calling the endpoint. Alternately, you can also set the override response in Azure and this changes the default for all knowledge bases deployed in a particular QnA Maker service.

Choosing Ranker type: By default, QnA Maker searches through questions and answers. If you want to search through questions only, to generate an answer, use the RankerType=QuestionOnly in the POST body of the GenerateAnswer request.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/concepts/best-practices>

Community vote distribution

D (86%)

14%

 **ManAtWorkAtNight** Highly Voted  2 years ago

Answer should be D. QnAMakerOptions.ScoreThreshold

Technical Requirements says "AI solution responses must have a confidence score that is equal to or greater than 70 percent" and "When the response confidence score is low, ensure that the chatbot can provide other response options to the customers"

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/concepts/confidence-score#set-threshold>

upvoted 19 times

 **Murtuza** Most Recent  2 months ago

Selected Answer: D

D is correct

upvoted 1 times

 **evangelist** 3 months, 2 weeks ago

Selected Answer: D

D. QnAMakerOptions.ScoreThreshold

This property allows you to set a minimum confidence score for answers. If an answer's confidence score is below this threshold, you can trigger alternative actions, such as providing different response options or escalating to a human agent, directly addressing the specified chatbot requirements.

upvoted 1 times

 **evangelist** 4 months ago

Selected Answer: D

QnAMakerOptions.ScoreThreshold is the property that controls the minimum confidence score required for a QnA Maker response to be considered valid. This directly addresses the requirement that AI responses must have a confidence score of at least 70%. By setting QnAMakerOptions.ScoreThreshold to 0.7 (or 70%), you ensure that the chatbot only provides answers from QnA Maker that have a high level of confidence, meeting the specified requirement.

upvoted 1 times

✉ **rdemontis** 6 months, 2 weeks ago

Selected Answer: D

<https://learn.microsoft.com/en-us/azure/ai-services/qnamaker/concepts/confidence-score>

upvoted 1 times

✉ **propanther** 7 months ago

Complete chatbot constraints from Case Study:

Contoso identifies the following requirements for the chatbot:

- * Provide customers with answers to the FAQs.
- * Ensure that the customers can chat to a customer service agent.
- * Ensure that the members of a group named Management-Accountants can approve the FAQs.
- * Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.
- * Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.
- * Ensure that access to the customer service agents is managed by using Omnichannel for Customer Service.
- * When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

upvoted 1 times

✉ **ExamDev** 8 months, 3 weeks ago

Selected Answer: D

James2033 is right. D (but he marked "C")

upvoted 2 times

✉ **james2033** 9 months, 1 week ago

Selected Answer: C

"AI solution responses must have a confidence score that is equal to or greater than 70 percent." --> So choose "D".
QnAMakerOptions.ScoreThreshold". See <https://learn.microsoft.com/en-us/dotnet/api/microsoft.bot.builder.ai.qna.qnamakeroptions.scorethreshold?view=botbuilder-dotnet-stable>

upvoted 2 times

✉ **Tin_Tin** 11 months ago

Selected Answer: D

should be D

upvoted 1 times

✉ **SSJA** 1 year, 5 months ago

Selected Answer: D

Correct answer is D

upvoted 2 times

✉ **Eltooth** 1 year, 10 months ago

Selected Answer: D

I'm going for answer D based on

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/concepts/best-practices>

upvoted 2 times

✉ **RamonKaus** 1 year, 10 months ago

Selected Answer: D

<https://www.google.com/url?sa=t&rct=j&q=&esrc=s&source=web&cd=&cad=rja&uact=8&ved=2ahUKEwilta3-tZD5AhXpGFkFHWGzAMAQFnoECAoQAw&url=https%3A%2F>

Choosing Ranker type

" By default, QnA Maker searches through questions and answers. If you want to search through questions only, to generate an answer, use the RankerType=QuestionOnly in the POST body of the GenerateAnswer request. "

from

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/concepts/best-practices?view=botbuilder-dotnet-stable>

vb3

I think it is D because "Gets or sets the minimum score threshold, used to filter returned results."

<https://docs.microsoft.com/en-us/dotnet/api/microsoft.bot.builder.ai.qna.qnamakeroptions.scorethreshold?view=botbuilder-dotnet-stable>

upvoted 1 times

Question #2

Introductory Info

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview -

General Overview -

Contoso, Ltd. is an international accounting company that has offices in France, Portugal, and the United Kingdom.

Contoso has a professional services department that contains the roles shown in the following table.

Name	Position	Office
Accountant	Manager	United Kingdom, France, Portugal
Accountant	Consultant	United Kingdom, France, Portugal
Customer Service	Manager	United Kingdom
Customer Service	Agent	United Kingdom
Bookkeeper	Manager	United Kingdom, France, Portugal
Bookkeeper	Consultant	United Kingdom, France, Portugal

Existing environment -

Infrastructure -

Contoso has the following subscriptions:

Azure

Microsoft 365

Microsoft Dynamics 365

Azure Active (Azure AD) Directory

Contoso has Azure Active Directory groups for securing role-based access. The company uses the following group naming conventions:

[Country]-[Level]-[Role]

[Level]-[Role]

Intellectual Property -

Contoso has the intellectual property shown in the following table.

Content	Format	Language	Content store	Domain
Weekly webinars	Video	English	Azure Blob storage	Vid.contoso.com
Blogs	Text	English, French, Portuguese	WordPress	Pt-blog.contoso.com Blog.contoso.com Fr-blog.contoso.com
Wikis	Text	English, French, Portuguese	Azure Cosmos DB	Internal.contoso.com/wiki
Monthly conference recordings	Video	English	SharePoint Online	Contoso.sharepoint.com
Frequently asked questions (FAQs)	Text	English	SharePoint Online	Contoso.sharepoint.com

Text-based content is provided only in one language and is not translated.

Requirements -

Planned Projects -

Contoso plans to develop the following:

A document processing workflow to extract information automatically from PDFs and images of financial documents

A customer-support chatbot that will answer questions by using FAQs

A searchable knowledgebase of all the intellectual property

Technical Requirements -

Contoso identifies the following technical requirements:

All content must be approved before being published.

All planned projects must support English, French, and Portuguese.

All content must be secured by using role-based access control (RBAC).

RBAC role assignments must use the principle of least privilege.

RBAC roles must be assigned only to Azure Active Directory groups.

AI solution responses must have a confidence score that is equal to or greater than 70 percent.

When the response confidence score of an AI response is lower than 70 percent, the response must be improved by human input.

Chatbot Requirements -

Contoso identifies the following requirements for the chatbot:

Provide customers with answers to the FAQs.

Ensure that the customers can chat to a customer service agent.

Ensure that the members of a group named Management-Accountants can approve the FAQs.

Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.

Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.

Ensure that access to the customer service agents is managed by using Omnichannel for Customer Service.

When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

Document Processing Requirements

Contoso identifies the following requirements for document processing:

The document processing solution must be able to process standardized financial documents that have the following characteristics:

- Contain fewer than 20 pages.

- Be formatted as PDF or JPEG files.

- Have a distinct standard for each office.

The document processing solution must be able to extract tables and text from the financial documents.

The document processing solution must be able to extract information from receipt images.

Members of a group named Management-Bookkeeper must define how to extract tables from the financial documents.

Members of a group named Consultant-Bookkeeper must be able to process the financial documents.

Knowledgebase Requirements -

Contoso identifies the following requirements for the knowledgebase:

Supports searches for equivalent terms

Can transcribe jargon with high accuracy

▪

Can search content in different formats, including video

Provides relevant links to external resources for further research

Question

You are developing the chatbot.

You create the following components:

- A QnA Maker resource
- A chatbot by using the Azure Bot Framework SDK

You need to add an additional component to meet the technical requirements and the chatbot requirements.

What should you add?

- A. Microsoft Translator
- B. Language Understanding
- C. Orchestrator
- D. Chatdown

Correct Answer: C

Scenario: All planned projects must support English, French, and Portuguese.

If a bot uses multiple LUIS models and QnA Maker knowledge bases (knowledge bases), you can use the Dispatch tool to determine which LUIS model or QnA

Maker knowledge base best matches the user input. The dispatch tool does this by creating a single LUIS app to route user input to the correct model.

Reference:

<https://docs.microsoft.com/en-us/azure/bot-service/bot-builder-tutorial-dispatch>

Community vote distribution



 **Gdavou**  11 months, 2 weeks ago

B. Language Understanding

Explanation:

To meet the requirement that the chatbot must support English, French, and Portuguese and provide other response options when the response confidence score is low, the Language Understanding component (often referred to as LUIS - Language Understanding Intelligent Service) is needed. LUIS allows the bot to understand the user's intent in multiple languages and can help to provide alternative responses.

Microsoft Translator could translate between languages, but it does not handle intent understanding, so it's not the best choice here. Orchestrator is more about managing multiple LUIS and QnA Maker models, which isn't specified in the requirements. Chatdown is a tool for creating mockups of chatbot conversations but doesn't provide the functionality needed to meet the requirements.

upvoted 8 times

 **M25** 8 months, 3 weeks ago

Correct! Thank you.

<https://learn.microsoft.com/en-us/azure/ai-services/luis/luis-language-support#multilingual-luis-apps>

If you need a multilingual LUIS client application such as a chatbot, you have a few options. If LUIS supports all the languages, you develop a LUIS app for each language. Each LUIS app has a unique app ID, and endpoint log. If you need to provide language understanding for a language LUIS does not support, you can use the Translator service to translate the utterance into a supported language, submit the utterance to the LUIS endpoint, and receive the resulting scores.

upvoted 1 times

 **rdemontis** 6 months, 2 weeks ago

thanks for reference

upvoted 1 times

 **rdemontis** 6 months, 2 weeks ago

I think you are right!
upvoted 1 times

 **Eltooth**  1 year, 10 months ago

Selected Answer: A

A is correct answer.

"If you need to support a knowledge base system, which includes several languages, you can:

1) Use the Translator service to translate a question into a single language before sending the question to your knowledge base. This allows you to focus on the quality of a single language and the quality of the alternate questions and answers.

2) Create a QnA Maker resource, and a knowledge base inside that resource, for every language. This allows you to manage separate alternate questions and answer text that is more nuanced for each language. This gives you much more flexibility but requires a much higher maintenance cost when the questions or answers change across all languages."

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/overview/language-support#supporting-multiple-languages-in-one-knowledge-base>

upvoted 8 times

 **AzureJobsTillRetire** 1 year, 2 months ago

Scenario: Text-based content is provided only in one language and is not translated.

Hence, translator may not be required

upvoted 1 times

 **Murtuza**  1 month, 3 weeks ago

Selected Answer: C

Orchestrator is the correct answer while the other choices are very close to being the right answer as well

upvoted 1 times

 **Murtuza** 2 months ago

Selected Answer: B

To meet both the technical and chatbot requirements, you should add Language Understanding (LUIS) to your chatbot solution.

upvoted 1 times

 **Murtuza** 2 months, 2 weeks ago

Orchestrator is a replacement for the now deprecated Bot Framework Dispatcher.

upvoted 1 times

 **evangelist** 3 months, 2 weeks ago

Selected Answer: C

C. Orchestrator:

Orchestrator is a component of the Bot Framework responsible for routing messages to different bots or processes based on the context or content of the conversation. It can help in managing complex scenarios where multiple bots or services are involved, such as switching between the FAQ bot and a live agent chat, or handling low-confidence responses by offering alternative options or escalating to human agents. This component is crucial for managing the flow of conversation based on AI confidence scores and ensuring that users can be directed to a customer service agent when needed.

upvoted 3 times

 **evangelist** 4 months ago

Selected Answer: A

A. Microsoft Translator

Here's why:

Technical requirements: All projects must support English, French, and Portuguese.

Chatbot requirements: Provide answers to FAQs and chat with customers in English, French, and Portuguese.

upvoted 1 times

 **rdemontis** 6 months, 2 weeks ago

Selected Answer: B

I think correct answer is B. Technical and Chatbot requirement have in common the behaviour to be taken in case of confidence score lower than 70%. Particularly one of the chatbot requirements is: "When the response confidence score is low, ensure that the chatbot can provide other response options to the customers". And LUIS fits the bill

<https://learn.microsoft.com/en-us/azure/ai-services/luis/luis-language-support#multilingual-luis-apps>

upvoted 4 times

 **sl_mslconsulting** 7 months ago

Selected Answer: C

FAQs is only available in English and there are no multi-language support required for the chat bot. You need dispatching capability however as you need to provide a way to talk to a agent and integrate with Omnichannel for Customer Service.

upvoted 3 times

✉️ **Salem2020s** 8 months, 1 week ago

Selected Answer: A

<https://learn.microsoft.com/en-us/azure/ai-services/qnamaker/overview/language-support#:~:text=to%20a%20query.-,Supporting%20multiple%20languages%20in%20one%20QnA%20Maker%20resource,maintenance%20cost%,when%20the%20questions%20or%20answers%20change%20across%20all%20languages.,-Languages%20supported>

upvoted 1 times

✉️ **halfway** 1 year, 5 months ago

Selected Answer: A

Use translator service to support multiple languages

upvoted 2 times

✉️ **not_a_robot** 1 year, 9 months ago

Selected Answer: C

QnA Maker resource can meet multi-language requirement. There is no point to add another translator. Orchestrator can provide multi-intent detection, which meets the chatbot requirements.

upvoted 5 times

✉️ **AzureJobsTillRetire** 1 year, 3 months ago

Orchestrator provides a multi-lingual model alongside English which provides the ability for a model trained with, for example English-only, data to process utterances in other languages.

<https://learn.microsoft.com/en-us/composer/concept-orchestrator?source=recommendations>

upvoted 2 times

✉️ **RamonKaus** 1 year, 10 months ago

I feel like this answer is fair.

upvoted 1 times

✉️ **g2000** 2 years ago

Selected Answer: A

There are six sources of IP. Qna supports multiple sources of urls. With that being said, there's no need to have one Qna for one source. Also one technical requirement is to support multiple languages. This is certainly feasible with the translation.

upvoted 3 times

Get IT Certification

Unlock free, top-quality video courses on ExamTopics with a simple registration. Elevate your learning journey with our expertly curated content. Register now to access a diverse range of educational resources designed for your success. Start learning today with ExamTopics!

Start Learning for free